C8 Copilot and AI in Fabric

Copilot across workloads, data agents, limits and verification.

C8-L1 Copilot in Fabric, where it runs and what it costs

Confirm Copilot is enabled on a paid capacity, record the release status of each experience, and budget its CU use.

Copilot in Fabric capacity requirement	Paid F2+ or P1+ in a supported region, with tenant settings enabled. Not after a trial expires.
Copilot release status	Per experience. Notebooks and Warehouse SQL preview. Data agents, Data Factory and Power BI GA (September 2026).
Copilot token rates per 1,000 tokens	Input 100, cached input 10, output 400 CU seconds. Background job, smoothed over 24 hours.
Copilot output	Sometimes inaccurate. A person able to judge it reviews it before use.

C8-L2 Data agents, questions in plain language

Narrow the schema, define your terms, add validated examples, and measure accuracy against known answers.

Data agent query languages	SQL for lakehouse and warehouse, DAX for semantic models, KQL for KQL databases. Up to five sources.
Data agent security	Read-only queries only. Runs with the asking user's permissions.
Example queries	Question and query pairs, validated against the schema. Not supported for semantic models.
Measure a data agent	Ground-truth questions with known answers, rerun after every change, accuracy recorded.

C8-L3 Verify what Copilot writes

Read it, run it on data with a hand-worked answer, compare it with a second method, and keep those checks as tests.

A Copilot query ran without error. What does that prove?	Syntax parses and names exist. Nothing about correctness.
Four checks for Copilot output	Read every line. Run on known data. Compare with a second method. Keep the checks as tests.
Average of averages	Weights small groups like large ones. Divide total by total instead.
Known data for verification	Five to ten hand-built rows with a null, a zero denominator, a duplicate key and an empty group.