

Full handbook

Generated from the same content files as the app. Each module lists its roadmap months. Solutions are in the appendix.

Contents

- Track A: Databricks pipeline engineering.....7**
 - A1 Code structure..... 8
 - A1-L1 Single-purpose functions and clear names 8
 - A1-L2 Type hints and docstrings a reviewer can trust..... 12
 - A1-L3 Config over hardcoded values, and PEP 8 with ruff..... 16
 - A2 Medallion architecture..... 21
 - A2-L1 Bronze: keep every delivery exactly as received 21
 - A2-L2 Silver: canonical names, checked types, one row per respondent per wave 25
 - A2-L3 Gold: state the grain, group by all of it, guard every ratio 30
 - A3 Delta Lake 35
 - A3-L1 Idempotent MERGE, and why you deduplicate the source first..... 35
 - A3-L2 Schema enforcement is on by default. Evolve on purpose. 41
 - A3-L3 Time travel, RESTORE, OPTIMIZE and VACUUM..... 45
 - A4 Unity Catalog 50
 - A4-L1 The three-level namespace, one catalog per environment 50
 - A4-L2 Grant to groups, grant the least, grant at the right level 54
 - A4-L3 Row filters and column masks in Unity Catalog 58
 - A5 PySpark performance 63
 - A5-L1 Lazy evaluation, and why every count() re-runs the plan 63
 - A5-L2 Built-in functions over Python UDFs 67
 - A5-L3 Joins, broadcast, partitions and reading the plan..... 71
 - A6 Data quality..... 76
 - A6-L1 Data contracts with severity levels and structured results 76
 - A6-L2 Statistical checks that survive empty, single-row and zero data 81
 - A6-L3 Reconcile counts and totals between layers 85
 - A7 Observability 90
 - A7-L1 Structured logging with the logging module, not print 90
 - A7-L2 A pipeline run log table with an explicit schema..... 94
 - A7-L3 Alert on failures, and on the silence that hides them..... 99
 - A8 Incremental processing 103

A8-L1 Process only the waves you have not processed, correctly.....	103
A8-L2 Auto Loader for file ingestion into bronze	107
A8-L3 Change data feed, so downstream layers read only what changed	111
A9 Testing	117
A9-L1 pytest basics, and a SparkSession fixture shared by every test	117
A9-L2 A fast local SparkSession and readable test data	121
A9-L3 Compare DataFrames properly, and test that MERGE is idempotent ...	125
A10 Deployment.....	130
A10-L1 Git folders, one branch per change, and no secrets in the repository	130
A10-L2 Declarative Automation Bundles and parameterised jobs	133
A10-L3 CI: lint, type-check, test and validate on every pull request	138
Track A best-practice checklist	144
Track B: Python for data science	145
B1 Project setup	146
B1-L1 One pyproject.toml and a src layout	146
B1-L2 One virtual environment per project, and python -m pip	150
B1-L3 ruff and mypy in daily work	153
B2 Code design	158
B2-L1 Pure functions, side effects at the edges	158
B2-L2 Dataclasses as configuration objects	162
B2-L3 When a class earns its place	166
B3 Clean pandas	171
B3-L1 Dtypes in pandas 3, nullable and string.....	171
B3-L2 Vectorise instead of apply	174
B3-L3 Method chaining and Copy-on-Write	179
B4 Polars and DuckDB	184
B4-L1 Polars expressions and lazy frames	184
B4-L2 DuckDB, SQL over files and DataFrames.....	187
B4-L3 Choosing pandas, Polars, DuckDB or Spark	191
B5 Data validation	196
B5-L1 DataFrame schemas with pandera.....	196
B5-L2 pydantic models for config and input records.....	200
B5-L3 Validate at the boundaries, with severity	204
B6 Reproducibility	210
B6-L1 Seeds and random generators.....	210
B6-L2 Record the data, config and code behind every result.....	214
B6-L3 Pinned dependencies and lock files	217
B7 Exploratory analysis done well.....	221

B7-L1 Profile a dataset before you analyse it.....	221
B7-L2 Plot before you summarise	224
B7-L3 Write assumptions down, then turn them into checks	228
B8 Statistical correctness	233
B8-L1 Survey weights and the effective sample size	233
B8-L2 Confidence intervals and effect sizes.....	236
B8-L3 Many tests, and the false positives they bring.....	240
B9 Experimental design	244
B9-L1 Power and sample size, before fieldwork	244
B9-L2 Analyse an A/B test the way it was randomised	248
B9-L3 Peeking, novelty and other ways tests mislead	252
B10 Feature engineering without leakage.....	256
B10-L1 Target leakage, features that know the answer	256
B10-L2 Fit on training data only	260
B10-L3 Time-safe features and point-in-time joins.....	264
B11 Machine learning workflow.....	268
B11-L1 Train, validation and test, split by the right unit	268
B11-L2 Pipeline and ColumnTransformer	271
B11-L3 Cross-validation and honest hyperparameter search.....	275
B12 Model evaluation.....	280
B12-L1 Choose the metric from the decision	280
B12-L2 Calibration, when the probability itself matters	283
B12-L3 Imbalanced classes and error analysis by segment.....	287
B13 Interpretability and fairness	291
B13-L1 Permutation importance on held-out data	291
B13-L2 SHAP values, and what they do not tell you.....	294
B13-L3 Performance by subgroup, and the fairness trade-offs	298
B14 Experiment tracking with MLflow	302
B14-L1 Track every run with MLflow.....	302
B14-L2 Registry versions and aliases, not stages	306
B14-L3 Reproduce a run from its record.....	310
B15 Embeddings and clustering	314
B15-L1 Sentence embeddings for open-ended answers	314
B15-L2 HDBSCAN, clusters with a noise label	317
B15-L3 Validate clusters before you name them.....	321
B16 PyTorch hygiene	325
B16-L1 Choose the device once, move model and data to it.....	325
B16-L2 A training loop that does what you think	329

B16-L3 Reproducible PyTorch runs and safe checkpoints.....	333
B17 Testing and performance.....	337
B17-L1 Unit tests for data code that catch real bugs	337
B17-L2 Property-based tests with Hypothesis	340
B17-L3 Measure before you optimise	344
B18 Documentation and communication	349
B18-L1 Docstrings and READMEs that answer the next person's questions ..	349
B18-L2 Model cards that state limits as clearly as scores	353
B18-L3 Decision logs, so the reasons outlive the meeting	356
Track B best-practice checklist	360
Track C: Microsoft Fabric	361
C1 Fabric architecture	362
C1-L1 Tenant, capacity, workspace and item	362
C1-L2 OneLake, one lake per tenant	365
C1-L3 Workspace roles, granted to groups	369
C2 Lakehouse, Warehouse and SQL analytics endpoint	374
C2-L1 Lakehouse Tables and Files, and schemas	374
C2-L2 Warehouse T-SQL, and where it differs from SQL Server.....	378
C2-L3 Lakehouse, Warehouse or SQL analytics endpoint.....	382
C3 Unifying datasets	388
C3-L1 OneLake shortcuts, and the delete that goes through.....	388
C3-L2 Mirroring operational databases into OneLake.....	391
C3-L3 Mirroring Azure Databricks Unity Catalog into Fabric	395
C4 Ingestion and orchestration	400
C4-L1 Choosing how data gets into Fabric	400
C4-L2 Incremental copy with a watermark, or with CDC	404
C4-L3 Pipelines that fail fast, retry sensibly and alert	407
C5 Spark notebooks in Fabric	412
C5-L1 From Databricks notebooks to Fabric notebooks.....	412
C5-L2 Environments, resource profiles and V-Order.....	416
C5-L3 Delta tables shared between Databricks and Fabric	420
C6 Semantic models and Direct Lake	424
C6-L1 Star schemas for semantic models	424
C6-L2 DAX measures that state their filter context.....	428
C6-L3 Direct Lake, framing and the fallback you did not plan.....	431
C7 Real-Time Intelligence	436
C7-L1 Eventstream, events in, routed out.....	436
C7-L2 Eventhouse and KQL for time-series questions.....	440

C7-L3 Activator, rules that act on events	443
C8 Copilot and AI in Fabric.....	448
C8-L1 Copilot in Fabric, where it runs and what it costs.....	448
C8-L2 Data agents, questions in plain language	452
C8-L3 Verify what Copilot writes	456
C9 Governance and security	460
C9-L1 Domains and endorsement, finding the right data	460
C9-L2 OneLake security, roles down to rows and columns.....	464
C9-L3 Sensitivity labels, Purview and lineage	468
C10 DevOps and cost.....	473
C10-L1 Git integration for Fabric workspaces.....	473
C10-L2 Deployment pipelines, dev to test to production	477
C10-L3 Capacity, smoothing and throttling.....	481
Track C best-practice checklist	486
Track D: Finding errors in code	487
D1 Reading tracebacks and stack traces.....	488
D1-L1 Read a Python traceback from the bottom up	488
D1-L2 Spark errors, the condition and the real cause	492
D1-L3 Py4J and Spark Connect errors	496
D2 Debugging tools.....	501
D2-L1 breakpoint() and pdb, stop and look.....	501
D2-L2 Logging levels that help at 3 a.m.	505
D2-L3 explain(), the Spark UI and the query profile	508
D3 Silent errors	513
D3-L1 Join row explosion and nulls in filters.....	513
D3-L2 Type coercion and timezone drift.....	517
D3-L3 Mutable defaults and the truthiness of zero	521
D4 Data errors.....	526
D4-L1 Row counts across layers.....	526
D4-L2 Distribution drift between waves	530
D4-L3 Weight totals and survey checks	533
D5 Errors typical of AI-generated code.....	537
D5-L1 Invented functions, parameters and syntax	537
D5-L2 Deprecated APIs and version mismatches.....	541
D5-L3 Five review checks for AI-written code.....	544
D6 Verifying AI output	550
D6-L1 Read before you run, and check each API.....	550
D6-L2 Small data with a known answer, failing test first	553

D6-L3 Ask for assumptions, then compare a second approach.....	557
D7 Review practice.....	561
D7-L1 Review practice: Python and pandas.....	561
D7-L2 Review practice: PySpark and Spark SQL	564
D7-L3 Review practice: Databricks SQL and DAX.....	569
D8 Prompting for correct code	573
D8-L1 State versions and constraints in every prompt	573
D8-L2 State expected outputs and tests in the prompt	576
D8-L3 Iterate on a failing answer.....	579
Track D best-practice checklist.....	583
Appendix: solutions and answer keys	584

Track A: Databricks pipeline engineering

Build Databricks pipelines that are correct, fast, observable and tested.

A1 Code structure

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Functions, type hints, docstrings, naming, config files and PEP 8.

A1-L1 Single-purpose functions and clear names

Concept

A function should do **one** of five jobs: read, transform, validate, write or log. A sixth kind, the **orchestrator**, calls the others in order and does nothing else.

Single-purpose functions pay off in three ways:

- **Tests.** A transform that takes a DataFrame and returns a DataFrame runs on a local SparkSession with three rows. A function that reads a table, transforms it and writes it needs a live catalog.
- **Reuse.** `filter_in_scope()` serves gold, QA and ad hoc analysis. A 200-line `process()` serves one notebook.
- **Change.** When the supplier renames a column, you edit one transform. Nothing else moves.

A useful test: describe the function in one sentence without "and". "Aggregate reach by wave and subgroup" passes. "Read silver and compute reach and write gold" fails three times.

Names carry the contract. Follow the standard from your reference project:

Object	Convention	Example
Function	verb_noun, snake_case	<code>check_table_health</code>
Class	PascalCase noun	<code>CheckResult</code>
Constant	UPPER_SNAKE_CASE	<code>MAX_WEIGHT_CV</code>
Bronze table	<code>[tracker]_raw</code>	<code>nat_raw</code>
Silver table	<code>[tracker]_respondents</code>	<code>nat_respondents</code>
Gold table	<code>[tracker]_metrics</code>	<code>nat_metrics</code>
Job	<code>[tracker]_[layer]_pipeline</code>	<code>nat_silver_pipeline</code>

A name must match what the function does. `get_` and `check_` promise no side effects. A `get_unprocessed_waves()` that also deletes rows breaks that promise, and the next reader calls it without knowing.

Why it matters

A gold notebook holds one 350-line cell: read silver, filter, weight, aggregate, write, log. A change request asks for a new subgroup. The developer edits the aggregation, and a variable named `df` later in the cell now points at the new frame. The write step saves the unfiltered data. Gold publishes reach figures that include out-of-scope respondents for two waves.

Nobody could test the cell without writing to production tables, so nobody tested it. With `filter_in_scope()` and `aggregate_reach()` as separate functions, a three-row unit test catches the mistake before the merge.

Worked example: weak code

```
❶ def process(t, w):
2   df = spark.table(t).filter(F.col("wave_id") == w)
❷   df = df.filter(F.col("in_scope") == True)
4   df = df.groupBy("wave_id",
    "region").agg(F.sum("final_weight").alias("weighted_base"))
❸   df.write.mode("overwrite").saveAsTable("research_gold.nat_metrics")
❹   print("done", df.count())
```

1. **Line 1.** The name says nothing, and `t` and `w` hide what the arguments are.
2. **Line 3.** `== True` trips ruff E712. Filter on the boolean column directly.
3. **Line 5.** The target table is hardcoded inside a function that also reads and transforms.
4. **Line 6.** Logging by `print`, and a second full scan to count rows.

Worked example: strong code

```
❶ def read_wave(spark: SparkSession, table: str, wave_id: str) -> DataFrame:
2   """Read one wave from a silver table."""
3   return spark.table(table).filter(F.col("wave_id") == wave_id)
4
5
❷ def filter_in_scope(df: DataFrame) -> DataFrame:
7   """Keep respondents flagged as in scope."""
❸   return df.filter(F.col("in_scope"))
9
10
11 def aggregate_base(df: DataFrame) -> DataFrame:
12   """Weighted base at wave x region grain."""
13   return df.groupBy("wave_id",
    "region").agg(F.sum("final_weight").alias("weighted_base"))
14
15
❹ def run_gold(spark: SparkSession, silver: str, gold: str, wave_id: str) -
> None:
17   """Orchestrate: read, filter, aggregate, write."""
```

```

18     metrics = aggregate_base(filter_in_scope(read_wave(spark, silver,
    wave_id)))
5     metrics.write.mode("overwrite").option("replaceWhere", f"wave_id =
    '{wave_id}'").saveAsTable(gold)

```

1. **Line 1.** Read only. The session and the table name come in as parameters.
2. **Line 6.** A pure transform: DataFrame in, DataFrame out. Testable with three rows.
3. **Line 8.** Boolean column used directly.
4. **Line 16.** The orchestrator calls the steps in order and holds no logic of its own.
5. **Line 19.** Overwrite only this wave, so a rerun never wipes other waves.

Each strong function passes the one-sentence test. The weak one needs "and" four times.

Recap card

- **The five jobs a pipeline function does, one at a time** Read, transform, validate, write, log. An orchestrator calls them in order.
- **The one-sentence test** Describe the function without using "and". If you need "and", split it.
- **Function, class and constant naming** `check_table_health`, `CheckResult`, `MAX_WEIGHT_CV`.
- **What does a get_ or check_ prefix promise?** No side effects. The function reads and returns. It never writes or deletes.

Practice

Exercise 1. Multiple choice (level 1)

You write a function that returns a list of `CheckResult` records for one table. Which name follows the standard?

- a) `CheckTableHealth`
- b) `check_table_health`
- c) `checkTableHealth`
- d) `CHECK_TABLE_HEALTH`

Why?

- r1) Functions use snake_case and start with a verb.
- r2) Shorter names run faster.

Exercise 2. Fill the gap (level 2)

Fill in the object names for tracker `nat`, following the naming standard.

```

1 BRONZE_TABLE = "research_bronze.____(1)"
2 SILVER_TABLE = "research_silver.____(2)"
3 GOLD_TABLE = "research_gold.____(3)"
4 JOB_NAME = "____(4)"

```

Exercise 3. Spot the bug (level 3)

This function has a name that promises no side effects. One line breaks that promise. Tap it, then fix the function so it only returns the waves.

```

1 def get_unprocessed_waves(spark: SparkSession, bronze: str, silver: str,
2   tracker_id: str) -> list[str]:
3     """Return wave ids in bronze but not yet in silver for one tracker."""
4     bronze_waves = spark.table(bronze).filter(F.col("tracker_id") ==
5       tracker_id).select("wave_id").distinct()
6     silver_waves = spark.table(silver).filter(F.col("tracker_id") ==
7       tracker_id).select("wave_id").distinct()
8     waves = [row.wave_id for row in
9       bronze_waves.subtract(silver_waves).collect()]
10    spark.sql(f"DELETE FROM {silver} WHERE wave_id IN ({','.join(repr(w)
11      for w in waves)})")
12    return waves

```

Write down the line number of the bug. Then write the fix.

Exercise 4. Refactor (level 4)

Split the transform logic out of `process`. Write two pure functions:

- `filter_in_scope(df)` keeps rows where `in_scope` is true.
- `aggregate_base(df)` returns `wave_id`, `region` and `weighted_base` (sum of `final_weight`).

Neither function may read or write a table. Add type hints and docstrings.

Starter code:

```

1 def process(t, w):
2     df = spark.table(t).filter(F.col("wave_id") == w)
3     df = df.filter(F.col("in_scope") == True)
4     df = df.groupBy("wave_id",
5       "region").agg(F.sum("final_weight").alias("weighted_base"))
6     df.write.mode("overwrite").saveAsTable("research_gold.nat_metrics")

```

Exercise 5. Architecture design (level 5)

You split a gold notebook into functions. Where does each piece of work belong?

Options: Config file, Read function, Transform function, Orchestrator.

1. The name of the gold table for tracker nat
2. `spark.table(silver).filter(wave_id == ...)`

3. Summing weights by wave and region
4. Calling read, filter, aggregate and write in turn, then logging the run

Sources

- ruff rule E712 flags comparisons to True or False. (verified: Reference README lists it; ruff 0.16.9 flags `x == True` under E712 in this repository.)
- `replaceWhere` overwrites only matching rows in a Delta table. (verified: Run locally on delta-spark 4.0.1 (VERSIONS.md, Track A facts).)

A1-L2 Type hints and docstrings a reviewer can trust

Concept

A type hint states a contract. `keys: list[str]` tells the caller, the reviewer and mypy that `keys` is a list of column names. A bare `list` says "some list". No hint says nothing.

Use the modern forms, valid from Python 3.10:

- Built-in generics: `list[str]`, `dict[str, int]`, `tuple[str, ...]`.
- Optional values: `str | None = None`. Never `str = None`, which claims the value is always a string.
- Fixed sets of strings: `Literal["PASS", "WARN", "FAIL"]` from `typing`. mypy then rejects `"pass"` or `"FAILED"`.
- A return type on every function, including `-> None`.

Run `mypy --strict` in CI. Strict mode refuses functions without hints, so gaps cannot creep back.

A docstring explains what hints cannot: meaning, units, grain and failure modes. Use the Google style:

```
1 def weighted_share(df: DataFrame, flag_col: str) -> float | None:
2     """Weighted share of rows where flag_col is 1.
3
4     Args:
5         df: Respondents with final_weight.
6         flag_col: 0/1 column to share.
7
8     Returns:
9         Share between 0 and 1, or None when the total weight is 0.
10
11     Raises:
12         ValueError: If flag_col is missing from df.
13     """
```

The first line fits on one line and says what the function returns or does. Args, Returns and Raises follow. A reviewer reads the docstring, then checks the body against it. A mismatch between the two is a bug in one of them. PySpark is not fully typed, and mypy cannot see column names. Hints still catch the most common mistake: passing the wrong kind of object.

Why it matters

The reference `upsert_silver` takes `merge_keys: list`. A caller passes the string `"respondent_id"` instead of `["respondent_id"]`. Python iterates a string character by character, so the join builds `t.r = s.r AND t.e = s.e AND ...`. The MERGE fails with a column-not-found error at 02:00, and the on-call engineer spends an hour reading the wrong part of the stack trace.

With `merge_keys: list[str]`, mypy flags the call in CI: a `str` is not a `list[str]`. The bug never reaches the scheduler.

Worked example: weak code

```
❷ def log_pipeline_run(pipeline_name, layer, rows_processed, status,
    error_message: str = None):
❸     """Append one record per pipeline run to the governance run log."""
    record = {
        "pipeline_name": pipeline_name,
        "layer": layer,
        "rows_processed": rows_processed,
❹     "status": status,
        "error_message": error_message,
    }
    return record
```

1. **Line 1.** Claims a string, defaults to None. Reference defect 7.
2. **Line 1.** No hints on four parameters and no return type. mypy --strict rejects the function.
3. **Line 2.** The docstring says it appends, but the function returns a dict. One of them is wrong.
4. **Line 7.** `status` accepts any string, so `"sucess"` slips into the run log.

Worked example: strong code

```
1 from typing import Literal
2
❶ RunStatus = Literal["SUCCESS", "FAILED"]
4 Layer = Literal["bronze", "silver", "gold"]
5
6
❷ def build_run_record(
    8     pipeline_name: str,
```

```

9     layer: Layer,
10     rows_processed: int,
11     status: RunStatus,
12     error_message: str | None = None,
13 ) -> dict[str, str | int | None]:
14     """Build one run log record.
15
16     Args:
17         pipeline_name: Job name, for example nat_silver_pipeline.
18         layer: Medallion layer the run wrote to.
19         rows_processed: Rows written by the run.
20         status: Final status of the run.
21         error_message: Error text for a failed run, otherwise None.
22
23     Returns:
24         A record ready to append to the run log table.
25     """
26     return {
27         "pipeline_name": pipeline_name,
28         "layer": layer,
29         "rows_processed": rows_processed,
30         "status": status,
31         "error_message": error_message,
32     }

```

1. **Line 3.** A named Literal type. mypy rejects any other string at every call site.
2. **Line 7.** The name now matches the job: it builds a record. Writing it is a separate function.
3. **Line 12.** Optional, stated honestly.
4. **Line 13.** The return type matches what the body returns.

The strong version tells mypy, the reviewer and the caller the same story.
The weak version tells three different ones.

Recap card

- **Optional parameter that defaults to None** param: str | None = None.
Never param: str = None.
- **A fixed set of string values** Literal["PASS", "WARN", "FAIL"] from typing.
- **Google docstring sections** One-line summary, then Args, Returns, Raises.
- **Why merge_keys: list[str] and not list?** mypy then rejects a string, which Python would split into characters.

Practice

Exercise 1. Multiple choice (level 1)

A function takes an optional error message that defaults to `None`. Which signature is correct for Python 3.11 with `mypy --strict`?

- a) `def log_run(error_message: str = None) -> None:`
- b) `def log_run(error_message: str | None = None) -> None:`
- c) `def log_run(error_message=None) -> None:`
- d) `def log_run(error_message: None = str) -> None:`

Why?

- r1) The hint must include `None`, because `None` is a value the parameter takes.
- r2) Type hints change how Python runs the function.

Exercise 2. Predict the output (level 2)

The reference code builds a MERGE condition like this. A caller passes a string instead of a list. What does this print?

```
1 merge_keys = "id"
2 condition = " AND ".join(f"t.{k} = s.{k}" for k in merge_keys)
3 print(condition)
```

Exercise 3. Fill the gap (level 2)

Complete the signature. `status` accepts only `"PASS"`, `"WARN"` or `"FAIL"`, and the function returns nothing.

```
1 from typing import ____ (1)
2
3 def record_check(name: str, status: ____ (2) ["PASS", "WARN", "FAIL"]) ->
   ____ (3):
4     ...
```

Exercise 4. Write the function (level 3)

Write `build_merge_condition(keys, target="t", source="s")`. It returns a condition such as `t.respondent_id = s.respondent_id AND t.wave_id = s.wave_id`.

- Raise `TypeError` when `keys` is a string, not a list.
- Raise `ValueError` when `keys` is empty.
- Add type hints and a Google-style docstring with Args, Returns and Raises.

Starter code:

```
1 def build_merge_condition(keys, target="t", source="s"):
```

Exercise 5. Refactor (level 4)

Type this function. Return one of "SKIP", "WARN" or "PASS" as a `Literal`, accept a missing `change`, and add a Google-style docstring. Keep the behaviour.

Starter code:

```
1 def classify_change(change, tolerance=0.10):
2     if change is None:
3         return "SKIP"
4     return "WARN" if change > tolerance else "PASS"
```

Sources

- Iterating a string in a join builds one clause per character. (verified: Run in Python 3.11 in this repository.)
- The `X | Y` union syntax works in annotations from Python 3.10. (verified: Python 3.10 release notes (PEP 604); used throughout this repository on Python 3.11.)

A1-L3 Config over hardcoded values, and PEP 8 with ruff

Concept

Pipeline code should contain logic. Everything that changes between environments, trackers or contract versions belongs in configuration:

- catalog, schema and table names
- column maps and type maps
- thresholds, such as `weight_sum_min` or a stability tolerance
- contract rules and their severity

Keep configuration in one of two places. Versioned JSON or YAML files in the repository suit values that change with a code release. Registry tables in Unity Catalog suit values that analysts maintain, such as the Variable Registry column map.

Load configuration once, at the start of the run, into a typed object. A frozen dataclass or a pydantic model works. Validate on load: every required key present, `weight_sum_min` below `weight_sum_max`, names non-empty. Fail at 00:01 with "missing key: silver_table", not at 02:47 with a `KeyError` deep in a transform.

Then pass the object, or the values from it, into your functions. A function that reads a global config dictionary is as hard to test as one with hardcoded strings.

Style is part of the same discipline. PEP 8 is the Python style guide. **ruff** checks it and formats code in one fast tool:

- `ruff check` . finds errors and style problems, such as undefined names (F821) or `== True` (E712).
- `ruff format` . rewrites layout consistently. `ruff format --check` . fails CI when a file needs formatting.

Settle the rules once in `pyproject.toml` under `[tool.ruff]`, run both commands in CI, and stop debating layout in reviews.

Why it matters

A developer copies the production silver notebook to test a change in the dev workspace. The notebook hardcodes `research_prod.research_silver.nat_respondents`. The dev run MERGEs half-finished test data into the production table. The morning report shows impossible reach figures, and the team restores the table from Delta history. With the catalog in a config file per environment, the dev run writes to `research_dev`. The same code runs in both places, and only the config changes.

Worked example: weak code

```
1 def check_weight_sum(spark, wave_id):
2     total = (
3         spark.table("research_prod.research_silver.nat_respondents")
4         .filter(F.col("wave_id") == wave_id)
5         .agg(F.sum("final_weight")).collect()[0][0]
6     )
7     if total < 45000000 or total > 55000000:
8         raise ValueError("weight sum out of range")
```

1. **Line 3.** The production table is typed into the function. A dev run writes, or here reads, production.
2. **Line 7.** Contract thresholds are buried in code. A new contract version needs a code change in every notebook that repeats them.
3. **Line 8.** The error hides the numbers. You cannot tell how far out of range the total is.

Worked example: strong code

```
1 import json
```

```

2 from dataclasses import dataclass
3 from pathlib import Path
4
5
6 ① @dataclass(frozen=True)
7 class WeightContract:
8     """Weight total limits from the data contract."""
9
10    silver_table: str
11    weight_sum_min: float
12    weight_sum_max: float
13
14 ②    def __post_init__(self) -> None:
15        if self.weight_sum_min >= self.weight_sum_max:
16            raise ValueError("weight_sum_min must be below
17                               weight_sum_max")
18
19    def load_contract(path: Path) -> WeightContract:
20        """Load and validate the contract file. Fails fast on a missing
21           key."""
22        raw = json.loads(path.read_text())
23        return WeightContract(
24            silver_table=raw["silver_table"],
25            weight_sum_min=float(raw["weight_sum_min"]),
26            weight_sum_max=float(raw["weight_sum_max"]),
27        )
28
29 ④ def weight_sum_in_range(total: float, contract: WeightContract) -> bool:
30     """True when the weight total is inside the contract limits."""
31     return contract.weight_sum_min <= total <= contract.weight_sum_max

```

1. **Line 6.** Frozen: no code can change a limit halfway through a run.
2. **Line 14.** Validation on load. An impossible contract fails before any data moves.
3. **Line 23.** A missing key raises `KeyError` here, at load, with the key's name.
4. **Line 29.** Pure logic with the limits passed in. Test it with plain numbers, no Spark needed.

The strong version runs unchanged in dev and prod. The contract file decides the table and the limits.

Recap card

- **What belongs in configuration?** Anything that changes between environments, trackers or contract versions: names, maps, thresholds, rules.
- **When do you validate configuration?** Once, on load, before any data moves. Fail fast with the key's name.

- **Why not `raw.get("catalog", "research_prod")`?** A default hides a missing key and can point a run at production.
- **The two ruff commands for CI** `ruff check .` and `ruff format --check .`

Practice

Exercise 1. Multiple choice (level 1)

Which of these belongs in configuration rather than in pipeline code?

- a) The formula for weighted reach: weighted reach divided by weighted base
- b) The catalog name, `research_prod` or `research_dev`
- c) The order of steps: read, validate, transform, write
- d) The decision to MERGE rather than append in silver

Why?

- r1) It changes between environments while the logic stays the same.
- r2) Short strings belong in configuration.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 from dataclasses import FrozenInstanceError, dataclass
2
3
4 @dataclass(frozen=True)
5 class RunConfig:
6     catalog: str
7
8
9 cfg = RunConfig("research_dev")
10 try:
11     cfg.catalog = "research_prod"
12 except FrozenInstanceError:
13     print("blocked")
14 print(cfg.catalog)
```

Exercise 3. Fill the gap (level 2)

Complete the two CI commands. The first finds lint errors. The second fails when a file needs reformatting, without changing it.

```
1 ruff ____ (1) .
2 ruff ____ (2) --check .
```

Exercise 4. Write the function (level 4)

Write `load_run_config(raw)`. It takes a dict parsed from JSON and returns a frozen `RunConfig` dataclass with `catalog` (str), `tracker` (str) and `stability_tolerance` (float).

- Raise `KeyError` naming the key when one is missing.
- Raise `ValueError` when `stability_tolerance` is not between 0 and 1.
- Add a property `silver_table` that returns `f"{catalog}.research_silver.{tracker}_respondents"`.

Starter code:

```
1 from dataclasses import dataclass
2
3
4 @dataclass(frozen=True)
5 class RunConfig:
6     ...
7
8
9 def load_run_config(raw: dict[str, object]) -> RunConfig:
10     ...
```

Exercise 5. Spot the bug (level 5)

The team moved every table name into `RunConfig`. A dev run still reads production. Tap the line, then fix the function.

```
1 def weighted_base_by_wave(spark: SparkSession, cfg: RunConfig) ->
    DataFrame:
2     """Weighted base per wave from the configured silver table."""
3     silver = spark.table(cfg.silver_table)
4     scope = spark.table("research_prod.research_governance.scope_rules")
5     in_scope = silver.join(scope, "region").filter(F.col("in_scope"))
6     return
    in_scope.groupBy("wave_id").agg(F.sum("final_weight").alias("weighted_base"))
```

Write down the line number of the bug. Then write the fix.

Sources

- ruff check reports F821 (undefined name) and E712 (comparison to True); ruff format --check fails when a file needs formatting. (verified: ruff 0.16.9 run in this repository.)
- Assigning to a frozen dataclass field raises `FrozenInstanceError`. (verified: Python 3.11 standard library, run in this repository.)

A2 Medallion architecture

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Bronze, silver and gold contracts.

A2-L1 Bronze: keep every delivery exactly as received

Concept

Bronze is your evidence. Its contract has four rules.

1. **Append only.** Every delivery adds rows. Nothing in bronze is updated, deleted or overwritten. A redelivery of wave W14 appends again, with a new `ingestion_id`.
2. **As received.** No cleaning, filtering, renaming or type casting. Read CSV columns as strings. `inferSchema` guesses types, and a guess can destroy data: the ID `00123` becomes the integer `123`, and the leading zeros are gone for good.
3. **Audit columns.** Add four columns to every row: `ingestion_id` (one UUID per load), `source_file`, `ingestion_timestamp` and `delivery_version`. Later layers use them to pick the newest version and to trace any number back to a file.
4. **Grain.** One row per respondent per wave per delivery file. Duplicates across deliveries are expected here. Silver resolves them.

Why keep the mess? Because silver logic changes. When you fix a mapping bug in March, you rebuild silver from bronze for every wave since January. You can only do that if bronze still holds the original deliveries.

Two practical points for Databricks Runtime 17.3 LTS:

- Count what you wrote from the Delta history instead of calling `count()` before the write. `DeltaTable.forName(spark, table).history(1)` returns `operationMetrics["numOutputRows"]`. A separate `count()` reads the source file a second time.
- Open-source Spark defaults `saveAsTable` to Parquet. Databricks defaults to Delta. Write `.format("delta")` explicitly, so the code behaves the same in local tests.

Why it matters

A bronze job reads supplier CSVs with `inferSchema=true`. Respondent IDs such as `00123` land as the integer `123`. Months later, the supplier sends a correction file keyed on the original IDs, and 18% of rows fail to match silver. The original IDs are not recoverable from bronze, because bronze never held them. The team asks the supplier to resend six months of deliveries.

Reading every column as a string keeps `00123` intact. Silver can then cast deliberately, with checks.

Worked example: weak code

```
1 def ingest(path, table):
2   ❶ df = spark.read.option("header", True).option("inferSchema",
      True).csv(path)
3   ❷ df = df.filter(F.col("final_weight").isNotNull())
4   ❸ print(df.count())
5   ❹ df.write.mode("overwrite").saveAsTable(table)
```

1. **Line 2.** Type guessing. `00123` becomes `123`, and the original value is lost.
2. **Line 3.** Cleaning in bronze. The rows with missing weights vanish before anyone can count them.
3. **Line 4.** Reads the whole file a second time, only to print a number.
4. **Line 5.** Overwrite destroys every earlier delivery.

Worked example: strong code

```
1 import uuid
2
3 from delta.tables import DeltaTable
4 from pyspark.sql import DataFrame, SparkSession
5 from pyspark.sql import functions as F
6
7
8 ❶ def add_audit_columns(df: DataFrame, ingestion_id: str, source_file: str,
      delivery_version: int) -> DataFrame:
9     """Add the four bronze audit columns."""
10    return df.withColumns({
11        "ingestion_id": F.lit(ingestion_id),
12        "source_file": F.lit(source_file),
13        ❷ "ingestion_timestamp": F.current_timestamp(),
14        "delivery_version": F.lit(delivery_version),
15    })
16
17
18 def ingest_to_bronze(spark: SparkSession, path: str, table: str,
      delivery_version: int) -> int:
19     """Append one delivery file to bronze as received. Returns rows
      written."""
```

```

3      raw = spark.read.option("header", True).csv(path)
21     audited = add_audit_columns(raw, str(uuid.uuid4()), path,
        delivery_version)
4      audited.write.format("delta").mode("append").saveAsTable(table)
5      metrics = DeltaTable.forName(spark, table).history(1).collect()[0]
        ["operationMetrics"]
24     return int(metrics["numOutputRows"])

```

1. **Line 8.** A pure transform. Test it with two rows and no files.
2. **Line 13.** When the row landed, from the cluster clock.
3. **Line 20.** No inferSchema: every column arrives as a string, exactly as delivered.
4. **Line 22.** Append only, and Delta stated explicitly.
5. **Line 23.** The row count comes from the commit metrics. No second read of the file.

The strong version keeps every byte the supplier sent, labels every row with its origin, and reads the file once.

Recap card

- **The four rules of the bronze contract** Append only. As received. Audit columns. Grain: respondent per wave per delivery file.
- **The four bronze audit columns** ingestion_id, source_file, ingestion_timestamp, delivery_version.
- **Why no inferSchema in bronze?** It guesses types and can destroy values, such as the leading zeros in 00123.
- **Rows written, without a second read** DeltaTable.forName(spark, t).history(1) and operationMetrics["numOutputRows"].

Practice

Exercise 1. Multiple choice (level 1)

A delivery contains 12 rows with a null `final_weight`. What should the bronze job do with them?

- a) Drop them, so silver never sees bad rows.
- b) Append them as received, with the audit columns.
- c) Replace the null weights with 1.0.
- d) Reject the whole file.

Why?

- r1) Bronze records deliveries as received. Silver and the data contract decide what to do with bad rows.
- r2) Delta tables cannot hold null values.

Exercise 2. Predict the output (level 2)

`ids.csv` holds a header line `respondent_id` and one data line `00123`. What does this print on Spark 4.0?

```
1 df = spark.read.option("header", True).option("inferSchema",  
    True).csv("ids.csv")  
2 print(df.collect()[0]["respondent_id"])
```

Exercise 3. Fill the gap (level 2)

Complete the bronze write and the timestamp audit column.

```
1 audited = raw.withColumn("ingestion_timestamp", F.____(1))  
2 audited.write.format("delta").mode("____(2)").saveAsTable(table)
```

Exercise 4. Write the function (level 3)

Write `add_audit_columns(df, ingestion_id, source_file, delivery_version)`. It returns `df` with four extra columns: `ingestion_id`, `source_file`, `ingestion_timestamp` (current timestamp) and `delivery_version`. Do not change any existing column. Use type hints and a docstring.

Starter code:

```
1 from pyspark.sql import DataFrame  
2 from pyspark.sql import functions as F  
3  
4  
5 def add_audit_columns(df: DataFrame, ingestion_id: str, source_file: str,  
    delivery_version: int) -> DataFrame:  
6     ...
```

Exercise 5. Spot the bug (level 4)

Run this job twice for two deliveries of the same wave. Bronze ends with only the second delivery. Tap the line, then fix the function.

```
1 def ingest_to_bronze(spark: SparkSession, path: str, table: str,  
    delivery_version: int) -> None:  
2     """Append one delivery file to bronze as received."""  
3     raw = spark.read.option("header", True).csv(path)  
4     audited = add_audit_columns(raw, str(uuid.uuid4()), path,  
        delivery_version)  
5     audited.write.format("delta").mode("overwrite").saveAsTable(table)
```

Write down the line number of the bug. Then write the fix.

Exercise 6. Architecture design (level 5)

Assign each piece of work to the layer that owns it.

Options: Bronze, Silver, Gold.

1. Store the supplier's CSV with every column as a string
2. Rename wgt to final_weight and cast it to DOUBLE
3. Keep only the newest delivery version per respondent and wave
4. Weighted reach by wave, age group and region for Power BI

Practical: Build a bronze loader and prove it keeps every delivery

Platform: Databricks Free Edition. Run this in Databricks Free Edition. It builds two small deliveries of the same wave in memory, loads both into bronze, and shows them side by side. In production the rows come from `spark.read.option("header", True).csv(path)`.

Download the notebook `A2-L1.ipynb` from the app.

Sources

- inferSchema turns 00123 into the integer 123. (verified: Run on PySpark 4.0.4 in this repository.)
- Delta history operationMetrics holds numOutputRows for an append. (verified: Run on delta-spark 4.0.1 in this repository.)
- DataFrame.withColumns exists (PySpark 3.3 and later). (verified: Checked on PySpark 4.0.4 in this repository.)
- Open-source Spark defaults saveAsTable to Parquet; Databricks defaults to Delta. (verified: OSS behaviour run locally; Databricks default per docs (VERSIONS.md).)

A2-L2 Silver: canonical names, checked types, one row per respondent per wave

Concept

Silver turns deliveries into a dependable table. Its contract:

- **Canonical names.** Rename supplier columns to the names in the Variable Registry, loaded at run time.
- **Declared types.** Cast each column to the type in the registry.
- **Grain.** One row per respondent per wave. Keep the newest delivery version per key, then MERGE (lesson A3-L1).

Two Spark behaviours make the first two steps fail silently or loudly, in the wrong place.

Renames skip missing columns. `withColumnRenamed("wgt", "final_weight")` does nothing when `wgt` is absent, and so does `withColumnsRenamed`. No error, no warning. If the supplier renames `wgt` to `weight`, silver loses its weight column, and the failure surfaces hours later in gold. Check the source columns yourself before you rename, and raise one error that lists every missing column.

Casts under ANSI mode. Databricks Runtime 17.3 LTS runs Spark 4.0, which turns ANSI mode on by default. `cast("abc" as int)` now raises `CAST_INVALID_INPUT` and stops the job at the first bad value. Older runtimes returned NULL without a word. Neither is what you want in silver. Use `try_cast`, which returns NULL for a bad value, then count the new NULLs and let a data-quality check decide: under 0.1% is a WARN, above it is CRITICAL.

```
1 F.col("final_weight").try_cast("double")
```

The pattern is: check columns, rename, `try_cast`, count cast failures, deduplicate, MERGE. Each step is a small function you can test.

Why it matters

A supplier renames `wgt` to `weight` without notice. The silver job calls `withColumnRenamed("wgt", "final_weight")`, which silently does nothing. The type-casting loop skips `final_weight` because the column is missing. Silver writes the wave with no weight column, since schema evolution is on. At 03:00 the gold job fails with `final_weight` not found. The morning report shows nothing for the new wave.

A column check before the rename fails the silver job at 01:00 with one clear message: "missing source columns: wgt". The on-call engineer knows exactly what to ask the supplier.

Worked example: weak code

```
1 def transform_to_silver(bronze_df, column_map, type_map, wave_id):
2     for src, tgt in column_map.items():
3         bronze_df = bronze_df.withColumnRenamed(src, tgt)
4     for col_name, col_type in type_map.items():
5         bronze_df = bronze_df.withColumn(col_name,
6                                           col(col_name).cast(col_type))
7     bronze_df = bronze_df.withColumn("wave_id", lit(wave_id))
8     return bronze_df
```

1. **Line 3.** A missing source column is skipped without an error.

2. **Line 5.** Under ANSI mode one bad value stops the job. Without ANSI, bad values become NULL with no count.
3. **Line 6.** Overwrites the `wave_id` from the delivery, hiding a file loaded against the wrong wave.

Worked example: strong code

```

1 def require_columns(df: DataFrame, columns: list[str]) -> None:
2     """Raise ValueError listing every column missing from df."""
3     missing = sorted(set(columns) - set(df.columns))
4     if missing:
5         raise ValueError(f"missing source columns: {missing}")
6
7
8 def rename_to_canonical(df: DataFrame, column_map: dict[str, str]) ->
    DataFrame:
9     """Rename supplier columns to registry names. Fails on any missing
    column."""
10    require_columns(df, list(column_map))
11    return df.withColumnsRenamed(column_map)
12
13
14 def cast_to_registry_types(df: DataFrame, type_map: dict[str, str]) ->
    DataFrame:
15
16    """Cast with try_cast, so a bad value becomes NULL instead of stopping
    the job."""
17    require_columns(df, list(type_map))
18    return df.withColumns({c: F.col(c).try_cast(t) for c, t in
19                          type_map.items()})
20
21
22 def count_cast_failures(df: DataFrame, type_map: dict[str, str]) ->
    dict[str, int]:
23
24    """Per column, rows whose non-null value fails try_cast. One pass over
    df."""
25    failures = [
26        F.count_if(F.col(c).isNotNull() &
27                  F.col(c).try_cast(t).isNull()).alias(c)
28        for c, t in type_map.items()
29    ]
30    return df.agg(*failures).collect()[0].asDict()

```

1. **Line 5.** One error that names every missing column.
2. **Line 11.** Safe to call now: every source column exists.
3. **Line 17.** `Column.try_cast` (PySpark 4.0) returns NULL for a bad value.
4. **Line 23.** Run on the uncast frame: a value that is present but fails the cast is a failure. All columns counted in one agg pass.

The strong version fails early and clearly on a changed delivery, and turns bad values into a number you can check instead of a crash or a silent NULL.

Recap card

- **What does `withColumnRenamed` do with a missing column?**
Nothing, silently. Check the source columns yourself and raise one error listing every missing one.
- **cast on Spark 4.0 (Runtime 17.3 LTS) with a malformed value**
Raises `CAST_INVALID_INPUT`, because ANSI mode is on by default.
- **Safe cast for silver** `F.col(c).try_cast(t)`, then count the new NULLs and check them against a threshold.
- **Silver dedup key** The full grain: `respondent_id` and `wave_id`.

Practice

Exercise 1. Predict the output (level 1)

What does this print on PySpark 4.0?

```
1 df = spark.createDataFrame([(1, 2.5)], "respondent_id INT, weight DOUBLE")
2 print(df.withColumnRenamed("wgt", "final_weight").columns)
```

Exercise 2. Multiple choice (level 2)

On Databricks Runtime 17.3 LTS, silver runs

`F.col("final_weight").cast("double")`. One row holds `"n/a"`. What happens?

- a) The row gets NULL in `final_weight`, and the job succeeds.
- b) The job fails with `CAST_INVALID_INPUT`.
- c) The row is dropped.
- d) The value becomes 0.0.

Why?

- r1) Spark 4.0 enables ANSI mode by default, and ANSI casts raise on malformed input.
- r2) Delta rejects NULL values in DOUBLE columns.

Exercise 3. Fill the gap (level 2)

Cast every column in `type_map` so that a bad value becomes NULL instead of stopping the job.

```
1 df = df.withColumns({c: F.col(c).____(1)(t) for c, t in
    type_map.items()})
```

Exercise 4. Write the function (level 3)

Write `rename_to_canonical(df, column_map)`. It raises `ValueError` naming every missing source column, then renames all columns in one call. Use type hints and a docstring.

Starter code:

```
1 from pyspark.sql import DataFrame
2
3
4 def rename_to_canonical(df: DataFrame, column_map: dict[str, str]) ->
    DataFrame:
5     ...
```

Exercise 5. Spot the bug (level 4)

Silver's grain is one row per respondent per wave. After loading waves W1 and W2, silver holds only one row per respondent. Tap the line, then fix the function.

```
1 def prepare_silver(df: DataFrame, column_map: dict[str, str], type_map:
    dict[str, str]) -> DataFrame:
2     """Rename, cast and keep the newest version per respondent and
    wave."""
3     renamed = rename_to_canonical(df, column_map)
4     typed = cast_to_registry_types(renamed, type_map)
5     return latest_per_key(typed, ["respondent_id"], "delivery_version")
```

Write down the line number of the bug. Then write the fix.

Exercise 6. Put in order (level 5)

Put the silver steps in order.

- Count cast failures and apply the WARN or CRITICAL threshold
- Rename to canonical names
- Check that every mapped source column exists
- MERGE into silver
- Keep the newest version per respondent and wave
- try_cast to registry types

Sources

- `withColumnRenamed` and `withColumnsRenamed` ignore missing columns without an error. (verified: Run on PySpark 4.0.4 in this repository.)
- Spark 4.0 enables ANSI mode by default; `cast` raises `CAST_INVALID_INPUT` on malformed input; `try_cast` returns `NULL`. (verified: Run on PySpark 4.0.4)

in this repository. Databricks Runtime 17.3 LTS uses Spark 4.0.0 (search-only, VERSIONS.md).

- Column `try_cast` exists in PySpark 4.0. (verified: Run on PySpark 4.0.4 in this repository.)

A2-L3 Gold: state the grain, group by all of it, guard every ratio

Concept

Gold tables feed Power BI and clients. The contract has three parts.

1. State the grain. Write it in the function docstring and on the table. "One row per tracker x wave x age group x gender x region." Store it as a table comment or a custom table property:

```
1 COMMENT ON TABLE research_gold.nat_metrics IS 'Grain: tracker x wave x
   age_group x gender x region';
2 ALTER TABLE research_gold.nat_metrics SET TBLPROPERTIES ('grain' =
   'tracker x wave x age_group x gender x region');
```

Avoid the property name `owner`. It is reserved, and Spark 4.0 rejects it with `UNSUPPORTED_FEATURE.SET_TABLE_PROPERTY`. Ownership is set with `ALTER TABLE ... OWNER TO`.

2. Group by every grain column. If the grain names `tracker_id`, the `groupBy` must include `tracker_id`. Miss one and two trackers' respondents fall into the same row. The numbers look plausible and are wrong.

3. Make aggregates safe.

- A `sum` over no matching rows returns NULL, not 0. `F.sum(F.when(F.col("reach_flag") == 1, 1))` gives NULL for a subgroup with no reach. Wrap it: `F.coalesce(F.sum(...), F.lit(0))`.
- Guard every ratio. `F.when(F.col("weighted_base") > 0, reach / base)` returns NULL for an empty base instead of an error or an infinite value.
- Keep the unweighted count next to every weighted figure. Analysts need it to judge whether a subgroup is too small to report.

Gold is also where you check totals against silver. The sum of `weighted_base` across all subgroups for a wave must equal the silver weight total for in-scope respondents. Lesson A6-L3 builds that check.

Why it matters

The reference `compute_gold_reach_metrics` promises "tracker x wave x subgroup grain" but groups by wave and subgroup only. When a second tracker starts writing to the same silver table, its respondents fold into the first tracker's rows. Weighted bases roughly double, and reach percentages blend two audiences. The dashboard shows smooth, believable figures for three weeks, until a client compares a subgroup against its published sample size.

Adding `tracker_id` to `groupBy`, and a test that checks the output has one row per grain key, prevents it.

Worked example: weak code

```
1 from pyspark.sql.functions import col, count, sum, when
2
3 def compute_gold_reach_metrics(silver_table, wave_id,
4     weight_col="final_weight"):
5     """Compute reach metrics at tracker x wave x subgroup grain."""
6     silver = spark.table(silver_table).filter(col("wave_id") == wave_id)
7     return (
8         silver.filter(col("in_scope") == True)
9         .groupBy("wave_id", "age_group", "gender", "region")
10        .agg(
11            count("*").alias("unweighted_n"),
12            sum(weight_col).alias("weighted_base"),
13            sum(when(col("reach_flag") == 1,
14                col(weight_col))).alias("weighted_reach"),
15        )
16        .withColumn("weighted_reach_pct", col("weighted_reach") /
17            col("weighted_base"))
18    )
```

1. **Line 1.** Imports `sum` by name, shadowing the Python built-in (reference defect 3).
2. **Line 8.** The docstring promises tracker grain, but `tracker_id` is missing.
3. **Line 12.** NULL, not 0, for a subgroup where nobody is reached.
4. **Line 14.** No guard on the base, and NULL reach gives a NULL percentage.

Worked example: strong code

```
1 from pyspark.sql import DataFrame
2 from pyspark.sql import functions as F
3
4 GRAIN = ["tracker_id", "wave_id", "age_group", "gender", "region"]
5
6
7 def reach_metrics(df: DataFrame, weight_col: str = "final_weight") ->
8     DataFrame:
```

```

8      """Weighted reach at tracker x wave x age_group x gender x region
      grain.
9
10     Args:
11         df: In-scope silver respondents with reach_flag and weights.
12         weight_col: Weight column to aggregate.
13
14     Returns:
15         One row per GRAIN key with unweighted_n, weighted_base,
16         weighted_reach
17         and weighted_reach_pct (NULL when the base is 0).
18         """
19     reached = F.when(F.col("reach_flag") == 1, F.col(weight_col))
20     metrics = df.groupBy(*GRAIN).agg(
21         F.count("*").alias("unweighted_n"),
22         F.coalesce(F.sum(weight_col), F.lit(0.0)).alias("weighted_base"),
23         F.coalesce(F.sum(reached), F.lit(0.0)).alias("weighted_reach"),
24     )
25     return metrics.withColumn(
26         "weighted_reach_pct",
27         F.when(F.col("weighted_base") > 0, F.col("weighted_reach") /
28             F.col("weighted_base")),
29     )

```

1. **Line 4.** The grain is one constant, used by the groupBy and by the tests.
2. **Line 19.** Groups by every grain column, tracker included.
3. **Line 22.** Zero, not NULL, when a subgroup has no reach.
4. **Line 26.** NULL for an empty base, never an error or infinity.

The strong version cannot merge two trackers, reports a true zero for zero reach, and never divides by nothing.

Recap card

- **Where do you record a gold table's grain?** In the function docstring and on the table: COMMENT ON TABLE or a custom property such as 'grain'.
- **sum over no matching rows** Returns NULL, not 0. Wrap it: F.coalesce(F.sum(...), F.lit(0)).
- **Safe ratio in PySpark** F.when(base > 0, part / base). NULL when the base is 0.
- **Why is SET TBLPROPERTIES ('owner' = ...) wrong?** owner is reserved. Spark 4.0 raises UNSUPPORTED_FEATURE.SET_TABLE_PROPERTY. Use ALTER TABLE ... OWNER TO.

Practice

Exercise 1. Spot the bug (level 1)

The docstring states the grain. The code does not match it. Tap the line, then fix the function.

```
1 def base_by_subgroup(df: DataFrame) -> DataFrame:
2     """Weighted base at tracker x wave x region grain."""
3     return df.groupBy("wave_id",
        "region").agg(F.sum("final_weight").alias("weighted_base"))
```

Write down the line number of the bug. Then write the fix.

Exercise 2. Predict the output (level 2)

Region S has one respondent, and nobody in S is reached. What does this print?

```
1 df = spark.createDataFrame([("S", 0)], "region STRING, reach_flag INT")
2 out = df.groupBy("region").agg(F.sum(F.when(F.col("reach_flag") == 1,
    1)).alias("reached"))
3 print(out.collect()[0]["reached"])
```

Exercise 3. Fill the gap (level 3)

Make reach a true zero when nobody is reached, and the percentage NULL when the base is zero.

```
1 F.____(1)(F.sum(reached), F.lit(0.0)).alias("weighted_reach")
2 ...
3 F.____(2)(F.col("weighted_base") > 0, F.col("weighted_reach") /
    F.col("weighted_base"))
```

Exercise 4. Write the function (level 4)

Write `reach_metrics(df, grain)`. It groups by every column in `grain` and returns `unweighted_n`, `weighted_base` (sum of `final_weight`, 0 when empty), `weighted_reach` (sum of `final_weight` where `reach_flag == 1`, 0 when none) and `weighted_reach_pct` (NULL when the base is 0).

Starter code:

```
1 from pyspark.sql import DataFrame
2 from pyspark.sql import functions as F
3
4
5 def reach_metrics(df: DataFrame, grain: list[str]) -> DataFrame:
6     ...
```

Exercise 5. Multiple choice (level 5)

You want the grain of `research_gold.nat_metrics` recorded on the table itself. Which statement fails on Spark 4.0?

- a) `COMMENT ON TABLE research_gold.nat_metrics IS 'Grain: tracker x wave x region'`
- b) `ALTER TABLE research_gold.nat_metrics SET TBLPROPERTIES ('grain' = 'tracker x wave x region')`
- c) `ALTER TABLE research_gold.nat_metrics SET TBLPROPERTIES ('owner' = 'data.engineering@example.org')`
- d) `ALTER TABLE research_gold.nat_metrics SET TBLPROPERTIES ('layer' = 'gold')`

Why?

- r1) `owner` is a reserved table property, so Spark rejects setting it this way.
- r2) Table properties cannot hold email addresses.

Sources

- `sum(when(cond, 1))` over rows where `cond` is false returns `NULL`. (verified: Run on PySpark 4.0.4 in this repository.)
- `'owner'` is a reserved table property; Spark 4.0 raises `UNSUPPORTED_FEATURE.SET_TABLE_PROPERTY`. (verified: Run on PySpark 4.0.4 with delta-spark 4.0.1 in this repository.)
- `COMMENT ON TABLE` and custom `TBLPROPERTIES` work on Delta tables. (verified: Run on delta-spark 4.0.1 in this repository.)

A3 Delta Lake

Roadmap: Months 16 to 18: MLOps and deployment; Current work. MERGE, idempotent writes, schema enforcement and evolution, time travel, maintenance.

A3-L1 Idempotent MERGE, and why you deduplicate the source first

Concept

A pipeline is **idempotent** when running it twice with the same input leaves the table in the same state as running it once. Retries, reruns and redeliveries happen every week in production, so idempotence is not optional.

An append is not idempotent. Run `write.mode("append")` twice and every row lands twice. A Delta `MERGE` is idempotent when two conditions hold:

1. The `ON` condition covers every column of the table's grain. For a silver table with one row per respondent per wave, that means `respondent_id` **and** `wave_id`.
2. The source holds at most one row per key.

The second condition catches people out. When two source rows match the same target row, Delta cannot tell which one should win. It stops the `MERGE` with the error `DELTA_MULTIPLE_SOURCE_ROW_MATCHING_TARGET_ROW_IN_MERGE`. When neither row matches a target row, both rows are inserted, and your grain is broken without any error.

So deduplicate the source before you merge. Rank the rows for each key by a version column, such as `delivery_version` or `ingestion_timestamp`, and keep the newest. Use a window with `row_number()`. Do not use `dropDuplicates()`: it keeps an arbitrary row per key, so a rerun keeps a different version with no warning.

The pattern for every silver load:

- Keep the newest source row per key.
- `MERGE` on the full grain.
- `whenMatchedUpdateAll()` for redelivered rows, `whenNotMatchedInsertAll()` for new ones.

Bronze stays append-only, because bronze records every delivery as received. Silver is where you resolve versions.

Why it matters

A supplier redelivers a wave with corrected weights. The file repeats 37 respondents, because the supplier's export joined two tables. The silver job uses the reference `upsert_silver`, which merges without deduplicating. The `MERGE` fails at 02:00 with

`DELTA_MULTIPLE_SOURCE_ROW_MATCHING_TARGET_ROW_IN_MERGE`. Gold does not refresh, and the morning Power BI report shows last week's numbers.

The on-call engineer "fixes" it by switching the job to append. The next run succeeds. Silver now holds the wave twice. Reach percentages look normal, because both the numerator and the denominator doubled. Weighted population estimates are twice their true size for a month, until a reconciliation check against the weight total catches it.

Deduplicating on the full grain before the `MERGE` prevents both failures.

Worked example: weak code

```
1 from delta.tables import DeltaTable
2 from pyspark.sql import DataFrame, SparkSession
3
4 spark = SparkSession.builder.getOrCreate()
5
6
7 def upsert_silver(new_df: DataFrame, target_table: str, merge_keys: list)
  -> None:
8     """Merge new data into silver table."""
9     delta_table = DeltaTable.forName(spark, target_table)
10    merge_condition = " AND ".join(
11        [f"target.{k} = source.{k}" for k in merge_keys]
12    )
13    (
14        delta_table.alias("target")
15        .merge(new_df.alias("source"), merge_condition)
16        .whenMatchedUpdateAll()
17        .whenNotMatchedInsertAll()
18        .execute()
19    )
```

1. **Line 4.** A session created at import time. Tests cannot inject a local session, and importing the module has a side effect.
2. **Line 7.** `list` with no element type. mypy cannot check that callers pass column names.

3. **Line 15.** The source goes straight into the MERGE. One repeated key in a redelivery fails the whole job.

Worked example: strong code

```
1 from delta.tables import DeltaTable
2 from pyspark.sql import DataFrame, SparkSession, Window
3 from pyspark.sql import functions as F
4
5
6 def latest_per_key(df: DataFrame, keys: list[str], order_col: str) ->
    DataFrame:
7     """Keep the newest row for each key.
8
9     Args:
10         df: Source rows, possibly with repeated keys.
11         keys: Columns that identify one target row.
12         order_col: Column that ranks versions. Highest wins.
13
14     Returns:
15         One row per key.
16     """
17
18     window = Window.partitionBy(*keys).orderBy(F.col(order_col).desc())
19     return (
20         df.withColumn("_rank", F.row_number().over(window))
21         .filter(F.col("_rank") == 1)
22         .drop("_rank")
23     )
24
25 def upsert_silver(
26     spark: SparkSession,
27     source: DataFrame,
28     target_table: str,
29     keys: list[str],
30     order_col: str,
31 ) -> None:
32     """MERGE the newest source row per key into a silver table.
33
34     Running it twice with the same source leaves the table unchanged.
35     """
36     condition = " AND ".join(f"t.{k} = s.{k}" for k in keys)
37     (
38         DeltaTable.forName(spark, target_table).alias("t")
39         .merge(latest_per_key(source, keys, order_col).alias("s"),
40             condition)
41         .whenMatchedUpdateAll()
42         .whenNotMatchedInsertAll()
43         .execute()
```

1. **Line 3.** Import the module as `F`. Importing `sum` or `max` by name shadows the Python built-ins.
2. **Line 17.** One partition per key, newest version first. The result is the same on every run.
3. **Line 19.** `row_number()` gives exactly one row rank 1 per key, even when versions tie. `rank()` would keep ties.
4. **Line 26.** The session is a parameter. A test passes a local session, a job passes the cluster session.
5. **Line 39.** Deduplicate inside the MERGE call, so no caller can skip it.

The strong version is idempotent for any source: repeated keys, a retry or a full redelivery. The weak version works only while the supplier never repeats a row.

Recap card

- **When is a pipeline idempotent?** When running it twice with the same input leaves the table as one run would.
- **Two conditions for an idempotent Delta MERGE** The ON condition covers the full grain, and the source has at most one row per key.
- **What happens when two source rows match one target row?** The MERGE fails with `DELTA_MULTIPLE_SOURCE_ROW_MATCHING_TARGET_ROW_IN_MERGE`.
- **Why not dropDuplicates() before a MERGE?** It keeps an arbitrary row per key. Use `row_number()` over a window ordered by version, newest first.

Practice

Exercise 1. Multiple choice (level 1)

The silver table holds one row for respondent `R1` in wave `W1`. A redelivery holds **two** rows for `R1` in `W1`. You run a MERGE on `respondent_id` and `wave_id` with `whenMatchedUpdateAll()` and `whenNotMatchedInsertAll()`. What happens?

- a) The target row takes the values of the last source row.
- b) The MERGE fails with an error.
- c) Both source rows are inserted as new rows.
- d) The target row takes the values of the first source row.

Why?

- r1) Two source rows match one target row, and Delta cannot tell which should win.
- r2) Delta enforces a primary key on the target table.
- r3) The source DataFrame must be cached before a MERGE.

Exercise 2. Fill the gap (level 2)

Complete the chain so that redelivered rows update, new rows insert, and the operation runs.

```
1 (
2     DeltaTable.forName(spark,
3         "research_silver.nat_respondents").alias("t")
4     .merge(source.alias("s"), "t.respondent_id = s.respondent_id AND
5         t.wave_id = s.wave_id")
6     ._____ (1) ()
7     ._____ (2) ()
8     ._____ (3) ()
9 )
```

Exercise 3. Predict the output (level 2)

`silver` starts empty. `source` holds 3 rows with 3 distinct (`respondent_id`, `wave_id`) keys. `load()` calls the strong `upsert_silver` from the worked example. What does the last line print?

```
1 load(source) # MERGE on respondent_id and wave_id
2 load(source) # the scheduler retries the same run
3 print(spark.table("silver").count())
```

Exercise 4. Spot the bug (level 3)

Silver has one row per respondent per wave. After wave `w2` loads, every wave `w1` row has disappeared. Tap the line with the bug, then fix the code.

```
1 def upsert_respondents(spark: SparkSession, source: DataFrame) -> None:
2     """MERGE one wave of respondents into silver. Grain: respondent per
3     wave."""
4     latest = latest_per_key(source, ["respondent_id", "wave_id"],
5         "delivery_version")
6     (
7         DeltaTable.forName(spark,
8             "research_silver.nat_respondents").alias("t")
9         .merge(latest.alias("s"), "t.respondent_id = s.respondent_id")
10        .whenMatchedUpdateAll()
11        .whenNotMatchedInsertAll()
12        .execute()
13    )
```

Write down the line number of the bug. Then write the fix.

Exercise 5. Write the function (level 4)

Write `latest_per_key(df, keys, order_col)`. It returns one row per key: the row with the highest `order_col`. Use type hints and a docstring. The checker reads your code without running Spark.

Starter code:

```
1 from pyspark.sql import DataFrame, Window
2 from pyspark.sql import functions as F
3
4
5 def latest_per_key(df: DataFrame, keys: list[str], order_col: str) ->
    DataFrame:
6     ...
```

Exercise 6. Put in order (level 5)

Put the steps of an idempotent silver load in order.

- Write one record to the pipeline run log
- Rename and cast columns from the Variable Registry
- Validate the bronze delivery against its data contract
- MERGE into silver on `respondent_id` and `wave_id`
- Run the table health check on silver
- Keep the newest row per respondent and wave

Practical: Prove a MERGE is idempotent

Platform: Databricks Free Edition. Run this in Databricks Free Edition on serverless compute. It creates a schema called `learning_a3` in your current catalog, loads the same data twice, and shows that the row count does not change. It then repeats a key and shows the error.

Download the notebook `A3-L1.ipynb` from the app.

Sources

- MERGE fails with `DELTA_MULTIPLE_SOURCE_ROW_MATCHING_TARGET_ROW_IN_MERGE` when several source rows match one target row. (verified: `practice/tests/test_a3_l1_merge.py` reproduces the error on delta-spark 4.0.1. Docs: <https://docs.databricks.com/aws/en/sql/language-manual/delta-merge-into>)
- `DeltaTable.forName`, `alias`, `merge`, `whenMatchedUpdateAll`, `whenNotMatchedInsertAll`, `whenNotMatchedBySourceDelete` and `execute`

exist in delta-spark 4.0. (verified: practice/tests/test_a3_l1_merge.py checks the methods on delta-spark 4.0.1.)

- Free Edition users can create schemas in the workspace catalog, whose name matches the workspace. (search-only: <https://docs.databricks.com/aws/en/catalogs/default>)

A3-L2 Schema enforcement is on by default. Evolve on purpose.

Concept

A Delta table has a schema, and writes must match it. This is **schema enforcement**, and it is on by default. Tested on delta-spark 4.0.1:

Write	Result
Append with an extra column	Fails: schema mismatch
Append with a column missing	Succeeds: the missing column is NULL
Append with a clashing type (INT into DOUBLE)	Fails: <code>DELTA_FAILED_TO_MERGE_FIELDS</code>
Overwrite with a different schema	Fails unless <code>overwriteSchema</code> is true
MERGE with <code>updateAll/insertAll</code> and an extra source column	Succeeds, and silently drops the extra column

Enforcement is your friend. A failing write is a cheap, early signal that a delivery changed.

Schema evolution changes the table's schema to accept the new shape. Delta offers it per write:

- Append: `.option("mergeSchema", "true")` adds new columns.
- MERGE: `.withSchemaEvolution()` on the merge builder adds new source columns.
- Planned change: `ALTER TABLE ... ADD COLUMN (region STRING)` in a reviewed migration.

Evolve on purpose, never by default. Two habits cause trouble:

- Leaving `mergeSchema` on every silver write. A supplier typo such as `final_wieght` becomes a new, permanent column, and the real weight is NULL for that wave.

- Setting `spark.databricks.delta.schema.autoMerge.enabled` for the session, which Databricks now labels legacy. It turns evolution on for every MERGE in the job, including ones you did not think about.

The safe pattern: validate the incoming columns against the registry first (lesson A2-L2). If a column is new and expected, add it with `ALTER TABLE` or evolve that one write. If it is unexpected, fail the run and ask the supplier.

Note the silent MERGE case in the table. Enforcement does not protect you there. Your column check does.

Why it matters

A silver job appends with `mergeSchema` always on, "to avoid schema errors". A supplier's export misspells `final_weight` as `final_wieght` for one wave. The append succeeds. Silver gains a permanent `final_wieght` column, and `final_weight` is NULL for 2,000 respondents. Gold sums NULL weights to a smaller base, and weighted reach for that wave is computed on the wrong population.

Without `mergeSchema`, the append fails with a schema mismatch naming the unexpected column. The fix takes minutes, and silver never holds the bad wave.

Worked example: weak code

```

1 spark.conf.set("spark.databricks.delta.schema.autoMerge.enabled", "true")
2
3
4 def append_wave(df, table):
5     df.write.format("delta").mode("append").option("mergeSchema",
        "true").saveAsTable(table)

```

1. **Line 1.** Turns evolution on for every MERGE in the session, not just the one you meant.
2. **Line 5.** Every append may now add columns. A typo in a delivery becomes part of the table.

Worked example: strong code

```

1 def check_columns_match(df: DataFrame, expected: list[str]) -> None:
2     """Fail when the delivery has unexpected or missing canonical
        columns."""
3     unexpected = sorted(set(df.columns) - set(expected))
4     missing = sorted(set(expected) - set(df.columns))
5     if unexpected or missing:
6         raise ValueError(f"schema drift: unexpected={unexpected}
        missing={missing}")
7

```

```

8
9 def append_wave(df: DataFrame, table: str, expected: list[str]) -> None:
10     """Append after an explicit column check. Delta enforcement stays
    on."""
11     check_columns_match(df, expected)
12     df.write.format("delta").mode("append").saveAsTable(table)
13
14
15 def add_planned_column(spark: SparkSession, table: str, column: str,
    sql_type: str) -> None:
16     """Planned schema change, run from a reviewed migration."""
17     spark.sql(f"ALTER TABLE {table} ADD COLUMN ({column} {sql_type})")

```

1. **Line 6.** Both directions checked. A missing column would otherwise land as NULL without an error.
2. **Line 12.** No mergeSchema. An unexpected shape fails loudly.
3. **Line 17.** New columns arrive through a reviewed change, not as a side effect of a write.

Schema changes become decisions a person reviews, not accidents a delivery causes.

Recap card

- **Delta append with an extra column, no options** Fails with a schema mismatch. Enforcement is on by default.
- **MERGE with updateAll and an extra source column, no evolution** Succeeds and silently ignores the extra column. Only your own column check catches it.
- **Evolve one write on purpose** Append: option mergeSchema. MERGE: .withSchemaEvolution(). Planned: ALTER TABLE ADD COLUMN.
- **Why not set autoMerge.enabled for the session?** It turns evolution on for every MERGE in the job, including ones you did not intend.

Practice

Exercise 1. Multiple choice (level 1)

Silver has columns `respondent_id`, `wave_id`, `final_weight`. You append a DataFrame with those three columns plus `region`, without any options. What happens on delta-spark 4.0?

- a) The append succeeds and region is added.
- b) The append succeeds and region is dropped.
- c) The append fails with a schema mismatch.
- d) Region is stored in a hidden rescue column.

Why?

- r1) Schema enforcement rejects writes with columns the table does not have.
- r2) Delta tables cannot change schema after creation.

Exercise 2. Predict the output (level 2)

Silver `s` has columns `id` and `w`. The source has `id`, `w` and `region`. What does the last line print on delta-spark 4.0?

```
1 (
2   DeltaTable.forName(spark, "s").alias("t")
3   .merge(source.alias("src"), "t.id = src.id")
4   .whenMatchedUpdateAll()
5   .whenNotMatchedInsertAll()
6   .execute()
7 )
8 print(spark.table("s").columns)
```

Exercise 3. Fill the gap (level 3)

You reviewed the supplier's new `region` column and want this one MERGE to add it to silver. Complete the builder.

```
1 (
2   DeltaTable.forName(spark, table).alias("t")
3   .merge(source.alias("s"), condition)
4   .____(1)()
5   .whenMatchedUpdateAll()
6   .whenNotMatchedInsertAll()
7   .execute()
8 )
```

Exercise 4. Write the function (level 4)

Write `schema_drift(actual, expected)`. It takes two lists of column names and returns a dict with sorted lists under `"unexpected"` (in actual, not expected) and `"missing"` (in expected, not actual). Then write `assert_no_drift(actual, expected)` that raises `ValueError` naming both lists when either is non-empty.

Starter code:

```
1 def schema_drift(actual: list[str], expected: list[str]) -> dict[str,
2   list[str]]:
3   ...
4
5 def assert_no_drift(actual: list[str], expected: list[str]) -> None:
6   ...
```

Exercise 5. Spot the bug (level 5)

This silver job evolved the schema on every run for a month, and silver now has 14 columns nobody asked for. Tap the line that causes it, then fix the function.

```
1 def upsert_silver(spark: SparkSession, source: DataFrame, table: str,
2   keys: list[str]) -> None:
3     """MERGE a deduplicated source into silver."""
4     spark.conf.set("spark.databricks.delta.schema.autoMerge.enabled",
5       "true")
6     condition = " AND ".join(f"t.{k} = s.{k}" for k in keys)
7     (
8       DeltaTable.forName(spark, table).alias("t")
9       .merge(source.alias("s"), condition)
10      .whenMatchedUpdateAll()
11      .whenNotMatchedInsertAll()
12      .execute()
13    )
```

Write down the line number of the bug. Then write the fix.

Sources

- Append with an extra column fails; with a missing column fills NULL; with a type clash fails with DELTA_FAILED_TO_MERGE_FIELDS; overwrite with a new schema needs overwriteSchema. (verified: Run on delta-spark 4.0.1 in this repository.)
- MERGE with updateAll/insertAll ignores extra source columns without evolution; DeltaMergeBuilder.withSchemaEvolution() adds them. (verified: Run on delta-spark 4.0.1 in this repository.)
- spark.databricks.delta.schema.autoMerge.enabled turns on schema evolution for MERGE across the session. (verified: <https://docs.databricks.com/aws/en/data-engineering/schema-evolution> (checked 2026-09-25))

A3-L3 Time travel, RESTORE, OPTIMIZE and VACUUM

Concept

Every Delta write creates a new table **version**. The transaction log keeps the list, and you can read any version still on storage.

- `DESCRIBE HISTORY t` lists versions with the operation, user and metrics.

- `spark.read.option("versionAsOf", 12).table("t")` or `SELECT * FROM t VERSION AS OF 12` reads an old version. `timestampAsOf` and `TIMESTAMP AS OF` work by time.
- `RESTORE TABLE t TO VERSION AS OF 12` makes version 12 current again, as a new version. History is kept, so you can undo the restore.

Time travel is your first response to a bad write: compare the versions, restore, then fix the job.

Two maintenance commands keep tables fast and storage bounded.

OPTIMIZE compacts many small files into fewer large ones. Frequent small appends leave thousands of tiny files, and every query pays to open them. `OPTIMIZE t ZORDER BY (wave_id)` also co-locates rows by a column. On Unity Catalog managed tables, Databricks recommends liquid clustering (`CLUSTER BY`, or `CLUSTER BY AUTO`) with predictive optimization, which runs `OPTIMIZE` and `VACUUM` for you.

VACUUM deletes data files that no current version needs and that are older than the retention window. The window comes from `delta.deletedFileRetentionDuration`, which defaults to 7 days. Delta blocks `VACUUM t RETAIN 0 HOURS` with a safety check, because a running reader or writer might still need those files.

`VACUUM` limits time travel. After it runs, versions whose files it removed can no longer be read. The log itself keeps entries for `delta.logRetentionDuration`, 30 days by default. From Databricks Runtime 18.0, time travel beyond `deletedFileRetentionDuration` is blocked outright. If you need 30 days of recoverable history, set both properties to at least 30 days.

Why it matters

A developer runs a gold rebuild against the wrong wave filter and overwrites three weeks of metrics. The team wants to restore, but a clean-up job ran `VACUUM ... RETAIN 0 HOURS` nightly, with the retention check switched off "to save storage". The overwritten files are gone. Three weeks of gold are rebuilt from silver overnight, and two client extracts go out late.

With the default 7-day retention, `RESTORE TABLE research_gold.nat_metrics TO VERSION AS OF` the last good version takes seconds.

Worked example: weak code

- 1 **SET** `spark.databricks.delta.retentionDurationCheck.enabled = false`;
- 2 **VACUUM** `research_gold.nat_metrics RETAIN 0 HOURS`;

```
3 -- gold was overwritten by mistake last night; rebuild everything
3 DROP TABLE research_gold.nat_metrics;
```

1. **Line 1.** Switches off the safety check that exists to stop the next line.
2. **Line 2.** Deletes every unreferenced file now. Time travel and RESTORE no longer work.
3. **Line 4.** With history intact, a RESTORE takes seconds. Dropping and rebuilding takes hours.

Worked example: strong code

```
1 -- 1. Find the last good version.
1 DESCRIBE HISTORY research_gold.nat_metrics;
3
4 -- 2. Compare before you act.
5 SELECT wave_id, SUM(weighted_base) AS base
2 FROM research_gold.nat_metrics VERSION AS OF 41
7 GROUP BY wave_id;
8
9 -- 3. Restore. This writes a new version, so it can be undone.
3 RESTORE TABLE research_gold.nat_metrics TO VERSION AS OF 41;
11
12 -- 4. Keep 30 days recoverable.
13 ALTER TABLE research_gold.nat_metrics SET TBLPROPERTIES (
4 'delta.deletedFileRetentionDuration' = 'interval 30 days',
15 'delta.logRetentionDuration' = 'interval 30 days'
16 );
```

1. **Line 2.** Each version with its operation, user and metrics such as numOutputRows.
2. **Line 6.** Read the old version side by side with the current one before restoring.
3. **Line 10.** Restore adds a version. The bad write stays in history for the post-mortem.
4. **Line 14.** VACUUM keeps unreferenced files for 30 days, so 30 days of versions stay readable.

History turns a data incident into a two-minute restore. Keep it long enough to be useful.

Recap card

- **Read an old Delta version** VERSION AS OF n or TIMESTAMP AS OF ts in SQL; option versionAsOf in PySpark.
- **What does RESTORE do to history?** Adds a new version with the old content. Nothing is deleted.
- **Default retention windows** Deleted files 7 days (delta.deletedFileRetentionDuration). Log 30 days (delta.logRetentionDuration).

- **Why is VACUUM ... RETAIN 0 HOURS dangerous?** It deletes the files that time travel and RESTORE need, and it breaks readers still using them.
- **Databricks' recommendation for layout and maintenance on UC managed tables** Liquid clustering (CLUSTER BY or CLUSTER BY AUTO) with predictive optimization.

Practice

Exercise 1. Multiple choice (level 1)

A job overwrote gold with bad data one hour ago. VACUUM has not run since. What is the fastest safe recovery?

- a) Drop the table and rebuild it from silver.
- b) RESTORE TABLE to the version before the bad write.
- c) Run VACUUM to remove the bad files.
- d) Run OPTIMIZE to rewrite the table.

Why?

- r1) The earlier version's files are still on storage, and RESTORE makes that version current again.
- r2) RESTORE deletes the bad version from history.

Exercise 2. Fill the gap (level 2)

Read version 41 of the gold table, then make it current.

```
1 SELECT * FROM research_gold.nat_metrics ____ (1) 41;
2 ____ (2) TABLE research_gold.nat_metrics TO VERSION AS OF 41;
```

Exercise 3. Predict the output (level 3)

A new Delta table gets these statements in order. What version number does the RESTORE create?

```
1 CREATE TABLE t (id INT) USING DELTA;           -- version 0
2 INSERT INTO t VALUES (1);                       -- version 1
3 INSERT INTO t VALUES (2);                       -- version 2
4 DELETE FROM t WHERE id = 1;                      -- version 3
5 RESTORE TABLE t TO VERSION AS OF 1;             -- version ?
```

Exercise 4. Spot the bug (level 4)

A nightly maintenance script. After an incident, the team cannot time travel even one day back. Tap the line that causes it, then fix the script.

```
1 OPTIMIZE research_silver.nat_respondents ZORDER BY (wave_id);
2 SET spark.databricks.delta.retentionDurationCheck.enabled = false;
3 VACUUM research_silver.nat_respondents RETAIN 0 HOURS;
```



```
4 ANALYZE TABLE research_silver.nat_respondents COMPUTE STATISTICS;
```

Write down the line number of the bug. Then write the fix.

Exercise 5. Architecture design (level 5)

The data governance policy says: "Any silver or gold change must be recoverable for 30 days." Choose the setting for each item.

Options: interval 7 days (default), interval 30 days, interval 0 hours.

1. delta.deletedFileRetentionDuration
2. delta.logRetentionDuration
3. Nightly VACUUM command

Sources

- DESCRIBE HISTORY, VERSION AS OF, versionAsOf, RESTORE TABLE and OPTIMIZE ... ZORDER BY work; RESTORE creates a new version; VACUUM RETAIN 0 HOURS fails the safety check. (verified: Run on delta-spark 4.0.1 in this repository.)
- deletedFileRetentionDuration defaults to 7 days and logRetentionDuration to 30 days; from Databricks Runtime 18.0 time travel beyond deletedFileRetentionDuration is blocked. (verified: <https://docs.databricks.com/aws/en/delta/vacuum> and <https://docs.databricks.com/aws/en/tables/history> (checked 2026-09-25))
- Databricks recommends liquid clustering and predictive optimization for Unity Catalog managed tables. (verified: <https://docs.databricks.com/aws/en/tables/clustering> (checked 2026-09-25))

A4 Unity Catalog

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Three-level namespace, grants, row filters and column masks.

A4-L1 The three-level namespace, one catalog per environment

Concept

Unity Catalog names every table with three parts: `catalog.schema.table`.

```
1 research_prod.research_silver.nat_respondents
```

Use each level for one purpose:

- **Catalog = environment** (or a major business domain): `research_dev`, `research_test`, `research_prod`. Permissions and storage attach at this level, so a dev job with no rights on `research_prod` cannot damage it.
- **Schema = layer or subject**: `research_bronze`, `research_silver`, `research_gold`, `research_governance`.
- **Table = content at a stated grain**: `nat_raw`, `nat_respondents`, `nat_metrics`.

Always write the full three-part name in pipeline code, built from configuration. A two-part name such as `research_silver.nat_respondents` resolves against the session's **current catalog**. That depends on the workspace default, a `USE CATALOG` somewhere earlier in the notebook, or a job setting. The same code then reads a different table in each place it runs. Build names in one function from config, so no code concatenates strings on its own:

```
1 def table_name(catalog: str, layer: str, tracker: str) -> str:
2     suffix = {"bronze": "raw", "silver": "respondents", "gold": "metrics"}
   [layer]
3     return f"{catalog}.research_{layer}.{tracker}_{suffix}"
```

In a notebook, `spark.catalog.currentCatalog()` tells you which catalog a two-part name would use. It is useful for checks, never as a substitute for naming the catalog.

Workspaces created before Unity Catalog also show `hive_metastore`, the legacy two-level store. Treat it as read-only history while you migrate.

Why it matters

A shared all-purpose cluster has `research_prod` as its default catalog. A developer tests a silver fix from a notebook that writes to `research_silver.nat_respondents`, a two-part name, believing the notebook "runs in dev". The MERGE lands in production. The fix was only half finished, and 1,400 rows get wrong region codes before anyone notices.

With a three-part name built from a dev config file, the same notebook writes to `research_dev.research_silver.nat_respondents`. The developer's lack of MODIFY on `research_prod` would also have blocked it.

Worked example: weak code

```
❶ spark.sql("USE CATALOG research_prod")
2
3 def load_silver(spark):
❷     return spark.table("research_silver.nat_respondents")
5
6 def write_gold(df):
❸     df.write.mode("append").saveAsTable("research_gold." + "nat" +
        "_metrics")
```

1. **Line 1.** Changes the current catalog for everything after it in the session.
2. **Line 4.** Two-part name. Which catalog it reads depends on the session, not the code.
3. **Line 7.** Names glued together in several places. One typo and a new table appears.

Worked example: strong code

```
❶ LAYER_SUFFIX = {"bronze": "raw", "silver": "respondents", "gold":
    "metrics"}
2
3
4 def table_name(catalog: str, layer: str, tracker: str) -> str:
5     """Three-level table name for a tracker and layer."""
6
7     Raises:
8         KeyError: If layer is not bronze, silver or gold.
9     """
❷     return f"{catalog}.research_{layer}.{tracker}_{LAYER_SUFFIX[layer]}"
11
12
13 def load_silver(spark: SparkSession, cfg: RunConfig) -> DataFrame:
14     """Read the silver table for the configured environment and
        tracker."""
❸     return spark.table(table_name(cfg.catalog, "silver", cfg.tracker))
```

1. **Line 1.** The naming convention lives in one constant.

2. **Line 10.** Always three parts. An unknown layer raises `KeyError` instead of inventing a table.
3. **Line 15.** The environment comes from config. The same code reads dev in dev and prod in prod.

The strong version reads exactly the table its config names, whatever the session's current catalog is.

Recap card

- **Unity Catalog name format** `catalog.schema.table`, for example `research_prod.research_silver.nat_respondents`.
- **What does each level hold in this project?** Catalog: environment. Schema: layer or governance. Table: tracker at a grain.
- **Why never write a two-part table name in pipeline code?** It resolves against the session's current catalog, which varies by cluster, notebook and job.
- **Where do table names come from?** One function built from config, such as `table_name(catalog, layer, tracker)`.

Practice

Exercise 1. Multiple choice (level 1)

Which level of the Unity Catalog namespace should separate dev from prod?

- a) Catalog
- b) Schema
- c) Table name prefix, such as `dev_nat_metrics`
- d) A column called environment

Why?

- r1) Permissions and storage attach at catalog level, so dev principals can be kept out of prod entirely.
- r2) Catalog names are shorter.

Exercise 2. Fill the gap (level 2)

Complete the fully qualified name of the gold table for tracker `nat` in production.

```
1 GOLD = "_____(1)._____(2).nat_metrics"
```

Exercise 3. Spot the bug (level 3)

This job runs in the dev workspace, but its writes appear in production. Tap the line, then fix the function.

```
1 def write_metrics(df: DataFrame, cfg: RunConfig) -> None:
2     """Write gold metrics for the configured environment."""
3     target = f"research_gold.{cfg.tracker}_metrics"
4     df.write.format("delta").mode("append").saveAsTable(target)
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 4)

Write `table_name(catalog, layer, tracker)`. It returns

`catalog.research_<layer>.<tracker>_<suffix>` where the suffix is `raw`, `respondents` or `metrics` for bronze, silver and gold. Raise `ValueError` naming the layer for any other layer, and raise `ValueError` if `catalog` contains a dot.

Starter code:

```
1 def table_name(catalog: str, layer: str, tracker: str) -> str:
2     ...
```

Exercise 5. Architecture design (level 5)

You lay out Unity Catalog for the research programme. Choose the level for each item.

Options: Catalog, Schema, Table.

1. Development versus production
2. Bronze, silver and gold
3. Run log, variable registry and contracts
4. One tracker's respondents at respondent x wave grain

Sources

- Unity Catalog uses a three-level namespace, `catalog.schema.table`, and two-part names resolve against the current catalog. (verified: <https://docs.databricks.com/aws/en/data-governance/unity-catalog/> (checked 2026-09-25))
- `spark.catalog.currentCatalog()` returns the current catalog (PySpark 3.4 and later). (verified: Present on PySpark 4.0.4 in this repository.)

A4-L2 Grant to groups, grant the least, grant at the right level

Concept

Unity Catalog controls access with privileges on securable objects. Four privileges cover most pipeline work:

Privilege	On	Lets the principal
USE CATALOG	catalog	see into the catalog. Grants no data access by itself.
USE SCHEMA	schema	see into the schema. Grants no data access by itself.
SELECT	table, or schema	read data
MODIFY	table, or schema	insert, update and delete data

To read a table, a principal needs all three: `USE CATALOG` on its catalog, `USE SCHEMA` on its schema and `SELECT` on the table. `MODIFY` also needs `SELECT` and both `USE` privileges.

Grant at the container. A privilege granted on a schema is inherited by every current and future table in it. `GRANT SELECT ON SCHEMA research_prod.research_gold` covers next month's new gold table without another ticket.

Grant to groups, never to individual users. People join, move and leave. A group grant is managed once, in one place. A user grant outlives the reason it was made, and nobody finds it until an audit.

Run pipelines as a service principal. Jobs that write bronze, silver and gold run as a service principal with `MODIFY` on those schemas. People get `SELECT` on the layers they need, and no person needs `MODIFY` on production.

```
1 GRANT USE CATALOG ON CATALOG research_prod TO `research_analysts`;  
2 GRANT USE SCHEMA, SELECT ON SCHEMA research_prod.research_gold TO  
  `research_analysts`;  
3 SHOW GRANTS ON SCHEMA research_prod.research_gold;
```

Avoid `ALL PRIVILEGES` outside sandboxes. It includes privileges you did not think about, and future privileges too.

Why it matters

During a deadline, an analyst gets `MODIFY` on `research_prod.research_silver` "for one day" to patch a mapping by hand. The grant is to their user account. Eight months later they run a clean-up notebook against what they believe is a dev copy and delete 12,000 silver rows. Nobody remembered the grant, and silver had to be rebuilt from bronze.

With the rule "people get `SELECT`, pipelines get `MODIFY`, grants go to groups", the patch would have gone through the pipeline, and the delete would have failed with a permission error.

Worked example: weak code

```
❶ GRANT ALL PRIVILEGES ON CATALOG research_prod TO `alex@example.org`;
❷ GRANT SELECT ON TABLE research_prod.research_gold.nat_metrics TO
  `research_analysts`;
❸ GRANT MODIFY ON SCHEMA research_prod.research_silver TO
  `research_analysts`;
```

1. **Line 1.** Everything, on all of production, to one person. Outlives the reason for it.
2. **Line 2.** `SELECT` alone is not enough, and a table-level grant misses every new gold table.
3. **Line 3.** Analysts can now change silver directly, bypassing the pipeline and its checks.

Worked example: strong code

```
1 -- Analysts read gold.
2 GRANT USE CATALOG ON CATALOG research_prod TO `research_analysts`;
❶ GRANT USE SCHEMA, SELECT ON SCHEMA research_prod.research_gold TO
  `research_analysts`;
4
5 -- The pipeline service principal writes silver and gold.
6 GRANT USE CATALOG ON CATALOG research_prod TO `research-pipeline-sp`;
❷ GRANT USE SCHEMA, SELECT, MODIFY ON SCHEMA research_prod.research_silver
  TO `research-pipeline-sp`;
8 GRANT USE SCHEMA, SELECT, MODIFY ON SCHEMA research_prod.research_gold TO
  `research-pipeline-sp`;
9
10 -- Check what is in place.
❸ SHOW GRANTS ON SCHEMA research_prod.research_gold;
```

1. **Line 3.** Schema level: inherited by every current and future gold table.
2. **Line 7.** Only the pipeline identity writes to silver.
3. **Line 11.** Review grants from SQL, and keep these statements in version control.

Three rules: groups not users, schemas not tables, SELECT for people and MODIFY for pipelines.

Recap card

- **Privileges needed to read a table** USE CATALOG on the catalog, USE SCHEMA on the schema, SELECT on the table or schema.
- **Why grant SELECT on the schema rather than the table?** Schema grants are inherited by every current and future table in it.
- **Who gets MODIFY on production silver and gold?** The pipeline service principal. People get SELECT.
- **Why never grant to a user account?** User grants outlive the reason for them. Group membership is managed in one place.

Practice

Exercise 1. Put in order (level 1)

Analysts must read every table in `research_prod.research_gold`. Put the grants in order from the outermost object inwards.

- GRANT USE SCHEMA ON SCHEMA `research_prod.research_gold` TO `research_analysts`
- GRANT SELECT ON SCHEMA `research_prod.research_gold` TO `research_analysts`
- GRANT USE CATALOG ON CATALOG `research_prod` TO `research_analysts`

Exercise 2. Multiple choice (level 2)

An analyst has SELECT on `research_prod.research_gold.nat_metrics` and nothing else. What happens when they query the table?

- a) The query succeeds.
- b) The query fails, because they lack USE CATALOG and USE SCHEMA.
- c) The query returns an empty result.
- d) The query succeeds for columns without a mask.

Why?

- r1) Reading needs USE on the parent catalog and schema as well as SELECT on the table.
- r2) SELECT only works when granted on a schema.

Exercise 3. Spot the bug (level 3)

A reviewer rejects this grant script. Tap the line that breaks the "groups, not users" rule, then rewrite the script.

```
1 GRANT USE CATALOG ON CATALOG research_prod TO `research_analysts`;
2 GRANT USE SCHEMA, SELECT ON SCHEMA research_prod.research_gold TO
  `research_analysts`;
3 GRANT USE SCHEMA, SELECT ON SCHEMA research_prod.research_silver TO
  sam.jones@example.org;
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 4)

Write the grants for the pipeline service principal `research-pipeline-sp`: it must read and write every table in `research_prod.research_silver`, now and in future. Give it only what it needs, on the right objects.

Starter code:

```
1 -- Grants for research-pipeline-sp
```

Exercise 5. Architecture design (level 5)

Assign each principal the least privilege that lets them do their job.

Options: USE CATALOG, USE SCHEMA, SELECT on the schema, USE CATALOG, USE SCHEMA, SELECT, MODIFY on the schema, ALL PRIVILEGES on the catalog, Nothing on this schema.

1. research_analysts build Power BI reports from gold
2. research-pipeline-sp writes silver
3. research_analysts and the bronze schema
4. Platform administrators managing the metastore

Sources

- Reading a table needs USE CATALOG, USE SCHEMA and SELECT; MODIFY needs SELECT and both USE privileges; privileges granted on a schema are inherited by current and future tables. (search-only: <https://docs.databricks.com/aws/en/data-governance/unity-catalog/manage-privileges/>)
- GRANT accepts a comma-separated list of privileges, and SHOW GRANTS lists grants on a securable. (verified: <https://docs.databricks.com/aws/en/sql/language-manual/security-grant> and <https://docs.databricks.com/aws/en/sql/language-manual/security-show-grant> (checked 2026-09-25))

A4-L3 Row filters and column masks in Unity Catalog

Concept

Grants decide who can read a table. **Row filters** and **column masks** decide what each reader sees inside it. Both are SQL user-defined functions that Unity Catalog applies at query time, for every query, from every tool.

Row filter. A SQL UDF that returns BOOLEAN. Unity Catalog passes it the columns you name, and drops each row where the function returns FALSE or NULL.

```
1 CREATE FUNCTION
  research_prod.research_governance.respondent_filter(tracker_id STRING)
2 RETURN is_account_group_member('research_approved_analysts') OR
  tracker_id = 'TRK-PUBLIC';
3
4 ALTER TABLE research_prod.research_silver.nat_respondents
5 SET ROW FILTER research_prod.research_governance.respondent_filter ON
  (tracker_id);
```

Column mask. A SQL UDF whose first parameter is the column value. It returns the value, or a masked version.

```
1 CREATE FUNCTION research_prod.research_governance.mask_postcode(postcode
  STRING)
2 RETURN CASE WHEN is_account_group_member('pii_readers') THEN postcode
  ELSE '***' END;
3
4 ALTER TABLE research_prod.research_silver.nat_respondents
5 ALTER COLUMN postcode SET MASK
  research_prod.research_governance.mask_postcode;
```

Three rules:

- Use `is_account_group_member()`. It checks account-level groups, which Unity Catalog uses. `is_member()` checks workspace-local groups.
- Test group membership, never `current_user() = 'someone@example.org'`. A named user in a function is a hidden grant nobody reviews.
- Keep filter and mask functions in a governance schema, so their owners and grants are separate from the data.

Your reference script uses `CREATE ROW ACCESS POLICY`. That is Snowflake syntax. Databricks rejects it. Where many tables share one rule, Databricks also offers attribute-based policies on governed tags.

Why it matters

The reference `unity_catalog.sql` was written with an AI assistant and never run. Its row-level security uses `CREATE ROW ACCESS POLICY`, which does not exist in Databricks. The deployment script stops at that statement, and the PII table goes live without a filter, because the grants earlier in the script had already succeeded. Every analyst with `SELECT` on silver can read restricted respondent rows until someone notices the failed step.

A SQL UDF with `SET ROW FILTER`, tested in dev, applies the rule where Databricks enforces it.

Worked example: weak code

```
❶ CREATE ROW ACCESS POLICY restrict_respondent_data
❷   ON research_silver.nat_respondents
3   USING (
❸       is_member('research_approved_analysts')
❹       OR current_user() = 'pipeline-svc@example.org'
6   );
```

1. **Line 1.** Snowflake syntax. Databricks raises a parse error. Reference defect 4.
2. **Line 2.** Two-part name: resolves against whatever catalog is current.
3. **Line 4.** Checks workspace-local groups. Unity Catalog groups are account-level.
4. **Line 5.** A named identity inside a function: a grant hidden from every access review.

Worked example: strong code

```
❶ CREATE OR REPLACE FUNCTION
    research_prod.research_governance.respondent_filter(tracker_id STRING)
2 RETURN is_account_group_member('research_approved_analysts')
❷   OR is_account_group_member('research_pipelines');
4
5 ALTER TABLE research_prod.research_silver.nat_respondents
❸ SET ROW FILTER research_prod.research_governance.respondent_filter ON
    (tracker_id);
7
8 CREATE OR REPLACE FUNCTION
    research_prod.research_governance.mask_postcode(postcode STRING)
9 RETURN CASE WHEN is_account_group_member('pii_readers') THEN postcode
    ELSE '***' END;
10
11 ALTER TABLE research_prod.research_silver.nat_respondents
❹ ALTER COLUMN postcode SET MASK
    research_prod.research_governance.mask_postcode;
```

1. **Line 1.** A SQL UDF in the governance schema. Its parameter receives the column named in ON (...).

2. **Line 3.** The pipeline service principal is in a group. No identity is named in the function.
3. **Line 6.** Binds the filter to the table. It applies to every query from every tool.
4. **Line 12.** Masks one column. Readers outside pii_readers see ***.

The strong script runs on Databricks, names groups instead of people, and uses three-part names throughout.

Recap card

- **Row filter in Databricks** A SQL UDF returning BOOLEAN, bound with ALTER TABLE t SET ROW FILTER f ON (col).
- **Column mask in Databricks** A SQL UDF taking the column value, bound with ALTER TABLE t ALTER COLUMN c SET MASK f.
- **is_account_group_member or is_member?** is_account_group_member: Unity Catalog groups are account-level. is_member checks workspace groups.
- **What is CREATE ROW ACCESS POLICY?** Snowflake syntax. Databricks rejects it. Reference defect 4.

Practice

Exercise 1. Multiple choice (level 1)

Which statement applies a row filter to a Unity Catalog table?

- a) CREATE ROW ACCESS POLICY p ON t USING (...)
- b) ALTER TABLE t SET ROW FILTER f ON (col)
- c) GRANT SELECT ON t TO g WHERE col = 'x'
- d) CREATE VIEW v AS SELECT * FROM t WHERE ... and nothing else

Why?

- r1) Databricks row filters are SQL UDFs bound to a table with SET ROW FILTER.
- r2) Row filters must be created before the table.

Exercise 2. Fill the gap (level 2)

Complete the mask so that only members of the account group pii_readers see real postcodes.

```
1 CREATE FUNCTION research_prod.research_governance.mask_postcode(postcode
  STRING)
2 RETURN CASE WHEN ____ (1)('pii_readers') THEN postcode ELSE '***' END;
3
4 ALTER TABLE research_prod.research_silver.nat_respondents
```

```
5 ALTER COLUMN postcode ____ (2)
   research_prod.research_governance.mask_postcode;
```

Exercise 3. Spot the bug (level 3)

An access review flags this filter. Tap the line that hides a grant from reviewers, then rewrite the function.

```
1 CREATE OR REPLACE FUNCTION
   research_prod.research_governance.respondent_filter(tracker_id STRING)
2 RETURN is_account_group_member('research_approved_analysts')
3 OR current_user() = 'pipeline-svc@example.org';
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 4)

Write a row filter for `research_prod.research_gold.nat_metrics`. Members of `research_all_regions` see every row. Everyone else sees only rows where `region = 'NATIONAL'`. Create the function `research_prod.research_governance.region_filter` and bind it to the table on the `region` column.

Starter code:

```
1 CREATE OR REPLACE FUNCTION
   research_prod.research_governance.region_filter(region STRING)
2 RETURN ...
```

Exercise 5. Architecture design (level 5)

Choose the control for each requirement on `nat_respondents`.

Options: Grants (or their absence), Row filter, Column mask.

1. The client team must not query the table at all
2. Analysts see only the trackers their group is approved for
3. Analysts see respondent rows, but postcodes appear as *** unless they are in `pji_readers`

Sources

- Row filters are SQL UDFs returning BOOLEAN bound with SET ROW FILTER ... ON (cols); column masks are bound with ALTER COLUMN ... SET MASK. (verified: <https://docs.databricks.com/aws/en/data-governance/unity-catalog/filters-and-masks/manually-apply> (checked 2026-09-25))
- `is_account_group_member` returns true when the session user is a direct or indirect member of an account-level group; `is_member` checks workspace groups. (verified: <https://docs.databricks.com/aws/en/sql/>)

language-manual/functions/is_account_group_member (checked 2026-09-25))

- Attribute-based access control (ABAC) policies apply filters and masks centrally through governed tags. (verified: <https://docs.databricks.com/aws/en/data-governance/unity-catalog/abac/policies> (checked 2026-09-25))

A5 PySpark performance

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Lazy evaluation, built-ins over UDFs, joins, partitioning, caching, explain plans.

A5-L1 Lazy evaluation, and why every `count()` re-runs the plan

Concept

Spark splits DataFrame operations into two kinds.

- **Transformations** build a plan and return a new DataFrame at once: `filter`, `select`, `withColumn`, `join`, `groupBy().agg()`. Nothing is read yet.
- **Actions** run the plan and return a result: `count()`, `collect()`, `show()`, `first()`, `toPandas()`, and every write.

This is **lazy evaluation**. It lets the optimiser see the whole plan before running it, so it can push filters into the file scan and skip columns you never use.

The cost is easy to miss: **each action runs the whole plan again**, from the source. Nothing is kept between actions unless you ask for it. A local test shows it plainly: a Python function counted its own calls inside a plan over 1,000 rows. After the first `count()`, it had run 1,000 times. After a second `count()`, 2,000.

Your reference data-contract validator calls `count()` more than ten times on the same uncached DataFrame: once for the row count, once per mandatory column, once per key and so on. Each call rereads the delivery from storage.

Three fixes, in order of preference:

1. **Combine the questions into one action.** One `agg` with many expressions answers them all in a single pass:

```
python df.agg( F.count("*").alias("rows"),  
              F.count_if(F.col("final_weight").isNull()).alias("null_weights"),  
              F.sum("final_weight").alias("weight_sum"), ).collect()[0]
```

1. **Materialise** an expensive intermediate result to a Delta table, then read that table for the follow-up actions.
2. **Cache** with `df.cache()` on classic clusters only. Serverless compute, which Free Edition uses, does not support cache APIs.

Why it matters

A contract validator checks a 40-million-row delivery with 23 separate `count()` calls, one per check. Each call rescans the source files in cloud storage. The validation step takes 38 minutes, longer than the transform it protects, and the team starts skipping it "for small redeliveries". A redelivery with duplicated keys then goes straight to silver.

Rewritten as one `agg` with 23 expressions, the same checks run in a single scan of about two minutes, and nobody has a reason to skip them.

Worked example: weak code

```
1 def check_delivery(df):
2     results = {}
3     results["rows"] = df.count()
4     results["null_ids"] =
5     df.filter(F.col("respondent_id").isNull()).count()
6     results["null_weights"] =
7     df.filter(F.col("final_weight").isNull()).count()
8     results["distinct_ids"] =
9     df.select("respondent_id").distinct().count()
10    return results
```

1. **Line 3.** Action 1: full scan of the source.
2. **Line 4.** Action 2: a second full scan.
3. **Line 5.** Action 3: a third.
4. **Line 6.** Action 4: a fourth scan, plus a shuffle.

Worked example: strong code

```
1 def check_delivery(df: DataFrame) -> dict[str, int]:
2     """Row, null and distinct counts for a delivery in one pass."""
3     row = df.agg(
4         F.count("*").alias("rows"),
5         F.count_if(F.col("respondent_id").isNull()).alias("null_ids"),
6         F.count_if(F.col("final_weight").isNull()).alias("null_weights"),
7         F.count_distinct("respondent_id").alias("distinct_ids"),
8     ).collect()[0]
9     return row.asDict()
```

1. **Line 3.** One action. Spark computes every expression during the same scan.
2. **Line 5.** `count_if` counts rows where a condition holds, inside the aggregation.
3. **Line 7.** Distinct count in the same pass. For very large data, `approx_count_distinct` is cheaper.

Four questions, one scan. The strong version reads the delivery once, however many checks you add.

Recap card

- **Transformation or action?** Transformations (filter, select, join, groupBy().agg) build a plan. Actions (count, collect, show, write) run it.
- **What happens on the second count() of an uncached DataFrame?** The whole plan runs again from the source.
- **Many checks, one scan** One df.agg(...) with count, count_if, sum, min, max expressions.
- **Reusing an expensive result on serverless** Write it to a Delta table. cache() is not supported on serverless compute.

Practice

Exercise 1. Multiple choice (level 1)

Which line starts a Spark job?

- a) `df2 = df.filter(F.col("in_scope"))`
- b) `df3 = df2.withColumn("w2", F.col("final_weight") * 2)`
- c) `n = df3.count()`
- d) `df4 = df3.groupBy("wave_id").agg(F.sum("w2"))`

Why?

- r1) count() is an action. The other lines are transformations that only extend the plan.
- r2) Only the last line of a cell runs.

Exercise 2. Predict the output (level 2)

tag is a Python UDF that adds 1 to an accumulator each time it runs. src has 1,000 rows and is not cached. What does the last line print?

```
1 df = src.select(tag("id").alias("id")).filter(F.col("id") % 2 == 0)
2 df.count()
3 df.count()
4 print(calls.value)
```

Exercise 3. Fill the gap (level 2)

Count rows and null weights in a single pass.

```
1 stats = df.____(1)(
2     F.count("*").alias("rows"),
3     F.____(2)(F.col("final_weight").isNull()).alias("null_weights"),
4 ).collect()[0]
```

Exercise 4. Refactor (level 4)

Refactor `table_health` to answer every question in **one** action. Keep the returned keys: `rows`, `null_ids`, `null_weights`, `min_weight`, `max_weight`.

Starter code:

```
1 def table_health(df):
2     return {
3         "rows": df.count(),
4         "null_ids": df.filter(F.col("respondent_id").isNull()).count(),
5         "null_weights": df.filter(F.col("final_weight").isNull()).count(),
6         "min_weight": df.agg({"final_weight": "min"}).collect()[0][0],
7         "max_weight": df.agg({"final_weight": "max"}).collect()[0][0],
8     }
```

Exercise 5. Architecture design (level 5)

You run the silver QA notebook on Databricks Free Edition (serverless). An expensive joined DataFrame feeds six different checks. Choose the approach for each situation.

Options: Combine them into one agg, Write the joined data to a Delta table and read it back, `df.cache()`, Leave each check as its own action.

1. Four checks are simple counts and sums over the joined data
2. Two checks need different groupings and joins over the same expensive result
3. A colleague suggests `cache()` to speed it all up on serverless

Sources

- Each action re-executes the plan: a counting UDF ran 1,000 times after one count and 2,000 after two, on 1,000 rows. (verified: Run on PySpark 4.0.4 in this repository.)
- `F.count_if` and `F.count_distinct` exist in PySpark 4.0. (verified: Run on PySpark 4.0.4 in this repository.)
- Serverless compute does not support DataFrame or SQL cache APIs. (verified: <https://docs.databricks.com/aws/en/compute/serverless/limitations> (checked 2026-09-25))

A5-L2 Built-in functions over Python UDFs

Concept

A **Python UDF** runs your Python function once per row. Spark must send each value from the JVM to a Python worker process, run the function, and send the result back. The physical plan shows this as a `BatchEvalPython` step.

A **built-in function** from `pyspark.sql.functions` runs inside the JVM, in Spark's generated code, with no round trip.

Measured on 2 million strings (PySpark 4.0.4, two local cores):

Approach	Time
<code>F.upper(col)</code>	0.30 s
pandas UDF (<code>ArrowEvalPython</code>)	1.16 s
Python UDF (<code>BatchEvalPython</code>)	1.82 s

The Python UDF was six times slower, and the gap grows with cluster size and data volume.

Speed is one reason. There are three more:

- **Nulls.** Built-ins return NULL for NULL input. A UDF receives `None` and must handle it. `lambda s: s.strip()` fails the whole job with a `PythonException` on the first NULL.
- **Optimisation.** The optimiser cannot see inside a UDF, so it cannot push a filter on the UDF's output into the file scan.
- **Readability.** `F.when(...).when(...)` states a rule in one expression a reviewer can check.

Most UDFs in pipelines replace functions that already exist: `F.upper`, `F.trim`, `F.regexp_replace`, `F.regexp_extract`, `F.split`, `F.to_date`, `F.coalesce`, `F.when`, `F.concat_ws`, `F.lpad`. Search the functions reference before you write one.

When no built-in fits, prefer a **pandas UDF**: it moves data in Arrow batches and runs vectorised pandas code. Use a row-by-row Python UDF only as the last resort, and handle `None` first.

Why it matters

A silver job cleans postcodes with a Python UDF: `lambda s: s.replace(" ", "").upper()`. It runs for months on deliveries where every postcode is present. Then a supplier sends a wave with 40 blank postcodes, which arrive as NULL. The first NULL raises `AttributeError: 'NoneType' object has no`

attribute 'replace' inside a Python worker, and the whole silver job fails at 02:00.

`F.upper(F.regexp_replace("postcode", r"\s+", ""))` returns NULL for NULL, runs faster, and would never have failed.

Worked example: weak code

```
1 from pyspark.sql.functions import udf
2
3
4 @udf("string")
5 def clean_postcode(s):
6     return s.replace(" ", "").upper()
7
8
9 @udf("string")
10 def age_band(age):
11     if age < 35:
12         return "16-34"
13     if age < 55:
14         return "35-54"
15     return "55+"
16
17
18 silver = bronze.withColumn("postcode",
19                             clean_postcode("postcode")).withColumn("age_group", age_band("age"))
```

1. **Line 6.** Crashes on NULL: None has no replace method.
2. **Line 11.** In Python, comparing None with 35 raises TypeError, so a NULL age crashes the job too.
3. **Line 18.** Two round trips to Python per row, invisible to the optimiser.

Worked example: strong code

```
1 from pyspark.sql import Column
2 from pyspark.sql import functions as F
3
4
5 def clean_postcode(col: Column) -> Column:
6     """Remove whitespace and upper-case. NULL stays NULL."""
7     return F.upper(F.regexp_replace(col, r"\s+", ""))
8
9
10 def age_band(age: Column) -> Column:
11     """16-34, 35-54 or 55+. NULL age gives NULL."""
12     return (
13         F.when(age < 35, "16-34")
14         .when(age < 55, "35-54")
15         .when(age.isNull(), "55+")
16     )
```

```

17
18
19 silver = bronze.withColumns({
20     "postcode": clean_postcode(F.col("postcode")),
21     "age_group": age_band(F.col("age")),
22 })

```

1. **Line 5.** A function that builds a Column expression. Reusable and testable, with no UDF.
2. **Line 7.** Built-ins pass NULL through as NULL.
3. **Line 15.** Explicit: a NULL age stays NULL instead of falling into 55+.

The strong version runs in the JVM, survives NULLs and reads as a rule. Functions that return Column expressions give you the reuse that UDFs promised.

Recap card

- **Why are Python UDFs slow?** Every value crosses from the JVM to a Python worker and back. The plan shows BatchEvalPython.
- **Measured cost on 2 million strings** Built-in 0.30 s, pandas UDF 1.16 s, Python UDF 1.82 s (PySpark 4.0.4, local).
- **What does a Python UDF receive for NULL?** None. `s.strip()` then raises and fails the job. Built-ins return NULL for NULL.
- **Reusable logic without a UDF** A Python function that takes and returns Column expressions, built from `F.when`, `F.regexp_replace` and friends.

Practice

Exercise 1. Multiple choice (level 1)

You need to trim and upper-case a string column on 50 million rows. Which approach is fastest?

- a) A Python UDF: `udf(lambda s: s.strip().upper() if s else None)`
- b) A pandas UDF using `s.str.strip().str.upper()`
- c) `F.upper(F.trim(F.col("s")))`
- d) `collect()` the column and process it in Python

Why?

- r1) Built-ins run inside the JVM with no data transfer to Python.
- r2) Built-ins run on the driver, which is faster than executors.

Exercise 2. Predict the output (level 2)

What does this print?

```

1 df = spark.createDataFrame([(" ab 1 ",), (None,)], "s STRING")
2 out = df.select(F.upper(F.regexp_replace(F.col("s"), r"\s+",
    "")).alias("u"))
3 print([r.u for r in out.collect()])

```

Exercise 3. Fill the gap (level 3)

Replace a UDF that pads respondent IDs to 8 characters with leading zeros and extracts the 2-letter region prefix from a code such as `NW-1234`.

```

1 df = df.withColumns({
2     "respondent_id": F.____(1)(F.col("respondent_id"), 8, "0"),
3     "region": F.____(2)(F.col("area_code"), r"^[A-Z]{2}" -, 1),
4 })

```

Exercise 4. Refactor (level 4)

Replace the UDF with a function `age_band(age: Column) -> Column` built from `F.when`. Keep the bands: under 35 is `16-34`, under 55 is `35-54`, anything else `55+`. A NULL age must give NULL.

Starter code:

```

1 @udf("string")
2 def age_band(age):
3     if age < 35:
4         return "16-34"
5     if age < 55:
6         return "35-54"
7     return "55+"

```

Exercise 5. Spot the bug (level 5)

This job ran for months, then failed on a wave where some postcodes are blank. Tap the line that fails, then rewrite `clean_postcodes` without a UDF.

```

1 from pyspark.sql.functions import udf
2
3
4 @udf("string")
5 def _clean(s):
6     return s.replace(" ", "").upper()
7
8
9 def clean_postcodes(df: DataFrame) -> DataFrame:
10     """Remove spaces from postcodes and upper-case them."""
11     return df.withColumn("postcode", _clean("postcode"))

```

Write down the line number of the bug. Then write the fix.

Sources

- Timings: built-in 0.30 s, pandas UDF 1.16 s, Python UDF 1.82 s on 2 million strings; plans show BatchEvalPython and ArrowEvalPython. (verified: Measured on PySpark 4.0.4, local[2], in this repository.)
- Built-ins return NULL for NULL input; a Python UDF calling a string method on None raises PythonException. (verified: Run on PySpark 4.0.4 in this repository.)

A5-L3 Joins, broadcast, partitions and reading the plan

Concept

A join between two large tables needs a **shuffle**: Spark moves rows across the network so equal keys meet on the same executor, then sorts and merges them. The plan shows this as `SortMergeJoin` with `Exchange` steps. Shuffles are the most expensive thing Spark does.

When one side is small, such as a lookup of 200 region codes, Spark can **broadcast** it: send a full copy to every executor and join locally, with no shuffle of the large side. The plan shows `BroadcastHashJoin`.

- Spark broadcasts automatically when it estimates a side is under `spark.sql.autoBroadcastJoinThreshold`, 10 MB by default in Spark 4.0.
- Adaptive query execution (AQE), on by default, can switch a join to broadcast at run time once it sees real sizes.
- You can force it with `F.broadcast(small_df)` or `small_df.hint("broadcast")`. Do this for lookups you know are small. Never broadcast something that grows with your data.

Read the plan before you tune. `df.explain()` prints the physical plan. `df.explain(mode="formatted")` adds details per step, including `PushedFilters`, which confirm that a filter reached the file scan.

Partitions decide parallelism.

- `repartition(n)` shuffles to exactly `n` partitions. `coalesce(n)` merges partitions without a shuffle, and only reduces the count.
- `repartition(1)` or `coalesce(1)` before a large write puts all the work on one task. Avoid it outside tiny outputs.
- For table layout on Unity Catalog managed tables, prefer liquid clustering over `partitionBy`. If you do partition a table, choose a low-cardinality column such as `wave_id`, never an ID.

Caching: `df.cache()` helps on classic clusters when several actions reuse a result. Serverless compute does not support it. There, write the intermediate result to a Delta table.

Why it matters

A gold job joins 40 million respondent rows to a 200-row region lookup. Auto-broadcast does not kick in, because the lookup is read through a view whose size Spark overestimates. Every run shuffles all 40 million rows across the cluster, and the join takes 22 minutes.

`explain()` shows `SortMergeJoin` with two `Exchange` steps. Wrapping the lookup in `F.broadcast()` turns it into `BroadcastHashJoin`, and the join drops to under two minutes. Five minutes reading the plan saved twenty minutes on every run.

Worked example: weak code

```
1 def enrich(fact, regions):
2     regions_local = regions.collect()
3     lookup = {r.region_code: r.region for r in regions_local}
4     label = F.udf(lambda code: lookup.get(code), "string")
5     out = fact.withColumn("region", label("region_code"))
6     return out.repartition(1)
```

1. **Line 2.** Pulls the lookup to the driver to build a Python dict. A join does this job.
2. **Line 4.** A Python UDF per row, where a broadcast join would stay in the JVM.
3. **Line 6.** All 40 million rows squeezed into one partition, so one task does all the writing.

Worked example: strong code

```
1 def enrich_with_regions(fact: DataFrame, regions: DataFrame) -> DataFrame:
2     """Add region names to fact rows. regions is a small lookup, so
3     broadcast it.
4
5     A left join keeps every fact row, even with an unknown region code.
6     """
7     2     return fact.join(F.broadcast(regions), on="region_code", how="left")
8
9     9 def show_plan(df: DataFrame) -> None:
10     """Print the formatted physical plan: check for BroadcastHashJoin."""
11     3     df.explain(mode="formatted")
```

1. **Line 6.** A left join keeps fact rows with unknown codes. An inner join would silently drop them.
2. **Line 6.** Broadcast the small side. The plan shows `BroadcastHashJoin` and no shuffle of the fact table.

3. **Line 11.** Confirm the strategy and the pushed filters before and after tuning. One broadcast join replaces a collect, a dict and a UDF, keeps the work distributed, and shows up clearly in the plan.

Recap card

- **Join strategy for a large table and a small lookup** Broadcast the lookup: `F.broadcast(small)`. The plan shows `BroadcastHashJoin`.
- **Default auto-broadcast threshold in Spark 4.0** 10 MB (`spark.sql.autoBroadcastJoinThreshold`). AQE can also switch to broadcast at run time.
- **repartition versus coalesce** `repartition(n)` shuffles to `n` partitions. `coalesce(n)` merges without a shuffle and only reduces.
- **How do you confirm a filter reached the file scan?**
`df.explain(mode="formatted")` and look for `PushedFilters` on the scan.

Practice

Exercise 1. Multiple choice (level 1)

You join a 40-million-row respondent table with a 200-row region lookup. Which strategy should the plan show?

- a) `SortMergeJoin`, shuffling both sides
- b) `BroadcastHashJoin`, with the lookup sent to every executor
- c) Collect the lookup and use a Python UDF
- d) A cross join, then a filter

Why?

- r1) Broadcasting the small side avoids shuffling the large side.
- r2) Broadcast joins always use less memory.

Exercise 2. Predict the output (level 2)

Which join node name appears in the physical plan printed by this code on Spark 4.0?

```
1 joined = fact.join(F.broadcast(regions), "region_code", "left")
2 joined.explain()
```

Exercise 3. Fill the gap (level 3)

Broadcast the lookup, keep every fact row, then print the detailed plan.

```
1 joined = fact.join(F.____(1)(regions), on="region_code", how="____(2)")
2 joined.explain(mode="____(3)")
```

Exercise 4. Spot the bug (level 4)

Writing gold takes 45 minutes on a 16-core cluster, and the Spark UI shows one busy task. Tap the line, then fix the function.

```
1 def write_gold(metrics: DataFrame, table: str) -> None:
2     """Write gold metrics for the wave."""
3     (
4         metrics.repartition(1)
5         .write.format("delta")
6         .mode("append")
7         .saveAsTable(table)
8     )
```

Write down the line number of the bug. Then write the fix.

Exercise 5. Write the function (level 5)

Write `enrich_with_regions(fact, regions)`. It left-joins `fact` to the small `regions` lookup on `region_code`, broadcasting the lookup. Then write `unmatched_codes(enriched)` that returns the number of rows where `region` is NULL, using a single aggregation.

Starter code:

```
1 from pyspark.sql import DataFrame
2 from pyspark.sql import functions as F
3
4
5 def enrich_with_regions(fact: DataFrame, regions: DataFrame) -> DataFrame:
6     ...
7
8
9 def unmatched_codes(enriched: DataFrame) -> int:
10     ...
```

Practical: Read the plan: sort-merge join versus broadcast join

Platform: Databricks Free Edition. Run this in Databricks Free Edition. It builds a fact table and a small lookup, then prints the physical plan with and without a broadcast hint, so you can find `SortMergeJoin`, `BroadcastHashJoin` and `Exchange` yourself. On serverless, Spark settings you can change are limited, so the notebook compares plans rather than timings.

Download the notebook `A5-L3.ipynb` from the app.

Sources

- `F.broadcast` and `hint('broadcast')` give `BroadcastHashJoin`; with broadcasting off, `SortMergeJoin`; `autoBroadcastJoinThreshold` defaults to

10 MB; AQE is on by default; explain supports mode='formatted' with PushedFilters. (verified: Run on PySpark 4.0.4 in this repository.)

- Databricks recommends liquid clustering over partitioning for UC managed tables; serverless does not support cache APIs. (verified: <https://docs.databricks.com/aws/en/tables/clustering> and <https://docs.databricks.com/aws/en/compute/serverless/limitations> (checked 2026-09-25))

A6 Data quality

Roadmap: Months 16 to 18: MLOps and deployment; Current work.
Expectations, health checks and reconciliation between layers.

A6-L1 Data contracts with severity levels and structured results

Concept

A **data contract** states what a delivery must look like before it may enter silver: row count range, mandatory columns, weight total, unique keys, allowed categories, expected distributions. Your reference project keeps each contract as versioned JSON, such as `nat_contract_v2.json`.

Every check produces a **structured result**, never a sentence:

```
1 CheckName = str
2 Status = Literal["PASS", "WARN", "FAIL", "SKIP"]
3 Severity = Literal["CRITICAL", "HIGH", "MEDIUM", "LOW"]
4
5 @dataclass(frozen=True)
6 class CheckResult:
7     check_name: CheckName
8     status: Status
9     severity: Severity
10    expected: str
11    actual: str
```

Records can be counted, stored in a Delta table, charted and compared across waves. A narrative such as "weights look a bit high" cannot.

Severity decides what happens next.

- A failed **CRITICAL** check stops the pipeline: raise an exception, so the job fails and silver stays untouched. Row count out of range, null keys, duplicate keys and a weight total outside limits are CRITICAL.
- A **WARN** logs the result and the run continues. Unexpected category values and distribution shifts are warnings: worth a look, not worth a missed report.

Evaluate every check first, store all the results, then decide. If you raise on the first failure, you learn about one problem per run.

Keep the contract in data and the logic in code. A new contract version changes JSON, not Python.

On Databricks, Lakeflow Spark Declarative Pipelines offer the same idea declaratively: `@dp.expect` records violations, `@dp.expect_or_drop` removes violating rows, and `@dp.expect_or_fail` stops the update. Use them when your pipeline runs in that framework. The contract pattern here works in any job.

Why it matters

A contract validator raises on the first failed check. A redelivery has a row count 2% below the minimum, so the run stops at check 1. The supplier fixes the count and resends. The next run stops at check 4, duplicate keys. The third run stops at check 6, a weight total outside the range. Three round trips with the supplier cost four days.

Evaluating every check, storing all results, and then raising once for all CRITICAL failures gives the supplier one complete list on the first day.

Worked example: weak code

```
1 def validate(df, contract):
2     if df.count() < contract["row_count_min"]:
3         raise Exception("Too few rows")
4     if df.filter("respondent_id is null").count() > 0:
5         raise Exception("Null ids found")
6     if df.select("gender").distinct().count() > 4:
7         print("warning: odd genders")
8     return "ok"
```

1. **Line 3.** Stops at the first failure, with no numbers. Later checks never run.
2. **Line 5.** A bare Exception with a sentence: nothing to store, count or compare.
3. **Line 7.** A warning that goes to stdout and is lost.
4. **Line 8.** The caller learns nothing about which checks ran.

Worked example: strong code

```
1 from dataclasses import dataclass
2 from typing import Literal
3
4 ① Status = Literal["PASS", "WARN", "FAIL", "SKIP"]
5 Severity = Literal["CRITICAL", "HIGH", "MEDIUM", "LOW"]
6
7
8 @dataclass(frozen=True)
9 ② class CheckResult:
10     """One contract check outcome, ready to store in the QA results
11     table."""
12     check_name: str
13     status: Status
```

```

14     severity: Severity
15     expected: str
16     actual: str
17
18
19 ③ class ContractViolation(Exception):
20     """Raised when one or more CRITICAL checks fail."""
21
22
23 def critical_failures(results: list[CheckResult]) -> list[CheckResult]:
24     """Failed checks with CRITICAL severity."""
25     return [r for r in results if r.status == "FAIL" and r.severity ==
26             "CRITICAL"]
27
28 ④ def enforce(results: list[CheckResult]) -> None:
29     """Raise once, listing every CRITICAL failure. WARN results pass
30     through."""
31     failures = critical_failures(results)
32     if failures:
33         names = ", ".join(f"{r.check_name} (expected {r.expected}, got
34         {r.actual})" for r in failures)
35         raise ContractViolation(f"{len(failures)} critical check(s)
36         failed: {names}")

```

1. **Line 4.** Typed statuses: a typo such as "FAILED" is a mypy error.
2. **Line 9.** A record with expected and actual values, stored for every check, pass or fail.
3. **Line 19.** A specific exception type the orchestrator can catch, log and re-raise.
4. **Line 28.** Decides after every check has run, and reports all CRITICAL failures at once.

Results are data, severity decides the action, and one run reports every problem.

Recap card

- **What does a contract check produce?** A structured record: check_name, status, severity, expected, actual. Never a sentence.
- **CRITICAL versus WARN** CRITICAL failures stop the pipeline by raising. WARN results are logged and the run continues.
- **Why evaluate every check before raising?** One run reports every problem, so the supplier fixes everything in one round trip.
- **Lakeflow Declarative Pipelines expectations** @dp.expect keeps and records, @dp.expect_or_drop drops, @dp.expect_or_fail stops the update.

Practice

Exercise 1. Architecture design (level 1)

Assign a severity to each contract check.

Options: CRITICAL: stop the pipeline, WARN: log and continue.

1. Duplicate respondent_id values within a wave
2. Weight total outside the contracted range
3. 12 rows with an unexpected gender label
4. Share of 16-34s at 19%, just below the expected 20-35%

Exercise 2. Multiple choice (level 2)

A validator finds three problems in a delivery: a low row count (CRITICAL), null keys (CRITICAL) and odd labels (WARN). What should one run report?

- a) Only the first CRITICAL failure, then stop.
- b) All three results stored, then one exception naming both CRITICAL failures.
- c) Nothing: warnings and failures are only printed.
- d) All three results stored, and the run continues to silver.

Why?

- r1) Evaluate everything, record everything, then act on severity once.
- r2) Exceptions should always be raised as early as possible.

Exercise 3. Fill the gap (level 2)

Type the result fields so mypy rejects any status or severity outside the allowed values.

```
1 Status = ____ (1) ["PASS", "WARN", "FAIL", "SKIP"]
2
3
4 @dataclass(____ (2)=True)
5 class CheckResult:
6     check_name: str
7     status: Status
```

Exercise 4. Write the function (level 4)

Results arrive as dicts with keys `check_name`, `status`, `severity`, `expected`, `actual`. Write `enforce(results)`:

- Return the list of WARN results.

- Raise `ContractViolation` (define it, subclassing `Exception`) when any result has status `FAIL` and severity `CRITICAL`. The message must name every such check.
- A `FAIL` with a non-`CRITICAL` severity does not raise.

Starter code:

```
1 class ContractViolation(Exception):
2     ...
3
4
5 def enforce(results: list[dict[str, str]]) -> list[dict[str, str]]:
6     ...
```

Exercise 5. Spot the bug (level 5)

This is the reference `ContractValidationResult.pass_rate`. It fails with a type error the first time it runs. Tap the line that causes it, then fix the imports.

```
1 from dataclasses import dataclass, field
2 from pyspark.sql.functions import col, sum
3
4
5 @dataclass
6 class ContractValidationResult:
7     checks: list = field(default_factory=list)
8
9     @property
10    def pass_rate(self) -> float:
11        passed = sum(1 for c in self.checks if c.status == "PASS")
12        return passed / len(self.checks) if self.checks else 0
```

Write down the line number of the bug. Then write the fix.

Exercise 6. Multiple choice (level 5)

Your silver table is built with Lakeflow Spark Declarative Pipelines. Rows with a `NULL` `respondent_id` must never reach the table, but the update should keep going for valid rows. Which decorator fits?

- a) `@dp.expect("valid_id", "respondent_id IS NOT NULL")`
- b) `@dp.expect_or_drop("valid_id", "respondent_id IS NOT NULL")`
- c) `@dp.expect_or_fail("valid_id", "respondent_id IS NOT NULL")`
- d) No decorator: filter in the next notebook

Why?

- r1) `expect_or_drop` removes violating rows and records how many were dropped.
- r2) `expect_or_drop` fails the update when any row violates.

Sources

- Importing `sum` from `pyspark.sql.functions` shadows the built-in, so `sum` over a generator fails. (verified: Run on PySpark 4.0.4: `PySparkTypeError [NOT_COLUMN_OR_STR] Argument 'col' should be a Column or str, got generator.`)
- Lakeflow Spark Declarative Pipelines provide `@dp.expect`, `@dp.expect_or_drop` and `@dp.expect_or_fail` in `pyspark.pipelines`. (verified: <https://docs.databricks.com/aws/en/ldp/developer/ldp-python-ref-expectations> (checked 2026-09-25))

A6-L2 Statistical checks that survive empty, single-row and zero data

Concept

Statistical QA runs on every wave, including the strange ones: an empty delivery, a subgroup with one respondent, a metric that drops to exactly zero. A check that crashes on those cases fails exactly when you need it. Know what Spark 4.0 returns at the edges (all tested on PySpark 4.0.4):

Case	Result
<code>F.mean</code> , <code>F.stddev</code> , <code>percentile_approx</code> on an empty <code>DataFrame</code>	NULL
<code>F.stddev</code> on one row	NULL (it is the sample standard deviation)
<code>F.col("a") / F.col("b")</code> where <code>b</code> is 0, ANSI mode on	Error: <code>DIVIDE_BY_ZERO</code> , even for <code>DOUBLE</code>
<code>F.try_divide(a, b)</code> where <code>b</code> is 0	NULL
Python <code>round(None, 4)</code>	<code>TypeError</code>

So a weight-spread check built as `round(stats["stddev"] / stats["mean"], 4)` fails three ways: NULL mean on an empty wave, NULL stddev on one row, and `round(None)` in Python.

Four habits make checks robust:

1. **Compute the count with the statistics**, in the same `agg`, and decide early. Fewer than two rows: return a `SKIP` result that says why.
2. **Divide with** `F.try_divide` or guard with `F.when(denominator > 0, ...)`. Never a bare `/` on data you do not control.

3. **Test for None with** `is None` in Python. `if cv:` treats a real CV of 0.0 as missing, and `if current_pct` and `prior_pct:` skips a share that fell to 0.0 (reference defect 8).
4. **Return a result for every check**, including SKIP. A check that silently returns nothing looks exactly like a check that passed.

Why it matters

A new tracker's first wave has one respondent in a small region, because of a routing test. The weight-spread check computes the sample standard deviation for that region: NULL. Python then runs `round(None, 4)` and raises `TypeError`. The QA job fails, and because QA gates publication, the whole wave's reporting is delayed by a day while someone reads the stack trace.

With the count checked first, that region gets a SKIP result, "fewer than 2 respondents", and every other region is checked as normal.

Worked example: weak code

```
1 def check_weight_distribution(df, weight_col):
2     stats = df.agg(
3         F.mean(weight_col).alias("mean"),
4         F.stddev(weight_col).alias("stddev"),
5     ).collect()[0]
6     cv = stats["stddev"] / stats["mean"] if stats["mean"] else None
7     return {
8         "mean": round(stats["mean"], 4),
9         "cv": round(cv, 4) if cv else None,
10        "cv_warning": cv > 1.5 if cv else False,
11    }
```

1. **Line 6.** `stddev` is NULL for one row, so `None / mean` raises `TypeError`.
2. **Line 8.** An empty frame gives `mean` `None`, and `round(None)` raises `TypeError`.
3. **Line 10.** Truthiness: a missing CV reports no warning, as if the check had passed.

Worked example: strong code

```
1 MIN_ROWS = 2
2 MAX_WEIGHT_CV = 1.5
3
4
5 def weight_spread(df: DataFrame, weight_col: str) -> dict[str, float |
6     int | None]:
7     """Count, mean, stddev and CV of weights in one pass. CV is NULL when
8     undefined."""
9     row = df.agg(
10        F.count(weight_col).alias("n"),
11        F.mean(weight_col).alias("mean"),
```

```

10     F.stddev(weight_col).alias("stddev"),
11     F.try_divide(F.stddev(weight_col),
12     F.mean(weight_col)).alias("cv"),
13     ).collect()[0]
14     return row.asDict()
15
16 def check_weight_spread(stats: dict[str, float | int | None]) ->
17     CheckResult:
18     """PASS, WARN, or SKIP with a reason. Never raises on edge cases."""
19     if stats["n"] < MIN_ROWS or stats["cv"] is None:
20         return CheckResult("weight_cv", "SKIP", "MEDIUM", f">= {MIN_ROWS}
21         rows", f"n={stats['n']}")
22     status = "WARN" if stats["cv"] > MAX_WEIGHT_CV else "PASS"
23     return CheckResult("weight_cv", status, "MEDIUM", f"<=
24     {MAX_WEIGHT_CV}", f"{stats['cv']:.4f}")

```

1. **Line 8.** The count travels with the statistics, in the same pass.
2. **Line 11.** CV computed in Spark. `try_divide` returns NULL for a zero or NULL mean instead of raising.
3. **Line 18.** Edge cases become an explicit SKIP with the reason, checked with `is None`.
4. **Line 20.** Only reached with a real number, so the comparison is safe.

The strong version returns a result for every wave and every subgroup, and says why when a statistic is undefined.

Recap card

- **F.stddev on one row, and mean on an empty frame** Both NULL. `stddev` is the sample standard deviation.
- **Division by zero on Spark 4.0 with ANSI on** Raises `DIVIDE_BY_ZERO`, even for `DOUBLE`. Use `F.try_divide` or `F.when(b > 0, a / b)`.
- **What should a check return when a statistic is undefined?** An explicit SKIP result with the reason. Never nothing, and never a crash.
- **Why is `if cv:` wrong?** A real CV of 0.0 is falsy. Test missing values with `is None`.

Practice

Exercise 1. Predict the output (level 1)

A region has exactly one respondent. What does this print?

```

1 one = spark.createDataFrame([(1.2,)], "final_weight DOUBLE")
2 print(one.agg(F.stddev("final_weight")).collect()[0][0])

```

Exercise 2. Multiple choice (level 2)

On Spark 4.0 (Runtime 17.3 LTS), `F.col("weighted_reach") / F.col("weighted_base")` meets a row where `weighted_base` is `0.0`. What happens?

- a) The result is Infinity.
- b) The result is NULL.
- c) The query fails with `DIVIDE_BY_ZERO`.
- d) The result is `0.0`.

Why?

- r1) ANSI mode is on by default in Spark 4.0, and it raises on division by zero, `DOUBLE` included.
- r2) Delta rejects infinite values.

Exercise 3. Fill the gap (level 3)

Compute the coefficient of variation in Spark so that a zero or NULL mean gives NULL instead of an error.

```
1 cv = F.____(1)(F.stddev("final_weight"), F.____(2)("final_weight"))
```

Exercise 4. Write the function (level 4)

Write `check_weight_spread(stats, max_cv=1.5)`. `stats` is a dict with `n` (int), `cv` (float or None). Return a tuple `(status, reason)`:

- `("SKIP", "fewer than 2 rows")` when `n < 2`.
- `("SKIP", "cv undefined")` when `cv` is None.
- `("WARN", ...)` when `cv > max_cv`, otherwise `("PASS", ...)`. A CV of `0.0` is a real value and passes.

Starter code:

```
1 def check_weight_spread(stats: dict[str, float | int | None], max_cv:
    float = 1.5) -> tuple[str, str]:
2     ...
```

Exercise 5. Spot the bug (level 5)

This is the reference wave-stability comparison. A share that dropped from 12% to 0% produced no result at all. Tap the line, then fix the function.

```
1 def compare_share(metric: str, current_pct: float | None, prior_pct:
    float | None, tolerance: float = 0.10) -> dict[str, object] | None:
2     """Return a stability result, or None when a share is missing."""
3     if current_pct and prior_pct:
4         change = abs(current_pct - prior_pct)
```

```

5         return {"metric": metric, "change_pp": round(change, 4),
"status": "WARN" if change > tolerance else "PASS"}
6     return None

```

Write down the line number of the bug. Then write the fix.

Sources

- On PySpark 4.0.4: stddev of one row and mean/stddev/percentile_approx of an empty frame are NULL; division by zero raises DIVIDE_BY_ZERO for DOUBLE under ANSI; try_divide returns NULL; Python round(None) raises TypeError. (verified: Run in this repository.)

A6-L3 Reconcile counts and totals between layers

Concept

Each layer transforms data, and each transformation can lose or duplicate rows. **Reconciliation** checks that what came out of one layer matches what went into the next, for every wave you load.

Compare quantities that must agree by design:

From	To	Must match
Bronze, wave W	Silver, wave W	Distinct (respondent_id, wave_id) keys in bronze = silver rows
Silver, wave W, in scope	Gold, wave W	Sum of final_weight = sum of weighted_base over all subgroups
Silver, wave W, in scope	Gold, wave W	Row count = sum of unweighted_n

Pick the right bronze figure. Bronze holds every delivery, redeliveries included, so its raw row count is larger than silver's by design. Compare **distinct keys**.

Compare totals with a **tolerance**. Floating-point sums depend on the order Spark adds values in, which varies with partitioning, so two correct totals can differ in the last digits. In Python, `0.1 + 0.2 == 0.3` is `False`. Use `math.isclose(a, b, rel_tol=1e-9, abs_tol=1e-6)`. The `abs_tol` handles totals near zero, where a relative tolerance means nothing.

Compute each side with one aggregation, record the result as a structured check (expected, actual, status), and treat a failed reconciliation as **CRITICAL**. If gold does not add up to silver, every published figure for that wave is suspect.

Run reconciliation after each layer loads and before anything downstream reads it. Store the results with the run log, so you can show that every published wave reconciled.

Why it matters

A gold job joins silver to a region lookup with an inner join. A supplier introduces a new region code, **NI**, that the lookup does not have. 312 respondents drop out of gold without an error. Weighted reach for the wave looks plausible, because the base shrank along with the reach.

A reconciliation of silver's in-scope weight total against the sum of gold's `weighted_base` catches the 1.8% gap on the same run. The alert names the wave and the size of the difference, and the fix is a lookup update, not a client complaint.

Worked example: weak code

```
1 def reconcile(bronze, silver, wave):
2     b = bronze.filter(F.col("wave_id") == wave).count()
3     s = silver.filter(F.col("wave_id") == wave).count()
4     if b != s:
5         print(f"mismatch {b} vs {s}")
```

1. **Line 2.** Counts every delivery, redeliveries included. A redelivered wave always mismatches.
2. **Line 4.** Exact equality: fine for counts, wrong for weight totals.
3. **Line 5.** A print that nobody reads, and the pipeline carries on.

Worked example: strong code

```
1 import math
2
3
4 def silver_totals(silver: DataFrame, wave_id: str) -> tuple[int, float]:
5     """In-scope row count and weight total for one wave, in one pass."""
6     row = (
7         silver.filter((F.col("wave_id") == wave_id) & F.col("in_scope")))
8         .agg(F.count("*").alias("n"), F.coalesce(F.sum("final_weight"),
9           F.lit(0.0)).alias("w"))
10    )
11    return int(row["n"]), float(row["w"])
12
13
14 def gold_totals(gold: DataFrame, wave_id: str) -> tuple[int, float]:
15     """Sum of unweighted_n and weighted_base over every subgroup of one
16     wave."""
17     row = (
```

```

17     gold.filter(F.col("wave_id") == wave_id)
18     .agg(F.coalesce(F.sum("unweighted_n"), F.lit(0)).alias("n"),
19          F.coalesce(F.sum("weighted_base"), F.lit(0.0)).alias("w"))
20     .collect()[0]
21     )
22     return int(row["n"]), float(row["w"])
23
24 def totals_match(expected: float, actual: float) -> bool:
25     """Equal within floating-point tolerance, including totals near
26     zero."""
27     return math.isclose(expected, actual, rel_tol=1e-9, abs_tol=1e-6)

```

1. **Line 7.** Same population as gold: one wave, in scope only.
2. **Line 8.** Count and total in one pass. Coalesce so an empty wave gives 0, not NULL.
3. **Line 26.** Tolerant comparison for floating-point totals.

Silver and gold totals are computed the same way and compared with tolerance. A mismatch becomes a CRITICAL result with numbers, not a print.

Recap card

- **Bronze figure that must match silver's row count for a wave**
Distinct (respondent_id, wave_id) keys, not raw bronze rows, because bronze keeps every delivery.
- **Silver-to-gold reconciliation** In-scope silver weight total = sum of gold weighted_base; in-scope silver rows = sum of unweighted_n.
- **How to compare floating-point totals** `math.isclose(a, b, rel_tol=1e-9, abs_tol=1e-6)`. Never `==`.
- **Severity of a failed reconciliation** CRITICAL. Every figure for the wave is suspect until it passes.

Practice

Exercise 1. Multiple choice (level 1)

Wave W14 was delivered twice to bronze. What should the silver row count for W14 equal?

- a) The bronze row count for W14
- b) The number of distinct (respondent_id, wave_id) keys in bronze for W14
- c) Half the bronze row count for W14
- d) The row count of the latest delivery only

Why?

- r1) Silver keeps one row per key, the newest version, so it matches the number of distinct keys across all deliveries.
- r2) Bronze deletes old deliveries after silver loads.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 import math
2
3 print(0.1 + 0.2 == 0.3, math.isclose(0.1 + 0.2, 0.3))
```

Exercise 3. Write the function (level 3)

Write `reconciliation_result(name, expected, actual)` that returns a dict with `check_name`, `status` ("PASS" or "FAIL"), `severity` (always "CRITICAL"), `expected` and `actual` (both formatted to 2 decimal places as strings), and `difference` (actual minus expected, a float). Compare with `math.isclose(expected, actual, rel_tol=1e-9, abs_tol=1e-6)`.

Starter code:

```
1 import math
2
3
4 def reconciliation_result(name: str, expected: float, actual: float) ->
    dict[str, object]:
5     ...
```

Exercise 4. Spot the bug (level 4)

This bronze-to-silver reconciliation fails every time a wave is redelivered, even when silver is correct. Tap the line, then fix the function.

```
1 def bronze_silver_keys(bronze: DataFrame, silver: DataFrame, wave_id:
    str) -> tuple[int, int]:
2     """Keys in bronze and rows in silver for one wave."""
3     b = bronze.filter(F.col("wave_id") == wave_id).count()
4     s = silver.filter(F.col("wave_id") == wave_id).count()
5     return b, s
```

Write down the line number of the bug. Then write the fix.

Exercise 5. Put in order (level 5)

Order the steps for one wave, from load to publication.

- MERGE the wave into silver
- Store reconciliation results with the run log
- Reconcile silver in-scope totals with gold totals

- Refresh the Power BI semantic model
- Build gold metrics for the wave
- Reconcile bronze distinct keys with silver rows

Practical: Reconcile silver with gold, and catch a dropped region

Platform: Databricks Free Edition. Run this in Databricks Free Edition. It builds a small silver table and two gold versions: one correct, one built with an inner join that drops an unknown region. The reconciliation passes for the first and fails for the second.

Download the notebook `A6-L3.ipynb` from the app.

Sources

- $0.1 + 0.2 == 0.3$ is False and `math.isclose(0.1 + 0.2, 0.3)` is True in Python. (verified: Run in Python 3.11 in this repository.)
- `F.count_distinct` exists in PySpark 4.0. (verified: Run on PySpark 4.0.4 in this repository.)

A7 Observability

Roadmap: Months 16 to 18: MLOps and deployment; Current work.
Structured logging and pipeline run log tables.

A7-L1 Structured logging with the logging module, not print

Concept

`print` writes text to standard output, with no level, no source and no context. In a scheduled job, that text lands in a driver log that nobody searches until something breaks.

Python's standard `logging` module fixes all three problems.

One logger per module. `log = logging.getLogger(__name__)` names the logger after the module, such as `pipelines.silver`. Every line then says where it came from.

Levels. `DEBUG` for detail while developing, `INFO` for normal milestones, `WARNING` for something odd that did not stop the run, `ERROR` for a failure, `CRITICAL` for a failure that stops everything. Loggers pass `WARNING` and above by default, so `log.info(...)` prints nothing until you configure a level.

Context on every line. Include the identifiers you will search for later: `run_id`, `tracker_id`, `wave_id`, `layer`. A `LoggerAdapter` attaches them once, and a format string prints them:

```
1 base = logging.getLogger(__name__)
2 log = logging.LoggerAdapter(base, {"run_id": run_id, "wave_id": wave_id})
3 log.info("merged rows=%d", n)
4 # INFO pipelines.silver run_id=7f3c wave_id=W14 merged rows=1832
```

Three more rules:

- Pass values as arguments, `log.info("rows=%d", n)`, not pre-built f-strings. The logger formats them only when the line is actually emitted.
- Inside `except`, call `log.exception("...")`. It logs at `ERROR` and appends the full traceback.
- Configure handlers and formats once, in the entry point, never inside library functions.

Logs are for people reading one run. Metrics you want to compare across runs, such as row counts and statuses, also go to the run log table (lesson A7-L2).

Why it matters

A silver job prints "done" at the end of each step. One night it finishes "successfully" with zero rows merged, because an upstream filter used the wrong wave ID. The driver log for that run holds 400 lines of "done" and Spark noise, with no wave ID and no row counts.

With `log.info("merged rows=%d", n)` and the wave ID on every line, the zero is visible in one search, and a WARNING for an empty merge would have flagged it at 02:00.

Worked example: weak code

```
1 def merge_wave(df, wave_id):
2     print("starting")
3     try:
4         n = do_merge(df)
5         print(f"done {n}")
6     except Exception as e:
7         print("error:", e)
```

1. **Line 2.** No level, no module, no wave ID: unsearchable.
2. **Line 5.** A count with no label. Zero rows looks the same as a normal run.
3. **Line 7.** Swallows the exception and loses the traceback. The job reports success.

Worked example: strong code

```
1 import logging
2
3 logger = logging.getLogger(__name__)
4
5
6 def merge_wave(df: DataFrame, run_id: str, wave_id: str) -> int:
7     """MERGE one wave, logging with run and wave context. Re-raises on
8     failure."""
9     log = logging.LoggerAdapter(logger, {"run_id": run_id, "wave_id":
10     wave_id})
11     log.info("merge started")
12     try:
13         rows = do_merge(df)
14     except Exception:
15         log.exception("merge failed")
16         raise
17     if rows == 0:
18         log.warning("merge wrote 0 rows")
```

```

17     else:
18         log.info("merge finished rows=%d", rows)
19     return rows

```

1. **Line 3.** One logger per module, named after it.
2. **Line 8.** Context attached once, printed on every line by the entry point's format.
3. **Line 13.** ERROR level plus the full traceback.
4. **Line 14.** Re-raise, so the job fails and alerts fire.
5. **Line 16.** A suspicious result gets its own level, not a buried number.

Every line says what, where, which run and which wave. Failures keep their traceback and still fail the job.

Recap card

- **Create a module logger** `logger = logging.getLogger(name)`
- **Default level threshold with no configuration** WARNING. `log.info(...)` prints nothing until the entry point sets a level.
- **Attach run_id and wave_id to every line**
`logging.LoggerAdapter(logger, {"run_id": ..., "wave_id": ...})`
- **Inside except** `log.exception("...")` for ERROR plus traceback, then a bare `raise`.

Practice

Exercise 1. Predict the output (level 1)

No logging configuration has run. What appears in the output?

```

1 import logging
2
3 log = logging.getLogger("silver")
4 log.info("loaded")
5 log.warning("3 unexpected labels")

```

Exercise 2. Fill the gap (level 2)

Create a module logger, then log the failure with its traceback and re-raise.

```

1 logger = logging.getLogger(____(1))
2
3 try:
4     rows = do_merge(df)
5 except Exception:
6     logger.____(2)("merge failed")
7     ____ (3)

```

Exercise 3. Spot the bug (level 3)

This job failed at the MERGE last night, yet the scheduler marked it as succeeded and nobody was alerted. Tap the line, then fix the function.

```
1 def merge_wave(df, run_id: str, wave_id: str) -> int:
2     """MERGE one wave. Logs failures."""
3     log = logging.LoggerAdapter(logger, {"run_id": run_id, "wave_id":
4         wave_id})
5     try:
6         return do_merge(df)
7     except Exception:
8         log.exception("merge failed")
9     return 0
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 4)

Write `wave_logger(run_id, wave_id)` that returns a `logging.LoggerAdapter` around the module logger `logging.getLogger("pipelines.silver")`, carrying `run_id` and `wave_id`. Then write `log_merge_result(log, rows)`: log at WARNING "merge wrote 0 rows" when rows is 0, otherwise at INFO "merge finished rows=%d" with the count passed as an argument.

Starter code:

```
1 import logging
2
3
4 def wave_logger(run_id: str, wave_id: str) -> logging.LoggerAdapter:
5     ...
6
7
8 def log_merge_result(log: logging.LoggerAdapter, rows: int) -> None:
9     ...
```

Exercise 5. Architecture design (level 5)

Choose the level for each event in a silver job.

Options: DEBUG, INFO, WARNING, ERROR.

1. Run started for wave W14
2. MERGE wrote 0 rows for a wave that should have data
3. MERGE raised
DELTA_MULTIPLE_SOURCE_ROW_MATCHING_TARGET_ROW_IN_MERGE
4. The full query plan of the MERGE source, for troubleshooting

Sources

- Without configuration, loggers pass WARNING and above, and the last-resort handler prints only the message; LoggerAdapter adds context; logger.exception includes the traceback. (verified: Run in Python 3.11 in this repository.)

A7-L2 A pipeline run log table with an explicit schema

Concept

Logs describe one run. A **run log table** lets you compare every run: which waves loaded, how many rows, how long, and which failed. It is a Delta table, append-only, with one row per pipeline run per layer.

A useful schema:

Column	Type	Notes
<code>run_id</code>	STRING, not null	<code>str(uuid.uuid4())</code> , shared by every log line of the run
<code>pipeline_name</code>	STRING, not null	<code>nat_silver_pipeline</code>
<code>tracker_id</code> , <code>wave_id</code> , <code>layer</code>	STRING	what the run touched
<code>rows_processed</code>	BIGINT, not null	from Delta commit metrics
<code>status</code>	STRING, not null	SUCCESS or FAILED
<code>error_message</code>	STRING, nullable	NULL on success
<code>started_at</code> , <code>finished_at</code>	TIMESTAMP, not null	UTC

Build the record against an **explicit** `StructType`. Your reference `log_pipeline_run` builds a `DataFrame` from a dict and lets Spark infer types. That fails on every successful run: `error_message` is `None`, and Spark raises `CANNOT_DETERMINE_TYPE` because it cannot infer a type from a column of only nulls. The explicit schema also rejects a wrong type, such as `"12"` for `rows_processed`, before anything is written.

Take timestamps with `datetime.now(timezone.utc)`. `datetime.utcnow()` is deprecated from Python 3.12 and returns a naive datetime that carries no timezone.

Import everything you use. The reference function calls `uuid` and `datetime` without importing them, a `NameError` on its first run. `ruff check` (rule F821) catches it in CI.

Write the log record in a `finally` block, or in both the success and failure paths, so failed runs are logged too. A run log that only records successes cannot tell you how often the pipeline fails.

Why it matters

The reference `log_pipeline_run` has never succeeded. It fails first with `NameError: name 'uuid' is not defined`. Once someone adds the imports, it fails on every successful run with `CANNOT_DETERMINE_TYPE`, because `error_message=None` is the only value in its column. The orchestrator catches that error "so logging never breaks the pipeline". Months later the team wants to show a client every load of 2025. The run log table is empty.

An explicit schema, imported names and a test that writes one record on a local `SparkSession` prevent all three problems.

Worked example: weak code

```
❶ def log_pipeline_run(pipeline_name, tracker_id, wave_id, layer,
❷     rows_processed, status, error_message: str = None):
❸     run_record = spark.createDataFrame([
❹         {"run_id": str(uuid.uuid4()),
❺          "pipeline_name": pipeline_name,
❻         "rows_processed": rows_processed,
❼         "status": status,
❽         "error_message": error_message,
❹         "run_timestamp": datetime.utcnow().isoformat(),
❺     }])
❺     run_record.write.format("delta").mode("append").saveAsTable("research_governance")
```

1. **Line 1.** Claims a string, defaults to `None` (defect 7).
2. **Line 2.** Types inferred from one dict. An all-`None` column raises `CANNOT_DETERMINE_TYPE`.
3. **Line 3.** `uuid` is never imported: `NameError` (defect 1).
4. **Line 8.** `datetime` not imported, `utcnow` deprecated (defect 6), and the timestamp stored as a string.

Worked example: strong code

```
1 import uuid
2 from datetime import datetime, timezone
3
4 from pyspark.sql import DataFrame, SparkSession
```

```

5 from pyspark.sql import types as T
6
7 RUN_LOG_SCHEMA = T.StructType([
8     T.StructField("run_id", T.StringType(), False),
9     T.StructField("pipeline_name", T.StringType(), False),
10    T.StructField("wave_id", T.StringType(), True),
11    T.StructField("rows_processed", T.LongType(), False),
12    T.StructField("status", T.StringType(), False),
13    T.StructField("error_message", T.StringType(), True),
14    T.StructField("started_at", T.TimestampType(), False),
15    T.StructField("finished_at", T.TimestampType(), False),
16 ])
17
18
19 def new_run_id() -> str:
20     """A unique id shared by the log lines and the run log row of one
21     run."""
22     return str(uuid.uuid4())
23
24 def run_record(
25     spark: SparkSession,
26     run_id: str,
27     pipeline_name: str,
28     wave_id: str | None,
29     rows_processed: int,
30     status: str,
31     started_at: datetime,
32     error_message: str | None = None,
33 ) -> DataFrame:
34     """One run log row, typed by RUN_LOG_SCHEMA."""
35     finished_at = datetime.now(timezone.utc)
36     row = (run_id, pipeline_name, wave_id, rows_processed, status,
37           error_message, started_at, finished_at)
38     return spark.createDataFrame([row], RUN_LOG_SCHEMA)

```

- Line 7.** The schema is declared once. Nothing is inferred.
- Line 13.** Nullable string: None is fine, because the type is declared.
- Line 35.** Timezone-aware UTC, stored as a real `TIMESTAMP`.
- Line 37.** A wrong type, such as "12" for `rows_processed`, fails here with a clear message.

The strong record builds on every run, success or failure, and its types never depend on the values in one row.

Recap card

- Why does the reference `log_pipeline_run` fail on success?**
`error_message=None` is inferred with no type: `CANNOT_DETERMINE_TYPE`.
 Declare a `StructType`.

- **UTC timestamp in modern Python** `datetime.now(timezone.utc)`. `utcnow()` is deprecated from 3.12 and naive.
- **Which ruff rule catches the missing `uuid` import?** F821, undefined name.
- **Where do you write the run record so failures are logged too?** In a finally block, after setting status in try and except.

Practice

Exercise 1. Multiple choice (level 1)

`spark.createDataFrame([{"run_id": "r1", "error_message": None}])` runs on PySpark 4.0. What happens?

- a) `error_message` becomes a `STRING` column holding `NULL`.
- b) It fails with `CANNOT_DETERMINE_TYPE`.
- c) `error_message` is dropped from the `DataFrame`.
- d) `error_message` becomes a `NULL`-typed column that Delta accepts.

Why?

- r1) Spark infers each column's type from the data, and a column holding only `None` gives it nothing to infer from.
- r2) Spark never allows `None` in a `DataFrame`.

Exercise 2. Fill the gap (level 2)

Complete the imports and the timestamp so the record code runs on Python 3.12 without deprecation warnings.

```
1 import ____ (1)
2 from datetime import datetime, ____ (2)
3
4 run_id = str(uuid.uuid4())
5 started_at = datetime.____ (3)(timezone.utc)
```

Exercise 3. Spot the bug (level 3)

This builds a run log row. It fails on every successful run. Tap the line, then fix the function so it declares its schema.

```
1 def run_record(spark: SparkSession, run_id: str, rows: int, status: str,
2               error_message: str | None) -> DataFrame:
3     """One run log row."""
4     return spark.createDataFrame([
5         {"run_id": run_id,
6          "rows_processed": rows,
7          "status": status,
8          "error_message": error_message,
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 4)

Write `run_record(spark, run_id, pipeline_name, rows_processed, status, started_at, error_message=None)` returning a one-row `DataFrame` typed by a module constant `RUN_LOG_SCHEMA` with columns `run_id`, `pipeline_name`, `rows_processed` (`LongType`), `status`, `error_message` (nullable), `started_at` and `finished_at` (`TimestampType`). Set `finished_at` with `datetime.now(timezone.utc)`.

Starter code:

```
1 from datetime import datetime, timezone
2
3 from pyspark.sql import DataFrame, SparkSession
4 from pyspark.sql import types as T
5
6 RUN_LOG_SCHEMA = ...
7
8
9 def run_record(spark: SparkSession, run_id: str, pipeline_name: str,
    rows_processed: int, status: str, started_at: datetime, error_message:
    str | None = None) -> DataFrame:
10     ...
```

Exercise 5. Put in order (level 5)

Order the orchestrator's steps so both successful and failed runs are logged, and a failure still fails the job.

- except: set status FAILED and error_message, log.exception, then raise
- Create run_id and started_at
- finally: append the run record to the run log table
- Set status SUCCESS
- try: run the layer and read rows_processed from the commit metrics

Practical: Write run log records for a successful and a failed run

Platform: Databricks Free Edition. Run this in Databricks Free Edition. It creates a run log table with an explicit schema, runs one step that succeeds and one that fails, and records both. It also shows the inference error the reference code hits.

Download the notebook `A7-L2.ipynb` from the app.

Sources

- `createDataFrame` from a dict with an all-None column raises `CANNOT_DETERMINE_TYPE`; an explicit `StructType` accepts None and rejects '12' for `LongType` with `FIELD_DATA_TYPE_UNACCEPTABLE_WITH_NAME`. (verified: Run on PySpark 4.0.4 in this repository.)
- `datetime.utcnow()` is deprecated from Python 3.12. (verified: Python 3.12 release notes (What's New in Python 3.12, deprecated APIs); CLAUDE.md standard.)

A7-L3 Alert on failures, and on the silence that hides them

Concept

An alert is useful when it fires for real problems, reaches the person who can act, and says enough to start. Pipelines need two kinds.

1. Loud failures: the job fails. Databricks job notifications send an email or a webhook message when a run fails. In a job definition or bundle:

```
1 email_notifications:
2   on_failure:
3     - data-oncall@example.org
```

These fire only if the job actually fails. Every `except` block that logs and carries on silences them. Log, record the run as `FAILED`, then re-raise (lesson A7-L2).

Job-level notifications are not sent when a failed task is retried. If you want a message for every failed attempt, add task-level notifications.

2. Silent failures: the job succeeds, but the data is wrong or missing. A job that never starts, a `MERGE` that writes 0 rows, a wave with a `WARN` on every contract check. No job fails, so no notification fires. Detect these with queries on your run log and QA results tables, run on a schedule, for example as Databricks SQL alerts:

- No `SUCCESS` run for a pipeline in the last 26 hours (a daily job plus slack).
- A run with `rows_processed = 0`.
- More `WARN` results for a wave than the last five waves' average.

Write each alert query against the exact values your tables hold. A query for `status = 'FAIL'` never fires on a table that stores `FAILED`.

Send alerts to a group address or channel, never to one person's inbox. Include the pipeline, the wave, the run ID and a link, so the reader can start without searching.

Why it matters

A nightly bronze job's schedule is paused during a workspace migration and never resumed. The job does not fail, because it never runs, so no failure email arrives. Silver and gold keep succeeding on the old data. Eleven days later a client asks why the dashboard has not moved.

A scheduled query, "no SUCCESS for nat_bronze_pipeline in 26 hours", would have alerted the on-call group on the second morning.

Worked example: weak code

```
1 -- Alert: failed runs today
2 SELECT COUNT(*) AS failures
3 FROM research_governance.pipeline_run_log
4 WHERE status = 'FAIL'
5 AND run_timestamp >= current_date();
```

- Line 2.** Counts failures only. A job that stopped running produces no rows and no alert.
- Line 4.** The table stores FAILED. This alert has never fired.

Worked example: strong code

```
1 -- Alert 1: pipelines with no successful run in the last 26 hours.
2 SELECT pipeline_name,
3     max(finished_at) AS last_success
4 FROM research_prod.research_governance.pipeline_run_log
5 WHERE status = 'SUCCESS'
6 GROUP BY pipeline_name
7 HAVING max(finished_at) < current_timestamp() - INTERVAL 26 HOURS;
8
9 -- Alert 2: successful runs that wrote nothing in the last day.
10 SELECT pipeline_name, wave_id, run_id
11 FROM research_prod.research_governance.pipeline_run_log
12 WHERE status = 'SUCCESS'
13 AND rows_processed = 0
14 AND finished_at >= current_timestamp() - INTERVAL 1 DAY;
```

- Line 7.** Detects silence: a pipeline whose last success is too old, whether it failed or never ran.
- Line 10.** The alert names what to look at: pipeline, wave and run.
- Line 13.** A success that did nothing is a silent failure worth an alert.

Job notifications catch runs that fail. Scheduled queries catch runs that do not happen, or that succeed at doing nothing.

Recap card

- **Why must an except block re-raise?** Otherwise the run succeeds and on_failure notifications never fire.
- **Job-level versus task-level failure notifications** Job-level notifications are not sent for retried task failures. Use task-level ones to hear about every failed attempt.
- **How do you detect a pipeline that stopped running?** A scheduled query: no SUCCESS in the run log within the expected window.
- **Where should alerts go?** A group address or channel, with pipeline, wave, run ID and a link.

Practice

Exercise 1. Multiple choice (level 1)

A job has `email_notifications.on_failure` configured. Its silver task catches every exception, logs it, and returns normally. The MERGE fails. What arrives in the on-call inbox?

- a) A failure email.
- b) Nothing.
- c) A warning email.
- d) A success email.

Why?

- r1) The task returns normally, so the run succeeds and on_failure never fires.
- r2) Notifications only work for SQL tasks.

Exercise 2. Fill the gap (level 2)

Complete the job settings so a failed run emails the on-call group.

```
1 email_notifications:
2   ____ (1):
3     - ____ (2)
```

Exercise 3. Spot the bug (level 3)

The run log stores status values SUCCESS and FAILED. This alert has never fired, even on nights with failures. Tap the line, then fix the query.

```
1 SELECT pipeline_name, run_id, error_message
2 FROM research_prod.research_governance.pipeline_run_log
3 WHERE status = 'FAIL'
4 AND finished_at >= current_timestamp() - INTERVAL 1 DAY;
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 4)

Write `stale_pipelines(runs, now, max_age_hours=26)`. `runs` is a list of dicts with `pipeline_name`, `status` and `finished_at` (timezone-aware datetime). Return the sorted names of pipelines whose most recent SUCCESS is older than `max_age_hours` before `now`. A pipeline with runs but no SUCCESS at all is stale too.

Starter code:

```
1 from datetime import datetime, timedelta
2
3
4 def stale_pipelines(runs: list[dict[str, object]], now: datetime,
5                     max_age_hours: int = 26) -> list[str]:
6     ...
```

Exercise 5. Architecture design (level 5)

Choose how each problem should be detected.

Options: Job failure notification, Scheduled query on the run log or QA tables, No alert needed.

1. The silver MERGE raises an exception
2. The bronze job's schedule was paused and nothing has run for two days
3. Gold succeeds but writes 0 rows for a new wave
4. A run took 9 minutes instead of the usual 8

Sources

- Jobs support `email_notifications` and `webhook_notifications` with `on_start`, `on_success` and `on_failure`; job-level notifications are not sent when failed tasks are retried; at most 3 destinations per `on_failure`. (verified: <https://docs.databricks.com/aws/en/jobs/notifications> (checked 2026-09-25))
- Databricks SQL alerts run a query on a schedule and notify when a condition is met. (verified: <https://docs.databricks.com/aws/en/sql/user/alerts/> (checked 2026-09-25))

A8 Incremental processing

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Auto Loader, change data feed and processing only new waves.

A8-L1 Process only the waves you have not processed, correctly

Concept

Incremental loading means each run handles only new work: waves in bronze that silver does not have yet. It saves compute, and it keeps reruns cheap. It is also easy to get subtly wrong.

The pattern:

```
1 def unprocessed_waves(spark, bronze, silver, tracker_id) -> list[str]:
2     source = spark.table(bronze).filter(F.col("tracker_id") ==
3         tracker_id).select("wave_id").distinct()
4     done = spark.table(silver).filter(F.col("tracker_id") ==
5         tracker_id).select("wave_id").distinct()
6     return sorted(r.wave_id for r in source.join(done, "wave_id",
7         "left_anti").collect())
```

Four rules:

1. **Filter both sides the same way.** Your reference `get_unprocessed_waves` filters bronze by tracker but reads silver for every tracker. If tracker B already has wave W14 in silver, tracker A's W14 is treated as done and never loads. Scope both sides with the same keys.
2. **Use a left anti join** (`"left_anti"`): rows on the left with no match on the right. It says exactly what you mean, and it works for multi-column keys.
3. **Handle the first run.** When silver does not exist yet, every bronze wave is new. Check with `spark.catalog.tableExists(silver)`.
4. **Keep each wave's load idempotent.** Detection can be wrong after a partial failure, for example if silver has half a wave. Because the silver MERGE is idempotent (A3-L1), reprocessing a wave is always safe. Design for "process again", not "process exactly once".

Deciding by "waves not in silver" misses redeliveries of a wave silver already has. Either compare the latest `ingestion_timestamp` per wave in bronze with the one recorded when silver loaded it, or read bronze incrementally with change data feed (A8-L3) or Auto Loader (A8-L2).

Why it matters

Two trackers share one silver table. Tracker B loads wave W14 on Monday. Tracker A's W14 arrives on Tuesday. The reference `get_unprocessed_waves` compares tracker A's bronze waves against silver waves for all trackers, finds W14 already present, and skips it. Tracker A's dashboard shows no W14, and no job fails. The gap is found nine days later, by a client.

Filtering silver by `tracker_id` too makes the comparison correct, and a reconciliation of bronze keys against silver rows (A6-L3) would have flagged it on Tuesday.

Worked example: weak code

```
❶ def get_unprocessed_waves(bronze_table, silver_table, tracker_id):
    2     bronze_waves = (
    3         spark.table(bronze_table)
    4         .filter(col("tracker_id") == tracker_id)
    5         .select("wave_id").distinct()
    6     )
    ❷     silver_waves = spark.table(silver_table).select("wave_id").distinct()
    ❸     return [row.wave_id for row in
        bronze_waves.subtract(silver_waves).collect()]
```

1. **Line 1.** Uses a global spark and has no type hints.
2. **Line 7.** No tracker filter: another tracker's wave counts as done (reference defect 10).
3. **Line 8.** Unsorted output, and it fails if silver does not exist yet.

Worked example: strong code

```
1 def unprocessed_waves(spark: SparkSession, bronze: str, silver: str,
    tracker_id: str) -> list[str]:
    2     """Wave ids in bronze for this tracker that silver does not hold for
        this tracker."""
    3     source = (
    4         spark.table(bronze)
    5         .filter(F.col("tracker_id") == tracker_id)
    6         .select("wave_id")
    7         .distinct()
    8     )
    ❶     if not spark.catalog.tableExists(silver):
    10         return sorted(r.wave_id for r in source.collect())
    11     done = (
    ❷         spark.table(silver)
    13         .filter(F.col("tracker_id") == tracker_id)
    14         .select("wave_id")
    15         .distinct()
    16     )
```



```
3 return sorted(r.wave_id for r in source.join(done, "wave_id",
"left_anti").collect())
```

1. **Line 9.** First run: silver does not exist, so every wave is new.
2. **Line 12.** Silver is read with the same tracker filter as bronze.
3. **Line 17.** Left anti join: waves with no match in silver. Sorted, so waves load in order and logs are predictable.

Same filters on both sides, an explicit anti join, a first-run branch, and an idempotent load behind it.

Recap card

- **Rule for incremental wave detection** Filter bronze and silver by the same keys, such as `tracker_id`, then left-anti-join.
- **Join type for 'in A but not in B'** `left_anti`.
- **First run, silver missing** `spark.catalog.tableExists(silver)` is `False`: every bronze wave is new.
- **What does wave-ID detection miss?** Redeliveries of waves silver already holds. Use ingestion timestamps, change data feed or Auto Loader.

Practice

Exercise 1. Multiple choice (level 1)

Bronze for tracker A holds waves W13 and W14. Silver holds W13 for tracker A and W14 for tracker B. The reference function compares A's bronze waves with all of silver's waves. Which waves does it return?

- a) W14
- b) None
- c) W13 and W14
- d) W13

Why?

- r1) Tracker B's W14 in silver makes tracker A's W14 look done.
- r2) subtract removes all waves when silver is not empty.

Exercise 2. Fill the gap (level 2)

Keep only bronze waves that silver does not have, and handle a missing silver table.

```
1 if not spark.catalog.__(1)(silver):
2     return sorted(r.wave_id for r in source.collect())
```

```
3 new = source.join(done, "wave_id", "_____(2)")
```

Exercise 3. Spot the bug (level 3)

Tracker A's wave W14 never loads when tracker B already has W14. Tap the line, then fix the function.

```
1 def unprocessed_waves(spark: SparkSession, bronze: str, silver: str,
2   tracker_id: str) -> list[str]:
3   """Waves in bronze for this tracker that silver does not hold for this
4   tracker."""
5   source = spark.table(bronze).filter(F.col("tracker_id") ==
6   tracker_id).select("wave_id").distinct()
7   done = spark.table(silver).select("wave_id").distinct()
8   return sorted(r.wave_id for r in source.join(done, "wave_id",
9   "left_anti").collect())
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 4)

Write `unprocessed_waves(spark, bronze, silver, tracker_id)`. Return the sorted wave IDs present in bronze for the tracker and absent from silver for the same tracker. When silver does not exist, return every bronze wave for the tracker. Use a left anti join.

Starter code:

```
1 from pyspark.sql import SparkSession
2 from pyspark.sql import functions as F
3
4
5 def unprocessed_waves(spark: SparkSession, bronze: str, silver: str,
6   tracker_id: str) -> list[str]:
7   ...
```

Exercise 5. Multiple choice (level 5)

Silver already holds wave W14. The supplier redelivers W14 with corrected weights. Your job processes "waves in bronze that are not in silver". What happens to the redelivery?

- a) It is processed, because bronze has a new ingestion_id.
- b) It is skipped, because W14 is already in silver.
- c) The job fails with a duplicate key error.
- d) It is processed twice.

Why?

- r1) The detection compares wave IDs only, and W14 exists on both sides.

- r2) Bronze rejects redeliveries.

Sources

- left_anti joins and spark.catalog.tableExists work as described on PySpark 4.0. (verified: spark_tests for A8-L1-E3 and E4 run on PySpark 4.0.4 with Delta.)

A8-L2 Auto Loader for file ingestion into bronze

Concept

Suppliers drop files into a landing folder. You want each file loaded into bronze exactly once, whenever it arrives, without keeping your own list of processed files. **Auto Loader** does this on Databricks. It is a Structured Streaming source called `cloudFiles`.

```
1 (
2   spark.readStream.format("cloudFiles")
3   .option("cloudFiles.format", "csv")
4   .option("cloudFiles.schemaLocation", schema_path)
5   .option("header", "true")
6   .load(landing_path)
7   .select("*", F.col("_metadata.file_path").alias("source_file"))
8   .writeStream
9   .option("checkpointLocation", checkpoint_path)
10  .trigger(availableNow=True)
11  .toTable("research_prod.research_bronze.nat_raw")
12 )
```

What each piece does:

- **Checkpoint location.** Auto Loader records every discovered file in a key-value store (RocksDB) under the checkpoint. After a failure it resumes where it stopped, and each file path is processed once. One checkpoint per stream, never shared, never deleted casually: deleting it reloads every file.
- **Schema location.** Auto Loader infers the schema once, stores it there, and tracks changes. Without a schema you provide, the default mode is `addNewColumns`: when a new column appears, the stream stops with `UnknownFieldException` after recording the new schema, and the next start includes the column. Run it from a job that retries, so the restart is automatic.

- **Rescued data.** When Auto Loader infers the schema, it adds a `_rescued_data` column that captures values that do not fit, such as a type mismatch. Count non-null `_rescued_data` in your contract checks.
- `trigger(availableNow=True)` processes every file that has arrived, then stops. Schedule it as a batch job: incremental logic, batch cost.
- `_metadata.file_path` gives each row its source file, for the bronze audit columns.

Keep checkpoints and schemas in a Unity Catalog volume, such as `/Volumes/research_prod/research_bronze/checkpoints/nat_raw`.

Why it matters

A bronze notebook lists the landing folder, compares file names with a hand-maintained "processed.txt", and loads the rest. A timeout kills the notebook after it has loaded a file but before it has updated the list. The next run loads the same file again. Bronze now holds the delivery twice with two ingestion IDs, and silver's newest-version logic cannot tell them apart, because both carry the same `delivery_version`.

Auto Loader's checkpoint records files and commits the Delta write together, so a restart after a failure neither skips nor repeats a file.

Worked example: weak code

```
1 import os
2
3 done = set(open("/dbfs/tmp/processed.txt").read().split())
4 for name in os.listdir("/dbfs/landing/nat"):
5     if name not in done:
6         df = spark.read.option("header", True).csv(f"/landing/nat/{name}")
7         df.write.mode("append").saveAsTable("research_bronze.nat_raw")
8         open("/dbfs/tmp/processed.txt", "a").write(name + "\n")
```

1. **Line 3.** Hand-built state in a temp file: easy to lose, and not linked to the table.
2. **Line 7.** The write and the list update are separate steps. A crash between them duplicates a file.
3. **Line 8.** State recorded after the write, with no transaction.

Worked example: strong code

```
1 def start_bronze_ingest(spark: SparkSession, landing: str, table: str,
2     checkpoint: str, schema_path: str) -> StreamingQuery:
3     """Load every new file in landing into bronze once, then stop."""
4     raw = (
5         spark.readStream.format("cloudFiles")
6         .option("cloudFiles.format", "csv")
7         .option("cloudFiles.schemaLocation", schema_path)
```

```

7      .option("header", "true")
8      .load(landing)
9      )
10     audited = raw.select(
11         "*",
12         F.col("_metadata.file_path").alias("source_file"),
13         F.current_timestamp().alias("ingestion_timestamp"),
14     )
15     return (
16         audited.writeStream.option("checkpointLocation", checkpoint)
17         .trigger(availableNow=True)
18         .toTable(table)
19     )

```

1. **Line 4.** Auto Loader discovers new files and tracks them in the checkpoint.
2. **Line 6.** Where the inferred schema and its changes are stored.
3. **Line 12.** Source file for every row, for the bronze audit columns.
4. **Line 17.** Process everything available, then stop: scheduled like a batch job.

Auto Loader replaces a fragile processed-files list with a checkpoint that commits together with the Delta write.

Recap card

- **What does Auto Loader's checkpoint store?** Discovered files (in RocksDB) and stream progress, so each file is processed exactly once across restarts.
- **Default schema evolution mode without a supplied schema**
addNewColumns: stop with UnknownFieldException, record the new schema, restart.
- **Incremental at batch cost** .trigger(availableNow=True) on a scheduled job.
- **Checkpoint rule** One per stream, in durable storage such as a UC volume. Never shared, never deleted casually.

Practice

Exercise 1. Multiple choice (level 1)

What lets Auto Loader process each landing file exactly once across restarts?

- a) The schema location
- b) The checkpoint location
- c) The `_rescued_data` column
- d) The file names in the target table

Why?

- r1) Auto Loader records discovered files in a key-value store in the checkpoint and resumes from it.
- r2) The checkpoint deletes files after loading them.

Exercise 2. Fill the gap (level 2)

Complete the Auto Loader read and a write that processes all available files, then stops.

```
1 raw = (  
2     spark.readStream.format("_____(1)")  
3     .option("cloudFiles.format", "csv")  
4     .option("cloudFiles.____(2)", schema_path)  
5     .load(landing)  
6 )  
7 raw.writeStream.option("checkpointLocation",  
    checkpoint).trigger(____(3)=True).toTable(table)
```

Exercise 3. Multiple choice (level 3)

An Auto Loader stream runs without a user-supplied schema. A new delivery adds a column `region`. What happens?

- a) The column is silently dropped.
- b) The stream stops with `UnknownFieldException` after recording the new schema. The next start includes `region`.
- c) The stream continues and `region` lands in `_rescued_data`, forever.
- d) The whole bronze table is rebuilt.

Why?

- r1) The default mode without a schema is `addNewColumns`, which evolves by failing once and restarting.
- r2) Auto Loader never changes a schema automatically.

Exercise 4. Spot the bug (level 4)

Two Auto Loader streams load different landing folders into two bronze tables. After a week, each table is missing files the other received. Tap the line, then fix the function.

```
1 CHECKPOINT = "/Volumes/research_prod/research_bronze/checkpoints/bronze"  
2  
3  
4 def start_ingest(spark: SparkSession, landing: str, table: str,  
    schema_path: str) -> StreamingQuery:  
5     """Load new files from landing into table."""
```

```

6     raw =
    spark.readStream.format("cloudFiles").option("cloudFiles.format",
    "csv").option("cloudFiles.schemaLocation", schema_path).load(landing)
7     return raw.writeStream.option("checkpointLocation",
    CHECKPOINT).trigger(availableNow=True).toTable(table)

```

Write down the line number of the bug. Then write the fix.

Exercise 5. Architecture design (level 5)

Design the Auto Loader job for tracker nat. Choose each setting.

Options: trigger(availableNow=True), scheduled, A continuously running stream.

1. Files arrive a few times a day, and cost matters more than latency
2. Checkpoint location
3. What to do with _rescued_data

Sources

- Auto Loader tracks files in RocksDB in the checkpoint for exactly-once processing; _metadata.file_path is available; checkpoints can live in Volumes paths. (verified: <https://docs.databricks.com/aws/en/ingestion/cloud-object-storage/auto-loader/> (checked 2026-09-25))
- addNewColumns is the default schema evolution mode when no schema is provided; the stream stops with UnknownFieldException and picks up new columns on restart; _rescued_data is added when the schema is inferred. (verified: <https://docs.databricks.com/aws/en/ingestion/cloud-object-storage/auto-loader/schema> (checked 2026-09-25))
- DataStreamWriter.trigger accepts availableNow=True and toTable exists in PySpark 4.0. (verified: Checked on PySpark 4.0.4 in this repository.)

A8-L3 Change data feed, so downstream layers read only what changed

Concept

When silver changes, gold should reprocess only the rows that changed, not the whole table. **Change data feed** (CDF) makes a Delta table record row-level changes for every write.

Turn it on per table:

```

1 ALTER TABLE research_prod.research_silver.nat_respondents
2 SET TBLPROPERTIES (delta.enableChangeDataFeed = true);

```

Read the changes since a version:

```
1 changes = (  
2     spark.read.option("readChangeFeed", "true")  
3     .option("startingVersion", last_version + 1)  
4     .table("research_prod.research_silver.nat_respondents")  
5 )
```

Each row carries three extra columns:

- `_change_type`: `insert`, `update_preimage` (the row before an update), `update_postimage` (the row after), or `delete`.
- `_commit_version` and `_commit_timestamp`: which commit made the change.

Two rules make CDF consumers correct:

1. **Drop** `update_preimage` **rows** before you use the changes. An update produces two rows, old and new. Keep the old one and every updated respondent is counted twice.
2. **Keep one row per key**: the latest by `_commit_version`. A respondent updated in three commits since your last run appears three times.

Track progress. Either store the last processed version in a state table, or read the feed as a stream (`readStream` with `readChangeFeed`) and let the checkpoint remember it. Change data only exists from the version where you enabled it. Asking for an earlier version fails with `DELTA_MISSING_CHANGE_DATA`. `VACUUM` also removes old change files after the retention window.

For gold, collect the waves and subgroups the changes touch, then recompute just those aggregates. Aggregates cannot be patched row by row.

Why it matters

A gold job rebuilds weighted metrics for every wave every night: 180 waves, 40 minutes. Most nights, only one wave changed. When a late correction arrives mid-afternoon, the team cannot refresh gold before the client call, because a full rebuild takes too long.

With CDF on silver, the job reads the day's changes, finds the one wave they touch, and recomputes it in under two minutes. The same job also stops double counting, because it drops `update_preimage` rows before aggregating.

Worked example: weak code

```
❶ def changed_rows(spark, table, since):  
    changes = (  
        3     spark.read.option("readChangeFeed", "true")  
        ❷     .option("startingVersion", since)  
        5     .table(table)
```



```

6 )
3 return changes.filter(F.col("_change_type") != "delete")

```

1. **Line 1.** A respondent updated in several commits still appears once per commit.
2. **Line 4.** If since is the last processed version, that version is read again.
3. **Line 7.** Keeps update_preimage rows: every updated respondent appears twice, once with old values.

Worked example: strong code

```

1 CURRENT_STATE = ["insert", "update_postimage"]
2
3
4 def read_changes(spark: SparkSession, table: str, last_processed_version:
  int) -> DataFrame:
5     """Row changes committed after last_processed_version."""
6     return (
7         spark.read.option("readChangeFeed", "true")
8         1 .option("startingVersion", last_processed_version + 1)
9         .table(table)
10    )
11
12
13 def latest_state(changes: DataFrame, keys: list[str]) -> DataFrame:
14     """One row per key: its newest change, excluding pre-images."""
15     2 window =
16     Window.partitionBy(*keys).orderBy(F.col("_commit_version").desc())
17     return (
18         3 changes.filter(F.col("_change_type") != "update_preimage")
19         .withColumn("_rank", F.row_number().over(window))
20         .filter(F.col("_rank") == 1)
21         .drop("_rank")
22    )
23
24 4 def touched_waves(changes: DataFrame) -> list[str]:
25     """Waves with any change: the only gold waves to recompute."""
26     return sorted(r.wave_id for r in
27         changes.select("wave_id").distinct().collect())

```

1. **Line 8.** Start after the last processed version, so nothing is read twice.
2. **Line 15.** Several changes to one key since the last run: keep the newest. Deletes stay, so the consumer can remove the key.
3. **Line 17.** Pre-images describe old state. Drop them.
4. **Line 24.** Gold recomputes whole aggregates, but only for the waves that changed.

Read after the last version, drop pre-images, keep the newest change per key, and recompute only what the changes touch.

Recap card

- **Turn on change data feed** ALTER TABLE t SET TBLPROPERTIES (delta.enableChangeDataFeed = true)
- **The four _change_type values** insert, update_preimage (old), update_postimage (new), delete.
- **Two rules for consuming CDF** Drop update_preimage. Keep the newest change per key by _commit_version.
- **Where does the next run start?** startingVersion = last processed version + 1 (it is inclusive), or a streaming checkpoint.

Practice

Exercise 1. Multiple choice (level 1)

An UPDATE changes one respondent's weight. Which change feed row holds the new weight?

- a) insert
- b) update_preimage
- c) update_postimage
- d) delete

Why?

- r1) An update records the row before (pre-image) and after (post-image). The post-image holds the new values.
- r2) Updates are recorded as a delete followed by an insert.

Exercise 2. Predict the output (level 2)

CDF is enabled when the table is created. How many rows does the change read return?

```
1 CREATE TABLE c1 (id INT, v STRING) USING DELTA
2   TBLPROPERTIES (delta.enableChangeDataFeed = true);    -- version 0
3 INSERT INTO c1 VALUES (1, 'a'), (2, 'b');              -- version 1
4 UPDATE c1 SET v = 'z' WHERE id = 1;                     -- version 2
5 DELETE FROM c1 WHERE id = 2;                            -- version 3
6
7 SELECT COUNT(*) FROM table_changes('c1', 2);
```

Exercise 3. Fill the gap (level 3)

Read silver's changes committed after the last processed version.

```
1 changes = (
2   spark.read.option("_____(1)", "true")
```

```

3     .option("startingVersion", _____.(2))
4     .table(silver)
5 )

```

Exercise 4. Write the function (level 4)

Write `latest_state(changes, keys)`. From a change feed DataFrame, drop `update_preimage` rows, then keep one row per key: the one with the highest `_commit_version`. Use a window with `row_number`.

Starter code:

```

1 from pyspark.sql import DataFrame, Window
2 from pyspark.sql import functions as F
3
4
5 def latest_state(changes: DataFrame, keys: list[str]) -> DataFrame:
6     ...

```

Exercise 5. Spot the bug (level 5)

After enabling CDF, updated respondents appear twice in the gold input, once with old values. Tap the line, then fix the function.

```

1 def current_rows(spark: SparkSession, table: str, last_processed_version:
  int) -> DataFrame:
2     """Inserted and updated rows since the last run, with their new
  values."""
3     changes = spark.read.option("readChangeFeed",
  "true").option("startingVersion", last_processed_version +
  1).table(table)
4     return changes.filter(F.col("_change_type") != "delete")

```

Write down the line number of the bug. Then write the fix.

Practical: Read a change data feed and keep the latest state

Platform: Databricks Free Edition. Run this in Databricks Free Edition. It enables change data feed on a small silver table, applies an insert, an update and a delete, then reads the changes and reduces them to the latest state per respondent.

Download the notebook `A8-L3.ipynb` from the app.

Sources

- CDF read semantics, the four change types, `table_changes`, and `DELTA_MISSING_CHANGE_DATA` for versions before CDF was enabled. (verified: Run on delta-spark 4.0.1 in this repository.)

- CDF on Databricks: `delta.enableChangeDataFeed`, `readChangeFeed`, `startingVersion`, `table_changes`. (verified: <https://docs.databricks.com/aws/en/tables/features/change-data-feed> (checked 2026-09-25))

A9 Testing

Roadmap: Months 16 to 18: MLOps and deployment; Current work. pytest, fixtures, local SparkSession and DataFrame equality checks.

A9-L1 pytest basics, and a SparkSession fixture shared by every test

Concept

pytest finds files named `test_*.py`, runs every function named `test_*`, and reports each failing `assert` with the values on both sides. No classes or special assert methods are needed.

```
1 def test_weighted_share_uses_weights() -> None:
2     assert weighted_share([1, 0], [3.0, 1.0]) == 0.75
```

Fixtures provide what tests need. A test names a fixture as a parameter, and pytest builds it and passes it in. For Spark, one fixture creates the session:

```
1 @pytest.fixture(scope="session")
2 def spark() -> Iterator[SparkSession]:
3     session = SparkSession.builder.master("local[2]").getOrCreate()
4     yield session
5     session.stop()
```

- `scope="session"` builds it once for the whole test run. Starting Spark takes several seconds. Once per test, and a 40-test suite spends minutes on start-up alone.
- `yield` hands the value to the tests. The code after `yield` runs once at the end, which is where you stop the session.
- Put shared fixtures in `conftest.py`. pytest loads it automatically for every test in that folder and below.

`@pytest.mark.parametrize` runs one test with several inputs, and each input is reported as its own test:

```
1 @pytest.mark.parametrize(("flags", "weights", "expected"), [([1], [2.0],
2     1.0), ([0, 0], [1.0, 1.0], 0.0)])
3 def test_share(flags, weights, expected): ...
```

A good test fails for the bug you fear. Write it against the known answer on small data, including the edge cases: empty input, zero totals, missing values. Check it by breaking the code on purpose and watching the test fail.

Why it matters

A team's test for `weighted_share` checks that the result is between 0 and 1. The function is later rewritten and accidentally returns the unweighted share. The result is still between 0 and 1, so the test passes, and `weighted` reach in the dashboard is wrong for a month.

A test with a known answer, `weighted_share([1, 0], [3.0, 1.0]) == 0.75`, fails the moment the weights are ignored, because the unweighted answer is 0.5.

Worked example: weak code

```
1 import pytest
2 from pyspark.sql import SparkSession
3
4
5 ① @pytest.fixture
6 def spark():
7     session = SparkSession.builder.master("local[2]").getOrCreate()
8     yield session
9     session.stop()
10
11
12 def test_share():
13     share = weighted_share([1, 0, 1], [2.0, 1.0, 1.0])
14 ② assert 0 <= share <= 1
```

1. **Line 5.** Default function scope: Spark starts and stops for every test that uses it.
2. **Line 14.** Passes for almost any wrong implementation. It does not test the rule.

Worked example: strong code

```
1 from collections.abc import Iterator
2
3 import pytest
4 from pyspark.sql import SparkSession
5
6
7 ① @pytest.fixture(scope="session")
8 def spark() -> Iterator[SparkSession]:
9     """One local session for the whole test run."""
10     session =
11     SparkSession.builder.master("local[2]").config("spark.ui.enabled",
12     "false").getOrCreate()
13 ② yield session
14     session.stop()
15
16 @pytest.mark.parametrize(
17     ("flags", "weights", "expected"),
```

```

17     [
18         ③      ([1, 0, 1], [2.0, 1.0, 1.0], 0.75),
19         ([0, 0], [1.0, 1.0], 0.0),
20         ④      ([1, 0], [0.0, 0.0], None),
21     ],
22 )
23 def test_weighted_share(flags: list[int], weights: list[float], expected:
    float | None) -> None:
24     assert weighted_share(flags, weights) == expected

```

1. **Line 7.** One Spark start-up for the whole run.
2. **Line 11.** Tests run here. The line after yield is the teardown.
3. **Line 18.** A known answer that differs from the unweighted one (0.67), so ignoring weights fails.
4. **Line 20.** The edge case: zero total weight must give None.

One session-scoped fixture, and tests with known answers that fail for the bugs you care about.

Recap card

- **What does pytest run?** Functions named test_ in files named test_.py. Plain assert statements.
- **SparkSession fixture scope** scope="session": start Spark once per test run.
- **yield in a fixture** Hands the value to tests. Code after yield is the teardown.
- **What makes a test worth having?** It fails for the bug you fear: known answers that separate right from wrong, plus edge cases.

Practice

Exercise 1. Multiple choice (level 1)

Which fixture scope suits a SparkSession used by 40 tests?

- a) function
- b) session
- c) No fixture: create a session inside each test
- d) class

Why?

- r1) Spark takes seconds to start, and tests can safely share one local session.
- r2) Only session-scoped fixtures can use yield.

Exercise 2. Predict the output (level 2)

How many tests does pytest report for this file?

```
1 import pytest
2
3
4 @pytest.mark.parametrize("wave", ["W13", "W14", "W15"])
5 def test_wave_loads(wave: str) -> None:
6     assert wave.startswith("W")
7
8
9 def test_empty_wave_list() -> None:
10     assert unprocessed([]) == []
```

Exercise 3. Fill the gap (level 2)

Complete the fixture so Spark starts once per test run and stops at the end.

```
1 @pytest.__(1)(scope="__(2)")
2 def spark() -> Iterator[SparkSession]:
3     session = SparkSession.builder.master("local[2]").getOrCreate()
4     __(3) session
5     session.stop()
```

Exercise 4. Write the function (level 4)

`weighted_share(flags, weights)` returns the share of weight where the flag is 1, or `None` when the total weight is 0. Write **the test**, `test_weighted_share()`, not the function.

The checker runs your test three times: against a correct implementation (it must pass), against one that ignores the weights, and against one that returns 0.0 instead of `None` for a zero total (both must fail).

Starter code:

```
1 def test_weighted_share() -> None:
2     ...
```

Exercise 5. Spot the bug (level 5)

The test suite takes six minutes for 45 small tests. Tap the line that causes most of it, then fix the fixture.

```
1 import pytest
2 from pyspark.sql import SparkSession
3
4
5 @pytest.fixture
6 def spark():
7     session = SparkSession.builder.master("local[2]").getOrCreate()
8     yield session
9     session.stop()
```


Write down the line number of the bug. Then write the fix.

Sources

- pytest collects one test per parametrize value; fixtures with yield run teardown after; scope='session' builds once per run. (verified: pytest 9.1.1 behaviour, used by this repository's own test suite.)

A9-L2 A fast local SparkSession and readable test data

Concept

Unit tests for PySpark run on your laptop and in CI, not on a cluster. A local session that starts in seconds and runs small jobs fast needs a few settings:

```
1 builder = (  
2     SparkSession.builder.master("local[2]")  
3     .config("spark.sql.shuffle.partitions", "2")  
4     .config("spark.ui.enabled", "false")  
5     .config("spark.sql.extensions",  
6         "io.delta.sql.DeltaSparkSessionExtension")  
7     .config("spark.sql.catalog.spark_catalog",  
8         "org.apache.spark.sql.delta.catalog.DeltaCatalog")  
9     .config("spark.sql.sources.default", "delta")  
10 )  
11 spark = configure_spark_with_delta_pip(builder).getOrCreate()
```

- `local[2]` runs Spark in one process with two worker threads. Two threads, not one, so bugs that depend on partitioning still show up.
- `spark.sql.shuffle.partitions` defaults to 200. On ten test rows, every shuffle creates 200 tiny tasks. Two is plenty.
- `configure_spark_with_delta_pip` from `delta-spark` adds the matching Delta JAR. The two `config` lines enable Delta SQL and the Delta catalog.
- Open-source Spark writes `saveAsTable` as Parquet by default. Databricks uses Delta. Set the default so tests behave like production.

Test data. Build it with an explicit schema, as a DDL string such as `"respondent_id STRING, final_weight DOUBLE"`. Inference guesses types from values, fails on a column of only `None` (`CANNOT_DETERMINE_TYPE`), and turns `2` into a `LONG` where production has `DOUBLE`.

A **test data builder** keeps tests short: one function returns a valid row with sensible defaults, and each test overrides only the field it is about.

```
1 def respondent(**overrides: object) -> dict[str, object]:  
2     row = {"respondent_id": "R1", "wave_id": "W1", "final_weight": 1.0,  
3         "in_scope": True}
```

```

3     row.update(overrides)
4     return row

```

A test that says `respondent(final_weight=None)` tells the reader exactly what it is testing.

Why it matters

A project's Spark tests take 14 minutes in CI, so developers stop running them locally. The session uses the default 200 shuffle partitions, and each of the 60 tests runs two or three shuffles on a handful of rows. Most of the time goes into scheduling empty tasks. A MERGE bug that a test would have caught ships to production, because nobody waited for the suite.

With `local[2]` and two shuffle partitions, the same suite runs in under two minutes, and people run it before every push.

Worked example: weak code

```

1 @pytest.fixture(scope="session")
2 def spark():
3     return SparkSession.builder.getOrCreate()
4
5
6 def test_null_weights_are_reported(spark):
7     df = spark.createDataFrame([{"respondent_id": "R1", "final_weight":
8         None, "wave_id": "W1", "in_scope": True}])
9     assert null_weight_count(df) == 1

```

1. **Line 3.** No master, 200 shuffle partitions, UI on, no Delta: slow and unlike production.
2. **Line 7.** Inferred schema: a column of only None raises CANNOT_DETERMINE_TYPE, and every field is spelt out in every test.

Worked example: strong code

```

1 RESPONDENT_SCHEMA = "respondent_id STRING, wave_id STRING, final_weight
2   DOUBLE, in_scope BOOLEAN"
3
4 def respondent(**overrides: object) -> dict[str, object]:
5     """A valid silver respondent. Override only the fields a test is
6     about."""
7     row: dict[str, object] = {"respondent_id": "R1", "wave_id": "W1",
8         "final_weight": 1.0, "in_scope": True}
9     unknown = set(overrides) - set(row)
10    if unknown:
11        raise KeyError(f"unknown fields: {sorted(unknown)}")
12    row.update(overrides)
13    return row

```

```

13
14 def respondents_df(spark: SparkSession, rows: list[dict[str, object]]) ->
    DataFrame:
15     """DataFrame with the silver schema, whatever the values."""
16     fields = [f.split()[0] for f in RESPONDENT_SCHEMA.split(", ")]
17     return spark.createDataFrame([tuple(r[f] for f in fields) for r in
    rows], RESPONDENT_SCHEMA)
18
19
20 def test_null_weights_are_reported(spark: SparkSession) -> None:
21     ③ df = respondents_df(spark, [respondent(),
    respondent(respondent_id="R2", final_weight=None)])
22     assert null_weight_count(df) == 1

```

- Line 1.** One DDL schema: types match production, and None is always allowed.
- Line 9.** A typo such as final_wieght fails loudly instead of silently adding nothing.
- Line 21.** The test reads as its intent: a second respondent with a missing weight.

A tuned local session makes tests fast enough to run on every change. Builders with explicit schemas make them short and exact.

Recap card

- **Local test session settings** master("local[2]"), shuffle partitions 2, UI off, Delta configured with configure_spark_with_delta_pip.
- **Default spark.sql.shuffle.partitions** 200: far too many for test data.
- **Why explicit schemas in test data?** Inference fails on all-None columns and can pick different types from production.
- **Test data builder** A function with valid defaults. Tests override only the fields they are about, and typos raise.

Practice

Exercise 1. Predict the output (level 1)

A fresh local session on PySpark 4.0 with no extra configuration. What does this print?

```
1 print(spark.conf.get("spark.sql.shuffle.partitions"))
```

Exercise 2. Fill the gap (level 2)

Complete the local test session: two threads, two shuffle partitions, with the Delta JAR added.

```

1 builder = (
2     SparkSession.builder.master("____(1)")

```

```

3     .config("spark.sql.shuffle.partitions", "____(2)")
4     .config("spark.sql.extensions",
5         "io.delta.sql.DeltaSparkSessionExtension")
6     .config("spark.sql.catalog.spark_catalog",
7         "org.apache.spark.sql.delta.catalog.DeltaCatalog")
8 )
9 spark = _____.getOrCreate()

```

Exercise 3. Write the function (level 3)

Write `respondent(**overrides)`. It returns a dict with defaults

`respondent_id="R1"`, `wave_id="W1"`, `final_weight=1.0`, `in_scope=True`, updated with the overrides. Raise `KeyError` naming any override that is not one of those four fields. Each call must return a new dict.

Starter code:

```

1 def respondent(**overrides: object) -> dict[str, object]:
2     ...

```

Exercise 4. Spot the bug (level 4)

This helper builds test data. It works for most tests, but fails with

`CANNOT_DETERMINE_TYPE` in the test for missing weights. Tap the line, then fix the function with an explicit schema.

```

1 def weights_df(spark: SparkSession, weights: list[float | None]) ->
    DataFrame:
2     """One respondent per weight."""
3     rows = [(f"R{i}", w) for i, w in enumerate(weights)]
4     return spark.createDataFrame(rows, ["respondent_id", "final_weight"])

```

Write down the line number of the bug. Then write the fix.

Exercise 5. Put in order (level 5)

Order the set-up of a local PySpark test harness.

- Write a session-scoped spark fixture in `conftest.py`
- Create a virtual environment and install `pyspark`, `delta-spark` and `pytest`
- Write tests with known answers against the builders
- Write test data builders with explicit schemas
- Install Java 17 or 21

Sources

- `spark.sql.shuffle.partitions` defaults to 200; `createDataFrame` with only column names infers types and fails on all-None columns. (verified: Run on PySpark 4.0.4 in this repository.)

- `configure_spark_with_delta_pip` adds the Delta package matching the installed delta-spark version. (verified: Used by `practice/conftest.py` with delta-spark 4.0.1.)

A9-L3 Compare DataFrames properly, and test that MERGE is idempotent

Concept

Comparing DataFrames by hand is harder than it looks. `actual.collect() == expected.collect()` fails when the rows come back in a different order, which Spark never guarantees. It also fails when a float is `0.30000000000000004` instead of `0.3`. Sorting both sides fixes the order but not the floats, and neither checks the schema.

`pyspark.testing` (PySpark 3.5 and later) handles all of this:

```
1 from pyspark.testing import assertDataFrameEqual, assertSchemaEqual
2
3 assertDataFrameEqual(actual, expected)
```

Tested on PySpark 4.0.4:

- Row order is ignored by default. Pass `checkRowOrder=True` when order is the point, as in a top-10 list.
- Floats compare with a tolerance (`rtol=1e-5`, `atol=1e-8` by default).
- Schemas must match: INT against BIGINT fails with `DIFFERENT_SCHEMA`.
- Duplicates count: two identical rows do not equal one.
- `expected` can be a list of `Row` objects, which is handy for snapshots.
- Failures raise `PySparkAssertionError`, a subclass of `AssertionError`, so `pytest` reports them as ordinary test failures, with a diff.

`assertSchemaEqual(actual.schema, expected.schema)` checks only the structure.

Testing MERGE. The property that matters is idempotence. The test:

1. Create an empty Delta table with the target schema.
2. Run the load once, and take a snapshot:
`first = spark.table(t).collect()`
3. Run the same load again.
4. `assertDataFrameEqual(spark.table(t), first)`.

Add cases for the failures you fear: a repeated key in the source, a redelivery with a newer version, and the same respondent in two waves. Each one takes a few rows and a second or two.

In the practice harness, PySpark 4.0.4 needs pandas 2.x. With pandas 3.x, `assertDataFrameEqual` fails to import.

Why it matters

A silver test compares `actual.collect() == expected.collect()`. It passes on a developer's laptop and fails one CI run in five, because the order of rows varies with partitioning. The team marks it flaky and skips it. A month later a real regression in the MERGE, duplicate rows after a retry, ships untested.

`assertDataFrameEqual` ignores order, so the test is stable, and an idempotence test that runs the load twice catches the regression.

Worked example: weak code

```
1 def test_silver_load(spark):
2     upsert_silver(spark, source, "silver", ["respondent_id", "wave_id"],
3         "delivery_version")
4     assert spark.table("silver").collect() == expected.collect()
```

1. **Line 2.** One run only. It never checks what a retry does.
2. **Line 3.** Order-dependent and exact on floats: flaky, and a failure shows no useful diff.

Worked example: strong code

```
1 from pyspark.testing import assertDataFrameEqual
2
3 KEYS = ["respondent_id", "wave_id"]
4 SCHEMA = "respondent_id STRING, wave_id STRING, final_weight DOUBLE,
5     delivery_version INT"
6
7 def test_silver_load_is_idempotent(spark: SparkSession, empty_silver:
8     str) -> None:
9     source = spark.createDataFrame(
10         [ ("R1", "W1", 1.2, 1), ("R1", "W1", 1.5, 2), ("R2", "W1", 0.8,
11             1)], SCHEMA
12     )
13     upsert_silver(spark, source, empty_silver, KEYS, "delivery_version")
14     first = spark.table(empty_silver).collect()
15
16     upsert_silver(spark, source, empty_silver, KEYS, "delivery_version")
17
18     assertDataFrameEqual(spark.table(empty_silver), first)
```

```
17     expected = spark.createDataFrame([("R1", "W1", 1.5, 2), ("R2", "W1",
    0.8, 1)], SCHEMA)
④     assertDataFrameEqual(spark.table(empty_silver), expected)
```

1. **Line 9.** The source includes a repeated key, the case that breaks a naive MERGE.
2. **Line 12.** A snapshot of the table after one run, as a list of Rows.
3. **Line 16.** Second run leaves the table unchanged: idempotence, order-independent.
4. **Line 18.** And the content is right: the newest version of R1 won.

One helper compares DataFrames the right way, and one small test proves a MERGE survives a retry and a repeated key.

Recap card

- **Compare two DataFrames in a test**
pyspark.testing.assertDataFrameEqual(actual, expected). Ignores order, tolerates float rounding, checks schema.
- **When should row order matter?** Only for ranked output. Pass `checkRowOrder=True`.
- **Idempotence test for a MERGE** Load, snapshot with `collect()`, load again, `assertDataFrameEqual(table, snapshot)`.
- **Why collect() for the snapshot?** A DataFrame is lazy and would re-read the table after the second run.

Practice

Exercise 1. Multiple choice (level 1)

`actual` holds rows (1, 0.3) and (2, 1.0). `expected` holds (2, 1.0) and (1, 0.1 + 0.2). Both have the schema `id INT, w DOUBLE`. What does `assertDataFrameEqual(actual, expected)` do on PySpark 4.0?

- a) Passes
- b) Fails because the rows are in a different order
- c) Fails because $0.1 + 0.2$ is not exactly 0.3
- d) Fails because it needs `checkRowOrder=True`

Why?

- r1) It ignores row order by default and compares floats with a relative tolerance.
- r2) It only compares schemas.

Exercise 2. Predict the output (level 2)

A hand-written comparison sorts both sides first. What does it print?

```
1 actual = [(1, 0.1 + 0.2), (2, 1.0)]
2 expected = [(2, 1.0), (1, 0.3)]
3 print(sorted(actual) == sorted(expected))
```

Exercise 3. Fill the gap (level 2)

Import the helpers, then check that the top-10 regions come back in the right order.

```
1 from pyspark.____(1) import assertDataFrameEqual, assertSchemaEqual
2
3 assertDataFrameEqual(top_regions, expected, ____ (2)=True)
```

Exercise 4. Write the function (level 4)

Write `check_idempotent(spark, load, table)`. `load` is a function that takes no arguments and loads a fixed source into `table`. Run it once, snapshot the table with `collect()`, run it again, and assert the table still equals the snapshot with `assertDataFrameEqual`.

Starter code:

```
1 from collections.abc import Callable
2
3 from pyspark.sql import SparkSession
4 from pyspark.testing import assertDataFrameEqual
5
6
7 def check_idempotent(spark: SparkSession, load: Callable[[], None],
8   table: str) -> None:
9     ...
```

Exercise 5. Spot the bug (level 5)

This test fails about one CI run in five, though the code under test is correct. Tap the line, then fix the test.

```
1 def test_gold_metrics(spark: SparkSession) -> None:
2     silver = spark.createDataFrame([(("R1", "W1", "N", 1.0), ("R2", "W1",
3     "S", 1.0)], "respondent_id STRING, wave_id STRING, region STRING,
4     final_weight DOUBLE")
5     expected = spark.createDataFrame([(("N", 1.0), ("S", 1.0)], "region
6     STRING, weighted_base DOUBLE")
7     actual = aggregate_base(silver).select("region", "weighted_base")
8     assert actual.collect() == expected.collect()
```

Write down the line number of the bug. Then write the fix.

Sources

- `assertDataFrameEqual` ignores row order by default, uses `rtol=1e-5` and `atol=1e-8`, fails on schema differences and duplicates, accepts a list of Rows as expected, and raises `PySparkAssertionError` (a subclass of `AssertionError`). (verified: Run on PySpark 4.0.4 with pandas 2.3.3 in this repository.)
- PySpark 4.0.4 fails to import `assertDataFrameEqual` with pandas 3.x. (verified: `ImportError: cannot import name '_builtin_table' from 'pandas.core.common', reproduced with pandas 3.0.6.`)

A10 Deployment

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Git folders, Databricks Asset Bundles, parameterised jobs and CI.

A10-L1 Git folders, one branch per change, and no secrets in the repository

Concept

Production pipeline code lives in Git, not in someone's workspace folder. Databricks **Git folders** connect a workspace folder to a remote repository, so you can edit, commit, pull and switch branches from the workspace.

The workflow Databricks documents:

- Each developer clones the repository into their own Git folder under `/Workspace/Users/<email>/` and works on their own branch.
- Use a **trunk-based** strategy: short-lived feature branches off `main`, merged through a pull request with review and CI checks.
- **Production Git folders** are different. An admin creates them outside user folders. They hold a deployment branch, and only automation updates them, after a pull request merges. Nobody edits code there by hand.
- For deploying jobs and their configuration together, Databricks recommends Declarative Automation Bundles (lesson A10-L2).

Store notebooks as source files (`.py` with `# Databricks notebook source` headers, or `.ipynb`). Put reusable functions in plain Python modules that notebooks import, so tests can import them too.

Secrets never go into the repository. Not in code, not in config files, not in notebook outputs. Git keeps every version forever, and removing a token in a later commit does not remove it from history. Rotate any secret that was ever committed.

In Databricks, read secrets at run time:

```
1 api_key = dbutils.secrets.get(scope="research", key="supplier_api_key")
```

Databricks redacts secret values read this way, showing `[REDACTED]` if you print them. Redaction does not stop deliberate transformations, so never try to display a secret. In CI, secrets come from the CI system's secret store as environment variables.

Add a `.gitignore` for local files (`.venv/`, `__pycache__`, `.env`), so credentials and environments never get staged by accident.

Why it matters

A developer tests a supplier API from a notebook with a hardcoded token, then commits the notebook to share the code. A week later they replace the token with a secret lookup in a new commit, believing the problem is fixed. The repository is later mirrored to a partner, and the token is still in the history. The supplier revokes the key after detecting unusual use, and every ingestion job that depends on it fails for two days.

With `dbutils.secrets.get` from the first line, there is nothing to leak.

Worked example: weak code

```
❶ # silver_nat transform (edited directly in the production folder)
❷ SUPPLIER_TOKEN = "example-token-not-real-1234"
❸ CATALOG = "research_prod"
4
5 def fetch_deliveries():
6     return requests.get(URL, headers={"Authorization": f"Bearer {SUPPLIER_TOKEN}"})
```

1. **Line 1.** Hand edits in production skip review and CI, and drift from Git.
2. **Line 2.** A token in code. Once committed, it stays in Git history for good.
3. **Line 3.** Environment hardcoded, so the same code cannot run in dev.

Worked example: strong code

```
1 # Databricks notebook source
❶ # silver_nat transform: edited on a feature branch, deployed by CI after review.
3 from pipelines.config import load_run_config
❷ from pipelines.supplier import fetch_deliveries
5
❸ cfg = load_run_config(dbutils.widgets.get("config_path"))
❹ token = dbutils.secrets.get(scope="research", key="supplier_api_key")
8 deliveries = fetch_deliveries(cfg.supplier_url, token)
```

1. **Line 2.** Changes reach production only through a reviewed, tested merge.
2. **Line 4.** Logic lives in importable modules that unit tests can reach.
3. **Line 6.** The environment comes from a job parameter, not from the code.
4. **Line 7.** Read at run time from a secret scope. Nothing sensitive in the repository.

Code flows from a branch, through review and CI, to production. Secrets flow from a secret store at run time, never through Git.

Recap card

- **Where do developers work in Databricks Git folders?** Their own clone under `/Workspace/Users/<email>/`, on their own branch.
- **Who updates production Git folders?** Only automation, after a pull request merges into the deployment branch.
- **Read a secret in a notebook** `dbutils.secrets.get(scope="...", key="...")`. Output shows [REDACTED].
- **A token was committed, then removed. Is it safe?** No. Git history keeps it. Rotate the token.

Practice

Exercise 1. Multiple choice (level 1)

Who should update the production Git folder?

- a) Any developer, after testing their change
- b) Automation, after a pull request merges into the deployment branch
- c) The team lead, by editing notebooks there directly
- d) Nobody: production reads from developers' user folders

Why?

- r1) Production changes must pass review and CI, and automation makes that the only path.
- r2) Git folders cannot be edited by people.

Exercise 2. Fill the gap (level 2)

Read the supplier API key from the `research` secret scope at run time.

```
1 token = dbutils._____(1).get(scope="research", key="_____(2)")
```

Exercise 3. Spot the bug (level 3)

A reviewer blocks this pull request. Tap the line that must never be committed, then rewrite the module.

```
1 import requests
2
3 SUPPLIER_URL = "https://supplier.example.org/api/deliveries"
4 SUPPLIER_TOKEN = "example-token-not-real-5678"
5
6
7 def fetch_deliveries() -> requests.Response:
8     """Fetch the delivery list from the supplier API."""
9     return requests.get(SUPPLIER_URL, headers={"Authorization": f"Bearer {SUPPLIER_TOKEN}"}, timeout=30)
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Put in order (level 4)

Order the steps for one change to the silver pipeline.

- Push and open a pull request
- Create a feature branch from main in your Git folder
- CI runs lint, type checks and tests
- Merge to main; automation deploys to production
- A colleague reviews and approves
- Change the code and add or update tests

Exercise 5. Architecture design (level 5)

Decide where each item lives.

Options: Secret scope, read at run time, Git repository, Job parameter.

1. The supplier API token
2. The functions that clean and merge silver
3. Which catalog a run writes to (research_dev or research_prod)
4. Your local .venv folder

Sources

- Git folders: per-user branches under /Workspace/Users/, trunk-based strategy, production Git folders updated only by automation, bundles recommended for deployment. (verified: <https://docs.databricks.com/aws/en/repos/ci-cd> (checked 2026-09-25))
- dbutils.secrets.get(scope, key) reads a secret; values are shown as [REDACTED]; redaction does not prevent deliberate transformations. (verified: <https://docs.databricks.com/aws/en/security/secrets/> (checked 2026-09-25))

A10-L2 Declarative Automation Bundles and parameterised jobs

Concept

A job built by clicking in the UI exists only in one workspace, with no history and no review. **Declarative Automation Bundles** describe jobs, their tasks, parameters and notifications as YAML files in your repository.

Databricks renamed them from Databricks Asset Bundles on 16 March 2026. The `databricks bundle` commands and `databricks.yml` did not change.

A bundle has one `databricks.yml` at its root:

```
1 bundle:
2   name: research_pipelines
3
4   variables:
5     catalog:
6       description: Catalog the pipelines write to.
7
8   resources:
9     jobs:
10      nat_silver_pipeline:
11        name: nat_silver_pipeline
12        parameters:
13          - name: catalog
14            default: ${var.catalog}
15        tasks:
16          - task_key: silver
17            notebook_task:
18              notebook_path: ./notebooks/silver_nat_transform.py
19
20   targets:
21     dev:
22       mode: development
23       default: true
24       variables:
25         catalog: research_dev
26     prod:
27       mode: production
28       variables:
29         catalog: research_prod
```

- **Targets** are environments. One code base deploys to each with different values.
- **Variables** hold what differs per target. They resolve from `--var`, then `BUNDLE_VAR` environment variables, then target mappings, then defaults.
- **Job parameters** reach notebook tasks, which read them with `dbutils.widgets.get("catalog")`. Dynamic references such as `{{job.run_id}}` fill in values at run time.
- `mode: development` prefixes job names with `[dev <your name>]` and pauses schedules, so your test deployment never runs on the production timetable.

The commands:

```
1 databricks bundle validate -t dev
2 databricks bundle deploy -t dev
```

```
3 databricks bundle run -t dev nat_silver_pipeline
```

Deploy to `prod` only from CI, as a service principal (lesson A10-L3).

Why it matters

A team manages its jobs in the UI. Someone changes the production silver job's notebook path to test a branch, meaning to change it back. They forget. For a week, production runs a half-finished branch. Nobody can say what changed or when, because the UI keeps no diff and nothing was reviewed. With the job in a bundle, the path is a line in `databricks.yml`. Changing it needs a pull request, CI validates it, and Git shows exactly what changed.

Worked example: weak code

```
1 resources:
2   jobs:
3     nat_silver_pipeline:
4       name: nat_silver_pipeline
5       tasks:
6         - task_key: silver
7           notebook_task:
8             notebook_path: /Workspace/Users/sam.jones@example.org/
9               silver_nat_transform
10            base_parameters:
11              catalog: research_prod
```

1. **Line 1.** No targets: there is no way to deploy the same job to dev.
2. **Line 8.** Points at one person's folder: unreviewed code, and it breaks when they leave.
3. **Line 10.** Production hardcoded: every deployment of this bundle writes to prod.

Worked example: strong code

```
1 bundle:
2   name: research_pipelines
3
4 variables:
5   catalog:
6     description: Catalog the pipelines write to.
7
8 resources:
9   jobs:
10    nat_silver_pipeline:
11      name: nat_silver_pipeline
12      parameters:
13        - name: catalog
14          1 default: ${var.catalog}
15      tasks:
16        - task_key: silver
```

```

17     notebook_task:
18     ②     notebook_path: ./notebooks/silver_nat_transform.py
19     email_notifications:
20     on_failure:
21         - data-oncall@example.org
22
23 targets:
24     dev:
25     ③     mode: development
26         default: true
27         variables:
28             catalog: research_dev
29     prod:
30         mode: production
31         variables:
32     ④     catalog: research_prod

```

1. **Line 14.** The job parameter takes its value from the target's variable.
2. **Line 18.** A path inside the bundle, deployed from the reviewed commit.
3. **Line 25.** Dev deployments get a [dev name] prefix and paused schedules.
4. **Line 32.** Production values live only in the prod target.

One reviewed file describes the job for every environment. The target decides where it writes.

Recap card

- **What are Declarative Automation Bundles?** Jobs and their settings as YAML in Git (formerly Databricks Asset Bundles). Same CLI: `databricks bundle`.
- **Targets and variables** Targets are environments. Variables hold what differs, such as the catalog, set per target.
- **What does mode: development do?** Prefixes resources with [dev <user>] and pauses schedules and triggers.
- **Read a job parameter in a notebook** `dbutils.widgets.get("name")`. Validate it before using it.

Practice

Exercise 1. Multiple choice (level 1)

What changed when Databricks Asset Bundles became Declarative Automation Bundles on 16 March 2026?

- a) The CLI command changed from `databricks bundle` to `databricks dab`.
- b) The name. The `databricks bundle` CLI and existing configuration keep working unchanged.

- c) databricks.yml was renamed automation.yml.
- d) Targets were removed in favour of workspaces.

Why?

- r1) Databricks describes the rename as non-breaking.
- r2) Bundles were replaced by a new product.

Exercise 2. Fill the gap (level 2)

Complete the targets so dev deployments are prefixed and paused, and each target sets its own catalog.

```

1 targets:
2   dev:
3     mode: ____ (1)
4     default: true
5     variables:
6       catalog: research_dev
7   prod:
8     mode: ____ (2)
9     variables:
10      ____ (3): research_prod

```

Exercise 3. Spot the bug (level 3)

A developer runs `databricks bundle deploy -t dev`, and the next morning production data has changed. Tap the line, then fix the job so the catalog comes from the target.

```

1 resources:
2   jobs:
3     nat_silver_pipeline:
4       name: nat_silver_pipeline
5       parameters:
6         - name: catalog
7           default: research_prod
8       tasks:
9         - task_key: silver
10          notebook_task:
11            notebook_path: ./notebooks/silver_nat_transform.py

```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 4)

The notebook for this job must read its `catalog` and `wave_id` job parameters. Write `read_job_parameters(widgets):` `widgets` behaves like `dbutils.widgets` (it has `.get(name)`). Return a dict with both values. Raise `ValueError` when `catalog` is not `research_dev` or `research_prod`, or when `wave_id` is empty.

Starter code:

```
1 def read_job_parameters(widgets: object) -> dict[str, str]:
2     ...
```

Exercise 5. Put in order (level 5)

Order the commands to try a bundle change in dev.

- databricks bundle deploy -t dev
- databricks bundle run -t dev nat_silver_pipeline
- Check the run log and QA results for the dev run
- databricks bundle validate -t dev

Sources

- Databricks Asset Bundles were renamed Declarative Automation Bundles on 16 March 2026; the change is non-breaking. (verified: <https://docs.databricks.com/aws/en/release-notes/dev-tools/bundles> (checked 2026-09-25))
- databricks.yml structure: bundle, variables, resources.jobs with parameters and tasks, targets with mode; variable resolution order. (search-only: <https://docs.databricks.com/aws/en/dev-tools/bundles/settings> and <https://docs.databricks.com/aws/en/dev-tools/bundles/variables>)
- Development mode prefixes resources with [dev \$ {workspace.current_user.short_name}] and pauses schedules and triggers. (verified: <https://docs.databricks.com/aws/en/dev-tools/bundles/deployment-modes> (checked 2026-09-25))
- Notebook tasks read job parameters with dbutils.widgets.get. (verified: <https://docs.databricks.com/aws/en/jobs/parameter-use> (checked 2026-09-25))

A10-L3 CI: lint, type-check, test and validate on every pull request

Concept

Continuous integration (CI) runs the same checks on every change, before a human reviews it. For a Databricks pipeline repository, run four checks on each pull request:

1. **Lint and format:** `ruff check .` and `ruff format --check .`
2. **Type-check:** `mypy`

3. **Unit tests:** `pytest`, with a local `SparkSession` (lesson A9-L2). Install Java in the CI runner first.
4. **Bundle validation:** `databricks bundle validate -t dev`. It catches YAML errors and broken references before anything is deployed.

On merge to `main`, a separate job deploys: `databricks bundle deploy -t prod`. Only CI deploys to production, and it does so as a **service principal**, never as a person.

Authenticating CI. Databricks recommends OAuth machine-to-machine authentication for a service principal. The CLI reads the workspace and credentials from environment variables, filled from the CI system's secret store:

```
1 env:
2   DATABRICKS_HOST: ${vars.DATABRICKS_HOST}
3   DATABRICKS_CLIENT_ID: ${secrets.DATABRICKS_CLIENT_ID}
4   DATABRICKS_CLIENT_SECRET: ${secrets.DATABRICKS_CLIENT_SECRET}
```

Workload identity federation goes a step further: with `DATABRICKS_AUTH_TYPE: github-oidc`, GitHub proves the workflow's identity, and no client secret is stored at all. Prefer it where your workspace supports it.

Two rules keep CI trustworthy:

- **CI must be able to fail.** A test step with `continue-on-error: true`, or a `|| true` after a command, reports success whatever happens.
- **Protect `main`.** Require the CI checks to pass before a pull request can merge. Otherwise CI only reports problems, and nothing stops them.

Install the CLI with the `databricks/setup-cli` action, and pin its version so a CLI release cannot change your pipeline's behaviour overnight.

Why it matters

A team adds CI but marks the test step `continue-on-error: true` "until the flaky tests are fixed". Six weeks later, every pull request shows a green tick. A change that breaks the silver MERGE key merges on a Friday. The failing test had run and failed, but the step's failure no longer failed the job. Production double-counts a wave over the weekend.

With the test step allowed to fail and `main` protected by required checks, the pull request could not have merged.

Worked example: weak code

```
1 on: push
2 jobs:
```

```

3  ci:
4    runs-on: ubuntu-latest
5    steps:
6      - uses: actions/checkout@v4
7      - run: pip install -r requirements.txt
8      - run: pytest
9  ❶  continue-on-error: true
10  ❷  - run: databricks bundle deploy -t prod
11    env:
12  ❸  DATABRICKS_TOKEN: dapi-REDACTED-EXAMPLE

```

1. **Line 9.** Failing tests no longer fail the job.
2. **Line 10.** Deploys every push, on every branch, straight to production.
3. **Line 12.** A credential written into the workflow file: committed to Git.

Worked example: strong code

```

1  name: CI
2  on:
3    pull_request:
4    push:
5      branches: [main]
6
7  jobs:
8    checks:
9      runs-on: ubuntu-24.04
10     steps:
11       - uses: actions/checkout@v4
12       - uses: actions/setup-python@v5
13         with:
14           python-version: "3.12"
15       - uses: actions/setup-java@v4
16         with:
17           distribution: temurin
18           java-version: "21"
19       - run: pip install -e . --group test
20  ❶  - run: ruff check . && ruff format --check .
21       - run: mypy
22       - run: pytest
23       - uses: databricks/setup-cli@main
24         with:
25           version: 0.270.0
26  ❷  - run: databricks bundle validate -t dev
27     env:
28       DATABRICKS_HOST: ${vars.DATABRICKS_HOST}
29       DATABRICKS_CLIENT_ID: ${secrets.DATABRICKS_CLIENT_ID}
30       DATABRICKS_CLIENT_SECRET: $
31       {{ secrets.DATABRICKS_CLIENT_SECRET }}
32
33  deploy:
34  ❸  if: github.event_name == 'push' && github.ref == 'refs/heads/main'
35  ❹  needs: checks

```

```

35     runs-on: ubuntu-24.04
36     steps:
37     - uses: actions/checkout@v4
38     - uses: databricks/setup-cli@main
39       with:
40         version: 0.270.0
41     5 - run: databricks bundle deploy -t prod
42       env:
43         DATABRICKS_HOST: ${vars.DATABRICKS_HOST}
44         DATABRICKS_CLIENT_ID: ${secrets.DATABRICKS_CLIENT_ID}
45         DATABRICKS_CLIENT_SECRET: $
         {{ secrets.DATABRICKS_CLIENT_SECRET }}

```

1. **Line 20.** Every check can fail the job. Nothing is marked continue-on-error.
2. **Line 26.** Validates the bundle on every pull request, before review.
3. **Line 33.** Deploys only after a merge to main.
4. **Line 34.** Deployment waits for every check to pass.
5. **Line 41.** Deploys as the service principal, with credentials from the CI secret store.

Every pull request is linted, typed, tested and validated. Only a merge to main deploys, as a service principal, with nothing secret in the file. The CLI version shown is an example: pin the version you tested.

Recap card

- **Four CI checks for a pipeline repository** ruff (lint and format), mypy, pytest with local Spark, databricks bundle validate.
- **Who deploys to production?** CI, after a merge to main, as a service principal.
- **OAuth M2M environment variables** DATABRICKS_HOST, DATABRICKS_CLIENT_ID, DATABRICKS_CLIENT_SECRET from the CI secret store.
- **Two ways CI lies** continue-on-error: true, and command || true. Both turn failures green.

Practice

Exercise 1. Multiple choice (level 1)

Which identity should CI use to run `databricks bundle deploy -t prod`?

- a) The personal access token of the most senior engineer
- b) A service principal, authenticated with OAuth M2M or workload identity federation
- c) Whoever opened the pull request

- d) An anonymous connection

Why?

- r1) A service principal is not tied to a person, and its access can be scoped to exactly what deployment needs.
- r2) Personal tokens cannot deploy bundles.

Exercise 2. Fill the gap (level 2)

Complete the environment for the bundle deploy step, using OAuth M2M with values from the CI secret store.

```
1 env:
2   DATABRICKS_HOST: ${vars.DATABRICKS_HOST}
3   DATABRICKS____(1): ${secrets.DATABRICKS_CLIENT_ID}
4   DATABRICKS____(2): ${secrets.DATABRICKS_CLIENT_SECRET}
```

Exercise 3. Spot the bug (level 3)

Every pull request shows a green tick, including ones that broke the MERGE. Tap the line, then fix the step.

```
1 steps:
2   - uses: actions/checkout@v4
3   - run: pip install -e . --group dev
4   - run: ruff check .
5   - run: pytest
6   continue-on-error: true
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Put in order (level 4)

Order the checks in the pull request job, cheapest and fastest first.

- ruff check . and ruff format --check .
- mypy
- databricks bundle validate -t dev
- pytest with a local SparkSession

Exercise 5. Architecture design (level 5)

Configure the repository and workflow. Choose each setting.

Options: On every pull request, and on push to main, Nightly, When someone remembers.

1. When do lint, types, tests and validation run?
2. When does the prod deployment run?
3. Branch protection on main

4. CI authentication

Sources

- Databricks recommends OAuth M2M for service principals in CI; environment variables `DATABRICKS_HOST`, `DATABRICKS_CLIENT_ID` and `DATABRICKS_CLIENT_SECRET`; github-oidc workload identity federation needs no client secret. (verified: <https://docs.databricks.com/aws/en/dev-tools/auth/oauth-m2m> and <https://docs.databricks.com/aws/en/dev-tools/bundles/authentication> (checked 2026-09-25))
- `databricks/setup-cli` installs the Databricks CLI in GitHub Actions and accepts a version input. (search-only: <https://github.com/databricks/setup-cli>)
- CLI version 0.270.0 in the example is illustrative, not a verified current release. (unverified: Pin the version you have tested; check <https://github.com/databricks/cli/releases>.)

Track A best-practice checklist

- ☐ Every function has type hints, a docstring and one job.
- ☐ No SparkSession is created at import time. Functions receive spark as a parameter.
- ☐ Functions are imported as F: `from pyspark.sql import functions as F`.
- ☐ Table names, column maps and thresholds come from config or registry tables.
- ☐ Every table states its grain in a comment or table property.
- ☐ Bronze appends. Silver deduplicates the source, then MERGEs. Gold aggregates by every grain column.
- ☐ Every run writes one record to the run log table with an explicit schema.
- ☐ CRITICAL checks raise. WARN checks log and continue.
- ☐ Every division has a guard for zero and null.
- ☐ Counts come from one agg pass, not repeated count() calls.
- ☐ Incremental loads compare source and target on the same keys.
- ☐ Row filters use a SQL UDF with SET ROW FILTER. Grants go to groups.
- ☐ Tests run on a local SparkSession in CI before every merge.

Track B: Python for data science

Write Python and run analyses that are correct, reproducible and easy to review.

B1 Project setup

Roadmap: Prerequisite for all months. pyproject.toml, virtual environments, ruff, mypy and folder layout.

B1-L1 One pyproject.toml and a src layout

Concept

A data science project starts as a notebook and a CSV. By month three it has helper modules, tests, a model, and three people asking how to run it. The fix is structure from day one.

pyproject.toml is the one configuration file for a modern Python project. It holds:

- [project]: name, version, requires-python, and runtime dependencies.
- [dependency-groups] (PEP 735): named groups for development only, such as test or lint. They never ship with the package. Install one with `pip install --group test` (pip 25.1 and later).
- [tool.*] tables for tools: [tool.ruff], [tool.mypy], [tool.pytest.ini_options].

One file replaces setup.py, requirements.txt, setup.cfg, .flake8 and pytest.ini.

The src layout puts importable code under `src/`:

```
1 reach_model/
2 └─ pyproject.toml
3 └─ src/reach_model/      # importable package
4   └─ __init__.py
5   └─ features.py
6   └─ train.py
7 └─ tests/                # pytest, imports reach_model
8 └─ notebooks/           # exploration; imports reach_model, holds no
    logic
9 └─ data/                # small samples only; real data stays out of
    Git
```

With `src/`, Python can only import your package after you install it (`pip install -e .`). Tests then run against the installed package, exactly as a user or a job would. Without it, tests silently import whatever sits in the working directory, and "works on my machine" follows.

Notebooks explore and present. Functions that matter move into `src/`, where tests and type checks reach them. A notebook that defines the feature pipeline inline cannot be tested or reused by the training job.

Pin `requires-python` to the versions you test, for example `">=3.11"`, so pip refuses an interpreter you never tried.

Why it matters

A team trains a churn model from `train.ipynb`, which defines `build_features()` in a cell. The scoring job copies the function into `score.py`. A month later someone fixes a leakage bug in the notebook version only. Training and scoring now compute features differently, and live accuracy drops 9 points while offline metrics look fine.

With `build_features()` in `src/churn/features.py`, imported by both the notebook and the job, one fix reaches both, and one test protects it.

Worked example: weak code

1	churn/	
❶	├─ train.ipynb	# defines build_features() in cell 7
❷	├─ score.py	# copy of build_features(), edited later
4	├─ utils.py	
❸	├─ requirements.txt	# pandas, scikit-learn (no versions)
6	├─ setup.cfg	# flake8 settings
❹	└─ data.csv	# 2 GB client extract, committed

1. **Line 2.** Logic trapped in a notebook: untestable, and copied instead of imported.
2. **Line 3.** Two versions of one function drift apart.
3. **Line 5.** No pins: every install resolves different versions.
4. **Line 7.** Client data in Git: a size problem and a data-protection problem.

Worked example: strong code

1	churn/	
❶	├─ pyproject.toml	# project, dependency groups, ruff, mypy, pytest
3	├─ src/churn/	
4	├─ └─ __init__.py	
❷	├─ └─ features.py	# build_features(), used by training and scoring
6	├─ └─ train.py	
7	├─ └─ score.py	
8	├─ tests/	
❸	├─ └─ test_features.py	
10	├─ notebooks/	
11	├─ └─ 01_explore.ipynb	# imports churn.features

```
12 └─ data/
    4 └─ sample_200_rows.csv
```

1. **Line 2.** One file configures the project and every tool.
2. **Line 5.** One definition, imported everywhere it is needed.
3. **Line 9.** Tests import the installed package, as the jobs do.
4. **Line 13.** A small, anonymised sample for tests. Real data stays in governed storage.

Logic lives once, in an installable package. Notebooks, jobs and tests import it.

Recap card

- **What goes in `pyproject.toml`?** `[project]` with dependencies, `[dependency-groups]` for development tools, and `[tool.*]` settings for `ruff`, `mypy` and `pytest`.
- **Install a dependency group** `pip install --group dev (pip 25.1 and later)`.
- **Why the `src` layout?** Only the installed package is importable, so tests run against it exactly as jobs do.
- **Where does shared logic live?** In the package under `src/`, imported by notebooks, jobs and tests. Never copied.

Practice

Exercise 1. Multiple choice (level 1)

Where do you declare `pytest` and `ruff` so they install for developers but never ship with the package?

- a) `[project]` dependencies
- b) `[dependency-groups]`
- c) `[tool.pytest.ini_options]`
- d) A comment in `README.md`

Why?

- r1) Dependency groups (PEP 735) are named lists that are not part of the built distribution.
- r2) Dependency groups are only for test tools.

Exercise 2. Fill the gap (level 2)

Complete the `pyproject.toml` so runtime needs and developer tools are separate.

```
1 [project]
```

```

2 name = "churn"
3 requires-python = ">=3.11"
4 ____ (1) = ["pandas>=3.0,<3.1", "scikit-learn>=1.8,<1.10"]
5
6 [____ (2)]
7 dev = ["pytest>=9", "ruff>=0.16", "mypy>=2.3"]

```

Exercise 3. Write the function (level 3)

Write `unpinned(pyproject_text)`. Parse the TOML with the standard library and return, sorted, every entry in `[project].dependencies` that has no version constraint at all (no `=`, `<`, `>`, `~` or `!`). Return an empty list when there is no dependencies key.

Starter code:

```

1 import tomllib
2
3
4 def unpinned(pyproject_text: str) -> list[str]:
5     ...

```

Exercise 4. Multiple choice (level 4)

With a `src` layout, a colleague runs `pytest` straight after cloning and gets `ModuleNotFoundError: No module named 'churn'`. What is the right fix?

- a) Add `sys.path.insert(0, 'src')` at the top of every test file.
- b) Run `pip install -e .` in the project's virtual environment.
- c) Move the package out of `src/` so imports work from the root.
- d) Copy the package into `tests/`.

Why?

- r1) The `src` layout is designed so that only the installed package is importable, which makes tests use it like any user would.
- r2) `pytest` cannot find packages in folders.

Exercise 5. Architecture design (level 5)

Place each piece of a modelling project.

Options: `src/<package>/`, A notebook cell, `tests/`, Outside the repository.

1. `build_features()`, used by training and by the nightly scoring job
2. Plots comparing feature distributions for a one-off question
3. A check that `build_features()` gives the right output for 5 hand-made rows
4. The full client data extract

Sources

- PEP 735 defines [dependency-groups]; pip 25.1 (April 2025) added --group. (search-only: <https://peps.python.org/pep-0735/> and <https://ichard26.github.io/blog/2025/04/whats-new-in-pip-25.1/>)
- pip 26.2.1 lists --group in `pip install --help`; tomllib is in the standard library from Python 3.11. (verified: Checked in this repository's build environment.)

B1-L2 One virtual environment per project, and python -m pip

Concept

A **virtual environment** is a folder with its own Python interpreter link and its own `site-packages`. Packages you install there do not touch the system Python or any other project.

```
1 python -m venv .venv
2 source .venv/bin/activate      # Windows: .venv\Scripts\activate
3 python -m pip install --upgrade pip
4 python -m pip install -e . --group dev
```

Rules that prevent most "it worked yesterday" problems:

- **One environment per project.** Two projects that share an environment share their version conflicts. Upgrading pandas for one breaks the other.
- **Never install into the system Python.** Operating systems and other tools depend on it.
- **Use** `python -m pip`, **not bare** `pip`. Bare `pip` runs whichever `pip` comes first on your PATH, which may belong to a different interpreter. `python -m pip` always installs into the interpreter you are about to run.
- **Keep** `.venv/` **out of Git.** It is large, machine-specific and rebuildable. Commit `pyproject.toml` and a lock file instead (lesson B6-L3).

In code, `sys.prefix != sys.base_prefix` is True inside a virtual environment. A script that must not run against the system Python can check it and stop.

Sometimes one project needs **two** environments, because two dependencies cannot live together. A real case from this course's own build: PySpark 4.0.4 fails to import its pandas integration with pandas 3.x, so the Spark test harness runs in its own environment pinned to pandas 2.x, while the rest uses pandas 3. Separate environments, each fully pinned, beat one environment held together by guesswork.

Tools such as `uv` create and sync environments much faster, and `conda` manages non-Python libraries. The principles are the same.

Why it matters

An analyst runs `pip install statsmodels` in a terminal, then `python analysis.py` in the same terminal, and gets `ModuleNotFoundError`. `pip` belonged to one Python, and `python` to another. They try `sudo pip install`, which upgrades NumPy in the system Python and breaks a command-line tool the operating system uses. Two hours are gone.

`python -m venv .venv`, activation, and `python -m pip install` make the interpreter that installs and the interpreter that runs the same one.

Worked example: weak code

```
❶ $ pip install pandas scikit-learn
❷ $ sudo pip install --upgrade numpy
3 $ python train.py
❸ ModuleNotFoundError: No module named 'sklearn'
❹ $ git add .venv/ data/ train.py
```

1. **Line 1.** Bare pip, no environment: installs into whichever Python owns this pip.
2. **Line 2.** Changes the system Python that the operating system relies on.
3. **Line 4.** The python that runs is not the one pip installed into.
4. **Line 5.** Commits a machine-specific environment folder and data.

Worked example: strong code

```
❶ $ python -m venv .venv
2 $ source .venv/bin/activate
3 (.venv) $ python -m pip install --upgrade pip
❷ (.venv) $ python -m pip install -e . --group dev
5 (.venv) $ python -c "import sys; print(sys.prefix != sys.base_prefix)"
❸ True
7 (.venv) $ python train.py
```

1. **Line 1.** A project-local environment, listed in `.gitignore`.
2. **Line 4.** The same interpreter installs and runs. Dependencies come from `pyproject.toml`.
3. **Line 6.** Proof that the active interpreter is the virtual environment.

One environment per project, created from `pyproject.toml`, with `python -m pip` so installing and running use the same interpreter.

Recap card

- **Create and use a project environment** `python -m venv .venv`, activate it, then `python -m pip install -e . --group dev`.

- **Why python -m pip?** It installs into the interpreter that python runs. Bare pip may belong to another Python.
- **Detect a virtual environment in code** `sys.prefix != sys.base_prefix`.
- **Two dependencies that cannot coexist** Two pinned environments and two dependency groups, each tested in CI.

Practice

Exercise 1. Multiple choice (level 1)

Why prefer `python -m pip install x` over `pip install x`?

- a) It is faster.
- b) It installs into the interpreter that `python` runs, whatever is first on PATH.
- c) It installs packages globally for all users.
- d) It skips dependency resolution.

Why?

- r1) `python -m pip` runs pip as a module of that exact interpreter.
- r2) Bare pip is deprecated.

Exercise 2. Predict the output (level 2)

You run this with `.venv/bin/python`, a virtual environment created by `python -m venv .venv`. What does it print?

```
1 import sys
2
3 print(sys.prefix != sys.base_prefix)
```

Exercise 3. Fill the gap (level 2)

Create the environment, then install the project in editable mode with its dev group.

```
1 python -m ____ (1) .venv
2 source .venv/bin/activate
3 python -m pip install ____ (2) . --group dev
```

Exercise 4. Write the function (level 3)

Write `require_virtualenv()`. It raises `RuntimeError` with the message "activate the project's virtual environment first" when the running interpreter is not inside a virtual environment, and returns `None` otherwise. Read `sys.prefix` and `sys.base_prefix` when the function is called, not at import time.

Starter code:

```
1 import sys
2
3
4 def require_virtualenv() -> None:
5     ...
```

Exercise 5. Architecture design (level 5)

Your project needs pandas 3.0 for analysis, and a Spark test harness on PySpark 4.0.4, which fails to import its pandas integration with pandas 3.x. Choose the set-up.

Options: Two environments: dev (pandas 3) and spark (pandas 2), each pinned, One environment with the latest of everything, One environment, everything downgraded to pandas 2.

1. Environments
2. Where the difference is recorded
3. CI

Sources

- `sys.prefix` differs from `sys.base_prefix` inside a `venv`. (verified: Python `venv` documentation; checked in this repository's virtual environments.)
- PySpark 4.0.4 fails to import `assertDataFrameEqual` with pandas 3.x. (verified: Reproduced in this repository (VERSIONS.md).)

B1-L3 ruff and mypy in daily work

Concept

Two tools catch most mistakes before a test runs.

ruff lints and formats. On a ten-line sample it reported three real problems:

```
1 m.py:1:1: I001 [*] Import block is un-sorted or un-formatted
2 m.py:1:8: F401 [*] `os` imported but unused
3 m.py:9:30: B905 `zip()` without an explicit `strict=` parameter
```

B905 matters for data work: `zip(flags, weights)` silently stops at the shorter list, so a missing weight drops a respondent without an error. `zip(..., strict=True)` raises instead.

mypy checks types. On the same file, strict mode found:

```
1 m.py:14: error: Unsupported operand types for * ("None" and "int")
   [operator]
2 m.py:17: error: Function is missing a type annotation [no-untyped-def]
```

Line 14 multiplies a `float | None` by 100. When the total weight is zero, the function returns `None` and the report crashes. `mypy` found the zero-weight bug without any test data.

Daily habits:

- Configure both once in `pyproject.toml` (`[tool.ruff]` with the rule sets you want, `[tool.mypy]` with `strict = true`).
- Run `ruff check --fix .` and `ruff format .` before each commit. `--fix` applies only safe fixes.
- Run `mypy` on `src/` and `tests/`. Fix errors by handling the case, not by adding `# type: ignore`.
- Run the same commands in CI (lesson A10-L3), so the rules are not optional.

Libraries without type information show `import-untyped` or `missing-stub` notes. Install the stub package if one exists, or add a targeted `[[tool.mypy.overrides]]` with `ignore_missing_imports = true` for that module only.

Why it matters

A reporting script formats a weighted share as `f"{share * 100:.1f}%"`. For eleven months every subgroup has respondents. Then a new region with no fieldwork yet appears in the sample frame, its total weight is zero, `weighted_share` returns `None`, and the monthly client pack fails to build on deadline day.

`mypy` flags `None * 100` the first time it runs on the file. The fix is two lines, written months before the failure.

Worked example: weak code

```
1 import os
2 import pandas as pd
3
4
5 def weighted_share(flags, weights):
6     total = sum(weights)
7     if total == 0:
8         return None
9
10    return sum(w for f, w in zip(flags, weights) if f == 1) / total
11
12 def report(flags, weights):
13     return f"{weighted_share(flags, weights) * 100:.1f}%"
```

1. **Line 1.** F401: unused import.

2. **Line 5.** No hints: mypy cannot see that the function can return None.
3. **Line 9.** B905: zip silently truncates to the shorter list.
4. **Line 13.** Crashes with TypeError when the share is None.

Worked example: strong code

```
❶ def weighted_share(flags: list[int], weights: list[float]) -> float |  
    None:  
    2     """Weighted share of flagged rows, or None when the total weight is  
        0."""  
    3     total = sum(weights)  
    4     if total == 0:  
    5         return None  
❷     return sum(w for f, w in zip(flags, weights, strict=True) if f ==  
        1) / total  
    7  
    8  
    9 def report(flags: list[int], weights: list[float]) -> str:  
    10     """Share as a percentage string, or 'n/a' when undefined."""  
    11     share = weighted_share(flags, weights)  
❸     if share is None:  
    13         return "n/a"  
    14     return f"{share * 100:.1f}%"
```

1. **Line 1.** The return type admits None, so mypy makes every caller handle it.
2. **Line 6.** Unequal lengths raise ValueError instead of dropping rows.
3. **Line 12.** The case mypy flagged, handled on purpose.

ruff and mypy found an unused import, a silent truncation and a crash on zero weight, before anyone wrote a test.

Recap card

- **Why does ruff flag zip() without strict= (B905)?** zip silently truncates to the shorter input. strict=True raises ValueError instead.
- **mypy says None * int is unsupported. The fix?** Handle the None case in code. Never # type: ignore it away.
- **Daily commands before a commit** ruff check --fix ., ruff format ., mypy.
- **Missing stubs for one library** Install types-<lib> if it exists, otherwise a per-module mypy override. Never global.

Practice

Exercise 1. Predict the output (level 1)

What does this print?

```
1 flags = [1, 0, 1]  
2 weights = [2.0, 1.0]
```

```
3 print(list(zip(flags, weights)))
```

Exercise 2. Multiple choice (level 2)

mypy reports `Unsupported operand types for * ("None" and "int") on share * 100`. What is the right fix?

- a) Add `# type: ignore` to the line.
- b) Handle the `None` case before the multiplication.
- c) Change the return type to `float` and keep returning `None`.
- d) Turn off strict mode.

Why?

- r1) mypy found a real path where the value is `None`. Handling it removes the crash.
- r2) mypy errors are style issues.

Exercise 3. Fill the gap (level 2)

Complete the pre-commit routine: apply safe lint fixes, format, then type-check.

```
1 ruff check ____ (1) .
2 ruff ____ (2) .
3 ____ (3)
```

Exercise 4. Refactor (level 4)

Make this pass `ruff` (no B905) and `mypy --strict`: add type hints, keep the `None` result for zero weight, make unequal lengths raise, and make `report` return `"n/a"` instead of crashing.

Starter code:

```
1 def weighted_share(flags, weights):
2     total = sum(weights)
3     if total == 0:
4         return None
5     return sum(w for f, w in zip(flags, weights) if f == 1) / total
6
7
8 def report(flags, weights):
9     return f"{weighted_share(flags, weights) * 100:.1f}%"
```

Exercise 5. Architecture design (level 5)

mypy reports `Skipping analyzing "shapely": module is installed, but missing library stubs or py.typed marker`. Choose the response for each case.

Options: Install the stubs in the dev group, Add `[[tool.mypy.overrides]]` with `ignore_missing_imports` for that module, Set `ignore_missing_imports = true` globally.

1. A stub package, `types-shapely`, exists on PyPI
2. No stubs exist for an internal library
3. A colleague proposes silencing all such messages for every library

Sources

- ruff reports `I001`, `F401` and `B905` and mypy `--strict` reports `[operator]` and `[no-untyped-def]` on the sample shown. (verified: ruff 0.16.9 and mypy 2.3.1 run in this repository.)
- `zip(strict=True)` raises `ValueError` on unequal lengths (Python 3.10+). (verified: Python standard library; run in Python 3.11.)

B2 Code design

Roadmap: Prerequisite for all months. Pure functions, dataclasses, when to use classes, configuration objects.

B2-L1 Pure functions, side effects at the edges

Concept

A **pure function** returns a value that depends only on its arguments, and changes nothing outside itself: no file writes, no global variables, no mutation of the inputs, no printing, no clock or random calls it does not receive.

```
1 def weighted_mean(values: list[float], weights: list[float]) -> float:
2     return sum(v * w for v, w in zip(values, weights, strict=True)) /
    sum(weights)
```

Pure functions are easy to test (inputs in, answer out), safe to reuse, and safe to run twice. Most analysis logic can be pure: cleaning, feature building, statistics, scoring.

Some work is not pure by nature: reading files, calling APIs, writing tables, logging, reading the clock. Push that work to the **edges** of the program, into a thin layer that gathers inputs, calls the pure core, and writes results.

```
1 read (impure) -> transform, compute (pure) -> write (impure)
```

Three side effects hide easily in data code:

- **Mutating an argument.** `df["x"] = ...` or `rows.append(...)` inside a function changes the caller's object. The caller did not ask for that. Return a new object instead: `df.assign(x=...)`, or build a new list.
- **Reading globals.** A function that reads a module-level `THRESHOLD` or `TODAY` gives different answers after someone changes it. Pass it in.
- **Hidden time and randomness.** `datetime.now()` or `np.random.rand()` inside a function makes its output untestable. Pass a date, or a `numpy.random.Generator`.

In pandas 3, Copy-on-Write is always on, so many operations already return new objects. Assigning a column with `df[col] = ...` still changes `df` itself, which is exactly the mutation to keep out of shared functions.

Why it matters

A feature function adds `age_group` to the DataFrame it receives, in place. The training script calls it, then passes the same DataFrame to a stratified split on `age_group`. Later, someone reorders the calls so the split runs first. The split fails with a `KeyError` on one laptop and works on another, depending on which notebook cells ran before.

A pure `add_age_group(df) -> DataFrame` that returns a new frame gives the same result in every order, and its test needs three rows.

Worked example: weak code

```
1 import datetime
2
3 THRESHOLD = 0.5
4
5
6 def score_respondents(df):
7     df["age"] = datetime.date.today().year - df["birth_year"]
8     df["high_reach"] = df["reach"] > THRESHOLD
9     df.to_csv("scored.csv", index=False)
10    print("scored", len(df))
11    return df
```

1. **Line 7.** Mutates the caller's frame and reads the clock: the answer changes on 1 January.
2. **Line 8.** Reads a module global: change `THRESHOLD` anywhere and every result changes.
3. **Line 9.** Writes a file from inside the computation.
4. **Line 10.** Output in the middle of logic, not in the program's edge.

Worked example: strong code

```
1 import pandas as pd
2
3
4 def score_respondents(df: pd.DataFrame, as_of_year: int, threshold:
5     float) -> pd.DataFrame:
6     """Add age and high_reach. Returns a new frame; the input is
7     unchanged."""
8     return df.assign(
9         age=as_of_year - df["birth_year"],
10        high_reach=df["reach"] > threshold,
11    )
12
13 def main(input_path: str, output_path: str, as_of_year: int, threshold:
14     float) -> None:
15     """The edge: read, call the pure core, write."""
```

```

14     scored = score_respondents(pd.read_csv(input_path), as_of_year,
    threshold)
15     scored.to_csv(output_path, index=False)

```

1. **Line 4.** Everything the result depends on is a parameter.
2. **Line 6.** assign returns a new frame. The caller's data is not touched.
3. **Line 12.** Reading, writing and choosing the year happen here, at the edge.

The pure core is tested with three rows and a fixed year. The edge is a few lines that are easy to read.

Recap card

- **What makes a function pure?** Its result depends only on its arguments, and it changes nothing outside itself.
- **Three hidden side effects in data code** Mutating arguments, reading globals, reading the clock or randomness inside the function.
- **Return a new DataFrame with an extra column** df.assign(col=...). df["col"] = ... changes the caller's frame.
- **Program shape** Read at the edge, compute in a pure core, write at the edge.

Practice

Exercise 1. Multiple choice (level 1)

Which function is pure?

- a) `def add(rows, x): rows.append(x); return rows`
- b) `def share(n, total): return n / total if total else None`
- c) `def stamp(df): return df.assign(ts=pd.Timestamp.now())`
- d) `def save(df, path): df.to_csv(path)`

Why?

- r1) Its result depends only on its arguments, and it changes nothing outside itself.
- r2) It is the shortest.

Exercise 2. Predict the output (level 2)

What does this print on pandas 3?

```

1 import pandas as pd
2
3
4 def add_flag(df):
5     df["flag"] = df["x"] > 1

```



```

6     return df
7
8
9 original = pd.DataFrame({"x": [1, 2]})
10 result = add_flag(original)
11 print(list(original.columns))

```

Exercise 3. Refactor (level 3)

Make `add_age` pure: take the reference year as a parameter, return a new DataFrame, and leave the input unchanged. Keep the name `add_age` and the new column `age`.

Starter code:

```

1 import datetime
2
3
4 def add_age(df):
5     df["age"] = datetime.date.today().year - df["birth_year"]
6     return df

```

Exercise 4. Spot the bug (level 4)

The second call returns a list containing the first call's flags. Tap the line, then fix the function without changing its behaviour otherwise.

```

1 def flag_low_bases(bases: dict[str, float], minimum: float, flags:
    list[str] = []) -> list[str]:
2     """Names of subgroups whose base is below minimum."""
3     for name, base in bases.items():
4         if base < minimum:
5             flags.append(name)
6     return flags

```

Write down the line number of the bug. Then write the fix.

Exercise 5. Architecture design (level 5)

Sort each step of a weekly scoring script into the pure core or the edge.

Options: Edge, Pure core.

1. Read the respondent extract from storage
2. Compute features from the extract and a reference date
3. Score respondents with a model object passed in
4. Write scores to the output table and log the row count

Sources

- In pandas 3 Copy-on-Write is always on, but assigning a column changes the DataFrame it is called on. (verified: Run on pandas 3.0.6 and 3.0.2 in this repository (predict-output exercise).)

B2-L2 Dataclasses as configuration objects

Concept

Analysis code needs settings: a minimum base size, a weight column, a list of waves. Passing them around as a loose `dict` has three costs. A typo in a key (`"min_bse"`) fails late or not at all. Nothing states which keys exist. Any function can change a value half way through a run.

A **dataclass** (Python 3.7 and later) turns the settings into a named, typed object.

```
1 from dataclasses import dataclass, field
2
3
4 @dataclass(frozen=True)
5 class ReportConfig:
6     weight_col: str
7     min_base: int = 50
8     waves: tuple[str, ...] = ()
9
10     def __post_init__(self) -> None:
11         if self.min_base <= 0:
12             raise ValueError(f"min_base must be positive, got {self.min_base}")
```

What you get:

- **Named fields.** `cfg.min_bse` raises `AttributeError` at once. Your editor and mypy see the fields.
- `frozen=True`. Assigning a field raises `dataclasses.FrozenInstanceError`. Settings stay fixed for the run. Use `dataclasses.replace(cfg, min_base=30)` to build a changed copy.
- `__post_init__`. Runs after the generated `__init__`. Put checks here so a bad config fails at load time, not in hour three.

Two traps:

- **Type hints are not checked at runtime.** `ReportConfig(weight_col="w", min_base="30")` builds without error. Convert or check types in `__post_init__`, or use `pydantic` (lesson B5-L2) when input comes from files.

- **Mutable defaults.** `waves: list[str] = []` raises `ValueError` when the class is defined. Use `field(default_factory=list)`, or an immutable tuple.

Load the config once at the edge, from YAML, JSON or arguments, and pass the object to the pure core.

Why it matters

A weighting script reads `config["min_base"]` from a dict built in a notebook cell. An analyst edits the cell to `config["minbase"] = 30` for a test and forgets to revert. The script keeps using the old default of 50, so 14 subgroups the client asked to see are suppressed in the published tables. Nobody notices for two waves.

With a frozen dataclass, `ReportConfig(minbase=30)` raises `TypeError: unexpected keyword argument` at the first run.

Worked example: weak code

```

1 CONFIG = {"weight_col": "weight", "min_base": 50, "waves": []}
2
3
4 def suppress_small(counts, config=CONFIG):
5     config["min_base"] = config.get("minbase", 50)
6     return {k: v for k, v in counts.items() if v >= config["min_base"]}

```

1. **Line 1.** A module-level dict: no fixed keys, no types, shared and mutable.
2. **Line 4.** A mutable default that every call shares.
3. **Line 5.** A typo in the key silently falls back to 50, and the write changes the shared dict.

Worked example: strong code

```

1 from dataclasses import dataclass
2
3
4 @dataclass(frozen=True)
5 class ReportConfig:
6     """Settings for one reporting run."""
7
8     weight_col: str
9     min_base: int = 50
10
11     def __post_init__(self) -> None:
12         if not isinstance(self.min_base, int) or self.min_base <= 0:
13             raise ValueError(f"min_base must be a positive int, got {self.min_base!r}")
14
15

```

```

3 def suppress_small(counts: dict[str, int], config: ReportConfig) ->
    dict[str, int]:
17     """Keep subgroups whose base reaches config.min_base."""
18     return {k: v for k, v in counts.items() if v >= config.min_base}

```

1. **Line 4.** Fixed fields, and no assignment after construction.
2. **Line 12.** Validation at load time. Hints alone are not checked at runtime.
3. **Line 16.** The function receives the config and reads it. It never changes it.

A frozen dataclass names every setting, rejects unknown ones, and checks values once, when you build it.

Recap card

- **Why a frozen dataclass for settings?** Named, typed fields, unknown keys rejected, and no assignment after construction.
- **Change one setting on a frozen config** `dataclasses.replace(cfg, min_base=30)` returns a new object.
- **Are dataclass type hints checked at runtime?** No. Validate or convert in **post_init**, or use `pydantic` for external input.
- **Empty list default on a dataclass field** `field(default_factory=list)`. A plain `[]` raises `ValueError` at class definition.

Practice

Exercise 1. Multiple choice (level 1)

`cfg` is an instance of a `@dataclass(frozen=True)` class with a field `min_base`. What does `cfg.min_base = 30` do?

- a) Sets `min_base` to 30.
- b) Raises `dataclasses.FrozenInstanceError`.
- c) Creates a copy with `min_base` 30.
- d) Does nothing, silently.

Why?

- r1) `frozen=True` makes the generated **setattr** raise on every field assignment.
- r2) Integer fields are read-only in every dataclass.

Exercise 2. Predict the output (level 2)

What does this print?

```

1 from dataclasses import dataclass
2
3

```

```

4 @dataclass
5 class Settings:
6     min_base: int = 50
7
8
9 print(Settings("30"))

```

Exercise 3. Fill the gap (level 2)

Give `waves` an empty list as its default, safely.

```

1 from dataclasses import dataclass, field
2
3
4 @dataclass
5 class Run:
6     waves: list[str] = field(____(1)=list)

```

Exercise 4. Write the function (level 3)

Complete `ReportConfig` and write `load_config(raw)`.

- `ReportConfig` is a frozen dataclass with `weight_col: str` and `min_base: int = 50`.
- In `__post_init__`, raise `ValueError` if `min_base` is not an `int` or is not positive.
- `load_config(raw: dict[str, object]) -> ReportConfig` builds the config from a dict. An unknown key must raise `TypeError`, a missing `weight_col` must raise `TypeError`.

Starter code:

```

1 from dataclasses import dataclass
2
3
4 class ReportConfig:
5     ...
6
7
8 def load_config(raw):
9     ...

```

Exercise 5. Architecture design (level 4)

A weighting pipeline reads its settings from `weighting.yaml`. Choose how the settings flow through the code.

Options: A frozen dataclass with typed fields, The dict returned by `yaml.safe_load`, Module-level constants set from the file at import.

1. Shape of the settings in code
2. When values are checked

3. How functions get the settings

Sources

- Assigning a field of a frozen dataclass raises `dataclasses.FrozenInstanceError`, a subclass of `AttributeError`. (verified: Python dataclasses documentation; run on Python 3.11 and 3.14 in this repository.)
- A list default on a dataclass field raises `ValueError` at class definition. (verified: Python dataclasses documentation; run on Python 3.11 in this repository.)
- Dataclasses do not check or convert field types at runtime. (verified: Python dataclasses documentation; predict-output exercise run in this repository.)

B2-L3 When a class earns its place

Concept

Python does not require classes. Many data science classes are functions in disguise: an `__init__` that stores arguments, and one method that uses them.

```
1 class Normaliser:
2     def __init__(self, col):
3         self.col = col
4
5     def run(self, df):
6         return df.assign(**{self.col: df[self.col] / df[self.col].sum()})
```

Every caller writes `Normaliser("spend").run(df)`. A function, `normalise(df, col)`, says the same thing in one call, with nothing to construct.

A class earns its place when at least one of these holds:

- **State that outlives one call.** A weighting model learns targets in `fit` and applies them in `transform`. scikit-learn's `StandardScaler` stores `mean_` and `scale_` between the two calls.
- **An invariant to protect.** A `Weights` object guarantees values are positive and sum to the sample size. Code that builds one through its constructor cannot break that rule.
- **Several operations on the same data.** A `SurveyDesign` with `mean()`, `ci()` and `subgroup()` that all need weights, strata and clusters.

- **A shared interface.** Several classes with the same method, such as `load()`, so the caller swaps one for another.

Settings with no behaviour belong in a dataclass (lesson B2-L2). A single step belongs in a function (lesson B2-L1).

Test a design by asking: if you delete the class and keep its methods as functions, what breaks? If nothing, keep the functions.

Why it matters

A team wraps each cleaning step in a class: `DropDuplicates`, `FixDates`, `RenameColumns`, each with `__init__` and `run`. A new analyst adds a `self.seen_ids` set to `DropDuplicates` to speed up a second call. The pipeline object is reused across two tracker waves, so wave 2 silently drops every respondent who also took part in wave 1: 38% of the sample.

Pure functions have no `self` to keep state state between calls.

Worked example: weak code

```

1 class AgeBander:
2     def __init__(self, bins, labels):
3         self.bins = bins
4         self.labels = labels
5
6     def run(self, df):
7         return df.assign(age_band=pd.cut(df["age"], self.bins,
8             labels=self.labels))
9
10 banded = AgeBander([0, 34, 54, 120], ["18-34", "35-54", "55+"]).run(df)

```

1. **Line 1.** No state outlives a call, no invariant, one operation: a function in disguise.
2. **Line 6.** A single method named `run` is the usual sign.
3. **Line 10.** Every caller builds an object only to call one method.

Worked example: strong code

```

1 import pandas as pd
2
3
4 def add_age_band(df: pd.DataFrame, bins: list[int], labels: list[str]) ->
    pd.DataFrame:
5     """Return df with age_band cut from age."""
6     return df.assign(age_band=pd.cut(df["age"], bins, labels=labels))
7
8
9 class RakingWeights:

```

```

10     """Learn raking targets on one wave, apply them to later waves."""
11
12     def __init__(self, targets: dict[str, dict[str, float]]) -> None:
13         self.targets = targets
14         self.factors_: dict[str, float] | None = None
15
16     def fit(self, df: pd.DataFrame) -> "RakingWeights":
17         ...
18         return self
19
20     def transform(self, df: pd.DataFrame) -> pd.Series:
21         if self.factors_ is None:
22             raise RuntimeError("call fit before transform")
23         ...

```

1. **Line 4.** One step, no state: a function.
2. **Line 14.** State that outlives one call, learned in fit and used in transform.
3. **Line 22.** The class protects an order of operations a function could not.

Functions by default. A class when state lives across calls, an invariant needs guarding, or several methods share the same data.

Recap card

- **A class with init and one run method** A function in disguise. Replace it with a function.
- **Four reasons a class earns its place** State across calls, an invariant to protect, several operations on the same data, a shared interface.
- **The deletion test** Turn the methods into functions. If nothing breaks, keep the functions.
- **Settings with no behaviour** A frozen dataclass.

Practice

Exercise 1. Multiple choice (level 1)

Which of these most needs a class?

- a) Rename columns using a mapping.
- b) Learn category frequencies on training data and encode test data with them.
- c) Convert a date column to UTC.
- d) Hold three settings with no behaviour.

Why?

- r1) State learned in one call must be reused in a later call.
- r2) It has the most lines of code.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 class Deduper:
2     def __init__(self) -> None:
3         self.seen: set[int] = set()
4
5     def run(self, ids: list[int]) -> list[int]:
6         out = [i for i in ids if i not in self.seen]
7         self.seen.update(ids)
8         return out
9
10
11 d = Deduper()
12 d.run([1, 2, 3])
13 print(d.run([2, 3, 4]))
```

Exercise 3. Refactor (level 3)

Replace the class with a function `add_share(df, col)` that returns a new DataFrame with a column `share` equal to `df[col] / df[col].sum()`. If the column sums to zero, set `share` to `0.0` for every row.

Starter code:

```
1 class ShareAdder:
2     def __init__(self, col):
3         self.col = col
4
5     def run(self, df):
6         return df.assign(share=df[self.col] / df[self.col].sum())
```

Exercise 4. Spot the bug (level 4)

The same `Scaler` object is fitted on wave 1 and used for waves 2 and 3. Wave 3 results match a scaler fitted on wave 2 only. Tap the line that causes it.

```
1 class Scaler:
2     """Standardise with the mean and sd learned in fit."""
3
4     def fit(self, values: list[float]) -> "Scaler":
5         n = len(values)
6         self.mean_ = sum(values) / n
7         self.sd_ = (sum((v - self.mean_) ** 2 for v in values) / n) ** 0.5
8         return self
9
10    def transform(self, values: list[float]) -> list[float]:
11        self.fit(values)
12        return [(v - self.mean_) / self.sd_ for v in values]
```

Write down the line number of the bug. Then write the fix.

Exercise 5. Architecture design (level 5)

Choose a form for each piece of a survey analysis package.

Options: Function, Frozen dataclass, Class with methods.

1. Recode a Likert column to numbers using a mapping
2. Confidence level, minimum base and weight column
3. Weights, strata and clusters used by mean, CI and subgroup methods

Sources

- scikit-learn's StandardScaler stores mean_ and scale_ after fit. (verified: scikit-learn 1.8 API reference; checked in this repository's dev environment.)

B3 Clean pandas

Roadmap: Prerequisite for all months. Dtypes, vectorisation, method chaining and Copy-on-Write in pandas 3.

B3-L1 Dtypes in pandas 3, nullable and string

Concept

Every pandas column has a **dtype**. The wrong dtype gives wrong answers without an error: IDs that lose leading zeros, counts that turn into floats, text that sorts as numbers.

Three facts about pandas 3.0:

- **Text columns default to the `str` dtype.** `pd.Series(["a", None]).dtype` prints `str`, not `object`. Missing values in it are `NaN`. Only strings are allowed, so a stray integer no longer hides in a text column.
- **A missing value turns an integer column into `float64`,** unless you ask for a nullable dtype. `pd.Series([1, None])` is `float64`. `pd.Series([1, None], dtype="Int64")` keeps integers and shows the gap as `<NA>`.
- **Setting a value that does not fit raises.** Assigning `1.5` into an `int64` Series raises `TypeError` instead of silently upcasting the column. (Assigning `NaN` still upcasts to `float64`.)

The capital letter matters. `int64` is NumPy's integer with no missing values. `Int64`, `Float64` and `boolean` are pandas **nullable** dtypes that hold `pd.NA`.

Set dtypes where data enters, not later:

```
1 df = pd.read_csv(  
2     path,  
3     dtype={"respondent_id": "str", "postcode": "str", "household_size":  
4         "Int64"},  
5 )
```

Without the `dtype` map, `read_csv` reads a postcode such as `01234` as the number `1234.0` once any row is blank. With it, `"01234"` stays a string.

Check the result with `df.dtypes` right after loading. `df.convert_dtypes()` switches every column to the nullable equivalents in one call, useful for exploration. An explicit map is better in pipelines, because it fails loudly when the source changes.

Why it matters

A panel file stores respondent IDs such as 00731. One wave has a blank ID row, so `read_csv` infers `float64` and the ID becomes 731.0. The join to the previous wave, where IDs were read as text, matches nothing. The longitudinal report shows 0% retention for a panel that kept 82% of its members.

`dtype={"respondent_id": "str"}` at load time keeps 00731 as written in every wave.

Worked example: weak code

```
❶ df = pd.read_csv("panel_w2.csv")
❷ df["respondent_id"] = df["respondent_id"].astype(int).astype(str)
❸ df["household_size"] = df["household_size"].fillna(0).astype(int)
```

1. **Line 1.** Dtypes inferred: a blank row turns IDs and counts into float64.
2. **Line 2.** Leading zeros were lost at load time. This cannot bring them back.
3. **Line 3.** Replaces 'unknown' with 0, a real value that biases every mean.

Worked example: strong code

```
1 import pandas as pd
2
3 PANEL_DTYPES: dict[str, str] = {
❶     "respondent_id": "str",
❷     "household_size": "Int64",
6 }
7
8
9 def read_panel(path: str) -> pd.DataFrame:
10     """Read one panel wave with explicit dtypes."""
❸     return pd.read_csv(path, dtype=PANEL_DTYPES)
```

1. **Line 4.** IDs are labels, not numbers: 00731 stays 00731.
2. **Line 5.** Nullable integer: a gap stays <NA>, the rest stay integers.
3. **Line 11.** Dtypes fixed where data enters. In a pipeline, load the map from config.

Decide dtypes at load time. IDs and codes as str, counts with gaps as Int64.

Recap card

- **Default text dtype in pandas 3.0** str, with NaN for missing values. Only strings are allowed.
- **int64 vs Int64** int64 is NumPy, no missing values. Int64 is the pandas nullable integer that holds pd.NA.
- **Keep leading zeros in IDs** Read the column with dtype str in `read_csv`.

- **Assign 1.5 into an int64 Series on pandas 3** Raises `TypeError` instead of silently upcasting.
- **Where to set dtypes** At load time, from a dtype map in config. Check `df.dtypes` right after.

Practice

Exercise 1. Multiple choice (level 1)

On pandas 3.0, what is the dtype of `pd.Series([4, None, 2])`?

- a) `int64`
- b) `Int64`
- c) `float64`
- d) `object`

Why?

- r1) NumPy `int64` cannot hold a missing value, so pandas stores `NaN` in a float column.
- r2) pandas 3 stores every number as float.

Exercise 2. Predict the output (level 2)

What does this print on pandas 3.0?

```
1 import io
2
3 import pandas as pd
4
5 csv = "id,postcode\n1,01234\n2,\n"
6 df = pd.read_csv(io.StringIO(csv))
7 print(df["postcode"].iloc[0])
```

Exercise 3. Fill the gap (level 2)

Keep IDs as text and let the count column hold missing values as integers.

```
1 df = pd.read_csv(path, dtype={"respondent_id": "____(1)",
    "household_size": "____(2)"})
```

Exercise 4. Write the function (level 3)

Write `read_panel(text: str) -> pd.DataFrame`. It reads CSV text with columns `respondent_id` and `household_size`. Keep `respondent_id` as text with leading zeros, and `household_size` as a nullable integer where blanks become `<NA>`. Take the dtype map from a module constant `PANEL_DTYPES`.

Starter code:

```

1 import io
2
3 import pandas as pd
4
5
6 def read_panel(text):
7     return pd.read_csv(io.StringIO(text))

```

Exercise 5. Spot the bug (level 4)

The mean household size reported for wave 3 dropped from 2.6 to 2.1. Nothing changed in the sample. Tap the line that caused it.

```

1 import pandas as pd
2
3
4 def mean_household(df: pd.DataFrame) -> float:
5     """Mean household size among respondents who answered."""
6     sizes = df["household_size"].astype("Int64")
7     sizes = sizes.fillna(0)
8     return float(sizes.mean())

```

Write down the line number of the bug. Then write the fix.

Sources

- In pandas 3.0 text columns default to the str dtype, with NaN as the missing value. (verified: pandas 3.0 release notes; run on pandas 3.0.6 and 3.0.2 in this repository.)
- In pandas 3.0 setting an incompatible value such as 1.5 into an int64 Series raises TypeError. (verified: pandas 3.0 release notes (setitem upcasting removed); run on pandas 3.0.6 and 3.0.2 in this repository.)
- read_csv infers float64 for an integer-like column with a blank value. (verified: Run on pandas 3.0.6 and 3.0.2 in this repository.)

B3-L2 Vectorise instead of apply

Concept

pandas stores each column as an array. A **vectorised** operation, such as `df["a"] * df["b"]`, `np.where()` or `pd.cut()`, runs one compiled loop over the whole array. `df.apply(func, axis=1)` does the opposite. It builds a Series for each row and calls your Python function once per row.

The gap is large. On one million rows, a row-wise `apply` that picks a weight when a flag is set took 4.0 seconds. `np.where` did the same job in 7.5

milliseconds, about 500 times faster. (pandas 3.0.6, NumPy 2.4.6, one laptop core.)

Speed is not the only gain. Vectorised code states the rule for the whole column in one line, so a reviewer checks one expression instead of a branching function.

Reach for these first:

Task	Vectorised tool
Arithmetic between columns	<code>+, -, *, /</code> on Series
Two-way choice	<code>np.where(cond, a, b)</code>
Many-way choice	<code>np.select(conditions, choices, default)</code>
Numeric bands	<code>pd.cut(values, bins, labels, right=False)</code>
Lookup from a dict	<code>Series.map(mapping)</code>
Limits	<code>Series.clip(lower, upper)</code>

Watch missing values. `NaN > 60` is `False`, so `np.where(minutes > 60, "heavy", "light")` labels a missing value "light". Handle missing values on purpose: `pd.cut` keeps `NaN` as missing, and `Series.where(values.notna())` puts it back after `np.where`.

Keep `apply` for work with no vectorised form, such as calling an external parser on each string. Even then, try `Series.map` or a list comprehension on one column before `apply(axis=1)`.

Why it matters

A weighting script derives 25 columns with `apply(axis=1)`. It ran in 8 minutes on 2 million rows, so nobody noticed. The programme then merged three trackers into one file of 40 million rows. At the measured rate of about 4 seconds per million rows per column, the derivations alone need roughly 67 minutes. The job hit its 60-minute timeout every night for a week, and weighted results for all three trackers were late.

The same script also labelled respondents with a missing `minutes` value as "light" viewers, because the row function used `if row["minutes"] > 60`. The vectorised rewrite fixed the speed, and review of the one-line rule exposed the missing-value bug.

Worked example: weak code

```
1 import pandas as pd
2
```

```

3
1 def derive(df):
5     def per_row(row):
6         weighted_reach = row["final_weight"] if row["reach_flag"] == 1
           else 0.0
2         if row["age"] < 35:
8             band = "16-34"
9             elif row["age"] < 55:
10                band = "35-54"
11            else:
12                band = "55+"
13            return pd.Series({"weighted_reach": weighted_reach, "age_group":
                band})
14
4     df[["weighted_reach", "age_group"]] = df.apply(per_row, axis=1)
16     return df

```

- Line 4.** No type hints and no docstring. The reader cannot tell what the frame must contain.
- Line 7.** `NaN < 35` is False, so a missing age falls through to `55+` without a warning.
- Line 15.** Builds a Series and calls Python for every row. About 500 times slower than a vectorised rule.
- Line 15.** Mutates the caller's frame. A pure function returns a new frame.

Worked example: strong code

```

1 import numpy as np
2 import pandas as pd
3
1 AGE_BINS = [16, 35, 55, np.inf]
5 AGE_LABELS = ["16-34", "35-54", "55+"]
6
7
8 def derive(df: pd.DataFrame) -> pd.DataFrame:
9     """Add weighted reach and age group.
10
11     Args:
12         df: Respondents with reach_flag, final_weight and age.
13
14     Returns:
15         A new frame. Missing ages give a missing age_group.
16     """
2     return df.assign(
3         weighted_reach=np.where(df["reach_flag"] == 1,
17         df["final_weight"], 0.0),
4         age_group=pd.cut(df["age"], bins=AGE_BINS, labels=AGE_LABELS,
18         right=False),

```


1. **Line 4.** Band edges live in one constant, so a change to the bands is a one-line change.
2. **Line 17.** `assign` returns a new frame, so the caller's data is untouched.
3. **Line 18.** One compiled pass over the column.
4. **Line 19.** `right=False` makes each band include its lower edge: 35 goes to 35-54. Missing ages stay missing.

The strong version is shorter, about 500 times faster on large frames, and handles missing ages on purpose.

Recap card

- **Why is `df.apply(func, axis=1)` slow?** It calls a Python function once per row. `np.where` measured about 500 times faster on one million rows.
- **Vectorised if-else in NumPy** `np.where(cond, value_if_true, value_if_false)`. For several conditions, `np.select`.
- **What does `NaN > 60` return?** False. `np.where` and `np.select` send missing values to a real category unless you handle them.
- **Numeric bands that keep missing values missing** `pd.cut(values, bins, labels, right=False)`.

Practice

Exercise 1. Multiple choice (level 1)

Which line computes `weighted_reach` for a frame with one million rows in a single compiled pass?

- a) `df.apply(lambda r: r["final_weight"] if r["reach_flag"] == 1 else 0.0, axis=1)`
- b) `[w if f == 1 else 0.0 for w, f in zip(df["final_weight"], df["reach_flag"])]`
- c) `np.where(df["reach_flag"] == 1, df["final_weight"], 0.0)`
- d) `for i, row in df.iterrows(): ...`

Why?

- r1) It evaluates the condition and the choice on whole arrays in compiled code.
- r2) It uses fewer characters than the other options.
- r3) NumPy runs the lambda on several cores.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 import pandas as pd
2
3 s = pd.Series([0.0, 1.5, None])
4 print((s > 1).tolist())
```

Exercise 3. Fill the gap (level 2)

Complete the line so rows with `reach_flag == 1` keep their weight and every other row gets zero.

```
1 df["weighted_reach"] = np.__(1)(df["reach_flag"] == 1,
    df["final_weight"], ____(2))
```

Exercise 4. Write the function (level 3)

Write `weighted_reach_pct(df)`. It returns the weighted share of rows with `reach_flag == 1`: the sum of `final_weight` where the flag is 1, divided by the sum of all `final_weight`. Return `None` when the total weight is zero or the frame is empty. Do not use `apply` or a loop over rows.

Starter code:

```
1 import numpy as np
2 import pandas as pd
3
4
5 def weighted_reach_pct(df: pd.DataFrame) -> float | None:
6     ...
```

Exercise 5. Refactor (level 4)

Refactor `add_age_group` to vectorised code. Keep the bands: under 35 is 16-34, under 55 is 35-54, anything else is 55+. A missing age must give a missing age group, not 55+. Do not change the input frame.

Starter code:

```
1 import pandas as pd
2
3
4 def add_age_group(df):
5     def band(row):
6         if row["age"] < 35:
7             return "16-34"
8         if row["age"] < 55:
9             return "35-54"
10        return "55+"
11
12    df["age_group"] = df.apply(band, axis=1)
```

Exercise 6. Spot the bug (level 5)

The docstring promises that missing minutes stay missing. A respondent with no minutes recorded shows as `"light"`. Tap the line with the bug, then fix the function.

```
1 def add_viewing_band(df: pd.DataFrame) -> pd.DataFrame:
2     """Band weekly minutes into heavy (60 or more) and light. Missing
   minutes stay missing."""
3     minutes = df["minutes"]
4     band = np.where(minutes >= 60, "heavy", "light")
5     return df.assign(viewing_band=band)
```

Write down the line number of the bug. Then write the fix.

Sources

- Row-wise apply was about 500 times slower than `np.where` on one million rows. (verified: Measured in this repo's build environment: pandas 3.0.6, NumPy 2.4.6, 4.03 s against 7.5 ms.)
- `pd.cut` keeps NaN as missing and supports `right=False`. (verified: Run in pandas 3.0.6.)

B3-L3 Method chaining and Copy-on-Write

Concept

Copy-on-Write (CoW) is always on in pandas 3.0. Every indexing result, such as `df[cols]`, `df[mask]` or `df["a"]`, behaves as a separate object. Changing it never changes `df`. pandas copies data only when you write to a shared object, so reads stay cheap.

One consequence: **chained assignment never works**.

```
1 df[df["age"] < 18]["weight"] = 0      # df unchanged,
   ChainedAssignmentError warning
2 df.loc[df["age"] < 18, "weight"] = 0  # one step: df changed
```

The first line writes to a temporary copy, which Python then discards. pandas 3 emits a `ChainedAssignmentError` warning, but the code keeps running. `SettingWithCopyWarning` no longer exists.

CoW rewards a second habit: **method chaining**. Build the result as one expression of steps that each return a new frame.

```
1 result = (
2     raw
```

```

3     .rename(columns=COLUMN_MAP)
4     .loc[lambda d: d["age"] >= 18]
5     .assign(weight=lambda d: d["weight"].clip(upper=5.0))
6     .pipe(add_age_band, bins=BINS, labels=LABELS)
7 )

```

- `assign(col=lambda d: ...)` refers to the frame at that step, not the original `raw`.
- `.loc[lambda d: ...]` filters on the current step.
- `.pipe(func, **kwargs)` calls your own pure function (lesson B2-L1) inside the chain.

Chains read top to bottom, leave no half-finished variables (`df2`, `df_tmp`, `df_final`) in the namespace, and do not mutate the input. Keep each chain to about ten steps. Split longer ones into named functions and join them with `pipe`.

To debug a chain, comment out steps from the bottom up, or add `.pipe(lambda d: print(d.shape) or d)` between steps.

Why it matters

An analyst caps extreme weights with `df[df["weight"] > 5]["weight"] = 5`. On pandas 2 it printed a `SettingWithCopyWarning` that everyone had learned to ignore. The weights were never capped. Three respondents with weights above 12 moved the reported brand share by 4 points. After the upgrade to pandas 3 the assignment still does nothing, with a different warning.

`df.assign(weight=df["weight"].clip(upper=5.0))` or a single `df.loc[mask, "weight"] = 5` does what the analyst meant.

Worked example: weak code

```

1 df = raw.rename(columns=COLUMN_MAP)
2 df2 = df[df["age"] >= 18]
3 df2[df2["weight"] > 5]["weight"] = 5.0
4 df3 = add_age_band(df2, BINS, LABELS)
5 result = df3

```

1. **Line 2.** Intermediate names that outlive their use and invite edits to the wrong one.
2. **Line 3.** Chained assignment: writes to a temporary copy. The weights are not capped.
3. **Line 5.** Which of `df`, `df2`, `df3` is the final state? The reader has to trace every line.

Worked example: strong code

```
1 import pandas as pd
2
3
4 def prepare(raw: pd.DataFrame, column_map: dict[str, str], bins:
    list[int], labels: list[str]) -> pd.DataFrame:
5     """Rename, keep adults, cap weights at 5 and add age bands."""
6     return (
7         raw
8         .rename(columns=column_map)
9         .loc[lambda d: d["age"] >= 18]
10        .assign(weight=lambda d: d["weight"].clip(upper=5.0))
11        .pipe(add_age_band, bins=bins, labels=labels)
12    )
```

1. **Line 9.** Filters the renamed frame at this step.
2. **Line 10.** Returns a new frame with capped weights. No assignment to a copy.
3. **Line 11.** Your own pure function, called as one step of the chain.

Under Copy-on-Write, write with `.loc` in one step, or better, build a new frame with a chain of `assign`, `loc` and `pipe`.

Recap card

- **Copy-on-Write in pandas 3.0** Always on. Indexing results behave as separate objects, copied only on write.
- **`df[mask]["col"] = value` on pandas 3** Changes nothing and emits `ChainedAssignmentError`. Use `df.loc[mask, "col"] = value`.
- **Refer to the current step inside a chain** A `lambda`: `assign(x=lambda d: ...)` or `loc[lambda d: ...]`.
- **Call your own function inside a chain** `.pipe(func, **kwargs)`.

Practice

Exercise 1. Multiple choice (level 1)

On pandas 3.0, `df[df["age"] < 18]["weight"] = 0` runs. What happens to `df`?

- a) Weights of under-18s become 0.
- b) Nothing. pandas emits a `ChainedAssignmentError` warning.
- c) It raises an exception and stops.
- d) Every weight becomes 0.

Why?

- r1) The first indexing step returns a new object, so the write lands on a temporary copy.

- r2) pandas 3 makes every DataFrame read-only.

Exercise 2. Predict the output (level 2)

What does this print on pandas 3.0?

```
1 import pandas as pd
2
3 df = pd.DataFrame({"a": [1, 2, 3]})
4 s = df["a"]
5 s.iloc[0] = 100
6 print(df["a"].tolist())
```

Exercise 3. Put in order (level 2)

Put the steps of the chain in order. The raw columns are `wt` and `age`. Weights must average 1.0 across the adults in the output.

- `.pipe(add_age_band, bins=BINS, labels=LABELS)`
- `.assign(weight=lambda d: d["weight"] / d["weight"].mean())`
- `.rename(columns={"wt": "weight"})`
- `.loc[lambda d: d["age"] >= 18]`

Exercise 4. Refactor (level 3)

Rewrite `cap_weights` as one method chain that returns a new frame. Keep only rows where `age >= 18`, then cap `weight` at `cap`. Leave the input unchanged. No chained assignment.

Starter code:

```
1 def cap_weights(df, cap):
2     adults = df[df["age"] >= 18]
3     adults[adults["weight"] > cap]["weight"] = cap
4     return adults
```

Exercise 5. Spot the bug (level 4)

Weights for invalid interviews should be zero. The check after the function shows they are unchanged. Tap the line to fix, and rewrite the function.

```
1 import pandas as pd
2
3
4 def zero_invalid(df: pd.DataFrame) -> pd.DataFrame:
5     """Set weight to 0 where valid is False."""
6     out = df.copy()
7     out[~out["valid"]]["weight"] = 0.0
8     return out
```

Write down the line number of the bug. Then write the fix.

Sources

- Copy-on-Write is always enabled in pandas 3.0 and the `mode.copy_on_write` option has no effect. (verified: pandas 3.0 release notes; Pandas4Warning text observed on pandas 3.0.6 and 3.0.2 in this repository.)
- Chained assignment never updates the original in pandas 3.0 and emits `ChainedAssignmentError` as a warning. (verified: pandas Copy-on-Write user guide; run on pandas 3.0.6 and 3.0.2 in this repository.)
- `SettingWithCopyWarning` no longer exists in pandas 3.0. (verified: pandas 3.0 release notes; pandas.errors has no `SettingWithCopyWarning` on 3.0.6.)

B4 Polars and DuckDB

Roadmap: Prerequisite for all months. Tools for larger local datasets.

B4-L1 Polars expressions and lazy frames

Concept

Polars is a DataFrame library written in Rust. It runs on all your cores and has no row index. You describe work with **expressions**:

`pl.col("weight").sum()` is a recipe, not a result. Polars runs it inside a context.

```
1 import polars as pl
2
3 shares = (
4     df.group_by("wave", "brand")
5     .agg(pl.col("weight").sum().alias("brand_weight"))
6     .with_columns(total=pl.col("brand_weight").sum().over("wave"))
7 )
```

- `select` returns only the expressions you list. `with_columns` adds or replaces columns. `filter` keeps rows. `group_by(...).agg(...)` aggregates.
- `.over("wave")` computes a group total and broadcasts it back to each row, like a window function in SQL.
- `pl.when(...).then(...).otherwise(...)` replaces `np.where`.

Polars has two modes. A `DataFrame` runs each step at once. A `LazyFrame`, from `df.lazy()` or `pl.scan_parquet(path)`, records the steps and runs nothing until `.collect()`. Before running, the optimiser rewrites the plan. It pushes filters down to the scan, so rows you drop are never read, and reads only the columns you use. `lf.explain()` prints the plan.

Two behaviours to know (Polars 1.33):

- `sum()` **of a group with only nulls returns 0**, not null. Check `null_count()` before you trust a total.
- `mean()` **skips nulls**. `sum() / pl.len()` counts null rows in the denominator and gives a different answer.

Use lazy mode for files. Scan, filter, select, aggregate, then collect once at the end.

Why it matters

A weekly job reads a 9 GB Parquet extract with `pl.read_parquet`, then filters to the current wave. Memory peaks at 14 GB and the job fails on a 16 GB runner on the week the extract grows.

```
pl.scan_parquet(path).filter(pl.col("wave") ==  
wave).select(cols).collect()
```

 lets Polars skip other waves' row groups and unused columns at read time, so memory scales with one wave, not the whole file.

Worked example: weak code

```
❶ df = pl.read_parquet("responses.parquet")  
2 df = df.filter(pl.col("wave") == "2026-09")  
3 totals = df.group_by("brand").agg(pl.col("weight").sum())  
❷ totals = totals.with_columns(mean_w=pl.col("weight") / df.height)
```

1. **Line 1.** Eager read: every wave and every column loads before the filter.
2. **Line 4.** Divides a brand total by the whole wave's row count: not the brand mean, and null rows count.

Worked example: strong code

```
1 import polars as pl  
2  
3  
4 def brand_weights(path: str, wave: str) -> pl.DataFrame:  
5     """Weighted totals and means per brand for one wave. Grain: one row  
   per brand."""  
6     return (  
❶         pl.scan_parquet(path)  
❷         .filter(pl.col("wave") == wave)  
9         .group_by("brand")  
10        .agg(  
11            total=pl.col("weight").sum(),  
12            mean_w=pl.col("weight").mean(),  
❸            missing=pl.col("weight").null_count(),  
14        )  
15        .sort("brand")  
❹        .collect()  
17    )
```

1. **Line 7.** Lazy scan: nothing is read yet.
2. **Line 8.** Pushed down to the scan: other waves are skipped at read time.
3. **Line 13.** Reports nulls, because an all-null group sums to 0.
4. **Line 16.** One collect at the end runs the optimised plan.

Scan lazily, filter and select early, aggregate with expressions, and collect once.

Recap card

- **Polars eager vs lazy** DataFrame runs each step at once. LazyFrame (scan_*, .lazy()) builds a plan and runs on collect().
- **Group total on every row in Polars** pl.col("x").sum().over("group").
- **Polars sum() of an all-null group** 0, not null. Report null_count() with totals.
- **Mean that skips nulls in Polars** pl.col("x").mean(). sum() / pl.len() counts null rows.

Practice

Exercise 1. Multiple choice (level 1)

lf = pl.scan_parquet("big.parquet").filter(pl.col("wave") == "w9"). What has happened after this line runs?

- a) The whole file is in memory, filtered to wave w9.
- b) Nothing is read yet. lf holds a query plan.
- c) Only the w9 rows are in memory.
- d) Polars wrote a filtered copy of the file.

Why?

- r1) A LazyFrame records steps and runs them only on collect().
- r2) Parquet files cannot be read without DuckDB.

Exercise 2. Predict the output (level 2)

What does this print on Polars 1.33?

```
1 import polars as pl
2
3 df = pl.DataFrame({"wave": ["w1", "w1", "w2"], "w": [1.0, 2.0, None]})
4 out = df.group_by("wave").agg(pl.col("w").sum()).sort("wave")
5 print(out["w"].to_list())
```

Exercise 3. Fill the gap (level 2)

Add a share column with each brand's weight over its wave's total, then run the lazy plan.

```
1 out = (
2     lf.with_columns(share=pl.col("weight") /
3                     pl.col("weight").sum().____(1)("wave"))
4     .____(2)()
5 )
```

Exercise 4. Write the function (level 3)

Write `brand_share(df: pl.DataFrame) -> pl.DataFrame`. Input columns: `wave`, `brand`, `weight`. Return one row per wave and brand with columns `wave`, `brand`, `share`, where `share` is the brand's weight total over the wave's weight total. When a wave's total is 0, `share` is null. Sort by `wave`, then `brand`.

Starter code:

```
1 import polars as pl
2
3
4 def brand_share(df):
5     ...
```

Exercise 5. Spot the bug (level 4)

The mean satisfaction score per region looks too low in regions with many skipped answers. Tap the line with the bug and fix it.

```
1 import polars as pl
2
3
4 def mean_score(df: pl.DataFrame) -> pl.DataFrame:
5     """Mean score per region among respondents who answered."""
6     return (
7         df.group_by("region")
8         .agg(mean_score=pl.col("score").sum() / pl.len())
9         .sort("region")
10    )
```

Write down the line number of the bug. Then write the fix.

Sources

- Polars `sum()` of a group with only nulls returns 0. (verified: Run on Polars 1.33.1 and 1.44.2 in this repository.)
- LazyFrame optimisation pushes filters and column selection down to the scan. (verified: Polars user guide, lazy API and optimisations; plan checked with `explain()` on Polars 1.33.1.)

B4-L2 DuckDB, SQL over files and DataFrames

Concept

DuckDB is an analytical SQL database that runs inside your Python process. There is no server to install. It reads Parquet, CSV and JSON files directly, and it queries pandas and Polars DataFrames that are in scope.

```
1 import duckdb
```

```

2
3 con = duckdb.connect() # in-memory database
4 rel = con.sql("""
5     SELECT wave, brand, sum(weight) AS total
6     FROM read_parquet('extracts/*.parquet')
7     GROUP BY wave, brand
8 """)
9 df = rel.df() # pandas; .pl() for Polars, .fetchall() for tuples

```

- A glob such as `'extracts/*.parquet'` reads many files as one table. Filters and column lists are pushed down to the reader, as in Polars lazy mode.
- `FROM responses` finds a DataFrame variable named `responses` in the calling scope. `con.register("responses", df)` names it explicitly, which is clearer inside functions.
- Results come back as a relation. Nothing is materialised until you call `.df()`, `.pl()`, `.fetchall()` or `.fetchone()`.

Pass values as parameters, never with f-strings.

```
1 con.execute("SELECT count(*) FROM responses WHERE wave = ?", [wave])
```

With `f"... WHERE wave = '{wave}'"`, a value that holds a quote changes the query. In this course's own test, the input `w1'` or `'1'='1` turned a filter for one wave into a count of every row. A `?` placeholder (or `$name` with a dict) sends the value as data, so quotes cannot change the SQL. File paths and table names cannot be parameters. Check those against a list of allowed values.

Create the connection at the edge of your program and pass `con` to functions. The module-level `duckdb.sql()` uses one shared default connection, a global that tests cannot replace.

Why it matters

A dashboard helper built its filter with `f"WHERE region = '{region}'"`. The region list came from a client file. One entry was `Hawke's Bay`. The apostrophe broke the SQL, and the job failed at 06:00 before the client's morning meeting. A second entry with a crafted value could have read other clients' rows.

`con.execute("... WHERE region = ?", [region])` handles `Hawke's Bay` as a plain string.

Worked example: weak code

```

1 import duckdb
2
3

```

```

4 def wave_total(wave):
1   q = f"SELECT sum(weight) FROM read_parquet('data/*.parquet') WHERE
    wave = '{wave}'"
2   return duckdb.sql(q).fetchone()[0]

```

1. **Line 5.** Value pasted into SQL: a quote in wave breaks or rewrites the query.
2. **Line 6.** Shared default connection, and None when no rows match.

Worked example: strong code

```

1 import duckdb
2
3 TOTAL_SQL = "SELECT sum(weight) FROM responses WHERE wave = ?"
4
5
6 def wave_total(con: duckdb.DuckDBPyConnection, wave: str) -> float:
7     """Total weight for one wave, 0.0 when the wave has no rows."""
8     row = con.execute(TOTAL_SQL, [wave]).fetchone()
9     if row is None or row[0] is None:
10         return 0.0
11     return float(row[0])

```

1. **Line 3.** The value goes in as a parameter. The SQL text never changes.
2. **Line 6.** The connection is passed in, so a test gives it a small in-memory table.
3. **Line 9.** SQL sum over no rows is NULL. Test with `is None`.

One connection passed in, parameters for values, and a None check on aggregates.

Recap card

- **Query many Parquet files in DuckDB** `FROM read_parquet('dir/*.parquet')`.
- **Pass a value into DuckDB SQL** `con.execute("... WHERE x = ?", [value])`. Never an f-string.
- **Table names from user input** Placeholders cannot bind identifiers. Check the name against an allowlist.
- **SQL sum over no rows** NULL, returned as None. Check with `is None`.

Practice

Exercise 1. Multiple choice (level 1)

Which call safely filters by a `region` value read from a client file?

- a) `con.execute(f"... WHERE region = '{region}')`
- b) `con.execute("... WHERE region = ?", [region])`
- c) `con.execute("... WHERE region = ' + region + '")`

- d) `con.execute("... WHERE region = '%s' " % region)`

Why?

- r1) The value is sent as data, separate from the SQL text.
- r2) It is shorter.

Exercise 2. Predict the output (level 2)

What does this print? (DuckDB 1.5)

```
1 import duckdb
2
3 con = duckdb.connect()
4 con.execute("CREATE TABLE t AS SELECT * FROM (VALUES ('w1', 1.0), ('w2',
    2.0)) v(wave, weight)")
5 print(con.execute("SELECT sum(weight) FROM t WHERE wave =
    'w9'").fetchone())
```

Exercise 3. Fill the gap (level 2)

Count rows for one wave from all Parquet files in `extracts/`, with the wave as a parameter.

```
1 con.execute(
2     "SELECT count(*) FROM ____ (1)('extracts/*.parquet') WHERE wave =
    ____ (2)",
3     [wave],
4 ).fetchone()
```

Exercise 4. Write the function (level 3)

Write `region_total(con, region)`. It returns the sum of `weight` in table `responses` for the given `region`, as a `float`. Return `0.0` when no rows match. Pass `region` as a parameter. The checker runs it with regions such as `Hawke's Bay`.

Starter code:

```
1 import duckdb
2
3
4 def region_total(con, region):
5     return con.execute(f"SELECT sum(weight) FROM responses WHERE region =
    '{region}'").fetchone()[0]
```

Exercise 5. Architecture design (level 4)

A reporting script lets users pick a table to summarise and a wave to filter on. Choose how each input reaches DuckDB.

Options: A ? parameter, An f-string inside the SQL.

1. The wave value

- 2. The table name
- 3. The connection

Sources

- DuckDB binds ? and \$name parameters with con.execute(sql, params). (verified: DuckDB Python API documentation; run on DuckDB 1.5.1 and 1.5.5 in this repository.)
- DuckDB queries pandas DataFrames in the calling scope by variable name and reads Parquet globs with read_parquet. (verified: DuckDB Python API documentation; run on DuckDB 1.5.1 and 1.5.5 in this repository.)
- Pasting w1' or '1'='1 into an f-string filter returned every row; the same value as a parameter matched none. (verified: Run on DuckDB 1.5.1 in this repository.)

B4-L3 Choosing pandas, Polars, DuckDB or Spark

Concept

All four tools answer the same questions. Pick by where the data lives and how big it is next to your memory, not by fashion.

A measurement from this course's build: a group-by with sum and mean over 10 million rows and 1,000 groups, best of three runs on a 4-core machine. pandas 3.0.6 took 0.22 s, Polars 1.44 took 0.11 s, DuckDB 1.5 took 0.06 s. At this size every tool is fast enough. The choice matters when data outgrows memory.

Situation	Good default
Data fits comfortably in memory, you need the scikit-learn, statsmodels or plotting ecosystem	pandas
Files larger than you want in memory, heavy transforms, one machine	Polars lazy (scan_parquet)
SQL is the clearest way to express it, or you join many files ad hoc	DuckDB
Data already in a lakehouse, or larger than one machine, or a scheduled production job	Spark (Track A)

The tools share the **Apache Arrow** memory format, so you mix them at the boundaries:

```
1 lazy_result = pl.scan_parquet("extracts/  
    *.parquet").filter(...).group_by(...).agg(...)
```

```

2 summary = lazy_result.collect().to_pandas()      # Polars to pandas for
statsmodels
3 rel = con.sql("SELECT ... FROM read_parquet('extracts/*.parquet')")
4 top = rel.pl()                                  # DuckDB to Polars

```

Rules for mixing:

- **Reduce first, convert last.** Filter and aggregate in the engine that reads the files. Convert the small result.
- **Watch the index.** Polars has no index. `pl.from_pandas(df)` drops a pandas index unless you pass `include_index=True`, or call `reset_index()` first.
- **Check dtypes after each hop.** Text becomes `str` in pandas 3 and `String` in Polars. Nullable integers need care in both directions.

One project, one primary engine. A second tool at a clear boundary is fine. Four tools in one function is not.

Why it matters

An analyst sets `respondent_id` as the pandas index to deduplicate, then hands the frame to a Polars function for a fast join. `pl.from_pandas` drops the index. The join key is gone, so the analyst joins on row position with `hstack`. Two waves sorted differently, and 11% of respondents got another person's scores.

`pl.from_pandas(df.reset_index())` keeps the ID as a column, and the join uses it.

Worked example: weak code

```

❶ big = pl.read_parquet("extracts/*.parquet").to_pandas()
❷ big = big[big["wave"] == WAVE].set_index("respondent_id")
❸ stats = pl.from_pandas(big).group_by("brand").agg(pl.col("weight").sum())

```

1. **Line 1.** Reads every wave eagerly, then copies all of it into pandas before filtering.
2. **Line 2.** The ID moves into the index.
3. **Line 3.** Back to Polars: the index, and so the ID, is dropped. Two conversions of the full data.

Worked example: strong code

```

1 import pandas as pd
2 import polars as pl
3
4
5 def brand_totals(path_glob: str, wave: str) -> pd.DataFrame:

```



```

6      """Weight per brand for one wave, as pandas for modelling. Grain: one
      row per brand."""
7      return (
8          pl.scan_parquet(path_glob)
9          ❶ .filter(pl.col("wave") == wave)
10         .group_by("brand")
11         .agg(total=pl.col("weight").sum())
12         .sort("brand")
13         .collect()
14         ❷ .to_pandas()
15     )

```

1. **Line 9.** Reduce in the engine that reads the files.
2. **Line 14.** Convert once, at the end, when the result has one row per brand.

One primary engine. Reduce where the data is read. Convert small results at a clear boundary.

Recap card

- **Group-by over 10M rows: pandas, Polars, DuckDB** 0.22 s, 0.11 s, 0.06 s on 4 cores. All fast enough. Choose by size against memory and where data lives.
- **Mixing engines** Reduce first in the engine that reads the files, convert the small result last.
- **pl.from_pandas and the index** Dropped by default. `reset_index()` first, or `include_index=True`.
- **When Spark** Data already in the lakehouse, larger than one machine, or a scheduled production job.

Practice

Exercise 1. Multiple choice (level 1)

You have 40 GB of Parquet on a laptop with 16 GB of RAM. You need one table of weighted totals by wave and brand, once. Which is the best first choice?

- a) pandas `read_parquet`, then `groupby`
- b) DuckDB or Polars lazy, aggregating while reading
- c) Set up a Spark cluster
- d) Split the files by hand and run pandas on each

Why?

- r1) Both read Parquet in chunks and push filters and aggregation into the scan, so the full data never sits in memory.

- r2) They are newer than pandas.

Exercise 2. Predict the output (level 2)

What does this print? (pandas 3.0, Polars 1.33)

```
1 import pandas as pd
2 import polars as pl
3
4 pdf = pd.DataFrame({"respondent_id": ["r1", "r2"], "score": [3,
    4]}).set_index("respondent_id")
5 print(pl.from_pandas(pdf).columns)
```

Exercise 3. Fill the gap (level 2)

Hand a DuckDB result to Polars, then a Polars result to pandas.

```
1 top_brands = con.sql("SELECT brand, sum(weight) AS total FROM responses
    GROUP BY brand").____(1)()
2 for_model = top_brands.____(2)()
```

Exercise 4. Spot the bug (level 3)

This function runs out of memory on the full extract. Tap the line that loads far more than the result needs, and fix the function.

```
1 import pandas as pd
2 import polars as pl
3
4
5 def wave_mean(path: str, wave: str) -> float:
6     """Mean weight for one wave."""
7     df = pl.read_parquet(path).to_pandas()
8     return float(df.loc[df["wave"] == wave, "weight"].mean())
```

Write down the line number of the bug. Then write the fix.

Exercise 5. Architecture design (level 4)

Choose the primary engine for each job.

Options: pandas, Polars lazy, DuckDB, Spark.

1. Fit a logistic regression with statsmodels on a 50,000-row survey file
2. Ad hoc SQL joins across 30 CSV and Parquet exports on your laptop
3. Nightly job over tables that already live in the Databricks lakehouse
4. Heavy column transforms over 25 GB of Parquet on one 32 GB machine

Sources

- Group-by sum and mean over 10M rows: pandas 3.0.6 0.22 s, Polars 1.44.2 0.11 s, DuckDB 1.5.5 0.06 s (best of 3, 4 cores). (verified: Measured in this repository's dev environment.)
- `pl.from_pandas` drops the pandas index unless `include_index=True`. (verified: Polars API reference; run on Polars 1.33.1 and 1.44.2 in this repository.)

B5 Data validation

Roadmap: Prerequisite for all months. Schema checks with pandera and input models with pydantic.

B5-L1 DataFrame schemas with pandera

Concept

A **schema** states what a DataFrame must look like: which columns exist, their dtypes, and rules on their values. **pandera** checks a pandas DataFrame against a schema and raises when it does not match. Examples here use pandera 0.33 with pandas 3.

```
1 import pandera.pandas as pa
2
3 RESPONSES = pa.DataFrameSchema(
4     {
5         "respondent_id": pa.Column(str, unique=True),
6         "age": pa.Column(int, pa.Check.in_range(16, 99)),
7         "weight": pa.Column(float, pa.Check.gt(0)),
8     },
9     strict=True,
10    coerce=True,
11 )
12 clean = RESPONSES.validate(raw, lazy=True)
```

- `pa.Column(dtype, checks, nullable=False)` is the default: nulls fail unless you set `nullable=True`.
- `unique=True` rejects duplicate keys.
- `strict=True` rejects columns the schema does not name. `coerce=True` converts dtypes before checking.
- `lazy=True` collects **every** failure into one `SchemaErrors` exception. Its `failure_cases` attribute is a DataFrame with the column, check and failing value. Without `lazy`, validation stops at the first `SchemaError`.

`SchemaErrors` is not a subclass of `SchemaError`. Catch both when you handle failures. A `strict=True` violation raises `SchemaErrors` even without `lazy=True`.

The class syntax, `pa.DataFrameModel`, does the same with type hints. Add `@pa.check_types` to a function whose parameters are annotated `DataFrame[Model]`, and pandera validates inputs and outputs on every call.

Import from `pandera.pandas`. The plain `import pandera` still works in 0.33, but the pandas-specific module is the documented path.

pandera does not run in this app's browser runtime. The exercises check your code's structure. Run `pytest` locally with the `dev` group to execute schemas for real.

Why it matters

A supplier changes a weight export so that non-responders get weight `0` instead of being left out. The weighting script divides by the sum of weights per cell. Three cells now contain only zeros. The script writes `NaN` shares to the client dashboard, and nobody sees an error.

`pa.Column(float, pa.Check.gt(0))` on the input fails the run on day one, with the failing rows listed in `failure_cases`.

Worked example: weak code

```
1 def load_weights(path):
2     df = pd.read_csv(path)
3     ❶ assert "weight" in df.columns
4     ❷ df = df.dropna()
5     return df
```

1. **Line 3.** Checks one column exists. Not its type, its values or other columns. And python -O strips asserts.
2. **Line 4.** Silently drops rows with any null, in any column. The count of dropped rows is never reported.

Worked example: strong code

```
1 import pandas as pd
2 import pandera.pandas as pa
3
4 WEIGHTS = pa.DataFrameSchema(
5     {
6     ❶     "respondent_id": pa.Column(str, unique=True),
7         "cell": pa.Column(str),
8     ❷     "weight": pa.Column(float, pa.Check.gt(0)),
9     },
10    ❸    strict=True,
11    coerce=True,
12 )
13
14
15 def load_weights(path: str) -> pd.DataFrame:
16     """Read and validate a weights file. Grain: one row per respondent.
17
18     Raises:
19         pa.errors.SchemaErrors: Listing every failing row and check.
20     """
```

```
4 return WEIGHTS.validate(pd.read_csv(path, dtype={"respondent_id":
    "str"}), lazy=True)
```

1. **Line 6.** The grain, enforced: one row per respondent.
2. **Line 8.** Not null (the default) and positive: zero weights fail loudly.
3. **Line 10.** A new or renamed column in the supplier file fails the run.
4. **Line 21.** `lazy=True` reports every failure at once, not just the first.

A schema at the point of entry turns a silent bad input into one loud, detailed failure.

Recap card

- **pandera Column defaults** `nullable=False`. Set `nullable=True` where nulls are expected.
- **`strict=True` and `coerce=True`** `strict` rejects unknown columns. `coerce` converts dtypes before checks.
- **`validate(df, lazy=True)`** Runs every check and raises `SchemaErrors` with `failure_cases`.
- **`SchemaError` vs `SchemaErrors`** Separate classes. `SchemaErrors` is not a subclass. Catch both, or better, let them raise.

Practice

Exercise 1. Multiple choice (level 1)

In `pandera 0.33`, what does `lazy=True` change in `schema.validate(df, lazy=True)` ?

- a) Validation runs later, when the `DataFrame` is first used.
- b) Every failure is collected and raised together as `SchemaErrors`.
- c) Failing rows are dropped instead of raising.
- d) Only dtypes are checked, not values.

Why?

- r1) You see all problems in one run, with `failure_cases` listing each one.
- r2) It is faster because it skips checks.

Exercise 2. Predict the output (level 2)

What does this print? (`pandera 0.33`)

```
1 import pandas as pd
2 import pandera.pandas as pa
3
4 schema = pa.DataFrameSchema({"w": pa.Column(float, pa.Check.gt(0))})
5 try:
```

```

6     schema.validate(pd.DataFrame({"w": [1.0, -2.0, -3.0]}), lazy=True)
7 except pa.errors.SchemaErrors as err:
8     print(err.failure_cases["failure_case"].tolist())

```

Exercise 3. Fill the gap (level 2)

Reject duplicate IDs, and reject any column the schema does not name.

```

1 schema = pa.DataFrameSchema(
2     {"respondent_id": pa.Column(str, ____ (1)=True), "score":
3     pa.Column(int)},
4     ____ (2)=True,

```

Exercise 4. Write the function (level 3)

Define a module constant `SCORES` as a pandera `DataFrameSchema` with:

- `respondent_id`: str, unique.
- `nps`: int, between 0 and 10 inclusive.
- `comment`: str, nulls allowed.
- Unknown columns rejected, dtypes coerced.

Then write `validate_scores(df: pd.DataFrame) -> pd.DataFrame` that returns `SCORES.validate(df, lazy=True)`.

pandera does not run in the browser, so the checker reads your code. Run it locally with `pytest` to see it validate.

Starter code:

```

1 import pandas as pd
2 import pandera.pandas as pa
3
4
5 def validate_scores(df):
6     return df

```

Exercise 5. Spot the bug (level 4)

A bad weights file passed through this loader and broke the dashboard. Tap the line that let it through.

```

1 import pandas as pd
2 import pandera.pandas as pa
3
4 WEIGHTS = pa.DataFrameSchema({"weight": pa.Column(float,
5     pa.Check.gt(0))}, strict=True)
6
7 def load_weights(path: str) -> pd.DataFrame:
8     """Read and validate weights."""
9     df = pd.read_csv(path)

```

```

10     try:
11         return WEIGHTS.validate(df, lazy=True)
12     except pa.errors.SchemaError:
13         return df

```

Write down the line number of the bug. Then write the fix.

Sources

- pandera 0.33: lazy=True raises SchemaErrors with a failure_cases DataFrame. (verified: pandera documentation; run on pandera 0.33.1 with pandas 3.0.6 in this repository.)
- SchemaErrors is not a subclass of SchemaError in pandera 0.33. (verified: issubclass check on pandera 0.33.1 in this repository.)
- A strict=True violation raises SchemaErrors even without lazy=True. (verified: Run on pandera 0.33.1 in this repository.)
- import pandera.pandas as pa is the documented import for pandas schemas. (search-only: pandera documentation (docs site not reachable from this environment).)

B5-L2 pydantic models for config and input records

Concept

A dataclass (lesson B2-L2) does not check types. **pydantic** does. A `BaseModel` validates and converts data when you build it, and raises one `ValidationError` that lists every problem. Examples use pydantic 2.12.

```

1  from pydantic import BaseModel, ConfigDict, Field
2
3
4  class WeightingConfig(BaseModel):
5      model_config = ConfigDict(extra="forbid", frozen=True)
6
7      weight_col: str
8      min_base: int = Field(default=50, gt=0)
9      cap: float = Field(default=5.0, gt=1.0)
10
11
12  cfg = WeightingConfig.model_validate(yaml.safe_load(text))

```

- `Field(gt=0, le=10)` adds bounds. `@field_validator("name")` adds your own rule.
- `extra="forbid"` rejects unknown keys, so `minbase: 30` in a YAML file fails instead of being ignored. The default, `"ignore"`, drops them silently.
- `frozen=True` rejects assignment after construction.

- `e.errors()` returns a list of dicts with `loc`, `type` and `msg`. Log it as structured data.

pydantic runs in **lax mode** by default. It converts `"30"` to `30` and `30.0` to `30`, but rejects `"thirty"` and `30.5`. It also converts `True` to `1` for an `int` field. Set `ConfigDict(strict=True)` when inputs must already have the right type.

One gap: `model_copy(update={...})` does **not** validate.

`cfg.model_copy(update={"min_base": -1})` returns a config with `min_base = -1`. To change a validated model, build a new one:

`WeightingConfig.model_validate(**cfg.model_dump(), "min_base": 30)`.

Use pydantic for records and settings that arrive from outside: YAML, JSON, API payloads. Use pandera (lesson B5-L1) for whole DataFrames. For a list of records, `TypeAdapter(list[Model]).validate_python(rows)` validates them in one call.

Why it matters

A weighting config file had `cap: 0.5` typed where `5` was meant, and `minbase` spelt without the underscore. The loader used a plain dict. The cap shrank every weight above `0.5` down to `0.5`, which flattened the weighting, and `min_base` stayed at its default. The report went out with unweighted-looking results.

`WeightingConfig.model_validate` with `gt=1.0` and `extra="forbid"` rejects both mistakes in one error message, before any data is read.

Worked example: weak code

```
1 def load_config(raw):
2     return {
3         "weight_col": raw.get("weight_col", "weight"),
4         "min_base": int(raw.get("min_base", 50)),
5         "cap": float(raw.get("cap", 5.0)),
6     }
```

1. **Line 3.** A missing key falls back to a guess.
2. **Line 4.** A typo such as `minbase` is ignored. `int()` accepts negatives.
3. **Line 5.** No bounds: `0.5` passes.

Worked example: strong code

```
1 from pydantic import BaseModel, ConfigDict, Field
2
3
4 class WeightingConfig(BaseModel):
5     """Settings for one weighting run, loaded from YAML."""
```

```

6
1 model_config = ConfigDict(extra="forbid", frozen=True)
8
2 weight_col: str
10 min_base: int = Field(default=50, gt=0)
3 cap: float = Field(default=5.0, gt=1.0)
12
13
14 def load_config(raw: dict[str, object]) -> WeightingConfig:
15     """Validate a parsed YAML mapping.
16
17     Raises:
18         pydantic.ValidationError: Listing every invalid or unknown key.
19     """
20     return WeightingConfig.model_validate(raw)

```

1. **Line 7.** Unknown keys fail. No changes after loading.
2. **Line 9.** Required: no default, so a missing key fails.
3. **Line 11.** A bound that encodes what a cap means.

Validate external settings and records once, at the boundary, with every rule in the model.

Recap card

- **pydantic extra keys, default behaviour** Ignored silently. Use `ConfigDict(extra="forbid")` for config.
- **pydantic lax mode** Converts "30" and 30.0 to 30, True to 1. Use `strict=True` to refuse conversions.
- **model_copy(update=...)** Does not validate. Rebuild with `model_validate(**m.model_dump(), ...)`.
- **pydantic or pandera** pydantic for records and settings from outside. pandera for whole DataFrames.

Practice

Exercise 1. Multiple choice (level 1)

A pydantic 2 model has the default config. A YAML file adds an unknown key `minbase: 30`. What happens on `model_validate`?

- a) `ValidationError: extra_forbidden`
- b) The key is ignored silently.
- c) `min_base` is set to 30.
- d) A warning is printed.

Why?

- r1) The default extra setting is "ignore". Set extra="forbid" to reject unknown keys.
- r2) pydantic never checks keys.

Exercise 2. Predict the output (level 2)

What does this print? (pydantic 2.12)

```
1 from pydantic import BaseModel, Field
2
3
4 class Cfg(BaseModel):
5     min_base: int = Field(default=50, gt=0)
6
7
8 cfg = Cfg(min_base="30")
9 print(cfg.model_copy(update={"min_base": -1}).min_base)
```

Exercise 3. Fill the gap (level 2)

Reject unknown keys, stop changes after loading, and require nps between 0 and 10.

```
1 class Response(BaseModel):
2     model_config = ConfigDict(extra="____(1)", ____ (2)=True)
3
4     respondent_id: str
5     nps: int = Field(ge=0, ____ (3)=10)
```

Exercise 4. Write the function (level 3)

Write a pydantic model `WeightingConfig` and a function `load_config(raw)`.

- Fields: `weight_col: str` (required), `min_base: int` default 50 and greater than 0, `cap: float` default 5.0 and greater than 1.0.
- Reject unknown keys. Reject assignment after construction.
- `load_config(raw: dict[str, object]) -> WeightingConfig` validates the mapping.

Starter code:

```
1 from pydantic import BaseModel, ConfigDict, Field
2
3
4 def load_config(raw):
5     return raw
```

Exercise 5. Spot the bug (level 4)

A job lowers the minimum base for a small market. The run then reports subgroups of size 0 as valid. Tap the line with the bug and fix the function so the new value is validated.

```
1 from pydantic import BaseModel, ConfigDict, Field
2
3
4 class Cfg(BaseModel):
5     """Run settings."""
6
7     model_config = ConfigDict(extra="forbid", frozen=True)
8
9     min_base: int = Field(default=50, gt=0)
10
11
12 def with_min_base(cfg: Cfg, value: int) -> Cfg:
13     """Return a copy of cfg with a new min_base."""
14     return cfg.model_copy(update={"min_base": value})
```

Write down the line number of the bug. Then write the fix.

Sources

- pydantic 2 default extra behaviour is ignore; extra="forbid" raises extra_forbidden. (verified: pydantic documentation; run on pydantic 2.12.5 and 2.13.5 in this repository.)
- Lax mode converts "30" and 30.0 to int 30 and True to 1, and rejects 30.5 and "thirty". (verified: pydantic conversion table; run on pydantic 2.12.5 in this repository.)
- model_copy(update=...) does not validate the update. (verified: pydantic documentation for model_copy; run on pydantic 2.12.5 in this repository.)

B5-L3 Validate at the boundaries, with severity

Concept

Validation belongs at **boundaries**, where data crosses from one owner to another:

- **On the way in.** After reading a file, API response or config. Check schema, keys and value ranges (lessons B5-L1 and B5-L2).
- **On the way out.** Before writing a table or sending a report. Check the promises you make to the next consumer: grain, row counts, totals, shares that sum to 1.

Checks in the middle of a pure transformation add noise. A function that receives validated input and is unit tested does not need to re-check it.

Not every failed check should stop a run. Give each check a **severity**:

- **CRITICAL** stops the run. Wrong grain, missing required column, empty input, weights that do not sum to the sample size.
- **WARN** records the problem and continues. Missing share above its usual level, a subgroup below its reporting base.

Record results as **structured data**, not prose, so a dashboard or alert filters them:

```
1 Severity = Literal["CRITICAL", "WARN"]
2
3
4 @dataclass(frozen=True)
5 class QaResult:
6     check: str
7     severity: Severity
8     passed: bool
9     observed: float | None
```

Run all checks, collect the results, log them, **then** raise if any CRITICAL check failed. The log keeps every result, even on the run that stops.

Two coding rules make checks trustworthy:

- **Guard every division.** An empty input has no missing share. Return `None`, not a `ZeroDivisionError`.
- **Test `None` with `is None`.** `if share:` treats a real share of `0.0` like a missing one, and the check silently never runs.

Why it matters

A tracker QA step printed "WARNING: low base in 3 cells" to the notebook output and carried on. Nobody reads notebook output for scheduled jobs. The same step never checked for an empty input. The week a supplier delivered an empty file, the job computed shares from zero rows, divided by zero, and wrote `NaN` to every cell of the client dashboard.

Structured results in a table get alerts. A CRITICAL `min_rows` check stops the run before the empty file reaches the dashboard.

Worked example: weak code

```
1 def qa(df):
2     share = df["weight"].isna().sum() / len(df)
3     if share:
4         if share > 0.1:
```

```

3         print("WARNING: many missing weights")
6     if len(df) == 0:
4         print("empty file")

```

1. **Line 2.** Divides by zero on an empty file, before the emptiness check runs.
2. **Line 3.** Truthiness: a share of 0.0 skips the check.
3. **Line 5.** Prose on stdout. Nothing stores or alerts on it.
4. **Line 7.** An empty file is critical, but the run carries on.

Worked example: strong code

```

1  from dataclasses import dataclass
2  from typing import Literal
3
4  import pandas as pd
5
6  Severity = Literal["CRITICAL", "WARN"]
7
8
9  @dataclass(frozen=True)
10 class QaResult:
11     """One check outcome, stored in the run log."""
12
13     check: str
14     severity: Severity
15     passed: bool
16     observed: float | None
17
18
19 def run_checks(df: pd.DataFrame, min_rows: int, max_missing: float) ->
    list[QaResult]:
20     """Evaluate every input check. Never raises."""
21     n = len(df)
22     share = None if n == 0 else float(df["weight"].isna().sum()) / n
23     return [
24         QaResult("min_rows", "CRITICAL", n >= min_rows, float(n)),
25         QaResult("missing_weight", "WARN", share is not None and share <=
    max_missing, share),
26     ]
27
28
29 def enforce(results: list[QaResult]) -> None:
30     """Raise when any CRITICAL check failed."""
31     failed = [r.check for r in results if r.severity == "CRITICAL" and
    not r.passed]
32     if failed:
33         raise RuntimeError(f"critical checks failed: {failed}")

```

1. **Line 6.** A closed set of severities, checked by mypy.
2. **Line 22.** Guarded division: no share on an empty frame.
3. **Line 25.** is None, so a real 0.0 passes the check instead of skipping it.

4. **Line 33.** Raised after all checks ran, so the log has the full picture.

Check at boundaries. Collect structured results. Log them all, then stop on CRITICAL.

Recap card

- **Where to validate** At boundaries: after reading input, before writing output.
- **CRITICAL vs WARN** CRITICAL stops the run (raise). WARN records and continues.
- **Order of QA operations** Run every check, log all results, then raise if any CRITICAL failed.
- **Two coding rules for checks** Guard every division. Test for None with is None, never truthiness.

Practice

Exercise 1. Multiple choice (level 1)

A pipeline reads a supplier file, runs five pure transforms, then writes a client table. Where do validation checks belong?

- a) After every transform.
- b) After reading, and before writing.
- c) Only before writing.
- d) Nowhere, if the transforms have unit tests.

Why?

- r1) Those are the boundaries where data changes owner. Tests cover the pure middle.
- r2) Fewer checks run faster.

Exercise 2. Predict the output (level 2)

No weights are missing this week. What does this print?

```
1 def missing_check(share: float | None, limit: float) -> str:
2     if share:
3         return "pass" if share <= limit else "fail"
4     return "no data"
5
6
7 print(missing_check(0.0, 0.1))
```

Exercise 3. Put in order (level 2)

Order the steps of one pipeline run.

- Run input checks and collect results
- Write the client table
- Run output checks on the result
- Run the pure transforms
- Log all results, then raise if any CRITICAL failed
- Read the supplier file

Exercise 4. Write the function (level 3)

Write `run_checks(df, min_rows, max_missing) -> list[QaResult]` and `enforce(results) -> None`. `QaResult` and `Severity` are given.

- `min_rows`, CRITICAL: passes when the row count is at least `min_rows`. `observed` is the row count as a float.
- `missing_weight`, WARN: the share of null `weight` values. Passes when the share is at most `max_missing`. On an empty frame the share is `None` and the check fails. `observed` is the share.
- `run_checks` never raises. `enforce` raises `RuntimeError` naming each failed CRITICAL check, and ignores WARN failures.

Starter code:

```
1 from dataclasses import dataclass
2 from typing import Literal
3
4 import pandas as pd
5
6 Severity = Literal["CRITICAL", "WARN"]
7
8
9 @dataclass(frozen=True)
10 class QaResult:
11     """One check outcome."""
12
13     check: str
14     severity: Severity
15     passed: bool
16     observed: float | None
17
18
19 def run_checks(df, min_rows, max_missing):
20     ...
21
22
23 def enforce(results):
```


Exercise 5. Architecture design (level 4)

Assign a severity to each check on a weekly brand-tracking run.

Options: CRITICAL, WARN.

1. respondent_id is unique in the input
2. A brand's subgroup base is below 50
3. Weights sum to the sample size within 0.1%
4. Share of missing age is above last month's level

B6 Reproducibility

Roadmap: Prerequisite for all months. Seeds, config, data versioning and pinned dependencies.

B6-L1 Seeds and random generators

Concept

A result you cannot rerun is a result you cannot defend. Randomness enters analysis through sampling, train/test splits, bootstrap intervals, model initialisation and synthetic data. Control each source.

Create a generator and pass it in. NumPy 1.17 introduced

`numpy.random.Generator`. Examples use NumPy 2.4.

```
1 rng = np.random.default_rng(20260925)      # at the edge, once
2 sample = draw_sample(ids, k=500, rng=rng)    # the function receives rng
```

- `np.random.default_rng(seed)` gives the same stream of numbers for the same seed, on every run and machine with the same NumPy version.
- Functions take `rng: np.random.Generator` as a parameter. They never call `np.random.seed()` or the legacy `np.random.rand()`, which share one hidden global state across your whole program and every library you import.
- One generator, used twice, gives two **different** draws. That is correct: two samples should differ. Rerunning the whole script gives the same pair.
- Need independent streams for parallel workers? `rng.spawn(n)` returns `n` child generators.

Libraries take the seed or generator as a parameter too:

```
train_test_split(..., random_state=42) in scikit-learn, df.sample(n,
random_state=rng) in pandas.
```

Two hidden sources of non-determinism:

- `hash()` **on strings changes between Python processes** (hash randomisation, `PYTHONHASHSEED`). `hash(respondent_id) % 10` puts the same respondent in a different bucket on each run. Use `hashlib.sha256` for stable buckets.
- **Set iteration order** for strings also changes between runs. Sort before you rely on order.

Record the seed in the run log with the results, so anyone reruns the exact analysis.

Why it matters

A holdout set was defined as `hash(respondent_id) % 10 == 0`. The model scored AUC 0.81 in the notebook. The scheduled retraining job, a new Python process, put a different 10% in the holdout. Some of the old test respondents were now in training, and the reported AUC rose to 0.84 with no change to the model. The team spent a week chasing an improvement that did not exist.

A bucket from `hashlib.sha256` of the ID is the same in every process.

Worked example: weak code

```
1 import numpy as np
2
3 np.random.seed(1)
4
5
6 def draw_sample(ids, k):
7     return list(np.random.choice(ids, size=k, replace=False))
8
9
10 def holdout(ids):
11     return [i for i in ids if hash(i) % 10 == 0]
```

1. **Line 3.** Global state, set at import: any other call to `np.random` changes what `draw_sample` returns.
2. **Line 7.** Reads the hidden global generator.
3. **Line 11.** String `hash()` is randomised per process: a different holdout every run.

Worked example: strong code

```
1 import hashlib
2
3 import numpy as np
4
5
6 def draw_sample(ids: list[str], k: int, rng: np.random.Generator) -> list[str]:
7     """Draw k distinct IDs using the generator passed in."""
8     return [str(x) for x in rng.choice(ids, size=k, replace=False)]
9
10
11 def bucket(respondent_id: str, n_buckets: int) -> int:
12     """Stable bucket for an ID, the same in every process."""
13     digest = hashlib.sha256(respondent_id.encode("utf-8")).hexdigest()
14     return int(digest, 16) % n_buckets
15
16
17 def main(seed: int) -> None:
18     """The edge: create the generator once and log the seed."""
```

```
3 rng = np.random.default_rng(seed)
20 ...
```

1. **Line 6.** The randomness is a parameter: tests pass a seeded generator.
2. **Line 13.** A cryptographic hash is stable across processes and machines.
3. **Line 19.** One generator, created at the edge from a logged seed.

Seed once at the edge, pass the generator in, and never use `hash()` for anything that must be stable.

Recap card

- **Reproducible randomness in NumPy 2** `rng = np.random.default_rng(seed)`, created once at the edge and passed in.
- **Why not `np.random.seed`?** It seeds one global generator that every library shares and advances.
- **Stable bucket for an ID** `int(hashlib.sha256(id.encode()).hexdigest(), 16) % n`. Never `hash()`, which is randomised per process.
- **Independent streams for workers** `rng.spawn(n)`.

Practice

Exercise 1. Multiple choice (level 1)

Which is the recommended way to get reproducible random numbers in NumPy 2?

- a) `np.random.seed(42)` at the top of the script
- b) `rng = np.random.default_rng(42)`, passed to the functions that need it
- c) `random.seed(42)` from the standard library
- d) No seed. Results average out.

Why?

- r1) Each Generator owns its state, so no other code can change the stream it gives your function.
- r2) `np.random.seed` was removed in NumPy 2.

Exercise 2. Predict the output (level 2)

What does this print? (NumPy 2.4)

```
1 import numpy as np
2
3 rng = np.random.default_rng(1)
4 first = rng.choice(10, size=3, replace=False).tolist()
5 second = rng.choice(10, size=3, replace=False).tolist()
6 print(first == second)
```

Exercise 3. Fill the gap (level 2)

Create a seeded generator and give four parallel workers independent streams.

```
1 rng = np.random.____(1)(20260925)
2 worker_rngs = rng.____(2)(4)
```

Exercise 4. Write the function (level 3)

Write `bucket(respondent_id: str, n_buckets: int) -> int`. It returns a number from 0 to `n_buckets - 1` that is the same for the same ID in every Python process. Use `hashlib.sha256` on the UTF-8 bytes of the ID, read the hex digest as an integer, and take it modulo `n_buckets`. Do not use `hash()`.

Starter code:

```
1 def bucket(respondent_id, n_buckets):
2     return hash(respondent_id) % n_buckets
```

Exercise 5. Spot the bug (level 4)

Two analysts run the same bootstrap with seed 7 and get different confidence intervals. One of them imports a plotting helper that draws random jitter with `np.random`. Tap the line that makes the result depend on other code, and fix the function.

```
1 import numpy as np
2
3
4 def bootstrap_means(values: list[float], n_boot: int, seed: int) ->
    list[float]:
5     """Bootstrap means of values."""
6     np.random.seed(seed)
7     out = []
8     for _ in range(n_boot):
9         draw = np.random.choice(values, size=len(values), replace=True)
10        out.append(float(draw.mean()))
11    return out
```

Write down the line number of the bug. Then write the fix.

Sources

- `np.random.default_rng(seed)` returns a Generator; `Generator.spawn(n)` returns `n` independent child generators. (verified: NumPy random documentation; run on NumPy 2.4.6 in this repository.)
- Python's `hash()` of a `str` differs between processes unless `PYTHONHASHSEED` is fixed. (verified: Python data model documentation (object.**hash**); observed across two processes in this repository.)

B6-L2 Record the data, config and code behind every result

Concept

To rerun a result you need three things exactly as they were: the **data**, the **config** and the **code**. Record all three with every run, in a small **run manifest**.

```
1 manifest = {
2     "run_id": run_id,
3     "started_utc": datetime.now(timezone.utc).isoformat(),
4     "input_sha256": file_sha256(input_path),
5     "config_sha256": config_sha256(config),
6     "git_commit": git_commit,
7     "seed": seed,
8     "packages": {"pandas": version("pandas"), "numpy": version("numpy")},
9 }
```

Data. A file name and a date do not identify content. Suppliers resend files under the same name. A **content hash** does: `hashlib.sha256` over the file bytes changes if any byte changes. `hashlib.file_digest(f, "sha256")` (Python 3.11 and later) reads the file in chunks, so large files do not load into memory. Keep raw inputs immutable: write each delivery to a new path and never overwrite it. For lakehouse tables, record the Delta table version instead (lesson A3-L3).

Config. Hash the config you loaded, not the file on disk, so defaults count. Serialise with `json.dumps(config, sort_keys=True)` first. Without `sort_keys`, the same settings in a different key order give a different hash.

Code. Record the Git commit. Refuse to produce client output from a working tree with uncommitted changes, or record that it was dirty.

Environment. `importlib.metadata.version("pandas")` returns the installed version. The lock file (lesson B6-L3) covers the rest.

Store the manifest next to the output, or as a row in a run log table. When someone asks "why did this number change?", compare two manifests. The field that differs is where to look.

Why it matters

A client queried a brand share that moved 3 points between two versions of the same wave's report. The analysis folder held `wave12.csv`, the config and the notebook, all edited since. Nobody knew whether the supplier had resent

the file, a setting had changed, or the code had. Reconstructing the answer took four days.

With manifests, the two runs differed in one field: `input_sha256`. The supplier had resent the file with 212 late completes.

Worked example: weak code

```
❶ df = pd.read_csv("data/wave12.csv")
  results = analyse(df, CONFIG)
❷ results.to_csv(f"out/wave12_{date.today()}.csv")
```

1. **Line 1.** The name says nothing about the content. A resent file has the same name.
2. **Line 3.** A date, but no record of the input, config, code or seed behind it.

Worked example: strong code

```
1 import hashlib
2 import json
3 from pathlib import Path
4
5
6 def file_sha256(path: Path) -> str:
7     """SHA-256 of a file's bytes, read in chunks."""
8     with path.open("rb") as f:
9         return hashlib.file_digest(f, "sha256").hexdigest()
10
11
12 def config_sha256(config: dict[str, object]) -> str:
13     """SHA-256 of a config, independent of key order."""
14     text = json.dumps(config, sort_keys=True)
15     return hashlib.sha256(text.encode("utf-8")).hexdigest()
```

1. **Line 9.** Content identity: any changed byte changes the hash.
2. **Line 14.** Stable serialisation: key order does not change the hash.

Hash the input and the loaded config, record the commit and the seed, and store the manifest with the output.

Recap card

- **Three things to record for a rerun** Data (content hash or table version), config (hash of the loaded config), code (Git commit).
- **Hash a large file** with `path.open("rb") as f: hashlib.file_digest(f, "sha256")` (Python 3.11+).
- **Stable config hash** `sha256 of json.dumps(config, sort_keys=True).encode("utf-8")`.
- **Raw inputs** Immutable: each delivery at a new path, never overwritten.

Practice

Exercise 1. Multiple choice (level 1)

Which value best identifies the exact input file a result came from?

- a) The file name
- b) The file's modified time
- c) The SHA-256 of the file's bytes
- d) The row count

Why?

- r1) Any change to the content changes the hash, and identical content gives the same hash everywhere.
- r2) Hashes are shorter than file names.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 import json
2
3 print(json.dumps({"min_base": 50, "cap": 5.0}, sort_keys=True))
```

Exercise 3. Fill the gap (level 2)

Hash a large file without loading it into memory.

```
1 with path.open("____(1)") as f:
2     digest = hashlib.____(2)(f, "sha256").hexdigest()
```

Exercise 4. Write the function (level 3)

Write two functions.

- `file_sha256(path: Path) -> str` returns the hex SHA-256 of the file's bytes.
- `config_sha256(config: dict[str, object]) -> str` returns the hex SHA-256 of `json.dumps(config, sort_keys=True)` encoded as UTF-8, so key order does not matter.

Starter code:

```
1 import hashlib
2 import json
3 from pathlib import Path
4
5
6 def file_sha256(path):
7     ...
8
```



```

9
10 def config_sha256(config):
11     return hashlib.sha256(str(config).encode()).hexdigest()

```

Exercise 5. Architecture design (level 4)

Choose what the run manifest records for a monthly client report built from a supplier CSV.

Options: SHA-256 of the input file, with the raw file kept unchanged, The input file name.

1. Input data
2. Settings
3. Code
4. Where the manifest lives

Sources

- `hashlib.file_digest` was added in Python 3.11. (verified: Python hashlib documentation; run on Python 3.11 in this repository.)

B6-L3 Pinned dependencies and lock files

Concept

`pyproject.toml` states what your project **accepts**: `pandas>=3.0,<4`. A **lock file** records what was **installed**: `pandas 3.0.6`, `numpy 2.4.6`, and every indirect dependency, each with a file hash. You need both.

- The ranges in `pyproject.toml` say which versions your code supports. Keep them honest: a lower bound you have tested, an upper bound at the next major version when the library breaks things across majors.
- The lock file makes every install identical: your laptop, a colleague's, CI and the scheduled job. Commit it.

With **uv** (examples use `uv 0.8`):

```

1 uv lock                                # resolve pyproject.toml ranges, write uv.lock
2 uv sync --locked                        # install exactly uv.lock; fail if it is out of
   date
3 uv lock --upgrade-package pandas      # move one package on purpose

```

`uv sync --locked` exits with an error when `pyproject.toml` has changed since the lock was written: "The lockfile at `uv.lock` needs to be updated, but `--locked` was provided." Use it in CI, so a dependency change without a new lock fails the build instead of resolving something new on the runner.

Without `uv`, `pip-compile --generate-hashes` (`pip-tools`) writes a fully pinned `requirements.txt` with hashes, and `pip install --require-hashes -r requirements.txt` refuses any file whose hash differs.

`pip freeze > requirements.txt` is a weaker substitute. It lists what happens to be in your environment, including packages you installed by hand for a one-off test, with no hashes and no record of why each is there.

Upgrade on purpose: one dependency at a time, on a branch, with tests. Record the reason in the commit message. Never let a fresh install pick up new versions silently.

Why it matters

A forecasting job ran on a fresh container every night. `requirements.txt` said `pandas>=2.0`. The night `pandas 3.0` was released, the container installed it. The job's text columns switched to the new `str` dtype, a `.str.contains` call behaved differently on missing values, and the job filtered out 6% of rows. Nobody had changed a line of code.

A committed lock file, installed with `uv sync --locked`, keeps the old `pandas` until someone upgrades it deliberately, with tests.

Worked example: weak code

```
1 # requirements.txt
❶ pandas
❷ scikit-learn>=1.0
4 numpy
❸ jupyter
```

1. **Line 2.** No bounds: any future major version installs on the next fresh build.
2. **Line 3.** A lower bound only. Every build resolves whatever is newest that day.
3. **Line 5.** A development tool mixed into the runtime requirements.

Worked example: strong code

```
1 # pyproject.toml
2 [project]
3 requires-python = ">=3.12"
❶ dependencies = ["pandas>=3.0,<4", "scikit-learn>=1.8,<2", "numpy>=2.4,<3"]
5
6 [dependency-groups]
❷ dev = ["pytest>=9", "ruff>=0.13", "jupyterlab>=4"]
8
9 # CI
❸ uv sync --locked --group dev
```

1. **Line 4.** Honest ranges: tested lower bounds, next major excluded.

- 2. **Line 7.** Development tools in their own group, not in the runtime install.
- 3. **Line 10.** Installs exactly uv.lock. A stale lock fails the build.

Ranges in pyproject.toml, exact versions and hashes in a committed lock file, and --locked in CI.

Recap card

- **pyproject.toml vs lock file** pyproject.toml: ranges you support. Lock: exact versions and hashes of everything installed.
- **uv sync --locked** Installs exactly uv.lock and fails if pyproject.toml changed without a new lock.
- **Hashes without uv** pip-compile --generate-hashes, then pip install --require-hashes -r requirements.txt.
- **Why not pip freeze?** It captures one environment's accidents, with no hashes and no reasons.

Practice

Exercise 1. Multiple choice (level 1)

What does a lock file add that `pandas>=3.0,<4` in pyproject.toml does not?

- a) The exact version of every direct and indirect dependency, with file hashes.
- b) A list of packages your code imports.
- c) Faster imports at runtime.
- d) Nothing. They hold the same information.

Why?

- r1) Every install then gets identical packages, on every machine.
- r2) It lets pip skip dependency resolution entirely on first install.

Exercise 2. Predict the output (level 2)

You add `polars>=1.30` to `dependencies` in pyproject.toml and push without running `uv lock`. CI runs `uv sync --locked`. What happens?

```
1 $ uv sync --locked
```

Exercise 3. Fill the gap (level 2)

Write a pinned requirements file with hashes using pip-tools, then install it so any mismatched file is refused.

```
1 pip-compile --____(1) requirements.in
2 python -m pip install --____(2) -r requirements.txt
```

Exercise 4. Put in order (level 3)

Order the steps to upgrade scikit-learn in a locked project.

- Commit pyproject.toml and uv.lock with the reason
- Create a branch
- `uv lock --upgrade-package scikit-learn`
- Raise the bound in pyproject.toml if the new version is outside the range
- Run the test suite and compare key outputs with the last release

Exercise 5. Architecture design (level 4)

Set up dependency management for a nightly scoring job.

Options: Tested lower bounds, next major excluded, in pyproject.toml, Package names only, no versions, Exact == pins in pyproject.toml, no lock.

1. Declared dependencies
2. Exact versions
3. Install in CI and the job

Sources

- `uv sync --locked` fails when uv.lock needs updating, with the message quoted in this lesson. (verified: Run with uv 0.8.17 in this repository's scratch directory.)
- `pip-compile --generate-hashes` and `pip install --require-hashes`. (search-only: pip-tools and pip documentation (not run here).)
- `Series.str.contains` returns `None` for a missing value on pandas 2.3 (object dtype) and `False` on pandas 3.0 (str dtype). (verified: Run on pandas 2.3.3 and 3.0.6 in this repository.)

B7 Exploratory analysis done well

Roadmap: Prerequisite for all months. Profiling, visual checks and documented assumptions.

B7-L1 Profile a dataset before you analyse it

Concept

Before any chart or model, answer five questions about a new dataset. Each takes one line of pandas, and each catches a class of error that no later step will.

1. **Shape and grain.** How many rows, and what does one row mean? Check the claimed key is unique: `df.duplicated(subset=["respondent_id", "wave"]).sum()` should be 0.
2. **Dtypes.** `df.dtypes`. IDs read as floats, dates as text and numbers as `str` all show here (lesson B3-L1).
3. **Missingness.** `df.isna().mean()` gives the share missing per column. Look for columns that are nearly empty, and for shares that differ by wave or region.
4. **Cardinality.** `df.nunique()`. A `gender` column with 14 distinct values holds typos. An ID column with fewer distinct values than rows holds duplicates.
5. **Ranges.** `df.describe()`. Ages of 999, negative weights, and dates in 1900 are codes for "missing" or errors.

Put the per-column answers in one table:

```
1 def profile(df: pd.DataFrame) -> pd.DataFrame:
2     return pd.DataFrame({
3         "dtype": df.dtypes.astype(str),
4         "missing_share": df.isna().mean(),
5         "n_unique": df.nunique(),
6     })
```

Two pandas defaults to know (pandas 3.0):

- `nunique()` **excludes** missing values. Pass `dropna=False` to count NaN as a value.
- `duplicated()` marks every repeat **except the first** (`keep="first"`). To count every row involved in a duplicate, use `keep=False`.

Save the profile with the analysis. When the next delivery arrives, compare the two profiles. A column whose missing share jumps from 2% to 40% tells you something changed at the source before it changes your results.

Libraries such as ydata-profiling build long HTML reports. They are useful for a first look. The five checks above, in code, are what you rerun on every delivery.

Why it matters

A satisfaction score was coded 1 to 10, with 99 for "don't know". Nobody profiled the file. The mean score for one region came out at 14.2, above the maximum possible value. The analyst assumed a weighting problem and spent a day on the weights. `describe()` would have shown a maximum of 99 in the first minute.

Worked example: weak code

```
1 df = pd.read_csv("wave12.csv")
2 df.head()
3 df.groupby("region")["score"].mean()
```

1. **Line 2.** Five rows show nothing about codes like 99, duplicates or missingness.
2. **Line 3.** Analysis before profiling: 99 codes are averaged in.

Worked example: strong code

```
1 import pandas as pd
2
3
4 def profile(df: pd.DataFrame) -> pd.DataFrame:
5     """One row per column: dtype, missing share and distinct values."""
6     return pd.DataFrame(
7         {
8             "dtype": df.dtypes.astype(str),
9             "missing_share": df.isna().mean(),
10            "n_unique": df.nunique(),
11        }
12    )
13
14
15 def duplicate_rows(df: pd.DataFrame, key: list[str]) -> int:
16     """Rows whose key appears more than once, all copies counted."""
17     return int(df.duplicated(subset=key, keep=False).sum())
```

1. **Line 9.** Share missing per column, comparable across deliveries.
2. **Line 17.** `keep=False` counts every copy, not only the repeats.

Shape and grain, dtypes, missingness, cardinality and ranges, in code, before any analysis.

Recap card

- **Five profiling questions** Shape and grain, dtypes, missingness, cardinality, ranges.
- **nunique() and missing values** Excluded by default. dropna=False counts NaN as a value.
- **Count every row in a duplicate group** df.duplicated(subset=key, keep=False).sum().
- **Why save the profile?** Compare it with the next delivery to spot source changes early.

Practice

Exercise 1. Multiple choice (level 1)

A satisfaction score should run from 1 to 10. Which call shows at once that some rows hold 99?

- a) df.head()
- b) df["score"].describe()
- c) df.dtypes
- d) df.shape

Why?

- r1) describe() reports min and max, so a max of 99 stands out.
- r2) describe() removes invalid values.

Exercise 2. Predict the output (level 2)

What does this print? (pandas 3.0)

```
1 import pandas as pd
2
3 ids = pd.Series(["a", "b", "b", None])
4 print(ids.nunique(), ids.duplicated().sum())
```

Exercise 3. Fill the gap (level 2)

Count every row that shares a (respondent_id, wave) key with another row, and the share of missing values per column.

```
1 dupes = df.duplicated(subset=["respondent_id", "wave"],
    keep=____(1)).sum()
2 missing = df.isna().____(2)()
```

Exercise 4. Write the function (level 3)

Write `profile(df: pd.DataFrame) -> pd.DataFrame`. Return one row per input column, indexed by column name, with columns:

- `dtype`: the dtype as a string.
- `missing_share`: share of missing values.
- `n_unique`: distinct non-missing values.
- `min` and `max`: for numeric columns, the minimum and maximum. For other columns, missing (`None` or `NaN`).

Starter code:

```
1 import pandas as pd
2
3
4 def profile(df):
5     return df.describe()
```

Exercise 5. Spot the bug (level 4)

The QA report says 1 duplicated respondent row, but 2 rows in the file share the same ID, and both must be sent back to the supplier. Tap the line with the bug and fix it.

```
1 import pandas as pd
2
3
4 def rows_with_duplicate_ids(df: pd.DataFrame) -> pd.DataFrame:
5     """Every row whose respondent_id appears more than once."""
6     mask = df.duplicated(subset=["respondent_id"])
7     return df[mask]
```

Write down the line number of the bug. Then write the fix.

Sources

- `pandas.nunique()` excludes missing values by default and `duplicated()` defaults to `keep="first"`. (verified: pandas API reference; run on pandas 3.0.6 and 3.0.2 in this repository.)

B7-L2 Plot before you summarise

Concept

Summary statistics hide shape. **Anscombe's quartet** (1973) is four small datasets with the same means, variances and correlation. Computed in this course's build for the first three:

Set	mean of y	variance of y	correlation with x
I	7.50	4.13	0.816
II	7.50	4.13	0.816
III	7.50	4.12	0.816

Plotted, set I is a noisy line, set II is a curve, and set III is a perfect line with one outlier. The numbers alone cannot tell them apart.

Four plots catch most problems:

- **Histogram** of each key numeric column. Shows skew, gaps, spikes at codes such as 99 or 0, and more than one peak.
- **Scatter plot** of each pair you will model or correlate. Shows curves, clusters and single points that drive a correlation.
- **Line over time** of counts and key rates by wave. Shows breaks in the series when a questionnaire or supplier changed.
- **Bar chart of missingness** by column and by subgroup.

Make each plot honest:

- Label both axes with units, and give every figure a title that states what it shows.
- Choose histogram bins on purpose. Too few hide a spike at 99. Try two or three bin counts.
- Use a log scale for heavily skewed values such as spend, and say so on the axis.
- Plot the raw data before you plot the smoothed or aggregated version.

Write plotting code as functions that return a `matplotlib.figure.Figure` (matplotlib 3.10). Build with `fig, ax = plt.subplots()`, not the global `plt.plot()` state. A function that returns a figure is testable: a test reads `ax.get_xlabel()` or counts the bars. The caller decides whether to show it or save it.

Why it matters

A team reported a correlation of 0.62 between ad recall and purchase intent across 40 markets, and the client planned a budget shift on it. A scatter plot showed 39 markets in a flat cloud and one market far in the top-right corner. Without that market the correlation was close to zero. The decision was about one market, not 40.

Worked example: weak code

```
❶ print(df["spend"].mean(), df["spend"].std())
❷ plt.hist(df["spend"])
❸ plt.show()
```

1. **Line 1.** Summaries first: a skewed spend column makes the mean misleading.
2. **Line 2.** Global pyplot state, default bins, no labels, no title.
3. **Line 3.** Nothing returned: nothing to test or save.

Worked example: strong code

```
1 import matplotlib.pyplot as plt
2 from matplotlib.figure import Figure
3
4
5 def plot_distribution(values: list[float], title: str, xlabel: str, bins:
    int = 30) -> Figure:
6     """Histogram of values with labelled axes. The caller shows or saves
    it."""
❶     fig, ax = plt.subplots(figsize=(6, 4))
❷     ax.hist(values, bins=bins)
9     ax.set_title(title)
10    ax.set_xlabel(xlabel)
11    ax.set_ylabel("Respondents")
❸    return fig
```

1. **Line 7.** An explicit figure and axes, not the global pyplot state.
2. **Line 8.** Bins are a parameter, so you try several.
3. **Line 12.** Returned for the caller to show, save or test.

Plot distributions, pairs and time series before summarising, with functions that return labelled figures.

Recap card

- **Anscombe's quartet** Same mean, variance and correlation (0.816), different shapes. Plot before summarising.
- **Four first plots** Histograms, scatter plots of modelled pairs, lines over waves, missingness bars.
- **Testable plotting code** fig, ax = plt.subplots(), set labels on ax, return fig. No plt.show() inside.
- **Histogram bins** Choose on purpose and try several. Too few hide spikes at codes.

Practice

Exercise 1. Multiple choice (level 1)

Two variables have correlation 0.816 in four different datasets. What can you conclude about the relationship without plotting?

- a) It is linear in all four.
- b) It is the same shape in all four.
- c) Very little about its shape.
- d) An outlier is present in all four.

Why?

- r1) Anscombe's quartet has identical correlations with a line, a curve and an outlier-driven fit.
- r2) Correlation is always misleading.

Exercise 2. Predict the output (level 2)

What does this print? (NumPy 2.4)

```
1 import numpy as np
2
3 x = [10, 8, 13, 9, 11, 14, 6, 4, 12, 7, 5]
4 y1 = [8.04, 6.95, 7.58, 8.81, 8.33, 9.96, 7.24, 4.26, 10.84, 4.82, 5.68]
5 y3 = [7.46, 6.77, 12.74, 7.11, 7.81, 8.84, 6.08, 5.39, 8.15, 6.42, 5.73]
6 print(round(np.corrcoef(x, y1)[0, 1], 2) == round(np.corrcoef(x, y3)[0,
    1], 2))
```

Exercise 3. Fill the gap (level 2)

Create an explicit figure and axes, and label the x axis.

```
1 fig, ax = plt.__(1)()
2 ax.hist(values, bins=30)
3 ax.__(2)("Monthly spend (GBP)")
```

Exercise 4. Write the function (level 3)

Write `plot_distribution(values, title, xlabel, bins=30)` -> Figure. Draw a histogram with `bins` bins on one axes, set the title and x label from the arguments, set the y label to "Respondents", and return the figure. Do not call `plt.show()`.

Starter code:

```
1 import matplotlib.pyplot as plt
2
3
4 def plot_distribution(values, title, xlabel, bins=30):
5     plt.hist(values)
```

Exercise 5. Architecture design (level 4)

Choose the first plot for each question.

Options: Histogram, Scatter plot, Line over waves, Bar chart of missing share.

1. Does the age column hold codes such as 999?
2. Is the recall and intent correlation driven by a few markets?
3. Did the supplier change in wave 9 shift the awareness rate?
4. Which questions did the new mobile layout leave unanswered?

Sources

- Anscombe's sets I to III share mean 7.50, variance about 4.12 to 4.13 and correlation 0.816. (verified: Anscombe (1973), The American Statistician 27(1); computed with NumPy 2.4.6 in this repository.)

B7-L3 Write assumptions down, then turn them into checks

Concept

Every analysis rests on assumptions: what a row means, how "don't know" is coded, which respondents are in scope, that weights sum to the sample size. Unwritten, they live in one analyst's head and break silently when the data changes.

Keep an **assumptions log** next to the analysis. One entry per assumption, with the same fields each time:

```
1 - id: A3
2 statement: "score uses 1 to 10; 98 = refused, 99 = don't know"
3 evidence: "Questionnaire v4, Q12; profile of waves 1 to 11 shows only
  1-10, 98, 99"
4 handling: "98 and 99 become missing before any statistic"
5 check: "score in 1..10 or missing, after recoding"
6 owner: "analysis team"
7 last_verified: 2026-09-01
```

Then make the assumption **executable**. The `handling` becomes code, and the `check` becomes a QA result (lesson B5-L3) that runs on every delivery.

```
1 MISSING_CODES: dict[str, list[int]] = {"score": [98, 99]}
  # from config, per A3
2
```

```

3
4 def recode_missing(df: pd.DataFrame, codes: dict[str, list[int]]) ->
    pd.DataFrame:
5     return df.assign(**{col: df[col].mask(df[col].isin(vals)) for col,
        vals in codes.items()})

```

Recode **per column**. `df.replace(99, np.nan)` on the whole frame also blanks a real age of 99. `Series.mask(cond)` sets values to missing where the condition holds and leaves the rest.

Three kinds of assumption to write down every time:

- **Structure.** Grain, keys, which columns are required.
- **Coding.** Missing codes, scales, units, category lists.
- **Scope.** Who is included: age limits, completes only, which markets and waves.

When a check fails, the log tells you which belief was wrong and who to ask. When a client asks "why is this respondent excluded?", the answer has an ID.

Why it matters

An analyst knew that 99 meant "don't know" on the satisfaction scale and removed those rows by hand in a notebook. The analyst left. The next wave went through the scripted pipeline, which did not know. The regional mean rose from 7.1 to 11.4. It was published, because 11.4 looked like a large improvement rather than an impossible value on a 1 to 10 scale.

A logged assumption with an executable check would have failed the run on the first 99.

Worked example: weak code

```

❶ # 99 = don't know, drop them (ask Sam)
❷ df = df[df["score"] != 99]
❸ df = df.replace(98, np.nan)

```

1. **Line 1.** An assumption in a comment, tied to one person, with no check.
2. **Line 2.** Drops whole rows, so other answers from those respondents are lost too.
3. **Line 3.** Whole-frame replace: any column with a real 98 is blanked.

Worked example: strong code

```

1 import pandas as pd
2
3
4 def recode_missing(df: pd.DataFrame, codes: dict[str, list[int]]) ->
    pd.DataFrame:

```

```

5      """Set listed codes to missing, per column (assumption A3). Other
      columns untouched."""
6  ①      return df.assign(**{col: df[col].mask(df[col].isin(vals)) for col,
      vals in codes.items()})
7
8
9  def score_in_range(df: pd.DataFrame, low: int, high: int) -> bool:
10     """Check for A3: every non-missing score lies in [low, high]."""
11     s = df["score"].dropna()
12  ②      return bool(s.between(low, high).all())

```

1. **Line 6.** Per-column recoding from config. The respondent keeps their other answers.
2. **Line 12.** The assumption as a check, run on every delivery.

Log each assumption with evidence and owner, implement its handling per column, and check it on every run.

Recap card

- **Fields of an assumptions log entry** id, statement, evidence, handling, check, owner, last_verified.
- **Recode missing codes safely** Per column:
df[col].mask(df[col].isin(codes)). Never a whole-frame replace.
- **Three kinds of assumption** Structure (grain, keys), coding (codes, scales, units), scope (who is included).
- **From assumption to code** Handling in config-driven code, and a QA check that runs on every delivery.

Practice

Exercise 1. Multiple choice (level 1)

Where should the assumption "98 and 99 mean refused and don't know on score" live?

- a) In a code comment where the rows are dropped.
- b) In the assumptions log, in config for the handling, and in a QA check.
- c) In the analyst's notes.
- d) Nowhere: everyone knows the convention.

Why?

- r1) Written, applied and checked in code, it survives staff changes and new deliveries.
- r2) Logs are required by law.

Exercise 2. Predict the output (level 2)

What does this print? (pandas 3.0)

```
1 import numpy as np
2 import pandas as pd
3
4 df = pd.DataFrame({"score": [3, 99, 7], "age": [99, 40, 30]})
5 print(df.replace(99, np.nan)["age"].tolist())
```

Exercise 3. Put in order (level 2)

Order the life of one assumption.

- Observe the coding in the questionnaire and the profile
- Implement the handling in code, driven by config
- Write the assumption, evidence and owner in the log
- Update last_verified when the check or the source is reviewed
- Add a QA check that runs on every delivery

Exercise 4. Write the function (level 3)

Write `recode_missing(df, codes) -> pd.DataFrame`. `codes` maps a column name to a list of values that mean "missing" in that column. Return a new frame where those values are missing in those columns only. Leave other columns and the input frame unchanged.

Starter code:

```
1 import numpy as np
2 import pandas as pd
3
4
5 def recode_missing(df, codes):
6     for vals in codes.values():
7         df = df.replace(vals, np.nan)
8     return df
```

Exercise 5. Architecture design (level 4)

Classify each assumption and pick its check.

Options: Structure: check key uniqueness, Coding: check values against the codebook, Scope: check inclusion rules.

1. One row per respondent per wave
2. region uses codes 1 to 12 from the lookup table
3. Only respondents aged 18 or over are reported
4. Only completed interviews are in the file

Sources

- `DataFrame.replace(99, np.nan)` changes every column holding 99; `Series.mask(cond)` sets values where `cond` is `True` to missing. (verified: pandas API reference; run on pandas 3.0.6 in this repository.)

B8 Statistical correctness

Roadmap: Prerequisite for all months. Survey weights, confidence intervals, effect sizes and multiple comparisons.

B8-L1 Survey weights and the effective sample size

Concept

A survey sample rarely matches its population. If young adults are 20% of the population but 10% of your completes, each young respondent gets a larger **weight** so the weighted sample matches the population.

The **weighted mean** divides by the sum of the weights, not by the row count:

```
1 weighted_mean = (y * w).sum() / w.sum()      # same as np.average(y,  
           weights=w)
```

Weights have a cost. Unequal weights make estimates less precise than the raw sample size suggests. Kish's **effective sample size** measures that cost:

```
1 n_eff = w.sum() ** 2 / (w ** 2).sum()  
2 design_effect = n / n_eff
```

Equal weights give $n_{\text{eff}} = n$. A sample of 5 with weights 0.5, 0.5, 1, 2 and 4 has $n_{\text{eff}} = 2.98$: it carries about as much information as three unweighted respondents. The design effect is 1.68.

Rules that prevent the common mistakes:

- **Use weights for population estimates:** shares, means and totals you report as "UK adults".
- **Use n_{eff} , not n , for precision:** confidence intervals, significance tests and minimum base rules (lesson B8-L2).
- **Report the unweighted base** next to every weighted figure, so readers see how many people stand behind it.
- **Drop rows where the value or the weight is missing** before you compute. A missing weight is not zero, and a missing answer is not a zero answer.
- **Guard the division.** An empty subgroup or a zero weight total has no mean. Return `None`.

Check weights before using them: all positive, none missing, and a sum equal to the sample size or the population total, whichever the weighting

scheme promises. Extreme weights (above about 5 times the mean) are a warning sign. Trimming them is a design decision, logged as an assumption (lesson B7-L3).

Why it matters

A brand tracker reported that awareness among 18 to 24s rose from 41% to 52%, "significant at 95%". The test used the 180 raw respondents in the group. Their weights were uneven, and the effective sample size was 71. With 71, the change was well inside the margin of error. The client moved budget on a result the data did not support.

Worked example: weak code

```
1 share = (df["aware"] * df["weight"]).mean()
2 se = (share * (1 - share) / len(df)) ** 0.5
```

1. **Line 1.** Divides by the row count, not the weight total. Only right if weights average exactly 1.
2. **Line 2.** Uses raw n, overstating precision whenever weights vary.

Worked example: strong code

```
1 import pandas as pd
2
3
4 def weighted_share(df: pd.DataFrame, value: str, weight: str) -> float |
  None:
5     """Weighted mean of value, ignoring rows with a missing value or
      weight."""
6     d = df[[value, weight]].dropna()
7     total = d[weight].sum()
8     if len(d) == 0 or total == 0:
9         return None
10    return float((d[value] * d[weight]).sum() / total)
11
12
13 def effective_n(weights: pd.Series) -> float | None:
14     """Kish effective sample size: (sum w)^2 / sum(w^2)."""
15     w = weights.dropna()
16     sum_sq = (w**2).sum()
17     if sum_sq == 0:
18         return None
19    return float(w.sum() ** 2 / sum_sq)
```

1. **Line 6.** Rows missing either the answer or the weight do not count.
2. **Line 10.** Divides by the weight total.
3. **Line 19.** The n to use for precision.

Weighted estimates divide by the weight total. Precision uses the effective sample size.

Recap card

- **Weighted mean** $\text{sum}(y * w) / \text{sum}(w)$, over rows with both present. Same as `np.average(y, weights=w)`.
- **Kish effective sample size** $(\text{sum } w)^2 / \text{sum}(w^2)$. Design effect = n / n_{eff} .
- **Which n for intervals and suppression?** n_{eff} . Report the raw base as well.
- **Missing weight** Not zero. Drop the row from weighted estimates, and investigate.

Practice

Exercise 1. Multiple choice (level 1)

A subgroup has 180 respondents and an effective sample size of 71. Which number belongs in the standard error of its weighted share?

- a) 180
- b) 71
- c) The sum of the weights
- d) 180 minus 71

Why?

- r1) Unequal weights reduce precision, and n_{eff} measures the information that remains.
- r2) Smaller numbers are always safer.

Exercise 2. Predict the output (level 2)

What does this print? (NumPy 2.4)

```
1 import numpy as np
2
3 y = np.array([1, 0, 1, 0, 1])
4 w = np.array([0.5, 0.5, 1.0, 2.0, 4.0])
5 print(np.average(y, weights=w), y.mean())
```

Exercise 3. Fill the gap (level 2)

Compute Kish's effective sample size and the design effect.

```
1 n_eff = w.sum() ** 2 / (w ** ____ (1)).sum()
2 deff = ____ (2) / n_eff
```

Exercise 4. Write the function (level 3)

Write two functions.

- `weighted_share(df, value, weight) -> float | None`: the weighted mean of `value`, ignoring rows where the value or the weight is missing. Return `None` when no rows remain or the weight total is 0.
- `effective_n(weights: pd.Series) -> float | None`: Kish's effective sample size, ignoring missing weights. Return `None` when the sum of squared weights is 0.

Starter code:

```
1 import pandas as pd
2
3
4 def weighted_share(df, value, weight):
5     return (df[value] * df[weight]).mean()
6
7
8 def effective_n(weights):
9     return len(weights)
```

Exercise 5. Architecture design (level 4)

Choose the right quantity for each part of a weighted subgroup report.

Options: Weighted estimate, Unweighted estimate, Raw base (n), Effective sample size.

1. The share shown in the chart
2. The n in the confidence interval
3. The base printed under the chart
4. The number compared with the minimum reporting base

Sources

- Kish effective sample size is $(\sum w)^2 / \sum(w^2)$. (verified: Kish (1965), Survey Sampling; computed with NumPy 2.4.6 in this repository.)

B8-L2 Confidence intervals and effect sizes

Concept

A share without an interval is a guess with no stated size. Report every estimate with a **confidence interval**, and every difference with an **effect size**.

Intervals for a proportion. The textbook Wald interval, $p \pm 1.96 * \sqrt{p(1-p)/n}$, fails at the edges. For 0 out of 20 it gives `[0, 0]`, as if you were certain. The **Wilson** interval behaves well for small `n` and for shares near 0 or 1. For 0 out of 20 it gives `[0, 0.161]`. In statsmodels 0.14:

```
1 from statsmodels.stats.proportion import proportion_confint
2 low, high = proportion_confint(count=40, nobs=100, alpha=0.05,
    method="wilson") # (0.309, 0.498)
```

For a **weighted** share, use the weighted estimate and the effective sample size from lesson B8-L1 in the same formula. The 18 to 24 group from that lesson, with a weighted share of 52% and `n_eff = 71`, has a Wilson interval of 40.6% to 63.2%. With the raw 180 it would be 44.7% to 59.2%, which looks far more precise than the data is.

Effect sizes. A p-value answers "is there any difference?", and with a large sample the answer is nearly always yes. Two groups of 100,000 with shares of 50.5% and 50.0% give $p = 0.025$, and the difference is too small to matter. State the size:

- The **difference in shares**, in percentage points, with its interval. This is what a client reads.
- **Cohen's h**, $2 * \arcsin(\sqrt{p_1}) - 2 * \arcsin(\sqrt{p_2})$, compares differences across different base rates. Rough guides: 0.2 small, 0.5 medium, 0.8 large.
- For means, the difference with its interval, or **Cohen's d** (difference over pooled standard deviation).

Decide before the analysis what size of difference matters for the decision. A significant difference below that size is reported as "no meaningful change".

Why it matters

A retailer's panel of 90,000 customers showed satisfaction at 7.62 against 7.58 the previous quarter, $p = 0.01$. The headline read "Satisfaction significantly up". The store managers' bonus scheme was tied to "significant improvement", and bonuses were paid. The change was 0.04 points on a 10-point scale, smaller than the rounding in the published chart.

Worked example: weak code

```
1 se = (p * (1 - p) / n) ** 0.5
2 low, high = p - 1.96 * se, p + 1.96 * se
3 if p_value < 0.05:
4     headline = "Significant increase"
```

1. **Line 1.** Wald: zero width at $p = 0$ or 1, and too narrow for small `n`.

2. **Line 4.** Significance without size: true for trivial differences in large samples.

Worked example: strong code

```
1 import math
2
3
4 def wilson_interval(p: float, n: float, z: float = 1.959964) ->
    tuple[float, float] | None:
5     """Wilson interval for a share p from n observations (use n_eff when
        weighted)."""
6
7     if n <= 0 or not 0 <= p <= 1:
8         return None
9     denom = 1 + z**2 / n
10    centre = (p + z**2 / (2 * n)) / denom
11    half = z / denom * math.sqrt(p * (1 - p) / n + z**2 / (4 * n**2))
12    return (max(0.0, centre - half), min(1.0, centre + half))
13
14 def cohens_h(p1: float, p2: float) -> float:
15     """Effect size for two shares, comparable across base rates."""
16
17     return 2 * math.asin(math.sqrt(p1)) - 2 * math.asin(math.sqrt(p2))
```

1. **Line 4.** n is a float, so n_eff from weights fits.

2. **Line 6.** No interval for an empty group or an impossible share.

3. **Line 16.** An effect size to report beside the difference in points.

Wilson intervals with n_eff for weighted shares, and an effect size beside every p-value.

Recap card

- **Why not the Wald interval?** Zero width at $p = 0$ or 1 , and too narrow for small n. Use Wilson.
- **Interval for a weighted share** Wilson with the weighted p and n_eff.
- **Significant vs meaningful** Large samples make trivial differences significant. Report the size with its interval.
- **Cohen's h** $2\text{asin}(\sqrt{p_1}) - 2\text{asin}(\sqrt{p_2})$. About 0.2 small, 0.5 medium, 0.8 large.

Practice

Exercise 1. Multiple choice (level 1)

0 of 20 respondents chose a new brand. What does the Wald interval $p \pm 1.96 * \sqrt{p(1-p)/n}$ give?

- a) [0, 0]

- b) [0, 0.161]
- c) [-0.05, 0.05]
- d) It raises an error.

Why?

- r1) With $p = 0$ the standard error is 0, so the interval collapses.
- r2) Twenty respondents is too few for any interval.

Exercise 2. Predict the output (level 2)

What does this print? (statsmodels 0.14)

```
1 from statsmodels.stats.proportion import proportions_ztest
2
3 stat, p_value = proportions_ztest([50500, 50000], [100000, 100000])
4 print(p_value < 0.05)
```

Exercise 3. Fill the gap (level 2)

Get a 95% Wilson interval for 40 out of 100 with statsmodels.

```
1 low, high = proportion_confint(count=40, nobs=100, alpha=____(1),
    method="____(2)")
```

Exercise 4. Write the function (level 3)

Write `wilson_interval(p: float, n: float, z: float = 1.959964) -> tuple[float, float] | None` with the formula:

- $\text{denom} = 1 + z^2 / n$
- $\text{centre} = (p + z^2 / (2 * n)) / \text{denom}$
- $\text{half} = z / \text{denom} * \sqrt{p * (1 - p) / n + z^2 / (4 * n^2)}$

Return $(\text{centre} - \text{half}, \text{centre} + \text{half})$, clipped to $[0, 1]$. Return `None` when $n \leq 0$ or p is outside $[0, 1]$. n is a float so you pass `n_eff` for weighted shares.

Starter code:

```
1 import math
2
3
4 def wilson_interval(p, n, z=1.96):
5     se = math.sqrt(p * (1 - p) / n)
6     return (p - z * se, p + z * se)
```

Exercise 5. Architecture design (level 4)

A panel of 90,000 shows satisfaction of 7.62 against 7.58 last quarter, $p = 0.01$. The team agreed beforehand that 0.2 points matters for decisions. Choose what to report.

Options: No meaningful change in satisfaction, Satisfaction significantly up.

1. Headline
2. Supporting detail
3. Rule for future quarters

Sources

- `statsmodels.proportion_confint(40, 100, method="wilson")` returns (0.3094, 0.4980); `method="normal"` gives (0, 0) for 0 of 20. (verified: Run on statsmodels 0.14.6 and 0.15.0 in this repository.)
- `proportions_ztest([50500, 50000], [100000, 100000])` gives $p = 0.025$. (verified: Run on statsmodels 0.14.6 in this repository.)
- Cohen's h thresholds of 0.2, 0.5 and 0.8 for small, medium and large. (verified: Cohen (1988), Statistical Power Analysis for the Behavioral Sciences.)

B8-L3 Many tests, and the false positives they bring

Concept

At a 5% significance level, each test of a true null hypothesis has a 5% chance of a false positive. Run 20 such tests and the chance of **at least one** false positive is $1 - 0.95^{**20} = 0.64$. A simulation in this course's build agreed: 20 t-tests on pure noise, repeated 2,000 times, gave at least one "significant" result 66% of the time.

Survey reports run many tests by default: every brand, every region, every age band, every wave. Some cells will light up by chance every time.

Adjust the p-values for the number of tests in the family. In statsmodels 0.14:

```
1 from statsmodels.stats.multitest import multipletests
2 reject, p_adj, _, _ = multipletests(p_values, alpha=0.05, method="holm")
```

- **Bonferroni** multiplies each p-value by the number of tests m . Simple and strict.
- **Holm** is a step-down version of Bonferroni. It controls the same error (any false positive in the family) and always rejects at least as many hypotheses. Prefer it to Bonferroni.
- **Benjamini-Hochberg** (`method="fdr_bh"`) controls the **false discovery rate**: the expected share of false positives among the results you flag. It is less strict and suits screening many cells for follow-up.

For p-values 0.01, 0.02, 0.03, 0.04 and 0.20, Bonferroni and Holm reject only the first. Benjamini-Hochberg rejects the first four.

Define the **family** before you look: "the 12 regional comparisons in this chart". Adjust within it. Choosing the family after seeing the results defeats the point.

Two habits that reduce the problem at the source:

- Pre-register the few comparisons that answer the business question. Test those. Label the rest as exploratory.
- Report how many tests you ran. "3 of 40 regions changed significantly" reads very differently from "Region 7 changed significantly".

Why it matters

A monthly tracker tested 14 brands in 9 regions, 126 tests, each at 5%. Every month about six cells were flagged in red on the dashboard. Regional managers were asked to explain their red cells. The following month, different cells turned red. Nobody could explain the pattern, because there was none: 126 tests at 5% produce about six false positives by chance.

Holm adjustment within each brand's family of 9 regions removed nearly all the monthly red cells. The few that remained repeated in the next month.

Worked example: weak code

```
1 for region, p in p_values.items():
2     if p < 0.05:
3         flag(region)
```

1. **Line 2.** Each test at 5%, with no adjustment for how many run: false positives are expected.

Worked example: strong code

```
1 from statsmodels.stats.multitest import multipletests
2
3
4 def flag_regions(p_values: dict[str, float], alpha: float = 0.05) ->
    dict[str, float]:
5     """Holm-adjusted p-values for one family of regional tests, flagged
    below alpha."""
6     regions = list(p_values)
7     reject, p_adj, _, _ = multipletests([p_values[r] for r in regions],
8     alpha=alpha, method="holm")
9     return {r: float(p) for r, p, rej in zip(regions, p_adj, reject,
10     strict=True) if rej}
```

1. **Line 7.** One family, adjusted together with Holm.

2. **Line 8.** Report adjusted p-values, so readers see the correction.

Define the family first, adjust within it with Holm or Benjamini-Hochberg, and report how many tests ran.

Recap card

- **Chance of at least one false positive in 20 tests at 5%** $1 - 0.95^{20} = 0.64$.
- **Holm vs Bonferroni** Same family-wise control. Holm is step-down and rejects at least as many. Prefer Holm.
- **Benjamini-Hochberg** Controls the false discovery rate. Suits screening many cells.
- **Define the family** Before looking at results. Report how many tests ran.

Practice

Exercise 1. Multiple choice (level 1)

You run 20 independent tests at the 5% level. All null hypotheses are true. What is the chance of at least one significant result?

- a) 5%
- b) About 64%
- c) 100%
- d) 0%, because the nulls are true

Why?

- r1) $1 - 0.95^{20} = 0.64$.
- r2) $20 \times 5\% = 100\%$, minus overlap.

Exercise 2. Predict the output (level 2)

What does this print? (statsmodels 0.14)

```
1 from statsmodels.stats.multitest import multipletests
2
3 p = [0.01, 0.02, 0.03, 0.04, 0.20]
4 print(multipletests(p, alpha=0.05, method="holm")[0].sum(),
       multipletests(p, alpha=0.05, method="fdr_bh")[0].sum())
```

Exercise 3. Fill the gap (level 2)

Adjust a family of p-values with Holm's method.

```
1 reject, p_adj, _, _ = ____ (1)(p_values, alpha=0.05, method="____ (2)")
```

Exercise 4. Write the function (level 3)

Write `holm_adjust(p_values: list[float]) -> list[float]` without `statsmodels`. Return adjusted p-values **in the input order**.

1. Sort the p-values ascending. For the k-th smallest (k from 0), compute $\min(1, (m - k) * p)$, where `m` is the number of p-values.
2. Walk the sorted list and keep a running maximum, so adjusted values never decrease.
3. Put each adjusted value back at its original position.

Starter code:

```
1 def holm_adjust(p_values):
2     m = len(p_values)
3     return [min(1.0, p * m) for p in p_values]
```

Exercise 5. Architecture design (level 4)

Choose the approach for each analysis.

Options: Holm, family of the three, Benjamini-Hochberg, No adjustment.

1. Three pre-registered comparisons that decide a campaign's budget
2. Screen 400 item-by-segment cells to pick a few for follow-up research
3. One primary hypothesis, stated before fieldwork

Sources

- multipletests on [0.01, 0.02, 0.03, 0.04, 0.20]: Holm and Bonferroni reject one, `fdr_bh` rejects four. (verified: Run on `statsmodels` 0.14.6 and 0.15.0 in this repository.)
- 20 t-tests on noise, 2,000 repetitions: at least one $p < 0.05$ in 66% of repetitions. (verified: Simulated with NumPy 2.4.6 and SciPy 1.17.1 in this repository.)

B9 Experimental design

Roadmap: Prerequisite for all months. Power, sample size and A/B test analysis.

B9-L1 Power and sample size, before fieldwork

Concept

A test with too few people cannot see the effect you care about. It returns "no significant difference", and the team concludes the change did nothing. Decide the sample size **before** fieldwork, from four inputs:

- **Baseline**: the current rate, for example 40% consideration.
- **Minimum detectable effect (MDE)**: the smallest change worth acting on, for example 5 points, to 45%.
- **Significance level** (alpha): usually 0.05, two-sided.
- **Power**: the chance of detecting the MDE if it is real. Usually 0.8.

For two proportions, the sample per group is:

```
1 z_a = NormalDist().inv_cdf(1 - alpha / 2)
2 z_b = NormalDist().inv_cdf(power)
3 p_bar = (p1 + p2) / 2
4 n = ((z_a * sqrt(2 * p_bar * (1 - p_bar)) + z_b * sqrt(p1 * (1 - p1) + p2
    * (1 - p2))) / (p2 - p1)) ** 2
```

Round up. For 40% to 45% at alpha 0.05 and power 0.8: **1,534 per group**. statsmodels 0.14 agrees:

`NormalIndPower().solve_power(effect_size=proportion_effectsize(0.45, 0.40), alpha=0.05, power=0.8)` returns 1,533.05.

The MDE drives everything. Required **n** grows with the inverse square of the effect:

Change	n per group
40% to 45%	1,534
40% to 42%	9,493

Halving the effect you want to detect needs roughly four times the sample.

Two adjustments for survey work:

- **Design effect.** Weighted or clustered samples carry less information (lesson B8-L1). Multiply n by the expected design effect: with a design effect of 1.5, $1,533.05 \times 1.5$ rounds up to 2,300.
- **Subgroups.** If the decision needs the effect within each of four regions, each region needs the full n , not a quarter of it.

Write the power calculation into the research plan, with its inputs. If the budget cannot buy the n , say so before fieldwork, and agree a larger MDE or a different design.

Why it matters

A packaging test used 300 respondents per cell to compare purchase intent. The new pack scored 43% against 40%. The report said "no significant difference, keep the current pack". With 300 per cell, the test had about a 1 in 4 chance of detecting a real 5-point lift. The result said nothing about the pack. It said the test was too small.

Worked example: weak code

```
❶ N_PER_CELL = 300 # what the budget allows
2 result = run_test(N_PER_CELL)
3 if result.p_value > 0.05:
❷     conclusion = "No effect"
```

1. **Line 1.** Sample size from budget, with no check of what it can detect.
2. **Line 4.** Absence of evidence from an underpowered test read as evidence of absence.

Worked example: strong code

```
1 import math
2 from statistics import NormalDist
3
4
❶ def n_per_group(p1: float, p2: float, alpha: float = 0.05, power: float =
    0.8) -> int:
5     """Sample per group to detect p1 vs p2 with a two-sided test."""
6     if p1 == p2:
7
❷         raise ValueError("p1 and p2 must differ")
8
9     z_a = NormalDist().inv_cdf(1 - alpha / 2)
10    z_b = NormalDist().inv_cdf(power)
11    p_bar = (p1 + p2) / 2
12    top = z_a * math.sqrt(2 * p_bar * (1 - p_bar)) + z_b * math.sqrt(p1 *
    (1 - p1) + p2 * (1 - p2))
```

```
3 return math.ceil((top / (p2 - p1)) ** 2)
```

1. **Line 5.** Every input to the decision is a named parameter.
2. **Line 8.** A zero effect needs infinite n. Fail clearly instead of dividing by zero.
3. **Line 13.** Round up: 1,533.05 means 1,534 people.

Fix baseline, MDE, alpha and power first, compute n, adjust for design effect and subgroups, then field.

Recap card

- **Four inputs to a sample size** Baseline, minimum detectable effect, alpha, power.
- **n per group, 40% vs 45%, alpha 0.05, power 0.8** 1,534.
- **Halve the effect you want to detect** About four times the sample.
- **Weighted or subgroup reads** Multiply n by the design effect. Each subgroup needs the full n.

Practice

Exercise 1. Multiple choice (level 1)

A test with 300 per cell finds no significant difference. Power for the effect you care about was 0.25. What can you conclude?

- a) There is no effect.
- b) Very little. The test was unlikely to detect the effect even if it were real.
- c) The effect is negative.
- d) The effect is exactly zero with 75% confidence.

Why?

- r1) Power 0.25 means a real effect of that size is missed 75% of the time.
- r2) Non-significant results are always meaningless.

Exercise 2. Predict the output (level 2)

Required n per group is 1,534 to detect 40% vs 45%. Roughly how many per group to detect 40% vs 42.5%, with the same alpha and power? The code prints the ratio of the two sample sizes, rounded.

```
1 import math
2 from statistics import NormalDist
3
4
5 def n_per_group(p1, p2, alpha=0.05, power=0.8):
6     z_a = NormalDist().inv_cdf(1 - alpha / 2)
7     z_b = NormalDist().inv_cdf(power)
```

```

8     p_bar = (p1 + p2) / 2
9     top = z_a * math.sqrt(2 * p_bar * (1 - p_bar)) + z_b * math.sqrt(p1 *
    (1 - p1) + p2 * (1 - p2))
10    return math.ceil((top / (p2 - p1)) ** 2)
11
12
13 print(round(n_per_group(0.40, 0.425) / n_per_group(0.40, 0.45)))

```

Exercise 3. Fill the gap (level 2)

Compute the sample size with statsmodels for 40% vs 45%.

```

1 h = proportion_effectsize(0.45, 0.40)
2 n = NormalIndPower().____(1)(effect_size=h, alpha=0.05, ____ (2)=0.8)

```

Exercise 4. Write the function (level 3)

Write `n_per_group(p1, p2, alpha=0.05, power=0.8, design_effect=1.0) -> int`, using the formula in the concept with `statistics.NormalDist`. Multiply by `design_effect` before rounding up. Raise `ValueError` if `p1 == p2`, or if either share is outside (0, 1).

Starter code:

```

1 import math
2
3
4 def n_per_group(p1, p2, alpha=0.05, power=0.8, design_effect=1.0):
5     return 300

```

Exercise 5. Architecture design (level 4)

A client wants to know whether a new ad lifts consideration from 40% by at least 5 points, overall and within each of 4 regions. Weighting will give a design effect of about 1.5. The budget buys 4,000 completes. Choose the plan.

Options: $2 \times 1,534 \times 1.5 = 4,602$: slightly above budget, 4,000 is plenty for any read.

1. Sample for the overall read
2. Regional reads
3. Recommendation

Sources

- `NormalIndPower().solve_power(effect_size=proportion_effectsize(0.45, 0.40), alpha=0.05, power=0.8)` returns 1533.05. (verified: Run on statsmodels 0.15.0 in this repository; matches the closed-form result of 1,534.)

- With 300 per cell, power to detect 40% vs 45% is 0.24. (verified: Computed with statsmodels NormalIndPower in this repository.)

B9-L2 Analyse an A/B test the way it was randomised

Concept

An A/B test compares groups that differ only by chance and by the change you made. Keep that property in the analysis.

1. Check the split first. If you assigned 50/50 and got 50,600 against 49,400, something went wrong: a redirect lost users, a bot filter hit one arm, a crash hit one version. A chi-square test on the counts detects this **sample ratio mismatch** (SRM). For those counts, $p = 0.00015$. A very small p-value here, below 0.001 by common practice, means stop and find the cause. The treatment effect is not trustworthy until you do.

```
1 from scipy.stats import chisquare
2 chisquare([50_600, 49_400]).pvalue      # 0.000148, SciPy 1.17
```

2. Analyse as randomised. Everyone assigned to B stays in B, even if they never saw the new page (intention to treat). Dropping "people who did not see it" from B, but not the matching people from A, breaks the randomisation.

3. Analyse at the unit of randomisation. If you randomised users, compute one outcome per user. Counting page views as independent inflates n and shrinks intervals.

4. Report the difference with its interval. For 1,150 of 10,000 in B against 1,000 of 10,000 in A, the lift is 1.5 points, with a 95% interval of 0.6 to 2.4 points. statsmodels 0.14:

```
1 from statsmodels.stats.proportion import confint_proportions_2indep
2 confint_proportions_2indep(1150, 10_000, 1000, 10_000)  # (0.0064,
0.0236)
```

5. Compare with the minimum effect you planned for (lesson B9-L1). An interval that lies wholly above the MDE is a clear win. One that straddles it needs more data or a judgement call, stated as such.

Fix the primary metric, the analysis and the stopping rule in writing before launch. Secondary metrics are exploratory and are reported as such (lesson B8-L3).

Why it matters

A checkout test showed a 4% lift in conversion. Nobody checked the split. Arm B had 3% fewer users than arm A, because a slow script in B timed out on older phones, and those users were logged as never arriving. They were also the users least likely to buy. The "lift" came from losing low-converting users from B. When the script was fixed, the lift disappeared.

An SRM check on day one would have flagged the missing 3%.

Worked example: weak code

```
1 b_seen = b[b["saw_new_page"]]
2 rate_a = a["converted"].mean()
3 rate_b = b_seen["converted"].mean()
4 print("lift", rate_b - rate_a)
```

1. **Line 1.** Filters one arm on a post-assignment event: no longer a fair comparison.
2. **Line 4.** No split check, no interval, no comparison with the planned effect.

Worked example: strong code

```
1 import math
2
3
4 def srm_p_value(n_a: int, n_b: int, expected_share_a: float = 0.5) ->
    float:
5     """Chi-square p-value (1 df) for the observed split against the
        planned one."""
6     n = n_a + n_b
7     e_a, e_b = n * expected_share_a, n * (1 - expected_share_a)
8     stat = (n_a - e_a) ** 2 / e_a + (n_b - e_b) ** 2 / e_b
9     return math.erfc(math.sqrt(stat / 2))
10
11
12 def lift_interval(x_a: int, n_a: int, x_b: int, n_b: int, z: float =
    1.959964) -> tuple[float, float, float]:
13     """Difference in conversion (B minus A) with a Wald interval, for
        large samples."""
14     p_a, p_b = x_a / n_a, x_b / n_b
15     se = math.sqrt(p_a * (1 - p_a) / n_a + p_b * (1 - p_b) / n_b)
16     diff = p_b - p_a
17     return diff, diff - z * se, diff + z * se
```

1. **Line 9.** Chi-square with 1 degree of freedom, in the standard library.
2. **Line 12.** Uses every assigned user, as randomised.

Check the split, keep everyone in their assigned arm, analyse per randomised unit, and report the lift with its interval.

Recap card

- **First step in any A/B analysis** Sample ratio mismatch check: chi-square on arm counts against the plan.
- **Intention to treat** Analyse everyone in the arm they were assigned to, whatever happened next.
- **Unit of analysis** The unit of randomisation, usually the user, not the page view.
- **Report a lift as** The difference with its interval, compared with the planned minimum effect.

Practice

Exercise 1. Multiple choice (level 1)

A 50/50 test ends with 50,600 users in A and 49,400 in B. The SRM check gives $p = 0.00015$. What next?

- a) Report the treatment effect. A 1% imbalance is small.
- b) Stop and find why B lost users before trusting any result.
- c) Reweight B up to 50,600 and continue.
- d) Drop 1,200 random users from A to balance.

Why?

- r1) A split that differs from the design means something removed users unevenly, and they may differ in outcome.
- r2) Any p-value below 0.05 invalidates a test.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 import math
2
3
4 def srm_p_value(n_a, n_b, expected_share_a=0.5):
5     n = n_a + n_b
6     e_a, e_b = n * expected_share_a, n * (1 - expected_share_a)
7     stat = (n_a - e_a) ** 2 / e_a + (n_b - e_b) ** 2 / e_b
8     return math.erfc(math.sqrt(stat / 2))
9
10
11 print(srm_p_value(5060, 4940) > 0.05, srm_p_value(50600, 49400) > 0.05)
```

Exercise 3. Put in order (level 2)

Order the analysis of a finished A/B test.

- Compare the interval with the planned minimum effect
- Check the split against the design
- Estimate the primary metric's lift with its interval
- Report secondary metrics as exploratory
- Aggregate to one outcome per randomised user

Exercise 4. Write the function (level 3)

Write `srn_p_value(n_a: int, n_b: int, expected_share_a: float = 0.5) -> float`. Compute the chi-square statistic with expected counts `n * expected_share_a` and `n * (1 - expected_share_a)`, and return the 1-degree-of-freedom p-value `math.erfc(math.sqrt(stat / 2))`. Raise `ValueError` if either count is negative, the total is 0, or `expected_share_a` is not strictly between 0 and 1.

Starter code:

```
1 def srm_p_value(n_a, n_b, expected_share_a=0.5):
2     return abs(n_a - n_b) / (n_a + n_b)
```

Exercise 5. Spot the bug (level 4)

The analysis reports a larger lift than the product team expected. Tap the line that breaks the randomisation, and fix the function.

```
1 import pandas as pd
2
3
4 def conversion_by_arm(users: pd.DataFrame) -> dict[str, float]:
5     """Conversion rate per assigned arm. One row per randomised user."""
6     analysed = users[(users["arm"] == "A") | users["saw_new_page"]]
7     return analysed.groupby("arm")["converted"].mean().to_dict()
```

Write down the line number of the bug. Then write the fix.

Sources

- `chisquare([50600, 49400]).pvalue` is 0.000148; for 5060 vs 4940 it is 0.230. (verified: Run on SciPy 1.17.1 in this repository; matches `math.erfc(sqrt(stat / 2))`.)
- `confint_proportions_2indep(1150, 10000, 1000, 10000)` gives (0.0064, 0.0236). (verified: Run on statsmodels 0.15.0 in this repository.)

- An SRM p-value below 0.001 is a common threshold for investigation. (search-only: Fabijan et al. (2019), Diagnosing Sample Ratio Mismatch in Online Controlled Experiments, KDD.)

B9-L3 Peeking, novelty and other ways tests mislead

Concept

A test's 5% false positive rate holds only if you analyse once, at the planned sample size. Common habits break it.

Peeking. Checking results every day and stopping the first time $p < 0.05$ gives the noise ten chances to cross the line. A simulation in this course's build: two identical arms with a 10% conversion rate, 1,000 runs, 200 users per arm per look. One look: 5.2% false positives. Five looks: 12.9%. Ten looks: 18.4%.

Fixes, in order of simplicity:

- **Fixed horizon.** Compute the sample size (lesson B9-L1), run to it, analyse once. Dashboards during the test show health metrics and the split check, not the lift.
- **Planned interim looks** with a group-sequential design, where each look uses a stricter threshold so the overall rate stays at 5%. Plan the number of looks and the boundaries before launch.
- **Always-valid methods** designed for continuous monitoring,, offered by some experimentation platforms. Use them as the platform documents them.

Novelty and learning effects. Users click a new button because it is new. The lift fades after a week. Run for at least one full weekly cycle, and compare the effect in week 1 and week 2.

Too many metrics. Twenty secondary metrics produce about one false positive per test (lesson B8-L3). Name one primary metric and a few guardrails before launch.

Stopping for a good result only. A team that stops early when B looks good but runs longer when B looks bad biases every conclusion towards B.

Interference. If treated users affect control users, through shared inventory, pricing or social features, the arms are not independent. Randomise at a level that contains the interaction, such as city or store.

Write the plan down: primary metric, MDE, sample size, duration, looks, and the decision rule. Then follow it.

Why it matters

A growth team ran 30 tests a quarter, checked each daily, and shipped any variant that reached $p < 0.05$. Their dashboard showed 11 winners, with a combined forecast lift of 9% in sign-ups. Sign-ups over the quarter were flat. With daily peeking over three-week tests, the false positive rate per test was far above 5%, and most of the winners were noise.

Worked example: weak code

```
❶ for day in range(1, 29):
    2     p = daily_p_value(day)
    ❷ if p < 0.05:
        4         ship("B")
        5         break
```

1. **Line 1.** Up to 28 looks at the same growing data.
2. **Line 3.** The same 5% threshold at every look: the overall false positive rate climbs well above 5%.

Worked example: strong code

```
1 from dataclasses import dataclass
2
3
4 @dataclass(frozen=True)
    ❶ class TestPlan:
        6     """Written before launch and stored with the results."""
        7
        8     primary_metric: str
        9     mde_points: float
       10     n_per_arm: int
       11     min_days: int
    ❷     looks: int = 1
       13
       14
       15 def ready_to_analyse(plan: TestPlan, n_per_arm_now: int, days_run: int) -
       16     > bool:
           16     """True only once the planned sample and duration are both reached."""
    ❸     return n_per_arm_now >= plan.n_per_arm and days_run >= plan.min_days
```

1. **Line 5.** The plan is data, fixed before launch.
2. **Line 12.** One analysis unless a sequential design was planned.
3. **Line 17.** The lift is looked at once, when the plan says so.

Plan the sample, duration and looks, analyse once at the horizon, and treat anything else as exploration.

Recap card

- **Why peeking breaks a test** Each look is another chance for noise to cross $p < 0.05$. Ten looks gave 18.4% false positives in simulation.
- **Safe alternatives to daily peeking** Fixed horizon, or planned group-sequential looks with stricter thresholds.
- **Novelty effect** Lift that fades after launch. Run full weeks and compare week 1 with later weeks.
- **Interference** Treated units affect controls. Randomise at a level that contains the interaction.

Practice

Exercise 1. Multiple choice (level 1)

You check an A/A test (no real difference) after each of 10 equal batches and stop at the first $p < 0.05$. In this lesson's simulation, how often did that declare a winner?

- a) About 5%
- b) About 18%
- c) About 50%
- d) Never, because the arms are identical

Why?

- r1) Each look is another chance for noise to cross the threshold.
- r2) Later batches have more users, so they are always significant.

Exercise 2. Predict the output (level 2)

A test ran 14 days. The lift in conversion was 3.1 points in week 1 and 0.4 points in week 2. What does this print?

```
1 week_lifts = {"week_1": 3.1, "week_2": 0.4}
2 fading = week_lifts["week_2"] < week_lifts["week_1"] / 2
3 print("novelty suspected" if fading else "stable")
```

Exercise 3. Fill the gap (level 2)

Allow analysis only when both the planned sample and the planned duration are reached.

```
1 return n_per_arm_now >= plan.n_per_arm ____ (1) days_run >= plan. ____ (2)
```

Exercise 4. Write the function (level 3)

Write `peeking_false_positive_rate(rng, n_sims, n_per_look, looks, rate, z_crit=1.959964)` -> float. It simulates `n_sims` A/A tests where both arms convert at `rate`. Each arm gets `n_per_look` new users per look. At each look, run a pooled two-proportion z-test on the cumulative data. Return the share of simulations where $|z| > z_crit$ at **any** look. Treat a zero standard error as $z = 0$. Use only the `rng` passed in.

Starter code:

```
1 import numpy as np
2
3
4 def peeking_false_positive_rate(rng, n_sims, n_per_look, looks, rate,
5     z_crit=1.959964):
6     return 0.05
```

Exercise 5. Architecture design (level 4)

Stakeholders want to "see results daily" on a four-week pricing test. Choose the plan.

Options: Split check, traffic and error rates. No lift until the horizon., Daily lift with p-values.

1. What the daily dashboard shows
2. Analyses of the primary metric
3. Randomisation unit, since prices are shared in store

Sources

- A/A simulation, 1,000 runs, 200 users per arm per look, rate 10%: 5.2% false positives with 1 look, 12.9% with 5, 18.4% with 10. (verified: Simulated with NumPy 2.4.6 in this repository (seed 1).)

B10 Feature engineering without leakage

Roadmap: Prerequisite for all months. Target leakage, contamination and time-safe features.

B10-L1 Target leakage, features that know the answer

Concept

Target leakage happens when a feature carries information that would not exist at the moment you make the prediction. The model learns the leak, scores brilliantly in validation, and fails in production, where the leak is absent.

Ask one question of every feature: **would I know this value at prediction time?**

Typical leaks in customer and survey data:

- **Consequences of the outcome.** `cancellation_reason` predicts churn perfectly, because only churners have one. `refund_amount` predicts complaints.
- **Values measured after the prediction date.** `calls_to_support_last_30d` computed on the extraction date, not the prediction date, includes calls made while the customer was leaving.
- **Proxies from the labelling process.** A `reviewed_by` column that is only filled in for fraud cases.
- **IDs and timestamps that encode order.** If churners were loaded in a later batch, a row ID predicts churn.

A quick screen catches the worst cases. Compute each feature's AUC on its own against the target. A single feature with an AUC near 1.0, or near 0.0 for an inverse relationship, deserves suspicion before celebration. With scikit-learn 1.8:

```
1 from sklearn.metrics import roc_auc_score
2 auc = roc_auc_score(y, X[col])
3 strength = max(auc, 1 - auc)
```

A screen does not prove leakage, and a leak spread across several features will pass it. The real defence is a **feature timeline**: for every feature, the date it becomes known, compared with the prediction date. Keep that table

in the project and review it with someone who knows how the data is produced.

When in doubt, leave the feature out and measure how much the model loses. A feature worth 0.15 AUC on its own that nobody can explain is more likely a leak than a discovery.

Why it matters

A churn model scored AUC 0.97 in validation. The top feature was `days_since_last_login`, computed on the day the training data was extracted. For churned customers, that included the weeks after they had already left. In production, scored on active customers, the model ranked customers little better than chance. The retention campaign built on it spent its budget on the wrong people for two months.

Worked example: weak code

```
❶ features = ["tenure_months", "plan", "days_since_last_login",
             "cancellation_reason"]
❷ model.fit(train[features], train["churned"])
❸ print(roc_auc_score(test["churned"], model.predict_proba(test[features])
                     [:, 1]))
```

1. **Line 1.** `cancellation_reason` exists only for churners. `days_since_last_login` was measured after the prediction date.
2. **Line 3.** The test set shares the leak, so the score looks superb.

Worked example: strong code

```
❶ import pandas as pd
❷ from sklearn.metrics import roc_auc_score
❸
❹ LEAK_THRESHOLD = 0.95
❺
❻
❼ def leakage_screen(X: pd.DataFrame, y: pd.Series, threshold: float =
    LEAK_THRESHOLD) -> list[str]:
    """Numeric features that alone separate the classes suspiciously
    well."""
    suspects = []
    ❶ for col in X.select_dtypes("number").columns:
    ❷     known = X[col].notna()
    ❸     if known.sum() == 0 or y[known].nunique() < 2:
    ❹         continue
    ❺     auc = roc_auc_score(y[known], X.loc[known, col])
    ❻     if max(auc, 1 - auc) >= threshold:
    ❼         suspects.append(col)
```

1. **Line 10.** One feature at a time: a single near-perfect feature is the classic leak.
2. **Line 12.** No AUC without both classes. Skip rather than crash.
3. **Line 15.** Catches inverse relationships too.

For every feature, know when it becomes available. Screen single-feature AUCs, and treat a near-perfect one as a suspect.

Recap card

- **The one question for every feature** Would I know this value at prediction time?
- **Single-feature leakage screen** AUC of each feature alone. Flag $\max(\text{auc}, 1 - \text{auc})$ near 1.
- **Common leaks** Consequences of the outcome, values measured after the prediction date, labelling proxies, order-encoding IDs.
- **Feature timeline** A table of when each feature becomes known, reviewed with the data owner.

Practice

Exercise 1. Multiple choice (level 1)

You predict on 1 March whether a customer will churn in March. Which feature leaks?

- a) Tenure in months on 1 March
- b) Number of support calls in February
- c) Number of support calls in March
- d) Plan type on 1 March

Why?

- r1) March calls happen after the prediction is made, and may be caused by leaving.
- r2) Support calls always leak.

Exercise 2. Predict the output (level 2)

What does this print? (scikit-learn 1.8)

```
1 from sklearn.metrics import roc_auc_score
2
3 churned = [0, 1, 0, 1]
4 days_inactive = [0.1, 0.9, 0.2, 0.8]
5 contract_months_left = [0.9, 0.1, 0.8, 0.2]
6 a = roc_auc_score(churned, days_inactive)
```

```
7 b = roc_auc_score(churned, contract_months_left)
8 print(a, max(b, 1 - b))
```

Exercise 3. Put in order (level 2)

Order the review of a proposed feature.

- Explain any high-AUC feature with the data owner, or drop it
- Drop features known only after the prediction date
- Find the date each feature value becomes known
- Write down the prediction date and what the model will see then
- Run the single-feature AUC screen on what remains

Exercise 4. Write the function (level 3)

Write `leakage_screen(X, y, threshold=0.95) -> list[str]`. For each numeric column of `X`, in column order, compute `roc_auc_score` on rows where the column is not missing. Return the names where `max(auc, 1 - auc) >= threshold`. Skip columns where the non-missing rows do not contain both classes.

Starter code:

```
1 import pandas as pd
2 from sklearn.metrics import roc_auc_score
3
4
5 def leakage_screen(X, y, threshold=0.95):
6     return [c for c in X.columns if roc_auc_score(y, X[c]) > threshold]
```

Exercise 5. Architecture design (level 4)

A fraud model predicts at transaction time. Decide what to do with each candidate feature.

Options: Keep, Drop: known only after the prediction, Rebuild as of transaction time.

1. Transaction amount
2. Chargeback filed (yes/no)
3. Customer's transaction count, computed on the extract date
4. Analyst who reviewed the case

Sources

- `roc_auc_score` returns 0.0 for a perfect inverse ranking. (verified: Run on scikit-learn 1.8.0 and 1.9.1 in this repository.)

B10-L2 Fit on training data only

Concept

Every step that **learns from data** must learn from the training rows only: imputation means, scaling parameters, category lists, target encodings, feature selection, even the choice of outliers to drop. If a step sees the test rows, information about them flows into the model, and the test score becomes optimistic. This is **train/test contamination**.

A demonstration from this course's build, with scikit-learn 1.8. Take 100 rows of pure noise: 5,000 random features and a random 0/1 target. Nothing can be predicted.

- Select the 20 features most correlated with the target **on all 100 rows**, then cross-validate a logistic regression on them: accuracy **0.85**.
- Put the same selection **inside** each cross-validation fold: accuracy **0.48**, the chance level it should be.

The first model did not learn anything real. The selection step had already looked at every test fold's labels and picked the features that happened to fit them.

The rule for manual preprocessing:

```
1 params = fit_preprocessing(train)           # learn from train only
2 train_ready = apply_preprocessing(train, params)
3 test_ready = apply_preprocessing(test, params) # reuse, never refit
```

Split **first**, before any step that learns. In cross-validation, that means refitting the preprocessing inside every fold. Doing that by hand is error-prone. scikit-learn's `Pipeline` does it for you (lesson B11-L2).

Steps that do **not** learn are safe on the full data: fixed recodes from a codebook, unit conversions, parsing dates. Steps that learn include anything with `fit`, anything that computes a mean, a quantile, a list of categories or a correlation with the target.

Why it matters

A propensity model imputed missing income with the median of the full dataset and scaled every feature on the full dataset before the train/test split. It then selected features by their correlation with the target, also on the full data. The validation AUC was 0.79. On the next month's customers it was 0.66. Most of the gap came from the feature selection: it had fitted the noise in the test rows.

Worked example: weak code

```
❶ df["income"] = df["income"].fillna(df["income"].median())
❷ df["income"] = (df["income"] - df["income"].mean()) / df["income"].std()
❸ train, test = train_test_split(df, test_size=0.2, random_state=1)
```

1. **Line 1.** The median includes test rows.
2. **Line 2.** Scaling parameters learned from test rows too.
3. **Line 3.** The split comes after the learning, too late.

Worked example: strong code

```
1 import pandas as pd
2
3
❶ def fit_income_prep(train: pd.DataFrame) -> dict[str, float]:
4     """Learn median, mean and sd of income from training rows only."""
5     filled = train["income"].fillna(train["income"].median())
6     return {"median": float(train["income"].median()), "mean":
7             float(filled.mean()), "sd": float(filled.std())}
8
9
10 def apply_income_prep(df: pd.DataFrame, params: dict[str, float]) ->
    pd.DataFrame:
11     """Impute and scale income with parameters learned elsewhere."""
❷     filled = df["income"].fillna(params["median"])
13     return df.assign(income=(filled - params["mean"]) / params["sd"])
```

1. **Line 4.** Learning happens in one place, on training rows.
2. **Line 12.** Applying reuses the learned values, for train and test alike.

Split first. Fit every learning step on training rows. Apply the fitted parameters to everything else.

Recap card

- **Train/test contamination** A learning step saw test rows: imputation, scaling, encoding, selection.
- **Selection on all rows, noise data** CV accuracy 0.85. Inside the folds: 0.48. The first number is a leak.
- **The manual rule** Split first, fit on train, apply the fitted parameters to train and test.
- **Safe before the split** Steps that learn nothing from the rows: fixed recodes, unit conversions, parsing.

Practice

Exercise 1. Multiple choice (level 1)

Which step is safe to run on the full dataset before the split?

- a) Filling missing income with the median
- b) Converting height from inches to centimetres
- c) Selecting the 20 features most correlated with the target
- d) Standardising features to mean 0 and sd 1

Why?

- r1) A unit conversion uses a fixed constant. It learns nothing from the rows.
- r2) Conversions happen before any modelling.

Exercise 2. Predict the output (level 2)

Pure noise: 100 rows, 5,000 random features, random labels. The 20 best features are selected on all rows, then 5-fold CV is run. Roughly what accuracy does it print? (scikit-learn 1.8, seed 0)

```
1 import numpy as np
2 from sklearn.feature_selection import SelectKBest, f_classif
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.model_selection import cross_val_score
5
6 rng = np.random.default_rng(0)
7 X = rng.normal(size=(100, 5000))
8 y = rng.integers(0, 2, 100)
9 X_sel = SelectKBest(f_classif, k=20).fit_transform(X, y)
10 print(round(cross_val_score(LogisticRegression(), X_sel, y, cv=5).mean(),
11                        2))
```

Exercise 3. Fill the gap (level 2)

Fit the scaler on training rows, then apply it to both sets.

```
1 scaler = StandardScaler().fit(____(1))
2 X_train_s = scaler.transform(X_train)
3 X_test_s = scaler.____(2)(X_test)
```

Exercise 4. Write the function (level 3)

Write `fit_prep(train) -> dict[str, float]` and `apply_prep(df, params) -> pd.DataFrame` for the column `income`.

- `fit_prep` learns `median` of `income`, then fills missing values with it and learns `mean` and `sd` of the filled column. Training rows only.

- `apply_prep` fills missing income with `params["median"]` and returns a new frame with `income` scaled as `(filled - mean) / sd`. If `sd` is 0, set scaled income to 0.0.

Starter code:

```
1 import pandas as pd
2
3
4 def fit_prep(train):
5     return {}
6
7
8 def apply_prep(df, params):
9     filled = df["income"].fillna(df["income"].median())
10    return df.assign(income=(filled - filled.mean()) / filled.std())
```

Exercise 5. Spot the bug (level 4)

Target encoding replaces each region with its churn rate. The model's test AUC is far above its production AUC. Tap the line that leaks, and fix the function so the rates come from training rows only.

```
1 import pandas as pd
2
3
4 def encode_region(train: pd.DataFrame, test: pd.DataFrame) ->
5     tuple[pd.DataFrame, pd.DataFrame]:
6     """Replace region with its churn rate."""
7     both = pd.concat([train, test])
8     rates = both.groupby("region")["churned"].mean()
9     return train.assign(region=train["region"].map(rates)),
10    test.assign(region=test["region"].map(rates))
```

Write down the line number of the bug. Then write the fix.

Sources

- SelectKBest(k=20) on all rows of 100 x 5,000 noise gives 0.85 CV accuracy; inside a Pipeline, 0.48. (verified: Run on scikit-learn 1.8.0 and 1.9.1 with NumPy 2.4.6 in this repository (seed 0).)
- scikit-learn provides TargetEncoder with internal cross fitting. (verified: scikit-learn 1.8 API reference; sklearn.preprocessing.TargetEncoder imported in this repository.)

B10-L3 Time-safe features and point-in-time joins

Concept

When rows have dates, every feature must be computed **as of** the row's own date, using only earlier data. Three tools in pandas 3.0 cover most cases.

Prior-only aggregates. "Total spend before this order" must exclude the current order. Take the running total and subtract the current value, or shift first:

```
1 orders = orders.sort_values(["customer", "ordered_at"])
2 prior = orders.groupby("customer")["spend"].cumsum() - orders["spend"]
```

Point-in-time joins. To attach the latest satisfaction score **before** each prediction date, use `pd.merge_asof`. It matches each left row with the last right row whose key is at or before it, within the same `by` group. Both frames must be sorted on the key.

```
1 pd.merge_asof(preds.sort_values("as_of"), scores.sort_values("scored_at"),
2               left_on="as_of", right_on="scored_at", by="customer",
3               allow_exact_matches=False)
```

The default `allow_exact_matches=True` also matches a score recorded **at the same timestamp** as the prediction. If that score could only be known after the prediction was made, set it to `False`. A plain `merge` on customer followed by "take the latest" silently uses scores from the future.

Time-based splits. Random splits put next month's rows in training and last month's in test. For anything you will use forward in time, split by date: train on data up to a cut-off, test on data after it. scikit-learn's `TimeSeriesSplit` gives expanding windows for cross-validation.

Build a small **as-of test** for every time feature: a customer with events before and after the prediction date, and an assertion that only the earlier events count. It catches most mistakes in a few lines.

Why it matters

A next-purchase model used "customer's total orders" as a feature, computed once over the full order table. For each training row, the count included orders placed after the row's date, including the very purchase the model was trying to predict. Validation AUC was 0.91. In production, where the future does not exist yet, it was 0.68.

Recomputing the count as prior orders only brought validation down to 0.70. That was the model's true performance, and the team finally had an honest number to improve.

Worked example: weak code

```
1 totals = orders.groupby("customer")["spend"].sum().rename("total_spend")
2 rows = orders.merge(totals, on="customer")
3 train, test = train_test_split(rows, test_size=0.2)
```

1. **Line 1.** One total per customer across all time: every row sees its own future.
2. **Line 3.** Random split: training rows come from after test rows.

Worked example: strong code

```
1 import pandas as pd
2
3
4 def add_prior_spend(orders: pd.DataFrame) -> pd.DataFrame:
5     """Spend and order count before each order, per customer. Grain: one
6     row per order."""
7     o = orders.sort_values(["customer", "ordered_at"])
8     grp = o.groupby("customer")["spend"]
9     return o.assign(
10         prior_spend=grp.cumsum() - o["spend"],
11         prior_orders=o.groupby("customer").cumcount(),
12     )
13
14 def time_split(df: pd.DataFrame, date_col: str, cutoff: pd.Timestamp) ->
15     tuple[pd.DataFrame, pd.DataFrame]:
16     """Train on rows before cutoff, test on rows at or after it."""
17     before = df[date_col] < cutoff
18     return df[before], df[~before]
```

1. **Line 9.** Running total minus the current order: only earlier orders count.
2. **Line 10.** cumcount starts at 0, so it counts earlier orders only.
3. **Line 16.** A time split: the test set is the future, as in production.

Compute every feature as of the row's date, join with `merge_asof`, and split by time.

Recap card

- **Prior-only running total** `groupby(key)[col].cumsum() - col`, after sorting by key and date.
- **Point-in-time join in pandas** `pd.merge_asof(left, right, left_on, right_on, by=..., allow_exact_matches=False)`, both sorted.

- **merge_asof default for equal timestamps**

allow_exact_matches=True: same-time rows match. Set False when they would not be known yet.

- **Split for forward-looking models** By time: train before a cut-off, test after. TimeSeriesSplit for CV.

Practice

Exercise 1. Multiple choice (level 1)

A model will score customers each month, going forward. How should you split the historical data for validation?

- a) Random 80/20 split
- b) Train on months up to a cut-off, test on later months
- c) Test on the earliest months, train on later ones
- d) No split: use training accuracy

Why?

- r1) It matches production, where the model predicts a future it has not seen.
- r2) Recent data is always more accurate.

Exercise 2. Predict the output (level 2)

What does this print? (pandas 3.0)

```
1 import pandas as pd
2
3 preds = pd.DataFrame({"customer": ["c1"], "as_of":
4     pd.to_datetime(["2026-04-01"])}))
5 scores = pd.DataFrame({
6     "customer": ["c1", "c1"],
7     "scored_at": pd.to_datetime(["2026-02-15", "2026-04-01"]),
8     "nps": [3, 9],
9 })
10 out = pd.merge_asof(preds, scores, left_on="as_of", right_on="scored_at",
11     by="customer")
12 print(out["nps"].iloc[0])
```

Exercise 3. Fill the gap (level 2)

Attach the latest earlier score to each prediction row, excluding same-timestamp scores.

```
1 pd.__(1)(preds, scores, left_on="as_of", right_on="scored_at",
2     by="customer", allow_exact_matches=__(2))
```

Exercise 4. Write the function (level 3)

Write `add_prior_spend(orders) -> pd.DataFrame`. `orders` has `customer`, `ordered_at` and `spend`, in any row order. Return the orders sorted by `customer` then `ordered_at`, with two new columns:

- `prior_spend`: total spend by the same customer on **earlier** orders (0 for the first).
- `prior_orders`: number of earlier orders by the same customer (0 for the first).

Starter code:

```
1 import pandas as pd
2
3
4 def add_prior_spend(orders):
5     totals = orders.groupby("customer")["spend"].transform("sum")
6     return orders.assign(prior_spend=totals)
```

Exercise 5. Spot the bug (level 4)

For each prediction, the feature should be the customer's most recent NPS score strictly **before** `as_of`. Validation looks too good. Tap the line that leaks and fix it.

```
1 import pandas as pd
2
3
4 def latest_nps(preds: pd.DataFrame, scores: pd.DataFrame) -> pd.DataFrame:
5     """Most recent NPS strictly before each prediction date."""
6     joined = preds.merge(scores, on="customer")
7     joined = joined.sort_values("scored_at").groupby(["customer",
8     "as_of"]).tail(1)
9     return joined[["customer", "as_of", "nps"]].sort_values(["customer",
10    "as_of"]).reset_index(drop=True)
```

Write down the line number of the bug. Then write the fix.

Sources

- `pd.merge_asof` defaults to `direction="backward"` and `allow_exact_matches=True`. (verified: pandas API reference; run on pandas 3.0.6 in this repository.)

B11 Machine learning workflow

Roadmap: Prerequisite for all months. Splits, Pipeline, ColumnTransformer, cross-validation and hyperparameter search.

B11-L1 Train, validation and test, split by the right unit

Concept

Three sets, three jobs:

- **Training** rows fit the model.
- **Validation** rows, or cross-validation folds within training, choose between models and settings.
- **Test** rows measure the final choice, **once**. Every extra look at the test set turns it into a second validation set, and its score stops being an honest estimate.

Split so the test set looks like the data the model will meet.

Stratify on the target when a class is rare. `train_test_split(X, y, stratify=y)` keeps a 3% positive rate at 3% in both sets. Without it, a small test set can end up with very few positives.

Group when rows are not independent. A panel respondent answering five waves appears in five rows. A random row split puts some of their waves in training and others in test. The model then recognises the **person** instead of learning the pattern, and the test score measures memory.

A demonstration from this course's build (scikit-learn 1.8, random forest): 200 simulated respondents, 5 waves each, a label that depends only on the person and is unrelated to the features. Random 5-fold CV reported accuracy **0.995**. `GroupKFold` by respondent reported **0.468**: chance, which is the truth.

```
1 from sklearn.model_selection import GroupKFold, GroupShuffleSplit
2
3 splitter = GroupShuffleSplit(n_splits=1, test_size=0.2, random_state=7)
4 train_idx, test_idx = next(splitter.split(df, groups=df["respondent_id"]))
```

The group is whatever repeats: respondent, household, store, patient, device. If it will be **new** in production, keep it out of training in validation too. For time-ordered use, split by time as well (lesson B10-L3).

Record the split: the seed, the group column, and the resulting row counts, in the run manifest (lesson B6-L2).

Why it matters

A model predicting survey drop-out was validated with a random row split over a panel dataset. It scored AUC 0.93. Deployed on a newly recruited panel, where every respondent was new, it scored 0.61. Each respondent's answering style had let the model match test rows to the same person's training rows. A group split would have shown 0.61 before deployment.

Worked example: weak code

```
❶ train, test = train_test_split(panel, test_size=0.2)
2 for params in grid:
❷     score = evaluate(fit(train, params), test)
```

1. **Line 1.** Random rows: the same respondent sits on both sides.
2. **Line 3.** Tuning on the test set: its score is no longer an estimate of anything.

Worked example: strong code

```
1 import pandas as pd
2 from sklearn.model_selection import GroupShuffleSplit
3
4
5 def split_by_group(
6     df: pd.DataFrame, group_col: str, test_size: float, seed: int
7 ) -> tuple[pd.DataFrame, pd.DataFrame]:
8     """Split so that no group appears in both train and test."""
❶     splitter = GroupShuffleSplit(n_splits=1, test_size=test_size,
        random_state=seed)
❷     train_idx, test_idx = next(splitter.split(df, groups=df[group_col]))
11    return df.iloc[train_idx], df.iloc[test_idx]
```

1. **Line 9.** test_size is the share of groups, not rows. The seed makes the split reproducible.
2. **Line 10.** Every row of a respondent lands on the same side.

Tune on validation, measure once on test, and split by the unit that will be new in production.

Recap card

- **Roles of train, validation and test** Fit, choose, measure once.
- **Rows repeat per respondent** Split by group (GroupKFold, GroupShuffleSplit), or the model learns people.
- **Group leakage demonstration** Random CV 0.995, GroupKFold 0.468 on a label unrelated to the features.
- **Rare class** stratify=y keeps the class share in each set.

Practice

Exercise 1. Multiple choice (level 1)

You tried 40 model settings and picked the one with the best test score. What is that test score now?

- a) An unbiased estimate of performance on new data.
- b) An optimistic estimate, because the test set chose the model.
- c) A pessimistic estimate.
- d) The same as the training score.

Why?

- r1) Picking the best of 40 on the same data selects for luck on that data.
- r2) Test sets always overstate performance.

Exercise 2. Predict the output (level 2)

What does this print? (scikit-learn 1.8)

```
1 import pandas as pd
2 from sklearn.model_selection import GroupShuffleSplit
3
4 df = pd.DataFrame({"rid": ["a", "a", "b", "b", "c", "c", "d", "d"], "x":
    range(8)})
5 splitter = GroupShuffleSplit(n_splits=1, test_size=0.25, random_state=0)
6 train_idx, test_idx = next(splitter.split(df, groups=df["rid"]))
7 print(len(test_idx), df.iloc[test_idx]["rid"].nunique())
```

Exercise 3. Fill the gap (level 2)

Keep the rare positive class at the same rate in train and test.

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ____ (1)=y, ____ (2)=42)
```

Exercise 4. Write the function (level 3)

Write

`split_by_group(df, group_col, test_size, seed) -> tuple[pd.DataFrame, pd.DataFrame]` with `GroupShuffleSplit`, so that no value of `group_col` appears in both returned frames. The same seed must give the same split.

Starter code:

```
1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3
4
5 def split_by_group(df, group_col, test_size, seed):
6     return train_test_split(df, test_size=test_size, random_state=seed)
```

Exercise 5. Architecture design (level 4)

Choose the split for each project.

Options: Group split by member, Time split, Stratified random split.

1. Predict drop-out for newly recruited panel members; data has 6 waves per member
2. Fraud flag at 0.8% prevalence, one row per independent transaction, no time drift
3. Forecast next month's store sales

Sources

- Random 5-fold CV gave 0.995 and GroupKFold gave 0.468 on 200 simulated respondents with 5 waves each and a person-level random label. (verified: Run with scikit-learn 1.8.0 and 1.9.1, RandomForestClassifier(n_estimators=100, random_state=0), in this repository.)

B11-L2 Pipeline and ColumnTransformer

Concept

Lesson B10-L2 fitted preprocessing on training rows by hand. scikit-learn's `Pipeline` does it for you, every time, including inside each cross-validation fold. Examples use scikit-learn 1.8.

A `Pipeline` chains steps. Calling `fit` fits each step on the training rows and passes its output to the next. Calling `predict` runs the fitted steps on new rows **without refitting**. Cross-validation, grid search and `cross_val_score` all refit the whole pipeline per fold, so no fold's test rows leak into its preprocessing.

A `ColumnTransformer` applies different preprocessing to different columns:

```
1 prep = ColumnTransformer([
2     ("num", make_pipeline(SimpleImputer(strategy="median"),
3     StandardScaler()), NUMERIC),
4     ("cat", make_pipeline(SimpleImputer(strategy="most_frequent"),
5     OneHotEncoder(handle_unknown="ignore"))),
6     CATEGORICAL),
7 ])
8 model = Pipeline([("prep", prep), ("clf",
9     LogisticRegression(max_iter=1000))])
```

Three defaults to know:

- `OneHotEncoder(handle_unknown="error")` is the default. A region that did not appear in training raises `ValueError` at prediction time. Set `handle_unknown="ignore"` to encode it as all zeros.
- `ColumnTransformer(remainder="drop")` is the default. Columns you do not list are **silently dropped**. List every column you mean to use, or set `remainder="passthrough"` on purpose.
- Output feature names come from `get_feature_names_out()`, prefixed with the step name: `num__age`, `cat__region_n`.

Keep the column lists in config (lesson A1-L3), and have one function build the pipeline. The same fitted object is what you save, test and deploy, so preprocessing in production is exactly what was validated.

Why it matters

A lead-scoring model was trained in a notebook. Preprocessing ran in a separate cell: median imputation and scaling by hand. The deployment team rewrote the preprocessing from the notebook for production and imputed with the **mean**. Income was skewed, so the mean was 40% above the median. Scores for every lead with missing income shifted, and the sales team's call list changed without anyone noticing.

A single fitted `Pipeline`, saved and deployed as one object, leaves no second copy of the preprocessing to get wrong.

Worked example: weak code

```
❶ X["age"] = X["age"].fillna(X["age"].median())
❷ X = pd.get_dummies(X, columns=["region"])
3 scores = cross_val_score(LogisticRegression(), X, y, cv=5)
```

1. **Line 1.** Imputation on all rows, before CV: every fold's test rows shape the median.
2. **Line 2.** Dummy columns depend on the data seen. A new region in production changes the columns.

Worked example: strong code

```
1 from sklearn.compose import ColumnTransformer
2 from sklearn.impute import SimpleImputer
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.pipeline import Pipeline, make_pipeline
5 from sklearn.preprocessing import OneHotEncoder, StandardScaler
6
7
```



```

8 def build_pipeline(numeric: list[str], categorical: list[str]) ->
  Pipeline:
9     """Preprocessing and model as one object, fitted per fold by CV."""
10    prep = ColumnTransformer(
11        [
12            ("num", make_pipeline(SimpleImputer(strategy="median"),
13                                  StandardScaler()), numeric),
14            ("cat",
15             make_pipeline(SimpleImputer(strategy="most_frequent"),
16                           OneHotEncoder(handle_unknown="ignore")),
17             categorical,
18        ]
19    )
20    return Pipeline([("prep", prep), ("clf",
                                     LogisticRegression(max_iter=1000))])

```

1. **Line 15.** An unseen category becomes all zeros instead of an error.
2. **Line 20.** One object to cross-validate, save and deploy.

Put every learning step in one Pipeline, split columns with a ColumnTransformer, and deploy the fitted object itself.

Recap card

- **Why Pipeline** Every learning step refits per CV fold, and one fitted object goes to production.
- **OneHotEncoder default for unseen categories**
handle_unknown="error": raises ValueError. Use "ignore" for production.
- **ColumnTransformer default for unlisted columns** remainder="drop": silently dropped.
- **Feature names after preprocessing** get_feature_names_out(), prefixed with step names such as num__age.

Practice

Exercise 1. Multiple choice (level 1)

Why does putting a scaler inside a Pipeline prevent leakage in cross_val_score?

- a) The Pipeline refits the scaler on each fold's training rows only.
- b) Pipelines do not scale test data.
- c) The scaler is fitted once on all rows, but faster.
- d) Pipelines remove outliers automatically.

Why?

- r1) cross_val_score clones and fits the whole pipeline per fold.
- r2) Scalers inside pipelines have no parameters.

Exercise 2. Predict the output (level 2)

What does this print? (scikit-learn 1.8)

```
1 import pandas as pd
2 from sklearn.preprocessing import OneHotEncoder
3
4 enc = OneHotEncoder().fit(pd.DataFrame({"region": ["n", "s"]}))
5 try:
6     enc.transform(pd.DataFrame({"region": ["e"]}))
7     print("encoded")
8 except ValueError:
9     print("ValueError")
```

Exercise 3. Fill the gap (level 2)

Keep unlisted columns instead of dropping them, and encode unseen categories as zeros.

```
1 prep = ColumnTransformer(
2     [("cat", OneHotEncoder(handle_unknown="____(1)"), ["region"])],
3     remainder="____(2)",
4 )
```

Exercise 4. Write the function (level 3)

Write `build_pipeline(numeric, categorical) -> Pipeline` with two named steps:

- "prep": a `ColumnTransformer` with "num" (median `SimpleImputer`, then `StandardScaler`) on `numeric`, and "cat" (most-frequent `SimpleImputer`, then `OneHotEncoder(handle_unknown="ignore")`) on `categorical`.
- "clf": `LogisticRegression(max_iter=1000)`.

Starter code:

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.pipeline import Pipeline
3
4
5 def build_pipeline(numeric, categorical):
6     return Pipeline([("clf", LogisticRegression())])
```

Exercise 5. Spot the bug (level 4)

The model ignores `tenure`, although it is the strongest predictor in exploration. Tap the line that causes it, and fix the function.

```

1 from sklearn.compose import ColumnTransformer
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.pipeline import Pipeline
4 from sklearn.preprocessing import OneHotEncoder
5
6
7 def build(categorical: list[str]) -> Pipeline:
8     """Encode categorical columns and fit a classifier on every
        feature."""
9     prep = ColumnTransformer([("cat",
10        OneHotEncoder(handle_unknown="ignore"), categorical)])
11     return Pipeline([("prep", prep), ("clf",
12        LogisticRegression(max_iter=1000))])

```

Write down the line number of the bug. Then write the fix.

Sources

- In scikit-learn 1.8, OneHotEncoder defaults to `handle_unknown="error"` and ColumnTransformer to `remainder="drop"`. (verified: scikit-learn 1.8 API reference; signatures checked on 1.8.0 and 1.9.1 in this repository.)

B11-L3 Cross-validation and honest hyperparameter search

Concept

Cross-validation (CV) splits the training data into k folds, trains on $k - 1$ and scores on the remaining one, k times. Report the **mean and the spread** across folds. A mean AUC of 0.81 with fold scores from 0.70 to 0.90 is a less certain result than 0.81 with scores from 0.79 to 0.83.

`cross_validate` (scikit-learn 1.8) returns several metrics in one pass:

```

1 res = cross_validate(pipe, X, y, cv=StratifiedKFold(5, shuffle=True,
2     random_state=7),
3     scoring=["roc_auc", "average_precision"])
4 res["test_roc_auc"].mean(), res["test_roc_auc"].std()

```

Hyperparameter search. `GridSearchCV` or `RandomizedSearchCV` tries settings and keeps the best by CV score. Inside a pipeline, name parameters as `step_param`: `{"clf_C": [0.01, 0.1, 1, 10]}`.

The catch: `best_score_` is the **best of many tries** on the same folds. It is optimistic, for the same reason as picking a model on the test set (lesson B11-L1). A simulation in this course's build: 20 datasets of pure noise, a k -nearest-neighbours grid of 60 settings. The mean `best_score_` was **0.598** accuracy. The honest estimate was **0.515**.

Get the honest estimate with **nested CV**: an outer CV loop scores the **whole search** as if it were one model.

```
1 search = GridSearchCV(pipe, grid, cv=inner, scoring="roc_auc")
2 honest = cross_val_score(search, X, y, cv=outer, scoring="roc_auc")
```

Or keep a held-out test set that the search never sees, and score the refitted best model on it once.

Practical rules:

- Fix `random_state` in every splitter, so reruns compare like with like.
- Search on a log scale for regularisation strengths: 0.001, 0.01, 0.1, 1, 10.
- Prefer `RandomizedSearchCV` when the grid has more than about 50 combinations.
- Report the chosen settings, the honest score and the fold spread in the run log.

Why it matters

A team compared 300 configurations of a gradient-boosted model and reported the top `best_score_`, AUC 0.842, to the client as expected performance. The model went live and scored 0.80 over the first quarter. Nothing had drifted. The 0.842 had been the luckiest of 300 draws on the same five folds. A nested CV estimate of 0.80 would have set the right expectation from the start.

Worked example: weak code

```
❶ search = GridSearchCV(pipe, {"clf__C": [0.001, 0.01, 0.1, 1, 10, 100]},
                          cv=5)
2 search.fit(X, y)
❷ print("expected AUC", search.best_score_)
```

1. **Line 1.** Default scoring is accuracy for classifiers, not the AUC being reported.
2. **Line 3.** The best of several tries on the same folds: optimistic, and the wrong metric.

Worked example: strong code

```
1 import numpy as np
2 from sklearn.model_selection import GridSearchCV, StratifiedKFold,
  cross_val_score
3 from sklearn.pipeline import Pipeline
4
5
6 def nested_cv_auc(pipe: Pipeline, X: np.ndarray, y: np.ndarray, grid:
  dict[str, list[float]], seed: int) -> float:
```

```

7     """Honest AUC of the whole search: the search is refitted inside each
   outer fold."""
8     inner = StratifiedKFold(n_splits=3, shuffle=True, random_state=seed)
9     outer = StratifiedKFold(n_splits=5, shuffle=True, random_state=seed +
   1)
10    search = GridSearchCV(pipe, grid, cv=inner, scoring="roc_auc")
11    return float(cross_val_score(search, X, y, cv=outer,
   scoring="roc_auc").mean())

```

1. **Line 10.** The metric you will report is the metric you tune on.
2. **Line 11.** Outer folds never take part in choosing the settings.

Report CV mean and spread, tune with the reported metric, and estimate performance with nested CV or an untouched test set.

Recap card

- **Report CV results as** Mean and spread across folds, for the metric you care about.
- **Pipeline parameter names in a grid** `step__param`, such as `clf__C`. `make_pipeline` names steps after their class.
- **Is best_score_ an honest estimate?** No. Noise simulation: 0.598 `best_score_` against 0.515 nested. Use nested CV or a held-out test set.
- **Nested CV in scikit-learn** `cross_val_score(GridSearchCV(pipe, grid, cv=inner), X, y, cv=outer)`.

Practice

Exercise 1. Multiple choice (level 1)

After `GridSearchCV` over 80 settings, what is `best_score_`?

- a) An unbiased estimate of performance on new data
- b) The mean CV score of the best setting, optimistic because it was selected as the best
- c) The test-set score
- d) The training score of the refitted model

Why?

- r1) Selecting the maximum of many noisy scores favours settings that were lucky on these folds.
- r2) `GridSearchCV` always overfits.

Exercise 2. Predict the output (level 2)

What does this print? (scikit-learn 1.8)

```

1 from sklearn.linear_model import LogisticRegression
2 from sklearn.pipeline import make_pipeline
3 from sklearn.preprocessing import StandardScaler
4
5 pipe = make_pipeline(StandardScaler(), LogisticRegression())
6 print("clf_C" in pipe.get_params(), "logisticregression_C" in
      pipe.get_params())

```

Exercise 3. Fill the gap (level 2)

Tune C of the step named clf, scoring by ROC AUC.

```

1 search = GridSearchCV(pipe, {"____(1)": [0.01, 0.1, 1, 10]}, cv=inner,
      ____ (2)="roc_auc")

```

Exercise 4. Write the function (level 3)

Write `nested_cv_auc(X, y, c_values, seed) -> float`. Build a pipeline with steps "scale" (StandardScaler) and "clf" (LogisticRegression(max_iter=1000)). Wrap it in GridSearchCV over {"clf_C": c_values} with an inner StratifiedKFold(n_splits=3, shuffle=True, random_state=seed) and scoring="roc_auc". Return the mean of cross_val_score of the search with an outer StratifiedKFold(n_splits=5, shuffle=True, random_state=seed + 1), also scored by ROC AUC.

Starter code:

```

1 from sklearn.linear_model import LogisticRegression
2 from sklearn.model_selection import GridSearchCV
3
4
5 def nested_cv_auc(X, y, c_values, seed):
6     search = GridSearchCV(LogisticRegression(), {"C": c_values},
7       cv=5).fit(X, y)
8     return search.best_score_

```

Exercise 5. Architecture design (level 4)

You must tune a model and tell the client what AUC to expect. You have 8,000 labelled rows. Choose the plan.

Options: Hold out 20%, untouched until the end, Use all rows for the search.

1. Held-out test set
2. Search on the other 80%
3. Number given to the client

Sources

- Across 20 noise datasets with a 60-setting kNN grid, mean best_score_ was 0.598 and mean nested CV accuracy 0.515. (verified: Simulated with scikit-learn 1.9.1 in this repository.)
- cross_validate returns test_<metric> keys for each scorer. (verified: Run on scikit-learn 1.8.0 in this repository.)

B12 Model evaluation

Roadmap: Prerequisite for all months; Months 4 to 6: CNNs; Months 10 to 12: Transformers applied. Metric choice, calibration, imbalanced classes and error analysis.

B12-L1 Choose the metric from the decision

Concept

A metric is a summary of the mistakes you care about. Choose it from the decision the model supports, before you train anything.

Accuracy fails on imbalanced data. A test set from this course's build has 2.4% positives. Predicting "negative" for everyone scores accuracy **0.976**. A logistic regression at the default 0.5 threshold scored 0.9755, with **recall 0.0**: it found no positives at all.

Better summaries, all in scikit-learn 1.8:

- **ROC AUC** (`roc_auc_score`): the chance a random positive ranks above a random negative. It ignores the threshold and the class balance. Good for comparing rankers. It hides how many false positives come with each true positive when positives are rare.
- **Average precision**, the area under the precision-recall curve (`average_precision_score`): sensitive to rare positives. A random ranking scores the **prevalence**, 0.025 here. The same logistic regression scored **0.30**, twelve times the baseline. Always quote the baseline next to it.
- **Precision and recall at a threshold**: precision is the share of flagged cases that are real. Recall is the share of real cases that are flagged. These map straight onto work: "of 100 customers we call, 30 were going to churn" and "we reach 64% of the churners".
- **Cost-weighted metrics**: if a missed fraud costs £400 and a false alarm costs £5 of analyst time, compute the total cost at each threshold and pick the cheapest.

Pick the **threshold** from the decision too. 0.5 is a default of the library, not a business rule. If the team can review 200 cases a day, the threshold is whatever flags 200 cases.

Report the metric with the prevalence and the number of positives in the test set. A precision of 0.30 from 12 positives is not the same evidence as from 1,200.

Why it matters

A claims triage model was reported as "97.6% accurate". It flagged nothing. Fraud prevalence was 2.4%, so a model that approved every claim scored the same. The number went into a board paper before anyone asked how many frauds it caught. The answer was zero.

Worked example: weak code

```
❶ pred = model.predict(X_test)
❷ print("accuracy", accuracy_score(y_test, pred))
```

1. **Line 1.** predict() uses a 0.5 threshold that nobody chose.
2. **Line 2.** Accuracy on rare positives rewards predicting the majority class.

Worked example: strong code

```
1 import numpy as np
2 from sklearn.metrics import average_precision_score, roc_auc_score
3
4
5 def evaluate(y_true: np.ndarray, scores: np.ndarray, threshold: float) ->
    dict[str, float | None]:
6     """Ranking and threshold metrics, with the baseline to compare
    against."""
7     flagged = scores >= threshold
8     tp = int((flagged & (y_true == 1)).sum())
9     n_flagged, n_pos = int(flagged.sum()), int((y_true == 1).sum())
10    return {
11        "prevalence": float(y_true.mean()),
12        "roc_auc": float(roc_auc_score(y_true, scores)),
13        "average_precision": float(average_precision_score(y_true,
    scores)),
14        "precision": tp / n_flagged if n_flagged > 0 else None,
15        "recall": tp / n_pos if n_pos > 0 else None,
16    }
```

1. **Line 11.** The baseline for average precision, reported alongside it.
2. **Line 14.** No flags, no precision: None, not a division error or a misleading 0.

Choose metrics and thresholds from the decision. Report prevalence and counts beside every number.

Recap card

- **Accuracy at 2.4% prevalence** 0.976 for predicting all negatives. Useless as a headline.
- **Baseline for average precision** The prevalence. Report it beside the score.
- **Where the threshold comes from** The decision: review capacity or costs. Not the 0.5 default.
- **Precision and recall** Precision: share of flags that are real. Recall: share of real cases flagged.

Practice

Exercise 1. Multiple choice (level 1)

Positives are 2.4% of the test set. A model has accuracy 0.976. What do you know?

- a) The model is excellent.
- b) Nothing useful yet. Predicting all negatives scores the same.
- c) The model finds 97.6% of positives.
- d) The model is calibrated.

Why?

- r1) Accuracy is dominated by the 97.6% majority class.
- r2) Accuracy is never useful.

Exercise 2. Predict the output (level 2)

What does this print, to two decimals? (scikit-learn 1.8, NumPy 2.4)

```
1 import numpy as np
2 from sklearn.metrics import average_precision_score
3
4 rng = np.random.default_rng(0)
5 y = (rng.random(20000) < 0.05).astype(int)
6 random_scores = rng.random(20000)
7 print(round(average_precision_score(y, random_scores), 2))
```

Exercise 3. Fill the gap (level 2)

Flag the top 200 cases a day, whatever their scores.

```
1 threshold = np.sort(scores)[____(1)]
2 flagged = scores >= ____ (2)
```

Exercise 4. Write the function (level 3)

Write `evaluate(y_true, scores, threshold) -> dict[str, float | None]` returning:

- `prevalence`: mean of `y_true`.
- `roc_auc` and `average_precision` from `scikit-learn`.
- `precision` and `recall` at `scores >= threshold`, computed from counts. Return `None` for precision when nothing is flagged, and for recall when there are no positives.

Starter code:

```
1 from sklearn.metrics import accuracy_score
2
3
4 def evaluate(y_true, scores, threshold):
5     return {"accuracy": accuracy_score(y_true, scores >= threshold)}
```

Exercise 5. Architecture design (level 4)

Choose the headline metric for each model.

Options: Precision among the top k flagged, Accuracy, ROC AUC, Total cost at the chosen threshold.

1. Fraud: 1% prevalence, analysts review 150 cases a day
2. Compare two ranking models before any threshold exists
3. Loan approval: a default costs £4,000, a wrongly refused customer costs £300 of margin

Sources

- On `make_classification` data with 2.4% positives, all-negative accuracy was 0.976, logistic regression at 0.5 had recall 0.0, average precision 0.30 against a random 0.025. (verified: Run with `scikit-learn` 1.8.0 and 1.9.1 in this repository (`random_state=0`).)

B12-L2 Calibration, when the probability itself matters

Concept

A model is **calibrated** when its probabilities mean what they say: of all customers scored 0.30, about 30% churn. Ranking metrics such as ROC AUC ignore calibration. A model can rank perfectly and still say 0.9 where the truth is 0.1.

Calibration matters whenever someone uses the number itself: expected revenue (`probability × value`), a forecast count of churners, a risk shown to a person, or thresholds agreed in probability terms.

Check it with a calibration table: group predictions into bins, then compare the mean predicted probability with the observed rate in each bin. `scikit-learn 1.8` gives the same with `calibration_curve(y, p, n_bins=10)`. Plot observed against predicted. A calibrated model follows the diagonal.

Summarise with the **Brier score** (`brier_score_loss`), the mean squared error of the probabilities. Compare it with a constant model that always predicts the prevalence.

A common cause of bad calibration is **reweighting for imbalance**. In this course's build, a logistic regression on data with 2.4% positives had a mean predicted probability of 0.024 and a Brier score of 0.021. The same model with `class_weight="balanced"` had a mean predicted probability of **0.42**, and a Brier score of **0.199**, worse than predicting 0.024 for everyone (0.024). Its AUC was slightly higher. Its probabilities were useless as probabilities.

Fix it with `CalibratedClassifierCV(model, method="sigmoid", cv=5)`, which fits a calibration map on held-out folds. Use `method="isotonic"` only with plenty of data, roughly a thousand positives or more, because it overfits small samples. Check calibration again on the test set after fitting.

If you only need a ranking or a top-k list, calibration does not matter. Say which one you deliver.

Why it matters

A retention team estimated campaign value as the sum of churn probabilities times customer margin. The model had been trained with balanced class weights. Its average churn probability was 0.41 on a base that churned at 3% a month. The business case promised ten times the saves the campaign could deliver, and budget was approved on it.

Worked example: weak code

```
❶ model = LogisticRegression(class_weight="balanced").fit(X_train, y_train)
❷ expected_churners = model.predict_proba(X_live)[: , 1].sum()
```

1. **Line 1.** Balanced weights push probabilities towards 0.5.
2. **Line 2.** Sums miscalibrated probabilities as if they were real rates.

Worked example: strong code

```
1 import numpy as np
```

```

2 import pandas as pd
3
4
5 def calibration_table(y: np.ndarray, p: np.ndarray, n_bins: int) ->
    pd.DataFrame:
6     """Mean predicted and observed rate per equal-width bin. Empty bins
    are left out."""
7     edges = np.linspace(0.0, 1.0, n_bins + 1)
8     idx = np.searchsorted(edges[1:-1], p, side="right")
9     df = pd.DataFrame({"bin": idx, "y": y, "p": p})
10    table = df.groupby("bin").agg(mean_predicted=("p", "mean"),
    observed=("y", "mean"), n=("y", "size"))
11    return table.reset_index(drop=True)

```

1. **Line 8.** Assigns each probability to an equal-width bin.

2. **Line 10.** Predicted against observed, with the count behind each row.

When the probability is used as a number, check calibration, compare the Brier score with a constant baseline, and recalibrate on held-out data.

Recap card

- **Calibrated means** Of cases scored p , about p are positive.
- **Check calibration** Calibration table or calibration_curve, plus Brier score against the constant prevalence baseline.
- **class_weight="balanced" and probabilities** Pushes them towards 0.5. Mean p went from 0.024 to 0.42 at 2.4% prevalence.
- **Recalibrate** CalibratedClassifierCV(model, method="sigmoid", cv=5). Isotonic only with plenty of data.

Practice

Exercise 1. Multiple choice (level 1)

A model has ROC AUC 0.85. Of customers it scores between 0.8 and 0.9, 20% churn. What is true?

- a) It ranks well but is poorly calibrated.
- b) It is well calibrated.
- c) The AUC must be wrong.
- d) It ranks badly.

Why?

- r1) AUC measures ordering only. The probabilities can be far off while the order is right.
- r2) Any model with AUC above 0.8 is calibrated.

Exercise 2. Predict the output (level 2)

What does this print? (scikit-learn 1.8)

```
1 from sklearn.metrics import brier_score_loss
2
3 y = [0, 0, 0, 1]
4 print(brier_score_loss(y, [0.25, 0.25, 0.25, 0.25]), brier_score_loss(y,
    [0.5, 0.5, 0.5, 0.5]))
```

Exercise 3. Fill the gap (level 2)

Recalibrate a classifier with a sigmoid map fitted on held-out folds.

```
1 calibrated = ____ (1) (model, method="____ (2)", cv=5).fit(X_train,
    y_train)
```

Exercise 4. Write the function (level 3)

Write `calibration_table(y, p, n_bins) -> pd.DataFrame`. Split `[0, 1]` into `n_bins` equal-width bins. Assign each probability with `np.searchsorted(edges[1:-1], p, side="right")`. Return one row per **non-empty** bin, in bin order, with columns `mean_predicted`, `observed` (mean of `y`) and `n`, and a default 0..k index.

Starter code:

```
1 import numpy as np
2 import pandas as pd
3
4
5 def calibration_table(y, p, n_bins):
6     return pd.DataFrame({"mean_predicted": [p.mean()], "observed":
    [y.mean()]})
```

Exercise 5. Architecture design (level 4)

A churn model will be used two ways. Decide what each use needs.

Options: Good ranking only, Calibrated probabilities.

1. Pick the top 500 customers to call each week
2. Forecast the number of churners next month for finance
3. Class imbalance handling for the forecast model

Sources

- Logistic regression at 2.4% prevalence: mean p 0.024 and Brier 0.021; with `class_weight="balanced"`, mean p 0.42 and Brier 0.199; constant baseline 0.024. (verified: Run with scikit-learn 1.8.0 and 1.9.1 in this repository (make_classification, random_state=0).)

- The table matches `sklearn.calibration.calibration_curve` with `strategy="uniform"` on test data. (verified: Compared with `np.allclose` on `scikit-learn 1.8.0` in this repository.)

B12-L3 Imbalanced classes and error analysis by segment

Concept

Imbalance. When positives are rare, a model trained as usual predicts low probabilities, and the default 0.5 threshold flags almost nothing. That is a threshold problem, not always a model problem.

In this course's build (2.4% positives, logistic regression, `scikit-learn 1.8`), the plain model at 0.5 flagged 3 cases and found none. Lowering the threshold to 0.1 flagged 284 cases with precision 0.36 and recall 0.42. At 0.02 it reached recall 0.63 with precision 0.04. The same model trained with `class_weight="balanced"` reached recall 0.64 and precision 0.05 at 0.5: a very similar trade-off, with probabilities that were no longer calibrated (lesson B12-L2).

Options, in the order to try them:

1. **Train normally and choose the threshold** from capacity or cost (lesson B12-L1). Probabilities stay calibrated.
2. **Class weights** when the model struggles to learn the minority class at all. Recalibrate afterwards if you need probabilities.
3. **Resampling** (over- or under-sampling) **inside the training folds only**.
Never resample the test set: its class balance must match production.

Error analysis. One overall metric hides where a model fails. Break the test results down by segment: region, device, customer tenure, survey mode. For each segment report the count, positives, precision and recall. A model with recall 0.60 overall can have recall 0.15 for new customers, the group the campaign most wants to reach.

Then read examples. Take the 20 most confident false positives and the 20 most confident false negatives, and look at them with someone who knows the domain. Patterns in those 40 rows often point to a missing feature, a label error or a leak.

Report segment results with their counts. A recall of 0.20 from 5 positives is a question, not a finding.

Why it matters

A complaint-routing model reported recall 0.71 overall. Complaints arriving by the new web chat were 18% of the volume and had recall 0.22: the model had been trained mostly on email text, and chat messages were shorter and informal. Urgent chat complaints waited in the general queue for three weeks before a supervisor noticed. A per-channel breakdown at validation would have shown it.

Worked example: weak code

```
1 model = LogisticRegression().fit(X_train, y_train)
2 print("recall", recall_score(y_test, model.predict(X_test)))
```

1. **Line 2.** The 0.5 default flags almost nothing on rare positives, and one overall number hides weak segments.

Worked example: strong code

```
1 import pandas as pd
2
3
4 def metrics_by_segment(df: pd.DataFrame, segment: str) -> pd.DataFrame:
5     """Count, positives, precision and recall per segment. Columns y and
6         flagged are 0/1."""
7     rows = []
8     for name, g in df.groupby(segment, sort=True):
9         tp = int((g["y"] == 1) & (g["flagged"] == 1)).sum()
10        n_flag, n_pos = int(g["flagged"].sum()), int(g["y"].sum())
11        rows.append({
12            "segment": name,
13            "n": len(g),
14            "positives": n_pos,
15            "precision": tp / n_flag if n_flag > 0 else None,
16            "recall": tp / n_pos if n_pos > 0 else None,
17        })
18     return pd.DataFrame(rows)
```

1. **Line 7.** One row per segment, in a stable order.
2. **Line 13.** Counts beside every rate, so small segments are read with care.
3. **Line 15.** No positives, no recall: None instead of a crash or a fake 0.

Treat rare positives with a chosen threshold first, then break results down by segment and read the worst errors.

Recap card

- **First response to rare positives** Train normally and choose the threshold from capacity or cost.
- **Resampling rule** Only inside training folds. Never the test set.

- **Error analysis** Metrics per segment with counts, then read the most confident errors.
- **Pooled vs segment recall** Pooling weights large segments. Overall 0.7 can hide a segment at 0.2.

Practice

Exercise 1. Multiple choice (level 1)

A model trained on 2% positives flags nothing at the default threshold. What should you try first?

- a) Oversample the test set to 50% positives.
- b) Lower the threshold to match review capacity or costs.
- c) Switch to a deep neural network.
- d) Report accuracy instead.

Why?

- r1) The model's ranking may be fine. Only the cut-off is wrong for rare positives.
- r2) Lower thresholds always improve models.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 segments = {"email": (80, 100), "chat": (4, 20)}
2 for name, (found, positives) in segments.items():
3     print(name, found / positives)
4 print("overall", sum(f for f, _ in segments.values()) / sum(p for _, p in
    segments.values()))
```

Exercise 3. Put in order (level 2)

Order the steps for handling a rare positive class.

- Read the most confident false positives and false negatives
- Break test results down by segment, with counts
- Train the model normally
- Choose the threshold from capacity or cost on validation data
- Split train and test, stratified, test left at its natural balance

Exercise 4. Write the function (level 3)

Write `metrics_by_segment(df, segment) -> pd.DataFrame`. `df` has 0/1 columns `y` and `flagged` and a segment column named by `segment`. Return one row

per segment, sorted by segment, with columns `segment` (named as the input column), `n`, `positives`, `precision` and `recall`. Use `None` for precision when a segment has no flags, and for recall when it has no positives.

Starter code:

```
1 import pandas as pd
2
3
4 def metrics_by_segment(df, segment):
5     return df.groupby(segment)["y"].mean().reset_index()
```

Exercise 5. Spot the bug (level 4)

Validation reports recall 0.78, but production recall is 0.41. Tap the line that makes the evaluation unrepresentative, and fix the function.

```
1 import pandas as pd
2
3
4 def prepare_sets(train: pd.DataFrame, test: pd.DataFrame, seed: int) ->
    tuple[pd.DataFrame, pd.DataFrame]:
5     """Oversample positives so the model sees more of them."""
6     pos_train = train[train["y"] == 1].sample(frac=5.0, replace=True,
        random_state=seed)
7     pos_test = test[test["y"] == 1].sample(frac=5.0, replace=True,
        random_state=seed)
8     return pd.concat([train, pos_train]), pd.concat([test, pos_test])
```

Write down the line number of the bug. Then write the fix.

Sources

- Plain logistic regression at 2.4% prevalence: threshold 0.1 gave precision 0.36, recall 0.42; threshold 0.02 gave recall 0.63, precision 0.04; balanced weights at 0.5 gave recall 0.64, precision 0.05. (verified: Run with scikit-learn 1.9.1 in this repository (make_classification, random_state=0).)

B13 Interpretability and fairness

Roadmap: Prerequisite for all months. Permutation importance, SHAP and subgroup performance.

B13-L1 Permutation importance on held-out data

Concept

"Which features matter?" has a trustworthy answer and a misleading one.

The misleading one is **impurity importance**, `feature_importances_` on scikit-learn tree ensembles. It is computed on training data and favours features with many distinct values. In this course's build (scikit-learn 1.9, random forest, 3,000 rows), a meaningless column of random IDs with 3,000 distinct values got an impurity importance of **0.335**, half as much as the real signal (0.661).

The trustworthy one is **permutation importance** on **held-out** data: shuffle one column, breaking its link with the target, and measure how much the score drops. A feature the model relies on causes a large drop. The same random ID column dropped the test AUC by **0.004**: nothing.

```
1 from sklearn.inspection import permutation_importance
2 result = permutation_importance(model, X_test, y_test, scoring="roc_auc",
3                               n_repeats=10, random_state=0)
4 result.importances_mean, result.importances_std
```

Read it with three cautions:

- **It measures the model, not the world.** A large drop means this model uses the feature. It does not mean the feature causes the outcome.
- **Correlated features share credit.** If `age` and `years_since_birth_cohort` carry the same information, shuffling one leaves the other, and both look unimportant. Group correlated features and permute them together, or report them as a group.
- **Use the spread.** Report the mean drop with its standard deviation across repeats. A drop of 0.003 ± 0.004 is zero.

Use the metric you report for the model. An importance in accuracy on imbalanced data says little (lesson B12-L1).

Why it matters

A retention model's slide showed "customer ID" as the third most important driver, from `feature_importances_`. The product team asked what it meant. It meant nothing: the ID had thousands of distinct values, so trees split on it to memorise training rows. Permutation importance on the test set gave it zero. The slide had already been shown to the client.

Worked example: weak code

```
1 importances = dict(zip(X_train.columns, model.feature_importances_))
2 top = sorted(importances, key=importances.get, reverse=True)[:5]
```

1. **Line 1.** Impurity importance from training data: biased towards high-cardinality features.

Worked example: strong code

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.base import ClassifierMixin
4 from sklearn.metrics import roc_auc_score
5
6
7 def permutation_drop(model: ClassifierMixin, X: pd.DataFrame, y:
  pd.Series, col: str, rng: np.random.Generator, n_repeats: int) ->
  float:
8     """Mean fall in held-out ROC AUC when col is shuffled."""
9     base = roc_auc_score(y, model.predict_proba(X[:, 1]))
10     drops = []
11     for _ in range(n_repeats):
12         shuffled = X.assign(**{col: rng.permutation(X[col].to_numpy())})
13         drops.append(base - roc_auc_score(y, model.predict_proba(shuffled)
14        [:, 1]))
14     return float(np.mean(drops))
```

1. **Line 9.** Score on held-out rows, with the reported metric.
2. **Line 12.** Breaks the link between one feature and the target, keeping its distribution.

Measure importance by shuffling on held-out data, with the reported metric, and report the spread.

Recap card

- **Impurity importance bias** Favours high-cardinality features, from training data. A random ID scored 0.335 in this lesson's example.
- **Permutation importance** Shuffle a column on held-out data and measure the score drop, with its spread.

- **Correlated features** Share credit. Permute and report them as a group.
- **What importance means** What the model relies on. Not what causes the outcome.

Practice

Exercise 1. Multiple choice (level 1)

A random ID column ranks high in a random forest's `feature_importances_`. What is the likely cause?

- a) IDs predict the outcome.
- b) Impurity importance favours features with many distinct values, measured on training data.
- c) The forest has too few trees.
- d) The ID was scaled incorrectly.

Why?

- r1) Many split points let trees memorise training rows, which impurity importance rewards.
- r2) Random forests cannot handle integers.

Exercise 2. Predict the output (level 2)

`age` and `birth_year` hold the same information. A model uses both. You permute only `age` on held-out data. What happens to its importance?

```
1 drop_age_alone = 0.01      # measured
2 drop_both_together = 0.12  # measured when both are shuffled together
3 print(drop_age_alone < drop_both_together)
```

Exercise 3. Fill the gap (level 2)

Compute permutation importance on the test set, scored by ROC AUC, with 10 repeats.

```
1 result = permutation_importance(model, _____(1), y_test,
    scoring="roc_auc", n_repeats=10, random_state=0)
2 print(result.importances_mean, result._____(2))
```

Exercise 4. Write the function (level 3)

Write `permutation_drop(model, X, y, col, rng, n_repeats) -> float`.

Compute the ROC AUC of `model.predict_proba(X)[: , 1]`. Then, `n_repeats` times, replace `col` with `rng.permutation(X[col].to_numpy())` in a **copy** of `X`, recompute the AUC, and record the drop. Return the mean drop. Leave `X` unchanged.

Starter code:

```
1 import numpy as np
2
3
4 def permutation_drop(model, X, y, col, rng, n_repeats):
5     return float(model.coef_[0][list(X.columns).index(col)])
```

Exercise 5. Architecture design (level 4)

A client asks for "the top drivers of churn" from your model. Choose what to deliver.

Options: Permutation importance on the test set, with the reported metric, feature_importances_ from training.

1. Importance method
2. tenure_months and contract_start_year are 0.97 correlated
3. Wording on the slide

Sources

- A random ID with 3,000 values got impurity importance 0.335 and permutation importance (ROC AUC, test) 0.004 in a 200-tree random forest. (verified: Run with scikit-learn 1.9.1 in this repository (random_state=0).)

B13-L2 SHAP values, and what they do not tell you

Concept

Permutation importance (lesson B13-L1) ranks features for the whole model.

SHAP values explain **one prediction**: how much each feature pushed this customer's score up or down from a baseline.

For one row, SHAP values are **additive**:

```
1 prediction = base_value + sum of the row's SHAP values
```

The base value is the model's average output over the background data. In shap 0.51:

```
1 import shap
2 explainer = shap.TreeExplainer(model)
3 explanation = explainer(X_test) # shap.Explanation
4 explanation.values, explanation.base_values
```

Know which **output space** the values live in. For a scikit-learn `GradientBoostingClassifier`, `TreeExplainer` explains the log-odds

(`decision_function`), not the probability. For a `RandomForestClassifier` it explains probabilities, with one set of values per class, so `.values` has shape `(rows, features, classes)`. This course's build checked both additivity identities on shap 0.51 with scikit-learn 1.9. Say which space you report.

Aggregate carefully. The mean **absolute** SHAP value per feature gives a global ranking. A beeswarm plot shows the direction and spread.

What SHAP does **not** tell you:

- **Causes.** A feature with a large SHAP value is one the model uses. Changing it in the real world may do nothing.
- **Robust truths about correlated features.** When two features carry the same information, the credit split between them depends on the model and the explainer's assumptions.
- **Anything about data the model never saw.** Explanations for out-of-range inputs describe extrapolation.

For customer-facing reasons ("your application was declined because..."), SHAP values are a starting point for wording. They are not a legal explanation on their own.

shap is not available in this app's browser runtime. The exercises use plain Python on SHAP outputs. Run shap itself locally.

Why it matters

A pricing team read a SHAP summary showing that `number_of_support_calls` pushed renewal probability **up**. They proposed encouraging customers to call support more. The model had learned that customers who called support were the engaged ones who had already decided to stay. More calls would not create engagement. An experiment on proactive calls later showed no effect on renewals.

Worked example: weak code

```
1 explanation = shap.TreeExplainer(rf_classifier)(X_test)
2 top = X_test.columns[np.argsort(explanation.values.mean(axis=0))[-3:]]
3 print("Top causes of churn:", list(top))
```

1. **Line 2.** Signed means cancel out, and a random forest's values have a class axis that is ignored here.
2. **Line 3.** Model reliance reported as causes.

Worked example: strong code

```
1 def top_contributions(names: list[str], values: list[float], k: int) ->
    list[tuple[str, float]]:
2     """The k largest contributions to one prediction, by absolute size,
    with their sign."""
3     if len(names) != len(values):
4         raise ValueError("one SHAP value per feature is required")
5     ranked = sorted(zip(names, values, strict=True), key=lambda pair:
    abs(pair[1]), reverse=True)
6     return ranked[:k]
7
8
9 def check_additivity(base_value: float, values: list[float], prediction:
    float, tol: float = 1e-6) -> bool:
10     """True when base value plus contributions reproduces the model
    output."""
11     return abs(base_value + sum(values) - prediction) <= tol
```

1. **Line 5.** Rank by size, keep the sign for the wording.
2. **Line 11.** A cheap test that you are reading the right output space.

SHAP explains one prediction additively in a stated output space. It describes the model, not the world.

Recap card

- **SHAP additivity** prediction = base_value + sum of the row's SHAP values, in the explainer's output space.
- **Output space for GradientBoostingClassifier in TreeExplainer** Log-odds (decision_function), not probability.
- **Global ranking from SHAP** Mean absolute SHAP value per feature.
- **What SHAP cannot show** Causes, a stable split between correlated features, behaviour outside the training range.

Practice

Exercise 1. Multiple choice (level 1)

A feature has a large positive SHAP value for most churners. What can you conclude?

- a) Increasing the feature causes churn.
- b) The model uses the feature to push churn scores up.
- c) The feature is measured wrongly.
- d) The model is overfitting.

Why?

- r1) SHAP values attribute the model's output. They say nothing about intervening in the world.
- r2) Positive SHAP values are always spurious.

Exercise 2. Predict the output (level 2)

A GradientBoostingClassifier row has base value -1.2 (log-odds) and SHAP values 0.8, 0.5 and -0.3. What does this print?

```
1 import math
2
3 log_odds = -1.2 + 0.8 + 0.5 - 0.3
4 print(round(log_odds, 2), round(1 / (1 + math.exp(-log_odds)), 2))
```

Exercise 3. Fill the gap (level 2)

Compute SHAP values for a tree model with shap 0.51.

```
1 explainer = shap.____(1)(model)
2 explanation = explainer(X_test)
3 contributions = explanation.____(2)
```

Exercise 4. Write the function (level 3)

Write `top_contributions(names, values, k) -> list[tuple[str, float]]`.

Return the `k` pairs `(name, value)` with the largest **absolute** values, largest first, keeping the sign. Raise `ValueError` if `names` and `values` differ in length.

Starter code:

```
1 def top_contributions(names, values, k):
2     return sorted(zip(names, values), key=lambda p: p[1], reverse=True)
   [:k]
```

Exercise 5. Architecture design (level 4)

Explain a declined loan application to the customer service team. Choose each part.

Options: State the space (log-odds) and translate the final score to a probability, Show log-odds contributions labelled as percentage points.

1. Output space shown
2. What the explanation lists
3. Wording

Sources

- shap 0.51 TreeExplainer: additivity holds in decision_function space for GradientBoostingClassifier and in probability space, per class, for RandomForestClassifier. (verified: Checked with np.allclose on shap 0.51.0 and scikit-learn 1.9.1 in this repository.)

B13-L3 Performance by subgroup, and the fairness trade-offs

Concept

A model can be accurate overall and fail one group. Before release, compute the key rates **per subgroup**: age band, region, sex, ethnicity where you hold it lawfully, and any group the decision affects differently.

Per group, report:

- **n** and **positives**, so readers can judge the evidence.
- **Selection rate**: the share flagged.
- **True positive rate** (recall): the share of real positives flagged.
- **False positive rate**: the share of real negatives flagged.
- **Precision**, and a calibration check (lesson B12-L2).

Then face a hard fact. When groups have **different base rates**, a calibrated score cannot also give equal false positive and true positive rates, except in trivial cases (Kleinberg, Mullainathan and Raghavan, 2016; Chouldechova, 2017).

A simulation in this course's build: a perfectly calibrated score for two groups with base rates of 30% and 10%, one threshold of 0.3. In both groups, cases scored 0.3 to 0.4 were positive about 34% of the time. Yet group A had a true positive rate of 0.65 and a false positive rate of 0.39. Group B had 0.15 and 0.03.

So choose, with the people accountable for the decision, which property matters most for this use:

- Equal **calibration**: a score means the same risk in every group.
- Equal **true positive rates**: each group's genuine cases are found at the same rate. Suits allocating help.
- Equal **false positive rates**: each group bears the same rate of wrong flags. Suits decisions with a cost to the flagged person.

Write the choice, the reasons and the measured gaps in the model card (lesson B18-L2).

Removing a protected attribute from the features does not remove its influence. Postcode, first name and shopping patterns act as **proxies**. Keep the attribute available for **measuring**, where lawful, even if the model never uses it.

Why it matters

A hardship-support model flagged customers for proactive help. Overall recall was 0.62. By age band, customers over 70 had recall 0.28: they used fewer digital channels, so the features the model relied on were sparse for them. The group most likely to need help was the least likely to be offered it. A per-group table at validation showed the gap in one line.

Worked example: weak code

```
❶ features = [c for c in df.columns if c not in ("age_band", "sex")]
❷ print("recall", recall_score(y_test, flagged))
```

1. **Line 1.** Dropping attributes does not remove proxies, and leaves nothing to measure gaps with.
2. **Line 2.** One overall number hides a group at a fraction of it.

Worked example: strong code

```
1 import pandas as pd
2
3
4 def group_rates(df: pd.DataFrame, group: str) -> pd.DataFrame:
5     """Selection rate, TPR and FPR per group, with counts. Columns y and
6         flagged are 0/1."""
7     rows = []
8     ❶ for name, g in df.groupby(group, sort=True):
9         pos, neg = g["y"] == 1, g["y"] == 0
10        flag = g["flagged"] == 1
11        rows.append({
12            "group": name,
13            "n": len(g),
14            "positives": int(pos.sum()),
15            "selection_rate": float(flag.mean()),
16            ❷ "tpr": float((flag & pos).sum() / pos.sum()) if pos.sum() > 0
17            else None,
18            "fpr": float((flag & neg).sum() / neg.sum()) if neg.sum() > 0
19            else None,
20        })
```

```
18 return pd.DataFrame( rows )
```

1. **Line 7.** The group attribute is used to measure, whether or not the model uses it.
2. **Line 15.** Guarded: a group with no positives has no TPR.

Measure rates per group with counts, choose the fairness property that fits the decision, and document the gaps.

Recap card

- **Rates to report per group** n, positives, selection rate, TPR, FPR, precision, calibration.
- **Impossibility result** With different base rates, calibration and equal error rates cannot all hold (Kleinberg et al. 2016; Chouldechova 2017).
- **Dropping a protected attribute** Proxies keep its influence. Keep it for measurement where lawful.
- **Choosing a fairness property** From the harm: missed help (TPR), wrong flags (FPR), numbers used directly (calibration).

Practice

Exercise 1. Multiple choice (level 1)

You remove `sex` from a model's features. What happens to sex-related disparities?

- a) They disappear.
- b) They can persist through proxies, and you lose the ability to measure them if you discard the column.
- c) They always get worse.
- d) The model becomes calibrated.

Why?

- r1) Other features correlate with sex, and measurement needs the attribute.
- r2) Removing features never changes a model.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 groups = {"A": {"tp": 65, "pos": 100, "fp": 150, "neg": 390}, "B": {"tp":  
    15, "pos": 100, "fp": 25, "neg": 900}}  
2 for name, g in groups.items():  
3     print(name, round(g["tp"] / g["pos"], 2), round(g["fp"] / g["neg"],  
    2))
```

Exercise 3. Fill the gap (level 2)

Compute one group's true positive rate and false positive rate from boolean Series `flag`, `pos` and `neg`.

```
1 tpr = (flag & _____.sum() / pos.sum()
2 fpr = (flag & _____.sum() / neg.sum()
```

Exercise 4. Write the function (level 3)

Write `group_rates(df, group) -> pd.DataFrame`. `df` has 0/1 columns `y` and `flagged`, and a group column. Return one row per group, sorted, with columns `<group>`, `n`, `positives`, `selection_rate`, `tpr` and `fpr`. Use `None` for `tpr` when a group has no positives, and for `fpr` when it has no negatives.

Starter code:

```
1 import pandas as pd
2
3
4 def group_rates(df, group):
5     return df.groupby(group)["flagged"].mean().reset_index()
```

Exercise 5. Architecture design (level 4)

Choose the fairness property to prioritise for each use.

Options: Equal true positive rates, Equal false positive rates, Equal calibration.

1. Offer proactive hardship support
2. Refer claims for a fraud investigation that delays payment
3. Risk scores used directly as probabilities in pricing

Sources

- Calibrated scores with base rates of 30% and 10% at threshold 0.3 gave TPR 0.65 and FPR 0.39 in one group, 0.15 and 0.03 in the other. (verified: Simulated with NumPy 2.4.6 in this repository (Beta(3,7) and Beta(1,9) risks, 200,000 rows each).)
- Calibration within groups and balance of error rates cannot all hold when base rates differ, except in trivial cases. (verified: Kleinberg, Mullainathan and Raghavan (2016), Inherent Trade-Offs in the Fair Determination of Risk Scores; Chouldechova (2017), Big Data 5(2).)

B14 Experiment tracking with MLflow

Roadmap: Prerequisite for all months; Months 16 to 18: MLOps and deployment. Tracking, registry and the model lifecycle.

B14-L1 Track every run with MLflow

Concept

After twenty experiments nobody remembers which settings produced the good score. **MLflow tracking** records each run: parameters, metrics, tags and files. Examples use MLflow 3.16. On Databricks the tracking server is built in. Locally, MLflow 3 stores runs in a SQLite file, `mlflow.db`, in the working directory by default.

```
1 import mlflow
2
3 mlflow.set_experiment("churn")
4 with mlflow.start_run(run_name="lr-c1"):
5     mlflow.log_params({"C": 1.0, "seed": 7, "features": "v3"})
6     mlflow.set_tags({"git_commit": commit, "input_sha256": data_hash})
7     model = train(X_train, y_train, C=1.0, seed=7)
8     mlflow.log_metric("cv_auc", cv_auc)
9     mlflow.sklearn.log_model(model, name="model", input_example=X_train[:
10     5])
```

What to record, every time:

- **Parameters:** every setting that changes the result, including the seed and the feature set version. MLflow stores them as **strings**: `1.0` comes back as `'1.0'`. A parameter cannot be changed once logged in a run. Logging `C` twice with different values raises `MLflowException`.
- **Tags:** the Git commit and the input data hash from the run manifest (lesson B6-L2). They link the run to exact code and data.
- **Metrics:** the honest estimate (lesson B11-L3) and its spread, named consistently across runs: `cv_auc_mean`, `cv_auc_std`.
- **The model:** `log_model` saves the fitted object with a `requirements.txt`, `conda.yaml` and `python_env.yaml` pinned to the versions used. For scikit-learn models MLflow 3.16 saves in the skops format.

Use the `with mlflow.start_run():` block. If training raises, the run ends with status `FAILED` instead of staying open and collecting the next run's metrics.

Compare runs with `mlflow.search_runs(experiment_names=["churn"])`, which returns a pandas DataFrame with `params.*`, `metrics.*` and `tags.*` columns. mlflow does not run in this app's browser runtime. The exercises check your code's structure.

Why it matters

A model with AUC 0.83 went to production from a notebook. Three months later it needed retraining. Nobody could reproduce 0.83: the notebook had been edited since, the seed had not been recorded, and the training extract had been refreshed. The team spent two weeks rebuilding a model they had already built once. One tracked run with parameters, commit and data hash would have made the retrain a single command.

Worked example: weak code

```
1 model = LogisticRegression(C=c).fit(X_train, y_train)
2 print(c, roc_auc_score(y_val, model.predict_proba(X_val)[: , 1]))
3 pickle.dump(model, open("model_final_v2.pkl", "wb"))
```

1. **Line 2.** Results live in cell output, lost on the next run.
2. **Line 3.** A file name as version control, with no code, data or environment attached.

Worked example: strong code

```
1 import mlflow
2 import numpy as np
3 from sklearn.linear_model import LogisticRegression
4
5
6 def tracked_fit(X_train: np.ndarray, y_train: np.ndarray, c: float, seed:
    int, commit: str, data_hash: str, cv_auc: float) -> str:
7     """Fit and log one run. Returns the run ID."""
8     with mlflow.start_run(run_name=f"lr-c{c}") as run:
9         mlflow.log_params({"C": c, "seed": seed})
10        mlflow.set_tags({"git_commit": commit, "input_sha256": data_hash})
11        model = LogisticRegression(C=c, random_state=seed,
12            max_iter=1000).fit(X_train, y_train)
13        mlflow.log_metric("cv_auc", cv_auc)
14        mlflow.sklearn.log_model(model, name="model",
15            input_example=X_train[:5])
16    return run.info.run_id
```

1. **Line 8.** The run ends cleanly even if fitting raises.
2. **Line 10.** Code and data identity, attached to the run.
3. **Line 13.** The model with pinned requirements and an input example.

One run per fit, inside a with block, with parameters, commit, data hash, metrics and the model.

Recap card

- **What to log per run** Params (incl. seed), tags (git_commit, input_sha256), metrics, the model.
- **MLflow parameters** Stored as strings and immutable within a run.
- **Why with mlflow.start_run()** The run ends with FAILED status if training raises, instead of staying open.
- **log_model in MLflow 3** mlflow.sklearn.log_model(model, name="model"), with pinned requirements saved alongside.

Practice

Exercise 1. Multiple choice (level 1)

You log `mlflow.log_param("C", 1.0)` and later read the run back. What type is the value?

- a) float
- b) str
- c) int
- d) Decimal

Why?

- r1) MLflow stores parameters as strings, so compare and convert carefully.
- r2) Python cannot store floats in a database.

Exercise 2. Predict the output (level 2)

What happens? (MLflow 3.16)

```
1 import mlflow
2
3 with mlflow.start_run():
4     mlflow.log_param("C", 1.0)
5     mlflow.log_param("C", 2.0)
```

Exercise 3. Fill the gap (level 2)

Log the Git commit as a tag and the model under the name `model`.

```
1 mlflow.__(1)("git_commit", commit)
2 mlflow.sklearn.log_model(model, __(2)="model")
```


Exercise 4. Write the function (level 3)

Write `tracked_fit(X_train, y_train, c, seed, commit, data_hash, cv_auc) -> str`. Inside with `mlflow.start_run(...)` as `run`: log params `C` and `seed`, set tags `git_commit` and `input_sha256`, fit `LogisticRegression(C=c, random_state=seed, max_iter=1000)`, log the metric `cv_auc`, log the model with `mlflow.sklearn.log_model(model, name="model")`, and return `run.info.run_id`. `mlflow` does not run in the browser. The checker reads your code.

Starter code:

```
1 import mlflow
2 import numpy as np
3 from sklearn.linear_model import LogisticRegression
4
5
6 def tracked_fit(X_train, y_train, c, seed, commit, data_hash, cv_auc):
7     mlflow.start_run()
8     model = LogisticRegression(C=c).fit(X_train, y_train)
9     print(cv_auc)
10    return model
```

Exercise 5. Architecture design (level 4)

Decide how each item is recorded in MLflow.

Options: Parameter, Tag, Metric, Artifact.

1. Random seed
2. Git commit
3. Cross-validated AUC
4. Calibration plot

Sources

- MLflow 3.16: default local tracking URI is `sqlite:///<cwd>/mlflow.db`; params come back as strings; re-logging a param with a new value raises `MLflowException`; `log_model` accepts `name=`; sklearn models are saved as `model.skops` with `requirements.txt`, `conda.yaml` and `python_env.yaml`. (verified: Run with MLflow 3.16.1 in a scratch environment for this repository.)

B14-L2 Registry versions and aliases, not stages

Concept

The **model registry** turns logged models into named, numbered versions. Each `log_model(..., registered_model_name="churn_model")` or `mlflow.register_model(...)` call creates a new **version**: 1, 2, 3. A version never changes. It points at one run, one set of parameters and one model file.

Aliases are movable labels on versions. MLflow 3.16:

```
1 from mlflow import MlflowClient
2
3 client = MlflowClient()
4 client.set_registered_model_alias("churn_model", "champion", "3")
5 model = mlflow.pyfunc.load_model("models:/churn_model@champion")
```

The scoring job always loads `@champion`. Promotion moves the alias to a new version. Rollback moves it back. The job's code never changes, and every prediction can be traced to a version and its run.

Stages (Staging, Production, Archived) are the older mechanism.

`transition_model_version_stage` has been deprecated since MLflow 2.9, and MLflow 3.16 warns that stages will be removed. Models in Unity Catalog on Databricks use aliases only, with three-level names such as `ml.churn.churn_model` after `mlflow.set_registry_uri("databricks-uc")`.

Promote on evidence, not by hand:

1. Train and register a candidate version.
2. Evaluate it on the same held-out data and metric as the current champion.
3. Move `champion` only if the candidate beats it by a margin agreed in advance.
4. Keep the old version reachable, for example with a `previous_champion` alias, so rollback is one call.

Record the decision: which versions were compared, the scores and who approved. On Databricks, Unity Catalog grants (lesson A4-L2) control who can move an alias.

Why it matters

A team kept production models as files named `model_prod.pkl`, overwritten on each release. A new version cut conversion forecasts by 12% overnight. Rolling back meant finding last month's file on someone's laptop. It took a

day. With versions and an alias, rollback is

`set_registered_model_alias("churn_model", "champion", "4")`, and the scoring job picks it up on its next run.

Worked example: weak code

```
❶ client.transition_model_version_stage("churn_model", new_version,
    "Production")
2 model = mlflow.pyfunc.load_model("models:/churn_model/Production")
```

1. **Line 1.** Deprecated since MLflow 2.9, not available in Unity Catalog, and promoted without a comparison.

Worked example: strong code

```
1 from __future__ import annotations
2
3 from typing import TYPE_CHECKING
4
5 if TYPE_CHECKING:
6     from mlflow import MlflowClient
7
8
9 def promote_if_better(client: MlflowClient, name: str, candidate: str,
    metric: str, min_gain: float) -> bool:
10     """Move the champion alias to candidate if it beats the champion by
        min_gain."""
    ❶ champ = client.get_model_version_by_alias(name, "champion")
12     champ_score = client.get_run(champ.run_id).data.metrics[metric]
13     cand_score = client.get_run(client.get_model_version(name,
        candidate).run_id).data.metrics[metric]
    ❷ if cand_score < champ_score + min_gain:
15         return False
    ❸ client.set_registered_model_alias(name, "previous_champion",
        champ.version)
17     client.set_registered_model_alias(name, "champion", candidate)
18     return True
```

1. **Line 11.** The current champion, resolved through its alias.
2. **Line 14.** An agreed margin, so noise alone does not trigger a release.
3. **Line 16.** Rollback stays one alias move away.

Versions are fixed. Aliases move. Promote by comparison on the same data and metric, and keep the previous champion addressable.

Recap card

- **Version vs alias** A version is fixed. An alias is a movable label on a version.

- **Load the current production model**

`mlflow.pyfunc.load_model("models:/name@champion")`.

- **Model stages** Deprecated since MLflow 2.9, not supported in Unity Catalog. Use aliases.
- **Safe promotion** Compare on the same data and metric, require a margin, keep `previous_champion` for rollback.

Practice

Exercise 1. Multiple choice (level 1)

What does the scoring job load in an alias-based setup?

- a) `models:/churn_model/7`
- b) `models:/churn_model@champion`
- c) `models:/churn_model/Production`
- d) The latest file in a shared folder

Why?

- r1) The alias moves on promotion and rollback, so the job's code never changes.
- r2) Aliases load faster than versions.

Exercise 2. Predict the output (level 2)

A stub registry in plain Python behaves like MLflow's aliases. What does this print?

```
1 aliases = {"champion": "3"}
2 aliases["previous_champion"] = aliases["champion"]
3 aliases["champion"] = "5"
4 aliases["champion"] = aliases["previous_champion"]
5 print(aliases)
```

Exercise 3. Fill the gap (level 2)

Point the champion alias at version 4, then load whatever it points to.

```
1 client.__(1)("churn_model", "champion", "4")
2 model = mlflow.pyfunc.load_model("models:/churn_model__(2)champion")
```

Exercise 4. Write the function (level 3)

Write `promote_if_better(client, name, candidate, metric, min_gain) -> bool`. The model already has a `champion` alias.

- Read the champion with `client.get_model_version_by_alias(name, "champion")`, and its score from `client.get_run(champ.run_id).data.metrics[metric]`.
- Read the candidate's score through `client.get_model_version(name, candidate)`.
- If the candidate's score is at least the champion's plus `min_gain`, set `previous_champion` to the champion's version, then `champion` to `candidate`, and return `True`. Otherwise change nothing and return `False`.

The checker passes a stub client with the same methods as `MLflowClient`.

Starter code:

```
1 def promote_if_better(client, name, candidate, metric, min_gain):
2     client.set_registered_model_alias(name, "champion", candidate)
3     return True
```

Exercise 5. Architecture design (level 4)

Set up model releases on Databricks with Unity Catalog.

Options: `ml_prod.churn.churn_model`, `churn_model_final`.

1. Model name
2. What the scoring job loads
3. Who can move the champion alias

Sources

- `set_registered_model_alias`, `get_model_version_by_alias` and `models:/name@alias` work as described; `transition_model_version_stage` warns it is deprecated since 2.9.0. (verified: Run with MLflow 3.16.1 in a scratch environment for this repository.)
- Models in Unity Catalog use three-level names with `mlflow.set_registry_uri("databricks-uc")` and do not support stages. (search-only: Databricks documentation, Manage model lifecycle in Unity Catalog (docs site not reachable from this environment).)

B14-L3 Reproduce a run from its record

Concept

A tracked run is only useful if you can rebuild its result from the record. Test that before you need it, not during an incident.

From a run ID, MLflow 3.16 gives you everything the run logged:

```
1 run = mlflow.get_run(run_id)
2 run.data.params                # {'C': '1.0', 'seed': '7', 'features': 'v3'}
   strings
3 run.data.tags["git_commit"], run.data.tags["input_sha256"]
4 run.data.metrics["cv_auc"]
5 reqs = mlflow.pyfunc.get_model_dependencies(f"runs://{run_id}/model") #
   path to requirements.txt
```

The procedure:

1. **Code.** Check out the logged commit: `git switch --detach <commit>`.
2. **Environment.** Create a fresh environment from the lock file at that commit (lesson B6-L3), or from the model's `requirements.txt`.
3. **Data.** Load the input and compare its SHA-256 with `input_sha256`. If it differs, stop: you are reproducing on different data.
4. **Settings.** Convert the parameters back to their types. `'1.0'` must become `1.0` and `'7'` must become `7`. Passing `'7'` as a seed fails, or worse, works differently.
5. **Rerun and compare.** Recompute the metric and compare with the logged value within a small tolerance. Identical code, data, settings and versions on the same hardware usually reproduce exactly. Different CPUs or thread counts can change the last digits of floating-point results.

If the numbers differ by more than the tolerance, the record is incomplete. Something that changes the result was not logged: a feature list read from a file, a default that changed between library versions, an unseeded shuffle. Add it to the run's parameters and try again.

Schedule a reproduction check: once a quarter, pick the champion model's run and rebuild it from scratch in CI.

Why it matters

An auditor asked a lender to show how a credit model's reported AUC of 0.78 had been produced. The run existed in MLflow with parameters and a metric. The input table had since been overwritten, and no data hash had been logged. The team could not show the model had been validated on the data

they claimed. The model was suspended for four weeks while it was rebuilt with full records.

Worked example: weak code

```
1 params = mlflow.get_run(run_id).data.params
2 ① model = LogisticRegression(C=params["C"], random_state=params["seed"])
```

1. **Line 2.** Parameters are strings. '1.0' and '7' are passed through unconverted.

Worked example: strong code

```
1 PARAM_TYPES: dict[str, type] = {"C": float, "seed": int, "features": str}
2
3
4 def parse_params(params: dict[str, str], types: dict[str, type]) ->
  dict[str, object]:
5     """Convert logged string parameters back to their types. Fails on
      missing or bad values."""
6     missing = sorted(set(types) - set(params))
7     if missing:
8         ② raise KeyError(f"run is missing parameters: {missing}")
9     return {key: cast(params[key]) for key, cast in types.items()}
10
11
12 def reproduced(logged: float, rerun: float, tol: float = 1e-9) -> bool:
13     """True when the rerun metric matches the logged one within tol."""
14     ③ return abs(logged - rerun) <= tol
```

1. **Line 1.** The expected parameters and their types, in one place.
2. **Line 8.** A run without a needed parameter cannot be reproduced. Say so.
3. **Line 14.** An explicit tolerance, recorded with the check.

Code from the commit, environment from the lock, data checked by hash, settings converted back to types, metric compared within a tolerance.

Recap card

- **Reproduce a run** Commit, environment from lock, data verified by hash, typed params, rerun, compare within tolerance.
- **Parameters from mlflow.get_run** Strings. Convert with the expected types before use.
- **Model dependencies**
mlflow.pyfunc.get_model_dependencies(model_uri) returns the path to requirements.txt.
- **Rerun differs beyond tolerance** Something that changes results was not logged. Find it and log it.

Practice

Exercise 1. Multiple choice (level 1)

You are reproducing a run. The input file's SHA-256 differs from the logged `input_sha256`. What next?

- a) Continue. The file name is the same.
- b) Stop. You would be reproducing on different data. Find the original input.
- c) Update the logged tag to the new hash.
- d) Rerun twice and average.

Why?

- r1) Any changed byte changes the hash, so the data is not the data the run used.
- r2) Hashes are unreliable.

Exercise 2. Predict the output (level 2)

Logged parameters come back as strings. What does this print?

```
1 params = {"C": "1.0", "seed": "7"}
2 print(params["seed"] + "1", int(params["seed"]) + 1)
```

Exercise 3. Put in order (level 2)

Order the steps to reproduce a logged run.

- Load the input and verify its SHA-256 against `input_sha256`
- Check out the logged git commit
- Convert the logged parameters back to their types
- Build a fresh environment from the lock file at that commit
- Rerun and compare the metric within a tolerance

Exercise 4. Write the function (level 3)

Write `parse_params(params: dict[str, str], types: dict[str, type]) -> dict[str, object]`. For each key in `types`, convert `params[key]` with the type, for example `float("1.0")`. Raise `KeyError` naming the missing keys (sorted) when any key in `types` is absent from `params`. Ignore extra keys in `params`. Let a bad value raise the type's own `ValueError`.

Starter code:

```
1 def parse_params(params, types):
2     return dict(params)
```


Exercise 5. Architecture design (level 4)

A quarterly check rebuilds the champion model's run in CI. The rerun AUC is 0.7812 against a logged 0.7851. Choose the response.

Options: Not reproduced: the gap is far above the tolerance of $1e-6$, Close enough.

1. Verdict
2. Next step
3. Lasting fix

Sources

- `mlflow.pyfunc.get_model_dependencies` returns a path to `requirements.txt`; `get_run(...).data.params` values are strings. (verified: Run with MLflow 3.16.1 in a scratch environment for this repository.)

B15 Embeddings and clustering

Roadmap: Months 7 to 9: NLP with deep learning. SentenceTransformers, HDBSCAN and cluster validation.

B15-L1 Sentence embeddings for open-ended answers

Concept

Open-ended survey answers ("What could we improve?") arrive as thousands of short texts. A **sentence embedding model** turns each text into a fixed-length vector, so that texts with similar meaning get nearby vectors. With the sentence-transformers library:

```
1 from sentence_transformers import SentenceTransformer
2
3 model = SentenceTransformer("sentence-transformers/all-MiniLM-L6-v2")
4 vectors = model.encode(answers, batch_size=64, normalize_embeddings=True)
```

`all-MiniLM-L6-v2` returns 384 numbers per text. With

`normalize_embeddings=True` every vector has length 1, so the **cosine similarity** of two texts is their dot product: 1 for the same direction, 0 for unrelated.

Rules that keep embeddings usable:

- **One model per analysis.** Vectors from different models, or different versions of one model, live in different spaces. Never compare or mix them. Record the model name and version with the vectors.
- **Cache the vectors.** Encoding is the slow step. Store vectors with the text ID, the model name and a hash of the text, and re-encode only new or changed answers.
- **Clean lightly.** Strip whitespace and drop empty answers. Do not stem or remove stop words: the model was trained on natural text.
- **Check the language.** Many small models are English-only. Use a multilingual model for multi-country surveys, and test it on your own answers.
- **Handle zero vectors.** An empty string can produce a vector you should not compare. Filter before encoding, and guard any division by a vector's length.

Embeddings reflect the text the model was trained on, including its biases. Read samples of what the model calls "similar" before you trust any grouping built on it.

sentence-transformers needs PyTorch, which this app cannot run in the browser. This course's build did not run sentence-transformers, so the API details above come from the library's documentation and are marked as not run. The exercises use NumPy on vectors.

Why it matters

An analyst embedded wave 1 answers with one model and, a year later, wave 2 answers with a newer model under the same short name. Similarity between the waves came out low across the board, and the report said themes had "shifted dramatically". The texts had not changed much. The two sets of vectors came from different spaces.

Worked example: weak code

```
1 def most_similar(query, matrix):
2     scores = matrix @ query
3     return int(scores.argmax())
```

1. **Line 2.** A dot product equals cosine similarity only when all vectors have length 1.

Worked example: strong code

```
1 import numpy as np
2
3
4 def top_k_similar(query: np.ndarray, matrix: np.ndarray, k: int) ->
    list[int]:
5     """Row indices of the k rows most cosine-similar to query, most
        similar first."""
6     q_norm = np.linalg.norm(query)
7     if q_norm == 0:
8         raise ValueError("query vector has zero length")
9     norms = np.linalg.norm(matrix, axis=1)
10    sims = np.full(len(matrix), -np.inf)
11    ok = norms > 0
12    sims[ok] = (matrix[ok] @ query) / (norms[ok] * q_norm)
13    return [int(i) for i in np.argsort(-sims, kind="stable")[:k]]
```

1. **Line 10.** Zero-length rows rank last instead of causing a division by zero.
2. **Line 12.** Cosine similarity that does not assume normalised input.

One model, recorded and cached. Compare with cosine similarity, and guard against zero vectors.

Recap card

- **Encode with sentence-transformers**
SentenceTransformer(name).encode(texts, normalize_embeddings=True).
Not run in this course's build.
- **Cosine similarity for unit vectors** The dot product. Otherwise divide by both lengths.
- **Mixing embedding models** Never. Each model defines its own space.
- **Cache key for a vector** Text ID, text hash, model name and version.

Practice

Exercise 1. Multiple choice (level 1)

You have vectors for 2025 answers from model A, and for 2026 answers from model B. Can you compare them?

- a) Yes, if both have 384 dimensions.
- b) No. Re-encode both years with the same model.
- c) Yes, after normalising both.
- d) Yes, by averaging them.

Why?

- r1) Each model defines its own space. The same position means different things in each.
- r2) Newer models are always better.

Exercise 2. Predict the output (level 2)

What does this print? (NumPy 2.4)

```
1 import numpy as np
2
3 a = np.array([3.0, 4.0])
4 b = np.array([1.0, 0.0])
5 print(float(a @ b), round(float(a @ b / (np.linalg.norm(a) *
    np.linalg.norm(b))), 2))
```

Exercise 3. Fill the gap (level 2)

Normalise each row of a matrix to length 1.

```
1 norms = np.linalg.norm(matrix, axis=____(1), keepdims=True)
2 unit = matrix / ____ (2)
```

Exercise 4. Write the function (level 3)

Write `top_k_similar(query, matrix, k) -> list[int]`. Return the row indices of the `k` rows with the highest cosine similarity to `query`, highest first, as Python ints. Rows with zero length rank last. Raise `ValueError` if `query` has zero length. Do not assume the inputs are normalised.

Starter code:

```
1 import numpy as np
2
3
4 def top_k_similar(query, matrix, k):
5     return list(np.argsort(matrix @ query)[:k])
```

Exercise 5. Architecture design (level 4)

Plan embeddings for a quarterly tracker's open-ended question, in five languages.

Options: One multilingual model, tested on a sample from each language, An English-only model for all languages, A different model per language.

1. Model
2. What is stored with each vector
3. Next quarter

Sources

- `SentenceTransformer("sentence-transformers/all-MiniLM-L6-v2").encode(texts, batch_size=..., normalize_embeddings=True)` returns 384-dimensional vectors. (unverified: sentence-transformers documentation and model card. Not run in this course's build.)

B15-L2 HDBSCAN, clusters with a noise label

Concept

k-means needs you to choose the number of clusters, and it puts every point in one. Open-ended answers rarely fit that: some themes are large, some small, and many answers belong to no theme at all. **HDBSCAN** finds dense groups of points, lets clusters differ in size and shape, and labels points in sparse regions as **noise**, `-1`.

scikit-learn has included it since version 1.3. With scikit-learn 1.8:

```
1 from sklearn.cluster import HDBSCAN
2
```

```
3 labels = HDBSCAN(min_cluster_size=15, copy=True).fit_predict(vectors)
```

- `min_cluster_size` (default 5) is the main setting: the smallest group you would call a theme. Raise it for fewer, larger themes.
- `min_samples` controls how conservative clustering is. Higher values label more points as noise.
- Pass `copy=True`. scikit-learn 1.8 warns that the default will change from `False` to `True` in 1.10, and `False` lets the estimator modify your input array.
- The default metric is Euclidean. For normalised embeddings (lesson B15-L1), Euclidean distance ranks pairs in the same order as cosine similarity.

Read the result before you use it:

- **Noise share.** On 600 points in four blobs plus 60 random points, this course's build labelled 12% as noise. For short survey answers a large noise share is normal, and it is information: those answers do not share a theme with enough others.
- **Number and size of clusters.** One giant cluster and many tiny ones usually means the settings or the embedding need work.
- **Samples.** Read 10 answers from each cluster before you name it.

High-dimensional embeddings are often reduced first, for example with UMAP or PCA, to make density easier to find. Record every setting and the random seed of any reduction step, because the clusters depend on them.

Why it matters

A team clustered 8,000 complaints with k-means and $k = 12$, because the dashboard had space for 12 tiles. Every complaint was forced into a theme. A cluster labelled "billing" contained 900 complaints about billing and 400 unrelated ones that were nearest to it. Its share of complaints was overstated by 40%. HDBSCAN on the same vectors marked most of the 400 as noise.

Worked example: weak code

```
❶ labels = KMeans(n_clusters=12).fit_predict(vectors)
❷ shares = pd.Series(labels).value_counts(normalize=True)
```

1. **Line 1.** k chosen by the dashboard, every answer forced into a cluster, no seed.
2. **Line 2.** Shares inflated by answers that belong nowhere.

Worked example: strong code

```
1 import numpy as np
```

```

2
3
4 def summarise_labels(labels: np.ndarray) -> dict[str, float | int | None]:
5     """Cluster count, noise share and largest cluster share (of clustered
6     points)."""
7     n = len(labels)
8     clustered = labels[labels != -1]
9     if n == 0:
10         return {"n_clusters": 0, "noise_share": None, "largest_share":
11         None}
12     sizes = np.unique(clustered, return_counts=True)[1]
13     return {
14         "n_clusters": int(len(sizes)),
15         "noise_share": float((labels == -1).mean()),
16         "largest_share": float(sizes.max() / len(clustered)) if
17         len(clustered) > 0 else None,
18     }

```

1. **Line 7.** Noise is reported separately, never counted as a cluster.

2. **Line 14.** Guarded: all noise means no clusters to compare.

Let density decide the clusters, keep noise as its own answer, and summarise counts, noise share and balance before naming themes.

Recap card

- **HDBSCAN noise label** -1: points in sparse regions. Report the share, never treat it as a cluster.
- **HDBSCAN main setting** `min_cluster_size` (default 5). `min_samples` makes it more conservative.
- **HDBSCAN in scikit-learn 1.8** `sklearn.cluster.HDBSCAN`. Pass `copy=True`: the default changes in 1.10.
- **Read a clustering** Noise share, number and size of clusters, then samples from each.

Practice

Exercise 1. Multiple choice (level 1)

HDBSCAN labels 30% of survey answers -1. What does that mean?

- a) The algorithm failed on those answers.
- b) Those answers are in sparse regions and share no dense theme with enough others.
- c) They form a 31st cluster.
- d) They are duplicates.

Why?

- r1) -1 is the noise label for points that no dense cluster claims.
- r2) HDBSCAN always drops 30%.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 import numpy as np
2
3 labels = np.array([0, 0, 1, -1, 1, 1, -1, 0, 0, 0])
4 clustered = labels[labels != -1]
5 print(len(np.unique(clustered)), float((labels == -1).mean()),
        float(np.bincount(clustered).max() / len(clustered)))
```

Exercise 3. Fill the gap (level 2)

Cluster embeddings with HDBSCAN, smallest theme 20 answers, without modifying the input array.

```
1 labels = HDBSCAN(____(1)=20, copy=____(2)).fit_predict(vectors)
```

Exercise 4. Write the function (level 3)

Write `summarise_labels(labels) -> dict[str, float | int | None]` returning:

- `n_clusters`: distinct labels other than -1.
- `noise_share`: share of all points labelled -1, or `None` for an empty array.
- `largest_share`: size of the largest cluster over the number of clustered (non-noise) points, or `None` if there are none.

Starter code:

```
1 import numpy as np
2
3
4 def summarise_labels(labels):
5     return {"n_clusters": len(np.unique(labels))}
```

Exercise 5. Architecture design (level 4)

HDBSCAN on 5,000 complaint embeddings gives one cluster with 70% of clustered points, 40 tiny clusters, and 45% noise. Choose the next steps.

Options: Read samples from the big cluster, a few tiny ones, and the noise, Publish the 41 themes.

1. First step
2. Setting to try
3. Record

Sources

- `sklearn.cluster.HDBSCAN` defaults to `min_cluster_size=5` and `metric="euclidean"`; `scikit-learn` 1.8 warns that `copy` will default to `True` in 1.10. (verified: Signature and warning observed on `scikit-learn` 1.8.0 and 1.9.1 in this repository.)
- Four blobs of 600 points plus 60 uniform points:
`HDBSCAN(min_cluster_size=15)` found 4 clusters and 12% noise. (verified: Run with `scikit-learn` 1.8.0 in this repository (`make_blobs random_state=0`).)

B15-L3 Validate clusters before you name them

Concept

Any clustering algorithm returns clusters, even on random data. Before a cluster becomes a "theme" in a report, check three things.

1. Separation. The **silhouette score** compares each point's distance to its own cluster with its distance to the nearest other cluster. It runs from -1 to 1, higher is better. `scikit-learn` 1.8: `silhouette_score(X, labels)`. Compute it **without noise points**. The function treats -1 as a cluster like any other. In this course's build, four blobs with `HDBSCAN` noise scored 0.70 with noise removed and 0.555 with noise counted as a cluster.

2. Stability. Rerun the whole pipeline on resampled data, or with different seeds, and compare the labels on the points both runs share. The **adjusted Rand index** (`adjusted_rand_score`) is 1 for identical groupings, around 0 for chance, and ignores label names: `[0, 0, 1, 1]` against `[1, 1, 0, 0]` scores 1.0. Themes that appear in only some runs are artefacts of the settings.

3. Meaning. Read 10 to 20 answers per cluster. Have two people name the clusters independently, then compare. A cluster nobody can name consistently is not a theme yet.

Report with each theme:

- its size and share, with noise reported separately,
- the silhouette and stability figures for the clustering as a whole,
- three verbatim example answers, chosen at random, not the most typical ones.

No single number proves clusters are real. Together, separation, stability and human reading make a defensible case. Keep the validation results in the run record with the settings (lesson B6-L2).

Why it matters

An insight team presented seven "customer mindsets" from a clustering of attitude statements. A year later a new team reran the same method on a fresh sample. Three of the seven mindsets reappeared. The other four had been unstable from the start: a stability check, rerunning on bootstrap samples, would have shown them splitting and merging from run to run. Two product launches had been planned around mindsets that did not exist.

Worked example: weak code

```
❶ score = silhouette_score(vectors, labels)
❷ print("7 clear segments, silhouette", round(score, 2))
```

1. **Line 1.** Noise (-1) counted as a cluster, and one run taken as the truth.
2. **Line 2.** A single number presented as proof.

Worked example: strong code

```
1 import numpy as np
2 from sklearn.metrics import adjusted_rand_score, silhouette_score
3
4
5 def silhouette_without_noise(X: np.ndarray, labels: np.ndarray) -> float
6     | None:
7     """Silhouette over clustered points only. None if fewer than two
8     clusters remain."""
9
10    ❶ keep = labels != -1
11    kept = labels[keep]
12    n_clusters = len(np.unique(kept))
13
14    ❷ if n_clusters < 2 or n_clusters == len(kept):
15        return None
16    return float(silhouette_score(X[keep], kept))
17
18 def stability(labels_a: np.ndarray, labels_b: np.ndarray) -> float:
19     """Agreement of two labelings of the same points, ignoring label
20     names."""
21
22    ❸ return float(adjusted_rand_score(labels_a, labels_b))
```

1. **Line 7.** Noise is not a cluster and does not enter the silhouette.
2. **Line 10.** silhouette_score needs between 2 and n - 1 clusters. Return None instead of raising.
3. **Line 17.** Compare runs on shared points. Label names do not matter.

Check separation without noise, stability across reruns, and meaning by reading, before any cluster gets a name.

Recap card

- **Three checks before naming clusters** Separation (silhouette), stability (ARI across reruns), meaning (read and name independently).
- **Silhouette and HDBSCAN noise** Exclude -1 first. silhouette_score treats it as a cluster (0.70 vs 0.555 in this lesson's example).
- **Adjusted Rand index** 1 for identical groupings, about 0 for chance, can be negative, ignores label names.
- **Examples on a theme slide** Random verbatims, not hand-picked ones.

Practice

Exercise 1. Multiple choice (level 1)

A clustering has silhouette 0.62. Is that enough to report its clusters as themes?

- a) Yes. Above 0.5 is proof.
- b) No. Also check stability across reruns and read samples from each cluster.
- c) No. Silhouette must be exactly 1.
- d) Yes, if the number of clusters matches the dashboard.

Why?

- r1) Separation in one run says nothing about whether clusters recur or mean anything.
- r2) Silhouette is never informative.

Exercise 2. Predict the output (level 2)

What does this print? (scikit-learn 1.8)

```
1 from sklearn.metrics import adjusted_rand_score
2
3 print(adjusted_rand_score([0, 0, 1, 1], [1, 1, 0, 0]),
      adjusted_rand_score([0, 0, 1, 1], [0, 1, 0, 1]))
```

Exercise 3. Fill the gap (level 2)

Compute the silhouette over clustered points only.

```
1 keep = labels ____ (1) -1
2 score = silhouette_score(X[keep], labels[____ (2)])
```

Exercise 4. Write the function (level 3)

Write `silhouette_without_noise(X, labels) -> float | None`. Drop points labelled -1. If fewer than 2 clusters remain, or as many clusters as remaining points, return `None`. Otherwise return `silhouette_score` on the remaining points as a float.

Starter code:

```
1 from sklearn.metrics import silhouette_score
2
3
4 def silhouette_without_noise(X, labels):
5     return silhouette_score(X, labels)
```

Exercise 5. Architecture design (level 4)

You will present six themes from open-ended answers. Choose what goes on the methods slide.

Options: Silhouette without noise, with the noise share stated, Silhouette with noise counted as a cluster.

1. Separation
2. Stability
3. Examples per theme

Sources

- Silhouette 0.70 excluding noise and 0.555 including it for HDBSCAN on four blobs plus uniform noise; `adjusted_rand_score([0,0,1,1],[1,1,0,0]) = 1.0` and `([0,0,1,1],[0,1,0,1]) = -0.5`. (verified: Run with scikit-learn 1.8.0 and 1.9.1 in this repository.)

B16 PyTorch hygiene

Roadmap: Months 1 to 3: PyTorch foundations. Device handling, training loops and reproducibility.

B16-L1 Choose the device once, move model and data to it

Concept

PyTorch runs on the CPU, on NVIDIA GPUs (`cuda`) and on Apple silicon (`mps`). Every tensor and every model parameter lives on one **device**, and an operation needs all its inputs on the same device. Examples use PyTorch 2.14.

Choose the device once, at the edge, and pass it in:

```
1 def pick_device(prefer_gpu: bool = True) -> torch.device:
2     if prefer_gpu and torch.cuda.is_available():
3         return torch.device("cuda")
4     if prefer_gpu and torch.backends.mps.is_available():
5         return torch.device("mps")
6     return torch.device("cpu")
```

Then move the model once, and each batch as you use it.

The trap: `.to()` **behaves differently on modules and tensors**.

- `model.to(device)` moves the parameters **in place** and returns the same model.
- `tensor.to(device)` returns a **new** tensor. The original stays where it was.
`batch.to(device)` on its own line does nothing useful. Write `batch = batch.to(device)`.

Mixing devices raises `RuntimeError`, with a message such as "Tensor on device cpu is not on the expected device". This course's build reproduced it with PyTorch's `meta` device, because no GPU was available.

Other habits:

- Never hardcode `"cuda"`. The same code must run on a laptop CPU, in CI and on a GPU cluster.
- Move tensors to the CPU before converting to NumPy:
`preds.cpu().numpy()`. NumPy cannot read GPU memory.
- Log the device with every run (lesson B14-L1). Results can differ slightly between devices.

PyTorch does not run in this app's browser runtime. The exercises check your code's structure, and the device logic uses plain Python.

Why it matters

A training script written on a GPU workstation had `.cuda()` calls on 14 lines. It had to run in CI on CPU machines for a smoke test. Every run failed at import of the first batch. The fix took an afternoon of search and replace, and one `.cuda()` hidden in a helper was missed and broke the nightly job a week later. One `device` chosen at the edge and passed in would have needed no changes.

Worked example: weak code

```
❶ model = Net().cuda()
  2 for x, y in loader:
❷     x.to("cuda")
  4     loss = loss_fn(model(x), y.cuda())
```

1. **Line 1.** Hardcoded GPU: fails on any machine without CUDA.
2. **Line 3.** Returns a new tensor that is thrown away. `x` stays on the CPU.

Worked example: strong code

```
1 from collections.abc import Iterable
2
3 import torch
4
5
❶ def pick_device(prefer_gpu: bool = True) -> torch.device:
7     """CUDA, then Apple MPS, then CPU."""
8     if prefer_gpu and torch.cuda.is_available():
9         return torch.device("cuda")
10    if prefer_gpu and torch.backends.mps.is_available():
11        return torch.device("mps")
12    return torch.device("cpu")
13
14
15 def train_epoch(
16     model: torch.nn.Module,
17     loader: Iterable[tuple[torch.Tensor, torch.Tensor]],
18     loss_fn: torch.nn.Module,
19     optimizer: torch.optim.Optimizer,
20     device: torch.device,
21 ) -> float:
22     """One pass over loader. Returns the mean batch loss."""
23     model.train()
24     total, batches = 0.0, 0
25     for x, y in loader:
❷         x, y = x.to(device), y.to(device)
```

```

27     optimizer.zero_grad()
28     loss = loss_fn(model(x), y)
29     loss.backward()
30     optimizer.step()
31     total += loss.item()
32     batches += 1
33     return total / batches if batches > 0 else float("nan")

```

1. **Line 6.** One decision, made once, testable on any machine.
2. **Line 26.** Reassigned: the moved tensors are the ones used.
3. **Line 31.** A Python float for logging, not a tensor holding the graph.

Pick the device once and pass it in. Reassign tensors after `.to()`. Never hardcode "cuda".

Recap card

- **Choose a device** Once, at the edge: cuda if available, then mps, then cpu. Pass it in.
- **Tensor.to vs Module.to** Tensor.to returns a new tensor (reassign). Module.to moves in place.
- **GPU tensor to NumPy** `t.detach().cpu().numpy()`.
- **Hardcoded .cuda()** Breaks on CPU machines and CI. Use `.to(device)`.

Practice

Exercise 1. Multiple choice (level 1)

After `x.to(device)` on its own line, where is `x`?

- a) On device.
- b) Where it was before. `.to()` on a tensor returns a new tensor.
- c) Copied to both devices.
- d) Deleted.

Why?

- r1) Tensor.to returns a new tensor. Only `nn.Module.to` changes the object in place.
- r2) `.to()` only works on models.

Exercise 2. Predict the output (level 2)

What does this print? (PyTorch 2.14; the meta device stands in for a GPU)

```

1 import torch
2
3 t = torch.ones(2)
4 t.to("meta")

```

```

5 model = torch.nn.Linear(2, 1)
6 model.to("meta")
7 print(t.device, next(model.parameters()).device)

```

Exercise 3. Fill the gap (level 2)

Move a batch to the chosen device, and bring predictions back for NumPy.

```

1 x = x.____(1)(device)
2 preds = model(x).detach().____(2)().numpy()

```

Exercise 4. Write the function (level 3)

Write `pick_device_name(prefer_gpu: bool, cuda_ok: bool, mps_ok: bool) -> str` that returns "cuda", "mps" or "cpu": CUDA first, then MPS, then CPU, and always "cpu" when `prefer_gpu` is False. It takes availability as arguments so it runs and tests anywhere. The caller passes `torch.cuda.is_available()` and `torch.backends.mps.is_available()`.

Starter code:

```

1 def pick_device_name(prefer_gpu, cuda_ok, mps_ok):
2     return "cuda"

```

Exercise 5. Spot the bug (level 4)

Validation runs without errors on the CPU but raises "Expected all tensors to be on the same device" on the GPU. Tap the line that causes it.

```

1 from collections.abc import Iterable
2
3 import torch
4
5
6 def validate(model: torch.nn.Module, loader: Iterable[tuple[torch.Tensor,
    torch.Tensor]], device: torch.device) -> float:
7     """Accuracy on the validation loader."""
8     model.eval()
9     correct, total = 0, 0
10    with torch.inference_mode():
11        for x, y in loader:
12            x.to(device)
13            preds = model(x).argmax(dim=1)
14            correct += int((preds.cpu() == y).sum())
15            total += len(y)
16    return correct / total if total > 0 else float("nan")

```

Write down the line number of the bug. Then write the fix.

Sources

- In PyTorch 2.14, `Tensor.to` returns a new tensor and leaves the original unchanged; `Module.to` moves parameters in place and returns the same

module; mixing devices raises RuntimeError. (verified: Run with torch 2.14.0 (CPU) in a scratch environment for this repository, using the meta device.)

B16-L2 A training loop that does what you think

Concept

PyTorch leaves the training loop to you. That gives control, and five ways to get it silently wrong. The standard loop, PyTorch 2.14:

```
1 for epoch in range(max_epochs):
2     model.train()
3     for x, y in train_loader:
4         x, y = x.to(device), y.to(device)
5         optimizer.zero_grad()
6         loss = loss_fn(model(x), y)
7         loss.backward()
8         optimizer.step()
9     val_loss = evaluate(model, val_loader, loss_fn, device) #
    model.eval() + torch.inference_mode()
```

The five traps:

1. **Missing** `optimizer.zero_grad()`. Gradients **add up** across `backward()` calls. In this course's build, three backward passes on `2 * w` left a gradient of 6.0 instead of 2.0. Training still runs, with steps that grow each batch.
2. **Missing** `model.eval()` **for validation**. Dropout stays active and batch norm keeps updating its statistics. Two predictions for the same input differed in train mode and matched in eval mode.
3. **Missing** `torch.inference_mode()` (or `torch.no_grad()`) during validation. PyTorch builds a graph it never uses, wasting memory.
4. **Softmax before** `CrossEntropyLoss`. The loss applies log-softmax itself and expects raw **logits**. For logits `[2.0, 0.5]` and class 0, the correct loss is 0.2014. Passing softmax outputs gave 0.4252, a different loss with weaker gradients.
5. **Summing loss tensors for logging**. `total += loss` keeps every batch's graph alive, and memory grows each step. Use `total += loss.item()`.

Early stopping. Track validation loss each epoch. Stop when it has not improved by at least `min_delta` for `patience` epochs, and keep a copy of the best weights (`copy.deepcopy(model.state_dict())`), not the last ones.

Log the loss curve for both training and validation (lesson B14-L1). A validation loss that rises while training loss falls is overfitting. Both flat from the start usually means a bug, such as frozen parameters or a learning rate of zero.

Why it matters

A text classifier trained for 30 epochs showed a smooth, falling training loss and a validation accuracy stuck at 51%. The code applied `softmax` in the model's `forward` and then used `CrossEntropyLoss`. Gradients were small and the model barely learned. Removing one line lifted validation accuracy to 78% in six epochs. The previous two weeks of hyperparameter tuning had been on a broken loss.

Worked example: weak code

```
1 for x, y in loader:
  ❶ out = torch.softmax(model(x), dim=1)
  3 loss = nn.CrossEntropyLoss()(out, y)
  ❷ loss.backward()
  5 optimizer.step()
  ❸ total += loss
```

1. **Line 2.** `CrossEntropyLoss` expects logits. `Softmax` here changes the loss and flattens gradients.
2. **Line 4.** No `zero_grad` before it: gradients accumulate across batches.
3. **Line 6.** Keeps each batch's graph alive. Use `loss.item()`.

Worked example: strong code

```
1 def should_stop(val_losses: list[float], patience: int, min_delta: float
  = 0.0) -> bool:
  2     """True when the last `patience` epochs did not beat the earlier best
    by min_delta."""
  3     if patience < 1:
  4         raise ValueError("patience must be at least 1")
  ❶ if len(val_losses) <= patience:
  6         return False
  7     best_before = min(val_losses[:-patience])
  ❷ return min(val_losses[-patience:]) > best_before - min_delta
```

1. **Line 5.** Too few epochs to judge: keep training.
2. **Line 8.** Stop only when recent epochs failed to improve on the earlier best.

`zero_grad`, `train` and `eval` modes, `inference_mode` for validation, logits into `CrossEntropyLoss`, `loss.item()` for logs, and early stopping on validation loss.

Recap card

- **The training iteration** `to(device)`, `zero_grad`, `forward` and `loss`, `backward`, `step`.
- **Validation mode** `model.eval()` and `torch.inference_mode()`. Back to `model.train()` after.
- **CrossEntropyLoss input** Raw logits. It applies `log-softmax` itself.
- **Early stopping** Stop when validation loss has not improved by `min_delta` for patience epochs. Keep the best weights.

Practice

Exercise 1. Multiple choice (level 1)

A loop calls `loss.backward()` and `optimizer.step()` but never `optimizer.zero_grad()`. What happens?

- a) Nothing. PyTorch clears gradients automatically.
- b) Gradients accumulate across batches, so steps use the sum of all earlier gradients.
- c) `backward()` raises an error on the second batch.
- d) The model stops learning.

Why?

- r1) `backward()` adds into `.grad`. Only `zero_grad()` resets it.
- r2) Optimizers ignore gradients without `zero_grad`.

Exercise 2. Predict the output (level 2)

What does this print? (PyTorch 2.14)

```
1 import torch
2
3 w = torch.tensor([1.0], requires_grad=True)
4 for _ in range(3):
5     loss = (w * 2).sum()
6     loss.backward()
7 print(w.grad.item())
```

Exercise 3. Put in order (level 2)

Order the steps inside one training iteration.

- `optimizer.step()`
- `loss.backward()`
- `loss = loss_fn(model(x), y)`
- `x, y = x.to(device), y.to(device)`

- `optimizer.zero_grad()`

Exercise 4. Write the function (level 3)

Write `should_stop(val_losses: list[float], patience: int, min_delta: float = 0.0) -> bool`. Return `True` when the best loss in the last `patience` epochs is not lower than the best loss before them by more than `min_delta`. Return `False` when there are `patience` or fewer losses. Raise `ValueError` if `patience < 1`.

Starter code:

```
1 def should_stop(val_losses, patience, min_delta=0.0):
2     return val_losses[-1] > val_losses[-2]
```

Exercise 5. Spot the bug (level 4)

Validation accuracy is stuck near chance while training loss falls slowly. Tap the line that breaks the loss.

```
1 import torch
2 from torch import nn
3
4
5 class Classifier(nn.Module):
6     """Two-layer classifier. Train with nn.CrossEntropyLoss."""
7
8     def __init__(self, n_in: int, n_classes: int) -> None:
9         super().__init__()
10        self.net = nn.Sequential(nn.Linear(n_in, 32), nn.ReLU(),
11                                nn.Linear(32, n_classes))
12
13    def forward(self, x: torch.Tensor) -> torch.Tensor:
14        return torch.softmax(self.net(x), dim=1)
```

Write down the line number of the bug. Then write the fix.

Sources

- Three backward passes on $2 \times w$ leave `w.grad = 6.0`; dropout outputs differ in train mode and match in eval mode; `CrossEntropyLoss` on logits `[2.0, 0.5]` with target 0 is 0.2014, and 0.4252 on their softmax. (verified: Run with torch 2.14.0 (CPU) in a scratch environment for this repository.)

B16-L3 Reproducible PyTorch runs and safe checkpoints

Concept

Neural network training draws random numbers everywhere: weight initialisation, data shuffling, dropout, augmentation. Control every source and two runs with the same seed give the same weights. In this course's build (PyTorch 2.14, CPU), a small network trained twice with seed 1 produced identical parameters. Seed 2 produced different ones.

Seed at the entry point, once. Your own code takes generators as parameters (lesson B6-L1). Libraries and PyTorch internals use global generators, so seed those too, at the start of the training script:

```
1 random.seed(seed)
2 np.random.seed(seed)
3 torch.manual_seed(seed)                                # CPU and all CUDA devices
4 loader_gen = torch.Generator().manual_seed(seed)
5 train_loader = DataLoader(ds, batch_size=64, shuffle=True,
                           generator=loader_gen)
```

- Pass a seeded generator to every shuffling `DataLoader`, so batch order is fixed.
- With `num_workers > 0`, the PyTorch documentation recommends a `worker_init_fn` that seeds NumPy and `random` in each worker from `torch.initial_seed()`.
- `torch.use_deterministic_algorithms(True)` makes PyTorch raise an error for operations that have no deterministic implementation, instead of silently varying. On CUDA it also needs the `CUBLAS_WORKSPACE_CONFIG` environment variable, as the documentation describes.

Exact equality holds for the same code, library versions and hardware. Across GPU models or library versions, expect small differences, and compare metrics within a tolerance (lesson B14-L3).

Checkpoints let a long run resume and let you keep the best epoch. Save a dict of state, not the model object:

```
1 torch.save({"model": model.state_dict(), "optimizer":
            optimizer.state_dict(),
2            "epoch": epoch, "torch_rng": torch.get_rng_state()}, path)
```

`torch.load` loads **weights only** by default in PyTorch 2.14. A checkpoint holding a pickled model object fails with `UnpicklingError`. That default protects you: loading an arbitrary pickle can execute code. Save state dicts, rebuild the model class in code, and load the weights into it.

Why it matters

A team fine-tuned a classifier and reported 0.86 F1. A reviewer reran the notebook and got 0.83, then 0.85. The DataLoader shuffled with the global generator, dropout drew from it too, and an augmentation library used its own. Nobody could say which number was "the" result. Fixing the seeds and the loader generator made the run repeatable: 0.84, every time. The earlier 0.86 had been a lucky draw.

Worked example: weak code

```
❶ model = torch.load("best_model.pt")
❷ loader = DataLoader(ds, batch_size=64, shuffle=True)
```

1. **Line 1.** Expects a pickled model object: refused by the weights-only default, and unsafe if forced.
2. **Line 2.** Shuffles with the global generator: batch order changes with any other random call.

Worked example: strong code

```
1 import random
2 from pathlib import Path
3
4 import numpy as np
5 import torch
6
7
8 def seed_everything(seed: int) -> torch.Generator:
9
10     """Seed Python, NumPy and PyTorch global state; return a generator for
11     DataLoaders."""
12     random.seed(seed)
13     np.random.seed(seed)
14     torch.manual_seed(seed)
15
16     ❶ return torch.Generator().manual_seed(seed)
17
18
19 def save_checkpoint(path: Path, model: torch.nn.Module, optimizer:
20     torch.optim.Optimizer, epoch: int) -> None:
21     """Save state dicts and RNG state, never the model object."""
22     torch.save(
23         {
24             ❷ "model": model.state_dict(),
25             "optimizer": optimizer.state_dict(),
26             "epoch": epoch,
27             ❸ "torch_rng": torch.get_rng_state(),
28         },
29         path,
```

1. **Line 13.** A dedicated generator for DataLoader shuffling.
2. **Line 20.** Weights as tensors: loads under the weights-only default.
3. **Line 23.** Resuming restores the random stream as well as the weights.

Seed every generator at the entry point, give DataLoaders their own generator, and checkpoint state dicts, not objects.

Recap card

- **Seed a PyTorch run** `random.seed`, `np.random.seed`, `torch.manual_seed` at the entry point, plus a seeded Generator for each DataLoader.
- **torch.load default in PyTorch 2.14** Weights only. Pickled objects fail with `UnpicklingError`.
- **Checkpoint contents** Model and optimizer state dicts, epoch, RNG state, best metric.
- **Deterministic operations** `torch.use_deterministic_algorithms(True)` raises on operations with no deterministic version.

Practice

Exercise 1. Multiple choice (level 1)

`torch.load("model.pt")` raises `UnpicklingError` in PyTorch 2.14. The file was saved with `torch.save(model, path)`. Best fix?

- a) Always pass `weights_only=False`.
- b) Save `model.state_dict()` instead, rebuild the model class in code, and load the weights into it.
- c) Downgrade PyTorch.
- d) Rename the file to `.pkl`.

Why?

- r1) State dicts hold only tensors and plain values, which load safely under the weights-only default.
- r2) Model objects cannot be saved.

Exercise 2. Predict the output (level 2)

What does this print? (PyTorch 2.14)

```
1 import torch
2
3 g1 = torch.Generator().manual_seed(0)
4 g2 = torch.Generator().manual_seed(0)
```

```

5 a = torch.randperm(5, generator=g1).tolist()
6 torch.rand(100)
7 b = torch.randperm(5, generator=g2).tolist()
8 print(a == b)

```

Exercise 3. Fill the gap (level 2)

Give a shuffling DataLoader its own seeded generator.

```

1 loader_gen = torch.Generator().____(1)(seed)
2 loader = DataLoader(ds, batch_size=64, shuffle=True, ____ (2)=loader_gen)

```

Exercise 4. Write the function (level 3)

Write `seed_everything(seed: int) -> torch.Generator`. Seed Python's `random`, NumPy's global generator and PyTorch with `torch.manual_seed`, then return `torch.Generator().manual_seed(seed)` for DataLoaders.

PyTorch does not run in the browser. The checker reads your code.

Starter code:

```

1 import torch
2
3
4 def seed_everything(seed):
5     torch.manual_seed(seed)

```

Exercise 5. Architecture design (level 4)

Design checkpointing for a 12-hour training job that must resume after interruptions and keep the best epoch.

Options: Model and optimizer state dicts, epoch, RNG state, best validation loss, The whole model object with `torch.save(model)`.

1. Checkpoint content
2. When to save
3. Loading

Sources

- Same-seed training on CPU gave identical parameters; `torch.load` of a pickled model raised `UnpicklingError`; a dict of state dicts, epoch and `torch.get_rng_state()` loaded with the default. (verified: Run with torch 2.14.0 (CPU) in a scratch environment for this repository.)
- PyTorch recommends a `worker_init_fn` seeding NumPy and random from `torch.initial_seed()`, and `CUBLAS_WORKSPACE_CONFIG` for deterministic CUDA. (search-only: PyTorch documentation, Reproducibility notes (not reachable from this environment).)

B17 Testing and performance

Roadmap: Prerequisite for all months. Unit tests, Hypothesis, profiling and memory use.

B17-L1 Unit tests for data code that catch real bugs

Concept

A test is worth writing if it fails when the code is wrong. For data code, most bugs sit at the edges: empty input, missing values, zero denominators, one row, duplicate keys. Test those on purpose.

With pytest 9, a test is a function whose name starts with `test_` and which uses plain `assert`:

```
1 import pytest
2
3
4 def test_share_zero_total_is_none():
5     assert share(3, 0) is None
6
7
8 @pytest.mark.parametrize(("part", "total", "expected"), [(1, 4, 0.25),
9     (0, 5, 0.0), (5, 5, 1.0)])
9 def test_share_values(part, total, expected):
10     assert share(part, total) == expected
```

Habits that make tests catch bugs:

- **Test pure functions** (lesson B2-L1). No files, clocks or databases, so tests are fast and exact.
- **Build tiny, explicit inputs** in the test itself: three rows, written out. A reader sees the input and the expected answer together.
- **One behaviour per test**, named for it: `test_zero_total_returns_none`, not `test_share_2`.
- **Compare DataFrames with** `pd.testing.assert_frame_equal`. It checks values, dtypes, column order and index, and prints a clear difference.
- **Compare floats with** `pytest.approx` when the value comes from arithmetic: `assert mean == pytest.approx(0.3)`.
- **Use** `tmp_path` for anything that writes files. pytest gives each test a fresh temporary directory.

Check that each test can fail. Break the code on purpose, once, and watch the test go red. A test that stays green whatever you do protects nothing.

This course's own quality gate does the same: every seeded bug must fail its checker.

Run the tests in CI on every push (lesson A10-L3), and fix a failing test before merging anything else.

Why it matters

A share function had a test, `assert share(1, 4) == 0.25`, and 100% line coverage. The zero-total branch returned `0.0` instead of `None`. Dashboards showed empty segments as "0% of respondents", which the client read as a real finding. A four-word test, `assert share(3, 0) is None`, would have failed on the first run.

Worked example: weak code

```
1 def test_share():
2     result = share(1, 4)
3     1 assert result
```

1. **Line 3.** Passes for any truthy value: 0.25, 99 or 'oops'. And 0.0 would fail a correct result.

Worked example: strong code

```
1 import pytest
2
3
4 def share(part: float, total: float) -> float | None:
5     """part / total, or None when total is 0."""
6     if total == 0:
7         return None
8     return part / total
9
10
1 @pytest.mark.parametrize(("part", "total", "expected"), [(1, 4, 0.25),
    (0, 5, 0.0), (5, 5, 1.0)])
12 def test_share_values(part: float, total: float, expected: float) -> None:
13     assert share(part, total) == pytest.approx(expected)
14
15
16 def test_share_zero_total_is_none() -> None:
2     assert share(3, 0) is None
```

1. **Line 11.** Several cases, including 0 and 1, in one readable table.
2. **Line 17.** The edge case that most often ships broken, tested explicitly.

Test edges on purpose, with tiny explicit inputs, one behaviour per test, and prove each test can fail.

Recap card

- **A test worth writing** One that fails when the code is wrong. Break the code once to prove it.
- **Edge cases for data code** Empty input, missing values, zero denominators, one row, duplicate keys.
- **Compare DataFrames in tests** `pd.testing.assert_frame_equal`: values, dtypes, columns and index.
- **Compare floats in tests** `pytest.approx(expected)`.

Practice

Exercise 1. Multiple choice (level 1)

Which assertion best tests that `share(3, 0)` returns `None`?

- a) `assert not share(3, 0)`
- b) `assert share(3, 0) is None`
- c) `assert share(3, 0) == 0`
- d) `assert share(3, 0) != 1`

Why?

- r1) It fails for every result except `None`, including `0.0`.
- r2) `is` is faster than `==`.

Exercise 2. Predict the output (level 2)

What does this print? (pandas 3.0)

```
1 import pandas as pd
2
3 a = pd.DataFrame({"n": [1, 2]})
4 b = pd.DataFrame({"n": [1.0, 2.0]})
5 print(a.equals(b), (a == b).all().all())
6 try:
7     pd.testing.assert_frame_equal(a, b)
8     print("equal")
9 except AssertionError:
10    print("AssertionError")
```

Exercise 3. Fill the gap (level 2)

Run one test over three cases.

```
1 @pytest.mark.____(1)(("part", "total", "expected"), [(1, 4, 0.25), (0,
2     5, 0.0), (5, 5, 1.0)])
3 def test_share_values(part, total, expected):
4     assert share(part, total) == pytest.____(2)(expected)
```

Exercise 4. Write the function (level 3)

`share(part, total)` should return `part / total`, or `None` when `total` is 0. Write two tests:

- `test_share_basic()`: `share(1, 4)` equals 0.25.
- `test_share_zero_total()`: `share(3, 0)` is `None`.

The checker runs your tests against a correct `share`, which must pass, and against two broken versions, which your tests must catch.

Starter code:

```
1 def share(part: float, total: float) -> float | None:
2     """part / total, or None when total is 0."""
3     if total == 0:
4         return None
5     return part / total
6
7
8 def test_share_basic():
9     assert share(1, 4)
10
11
12 def test_share_zero_total():
13     assert not share(3, 0)
```

Exercise 5. Architecture design (level 4)

Choose test cases for `weighted_share(df, value, weight)` from lesson B8-L1.

Options: Four rows written out, expected value computed by hand, A 10,000-row sample of production data, expected value from the function itself.

1. Normal case
2. Edge cases
3. Float comparison

Sources

- `DataFrame.equals` returns `False` and `assert_frame_equal` raises for `int64` vs `float64` columns with equal values, while `==` reports `True`. (verified: Run on pandas 3.0.6 in this repository.)

B17-L2 Property-based tests with Hypothesis

Concept

Example tests check the cases you thought of. **Property-based tests** check a rule against hundreds of generated inputs, including ones you did not think

of. **Hypothesis** generates the inputs, and when a test fails it **shrinks** the input to the smallest failing case. Examples use Hypothesis 6.168.

```
1 from hypothesis import given, strategies as st
2
3 @given(st.lists(st.tuples(st.floats(-1e6, 1e6), st.floats(0, 10)),
4           min_size=1))
5 def test_weighted_mean_within_range(pairs):
6     values = [v for v, _ in pairs]
7     weights = [w for _, w in pairs]
8     m = weighted_mean(values, weights)
9     assert m is None or min(values) - 1e-6 <= m <= max(values) + 1e-6
```

Run against a `weighted_mean` without a zero guard, this test failed in this course's build with `ZeroDivisionError`, and Hypothesis shrank the input to `pairs=[(0.0, 0.0)]`: one value, one zero weight.

Properties that suit data code:

- **Bounds.** A mean lies between the minimum and maximum. A share lies in $[0, 1]$.
- **Invariance.** Shuffling rows does not change a total. Renaming a category does not change counts.
- **Round trip.** Writing then reading a file gives back the same frame.
- **Idempotence.** Deduplicating twice equals deduplicating once.
- **Grain.** The output has exactly one row per key.
- **Agreement.** A fast implementation equals a slow, obviously correct one.

Choose strategies deliberately. `st.floats()` generates `nan` and infinities by default. Bound it (`st.floats(-1e6, 1e6)`) or pass `allow_nan=False` unless your function must handle them, and then test that behaviour on purpose.

`hypothesis.extra.pandas` builds DataFrames with `data_frames` and `column`.

Use `@settings(max_examples=...)` to trade speed for coverage. Hypothesis saves failing examples in a local `.hypothesis/` directory and replays them first on the next run. Add it to `.gitignore`.

Hypothesis does not run in this app's browser runtime. The exercises check your code's structure.

Why it matters

A deduplication step kept the latest record per respondent. Its example tests used three tidy rows. In production, two records for one respondent had the **same** timestamp, and the function kept both, so the respondent was counted twice. A Hypothesis test of the grain, "the output has exactly one

row per respondent", fails on this bug at once. In this course's build it shrank the failure to `rows=[('a', 0), ('a', 0)]`: two identical rows.

Worked example: weak code

```
1 @given(st.floats(), st.floats())
2 def test_share_in_range(part, total):
3     assert 0 <= share(part, total) <= 1
```

1. **Line 1.** Unbounded floats include nan, infinities and negatives the function never promised to handle.
2. **Line 3.** Ignores the None branch, so the property is false for a correct function.

Worked example: strong code

```
1 from hypothesis import given, strategies as st
2
3 counts = st.integers(min_value=0, max_value=10_000)
4
5
6 @given(counts, counts)
7 def test_share_in_unit_interval(a: int, b: int) -> None:
8     part, total = min(a, b), max(a, b)
9     result = share(part, total)
10    if total == 0:
11        assert result is None
12    else:
13        assert result is not None and 0.0 <= result <= 1.0
```

1. **Line 3.** Inputs the function promises to handle: non-negative counts.
2. **Line 8.** Builds valid inputs (`part <= total`) instead of filtering most examples away.
3. **Line 11.** The zero branch is part of the property.

State a rule that must always hold, generate inputs from the domain the function promises, and let Hypothesis find and shrink the counterexample.

Recap card

- **Property-based test** A rule checked against many generated inputs, with failing inputs shrunk to a minimal case.
- **Four data-code properties** Bounds, invariance, round trip, idempotence, grain, and agreement with a slow version.
- **st.floats() default** Includes nan and infinities. Bound it or pass `allow_nan=False` on purpose.
- **Hypothesis example database** `.hypothesis/` replays past failures. Keep it out of Git.

Practice

Exercise 1. Multiple choice (level 1)

A Hypothesis test fails. What does Hypothesis report?

- a) The first random input that failed.
- b) A shrunk input: the simplest example it found that still fails.
- c) Every failing input it tried.
- d) Only the exception, without the input.

Why?

- r1) Shrinking turns a messy random failure into a minimal case you can read.
- r2) Hypothesis only ever generates simple inputs.

Exercise 2. Predict the output (level 2)

This property test runs with Hypothesis 6.168. Does it pass or fail?

```
1 from hypothesis import given, strategies as st
2
3
4 @given(st.floats())
5 def test_equal_to_itself(x):
6     assert x == x
```

Exercise 3. Fill the gap (level 2)

Generate bounded floats without nan, and run 500 examples.

```
1 @____(1)(max_examples=500)
2 @given(st.floats(min_value=-1e6, max_value=1e6, ____ (2)=False))
3 def test_round_trip(x):
4     assert parse(format(x)) == x
```

Exercise 4. Write the function (level 3)

Write a Hypothesis test `test_dedupe_idempotent` for `dedupe(rows: list[tuple[str, int]]) -> list[tuple[str, int]]`. Generate `rows` as lists of `(id, timestamp)` tuples with `st.lists(st.tuples(st.sampled_from(["a", "b", "c"]), st.integers(0, 5)))`. Assert that `dedupe(dedupe(rows)) == dedupe(rows)`.

Hypothesis does not run in the browser. The checker reads your code. Run it locally with `pytest`.

Starter code:

```
1 from hypothesis import given, strategies as st
2
```

```

3
4 def test_dedupe_idempotent():
5     rows = [("a", 1), ("a", 1)]
6     assert dedupe(rows) == [("a", 1)]

```

Exercise 5. Architecture design (level 4)

Pick a property for each function.

Options: Shuffling rows does not change the result, Write then read returns the same frame, Result lies between the minimum and maximum input, Equals a slow, obviously correct version.

1. `weighted_total(df)`: sum of value $\times$ weight
2. `save_parquet` and `load_parquet`
3. `weighted_mean(values, weights)` with positive weights
4. A vectorised rewrite of a row-by-row scoring loop

Sources

- A grain property found a dedupe that keeps timestamp ties and shrank it to `rows=[('a', 0), ('a', 0)]`. (verified: Run with Hypothesis 6.168.1 and pytest in this repository.)
- Hypothesis 6.168 found a `ZeroDivisionError` in an unguarded weighted mean and shrank it to `pairs=[(0.0, 0.0)]`; `st.floats()` generated nan and failed `x == x`. (verified: Run with Hypothesis 6.168.1 and pytest in this repository.)

B17-L3 Measure before you optimise

Concept

Most slow data code spends its time in one or two places, and they are rarely where you guess. Measure first, change one thing, measure again. All tools here are in the Python standard library.

Profile to find where time goes. `cProfile` counts calls and time per function:

```

1 import cProfile, pstats
2 cProfile.run("main()", "profile.out")
3 pstats.Stats("profile.out").sort_stats("cumulative").print_stats(10)

```

Sort by `cumulative` to see which of your functions, including everything they call, takes longest. In this course's build, a job calling a pure-Python sum of squares and a NumPy version over 200,000 numbers showed 0.037 s

cumulative for the Python one and 0.002 s for NumPy. In a notebook, `%prun main()` does the same.

Time a candidate change with `timeit`. `timeit.repeat(fn, repeat=5, number=100)` returns five totals, each for 100 calls. Report the **minimum** divided by `number`: other processes can only slow a run down, so the fastest run is the best estimate of the code's own cost. Use `time.perf_counter()`, not `time.time()`, for manual timing.

Measure memory when that is the limit. `tracemalloc` reports current and peak allocation: one million float64 values showed a peak of 8.0 MB, as expected at 8 bytes each. For pandas, `df.memory_usage(deep=True)` counts the contents of text columns, not just the pointers.

Rules:

- Profile with realistic data sizes. A 100-row sample hides what dominates at 10 million rows.
- Keep setup out of the timed region: reading files, building inputs, copying frames.
- Fix the biggest cost first. Speeding up a function that takes 2% of the runtime cannot save more than 2%.
- Keep a test around the code you optimise (lesson B17-L1), so faster does not become wrong.

Why it matters

An analyst spent two days rewriting a weighting function in Numba because "loops are slow in Python". The job still took 40 minutes. A profile, run afterwards, showed 36 of those minutes in reading 300 CSV files one at a time with default type inference. Switching to one Parquet read with explicit dtypes cut the job to 5 minutes. The Numba rewrite saved 20 seconds.

Worked example: weak code

```
❶ start = time.time()
❷ for _ in range(10):
❸     df = pd.read_csv("big.csv")
❹     result = transform(df)
❺ print("transform takes", (time.time() - start) / 10)
```

1. **Line 1.** Wall-clock time with low resolution. Use `time.perf_counter()`.
2. **Line 3.** The file read sits inside the timed loop, so the number is mostly I/O.
3. **Line 5.** An average over runs, inflated by any background activity.

Worked example: strong code

```
1 import time
2 from collections.abc import Callable
3
4
5 def best_time(fn: Callable[[], object], repeat: int = 5, number: int =
  10) -> float:
6     """Best per-call time in seconds over `repeat` rounds of `number`
    calls."""
7     if repeat < 1 or number < 1:
8         raise ValueError("repeat and number must be at least 1")
9     best = float("inf")
10    for _ in range(repeat):
11        ① start = time.perf_counter()
12        for _ in range(number):
13            fn()
14        ② best = min(best, (time.perf_counter() - start) / number)
15    return best
```

1. **Line 11.** A high-resolution clock meant for measuring intervals.
2. **Line 14.** Keep the fastest round: noise only ever adds time.

Profile to find the cost, time the change with `perf_counter` or `timeit`, keep setup out of the timing, and fix the biggest cost first.

Recap card

- **Find where time goes** `cProfile` with `pstats` sorted by cumulative, or `%prun` in a notebook.
- **Time a small change** `timeit.repeat` or `perf_counter`. Report the minimum per call.
- **Measure memory** `tracemalloc` peak, and `df.memory_usage(deep=True)` for pandas.
- **The order of work** Profile, fix the biggest cost, re-measure, keep a test around the change.

Practice

Exercise 1. Multiple choice (level 1)

A job takes 40 minutes. A profile shows 36 minutes in file reading and 1 minute in the weighting function. What should you optimise first?

- a) The weighting function, because loops are slow.
- b) File reading.
- c) Both equally.
- d) Neither: buy a faster machine.

Why?

- r1) A 1-minute function can save at most 1 minute. Reading holds 90% of the time.
- r2) I/O is always the bottleneck.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 import timeit
2
3 rounds = timeit.repeat(lambda: sum(range(1000)), repeat=5, number=100)
4 print(len(rounds))
```

Exercise 3. Fill the gap (level 2)

Profile `main()` and print the 10 functions with the highest cumulative time.

```
1 cProfile.run("main()", "profile.out")
2 pstats.Stats("profile.out").sort_stats("____(1)").____(2)(10)
```

Exercise 4. Write the function (level 3)

Write `best_time(fn, repeat=5, number=10) -> float`. Run `repeat` rounds. In each round, call `fn()` `number` times, timed with `time.perf_counter()`, and compute the per-call time. Return the smallest per-call time. Raise `ValueError` if `repeat` or `number` is below 1.

Starter code:

```
1 import time
2
3
4 def best_time(fn, repeat=5, number=10):
5     start = time.time()
6     fn()
7     return time.time() - start
```

Exercise 5. Spot the bug (level 4)

This benchmark says `add_features` takes 1.8 seconds. The function itself takes about 0.1 seconds. Tap the line that distorts the measurement.

```
1 import time
2
3 import pandas as pd
4
5
6 def benchmark(path: str, rounds: int) -> float:
7     """Best time of add_features over rounds."""
8     best = float("inf")
9     for _ in range(rounds):
```

```
10     start = time.perf_counter()
11     df = pd.read_parquet(path)
12     add_features(df)
13     best = min(best, time.perf_counter() - start)
14     return best
```

Write down the line number of the bug. Then write the fix.

Sources

- cProfile showed 0.037 s cumulative for a pure-Python sum of squares over 200,000 numbers and 0.002 s for NumPy; tracemalloc peak for one million float64 values was 8.0 MB. (verified: Run on Python 3.11 with NumPy 2.4.6 in this repository.)

B18 Documentation and communication

Roadmap: Prerequisite for all months; Every month. Docstrings, READMEs, model cards and decision logs.

B18-L1 Docstrings and READMEs that answer the next person's questions

Concept

The next person to read your code, often you in six months, has three questions: what does this do, how do I run it, and what will surprise me? Docstrings answer them for a function. The README answers them for the project.

Docstrings. This course uses the Google style (CLAUDE.md), which tools such as Sphinx and makedocstrings read:

```
1 def weighted_share(df: pd.DataFrame, value: str, weight: str) -> float |  
    None:  
2     """Weighted mean of a 0/1 column. Grain: one row per respondent.  
3  
4     Args:  
5         df: Respondent-level data.  
6         value: 0/1 column to average.  
7         weight: Non-negative weight column. Rows with a missing value or  
            weight are dropped.  
8  
9     Returns:  
10         The weighted share, or None when no rows remain or weights sum to  
            0.  
11  
12     Raises:  
13         KeyError: If a named column is missing.  
14     """
```

Put in the docstring what the signature cannot say: the **grain**, **units**, how **missing values** and **zeros** are handled, and what **None** means. Leave out what the code already says plainly, such as "loops over the rows".

`help(weighted_share)` and `weighted_share.__doc__` show the docstring. Editors show it on hover.

README. One page at the project root, in this order:

1. **Purpose:** one paragraph, including who uses the output.

2. **Quick start:** the exact commands to install and run, copied from a working terminal.
3. **Inputs and outputs:** sources, tables or files, their grain and schedule.
4. **Configuration:** where settings live and which ones matter.
5. **Tests and checks:** how to run them, and what CI enforces.
6. **Owners and support:** a team, not a person who might leave.

Keep both close to the truth. Update the docstring in the same commit as the code, and test the README's quick start in CI where you can.

Why it matters

A weighting function returned `None` for empty cells. Its docstring said "Returns the weighted share." A dashboard developer read that, assumed a float, and wrote `round(share * 100)`. The first empty cell crashed the nightly refresh. One line in the docstring, "or None when no rows remain", would have told them to handle it.

Worked example: weak code

```

1 def calc(df, v, w):
2     # calculate
3     d = df[[v, w]].dropna()
4     return (d[v] * d[w]).sum() / d[w].sum() if d[w].sum() else None

```

1. **Line 1.** Name, parameters and types say nothing about purpose or inputs.
2. **Line 2.** A comment that restates the obvious and hides what matters.

Worked example: strong code

```

1 import pandas as pd
2
3
4 def weighted_share(df: pd.DataFrame, value: str, weight: str) -> float |
  None:
5     """Weighted mean of a 0/1 column. Grain: one row per respondent.
6
7     Args:
8         df: Respondent-level data.
9         value: 0/1 column to average.
10        weight: Non-negative weight column. Rows missing value or weight
        are dropped.
11
12    Returns:
13        The weighted share, or None when no rows remain or weights sum to
        0.
14
15    Raises:
16        KeyError: If a named column is missing.

```

```

17     """
18     d = df[[value, weight]].dropna()
19     total = d[weight].sum()
20     if len(d) == 0 or total == 0:
21         return None
22     return float((d[value] * d[weight]).sum() / total)

```

1. **Line 5.** One-line summary with the grain.
2. **Line 13.** What None means, so callers handle it.
3. **Line 16.** The failure a caller might need to catch.

Docstrings state grain, units, missing values and what None means. The README states purpose, quick start, inputs, config, tests and owners.

Recap card

- **Google-style docstring sections** One-line summary, then Args:, Returns:, Raises:.
- **What a data docstring must state** Grain, units, missing and zero handling, and what None means.
- **README sections** Purpose, quick start, inputs and outputs, configuration, tests, owners.
- **Keeping docs true** Update docstrings in the same commit as the code. Test the quick start where possible.

Practice

Exercise 1. Multiple choice (level 1)

Which line adds the most value to a function's docstring?

- a) "Loops over rows and sums values."
- b) "Returns None when the weights sum to 0."
- c) "Written by the analytics team."
- d) "This is a function."

Why?

- r1) It tells callers about a behaviour the signature and name cannot show.
- r2) Short lines are better.

Exercise 2. Predict the output (level 2)

What does this print?

```

1 def share(part: float, total: float) -> float | None:
2     """part / total, or None when total is 0."""
3     return None if total == 0 else part / total

```

```
4
5
6 print(share. __doc__)
```

Exercise 3. Put in order (level 2)

Order the sections of a project README.

- Tests and checks
- Inputs and outputs, with grain and schedule
- Owners and support
- Quick start: install and run commands
- Purpose and users
- Configuration

Exercise 4. Write the function (level 3)

Add a Google-style docstring to `cap_weights`. It must have a one-line summary, an `Args:` section naming `weights` and `cap`, a `Returns:` section, and a `Raises:` section with `ValueError` for a cap below 1. Mention that missing weights stay missing. Keep the code as it is.

Starter code:

```
1 import pandas as pd
2
3
4 def cap_weights(weights: pd.Series, cap: float) -> pd.Series:
5     if cap < 1:
6         raise ValueError("cap must be at least 1")
7     return weights.clip(upper=cap)
```

Exercise 5. Architecture design (level 4)

Decide where each piece of information belongs.

Options: Docstring, README, Inline comment, Config file.

1. The function returns None for empty cells
2. The command to run the monthly job
3. Why one line uses a workaround for a library bug, with the issue link
4. The minimum reporting base of 50

B18-L2 Model cards that state limits as clearly as scores

Concept

A **model card** is a short document that travels with a model. It tells users what the model is for, how well it works, for whom, and where it should not be used. The format comes from Mitchell et al., "Model Cards for Model Reporting" (2019). Keep it in the repository next to the training code, and link it from the registered model version (lesson B14-L2).

Sections to include:

- **Model details:** name, version, registry alias, owners, training date, run ID.
- **Intended use:** the decision it supports and who uses it.
- **Out-of-scope uses:** uses it must not serve, stated plainly.
- **Training data:** source, period, grain, known gaps.
- **Evaluation:** the honest metric (lesson B11-L3), with the test period and prevalence.
- **Subgroup results:** rates per group with counts (lesson B13-L3), including the weakest group.
- **Limitations and risks:** known failure modes, drift risks, proxies.
- **Monitoring and review:** what is tracked, alert thresholds, next review date.

Write numbers with their context: "AUC 0.78 on October to December 2026 customers, 2.4% positive rate, 51,000 rows". A bare "AUC 0.78" invites misuse.

State limits as clearly as strengths. "Recall for customers over 70 is 0.28 against 0.62 overall. Do not use for outreach targeting of this group without human review" protects the people affected and the team that built the model.

Treat the card as code: a template with required sections, reviewed in the same pull request as the model change. A check in CI can refuse a card with missing sections.

Why it matters

A propensity model built for marketing emails was later reused by another team to prioritise collections calls. Nobody had written down that it was trained only on customers who had opted in to marketing, a group with very different payment behaviour. Collections priorities were wrong for months.

An "Out-of-scope uses" line, "Not validated for credit or collections decisions", would have stopped the reuse at the first review.

Worked example: weak code

```
1 Churn model v4
❶ AUC: 0.78
❷ Owner: Sam
❸ Works well, use for retention.
```

1. **Line 2.** No test period, prevalence or subgroup results.
2. **Line 3.** A person, not a team. Sam may leave.
3. **Line 4.** No limits, no out-of-scope uses, no monitoring.

Worked example: strong code

```
1 Model details: churn_model v4 (@champion), owner Retention Analytics, run
  4da8..., trained 2026-09-01
2 Intended use: rank active customers for the monthly retention call list
  (top 500)
❶ Out of scope: credit, collections or pricing decisions; customers with
  under 3 months' tenure
4 Training data: CRM extract Jan 2024 to Jun 2026, one row per customer-
  month
❷ Evaluation: AUC 0.78, average precision 0.21 (prevalence 2.4%), test Oct
  to Dec 2026, 51,000 rows
❸ Subgroups: recall at top 500 is 0.62 overall, 0.28 for customers over 70
  (n positives = 64)
7 Limitations: sparse digital activity lowers recall for older customers;
  tenure proxies contract type
8 Monitoring: monthly AUC and calibration, alert below 0.72; next review
  2027-03-01
```

1. **Line 3.** Stops reuse the model was never validated for.
2. **Line 5.** Each number with its period, baseline and size.
3. **Line 6.** The weakest group, stated with its count.

A model card states purpose, limits, data, honest metrics, subgroup results and monitoring, and is reviewed with every model change.

Recap card

- **Model card sections** Details, intended use, out-of-scope uses, training data, evaluation, subgroups, limitations, monitoring.
- **How to report a metric on a card** With test period, prevalence and row count.
- **Why out-of-scope uses** They stop reuse for decisions the model was never validated for.

- **Keep cards honest** Template with required sections, reviewed in the model's pull request, checked in CI.

Practice

Exercise 1. Multiple choice (level 1)

Which model card section would most directly have prevented a marketing model being reused for collections?

- a) Model details
- b) Out-of-scope uses
- c) Evaluation
- d) Monitoring

Why?

- r1) It names the uses the model was not built or validated for.
- r2) It is the longest section.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 REQUIRED = ["intended_use", "out_of_scope", "evaluation", "subgroups",
             "limitations"]
2 card = {"intended_use": "retention list", "evaluation": "AUC 0.78",
          "limitations": ""}
3 print([s for s in REQUIRED if not str(card.get(s, "")).strip()]])
```

Exercise 3. Fill the gap (level 2)

Complete the evaluation line so the metric has its context.

```
1 Evaluation: AUC 0.78, average precision 0.21 (____(1) 2.4%), test Oct to
   Dec 2026, 51,000 ____ (2)
```

Exercise 4. Write the function (level 3)

Write `missing_sections(card: dict[str, str], required: list[str]) -> list[str]`. Return the required section names, in the order of `required`, that are absent from `card` or contain only whitespace. CI will fail the build when the list is not empty.

Starter code:

```
1 def missing_sections(card, required):
2     return [s for s in required if s not in card]
```

Exercise 5. Architecture design (level 4)

Write the subgroup and limitations parts of a card for a hardship-support model.

Options: Recall and false positive rate per age band, with counts, weakest group first, Overall recall only.

1. Subgroup results
2. Known weakness for customers over 70
3. Review

Sources

- The model card format comes from Mitchell et al. (2019), Model Cards for Model Reporting. (verified: Mitchell et al., Proceedings of FAT* 2019, pages 220 to 229.)

B18-L3 Decision logs, so the reasons outlive the meeting

Concept

Every analysis rests on choices: the weighting scheme, the outlier rule, the metric, the cut-off, the fairness property. Six months later someone asks "why did we do it this way?". Without a record, the answer is a guess, and the choice gets reversed for the wrong reasons or kept for no reason.

A **decision log** is a list of short, dated entries, one per decision, kept in the repository (for example `docs/decisions/`). The format follows software architecture decision records:

```
1 ## D-007: Trim survey weights at 5
2 Date: 2026-09-12. Status: accepted.
3 Context: Four respondents had weights above 12 and moved brand share by 4
  points.
4 Options considered:
5 - No trimming
6 - Trim at 5, redistribute to keep the weighted total
7 - Rake with bounded weights
8 Decision: Trim at 5 with redistribution. Raking deferred until the new
  panel design.
9 Consequences: Design effect falls from 2.1 to 1.6. Small bias towards the
  sample profile. Review at wave 20.
```

Rules:

- **Write it when you decide**, in the same pull request as the code or config change.
- **List real options**, at least two, including "do nothing". A decision with one option was not a decision.
- **State the consequences**, including the costs you accept.
- **Never edit a past decision's reasoning**. When things change, add a new entry and mark the old one `superseded`, pointing to its replacement. The history of thinking is the value.
- **Link it**: from the config value it explains, from the model card (lesson B18-L2) and from the assumptions log (lesson B7-L3).

Keep entries short. Five to fifteen lines is enough for most analytical choices. A long, perfect document written after the fact helps less than a short one written on the day.

Why it matters

A new lead analyst removed the weight trimming from a tracker because it "looked arbitrary". The next wave's brand share jumped 5 points on four heavily weighted respondents, and the client asked why the trend broke. The original team had discussed exactly this case in a meeting two years earlier. Nothing was written down. A three-line decision entry linked from the config would have explained the rule in a minute.

Worked example: weak code

```
1 # config.yaml
❶ weight_cap: 5 # don't change!!
```

1. **Line 2**. A warning with no reason invites either blind obedience or careless removal.

Worked example: strong code

```
1 # config.yaml
❶ weight_cap: 5 # D-007, docs/decisions/D-007-trim-weights.md
3
4 # docs/decisions/D-007-trim-weights.md
5 ## D-007: Trim survey weights at 5
6 Date: 2026-09-12. Status: accepted.
❷ Options considered: no trimming; trim at 5 with redistribution; bounded
  raking.
8 Decision: trim at 5 with redistribution.
```

1. **Line 2.** The value links to its reasons.
2. **Line 7.** Real alternatives, including doing nothing.
3. **Line 9.** The cost accepted, and when to revisit.

One short, dated entry per decision, with options, consequences and status, linked from the code it explains, and superseded rather than edited.

Recap card

- **Parts of a decision entry** ID and title, date and status, context, options considered, decision, consequences.
- **When a decision changes** Add a new entry. Mark the old one superseded and link to the new one.
- **Where and when to write** In the repository, in the same pull request as the change.
- **Make it findable** Link from the config value, the model card and the assumptions log.

Practice

Exercise 1. Multiple choice (level 1)

A decision from last year no longer holds. What should you do with its log entry?

- a) Edit it to describe the new decision.
- b) Delete it.
- c) Add a new entry, and mark the old one superseded with a link to the new one.
- d) Leave both without links.

Why?

- r1) The log keeps the history of reasoning. Each entry records what was known when it was written.
- r2) Logs must never shrink for storage reasons.

Exercise 2. Predict the output (level 2)

What does this print?

```
1 log = {  
2   "D-003": {"status": "superseded", "superseded_by": "D-007"},  
3   "D-007": {"status": "accepted"},  
4 }
```

```
5 current = [k for k, v in log.items() if v["status"] == "accepted"]
6 print(current, log["D-003"]["superseded_by"] in log)
```

Exercise 3. Put in order (level 2)

Order the parts of a decision entry.

- Decision
- ID and title
- Consequences and review point
- Context: the problem and the evidence
- Date and status
- Options considered

Exercise 4. Write the function (level 3)

Write `decision_problems(entry: dict[str, object], log_ids: set[str]) -> list[str]` that checks one decision entry and returns problem messages, in this order:

1. "status" must be one of "proposed", "accepted", "superseded". Otherwise add "invalid status".
2. "options" must be a list with at least 2 items. Otherwise add "fewer than two options".
3. "consequences" must be non-blank text. Otherwise add "no consequences".
4. If status is "superseded", "superseded_by" must be in log_ids. Otherwise add "superseded without a valid replacement".

Starter code:

```
1 def decision_problems(entry, log_ids):
2     return []
```

Exercise 5. Architecture design (level 4)

Decide how the team handles decision records.

Options: docs/decisions/ in the repository, one file per decision, Slides from the steering meeting.

1. Where entries live
2. When an entry is written
3. How readers find it

Track B best-practice checklist

- ☐ The project has a pyproject.toml, a virtual environment and pinned dependencies.
- ☐ ruff and mypy pass before every commit.
- ☐ Functions are pure where possible. Side effects sit at the edges.
- ☐ Dtypes are explicit. pandas code is vectorised, with no row-wise apply on large frames.
- ☐ Inputs are validated with pandera or pydantic at every boundary.
- ☐ Seeds, config and data versions are recorded for every result.
- ☐ Weighted estimates use the survey weights. Intervals and effect sizes go with every estimate.
- ☐ Preprocessing sits inside a scikit-learn Pipeline, fitted on training folds only.
- ☐ The metric matches the decision. Calibration and subgroup performance are checked.
- ☐ Every experiment is tracked in MLflow.
- ☐ Tests cover edge cases, including empty inputs and zeros.
- ☐ Every model has a model card. Every decision has a log entry.

Track C: Microsoft Fabric

Unify data, analytics and Copilot in Fabric, and verify what Copilot produces.

C1 Fabric architecture

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Workspaces, capacities, items and OneLake as one logical lake.

C1-L1 Tenant, capacity, workspace and item

Concept

Everything in Fabric sits in one hierarchy. Learn it first, because cost, access and deployment all follow it.

- **Tenant.** Your organisation's Microsoft Entra ID tenant. Tenant settings, set by Fabric administrators, apply to everything below.
- **Capacity.** A pool of compute you buy. Every Fabric workload, from Spark to Power BI, draws from the capacity its workspace is assigned to.
- **Workspace.** A container where a team creates and shares items. Access is granted per workspace, through roles (lesson C1-L3).
- **Item.** A lakehouse, warehouse, notebook, pipeline, semantic model, report, eventhouse and so on. Items hold or process the data.

Capacities come as SKUs. F SKUs are the recommended Fabric capacities. The number is the compute in **capacity units** (CUs): an F2 has 2 CUs, an F64 has 64, up to F2048. A trial capacity has 64 CUs. Microsoft 365 P SKUs also run Fabric if Fabric is enabled on them.

Size affects licensing, not only speed. On **F64 or larger** (or a P capacity), users with a free Fabric license can view Power BI content. On smaller F SKUs, each viewer needs a Pro, Premium Per User or individual trial licence.

Pausing a capacity stops compute billing while it is paused. It is not free of consequences: when you pause, any smoothed and overage usage still owed is added to your Azure bill at once, and items on that capacity stop working until you resume.

Two practical rules follow from the hierarchy:

- Put dev, test and production in **separate workspaces**, and assign production to a capacity that experiments cannot exhaust.
- Before you design anything, find out which capacity SKU you have. It decides what users need to pay for and how much work runs at once.

Facts in this lesson come from Microsoft Learn's licensing and capacity pages, checked September 2026.

Why it matters

A team built a Power BI report for 300 frontline managers on an F32 capacity, assuming that managers with free licences could view it, as they had on the old Premium P1. Launch day produced 300 "You need a license" messages. The fix was either 300 Pro licences or a move to F64. Neither was in the budget, and the launch slipped a quarter.

Worked example: weak code

- | | |
|---|--|
| 1 | Workspace: Analytics (all projects) |
| 1 | Capacity: F32 shared with the data science sandbox |
| 2 | Viewers: 300 managers with free licences |
| 3 | Environments: dev and prod items in the same workspace |

1. **Line 2.** Production shares compute with experiments. One heavy notebook throttles the reports.
2. **Line 3.** Below F64, every viewer needs a Pro, PPU or trial licence.
3. **Line 4.** No separation: a change to a dev item can break production.

Worked example: strong code

- | | | |
|---|--|--------------------------------------|
| 1 | Workspace: sales_reporting_prod | -> capacity: F64 (production only) |
| 2 | Workspace: sales_reporting_test | -> capacity: F8 |
| 3 | Workspace: sales_reporting_dev | -> capacity: F8 |
| 4 | Viewers: 300 managers with free licences, reading the prod workspace's app | |
| 2 | Sandbox workspaces | -> capacity: F4, paused out of hours |

1. **Line 1.** F64 lets free-licence users view Power BI content, and production has its own compute.
2. **Line 5.** Experiments cannot starve production, and pausing saves money when nobody works.

The capacity decides who can view, how much runs at once and what it costs. Separate environments into workspaces, and give production its own capacity.

Recap card

- **Fabric hierarchy** Tenant, capacity, workspace, item.
- **What F64 means** An F SKU with 64 capacity units. F SKUs run from F2 to F2048.
- **Free-licence viewers of Power BI content** Need an F64 or larger (or P) capacity. Smaller F SKUs need Pro, PPU or trial licences for viewers.
- **Pausing a capacity** Stops compute billing from then on, bills remaining smoothed and overage usage at once, and stops its items.

Practice

Exercise 1. Multiple choice (level 1)

What does the number in the SKU name F32 tell you?

- a) 32 users can use the capacity.
- b) The capacity has 32 capacity units (CUs) of compute.
- c) 32 GB of OneLake storage is included.
- d) 32 workspaces can be assigned to it.

Why?

- r1) For F SKUs the number equals the capacity units: F2 is 2 CUs, F64 is 64.
- r2) SKU numbers are marketing names with no meaning.

Exercise 2. Put in order (level 2)

Order the Fabric hierarchy from the broadest to the most specific.

- Workspace
- Capacity
- Tenant
- Item, such as a lakehouse

Exercise 3. Architecture design (level 2)

A report will serve 200 viewers who have only free Fabric licences. Choose the set-up.

Options: F64 or larger, F16, cheaper, same features.

1. Capacity for the production workspace
2. What else runs on that capacity

Exercise 4. Write the function (level 3)

Write `cu_hours(sku: str, hours: float) -> float` for a cost estimate. The SKU is a string such as "F64". Return the capacity units multiplied by the hours running. Accept only the F SKUs F2, F4, F8, F16, F32, F64, F128, F256, F512, F1024 and F2048, and raise `ValueError` for any other SKU or for negative hours.

Starter code:

```
1 def cu_hours(sku, hours):  
2     return int(sku[1:]) * hours
```

Exercise 5. Multiple choice (level 4)

A team pauses its F64 capacity at 18:00 to save money after a heavy afternoon of Spark jobs. What happens?

- a) Nothing is billed for the afternoon's work, because the capacity is paused.
- b) Smoothed and overage usage still owed is added to the Azure bill at once, and items on the capacity stop working until it resumes.
- c) The capacity keeps running reports but stops Spark.
- d) OneLake data on the capacity is deleted after 24 hours.

Why?

- r1) Fabric smooths usage over time. Pausing settles what is still owed and makes the capacity's items unavailable.
- r2) Paused capacities are billed at double rate.

Sources

- F SKUs from F2 to F2048, where the number equals capacity units; a trial capacity has 64 CUs. (verified: <https://learn.microsoft.com/en-us/fabric/enterprise/licenses> (updated 2026-06-15))
- F64+ or P capacities support viewing Power BI content by users with a free Fabric licence; smaller F SKUs require Pro, PPU or trial licences for viewers. (verified: <https://learn.microsoft.com/en-us/fabric/enterprise/licenses> (updated 2026-06-15))
- When a capacity is paused, remaining cumulative overages and smoothed operations are added to the Azure bill. (verified: <https://learn.microsoft.com/en-us/fabric/enterprise/pause-resume> (updated 2026-04-20): pausing can prevent content from being available until resumed)

C1-L2 OneLake, one lake per tenant

Concept

Every Fabric tenant has exactly **one OneLake**. You cannot create a second one or delete it. It is built on Azure Data Lake Storage and stores tables in open formats: **Delta Parquet** or **Iceberg**. Every Fabric engine (Spark, the SQL engines, Power BI's Direct Lake, KQL) reads the same files, so one copy of the data serves them all.

OneLake mirrors the Fabric hierarchy: a folder per **workspace**, and inside it a folder per **item**, named with its type. A lakehouse holds two areas: **Tables** for managed Delta tables and **Files** for anything else.

You address data with an ADLS-style path:

```
1 https://onelake.dfs.fabric.microsoft.com/<workspace>/<item>.<itemtype>/<path>
2 abfss://<workspace>@onelake.dfs.fabric.microsoft.com/<item>.<itemtype>/<path>
```

For example `abfss://sales_prod@onelake.dfs.fabric.microsoft.com/lh_sales.lakehouse/Tables/orders`. The item type is lower case: `.lakehouse`, `.warehouse`.

Names change. **GUIDs do not**. OneLake accepts the workspace and item GUIDs in place of names, and they stay fixed when someone renames the workspace or item. With GUIDs you give both the workspace and the item as GUIDs, and you drop the item type. You find them in the item's URL in the Fabric portal.

Use GUIDs in pipelines, jobs and code that other people will run. Two more reasons from Microsoft Learn:

- The `abfs` driver URI does not allow special characters such as spaces in the workspace name. GUIDs avoid the problem.
- The global endpoint `onelake.dfs.fabric.microsoft.com` can route through another region. If data residency matters, use your workspace's regional endpoint, `https://<region>-onelake.dfs.fabric.microsoft.com`.

Because tools that speak ADLS APIs can read OneLake, Azure Databricks, Azure Storage Explorer or Python's Azure SDK can read Fabric data with the right permissions.

Why it matters

A Databricks job read a Fabric lakehouse through `abfss://Sales Analytics@onelake.dfs.fabric.microsoft.com/...`. It failed on the first run, because the workspace name held a space. The engineer renamed the workspace to `SalesAnalytics`. The job worked, and four Power BI dataflows and two pipelines that referred to the old name broke that night. Using the workspace and item GUIDs would have avoided both failures.

Worked example: weak code

```
❶ path = "abfss://Sales Analytics@onelake.dfs.fabric.microsoft.com/Sales
        LH.Lakehouse/Tables/orders"
```

```
2 df = spark.read.format("delta").load(path)
```

1. **Line 1.** Spaces break the abfs URI, the item type is not lower case, and a rename breaks the path.

Worked example: strong code

```
1 ONELAKE_HOST = "onelake.dfs.fabric.microsoft.com"
2
3
4 def onelake_table_path(workspace_id: str, item_id: str, table: str) ->
    str:
5     """abfss path to a lakehouse table, by GUID, stable across renames."""
6     ❶ return f"abfss://{workspace_id}@{ONELAKE_HOST}/{item_id}/Tables/
    {table}"
7
8
9     ❷ path = onelake_table_path(CONFIG["workspace_id"], CONFIG["lakehouse_id"],
    "orders")
10 df = spark.read.format("delta").load(path)
```

1. **Line 6.** With GUIDs, both workspace and item are IDs and the item type is dropped.
2. **Line 9.** IDs come from config, not from code (lesson A1-L3).

One lake, one copy, many engines. Address data by GUID in anything automated.

Recap card

- **How many OneLakes per tenant?** Exactly one. You cannot create or delete OneLakes.
- **OneLake table formats** Delta Parquet or Iceberg, open formats every Fabric engine reads.
- **OneLake path by name** abfss://
<workspace>@onelake.dfs.fabric.microsoft.com/<item>.<itemtype>/
Tables/<table>.
- **Why GUIDs in automation** They survive renames, avoid special-character problems in abfs URIs, and drop the item type.

Practice

Exercise 1. Multiple choice (level 1)

Your company has three business units in one Fabric tenant. How many OneLakes are there?

- a) One per business unit

- b) One for the whole tenant
- c) One per capacity
- d) One per workspace

Why?

- r1) Every tenant has a single OneLake. Workspaces organise it; they do not create new lakes.
- r2) OneLake licences are sold per tenant.

Exercise 2. Fill the gap (level 2)

Complete the path to table `orders` in lakehouse `lh_sales`, workspace `sales_prod`.

```
1 abfss://____(1)@onelake.dfs.fabric.microsoft.com/lh_sales.____(2)/
   ____ (3)/orders
```

Exercise 3. Spot the bug (level 3)

This path fails with an invalid URI error. Tap the line with the problem and rewrite the path so it survives renames too.

```
1 WORKSPACE = "Sales Analytics"
2 LAKEHOUSE = "lh_sales"
3 path = f"abfss://{WORKSPACE}@onelake.dfs.fabric.microsoft.com/
   {LAKEHOUSE}.lakehouse/Tables/orders"
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 3)

Write `onelake_path(workspace: str, item: str, item_type: str, relative: str) -> str` returning `abfss://<workspace>@onelake.dfs.fabric.microsoft.com/<item>.<item_type>/<relative>`.

- `item_type` must be `"lakehouse"` or `"warehouse"`. Otherwise raise `ValueError`.
- Raise `ValueError` if `workspace` or `item` contains a space.
- Strip leading slashes from `relative`.

Starter code:

```
1 def onelake_path(workspace, item, item_type, relative):
2     return f"abfss://{workspace}@onelake.dfs.fabric.microsoft.com/{item}.
   {item_type}/{relative}"
```

Exercise 5. Architecture design (level 4)

An Azure Databricks job in the UK will read a Fabric lakehouse in a UK South workspace every night. Data must stay in the UK. Choose the set-up.

Options: GUIDs from config, Display names in the code.

- 1. How the job refers to workspace and lakehouse
- 2. Endpoint
- 3. Data movement

Sources

- Every tenant has a single OneLake, built on ADLS, storing tables in Delta Parquet or Iceberg; you cannot create multiple OneLakes. (verified: <https://learn.microsoft.com/en-us/fabric/onelake/onelake-overview> (updated 2026-03-09))
- OneLake URI formats by name and by GUID; GUIDs do not change on rename; abfs URIs do not allow special characters such as spaces in the workspace name; the global endpoint may route data out of region, and regional endpoints avoid this. (verified: <https://learn.microsoft.com/en-us/fabric/onelake/onelake-access-api> (updated 2026-05-18))

C1-L3 Workspace roles, granted to groups

Concept

A workspace has four roles: **Admin**, **Member**, **Contributor** and **Viewer**. Each role includes the rights of the ones below it, with a few rights reserved for the top.

From Microsoft Learn's role table (checked September 2026):

Right	Admin	Member	Contributor	Viewer
Update or delete the workspace, add or remove admins	yes			
Connect the workspace to Git, create a workspace identity	yes			
Add members or lower roles, allow resharing	yes	yes		
Create and change items, run notebooks and pipelines	yes	yes	yes	
Read lakehouse and warehouse data through Spark and OneLake APIs (ReadAll)	yes	yes	yes	
	yes	yes	yes	yes

Right	Admin	Member	Contributor	Viewer
Read lakehouse and warehouse data with T-SQL (ReadData)				

Two consequences surprise people coming from Databricks:

- A **Viewer** can query a lakehouse through its SQL analytics endpoint, but cannot read the same tables from a Spark notebook or through OneLake APIs.
- A **Contributor** can change and delete every item in the workspace, including production lakehouse tables. Contributor in production is powerful.

Rules that keep access manageable:

- **Assign roles to Microsoft Entra security groups**, not to individuals (as in lesson A4-L2). People join and leave groups. Role assignments stay the same.
- **Least privilege per environment.** Engineers are Contributors in dev. In production, deployment runs under an automation identity, and people are Viewers unless they are on call.
- **Keep at least two admins**, as groups or named people, so a departure does not lock the workspace.
- Workspace roles are coarse. For finer control inside an item, use item sharing, SQL permissions, or OneLake security (lesson C9-L2).

Why it matters

A contractor was added to a production workspace as Contributor "to look at a notebook". Two weeks later they ran a clean-up notebook against the wrong lakehouse and dropped three tables that reports depended on. Viewer would have let them read the notebook and query the data with T-SQL, and would not have let them change anything.

Worked example: weak code

```

1 workspace: sales_prod
2 roles:
  ① Admin: [sam@contoso.com]
  ② Contributor: [grp-sales-analysts, contractor@partner.com]
5 Viewer: []

```

1. **Line 3.** One named admin: if Sam leaves, nobody can manage the workspace.
2. **Line 4.** Write access to production for people who need to read, including an external individual.

Worked example: strong code

```
1 workspace: sales_prod
2 roles:
3   ① Admin: [grp-fabric-platform-admins]
4   ② Member: [grp-sales-data-leads]
5   ③ Contributor: [sp-sales-deployment]
6   ④ Viewer: [grp-sales-analysts, grp-sales-contractors]
```

1. **Line 3.** A group of admins: no single point of failure.
2. **Line 5.** Only the deployment identity writes to production.
3. **Line 6.** Readers query with T-SQL and reports, and cannot change items.

Roles go to groups, write access in production goes to automation, and people read.

Recap card

- **Fabric workspace roles** Admin, Member, Contributor, Viewer.
- **Admin-only rights** Update or delete the workspace, manage admins, connect Git, create a workspace identity.
- **Viewer and lakehouse data** Reads with T-SQL (ReadData). Not through Spark or OneLake APIs (ReadAll starts at Contributor).
- **Who writes to production?** An automation identity as Contributor. People get Viewer, through groups.

Practice

Exercise 1. Multiple choice (level 1)

Which workspace role can connect a workspace to a Git repository?

- a) Admin only
- b) Admin and Member
- c) Any role except Viewer
- d) Any role

Why?

- r1) Connecting Git, like deleting the workspace or creating a workspace identity, is reserved for Admins.
- r2) Git is a paid feature for Admins.

Exercise 2. Multiple choice (level 2)

An analyst with the Viewer role runs `spark.read.table("lh_sales.orders")` in a notebook in that workspace. What happens?

- a) It works: Viewers read everything.
- b) It fails: Viewers read lakehouse data with T-SQL, not through Spark or OneLake APIs.
- c) It works, but returns only the first 1,000 rows.
- d) It fails because Viewers cannot see notebooks at all.

Why?

- r1) Viewer includes ReadData (T-SQL) but not ReadAll (Spark and OneLake APIs).
- r2) Spark is disabled for all users in production workspaces.

Exercise 3. Architecture design (level 2)

Assign roles in the production workspace `sales_prod`.

Options: Admin, Member, Contributor, Viewer.

1. Fabric platform team (group)
2. Deployment service principal
3. Sales analysts (group)

Exercise 4. Write the function (level 3)

Write `min_role(action: str) -> str` returning the lowest workspace role that allows the action. Use these actions from the role table:

- `"connect_git"`, `"delete_workspace"`: Admin
- `"add_members"`: Member
- `"run_notebook"`, `"spark_read"`: Contributor
- `"tsql_read"`: Viewer

Raise `KeyError` for any other action.

Starter code:

```
1 def min_role(action):  
2     return "Viewer"
```

Exercise 5. Spot the bug (level 4)

This role file for the production workspace breaks two rules of the lesson. Tap the line with the more dangerous one, then rewrite the file.

```
1 workspace: sales_prod
```

```
2 roles:
3   Admin: [grp-fabric-platform-admins]
4   Contributor: [sp-sales-deployment, grp-sales-analysts]
5   Viewer: [jo.bloggs@contoso.com]
```

Write down the line number of the bug. Then write the fix.

Sources

- The four workspace roles and their rights, including Admin-only Git connection and Viewer having ReadData but not ReadAll. (verified: <https://learn.microsoft.com/en-us/fabric/fundamentals/roles-workspaces> (updated 2026-04-24))

C2 Lakehouse, Warehouse and SQL analytics endpoint

Roadmap: Months 16 to 18: MLOps and deployment; Current work. When to use each.

C2-L1 Lakehouse Tables and Files, and schemas

Concept

A lakehouse has two top-level areas in OneLake:

- **Tables** holds **managed Delta tables**. Fabric registers them, the SQL analytics endpoint exposes them, and Direct Lake semantic models read them.
- **Files** holds anything else: CSV, JSON, Parquet, images, raw deliveries.

Only Delta tables in Tables reach the SQL engines. A Parquet or CSV file is invisible to the SQL analytics endpoint, wherever it sits. If you put a folder that is not a valid Delta table under Tables, the lakehouse explorer shows it in an **Unidentified** area. A broken shortcut also lands there.

So the flow for a delivery is: land the raw file in **Files** (your bronze copy, as received), then read it with Spark and write a Delta table to **Tables** with `saveAsTable`.

Schemas. New lakehouses created in the portal have schemas enabled by default. Every schema-enabled lakehouse has a `dbo` schema that you cannot rename or remove. A table written without a schema goes to `dbo`. Create schemas by domain (`sales`, `hr`) to keep a growing lakehouse navigable.

In Spark SQL you reference tables with a **four-part name**:

`workspace.lakehouse.schema.table`. Parts you leave out come from defaults: a bare table name means the `dbo` schema of the notebook's default lakehouse. For code that others run, name at least the lakehouse and schema, so it does not depend on which lakehouse happens to be attached.

One change to note if you learned Fabric earlier: since 5 September 2025, Fabric no longer creates a default semantic model automatically with each new lakehouse. Create semantic models on purpose (lesson C6-L1).

Why it matters

An engineer wrote cleaned survey data as Parquet into `Files/clean/respondents/`, the way they had written to cloud storage before. The Power BI team could not find the table in the SQL analytics endpoint, so they imported the Parquet into a separate dataflow, a second copy with its own refresh. Writing a Delta table to Tables with `saveAsTable` would have exposed it to every engine with no copy.

Worked example: weak code

```
❶ raw = spark.read.option("header", True).csv("Files/landing/
    respondents_w14.csv")
❷ raw.write.mode("overwrite").parquet("Files/clean/respondents")
```

1. **Line 1.** Every column read as a string, types unchecked.
2. **Line 2.** Parquet under Files: invisible to the SQL analytics endpoint and Direct Lake.

Worked example: strong code

```
1 from pyspark.sql import DataFrame, SparkSession
2 from pyspark.sql import types as T
3
4 RESPONDENT_SCHEMA = T.StructType([
5     T.StructField("respondent_id", T.StringType(), False),
6     T.StructField("wave", T.IntegerType(), False),
7     T.StructField("weight", T.DoubleType(), True),
8 ])
9
10
11 def load_delivery(spark: SparkSession, path: str, table: str) ->
    DataFrame:
12     """Read a CSV delivery from Files and append it to a managed Delta
        table in Tables."""
    ❶ df = spark.read.option("header",
        True).schema(RESPONDENT_SCHEMA).csv(path)
    ❷ df.write.mode("append").saveAsTable(table)
15     return df
16
17
    ❸ load_delivery(spark, "Files/landing/respondents_w14.csv",
        "lh_survey.sales.respondents")
```

1. **Line 13.** An explicit schema: types are decided, not inferred.
2. **Line 14.** A managed Delta table in Tables: visible to Spark, T-SQL and Direct Lake.
3. **Line 18.** Lakehouse and schema named, so the code does not depend on the attached default.

Raw files in Files, Delta tables in Tables, and schema-qualified names in shared code.

Recap card

- **Lakehouse Tables vs Files** Tables: managed Delta tables, visible to SQL and Direct Lake. Files: raw and non-Delta data.
- **Non-Delta folder under Tables** Shown in the Unidentified area, not queryable.
- **Lakehouse schemas** On by default for new lakehouses. dbo is the fixed default schema.
- **Spark four-part name** workspace.lakehouse.schema.table. A bare name means dbo in the default lakehouse.

Practice

Exercise 1. Multiple choice (level 1)

A Parquet file sits at `Files/exports/orders.parquet` in a lakehouse. What can the SQL analytics endpoint do with it?

- a) Query it as a table named orders.
- b) Nothing. Only Delta tables in Tables appear in the endpoint.
- c) Query it after a metadata refresh.
- d) Convert it to Delta automatically.

Why?

- r1) The endpoint exposes Delta tables only. Parquet and CSV cannot be queried through it.
- r2) Files is a private folder only the owner can see.

Exercise 2. Multiple choice (level 2)

In a notebook attached to schema-enabled lakehouse `lh_survey`, you run `spark.read.table("respondents")`. Which table does Spark read?

- a) `lh_survey.dbo.respondents`
- b) The first schema that has a respondents table
- c) It raises an error: schemas require a four-part name
- d) respondents in every schema, unioned

Why?

- r1) A bare table name resolves to the dbo schema of the default lakehouse.

- r2) dbo is the only schema Spark can read.

Exercise 3. Fill the gap (level 2)

Join two tables in different workspaces in one Spark SQL query.

```
1 SELECT e.name, d.dept_name
2 FROM ops_ws.hr_lakehouse.____(1).employees AS e
3 JOIN global_ws.corporate_lakehouse.company.departments AS d
4   ON e.deptno = d.deptno
```

Exercise 4. Write the function (level 3)

Write `load_delivery(spark, path, table)`. Read the CSV at `path` with a header and the given `RESPONDENT_SCHEMA`, append it to the managed table `table` with `saveAsTable`, and return the DataFrame you read. Do not write Parquet to Files.

Starter code:

```
1 from pyspark.sql import types as T
2
3 RESPONDENT_SCHEMA = T.StructType([
4     T.StructField("respondent_id", T.StringType(), False),
5     T.StructField("wave", T.IntegerType(), False),
6     T.StructField("weight", T.DoubleType(), True),
7 ])
8
9
10 def load_delivery(spark, path, table):
11     df = spark.read.option("header", True).csv(path)
12     df.write.mode("overwrite").parquet("Files/clean/respondents")
13     return df
```

Exercise 5. Architecture design (level 4)

Lay out a lakehouse for a survey programme with raw CSV deliveries, cleaned respondent data and reporting tables.

Options: Files, Tables, dbo schema, Tables, a named schema per domain.

1. Raw CSV deliveries, kept as received
2. Cleaned respondent-level data
3. Reporting tables read by Power BI

Sources

- Lakehouse Tables holds managed Delta tables and Files holds unstructured or non-Delta data; only Delta tables appear in the SQL analytics endpoint; default semantic models are no longer created automatically since 5 September 2025. (verified: <https://>

learn.microsoft.com/en-us/fabric/data-engineering/lakehouse-overview (updated 2026-05-07))

- Non-Delta folders under Tables and broken shortcuts appear in the Unidentified area. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/navigate-lakehouse-explorer> (updated 2026-05-07))
- Schemas are enabled by default for new lakehouses; dbo cannot be renamed or removed; Spark uses workspace.lakehouse.schema.table and resolves omitted parts from defaults. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/lakehouse-schemas> (updated 2026-02-24))

C2-L2 Warehouse T-SQL, and where it differs from SQL Server

Concept

Fabric Warehouse speaks T-SQL, and most SQL Server habits carry over: `CREATE TABLE`, `INSERT`, `UPDATE`, `DELETE`, stored procedures, views, `TRUNCATE TABLE`, identity columns and explicit transactions. `MERGE` is generally available. Its tables are Delta tables in OneLake, so Spark and Power BI read them too.

The differences cause real bugs. Microsoft Learn, checked September 2026:

Keys are not enforced. `PRIMARY KEY` is supported only as `NONCLUSTERED NOT ENFORCED`. `UNIQUE` likewise, and `FOREIGN KEY` only `NOT ENFORCED`. The engine uses them as hints for query plans. It does **not** stop duplicates or orphans. Deduplicate and check keys yourself (lesson A3-L1 for `MERGE` sources, lesson A6-L1 for contracts).

Some data types are missing, with documented alternatives:

Not supported	Use
<code>nchar</code> , <code>nvarchar</code>	<code>char</code> , <code>varchar</code> (UTF-8 collation)
<code>money</code> , <code>smallmoney</code>	<code>decimal</code>
<code>datetime</code> , <code>smalldatetime</code>	<code>datetime2</code> (up to 6 fractional digits)
<code>tinyint</code>	<code>smallint</code>
<code>text</code> , <code>ntext</code> , <code>json</code>	<code>varchar</code>

`varchar(max)` is limited to 16 MB per value.

Some features are absent or in preview. Triggers are not supported.

`ALTER COLUMN` and nested CTEs are in preview. Check the T-SQL surface area page before you port a procedure that relies on them.

Porting a SQL Server script is usually a type-mapping job plus one design change: replace constraint enforcement with explicit checks.

Why it matters

A team ported a customer dimension from SQL Server, where `PRIMARY KEY` had rejected duplicate loads for years. In Fabric the same key was created `NOT ENFORCED`, so a pipeline that ran twice on the same file loaded every customer twice. Sales per customer halved in the joined reports for a week before anyone traced it. A duplicate check after each load would have failed the run on the first rerun.

Worked example: weak code

```
1 CREATE TABLE dbo.dim_customer (  
  ① customer_id INT PRIMARY KEY,  
  ② name NVARCHAR(200),  
  ③ credit_limit MONEY,  
  ④ created_at DATETIME  
6 );
```

1. **Line 2.** An enforced clustered primary key is not supported in Fabric Warehouse.
2. **Line 3.** `nvarchar` is not supported. Use `varchar` with the UTF-8 collation.
3. **Line 4.** `money` is not supported. Use `decimal`.
4. **Line 5.** `datetime` is not supported. Use `datetime2`.

Worked example: strong code

```
1 CREATE TABLE dbo.dim_customer (  
2   customer_id INT NOT NULL,  
3   name VARCHAR(200),  
  ① credit_limit DECIMAL(19, 4),  
4   created_at DATETIME2(6)  
5 );  
6  
7 ALTER TABLE dbo.dim_customer  
  ② ADD CONSTRAINT pk_dim_customer PRIMARY KEY NONCLUSTERED (customer_id)  
    NOT ENFORCED;  
8  
9  
10 -- The key is a hint, so check it after every load.  
11 SELECT customer_id, COUNT(*) AS n  
12 FROM dbo.dim_customer  
13 GROUP BY customer_id  
  ③ HAVING COUNT(*) > 1;
```

1. **Line 4.** `decimal` replaces `money`, with the scale stated.

2. **Line 8.** The only supported form of primary key: a hint for the optimiser, not a guard.
3. **Line 14.** The duplicate check the key no longer performs. Any row returned fails the load.

Map the types, declare keys as NOT ENFORCED hints, and check uniqueness yourself after every load.

Recap card

- **Primary keys in Fabric Warehouse** Only PRIMARY KEY NONCLUSTERED ... NOT ENFORCED. Duplicates are not blocked.
- **nvarchar, money, datetime in Fabric Warehouse** Not supported. Use varchar, decimal, datetime2.
- **MERGE and TRUNCATE in Fabric Warehouse** Both supported. MERGE is generally available.
- **Triggers in Fabric Warehouse** Not supported. Put the logic in the loading procedure.

Practice

Exercise 1. Multiple choice (level 1)

A Fabric Warehouse table has PRIMARY KEY NONCLUSTERED (customer_id) NOT ENFORCED. You insert two rows with customer_id = 7. What happens?

- a) The second insert fails with a key violation.
- b) Both rows are stored. The key is not enforced.
- c) The second row replaces the first.
- d) The table is locked until you remove a duplicate.

Why?

- r1) NOT ENFORCED keys are optimiser hints. The engine does not check them.
- r2) Fabric deduplicates later in a background job.

Exercise 2. Fill the gap (level 2)

Add the only supported form of primary key to dbo.dim_customer.

```
1 ALTER TABLE dbo.dim_customer
2     ADD CONSTRAINT pk_dim_customer PRIMARY KEY ____ (1) (customer_id)
   ____ (2);
```

Exercise 3. Spot the bug (level 3)

This table definition, ported from SQL Server, fails in Fabric Warehouse. Tap one failing line, then rewrite the definition with supported types.

```
1 CREATE TABLE dbo.fact_payment (  
2     payment_id BIGINT NOT NULL,  
3     customer_id INT NOT NULL,  
4     amount MONEY NOT NULL,  
5     note NVARCHAR(500),  
6     paid_at DATETIME2(6) NOT NULL  
7 );
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 3)

Write a query that returns every `customer_id` appearing more than once in `dbo.dim_customer`, with its count as `n`. The load job fails if it returns any row.

Starter code:

```
1 SELECT DISTINCT customer_id FROM dbo.dim_customer;
```

Exercise 5. Architecture design (level 4)

You port a SQL Server data mart to Fabric Warehouse. Decide how to handle each feature.

Options: NOT ENFORCED keys plus an explicit duplicate check after each load, NOT ENFORCED keys, trusting the source.

1. Primary keys that used to block duplicates
2. An audit trigger on the orders table
3. nvarchar and money columns

Sources

- Identity columns, TRUNCATE TABLE and MERGE (GA) are supported; triggers are not; ALTER COLUMN and nested CTEs are in preview. (verified: <https://learn.microsoft.com/en-us/fabric/data-warehouse/tsql-surface-area> (updated 2026-08-26))
- PRIMARY KEY and UNIQUE are supported only as NONCLUSTERED NOT ENFORCED, FOREIGN KEY only as NOT ENFORCED. (verified: <https://learn.microsoft.com/en-us/fabric/data-warehouse/table-constraints> (updated 2026-09-08))
- Unsupported types and alternatives: nchar/nvarchar to char/varchar, money to decimal, datetime to datetime2, tinyint to smallint; varchar(max) limited to 16 MB; datetime2 limited to 6 fractional digits.

(verified: <https://learn.microsoft.com/en-us/fabric/data-warehouse/data-types> (updated 2026-06-11))

C2-L3 Lakehouse, Warehouse or SQL analytics endpoint

Concept

Fabric gives you three ways to hold and query tables. All three store data as Delta tables in OneLake. They differ in who writes the data and in which language.

Lakehouse. You write with Spark: notebooks, Spark job definitions, Dataflows Gen2 or pipelines. It holds files and Delta tables side by side. Choose it for raw files, PySpark transformations and data science. This is the closest match to your Databricks work.

SQL analytics endpoint. Every lakehouse gets one automatically. It gives a **read-only** T-SQL surface over the lakehouse's Delta tables. You can `SELECT`, and you can create views, functions and stored procedures. You cannot `INSERT`, `UPDATE` or `DELETE` table data through it. It picks up new or changed tables through a metadata sync. Microsoft documents the usual lag as under one minute, but it varies.

Warehouse. You write with T-SQL. It supports full DML and DDL, and multi-table transactions. Choose it when a team works in T-SQL, loads data with stored procedures, or needs several tables to change together or not at all.

One rule decides most cases: **who writes the data, and in which language?** Spark writers go to a lakehouse. T-SQL writers go to a warehouse. T-SQL readers of lakehouse data use the SQL analytics endpoint.

You do not have to copy data between them. A warehouse query reads lakehouse tables in the same workspace with a three-part name:

`lakehouse_name.schema.table`. Direct Lake semantic models read from either.

Check the status of each feature you rely on in `VERSIONS.md` before you build. Fabric changes monthly.

Why it matters

A team moves a finance reconciliation from SQL Server to Fabric. They land every table in a lakehouse, because the Databricks team already uses lakehouses. The month-end stored procedure updates three tables. On the SQL analytics endpoint it fails on its first `UPDATE`, because the endpoint is read-only.

Under deadline, someone rewrites it as three separate Spark notebook writes. On one month-end the second notebook fails. Table one holds the new month and table three holds the old one, and the reconciliation report shows a balance that does not tie. A warehouse transaction would have rolled back all three tables together.

Worked example: weak code

```
❶ -- Connected to the SQL analytics endpoint of lakehouse lh_finance
2 CREATE PROCEDURE dbo.close_month @period CHAR(7) AS
3 BEGIN
❷ UPDATE dbo.ledger SET status = 'closed' WHERE period = @period;
5 INSERT INTO dbo.balances
6 SELECT account_id, @period, SUM(amount) FROM dbo.ledger
7 WHERE period = @period GROUP BY account_id;
❸ DELETE FROM dbo.open_items WHERE period = @period;
9 END;
```

1. **Line 1.** The endpoint gives read-only access to lakehouse tables. Every write below fails.
2. **Line 4.** UPDATE, INSERT and DELETE on lakehouse tables are not supported through the endpoint.
3. **Line 8.** Three tables must change together. Without a transaction, a failure halfway leaves them inconsistent.

Worked example: strong code

```
❶ -- Connected to warehouse wh_finance
2 CREATE PROCEDURE dbo.close_month @period CHAR(7) AS
3 BEGIN
❷ BEGIN TRANSACTION;
5 UPDATE dbo.ledger SET status = 'closed' WHERE period = @period;
❸ INSERT INTO dbo.balances (account_id, period, balance)
7 SELECT account_id, @period, SUM(amount) FROM dbo.ledger
8 WHERE period = @period GROUP BY account_id;
9 DELETE FROM dbo.open_items WHERE period = @period;
10 COMMIT TRANSACTION;
11 END;
```

1. **Line 1.** The T-SQL writer works in a warehouse, which supports full DML.
2. **Line 4.** All three changes commit together or not at all. Warehouses support multi-table transactions.
3. **Line 6.** Name the target columns. An INSERT without a column list breaks when someone adds a column.

The data and the logic are the same. The item type decides whether the procedure runs at all, and whether a failure leaves the tables consistent.

Recap card

- **Lakehouse, Warehouse or endpoint: the deciding question** Who writes the data, and in which language? Spark writes to a lakehouse, T-SQL writes to a warehouse.
- **What can you do on the SQL analytics endpoint?** SELECT and create views, functions and procedures. No INSERT, UPDATE or DELETE on table data.
- **When do you need a Warehouse over a Lakehouse?** T-SQL writes, stored procedure loads, and multi-table transactions.
- **Read a lakehouse table from a warehouse in the same workspace** Use a three-part name: lakehouse.schema.table. No copy.
- **A new lakehouse table is missing from the endpoint** Metadata sync lag. Usually under a minute. Refresh or retry before dependent steps.

Practice

Exercise 1. Multiple choice (level 1)

Which statement about a lakehouse's SQL analytics endpoint is true?

- a) You can run UPDATE on lakehouse tables through it.
- b) You can create a view over lakehouse tables in it.
- c) You create it by hand after you create the lakehouse.
- d) It stores its own copy of the lakehouse tables.

Why?

- r1) Views are T-SQL objects that read data. The endpoint allows objects that read, not statements that change table data.
- r2) Views are stored inside the Delta tables.

Exercise 2. Architecture design (level 2)

Pick the Fabric item for each workload.

Options: Lakehouse, Warehouse, SQL analytics endpoint.

1. Land raw JSON survey deliveries and clean them with PySpark
2. Finance loads budgets with T-SQL stored procedures that change several tables together
3. Analysts query the cleaned lakehouse tables with T-SQL and save a few views

Exercise 3. Spot the bug (level 3)

This script runs against the SQL analytics endpoint of `lh_research`. It fails partway. Tap the failing line, then rewrite the script so it runs on the endpoint.

```
1 -- Connected to the SQL analytics endpoint of lh_research
2 CREATE VIEW dbo.v_wave_reach AS
3 SELECT wave_id, SUM(weighted_reach) / SUM(weighted_base) AS reach_pct
4 FROM dbo.nat_metrics
5 GROUP BY wave_id;
6 GO
7 UPDATE dbo.nat_metrics SET weighted_reach = 0 WHERE wave_id = 'W00';
```

Write down the line number of the bug. Then write the fix.

Exercise 4. Write the function (level 3)

On the SQL analytics endpoint, write a view `dbo.v_wave_reach` with one row per `wave_id` and a column `reach_pct`: the sum of `weighted_reach` divided by the sum of `weighted_base` from `dbo.nat_metrics`. A wave with a zero base must return NULL, not a divide-by-zero error.

Starter code:

```
1 CREATE VIEW dbo.v_wave_reach AS
```

Exercise 5. Fill the gap (level 4)

You are connected to warehouse `wh_finance`. Lakehouse `lh_research` sits in the same workspace. Complete the query so it reads the lakehouse table without copying it.

```
1 INSERT INTO dbo.wave_summary (wave_id, weighted_base)
2 SELECT wave_id, SUM(weighted_base)
3 FROM _____.(1).dbo.nat_metrics
4 GROUP BY wave_id;
```

Exercise 6. Multiple choice (level 5)

A pipeline runs a Spark notebook that creates a new gold table in `lh_research`. The next pipeline step immediately queries the table through the SQL analytics endpoint and fails with "Invalid object name". An hour later the same query works. What is the most likely cause?

- a) The notebook wrote the table in a format the endpoint cannot read.
- b) The endpoint had not yet synced the new table's metadata.
- c) The endpoint needs a manual refresh every time a table changes.
- d) The pipeline's identity lacks permission on new tables for one hour.

Why?

- r1) The endpoint discovers new Delta tables through an automatic sync that usually takes under a minute but has no fixed limit.
- r2) The endpoint copies data from the lakehouse on a schedule.

Practical: Lakehouse, endpoint and warehouse side by side

Platform: Fabric trial workspace. Run this in a Fabric trial workspace. You create a lakehouse and a warehouse, write one table with Spark, read it through the endpoint, and see a write fail there.

1. In your trial workspace, create a lakehouse named `lh_research`.
2. Open a new notebook attached to `lh_research`. Run:

```
spark.createDataFrame([("W01", 120.0, 400.0), ("W02", 0.0, 0.0)],  
"wave_id STRING, weighted_reach DOUBLE, weighted_base  
DOUBLE").write.mode("overwrite").saveAsTable("nat_metrics").
```
3. Switch the lakehouse view to its SQL analytics endpoint. Wait for `nat_metrics` to appear. Note how long it takes.
4. Open a new SQL query on the endpoint and create the `dbo.v_wave_reach` view from exercise 4. Query it. Wave W02 returns NULL.
5. In the same query window, run `UPDATE dbo.nat_metrics SET weighted_reach = 0 WHERE wave_id = 'W01'`; and read the error message.
6. Create a warehouse named `wh_finance`. In a new query, run `SELECT * FROM lh_research.dbo.nat_metrics;`. It returns the rows without a copy.
7. Write down what you saw in each step. Compare it with the recap card.

Sources

- The SQL analytics endpoint is read-only for Delta table data. No INSERT, UPDATE or DELETE. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/lakehouse-sql-analytics-endpoint> (updated 2026-05-19))
- You can create views, functions and stored procedures in the SQL analytics endpoint. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/lakehouse-sql-analytics-endpoint> (updated 2026-05-19))
- Fabric Warehouse supports explicit multi-table transactions with BEGIN TRAN, COMMIT TRAN and ROLLBACK TRAN. (verified: <https://learn.microsoft.com/en-us/fabric/data-warehouse/transactions> (updated 2026-06-03))
- Metadata sync runs in the background, and you can refresh it from the portal, the REST API or `sys.sp_dw_refresh_ext_table`. (verified: <https://>

learn.microsoft.com/en-us/fabric/data-engineering/sql-analytics-endpoint-metadata-sync (updated 2026-05-29))

- A warehouse reads lakehouse tables in the same workspace with three-part names, and transactions can span them. (verified: <https://learn.microsoft.com/en-us/fabric/data-warehouse/query-warehouse> and .../transactions)
- Under normal conditions the lag between a lakehouse and its SQL analytics endpoint is under one minute; the background sync stops after 15 minutes without query activity. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/sql-analytics-endpoint-performance> (updated 2026-08-07))

C3 Unifying datasets

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Shortcuts, mirroring (including Azure Databricks Unity Catalog) and avoiding duplicate copies.

C3-L1 OneLake shortcuts, and the delete that goes through

Concept

A **shortcut** is a folder in a lakehouse (or KQL database) that points to data somewhere else. It behaves like a symbolic link. The target can be inside OneLake, such as a table in another workspace, or outside it: ADLS Gen2, Azure Blob Storage, Amazon S3 and S3-compatible storage, Google Cloud Storage, Dataverse, OneDrive or SharePoint.

Shortcuts let every Fabric engine read the data **without copying it**. OneLake handles the connection and credentials, so each workload does not need its own.

Where you create the shortcut matters:

- Under **Tables**, point it at Delta data. The lakehouse recognises the folder as a table, and the SQL analytics endpoint and Direct Lake can use it.
- Under **Files**, it can point at data in any format.

Deletes are the trap. Microsoft Learn states the rule plainly:

- Deleting the **shortcut itself** deletes only the shortcut. The target data stays.
- Deleting a **file or folder inside** the shortcut deletes it **in the target**, if your identity has write permission there.

So a clean-up notebook that runs `rm -r` on `Files/partner_feed/2025/` removes the partner's 2025 files from their storage account, not a local copy.

Other points from Learn (checked September 2026):

- Moving, renaming or deleting the target breaks the shortcut. Broken shortcuts appear in the Unidentified area.
- For external shortcuts you can turn on **caching**. OneLake keeps files it has read for 1 to 28 days, which cuts cross-cloud egress costs.

- With Direct Lake over SQL or T-SQL in delegated identity mode, the calling user's identity is not passed through to the target. Check which identity reaches the source before relying on its permissions.

Why it matters

A lakehouse held a shortcut `Files/finance_share` to a folder in the finance team's storage account. An analyst ran a notebook to free up space: it deleted every subfolder older than two years under `Files/`. Their identity had write access to the finance account for another project. Seven years of finance exports disappeared from the source. The storage account had no soft delete, and the files were gone.

Worked example: weak code

```

1 for path in notebookutils.fs.ls("Files/"):
2     if path.name.startswith("2023"):
3         notebookutils.fs.rm(path.path, True)

```

1. **Line 1.** Walks every folder under Files, shortcuts included.
2. **Line 3.** Recursive delete inside a shortcut deletes the target data, if the identity can write there.

Worked example: strong code

```

1 OWNED_PREFIXES = ("Files/landing/", "Files/tmp/")
2
3
4 def safe_to_delete(path: str, shortcut_roots: set[str]) -> bool:
5     """True only for paths this lakehouse owns, never under a shortcut."""
6     if any(path == root or path.startswith(root.rstrip("/") + "/") for
7           root in shortcut_roots):
8         return False
9     return path.startswith(OWNED_PREFIXES)

```

1. **Line 1.** An allowlist of folders this lakehouse owns.
2. **Line 6.** Anything under a shortcut is refused, whatever its name.

Shortcuts save copies. Treat every path under a shortcut as someone else's data.

Recap card

- **What a shortcut is** A folder that points to data in OneLake or external storage, without copying it.
- **Delete the shortcut vs delete inside it** Deleting the shortcut keeps the data. Deleting files inside it deletes them in the target, if you can write there.

- **Shortcut as a table** Create it under Tables, pointing at Delta data.
- **External shortcut caching** OneLake caches read files for 1 to 28 days to cut egress.

Practice

Exercise 1. Multiple choice (level 1)

`Files/partner_feed` is a shortcut to an ADLS Gen2 folder. You delete the shortcut `Files/partner_feed` itself. What happens to the data in ADLS?

- a) It is deleted too.
- b) It stays. Deleting a shortcut deletes only the shortcut object.
- c) It is moved to the lakehouse recycle bin.
- d) It is copied into OneLake before the shortcut goes.

Why?

- r1) Shortcuts do not cascade deletes. The shortcut is an independent object.
- r2) ADLS data can never be deleted from Fabric.

Exercise 2. Multiple choice (level 2)

Your identity has write access to the target storage. You delete `Files/partner_feed/2023/` where `partner_feed` is a shortcut. What happens?

- a) Only the lakehouse view of 2023 is hidden.
- b) The 2023 folder is deleted in the target storage account.
- c) The delete fails: shortcuts are read-only.
- d) The shortcut breaks, but the data stays.

Why?

- r1) Deleting a file or folder within a shortcut deletes it in the target when you have permission there.
- r2) Fabric syncs deletes back every night.

Exercise 3. Fill the gap (level 2)

Complete the rule for where to create a shortcut so the SQL analytics endpoint sees it as a table.

1 Create the shortcut under ____ (1), pointing at data in ____ (2) format.

Exercise 4. Write the function (level 3)

Write `safe_to_delete(path: str, shortcut_roots: set[str]) -> bool` for a clean-up job. Return `False` if `path` is a shortcut root or lies under one. Otherwise return `True` only if it starts with `"Files/landing/"` or `"Files/tmp/"`. Treat `"Files/partner"` and `"Files/partner_old/x"` as different folders.

Starter code:

```
1 def safe_to_delete(path, shortcut_roots):
2     return path.startswith("Files/")
```

Exercise 5. Architecture design (level 4)

Marketing's campaign data lives in an S3 bucket, updated daily. Your lakehouse needs it for Spark and SQL. Choose the set-up.

Options: An S3 shortcut, A nightly copy pipeline into Files.

1. How the data reaches the lakehouse
2. Where the shortcut sits, given the data is Delta
3. Cross-cloud reads
4. Identity used by the connection

Sources

- Shortcut targets, Tables vs Files behaviour, no cascading deletes, deletes inside a shortcut reach the target with permission, caching 1 to 28 days, delegated identity not passed through. (verified: <https://learn.microsoft.com/en-us/fabric/onelake/onelake-shortcuts> (updated 2026-07-13))
- `notebookutils.fs.ls(path)` lists a directory and `notebookutils.fs.rm(path, recurse)` deletes; `MSSparkUtils` was renamed `NotebookUtils`. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/notebookutils/notebookutils-file-system> and [../notebook-utilities](#))

C3-L2 Mirroring operational databases into OneLake

Concept

Mirroring brings an external database into Fabric as a unit. For sources that store data in their own formats, such as Azure SQL Database, Azure SQL Managed Instance, Azure Cosmos DB, Snowflake or Google BigQuery, Fabric **continuously replicates** the chosen tables into OneLake as Delta tables, in

near real time. You get a mirrored database item and a **SQL analytics endpoint** over it, read-only, like a lakehouse's.

Shortcut or mirroring? Microsoft Learn's rule:

- **Shortcut** when you want specific data, already in an open format, in your OneLake view without copying it.
- **Mirroring** when you want an external database or catalog, and the data behind it, in Fabric as a unit. It replicates when the source uses a proprietary format.

Cost. For database mirroring and open mirroring, the background compute that replicates data does not consume capacity. Mirrored storage is free up to **1 TB per capacity unit**: an F64 includes 64 TB of mirroring storage. Queries on the mirrored data (Spark, SQL, Power BI) use capacity as usual.

Limits come per source. For Azure SQL Database, checked September 2026:

- at most 1,000 mirrored tables per database,
- tables whose primary key uses `datetime2(7)`, `datetimeoffset(7)` or `time(7)` cannot be mirrored,
- since April 2025 tables without a primary key can be mirrored, but existing ones are not added automatically, even with "Automatically mirror future tables" turned on.

Status varies too. Most sources are generally available. Some, such as Azure Database for MySQL, Dremio and SharePoint List, were in preview in September 2026. Check the mirroring overview page before you build on a source.

After you set up a mirror, compare row counts per table between source and mirror (lesson A6-L3). The limits above fail quietly: an ineligible table is simply not there.

Why it matters

A team mirrored an Azure SQL order database and built a dashboard on it. Weeks later, finance noticed that returns were missing. The returns table's primary key included a `datetime2(7)` column, so it had never been mirrored. Nothing errored. A per-table row-count check between source and mirror on the first day would have listed it as absent.

Worked example: weak code

```
❶ -- Mirror set up for all tables. Dashboard built on the mirror.
```



```
2 SELECT COUNT(*) FROM mirror_orders.dbo.orders;
```

1. **Line 1.** Assumes every table was eligible and replicated.

Worked example: strong code

```
1 -- Run in the source database: tables and row counts
1 SELECT s.name AS schema_name, t.name AS table_name, SUM(p.rows) AS
   source_rows
3 FROM sys.tables AS t
4 JOIN sys.schemas AS s ON s.schema_id = t.schema_id
5 JOIN sys.partitions AS p ON p.object_id = t.object_id AND p.index_id IN
   (0, 1)
6 GROUP BY s.name, t.name;
7
8 -- Run on the mirror's SQL analytics endpoint: which tables arrived
9 SELECT TABLE_SCHEMA, TABLE_NAME
10 FROM INFORMATION_SCHEMA.TABLES
11 WHERE TABLE_TYPE = 'BASE TABLE';
```

1. **Line 2.** The expected list and sizes, from the source.

2. **Line 10.** The actual list in the mirror. Any table missing here was not mirrored.

Mirroring replicates operational databases for free compute and generous storage, within per-source limits. Reconcile the table list after setup.

Recap card

- **Shortcut or mirroring?** Shortcut for open-format data in storage. Mirroring for databases and catalogs, replicated when proprietary.
- **Mirroring cost** Replication compute is free. Storage is free up to 1 TB per CU. Queries use capacity.
- **Azure SQL Database mirroring limits** Up to 1,000 tables. Primary keys on datetime2(7), datetimeoffset(7) or time(7) block a table.
- **After setting up a mirror** Reconcile table lists and row counts with the source.

Practice

Exercise 1. Multiple choice (level 1)

You need an Azure SQL Database's tables in Fabric, kept up to date, for reporting. Shortcut or mirroring?

- a) Shortcut: shortcuts connect to Azure SQL Database.
- b) Mirroring: it replicates the database into OneLake continuously.
- c) Neither: export CSV files nightly.
- d) A shortcut under Tables.

Why?

- r1) Azure SQL stores data in a proprietary format, so Fabric must replicate it, which is what mirroring does.
- r2) Mirroring is always cheaper than shortcuts.

Exercise 2. Predict the output (level 2)

Using the documented rule of 1 TB of free mirroring storage per capacity unit, what does this print?

```
1 capacity_units = 32
2 mirrored_tb = 45
3 free_tb = capacity_units * 1
4 print(free_tb, max(0, mirrored_tb - free_tb))
```

Exercise 3. Fill the gap (level 2)

List the base tables that arrived in the mirror, on its SQL analytics endpoint.

```
1 SELECT TABLE_SCHEMA, TABLE_NAME
2 FROM _____.TABLES
3 WHERE TABLE_TYPE = '_____'
```

Exercise 4. Write the function (level 3)

Write

`missing_from_mirror(source: dict[str, int], mirror: dict[str, int]) -> list[str]` for a reconciliation job. Both map `schema.table` to a row count. Return, sorted, the tables that are absent from the mirror, or whose mirror count differs from the source count.

Starter code:

```
1 def missing_from_mirror(source, mirror):
2     return [t for t in source if t not in mirror]
```

Exercise 5. Architecture design (level 4)

Plan mirroring for an Azure SQL order database with 1,400 tables. Some primary keys use `datetime2(7)`.

Options: The tables reporting needs, under the 1,000-table limit, All 1,400.

1. Which tables to mirror
2. Tables with `datetime2(7)` primary keys
3. After setup

Sources

- Shortcuts vs mirroring guidance, and mirroring replicates when the source stores data in a proprietary format. (verified: <https://learn.microsoft.com/en-us/fabric/onelake/unify-data> (updated 2026-07-15))
- Replication compute is free and does not consume capacity; mirroring storage is free at 1 TB per CU; MySQL, Dremio and SharePoint List sources are in preview; BigQuery is GA. (verified: <https://learn.microsoft.com/en-us/fabric/mirroring/overview> (updated 2026-08-28))
- Azure SQL Database mirroring: 1,000 tables maximum, datetime2(7)/datetimeoffset(7)/time(7) primary keys unsupported, tables without primary keys supported since April 2025 but not added automatically. (verified: <https://learn.microsoft.com/en-us/fabric/mirroring/azure-sql-database-limitations> (updated 2026-02-26))

C3-L3 Mirroring Azure Databricks Unity Catalog into Fabric

Concept

Your Track A pipelines write gold tables into Unity Catalog. **Mirroring Azure Databricks Unity Catalog** makes them readable in Fabric without copying them. Unlike database mirroring, **no data moves**. Fabric mirrors only the catalog structure (catalogs, schemas, tables) and reads the underlying Delta files through shortcuts.

You get a mirrored Azure Databricks item with a SQL analytics endpoint, and you can build Power BI reports on it in Direct Lake mode. By default, **Automatically sync future catalog changes for the selected schema** is on, so new tables appear. Propagation takes from seconds to several minutes.

Three rules decide whether it works and whether it is safe (Microsoft Learn, checked September 2026):

1. Some table types do not mirror: tables with row filters or column masks (RLS/CLM policies), Lakehouse federated tables, Delta Sharing tables, streaming tables, views and materialized views. The ADLS Gen2 account behind the Databricks workspace must also be reachable from Fabric.

2. Unity Catalog permissions do not carry over. Grants on catalogs, schemas and tables in Databricks do not apply in Fabric. You rebuild access

with Fabric's model: workspace roles, item permissions, SQL permissions, OneLake security.

3. One credential reads everything. The credential used to create the connection is used for all data queries. Unity Catalog policies such as row filters, column masks or ABAC are not enforced at the storage layer, so they would not apply to direct storage access anyway.

Put together: mirror tables that are safe to share at the level you grant in Fabric. For data that needs row-level rules, publish a gold table without the policy that already contains only what the audience may see, or apply the rules again in Fabric.

Why it matters

A team mirrored their gold schema so that regional managers could build Power BI reports. In Databricks, a row filter limited each manager to their region (lesson A4-L3). The filtered table did not appear in Fabric at all, so the team created an unfiltered copy for the mirror, with a connection that could read every row. Every manager with Viewer access now saw every region. The rules had lived in Unity Catalog, and nothing had been rebuilt in Fabric.

Worked example: weak code

- ❶ Mirror: catalog gold_prod, schema sales, all tables, auto-sync on
- ❷ Connection: service principal with SELECT on the whole catalog
- 3 Fabric workspace: regional managers as Viewers
- ❸ Assumption: Unity Catalog row filters still apply

1. **Line 1.** Tables with row filters are not mirrored, and any new table appears automatically.
2. **Line 2.** This credential reads every mirrored table for every user.
3. **Line 4.** Unity Catalog permissions and policies do not carry over to Fabric.

Worked example: strong code

- ❶ Mirror: catalog gold_prod, schema sales_shareable (tables approved for all managers)
- ❷ Connection: service principal with SELECT on sales_shareable only
- ❸ Regional detail: rebuilt in Fabric with row-level security on the semantic model
- 4 Check: row counts per mirrored table match Databricks after each release

1. **Line 1.** A schema of tables that are safe for the whole Fabric audience.
2. **Line 2.** Least privilege on the one credential that reads everything.
3. **Line 3.** Access rules recreated where the users are.

Mirroring Unity Catalog copies no data and no permissions. Decide what is shareable, scope the connection, and rebuild access rules in Fabric.

Recap card

- **What mirroring Unity Catalog moves** Only catalog metadata. Data is read in place through shortcuts.
- **Unity Catalog tables that do not mirror** Tables with row filters or column masks, federated, Delta Sharing and streaming tables, views and materialized views.
- **Unity Catalog permissions in Fabric** Not carried over. Rebuild access with Fabric's permission model.
- **Which identity reads mirrored data?** The connection's credential, for all queries. Scope it tightly.

Practice

Exercise 1. Multiple choice (level 1)

When Fabric mirrors an Azure Databricks Unity Catalog, what moves into OneLake?

- a) A full copy of every table's data
- b) Only the catalog structure. Data is read through shortcuts.
- c) Nightly snapshots of changed partitions
- d) Nothing. Fabric queries Databricks SQL warehouses.

Why?

- r1) Unity Catalog tables are already open Delta files, so Fabric mirrors metadata and reads the files in place.
- r2) Copying is blocked by Databricks licensing.

Exercise 2. Multiple choice (level 2)

A gold table in Unity Catalog has a row filter. You select its schema for mirroring. What happens to that table?

- a) It is mirrored with the row filter enforced in Fabric.
- b) It is mirrored without the filter.
- c) It is not mirrored: tables with RLS or CLM policies are unsupported.
- d) It is mirrored as a view.

Why?

- r1) The limitations page lists tables with RLS/CLM policies as unsupported.

- r2) Row filters make tables read-only.

Exercise 3. Put in order (level 2)

Order the steps to publish Databricks gold tables to Fabric users safely.

- Create a connection identity with SELECT on that schema only
- Compare table lists and row counts with Databricks
- Create the mirrored Azure Databricks item for that schema
- Choose tables safe for the whole Fabric audience, in a dedicated schema
- Grant Fabric access with workspace roles, SQL permissions or model RLS

Exercise 4. Write the function (level 3)

Write `mirror_eligible(tables: list[dict[str, object]]) -> list[str]`. Each dict has `name` (str), `table_type` (str) and `has_row_filter_or_mask` (bool). Return, sorted, the names that can be mirrored: `table_type` must be `"MANAGED"` or `"EXTERNAL"`, and the table must not have a row filter or column mask. Streaming tables, views, materialized views, federated and Delta Sharing tables arrive with other `table_type` values.

Starter code:

```
1 def mirror_eligible(tables):
2     return [t["name"] for t in tables]
```

Exercise 5. Architecture design (level 4)

Regional managers must each see only their region's sales in Power BI. The data is a Unity Catalog gold table with a row filter. Choose the design.

Options: An unfiltered copy of the table, for all managers, An unfiltered table, with row-level security rebuilt in the Fabric semantic model.

1. What Fabric mirrors
2. Connection identity
3. Manager access to the mirror's SQL endpoint

Sources

- Only the Azure Databricks catalog structure is mirrored and data is accessed through shortcuts, with no data movement; auto-sync of future catalog changes is on by default; propagation takes seconds to several minutes. (verified: <https://learn.microsoft.com/en-us/fabric/mirroring/azure-databricks> (updated 2025-10-24))
- Unsupported table types: RLS/CLM tables, Lakehouse federated tables, Delta Sharing tables, streaming tables, views and materialized views; the

Databricks ADLS account must be accessible to Fabric. (verified: <https://learn.microsoft.com/en-us/fabric/mirroring/azure-databricks-limitations> (updated 2025-12-19))

- Unity Catalog permissions are not mirrored; the connection credential is used for all data queries; UC policies are not enforced at the storage layer. (verified: <https://learn.microsoft.com/en-us/fabric/mirroring/azure-databricks-security> (updated 2026-05-01))
- Unity Catalog `information_schema.tables` reports `table_type` values including `MANAGED`, `EXTERNAL`, `VIEW`, `MATERIALIZED_VIEW`, `STREAMING_TABLE` and `FOREIGN`. (verified: <https://docs.databricks.com/aws/en/sql/language-manual/information-schema/tables>)

C4 Ingestion and orchestration

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Data Factory pipelines, Dataflows Gen2 and copy jobs.

C4-L1 Choosing how data gets into Fabric

Concept

Fabric offers several ways to move data. They overlap, and picking by habit creates pipelines nobody needs. Microsoft Learn's data movement guide (updated April 2026) gives a starting rule by layer:

- **Curated data from an existing system, for reporting: Mirroring.** If the source already does its own ETL and you need the result in Fabric, mirroring is the simplest and cheapest choice (lesson C3-L2).
- **Raw ingestion into bronze: Copy job.** It handles bulk copy, incremental copy and change data capture (CDC), with scheduling and state tracking built in, without building a pipeline.
- **Low-latency events: Eventstream** (lesson C7-L1).

Then add tools only where they earn their place:

- **Pipeline** (Data Factory): orchestration. Run activities in order, branch on success or failure, loop, retry, and call notebooks, stored procedures, dataflows and Copy jobs. A pipeline with a copy activity also moves data, but you manage incremental state yourself.
- **Dataflow Gen2:** Power Query, low-code transformation for analysts, with a data destination per query.
- **Spark notebook:** code-first transformation at scale, the closest match to your Databricks work (lesson C5-L1).

A simple way to decide:

1. Does the source need to be replicated as a whole database? Mirroring.
2. Is it a copy of tables or files into bronze? Copy job.
3. Are there events that must arrive within seconds? Eventstream.
4. Is there transformation? Spark for engineers, Dataflow Gen2 for low-code authors.
5. Do several steps need ordering, branching and alerts? Wrap them in a pipeline.

Two warnings. A pipeline that only copies tables re-creates what a Copy job does, with more to maintain. And Dataflow Gen2 is not a replacement for Spark on large volumes. Check your capacity usage (lesson C10-L3) before you standardise on one tool.

Why it matters

A team built 40 pipelines, each with a copy activity, a lookup of the last loaded timestamp from a control table, and a stored procedure to update it. Two of them had an off-by-one error in the watermark and silently skipped rows every night. A Copy job per source would have tracked incremental state itself, with nothing for the team to get wrong.

Worked example: weak code

- ❶ Source: Azure SQL CRM (200 tables) -> 200 pipelines, copy activity + control table
- ❷ Source: Clickstream events -> hourly copy activity from storage
- ❸ Transform: Dataflow Gen2 on 400 M rows per day

1. **Line 1.** Hand-built incremental state for a whole database: mirroring or a Copy job does this for you.
2. **Line 2.** Events that need low latency, handled as hourly batches.
3. **Line 3.** A low-code tool on engineering-scale volume. Spark suits this better.

Worked example: strong code

- ❶ Source: Azure SQL CRM (200 tables) -> Mirroring (curated reporting copy)
- ❷ Source: Supplier files in ADLS -> Copy job, incremental, into bronze
- 3 Source: Clickstream events -> Eventstream into an eventhouse
- 4 Transform: Spark notebooks, bronze to gold
- ❸ Orchestration: one pipeline runs the notebooks in order, retries, alerts on failure

1. **Line 1.** Free replication compute, no incremental logic to maintain.
2. **Line 2.** State tracking built in.
3. **Line 5.** The pipeline orchestrates. It does not re-create copying.

Mirroring for curated databases, Copy job for bronze, Eventstream for events, Spark or Dataflow Gen2 for transformation, and pipelines to orchestrate.

Recap card

- **Microsoft's starting rule by layer** Gold from existing systems:
Mirroring. Bronze ingestion: Copy job. Low-latency events: Eventstream.

- **What a pipeline is for** Orchestration: ordering, branching, loops, retries and alerts across activities.
- **Spark or Dataflow Gen2?** Spark for code-first, large-volume engineering. Dataflow Gen2 for low-code Power Query authors.
- **Copy activity vs Copy job for incremental loads** Copy activity needs hand-built state. Copy job tracks the last successful run itself.

Practice

Exercise 1. Multiple choice (level 1)

You need raw CSV deliveries from ADLS copied into a bronze lakehouse every hour, only new files each time. Which tool does Microsoft Learn recommend starting with?

- a) Mirroring
- b) Copy job
- c) Dataflow Gen2
- d) A pipeline with a copy activity and a control table

Why?

- r1) Copy job supports incremental copy and scheduling natively for bronze ingestion.
- r2) Copy job is the only tool that reads ADLS.

Exercise 2. Architecture design (level 2)

Pick the tool for each need.

Options: Mirroring, Copy job, Eventstream, Spark notebook, Pipeline.

1. An ERP database's curated tables, kept current for reporting
2. Sensor events that must trigger alerts within seconds
3. Joining and cleaning 300 M rows a day into silver
4. Run copy, then three notebooks in order, and email on failure

Exercise 3. Put in order (level 2)

Order the questions in the decision sequence from this lesson.

- Transformation needed, and by whom?
- Several steps that need ordering and alerts?
- Copy tables or files into bronze?
- Replicate a whole database or catalog?
- Events that must arrive within seconds?

Exercise 4. Write the function (level 3)

Write `recommend(need: str) -> str` encoding the decision rule. Accepted needs and answers:

- `"replicate_database"` → `"Mirroring"`
- `"bronze_copy"` → `"Copy job"`
- `"low_latency_events"` → `"Eventstream"`
- `"code_transform"` → `"Spark notebook"`
- `"lowcode_transform"` → `"Dataflow Gen2"`
- `"orchestrate"` → `"Pipeline"`

Raise `ValueError` for anything else, with the unknown need in the message.

Starter code:

```
1 def recommend(need):  
2     return "Pipeline"
```

Exercise 5. Multiple choice (level 4)

A team copies 60 tables from Azure SQL into bronze with 60 pipelines, each using a copy activity, a control table for the last watermark, and a stored procedure to update it. Two tables lose rows every night. What change removes the whole class of bug?

- a) Add a retry to each copy activity.
- b) Move the tables to Copy jobs with incremental copy, which track state themselves.
- c) Run the pipelines twice a night.
- d) Replace the stored procedure with a Dataflow Gen2.

Why?

- r1) Copy job manages the state of the last successful run, so there is no hand-written watermark logic to get wrong.
- r2) Pipelines cannot copy incrementally.

Sources

- Data movement guide: Mirroring for gold from existing systems, Copy job for bronze ingestion, Eventstream for low-latency events; copy activity requires managing incremental state yourself. (verified: <https://learn.microsoft.com/en-us/fabric/data-factory/decision-guide-data-movement> (updated 2026-04-13))

- Copy job supports full copy, watermark-based and CDC-based incremental copy, and tracks the state of the last successful run. (verified: <https://learn.microsoft.com/en-us/fabric/data-factory/what-is-copy-job> (updated 2026-09-17))

C4-L2 Incremental copy with a watermark, or with CDC

Concept

A **Copy job** copies data in one of two modes. **Full copy** moves everything on every run. **Incremental copy** moves everything on the first run, then only new or changed data. The job tracks the state of the last successful run itself (Microsoft Learn, checked September 2026).

Incremental copy works in one of two ways:

- **Watermark.** You pick an incremental column, such as a `ROWVERSION`, a datetime, a date or an increasing integer. Each subsequent load copies rows whose value is **larger than any value copied before**.
- **CDC** (preview in Copy job). If change data capture is enabled on the source database, the job reads its change stream and copies every insert, update and delete since the last run. You do not choose a column.

The watermark method has two blind spots, and both follow from "larger than":

- It **cannot see deletes**. A deleted row simply stops appearing, and the destination keeps it.
- It **misses late rows**. A row committed after the run, with a watermark value at or below the one already copied, is never picked up. A `modified_at` set by the application clock, or several rows written in the same second, causes this.

Prefer CDC when you need deletes, or a `ROWVERSION` column, which the database increments on every change. Microsoft also lets you reset incremental copy to a full copy, for the whole job or per table, to recover.

Choose the **update method** for the destination on purpose. The default is **Append**, which keeps full history, and appends every changed row again as a new row. **Merge** updates existing rows by key (the primary key by default), **Overwrite** replaces the data, and **SCD Type 2** keeps versioned history. For a silver table that must hold one row per key, choose Merge.

Why it matters

An order table was copied incrementally on `modified_at`, set by the ordering app. At 23:59:59 one batch committed after the nightly run had read up to 23:59:59. Those rows carried the same timestamp as the last copied row, so they were never picked up. Finance found 212 missing orders at month end. CDC, or a `ROWVERSION` column, would have caught them.

Worked example: weak code

```
1 def next_batch(rows, last_wm):
  ❶ batch = [r for r in rows if r["modified_at"] >= last_wm]
  ❷ return batch, max(r["modified_at"] for r in batch)
```

1. Line 2.

= re-copies the last run's boundary rows every time, which Append turns into duplicates.

2. Line 3. Fails with ValueError on a run with no new rows.

Worked example: strong code

```
1 Row = dict[str, int]
2
3
4 def next_batch(rows: list[Row], last_wm: int | None) -> tuple[list[Row],
  int | None]:
5     """Rows with a watermark above last_wm, and the new watermark."""
  ❶ batch = [r for r in rows if last_wm is None or r["rv"] > last_wm]
7     if not batch:
  ❷ return [], last_wm
9     return batch, max(r["rv"] for r in batch)
```

1. Line 6. Strictly greater, as Copy job does, on a ROWVERSION-style column.

2. Line 8. An empty run keeps the old watermark instead of failing.

Watermarks copy rows above the last value seen: no deletes, no late rows. Use CDC or a `ROWVERSION` column, and Merge for one row per key.

Recap card

- **Copy job incremental copy** First run copies everything. Later runs copy new or changed data. State is tracked by the job.
- **Watermark blind spots** Cannot see deletes. Misses late rows at or below the stored watermark.
- **When to use CDC** When deletes must replicate, or changes arrive out of order.

- **Copy job update methods** Append (default), Merge, Overwrite, SCD Type 2. Merge for one row per key.

Practice

Exercise 1. Multiple choice (level 1)

A Copy job uses watermark-based incremental copy on `modified_at`. A row is deleted in the source. What happens in the destination?

- a) The row is deleted on the next run.
- b) Nothing. Watermark-based incremental copy cannot detect deletes.
- c) The row is marked with a soft-delete flag.
- d) The job fails until you reset it.

Why?

- r1) A watermark only finds rows with a larger value. A deleted row has no value to find.
- r2) Deletes are never supported by Copy job.

Exercise 2. Predict the output (level 2)

Run 1 copied up to watermark 105. A late row with `rv` 104 commits afterwards. What does run 2 copy?

```
1 rows = [{"id": 1, "rv": 101}, {"id": 2, "rv": 105}, {"id": 3, "rv": 104}, {"id": 4, "rv": 107}]
2 last_wm = 105
3 batch = [r["id"] for r in rows if r["rv"] > last_wm]
4 print(batch)
```

Exercise 3. Fill the gap (level 2)

Complete the Copy job settings for a silver table that must hold one current row per `order_id` and reflect source deletes.

```
1 Incremental method: ____ (1)
2 Update method: ____ (2)
```

Exercise 4. Write the function (level 3)

Write `next_batch(rows, last_wm)` for a hand-built incremental load that behaves like Copy job's watermark method. `rows` is a list of dicts with an integer `rv`. `last_wm` is `None` before the first run. Return `(batch, new_wm)`: rows with `rv` strictly greater than `last_wm` (all rows when `last_wm` is `None`), and the largest `rv` in the batch. On an empty batch return `([], last_wm)`.

Starter code:

```

1 def next_batch(rows, last_wm):
2     batch = [r for r in rows if r["rv"] >= last_wm]
3     return batch, max(r["rv"] for r in batch)

```

Exercise 5. Architecture design (level 4)

Configure incremental copy of an orders table into silver. Orders are updated and sometimes deleted. The source has CDC enabled.

Options: CDC, Watermark on modified_at.

1. Incremental method
2. Update method
3. After an outage caused missed changes

Sources

- Copy job modes, watermark column types, strictly larger watermark rule, CDC captures deletes, watermark cannot detect deletes, reset to full copy per job or table. (verified: <https://learn.microsoft.com/en-us/fabric/data-factory/incremental-copy-job> (updated 2026-09-15) and [.../what-is-copy-job](https://learn.microsoft.com/en-us/fabric/data-factory/what-is-copy-job) (updated 2026-09-17))
- Update methods Append (default), Merge (primary key by default), Overwrite and SCD Type 2. (verified: <https://learn.microsoft.com/en-us/fabric/data-factory/what-is-copy-job> (updated 2026-09-17))
- CDC replication in Copy job is in preview. (verified: <https://learn.microsoft.com/en-us/fabric/data-factory/what-is-copy-job> (updated 2026-09-17))

C4-L3 Pipelines that fail fast, retry sensibly and alert

Concept

A Data Factory **pipeline** runs activities: Copy job, notebook, stored procedure, Dataflow, If Condition, ForEach, Until, and more. By default a pipeline holds up to 120 activities. Three settings on each activity decide how a failure behaves (Microsoft Learn, checked September 2026):

- **Timeout.** How long an activity may run before it is stopped. The default is **12 hours**, and the maximum is seven days, written D.HH:MM:SS. A notebook that normally takes 10 minutes and hangs will hold its slot, and its capacity, for 12 hours unless you lower this.
- **Retry.** Retries happen only if you enable them. The retry count then defaults to 1, with **30 seconds** between attempts. Retry steps that fail for

passing reasons, such as a busy source or a network drop. Do not retry a step that failed because its data is wrong. It will fail again, later.

- **Dependency branches.** Each activity connects to the next through a condition: **Upon Success**, **Upon Failure**, **Upon Completion** (either way) or **Upon Skip**. Put alerts and clean-up on Upon Failure. Upon Completion runs even when the step succeeded.

Parameters and expressions keep one pipeline reusable across sources and environments. Expressions start with `@:` `@pipeline().parameters.wave`. Inside a string, `@{...}` interpolates: `"landing/@{pipeline().parameters.wave}/"`. The result of interpolation is always a **string**, so `"@{pipeline().parameters.max_rows}"` passes text, not a number. When a setting needs a number, use the bare expression `@pipeline().parameters.max_rows`.

Keep business logic out of the pipeline. Transformations belong in notebooks and SQL, where they can be tested (lesson A9-L1). The pipeline decides order, retries and alerts.

Why it matters

A nightly pipeline ran a 10-minute notebook that sometimes hung on a locked source. With the default 12-hour timeout and no failure branch, each hang held its capacity until mid-morning. Downstream reports showed yesterday's data, and nobody was told. A 30-minute timeout, one retry and an Upon Failure email would have raised the alarm at 02:30 instead of 10:00.

Worked example: weak code

```
1 Activity: nb_silver_orders (notebook)
❶ Timeout: default
3 Retries: not enabled
❷ Activity: nb_gold_sales depends on nb_silver_orders: Upon Completion
❸ Activity: send_failure_email depends on nb_gold_sales: Upon Completion
```

1. **Line 2.** 12 hours for a 10-minute job: a hang goes unnoticed until morning.
2. **Line 4.** Gold runs even when silver failed, publishing stale or partial data.
3. **Line 5.** The 'failure' email is sent every night, so people learn to ignore it.

Worked example: strong code

```
1 Activity: nb_silver_orders (notebook)
❶ Timeout: 0.00:30:00
3 Retries: enabled, 1 retry, 60 seconds apart
❷ Activity: nb_gold_sales depends on nb_silver_orders: Upon Success
❸ Activity: send_failure_email depends on nb_silver_orders: Upon Failure
```


1. **Line 2.** Three times the normal run time: a hang fails within half an hour.
2. **Line 4.** Gold runs only on good silver.
3. **Line 5.** Alerts fire only when something failed.

Set timeouts near the real run time, retry only transient failures, chain on Upon Success and alert on Upon Failure.

Recap card

- **Activity timeout default** 12 hours, maximum seven days, format D.HH:MM:SS. Set it near the real run time.
- **Activity retries** Only when enabled. Count defaults to 1, interval 30 seconds. Retry transient failures only.
- **Dependency branches** Upon Success, Upon Failure, Upon Completion (either way), Upon Skip.
- **@{...} vs @...** @{...} inside text always yields a string. A bare @ expression keeps the value's type.

Practice

Exercise 1. Multiple choice (level 1)

You leave an activity's timeout at its default. How long can it run before Fabric stops it?

- a) 30 minutes
- b) 1 hour
- c) 12 hours
- d) 7 days

Why?

- r1) The documented default timeout is 12 hours. Seven days is the maximum you can set.
- r2) Timeouts match the capacity's pause schedule.

Exercise 2. Multiple choice (level 2)

An alert activity is attached to a notebook with **Upon Completion**. When does it run?

- a) Only when the notebook fails
- b) Only when the notebook succeeds
- c) Whether the notebook succeeds or fails

- d) Only when the notebook is skipped

Why?

- r1) Upon Completion runs after success or failure. Upon Failure is the failure-only branch.
- r2) Completion means the capacity finished all jobs.

Exercise 3. Fill the gap (level 2)

Build a folder path from the pipeline parameter `wave`, and pass the numeric parameter `max_rows` as a number.

```
1 Folder path: landing/____(1)pipeline().parameters.wave}/
2 Max rows: ____ (2)pipeline().parameters.max_rows
```

Exercise 4. Write the function (level 3)

Write `parse_timeout(value: str) -> int` for a pipeline linter. It parses the D.HH:MM:SS format and returns the total seconds. Raise `ValueError` if the format is wrong, if hours are above 23 or minutes or seconds above 59, or if the total exceeds 7 days.

Starter code:

```
1 def parse_timeout(value):
2     days, rest = value.split(".")
3     h, m, s = rest.split(":")
4     return int(days) * 86400 + int(h) * 3600 + int(m) * 60 + int(s)
```

Exercise 5. Spot the bug (level 4)

This pipeline publishes gold tables even when silver fails. Tap the line that causes it, then rewrite the dependencies so gold runs only on success and an alert fires on failure.

```
1 Activity: nb_silver_orders
2 Timeout: 0.00:30:00
3 Activity: nb_gold_sales depends on nb_silver_orders: Upon Completion
4 Activity: email_on_failure depends on nb_silver_orders: Upon Failure
```

Write down the line number of the bug. Then write the fix.

Sources

- Activity timeout defaults to 12 hours with a seven-day maximum in D.HH:MM:SS; retries when enabled default to 1 with a 30-second interval; up to 120 activities per pipeline; UponSuccess, UponFailure, UponCompletion and UponSkip branches. (verified: <https://learn.microsoft.com/en-us/fabric/data-factory/activity-overview> (updated 2026-08-10))

- Pipeline expressions use `@pipeline().parameters`; string interpolation with `@{...}` always returns a string. (verified: <https://learn.microsoft.com/en-us/fabric/data-factory/expression-language> (updated 2024-12-18))

C5 Spark notebooks in Fabric

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Differences from Databricks, environments, V-Order and Delta interoperability.

C5-L1 From Databricks notebooks to Fabric notebooks

Concept

Fabric notebooks run the same PySpark you wrote in Track A. The surroundings differ, and those differences break ported code (Microsoft Learn, checked September 2026).

Utilities. There is no `dbutils`. Fabric has **NotebookUtils**, formerly `MSSparkUtils`. The `mssparkutils` namespace still works but is being retired, so write `notebookutils`:

Databricks	Fabric
<code>dbutils.fs.ls(path)</code>	<code>notebookutils.fs.ls(path)</code>
<code>dbutils.secrets.get(scope, key)</code>	<code>notebookutils.credentials.getSecret("https://<vault>.vault.azure.net/", "name")</code>
<code>dbutils.notebook.run(path, timeout, args)</code>	<code>notebookutils.notebook.run(path, timeout_seconds, arguments)</code>
<code>dbutils.notebook.exit(value)</code>	<code>notebookutils.notebook.exit(value)</code>

Secrets live in **Azure Key Vault**, not in Databricks secret scopes. You need Get permission on the vault. Output that prints a secret is redacted, but do not print secrets.

Child notebooks time out at 90 seconds by default.

`notebookutils.notebook.run(path, timeout_seconds=90, ...)`. A child notebook that runs for five minutes fails unless you pass a longer timeout. For several notebooks, `notebookutils.notebook.runMultiple` runs a DAG of them in parallel.

Default lakehouse. A notebook can attach a default lakehouse. Relative paths such as `Files/landing` and bare table names resolve against it. Name the lakehouse and schema in shared code (lesson C2-L1), so a changed attachment does not change what the code reads.

Inline installs. `%pip install` affects only the current session, and inline commands are **turned off by default in pipeline runs**. Put libraries in an environment instead (lesson C5-L2).

Runtime. Runtime 2.0 (Spark 4.1, Delta Lake 4.2, Python 3.13, Java 21) is GA and recommended for production. New workspaces still default to Runtime 1.3 (Spark 3.5), whose end of support has been announced. Set the runtime deliberately. It changes the Spark version your code meets.

Why it matters

A team ported a Databricks orchestration notebook that called six child notebooks with `mssparkutils.notebook.run(path)`. In development each child took under a minute. On month-end data, the largest took four minutes. Every run failed at 90 seconds with a timeout, and the error pointed at the parent notebook. Passing `timeout_seconds` sized to the real run time fixed it.

Worked example: weak code

```
❶ password = dbutils.secrets.get("prod", "sql-password")
❷ %pip install great_expectations
❸ result = mssparkutils.notebook.run("nb_silver_orders")
❹ df = spark.read.table("orders")
```

1. **Line 1.** `dbutils` does not exist in Fabric. Secrets come from Azure Key Vault.
2. **Line 2.** Inline installs are turned off by default in pipeline runs, and give unpinned dependencies.
3. **Line 3.** Old namespace, and a 90-second default timeout for the child.
4. **Line 4.** Depends on whichever lakehouse is attached as default, and the `dbo` schema.

Worked example: strong code

```
1 VAULT_URL = "https://kv-sales-prod.vault.azure.net/"
2
❶ password = notebookutils.credentials.getSecret(VAULT_URL, "sql-password")
4 result = notebookutils.notebook.run(
❷     "nb_silver_orders", timeout_seconds=1800, arguments={"wave":
        "2026-09"}
6 )
❸ df = spark.read.table("lh_sales.silver.orders")
```

1. **Line 3.** Key Vault through `NotebookUtils`. The vault URL belongs in config.
2. **Line 5.** A timeout sized to the child's real run time, and explicit arguments.
3. **Line 7.** Lakehouse and schema named, independent of the default attachment.

Replace `dbutils` with `notebookutils`, secrets with Key Vault, inline installs with environments, and pass real timeouts to child notebooks.

Recap card

- **Fabric replacement for dbutils** notebookutils (formerly mssparkutils).
- **Secrets in Fabric notebooks**
notebookutils.credentials.getSecret(vault_url, name), from Azure Key Vault.
- **notebookutils.notebook.run default timeout** 90 seconds. Pass timeout_seconds.
- **%pip in pipeline runs** Turned off by default. Use an environment.

Practice

Exercise 1. Multiple choice (level 1)

Which call reads a secret in a Fabric notebook?

- a) `dbutils.secrets.get("prod", "sql-password")`
- b) `notebookutils.credentials.getSecret("https://kv-sales.vault.azure.net/", "sql-password")`
- c) `os.environ["SQL_PASSWORD"]` set in the workspace settings
- d) `spark.conf.get("secret.sql-password")`

Why?

- r1) Fabric reads secrets from Azure Key Vault through notebookutils.credentials.
- r2) getSecret caches the secret in the notebook.

Exercise 2. Multiple choice (level 2)

`notebookutils.notebook.run("nb_big_load")` calls a child notebook that takes four minutes. What happens?

- a) It completes after four minutes.
- b) It fails after 90 seconds, the default timeout.
- c) It fails after 12 hours.
- d) It runs in the background and returns at once.

Why?

- r1) `run()` has `timeout_seconds` with a default of 90.
- r2) Child notebooks are limited to one Spark job.

Exercise 3. Fill the gap (level 2)

Port this Databricks line to Fabric.

```
1 files = _____.fs.ls("Files/landing")
2 value = _____.notebook.run("nb_child", _____.(2)=600)
```

Exercise 4. Write the function (level 3)

Write `port_line(line: str) -> str` for a migration helper. It rewrites Databricks utility calls to Fabric:

- `dbutils.fs.` → `notebookutils.fs.`
- `dbutils.notebook.` → `notebookutils.notebook.`
- `mssparkutils.` → `notebookutils.`

If the line contains `dbutils.secrets.`, raise `ValueError` with a message that mentions Key Vault, because it needs a manual change. Leave other text unchanged.

Starter code:

```
1 def port_line(line):
2     return line.replace("dbutils", "notebookutils")
```

Exercise 5. Spot the bug (level 4)

This notebook runs from a nightly pipeline. It worked when run by hand and fails in the pipeline. Tap the line that fails there, and rewrite the notebook without it.

```
1 %pip install pandera==0.33.1
2 import pandera.pandas as pa
3 df = spark.read.table("lh_sales.silver.orders").toPandas()
4 SCHEMA.validate(df, lazy=True)
```

Write down the line number of the bug. Then write the fix.

Sources

- MsSparkUtils is renamed NotebookUtils; the mssparkutils namespace will be retired. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/notebook-utilities>)
- `notebookutils.notebook.run(path, timeout_seconds=90, arguments=None, workspace=")` returns the exit value; `runMultiple` runs a DAG. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/notebookutils/notebookutils-notebook-run>)
- `notebookutils.credentials.getSecret` reads from Azure Key Vault and needs Get permission; output redaction is attempted. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/notebookutils/notebookutils-credentials> (updated 2026-09-23))

- Inline installs are turned off by default in pipeline runs. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/library-management> (updated 2026-03-25))
- Runtime 2.0 (Spark 4.1, Delta 4.2, Python 3.13, Java 21) is GA and recommended; new workspaces default to Runtime 1.3, which is at end-of-support-announced stage. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/runtime> (updated 2026-07-24))

C5-L2 Environments, resource profiles and V-Order

Concept

An **environment** is a Fabric item that fixes what a Spark session starts with: the runtime version, compute settings, libraries and small resource files. Attach it to notebooks and Spark job definitions, or make it the **workspace default** so every notebook set to "Workspace default" inherits it. Once an environment is the workspace default, only workspace admins can change its contents.

Libraries publish in one of two modes (Microsoft Learn, checked September 2026):

- **Quick mode** installs libraries when each session starts. It is flexible, and it adds one to three minutes to start-up.
- **Full mode** builds a snapshot of the dependencies once. Microsoft recommends Full mode for pipelines, scheduled runs and shared workloads that need reproducible environments. It is the Fabric equivalent of the lock file in lesson B6-L3.

Resource profiles set a tested group of Spark and Delta settings with one property. New workspaces default to `writeHeavy`, tuned for ingestion. `readHeavyForSpark` and `readHeavyForPBI` suit read-heavy work.

V-Order is a Fabric write-time layout for Parquet files. It keeps files standard, open Parquet, and speeds up reads, especially for Power BI Direct Lake and the SQL engines. The cost is slower writes: Learn cites **around 15%** on average. Since the `writeHeavy` default, **V-Order is off by default in new workspaces** (`spark.sql.parquet.vorder.default=false`).

Turn it on where reads dominate, at the narrowest useful level:

- per table: `TBLPROPERTIES("delta.parquet.vorder.enabled" = "true")`, the best choice for gold tables that Power BI reads,
- per write: the DataFrame option `parquet.vorder.enabled`,

- per session: `spark.sql.parquet.vorder.default`, or the `readHeavyForPBI` profile, which enables it.

In Runtime 1.3 and later the older `spark.sql.parquet.vorder.enable` setting is removed. Code copied from older posts that sets it does nothing.

Why it matters

A team moved its gold layer to a new Fabric workspace and noticed Power BI reports running noticeably slower than before. The old workspace had V-Order on for every write. The new one used the `writeHeavy` default, with V-Order off. Setting `delta.parquet.vorder.enabled` on the dozen gold tables, and optimising them once, restored read performance without slowing bronze and silver ingestion.

Worked example: weak code

```
1 spark.conf.set("spark.sql.parquet.vorder.enable", "true")
2 bronze.write.mode("append").saveAsTable("lh.bronze.orders")
3 gold.write.mode("overwrite").saveAsTable("lh.gold.sales_by_region")
```

1. **Line 1.** Removed setting in Runtime 1.3 and later. It has no effect.
2. **Line 2.** Had the setting worked, bronze would pay the write cost for reads nobody makes.

Worked example: strong code

```
1 bronze.write.mode("append").saveAsTable("lh.bronze.orders")
2
3 spark.sql("""
4     CREATE TABLE IF NOT EXISTS lh.gold.sales_by_region (region STRING,
5     sales DECIMAL(19, 4))
6     USING DELTA
7     TBLPROPERTIES ("delta.parquet.vorder.enabled" = "true")
8 """)
9 gold.write.mode("overwrite").saveAsTable("lh.gold.sales_by_region")
```

1. **Line 1.** Write-heavy bronze keeps the fast default.
2. **Line 6.** V-Order only on the gold table that Power BI reads.

Pin libraries in a Full mode environment, keep the `writeHeavy` default for ingestion, and turn V-Order on for read-heavy gold tables.

Recap card

- **Environment publishing modes** Quick: installs at session start (adds minutes). Full: reproducible snapshot, recommended for pipelines.
- **V-Order in new workspaces** Off by default (`writeHeavy` profile). Write cost around 15%, faster reads.

- **Turn V-Order on for one table**

TBLPROPERTIES("delta.parquet.vorder.enabled" = "true").

- **spark.sql.parquet.vorder.enable** Removed in Runtime 1.3 and later. Use spark.sql.parquet.vorder.default or table and writer options.

Practice

Exercise 1. Multiple choice (level 1)

In a workspace created in 2026 with default settings, is V-Order applied when a notebook writes a Delta table?

- a) Yes, V-Order is always on in Fabric.
- b) No. New workspaces default to the writeHeavy profile, with V-Order off.
- c) Only for tables under Files.
- d) Only when the table is larger than 1 GB.

Why?

- r1) spark.sql.parquet.vorder.default is false in new workspaces to favour write performance.
- r2) V-Order has been retired.

Exercise 2. Multiple choice (level 2)

A nightly pipeline runs notebooks that need pandera and a pinned version of a statistics library. How should the libraries be delivered?

- a) %pip install at the top of each notebook
- b) An environment in Full mode, attached to the notebooks or set as workspace default
- c) An environment in Quick mode
- d) Upload wheel files to Files and install them at runtime

Why?

- r1) Full mode builds a reproducible snapshot, and Microsoft recommends it for pipelines and scheduled runs.
- r2) Full mode is the only mode that supports pandera.

Exercise 3. Fill the gap (level 2)

Turn on V-Order for one gold table only.

```
1 ALTER TABLE lh.gold.sales_by_region
2 SET TBLPROPERTIES ("_____(1)" = "_____(2)");
```

Exercise 4. Write the function (level 3)

Write `write_gold(df, table)` that overwrites a gold Delta table with V-Order enabled for this write, using the DataFrame writer option `parquet.vorder.enabled` set to `"true"`. Do not change the session-wide default.

Starter code:

```
1 def write_gold(df, table):
2     spark.conf.set("spark.sql.parquet.vorder.default", "true")
3     df.write.mode("overwrite").saveAsTable(table)
```

Exercise 5. Architecture design (level 4)

Set up Spark for a workspace with bronze and silver ingestion and a gold layer read by Power BI Direct Lake.

Options: One Full mode environment, set as workspace default, %pip in each notebook.

1. Libraries for scheduled notebooks
2. Session-wide resource profile
3. Gold tables

Sources

- Environments set runtime, compute, libraries and resources; a workspace default environment can be changed only by workspace admins. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/create-and-use-environment> (updated 2026-03-25))
- Quick mode installs at session start adding one to three minutes; Full mode is recommended for pipelines and scheduled runs. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/library-management> (updated 2026-03-25))
- V-Order is off by default in new workspaces (writeHeavy); write cost around 15%; table property, writer option and session default; `spark.sql.parquet.vorder.enable` removed in Runtime 1.3+. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/delta-optimization-and-v-order> (updated 2026-03-01))
- New workspaces default to writeHeavy; readHeavyForPBI enables V-Order. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/resource-profiles-overview> (updated 2026-07-23))

C5-L3 Delta tables shared between Databricks and Fabric

Concept

Delta is an open format, so a table written by Databricks can be read by Fabric through a shortcut or Unity Catalog mirroring (lessons C3-L1 and C3-L3). Open does not mean every reader supports every **table feature**. A table's features decide which Fabric experiences can read it.

From Microsoft Learn's Delta interoperability matrix (updated May 2026):

Feature	Spark Runtime 2.0	SQL analytics endpoint	Direct Lake	Python notebooks
Deletion vectors	yes	yes	yes	no
Column mapping, name mode	yes	yes	yes	no
Column mapping, id mode	yes	no	yes	no
<code>TIMESTAMP_NTZ</code>	yes	no	no	yes
Liquid clustering	yes	yes	yes	no

Three more limits from the same page:

- **V2 checkpoints**: only Spark notebooks and Spark jobs read and write them. The lakehouse and the SQL analytics endpoint do not list such tables correctly.
- **Databricks-only features**, such as writing identity columns and Lakeflow Spark Declarative Pipelines tables, are not supported.
- Newer **Delta 4.x features**, such as collations and the variant type, work only in Spark notebooks and jobs, not in the SQL endpoint or Direct Lake.

Fabric's own writers differ too. Runtime 2.0 writes new tables at reader version 3 and writer version 7 with **deletion vectors** on by default. Runtime 1.3 writes reader 1, writer 2.

The practical rule: **the table's writer decides which features the table has, and every reader must support all of them**. For gold tables that Power BI or T-SQL users read, use timestamps with a time zone (`TIMESTAMP`) instead of `TIMESTAMP_NTZ`, avoid column mapping in id mode, and check `DESCRIBE DETAIL` for the protocol versions and table features before you publish.

Why it matters

A Databricks team stored event times as `TIMESTAMP_NTZ` in a gold table and shared it with Fabric through a shortcut. Spark in Fabric read it fine. The SQL analytics endpoint and the Direct Lake report did not show the table's time columns, so the Power BI team built a workaround copy with converted columns. One type choice in the writer created a second dataset.

Worked example: weak code

```
1 CREATE TABLE gold.events (  
2     event_id STRING,  
3     event_time TIMESTAMP_NTZ  
4 )  
5 TBLPROPERTIES ("delta.columnMapping.mode" = "id");
```

1. **Line 3.** Not supported by the SQL analytics endpoint or Direct Lake.
2. **Line 5.** Column mapping in id mode is not supported by the SQL analytics endpoint.

Worked example: strong code

```
1 CREATE TABLE gold.events (  
2     event_id STRING,  
3     event_time TIMESTAMP  
4 );  
5  
6 -- Before sharing: check protocol versions and table features  
7 DESCRIBE DETAIL gold.events;
```

1. **Line 3.** A time-zone-aware timestamp that every Fabric reader supports.
2. **Line 7.** Shows `minReaderVersion`, `minWriterVersion` and table features to compare with the readers' support.

Open format, but not every feature for every reader. Design shared gold tables for the weakest reader that needs them.

Recap card

- **Delta interoperability rule** Every reader must support every feature the writer put in the table.
- **TIMESTAMP_NTZ in Fabric** Spark yes. SQL analytics endpoint and Direct Lake no.
- **Runtime 2.0 default table protocol** Reader 3, writer 7, deletion vectors on. Python notebooks cannot read deletion vectors.
- **Check a table's features** `DESCRIBE DETAIL`: `minReaderVersion`, `minWriterVersion`, table features.

Practice

Exercise 1. Multiple choice (level 1)

A Databricks gold table uses `TIMESTAMP_NTZ` columns. Which Fabric reader handles them?

- a) The SQL analytics endpoint
- b) Direct Lake semantic models
- c) Fabric Spark Runtime 2.0
- d) All of them

Why?

- r1) The interoperability matrix lists `TIMESTAMP_NTZ` as supported in Spark but not in the SQL endpoint or Direct Lake.
- r2) Only Spark can read Delta tables.

Exercise 2. Multiple choice (level 2)

You create a new Delta table with a Spark notebook on Fabric Runtime 2.0 and default settings. Which reader might fail to read it?

- a) The SQL analytics endpoint
- b) A Python notebook (non-Spark) using the default Delta reader
- c) Direct Lake
- d) Fabric Spark Runtime 2.0

Why?

- r1) Runtime 2.0 writes deletion vectors by default, and Python notebooks do not support deletion vectors.
- r2) Python notebooks cannot read OneLake.

Exercise 3. Fill the gap (level 2)

Inspect a table's protocol versions and features before sharing it.

```
1 _____.DETAIL gold.events;
```

Exercise 4. Write the function (level 3)

Write `unsupported(reader: str, features: set[str]) -> list[str]` using this support table from the lesson. Return the table's features the reader does not support, sorted. Raise `KeyError` for an unknown reader.

- `spark_2_0`: supports all of `deletion_vectors`, `column_mapping_name`, `column_mapping_id`, `timestamp_ntz`, `liquid_clustering`.
- `sql_endpoint`: all except `column_mapping_id` and `timestamp_ntz`.

- `direct_lake`: all except `timestamp_ntz`.
- `python_notebook`: only `timestamp_ntz`.

Starter code:

```
1 def unsupported(reader, features):  
2     return []
```

Exercise 5. Architecture design (level 4)

A Databricks team publishes a gold table that Fabric users read through the SQL endpoint and Direct Lake. Choose the table design.

Options: `TIMESTAMP`, `TIMESTAMP_NTZ`.

1. Event time column
2. Column mapping
3. Before each release

Sources

- Feature support matrix for deletion vectors, column mapping, `TIMESTAMP_NTZ` and liquid clustering across Spark runtimes, SQL analytics endpoint, Direct Lake and Python notebooks; Runtime 2.0 writes reader 3, writer 7 with deletion vectors; V2 checkpoints, identity columns writing, Lakeflow Spark Declarative Pipelines and Delta 4.x features outside Spark are limitations. (verified: <https://learn.microsoft.com/en-us/fabric/fundamentals/delta-lake-interoperability> (updated 2026-05-08))

C6 Semantic models and Direct Lake

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Star schemas, DAX best practices and Power BI reports.

C6-L1 Star schemas for semantic models

Concept

A Power BI semantic model works best on a **star schema**: fact tables surrounded by dimension tables. Microsoft's Power BI guidance (checked September 2026) defines the parts:

- **Fact tables** store observations or events: survey responses, orders, balances. They hold dimension key columns and numeric measure columns. The dimension keys decide the fact table's **grain**, for example one row per respondent per wave per brand.
- **Dimension tables** describe business entities: respondent, brand, region, date. Each has a unique key column and descriptive columns used to filter and group. A **date dimension** belongs in almost every model.

Relationships run **one-to-many** from each dimension to the fact table, and filters flow from the dimension to the fact. Microsoft recommends that you **minimise bi-directional relationships**, which can slow queries and confuse report users with unexpected filtering. Many-to-many relationships between facts and dimensions usually mean the grain is unclear.

Build the star in the gold layer, not in the model:

- Compute derived columns such as age bands or weighted flags in Spark or SQL, in gold tables. Direct Lake on SQL analytics endpoints does not support calculated columns, and on OneLake they are in preview with restrictions. Materialised calculated columns in Import models slow refresh.
- Give every fact table a clear grain in its docstring and table comment (lesson A2-L3).
- Use surrogate or natural keys consistently, and check they are unique in each dimension (lesson C2-L2 showed warehouse keys are not enforced).

A good model is boring: one fact table per grain, conformed dimensions shared between facts, single-direction one-to-many relationships, and measures that do the arithmetic (lesson C6-L2).

Why it matters

A brand tracker model joined a respondent table and a brand-awareness table with a bi-directional many-to-many relationship, because both held respondent_id and brand. Selecting a brand in a slicer filtered respondents, which filtered back to every brand those respondents had mentioned. Awareness for small brands showed as much higher than the survey measured. A single fact table at respondent-brand grain with one-way dimensions returned the correct figures.

Worked example: weak code

```
❶ respondents (respondent_id, age, region, weight, brand_list)
❷ brand_awareness (respondent_id, brand, aware)
❸ Relationship: respondents[respondent_id] <->
  brand_awareness[respondent_id], many-to-many, both directions
❹ Calculated column: respondents[age_band] = SWITCH(TRUE(), ...)
```

1. **Line 1.** A list column mixes grains: one respondent, many brands.
2. **Line 3.** Bi-directional many-to-many: filters flow both ways and inflate results.
3. **Line 4.** Derived column in the model. Not supported in Direct Lake on SQL endpoints.

Worked example: strong code

```
❶ fact_awareness (respondent_key, brand_key, date_key, region_key, aware,
  weight) grain: respondent x brand x wave
❷ dim_respondent (respondent_key, age_band, gender)
❸ dim_brand (brand_key, brand, category)
❹ dim_date (date_key, date, wave, month)
❺ dim_region (region_key, region, nation)
❻ Relationships: each dim[key] 1 -> * fact_awareness[key], single direction
```

1. **Line 1.** One fact table with a stated grain.
2. **Line 2.** age_band computed in gold, not in the model.
3. **Line 6.** Filters flow one way, from dimensions to facts.

Facts with a stated grain, dimensions with unique keys, one-to-many single-direction relationships, and derived columns computed upstream.

Recap card

- **Star schema parts** Fact tables (keys and measures, with a grain) and dimension tables (unique key, descriptive columns).
- **Relationship guidance** One-to-many, single direction. Minimise bi-directional relationships.

- **Calculated columns and Direct Lake** Not supported on SQL endpoints, preview on OneLake. Compute derived columns in gold.
- **Grain** Set by the fact table's dimension keys. State it in the table's docstring and comment.

Practice

Exercise 1. Multiple choice (level 1)

In a star schema, what decides the grain of a fact table?

- a) The number of measure columns
- b) The dimension keys it holds, and their values
- c) The number of rows
- d) The storage mode

Why?

- r1) Dimension key columns set the dimensionality, and their values set the granularity.
- r2) Grain is chosen in Power BI settings.

Exercise 2. Multiple choice (level 2)

What does Microsoft's Power BI guidance recommend about bi-directional relationships?

- a) Use them by default so slicers work everywhere.
- b) Minimise them. They can hurt performance and confuse report users.
- c) Use them only on date tables.
- d) They are not supported in Direct Lake.

Why?

- r1) Bi-directional filters can slow queries and filter in ways users do not expect.
- r2) They are deprecated.

Exercise 3. Put in order (level 2)

Order the steps to build a model for a brand tracker.

- Build the fact table with dimension keys and measure columns
- Write DAX measures
- Build dimension tables with unique keys in gold
- Create one-to-many, single-direction relationships

- State the fact table's grain

Exercise 4. Write the function (level 3)

Write `model_warnings(relationships: list[dict[str, str]]) -> list[str]` for a model review. Each relationship has `from_table`, `to_table`, `cardinality` ("one_to_many" or "many_to_many") and `direction` ("single" or "both"). Return, in input order, a warning for each problem:

- "<from>-<to>: many-to-many" for a many-to-many relationship,
- "<from>-<to>: bi-directional" for direction both.

A relationship with both problems produces both warnings, many-to-many first.

Starter code:

```
1 def model_warnings(relationships):
2     return []
```

Exercise 5. Architecture design (level 4)

A Direct Lake model on the SQL analytics endpoint needs an age band on respondents. Choose where it is built.

Options: In the gold `dim_respondent` table, in Spark or SQL, A calculated column in the semantic model.

1. Where `age_band` is computed
2. Awareness fact table grain
3. Relationship from `dim_respondent` to the fact

Sources

- Fact and dimension definitions; dimension keys set dimensionality and their values granularity. (verified: <https://learn.microsoft.com/en-us/power-bi/guidance/star-schema>)
- Microsoft recommends minimising bi-directional relationships because of performance and confusing experiences. (verified: <https://learn.microsoft.com/en-us/power-bi/guidance/relationships-bidirectional-filtering>)
- Calculated columns: Direct Lake on SQL analytics endpoints no; Direct Lake on OneLake user-context only, preview. (verified: <https://learn.microsoft.com/en-us/fabric/fundamentals/direct-lake-overview> (updated 2026-09-02))

C6-L2 DAX measures that state their filter context

Concept

In a semantic model, arithmetic belongs in **measures**. A measure is evaluated at query time, in the **filter context** of the visual: the slicers, rows and columns that apply to that cell. A calculated column is evaluated row by row when the model refreshes and stored, which costs memory and refresh time and cannot react to slicers.

Three habits make measures correct and readable (Microsoft's DAX guidance, checked September 2026):

- **Use DIVIDE, not /.** `DIVIDE(numerator, denominator)` returns BLANK when the denominator is zero or BLANK, instead of an error or infinity. You can pass an alternate result as a third argument. Microsoft notes it is also faster than testing for zero yourself.
- **Use variables.** `VAR ... RETURN` evaluates an expression once and names it. That improves performance, readability and debugging.
- **Change filter context only with CALCULATE,** and say why. `CALCULATE([Weighted base], ALL(dim_brand))` removes the brand filter, which is what a share-of-total denominator needs.

A weighted share, the survey version of lesson B8-L1:

```
1 Weighted aware % =  
2 VAR aware_w = SUMX ( fact_awareness, fact_awareness[aware] *  
   fact_awareness[weight] )  
3 VAR base_w = SUM ( fact_awareness[weight] )  
4 RETURN DIVIDE ( aware_w, base_w )
```

Put the unweighted base next to it as a separate measure, `Base (n) = DISTINCTCOUNT(fact_awareness[respondent_key])`, so report users see how many respondents stand behind each percentage.

Name measures for what they return, store them in one place (a measures table or the fact table), and set their format strings. A measure that nobody can read is a measure nobody can check.

Why it matters

A report computed awareness as `SUM(aware_weight) / SUM(weight)` in a measure. In a new market with no fieldwork yet, the base was zero, and the visual showed "Infinity" and "NaN" in cells that the board pack exported to PDF. `DIVIDE` would have returned BLANK, and the empty cells would have been left empty.

Worked example: weak code

```
1 Aware % = SUM(fact_awareness[aware_weight]) / SUM(fact_awareness[weight])
2 Brand share = [Aware %] / CALCULATE([Aware %], ALL(fact_awareness))
```

1. **Line 1.** The / operator breaks on a zero base.
2. **Line 2.** ALL on the whole fact table removes every filter, including date and region. The share compares against the wrong total.

Worked example: strong code

```
1 Weighted aware % =
2 VAR aware_w = SUMX ( fact_awareness, fact_awareness[aware] *
    fact_awareness[weight] )
3 VAR base_w = SUM ( fact_awareness[weight] )
4 RETURN DIVIDE ( aware_w, base_w )
5
6 Share of aware across brands =
7 VAR this_brand = [Weighted aware %]
8 VAR all_brands = CALCULATE ( [Weighted aware %], ALL ( dim_brand ) )
9 RETURN DIVIDE ( this_brand, all_brands )
```

1. **Line 4.** BLANK on a zero base, not an error.
2. **Line 8.** Removes only the brand filter. Date, region and wave filters still apply.

Measures, not calculated columns. DIVIDE for every ratio, variables for clarity, and CALCULATE with the narrowest filter change.

Recap card

- **Measure vs calculated column** Measures compute at query time in the visual's filter context. Calculated columns are stored at refresh.
- **DIVIDE in DAX** Returns BLANK (or an alternate result) when the denominator is zero or BLANK.
- **VAR ... RETURN** Evaluates once, names parts, improves performance, readability and debugging.
- **Share of total across brands** CALCULATE([measure], ALL(dim_brand)): remove only the brand filter.

Practice

Exercise 1. Multiple choice (level 1)

What does `DIVIDE(10, 0)` return in DAX?

- a) An error
- b) BLANK
- c) 0

- d) Infinity

Why?

- r1) DIVIDE returns BLANK when the denominator is zero or BLANK, unless an alternate result is given.
- r2) DAX has no concept of division by zero.

Exercise 2. Multiple choice (level 2)

A share-of-total measure uses `CALCULATE([Weighted aware %], ALL(fact_awareness))` as its denominator. A report filters to wave 12 and region North. What does the denominator ignore?

- a) Only the brand filter
- b) Every filter on fact_awareness, including wave and region
- c) Nothing: ALL only affects slicers
- d) Only filters from other pages

Why?

- r1) ALL on the fact table removes all its filters, including those that reach it through dimensions.
- r2) ALL always removes every filter in the model.

Exercise 3. Fill the gap (level 2)

Complete the weighted share measure.

```
1 Weighted aware % =
2 ____ (1) aware_w = SUMX ( fact_awareness, fact_awareness[aware] *
   fact_awareness[weight] )
3 VAR base_w = SUM ( fact_awareness[weight] )
4 RETURN ____ (2) ( aware_w, base_w )
```

Exercise 4. Write the function (level 3)

Write a measure `Share of aware across brands` that divides `[Weighted aware %]` by the same measure with the brand filter removed through `ALL ( dim_brand )`. Use variables named `this_brand` and `all_brands`, and `DIVIDE`.

Starter code:

```
1 Share of aware across brands = [Weighted aware %] / CALCULATE ( [Weighted
   aware %], ALL ( fact_awareness ) )
```

Exercise 5. Architecture design (level 4)

Decide where each calculation lives in a Direct Lake model on the SQL analytics endpoint.

Options: Gold table column, DAX measure, Calculated column.

1. Respondent age band
2. Weighted awareness percentage
3. Unweighted base (number of respondents)

Sources

- DIVIDE returns BLANK for a zero or BLANK denominator unless an alternate result is passed; it outperforms explicit zero checks. (verified: <https://learn.microsoft.com/en-us/dax/best-practices/dax-divide-function-operator>)
- Variables improve performance, reliability, readability and debugging. (verified: <https://learn.microsoft.com/en-us/dax/best-practices/dax-variables>)
- Materialised calculated columns can hurt refresh performance because values are derived at refresh time. (verified: <https://learn.microsoft.com/en-us/power-bi/transform-model/desktop-calculated-columns> (updated 2026-08-23))

C6-L3 Direct Lake, framing and the fallback you did not plan

Concept

Direct Lake lets a semantic model read Delta tables in OneLake directly, loading only the columns a query needs into memory. It combines Import-like query speed with no copy of the data. A refresh does not reload data. It performs **framing**: it records which Delta table versions the model uses. Framing takes seconds, and afterwards queries see the data as of that frame. Automatic updates, a schedule or an API call can trigger framing.

There are two kinds (Microsoft Learn, updated September 2026):

- **Direct Lake on OneLake** reads Delta tables from one or more Fabric items. It never falls back to DirectQuery. If it cannot load data, the query fails.
- **Direct Lake on SQL analytics endpoints** reads one Fabric item through its endpoint. When it cannot load a table directly, it **falls back to DirectQuery**: for example when the table is a SQL view, or when the warehouse uses SQL-based granular access control. Fallback queries are slower.

The model property **Direct Lake behavior** controls fallback for the SQL kind:

- `Automatic` (default): silently fall back to `DirectQuery`. Reports keep working, more slowly.
- `DirectLakeOnly`: fail with an error instead. Use it in development to find fallbacks.
- `DirectQueryOnly`: always use `DirectQuery`, to measure the difference.

Guardrails depend on the SKU. For each table, Fabric limits Parquet files, row groups and rows. Over a guardrail, framing can fail (or, on the SQL kind, queries fall back). From Learn's table:

SKU	Parquet files per table	Rows per table	Max memory
F2 to F32	1,000	300 million	3 to 10 GB
F64	5,000	1.5 billion	25 GB
F512 and above	10,000	12 to 24 billion	200 to 400 GB

Small files hit the file guardrail first. Compact gold tables with `OPTIMIZE` (lesson A3-L3), and give Direct Lake tables V-Order (lesson C5-L2).

Why it matters

A sales model on an F32 capacity was fast in testing. After six months of hourly appends, the fact table held 2,400 small Parquet files, above the F32 guardrail of 1,000. The model used Direct Lake on the SQL endpoint with the Automatic behaviour, so queries quietly fell back to `DirectQuery`. Reports took 40 seconds instead of two. Nobody saw an error. A nightly `OPTIMIZE` kept the table at a few dozen files and restored Direct Lake.

Worked example: weak code

```
1 Model: sales (Direct Lake on SQL analytics endpoint), behaviour: Automatic
2 Table fact_sales: 2,400 Parquet files, appended hourly, never optimised
3 Table v_sales_clean: SQL view over fact_sales
4 Capacity: F32
```

1. **Line 1.** Automatic hides both problems: queries silently fall back.
2. **Line 2.** Above the F32 guardrail of 1,000 files per table.
3. **Line 3.** A SQL view always forces `DirectQuery` fallback on the SQL kind.

Worked example: strong code

```
1 Model: sales (Direct Lake on SQL analytics endpoint), behaviour:
  DirectLakeOnly in dev, Automatic in prod
2 Table fact_sales: nightly OPTIMIZE, V-Order on, about 40 Parquet files
```


- 2 Table gold.sales_clean: Delta table written by Spark, replacing the SQL view
- 3 Capacity: F32, fact rows 120 million (limit 300 million)
- 5 Check: framing failures reviewed in refresh history each morning

1. **Line 1.** Development fails loudly on fallback. Production keeps working if one slips through.
2. **Line 3.** Materialise the view's logic as a Delta table so Direct Lake reads it directly.
3. **Line 4.** Row count checked against the SKU's guardrail.

Direct Lake stays fast when tables are Delta (not views), compact, within the SKU's guardrails, and fallback is made visible in development.

Recap card

- **Direct Lake refresh** Framing: records Delta table versions in seconds. No data copy.
- **Direct Lake on OneLake vs on SQL** OneLake: never falls back. SQL: falls back to DirectQuery for views and SQL-based security.
- **Direct Lake behavior values** Automatic (default, silent fallback), DirectLakeOnly (error), DirectQueryOnly (always DirectQuery).
- **F64 per-table guardrails** 5,000 Parquet files, 1.5 billion rows, 25 GB max memory.

Practice

Exercise 1. Multiple choice (level 1)

What does a refresh of a Direct Lake semantic model do?

- a) Reloads all data into the model, like Import.
- b) Frames the model: records which Delta table versions queries will use.
- c) Rewrites the Delta tables.
- d) Switches the model to DirectQuery.

Why?

- r1) Framing copies metadata only, takes seconds, and sets the point in time queries see.
- r2) Direct Lake models never refresh.

Exercise 2. Multiple choice (level 2)

A Direct Lake on SQL analytics endpoint model includes a SQL view. The Direct Lake behavior property is Automatic. What happens when a report queries the view?

- a) The query fails.
- b) The query falls back to DirectQuery and runs more slowly.
- c) The view is materialised automatically.
- d) The view is read directly from OneLake.

Why?

- r1) Views cannot be loaded directly from Delta, and Automatic falls back silently.
- r2) Views are always faster than tables.

Exercise 3. Fill the gap (level 2)

Set the Direct Lake behaviour so that development fails loudly on fallback, and production keeps working.

```
1 Development model: Direct Lake behavior = ____ (1)
2 Production model: Direct Lake behavior = ____ (2)
```

Exercise 4. Write the function (level 3)

Write `guardrail_problems(sku: str, parquet_files: int, rows_millions: float) -> list[str]` using these limits from the lesson: F2 to F32 allow 1,000 files and 300 million rows per table, F64 allows 5,000 and 1,500, F128 5,000 and 3,000, F256 5,000 and 6,000. Return `"files"` and/or `"rows"` for each limit exceeded, in that order. Raise `KeyError` for other SKUs.

Starter code:

```
1 def guardrail_problems(sku, parquet_files, rows_millions):
2     return []
```

Exercise 5. Spot the bug (level 4)

A gold table read by Direct Lake on an F32 grows by one small file per hour and has passed the file guardrail. Tap the line that causes the build-up, and rewrite the job so the table is compacted after the append.

```
1 def publish_hourly(df, table):
2     """Append the latest hour to a Direct Lake gold table."""
3     df.repartition(1).write.mode("append").saveAsTable(table)
4     return None
```

Write down the line number of the bug. Then write the fix.

Sources

- Direct Lake on OneLake never falls back; Direct Lake on SQL falls back to DirectQuery for SQL views and SQL-based granular access control; framing copies metadata in seconds; guardrails per SKU table. (verified: <https://learn.microsoft.com/en-us/fabric/fundamentals/direct-lake-overview> (updated 2026-09-02))
- Direct Lake behavior values Automatic (default), DirectLakeOnly and DirectQueryOnly; framing can fail when a table exceeds guardrails. (verified: <https://learn.microsoft.com/en-us/fabric/fundamentals/direct-lake-how-it-works> (updated 2025-09-22))

C7 Real-Time Intelligence

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Eventstream, Eventhouse, KQL and Activator.

C7-L1 Eventstream, events in, routed out

Concept

An **eventstream** brings real-time events into Fabric, optionally transforms them, and routes them to destinations, without code. It is Microsoft's recommended path for low-latency, event-driven ingestion (lesson C4-L1).

Sources include Azure Event Hubs, IoT Hub, custom apps, and change data capture (CDC) connectors for Azure SQL Database, PostgreSQL, MySQL and SQL Server on a VM. A CDC source first captures a snapshot of the current data, then streams the changes.

Transformations are no-code operators on the stream: filter, manage fields, aggregate over a time window, group by, union, join, expand. You can split one stream into **derived streams** for different consumers.

Destinations (Microsoft Learn, checked September 2026):

- **Eventhouse**, for KQL analysis (lesson C7-L2). Two ingestion modes: **direct ingestion**, or **event processing before ingestion** when you apply transformations in the eventstream first.
- **Lakehouse**. Events are converted to Delta Lake format and written to lakehouse tables.
- **Activator**, to trigger actions on conditions (lesson C7-L3).
- A **custom endpoint**, to route events to your own application.

Writing to an eventhouse, lakehouse or Activator destination needs

Contributor or higher on the destination's workspace.

Design choices that matter:

- **Filter early**. Drop events nobody needs in the eventstream, before they cost storage and query time.
- **Keep the raw stream**. Route an unfiltered copy to an eventhouse or lakehouse table as your bronze layer, and derive cleaned streams from it. You cannot replay events you dropped.

- **Decide the event time.** Events carry their own timestamp. Use it for windows and analysis, not the time of arrival, which drifts when senders buffer or reconnect.

Why it matters

A fleet team filtered its eventstream to keep only engine-fault events and wrote those to an eventhouse. Three months later, engineers needed the normal readings before each fault to build a failure model. They had never been stored. Keeping the raw stream in a cheap table, and deriving the fault stream from it, would have preserved the history the model needed.

Worked example: weak code

```
1 Source: IoT Hub (all vehicle telemetry)
❶ Transform: Filter event_type == "fault"
3 Destination: Eventhouse table faults
```

1. **Line 2.** The only copy of the data is filtered. Normal readings are lost for good.

Worked example: strong code

```
1 Source: IoT Hub (all vehicle telemetry)
❶ Destination 1: Eventhouse table telemetry_raw (direct ingestion,
    every event)
❷ Derived stream: faults = Filter event_type == "fault"
❸ Destination 2: Activator rule on faults (alert within
    seconds)
5 Destination 3: Lakehouse table bronze.vehicle_faults (Delta, for Spark
    models)
```

1. **Line 2.** The raw stream is kept as bronze. Anything else can be rebuilt from it.
2. **Line 3.** Consumers get a filtered stream without losing the source.
3. **Line 4.** Alerts come from the stream, not from a scheduled query.

Land the raw stream once, derive filtered streams for consumers, and route each to the destination that suits it.

Recap card

- **What an eventstream does** Ingests real-time events, transforms them without code, and routes them to destinations.
- **Eventstream destinations** Eventhouse (direct or processed ingestion), Lakehouse (Delta tables), Activator, custom endpoint.
- **Permission for eventhouse, lakehouse or Activator destinations** Contributor or higher on the destination workspace.
- **Protect replay** Land the raw stream first. Derive filtered streams from it.

Practice

Exercise 1. Multiple choice (level 1)

An eventstream sends events to a lakehouse destination. In what format are they stored?

- a) JSON files under Files
- b) Delta Lake tables
- c) CSV files
- d) KQL tables

Why?

- r1) The lakehouse destination converts real-time events to Delta Lake format in lakehouse tables.
- r2) Eventstream only writes Delta.

Exercise 2. Multiple choice (level 2)

Which permission does the person setting up an eventstream need on the workspace of an eventhouse destination?

- a) Viewer
- b) Contributor or higher
- c) Admin only
- d) None: eventstreams write as the service

Why?

- r1) Microsoft Learn requires Contributor or higher on the destination workspace for eventhouse, lakehouse and Activator destinations.
- r2) Contributors can create any item in any workspace.

Exercise 3. Put in order (level 2)

Order the design steps for a new telemetry eventstream.

- Connect the source
- Connect each derived stream to its destination
- Create derived streams with filters for each consumer
- Route every event to a raw table
- Compare event counts at source and destinations

Exercise 4. Write the function (level 3)

Write `route(event: dict[str, object]) -> list[str]` modelling the strong design. Every event goes to `"telemetry_raw"`. Events with `event_type == "fault"` also go to `"activator_faults"` and `"bronze.vehicle_faults"`. Events without an `event_type` key go to `"telemetry_raw"` and `"quarantine"`. Return destinations in that order.

Starter code:

```
1 def route(event):
2     if event["event_type"] == "fault":
3         return ["activator_faults"]
4     return []
```

Exercise 5. Architecture design (level 4)

Stream order changes from an Azure SQL database into Fabric for a live dashboard and for nightly Spark models.

Options: Azure SQL Database CDC source, A pipeline polling the table every minute.

1. Source connector
2. Destination for the live dashboard
3. Destination for nightly Spark models

Sources

- Eventstreams ingest, transform without code and route events; CDC source connectors for Azure SQL, PostgreSQL, MySQL and SQL Server on VM capture a snapshot then changes. (verified: <https://learn.microsoft.com/en-us/fabric/real-time-intelligence/event-streams/overview> (updated 2026-04-29))
- Destinations include custom endpoint, eventhouse (direct ingestion or processing before ingestion), lakehouse (Delta) and Activator; Contributor or higher is required for eventhouse, lakehouse and Activator destinations. (verified: <https://learn.microsoft.com/en-us/fabric/real-time-intelligence/event-streams/add-manage-eventstream-destinations> (updated 2026-03-19))

C7-L2 Eventhouse and KQL for time-series questions

Concept

An **eventhouse** holds one or more **KQL databases**, built for high-volume, time-stamped data such as telemetry, logs and clickstreams. You query them with the **Kusto Query Language (KQL)**: a pipeline of operators joined by `|`, read top to bottom.

```
1 telemetry_raw
2 | where event_time > ago(1h)
3 | where event_type == "reading"
4 | summarize avg_rpm = avg(rpm), n = count() by vehicle_id, bin(event_time, 5m)
```

- Put `where` filters first, especially on time. They cut the data every later operator touches.
- `ago(1h)` is one hour before now. `bin(event_time, 5m)` rounds times **down** to five-minute buckets.
- `summarize ... by` aggregates per group.

Three KQL behaviours catch people out (Microsoft Learn, checked September 2026):

- `take` **returns arbitrary rows**. There is no guarantee which records come back, and repeated runs can differ, unless the data is sorted. For "the latest 10", use `top 10 by event_time desc`.
- `summarize` **leaves gaps**. A five-minute bin with no events produces no row, so a chart draws a straight line across an outage. `make-series ... default = 0 ... step 5m` fills missing bins.
- **Empty input**. `summarize count()` without `by` returns one row. With `by`, empty input returns no rows at all.

Data policies are set per database or table. The **retention policy** removes old data automatically. The default is 3,650 days, the maximum 36,500, the minimum one day, and you can set it to unlimited. The **caching policy** decides how much recent data stays in the hot cache on local SSD: faster queries, higher cost. Set retention from the use: raw telemetry kept 90 days, daily aggregates kept for years.

Turn on **OneLake availability** for a database or table to expose its data as Delta in OneLake, so Spark and the SQL analytics endpoint can read it without a copy.

Why it matters

An on-call dashboard showed the "latest readings" from a KQL query ending in `| take 10`. During an incident, the panel showed readings from two hours earlier, mixed with current ones, and the engineer concluded the sensor had recovered. It had not. `take` returned arbitrary rows. `top 10 by event_time desc` would have shown the true latest readings.

Worked example: weak code

```
1 telemetry_raw
① | summarize avg_rpm = avg(rpm) by vehicle_id, bin(event_time, 5m)
3 | where event_time > ago(1h)
② | take 10
```

1. **Line 2.** Aggregates the whole table before filtering to the last hour.
2. **Line 4.** Arbitrary rows, different on each run.

Worked example: strong code

```
1 telemetry_raw
① | where event_time > ago(1h)
3 | where event_type == "reading"
② | make-series avg_rpm = avg(rpm) default = 0 on event_time step 5m by
   vehicle_id
5
6 telemetry_raw
7 | where event_time > ago(15m)
③ | top 10 by event_time desc
```

1. **Line 2.** Time filter first: less data for every later step.
2. **Line 4.** Every five-minute bin exists, so outages show as zeros, not as smooth lines.
3. **Line 8.** The true latest ten, every time.

Filter on time first, use top for "latest", make-series for continuous charts, and set retention and caching per table on purpose.

Recap card

- **take vs top in KQL** take: arbitrary rows, not repeatable. top N by col desc: ordered, repeatable.
- **Gaps in time series** summarize omits empty bins. make-series ... default = 0 fills them.
- **KQL retention policy defaults** Default 3,650 days, maximum 36,500, minimum 1 day, or unlimited.
- **Caching policy** How much recent data stays in the hot SSD cache: faster queries, higher cost.

Practice

Exercise 1. Multiple choice (level 1)

A query ends with `| take 10`. What do you get?

- a) The 10 most recent rows
- b) The first 10 rows ingested
- c) 10 arbitrary rows, possibly different on each run
- d) An error unless the table is sorted

Why?

- r1) take guarantees nothing about which records it returns unless the data is sorted.
- r2) take samples randomly with a fixed seed.

Exercise 2. Multiple choice (level 2)

A sensor sent nothing between 10:05 and 10:20. What does `summarize n = count() by bin(event_time, 5m)` return for those bins?

- a) Rows with `n = 0`
- b) No rows for those bins
- c) Rows with `n = null`
- d) An error

Why?

- r1) summarize groups existing rows. Bins with no rows do not appear. make-series fills them.
- r2) KQL drops zero counts to save space.

Exercise 3. Fill the gap (level 2)

Return the 10 most recent readings from the last hour.

```
1 telemetry_raw
2 | where event_time > ____ (1)(1h)
3 | ____ (2) 10 by event_time desc
```

Exercise 4. Write the function (level 3)

Write a KQL query on `telemetry_raw` that returns, for the last 24 hours, the average `rpm` per `vehicle_id` in one-hour steps, with **every** hour present and missing hours as 0. Filter on time first. Name the series `avg_rpm`.

Starter code:

```
1 telemetry_raw
```

```
2 | summarize avg_rpm = avg(rpm) by vehicle_id, bin(event_time, 1h)
```

Exercise 5. Architecture design (level 4)

Set data policies for a telemetry eventhouse.

Options: 90 days, Leave the default of 3,650 days, 7 days.

1. Retention for telemetry_raw (queried for 30 days, needed for audits for 90)
2. Hot cache for telemetry_raw
3. Retention for daily aggregates used in yearly trends

Sources

- take guarantees nothing about which records it returns unless the data is sorted; bin rounds down; summarize on empty input with a by clause returns no rows; make-series assigns a default to absent values. (verified: <https://learn.microsoft.com/en-us/kusto/query/> (take-operator, bin-function, summarize-operator, make-series-operator; view=microsoft-fabric))
- KQL database retention defaults to 3,650 days, maximum 36,500, minimum one day, unlimited option; caching policy controls data kept in the hot cache. (verified: <https://learn.microsoft.com/en-us/fabric/real-time-intelligence/data-policies> (updated 2026-03-22))
- OneLake availability can be enabled per database or table. (verified: <https://learn.microsoft.com/en-us/fabric/real-time-intelligence/eventhouse> (updated 2026-06-15))

C7-L3 Activator, rules that act on events

Concept

Activator is Fabric's no-code event detection engine. It watches data and takes an action when a condition is met. Its sources include eventstreams, Fabric and Azure events, Power BI reports and Real-Time dashboards.

Activator groups **events** into **objects** by a shared ID, such as `vehicle_id`. It evaluates each rule **per object** and keeps state for each one. "Vehicle 17 overheated" is a separate fact from "vehicle 42 overheated".

Two kinds of condition behave differently (Microsoft Learn, checked September 2026):

- **Change conditions** fire on a transition. `Increases above 100` fires when a value crosses 100 from below. `Becomes true` fires when a boolean changes from false to true.
- **State conditions** fire on every event that matches. `Is greater than 100` fires for each reading above 100.

For alerts, you want a change condition, or a state condition with an **occurrence** setting: "When it has been true for 10 minutes" fires only if the condition holds for the whole period. **No presence of data** (heartbeat) fires when no new events arrive within a set time. Use it to detect a silent sensor.

Actions include an email, a Teams message, a Power Automate flow, or running a Fabric item such as a pipeline or notebook.

Limits to design around:

- Up to 10,000 incoming events per second per rule. Above that, Activator stops the rule.
- Emails and Teams messages: 30 per rule per recipient per hour.
- Late events: by default, Activator drops events that arrive after their time window closes. The **Wait time for late-arriving events** setting (default two minutes) lets them in, at the cost of added latency.

Why it matters

A plant set a rule `Temperature Is greater than 80` on a stream that reports every 10 seconds. One oven ran hot for 20 minutes. The rule tried to send 120 emails to the shift lead. After 30, the rate limit throttled the rest, and the lead muted the alerts. An `Increases above 80` rule, or `Is greater than 80` with an occurrence of five minutes, would have sent one useful alert.

Worked example: weak code

- 1 Object: all events (no ID)
- 2 Condition: `temperature Is greater than 80`
- 3 Occurrence: Every time the condition is met
- 4 Action: Email shift lead

1. **Line 1.** One oven's reading clears the state of another. Group by oven.
2. **Line 2.** A state condition fires on every hot reading.

Worked example: strong code

- 1 Object: oven, keyed by `oven_id`

2 Condition: temperature Is greater than 80

② Occurrence: When it has been true for 5 minutes

4 Action: Teams message to the line channel

③ Rule 2: No presence of data for 2 minutes per oven_id

1. **Line 1.** State is tracked per oven.
2. **Line 3.** One alert for a sustained problem, not one per reading.
3. **Line 5.** A silent sensor is a problem too. Heartbeat detects it.

Key each rule by the object it describes. Fire on change or sustained state, not on every event. Add a heartbeat for silence.

Recap card

- **Events and objects in Activator** Events are grouped into objects by a shared ID. Rules run and keep state per object.
- **Increases above vs Is greater than** Increases above fires on the crossing. Is greater than fires on every matching event.
- **Detect a silent sensor** No presence of data (heartbeat) fires when no events arrive within a set time.
- **Activator limits to design around** 10,000 events per second per rule. 30 emails or Teams messages per rule per recipient per hour.

Practice

Exercise 1. Multiple choice (level 1)

A stream reports `temperature` every 10 seconds. It rises above 100 and stays there for 3 minutes. How many times does `Increases above 100` fire?

- a) Once
- b) 18 times, once per reading
- c) Never, because the value never decreases
- d) Once every minute

Why?

- r1) Increases above is a change condition. It fires when the value crosses the threshold from below.
- r2) Activator sends one alert per rule per day.

Exercise 2. Multiple choice (level 2)

A sensor stops sending data. Which Activator condition detects it?

- a) Is less than 1
- b) No presence of data

- c) Becomes false
- d) Changes

Why?

- r1) No presence of data fires when no new events arrive within a set time. Other conditions need an event to evaluate.
- r2) A silent sensor sends zeros.

Exercise 3. Fill the gap (level 2)

Alert once when a freezer stays above -15 for 10 minutes. Complete the rule.

```
1 Object: freezer, keyed by ____ (1)
2 Condition: temperature Is greater than -15
3 Occurrence: When it has been ____ (2) for 10 minutes
```

Exercise 4. Write the function (level 3)

Model an Activator Increases above rule. Write `activations(events: list[tuple[str, float]], threshold: float) -> list[int]`. Each event is `(object_id, value)` in arrival order. Return the indexes of events where that object's value goes above `threshold` and its previous value was at or below it. An object's first event counts as crossing if it is above the threshold. Track state per object.

Starter code:

```
1 def activations(events, threshold):
2     return [i for i, (_, value) in enumerate(events) if value > threshold]
```

Exercise 5. Architecture design (level 4)

Design alerts for 400 delivery vans that send GPS and engine data every 5 seconds. The fleet manager wants one Teams message when a van's engine runs hot for a sustained period, and to know when a van goes silent.

Options: `van_id`, One object for the whole fleet.

1. Object key
2. Hot-engine condition
3. Silent van

Sources

- Activator is a no-code event detection engine; events are grouped into objects by a shared identifier; rules are evaluated per object with state tracking; subsecond latency for stateless rules on streaming data.

(verified: <https://learn.microsoft.com/en-us/fabric/real-time-intelligence/data-activator/activator-introduction> (updated 2026-04-17))

- Conditions include Increases above, Decreases below, Is greater than, Becomes true/false, Changes and No presence of data; occurrence options are every time, n times, and true for a duration. (verified: <https://learn.microsoft.com/en-us/fabric/real-time-intelligence/data-activator/activator-detection-conditions> (updated 2026-05-04))
- Actions include email, Teams message (individual, group chat, channel), Power Automate and running Fabric items such as pipelines and notebooks. (verified: <https://learn.microsoft.com/en-us/fabric/real-time-intelligence/data-activator/rule-actions> (updated 2026-05-11))
- Up to 10,000 incoming events per second per rule, after which the rule stops; 30 email or Teams messages per rule per recipient per hour. (verified: <https://learn.microsoft.com/en-us/fabric/real-time-intelligence/data-activator/activator-limitations> (updated 2026-04-17))
- Late events are dropped by default; the wait time for late-arriving events defaults to two minutes and adds latency. (verified: <https://learn.microsoft.com/en-us/fabric/real-time-intelligence/data-activator/activator-latency> (updated 2026-04-30))

C8 Copilot and AI in Fabric

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Copilot across workloads, data agents, limits and verification.

C8-L1 Copilot in Fabric, where it runs and what it costs

Concept

Copilot in Fabric is an AI assistant built into each workload. It writes notebook code, pipelines, SQL, KQL and DAX, and summarises reports. Before you rely on it, check three things.

1. Is it switched on for you? Copilot needs a paid Fabric capacity of **F2 or higher**, or Power BI Premium **P1 or higher**, in a supported region. Fabric administrators control it through tenant settings, scoped to security groups. If your capacity's region lacks Azure OpenAI, an admin must also allow data to be processed outside that region. Copilot stops working when a Fabric trial expires.

2. What is its status? Release status differs by workload (Microsoft Learn, checked September 2026):

Experience	Status
Copilot in notebooks (Data Engineering, Data Science)	Preview
Data agents	GA
Copilot in Data Factory (Dataflow Gen2, pipelines)	GA
Copilot for SQL queries in a Warehouse	Preview
Copilot for semantic models and reports in Power BI	GA

3. What does it cost? Copilot requests consume capacity units, metered per 1,000 tokens. About 1,000 tokens equal 750 words. Input costs 100 CU seconds per 1,000 tokens, cached input 10 and output 400. A request with 2,000 input tokens and 500 output tokens costs $(2,000 \times 100 + 500 \times 400) / 1,000 = 400$ CU seconds. Copilot runs as a **background** operation, so Fabric smooths its cost over 24 hours.

On data handling, Microsoft states your data is not used to train models. It also states that Copilot output is sometimes inaccurate, and that people able to judge its accuracy should review it. Lesson C8-L3 shows you how.

Why it matters

An analytics team put Copilot-generated SQL into a scheduled report during the Warehouse Copilot preview. Nobody recorded which features were preview. When a reviewer asked whether the process relied on anything outside GA, the team spent a week tracing it. Check status first, record it, and keep production logic reviewed and owned by a person.

Worked example: weak code

1 Plan: use Copilot to write all transformation SQL

- ① Capacity: trial
- ② Review: none, Copilot output goes straight to the scheduled job
- ③ Cost estimate: Copilot is included

1. **Line 2.** Copilot stops working when the trial expires.
2. **Line 3.** Microsoft states output is sometimes inaccurate. Someone qualified must review it.
3. **Line 4.** Every request consumes CUs from your capacity.

Worked example: strong code

- ① Plan: Copilot drafts SQL, an engineer reviews, tests and owns it
- 2 Capacity: F8, Copilot tenant setting scoped to the data-engineering group
- ② Status check: Warehouse Copilot = preview (Learn, September 2026), noted in the design record
- ③ Cost estimate: about 400 CU seconds per request, smoothed over 24 hours

1. **Line 1.** Copilot drafts. A person decides.
2. **Line 3.** Preview features change. Record the status and the date.
3. **Line 4.** Estimated from the published token rates.

Confirm Copilot is enabled on a paid capacity, record the release status of each experience, and budget its CU use.

Recap card

- **Copilot in Fabric capacity requirement** Paid F2+ or P1+ in a supported region, with tenant settings enabled. Not after a trial expires.
- **Copilot release status** Per experience. Notebooks and Warehouse SQL preview. Data agents, Data Factory and Power BI GA (September 2026).
- **Copilot token rates per 1,000 tokens** Input 100, cached input 10, output 400 CU seconds. Background job, smoothed over 24 hours.
- **Copilot output** Sometimes inaccurate. A person able to judge it reviews it before use.

Practice

Exercise 1. Multiple choice (level 1)

Which capacity runs Copilot in Fabric?

- a) Any capacity, including an expired trial
- b) A paid F2 or higher, or P1 or higher, in a supported region
- c) Only F64 or higher
- d) No capacity: Copilot runs on the user's licence

Why?

- r1) Microsoft Learn lists a paid F2+ or P1+ capacity in a supported region, with tenant settings enabled.
- r2) Copilot is part of every Power BI Pro licence.

Exercise 2. Multiple choice (level 2)

Your design uses Copilot in notebooks and Copilot in Data Factory pipelines. What do you record about their status?

- a) Both GA
- b) Notebooks preview, Data Factory GA, with the date you checked
- c) Both preview
- d) Nothing: Copilot has one status across Fabric

Why?

- r1) Microsoft Learn's release-status page lists each Copilot experience separately.
- r2) Everything in Data Factory is GA.

Exercise 3. Fill the gap (level 2)

Complete the Copilot cost for 1,000 input tokens and 1,000 output tokens, with no cache.

```
1 input: 1,000 tokens x 100 / 1,000 = 100 CU seconds
2 output: 1,000 tokens x ____ (1) / 1,000 = ____ (2) CU seconds
3 total: 500 CU seconds
```

Exercise 4. Write the function (level 3)

Write `copilot_cu_seconds(input_tokens: int, output_tokens: int, cached_input_tokens: int = 0) -> float`. `cached_input_tokens` is the part of `input_tokens` served from the prompt cache. Rates per 1,000 tokens: input

100, cached input 10, output 400 CU seconds. Raise `ValueError` if any count is negative or if cached tokens exceed input tokens.

Starter code:

```
1 def copilot_cu_seconds(input_tokens, output_tokens,  
    cached_input_tokens=0):  
2     return (input_tokens + output_tokens) * 100 / 1000
```

Exercise 5. Architecture design (level 4)

Your organisation wants Copilot for its data-engineering team only, on an F16 capacity, with an audit trail of which Copilot features the pipelines depend on.

Options: Enabled for the data-engineering security group, Enabled for the whole organisation.

1. Tenant setting scope
2. Design record
3. Copilot output in production

Sources

- Copilot requires a paid F2+ or P1+ capacity in a supported region; tenant settings scoped to security groups; cross-geo processing settings; stops after trial expiry. (verified: <https://learn.microsoft.com/en-us/fabric/fundamentals/copilot-enable-fabric> (updated 2026-08-10); [troubleshoot-copilot-unavailable-after-trial](#) (updated 2026-07-10))
- Release status: notebooks preview, Data Factory GA, Warehouse SQL Copilot preview, Power BI semantic models and reports GA. Data agent GA per its concept page; the release-status page lists it in a Data Science row marked preview alongside AI functions. (verified: <https://learn.microsoft.com/en-us/fabric/fundamentals/copilot-ai-feature-state> (updated 2026-06-19); [data-science/concept-data-agent](#) (updated 2026-05-11))
- Copilot consumption: 100 CU seconds per 1,000 input tokens, 10 cached input, 400 output; about 1,000 tokens = 750 words; background operation smoothed over 24 hours; worked example 400 CU seconds. (verified: <https://learn.microsoft.com/en-us/fabric/fundamentals/copilot-fabric-consumption> (updated 2026-05-22))
- Data not used to train models; output is sometimes inaccurate and should be reviewed by people able to evaluate it. (verified: [https://](#))

C8-L2 Data agents, questions in plain language

Concept

A **Fabric data agent** (GA, Microsoft Learn, checked September 2026) answers questions in plain language over your data. It picks a data source, writes a query and runs it:

- **SQL** for lakehouses and warehouses.
- **DAX** for Power BI semantic models.
- **KQL** for KQL databases in an eventhouse.

One agent takes up to **five data sources**. Users do not write the query, but the query is still what decides the answer.

Three rules keep it safe:

- **Read-only.** The agent generates only read queries. It never creates, updates or deletes data, and it does not run notebooks or other write workflows.
- **The user's permissions.** The agent runs with the credentials of the person asking, so each user only reaches data they already have access to.
- **No unstructured files.** PDFs, Word documents and text files are not supported sources.

The agent is only as good as the context you give it. Microsoft's guidance:

- **Select only the tables and columns** the agent needs. Every extra table is another chance to pick the wrong one.
- **Use clear names** such as `order_submission_date`, not `date1`.
- **Write instructions** that explain business terms and which source answers which kind of question.
- **Add example queries:** question and query pairs. Each example is validated against the schema, and invalid ones are not used. Lakehouses, warehouses and KQL databases support examples. Semantic models do not.

Then **evaluate** the agent. Build a ground-truth set of questions with known answers, run them through the agent, and measure how many it gets right.

The `fabric-data-agent-sdk` package provides an evaluation API for this. Rerun the set after every change to instructions, examples or schema. Every query costs tokens. Instructions, example queries and conversation history all add to the input and so to CU use.

Why it matters

A sales team published a data agent over a lakehouse with 60 tables, including three with a column called `revenue`: gross, net and forecast. "What was revenue last quarter?" returned the forecast. Nobody noticed for a month. Selecting only the net-revenue table, plus an instruction defining revenue as net of returns, fixed it. A 20-question ground-truth test would have caught it on day one.

Worked example: weak code

- | | |
|---|---|
| 1 | Sources: <code>sales_lakehouse</code> (all 60 tables) |
| 2 | Instructions: (none) |
| 3 | Example queries: (none) |
| 4 | Testing: tried three questions in the chat, looked fine |

1. **Line 1.** Three revenue columns. The agent picks one.
2. **Line 4.** Looking fine is not a measured accuracy.

Worked example: strong code

- | | |
|---|---|
| 1 | Sources: <code>sales_lakehouse</code> (<code>fact_sales</code> , <code>dim_customer</code> , <code>dim_date</code> only) |
| 2 | Instructions: "Revenue" means <code>net_revenue_gbp</code> , after returns. Quarters are calendar quarters. |
| 3 | Example queries: 12 question and SQL pairs, all passing schema validation |
| 4 | Testing: 40 ground-truth questions, rerun after every change, accuracy recorded |

1. **Line 1.** Fewer tables, fewer wrong choices.
2. **Line 2.** Business terms defined once, for every question.
3. **Line 4.** A number you track, not a feeling.

Narrow the schema, define your terms, add validated examples, and measure accuracy against known answers.

Recap card

- **Data agent query languages** SQL for lakehouse and warehouse, DAX for semantic models, KQL for KQL databases. Up to five sources.
- **Data agent security** Read-only queries only. Runs with the asking user's permissions.

- **Example queries** Question and query pairs, validated against the schema. Not supported for semantic models.
- **Measure a data agent** Ground-truth questions with known answers, rerun after every change, accuracy recorded.

Practice

Exercise 1. Multiple choice (level 1)

A user asks a data agent to "delete the test orders". What happens?

- a) The agent writes and runs a DELETE statement
- b) The agent generates only read queries, so it does not delete data
- c) The agent deletes them if the user is a workspace Admin
- d) The agent asks the data owner for approval

Why?

- r1) Microsoft Learn states the data agent generates only read queries and keeps read-only connections.
- r2) The agent has no permissions of its own.

Exercise 2. Multiple choice (level 2)

Ana has access to the UK sales table but not the EU one. Both are sources of the same data agent. Ana asks for total EU sales. What does the agent do?

- a) Returns EU sales, because the agent's owner has access
- b) Answers only from data Ana has permission to read
- c) Returns EU sales with the numbers masked
- d) Returns a random estimate

Why?

- r1) The agent uses the requesting user's credentials and permissions.
- r2) The agent filters rows by country automatically.

Exercise 3. Put in order (level 2)

Order the steps to build a data agent over a sales lakehouse.

- Select only the tables and columns the questions need
- Add example question and SQL pairs
- Publish and share the agent
- Run the ground-truth question set and record accuracy
- Write instructions that define business terms

Exercise 4. Write the function (level 3)

Write

`evaluate(expected: dict[str, float], actual: dict[str, float | None], rel_tol: float = 0.001) -> tuple[float, list[str]]`. `expected` maps question ids to ground-truth answers. `actual` maps them to the agent's answers, where `None` means the agent gave no number. Return the share of questions answered correctly (within `rel_tol`, and exactly for an expected 0.0) and a sorted list of the failed ids. A question missing from `actual` fails. Raise `ValueError` if `expected` is empty.

Starter code:

```
1 def evaluate(expected, actual, rel_tol=0.001):
2     failed = [q for q in expected if not actual.get(q)]
3     return 1 - len(failed) / len(actual), failed
```

Exercise 5. Architecture design (level 4)

Configure a data agent for finance analysts. Sources: a finance warehouse and a Power BI semantic model of the same numbers. Contracts in PDF live in a lakehouse Files folder.

Options: Leave them out: unstructured files are not supported, Add the Files folder as a source.

1. The PDF contracts
2. Where to add example queries
3. Instruction to add

Sources

- Data agent is GA per its concept page; supports up to five sources (lakehouse, warehouse, semantic model, KQL database, and others); generates SQL, DAX, KQL; read-only; uses the requesting user's credentials; no unstructured data such as PDF, docx, txt. (verified: <https://learn.microsoft.com/en-us/fabric/data-science/concept-data-agent> (updated 2026-05-11); copilot-ai-feature-state (updated 2026-06-19))
- Select only needed tables and columns; use descriptive names; instructions and descriptions supply business meaning. (verified: <https://learn.microsoft.com/en-us/fabric/data-science/data-agent-configuration-best-practices> (updated 2026-08-27))
- Example queries are validated against the schema and supported for lakehouse, warehouse and KQL databases, not semantic models or

ontology. (verified: <https://learn.microsoft.com/en-us/fabric/data-science/data-agent-example-queries> (updated 2025-09-11))

- fabric-data-agent-sdk provides evaluation against ground-truth question sets. (verified: <https://learn.microsoft.com/en-us/fabric/data-science/evaluate-data-agent> (updated 2026-09-01))
- Instructions, example queries and conversation history add to token usage and CU consumption. (verified: <https://learn.microsoft.com/en-us/fabric/fundamentals/data-agent-consumption> (updated 2025-03-20))

C8-L3 Verify what Copilot writes

Concept

Microsoft's own guidance for Copilot in notebooks (preview, Microsoft Learn, checked September 2026) says: always review AI-generated code before you apply it. Output might be inaccurate, incomplete, or based on outdated library syntax. The chat pane shows changes as an approval diff. Approving the diff is your review, so make it a real one.

Code that runs without error proves only that the names exist and the syntax parses. Use four checks, in this order.

1. **Read every line.** Check each table, column and function against the schema and the docs. Check the grain: one row per what? Check joins: inner joins drop unmatched rows.
2. **Run it on known data.** Build five to ten rows where you have worked out the answer by hand. Include the awkward cases: a null, a zero denominator, a duplicate key, an empty group.
3. **Compare with a second method.** Compute the same number another way: a plain `count()` against the generated aggregate, SQL against PySpark, the new query against last month's report.
4. **Keep the checks.** Turn steps 2 and 3 into tests that run with the pipeline, so the next edit, by you or by Copilot, is checked too.

Errors to look for first, because they run cleanly and give plausible numbers:

- **Average of averages.** `AVG(avg_basket)` over stores weights a small store the same as a large one. Divide total sales by total baskets instead.
- **Wrong join type.** An inner join to a dimension drops facts with no match.
- `COUNT(*)` **for distinct things.** Customers with two orders count twice.
- **Unguarded division.** A zero denominator fails in some engines and returns null or infinity in others.

- **Invented names.** A column or function that does not exist. This one fails loudly.

Why it matters

An analyst asked Copilot for average basket value by region. The query ran and the numbers looked sensible. It averaged each store's average, so a pop-up store with four baskets counted as much as a flagship with 40,000. The regional figure was 11% too high. A five-row test with two stores of different sizes shows the difference at once.

Worked example: weak code

```

1 # Copilot suggestion, applied after it ran without error
2 region_basket = (
3     store_summary.join(dim_region, "store_id")
4     .groupBy("region")
5     .agg(F.avg("avg_basket").alias("avg_basket"))
6 )

```

1. **Line 1.** Running is not a check.
2. **Line 5.** Average of averages: every store weighs the same.

Worked example: strong code

```

1 from pyspark.sql import functions as F
2
3 region_basket = (
4     store_summary.join(dim_region, "store_id", "left")
5     .groupBy("region")
6     .agg(F.try_divide(F.sum("total_sales"),
7         F.sum("basket_count")).alias("avg_basket"))
8 )
9 # Known data: store A 4 baskets, 40.00 total. Store B 400 baskets,
10    2,000.00 total.
11 # Hand answer: 2,040.00 / 404 = 5.0495... Average of averages would give
12    7.50.
13 assert abs(result_for("North") - 2040.0 / 404) < 1e-9

```

1. **Line 6.** Total over total: each basket counts once. `try_divide` returns null for a zero denominator, where plain division fails under ANSI mode.
2. **Line 10.** Known data picked so the wrong method gives a visibly different answer.

Read it, run it on data with a hand-worked answer, compare it with a second method, and keep those checks as tests.

Recap card

- **A Copilot query ran without error. What does that prove?** Syntax parses and names exist. Nothing about correctness.
- **Four checks for Copilot output** Read every line. Run on known data. Compare with a second method. Keep the checks as tests.
- **Average of averages** Weights small groups like large ones. Divide total by total instead.
- **Known data for verification** Five to ten hand-built rows with a null, a zero denominator, a duplicate key and an empty group.

Practice

Exercise 1. Multiple choice (level 1)

A Copilot-generated query runs without an error. What has that shown?

- a) The results are correct
- b) The syntax parses and the referenced names exist
- c) The query is efficient
- d) The grain matches the requirement

Why?

- r1) An engine checks syntax and names before it runs a query. It does not know what answer you wanted.
- r2) Copilot tests its own queries before it shows them.

Exercise 2. Spot the bug (level 2)

Copilot wrote this T-SQL for average basket value by region. `store_summary` has one row per store with `total_sales`, `basket_count` and `avg_basket`. Tap the line that gives the wrong number, then fix it. Guard the division.

```
1 SELECT r.region,  
2        AVG(s.avg_basket) AS avg_basket  
3 FROM   dbo.store_summary AS s  
4 LEFT JOIN dbo.dim_region AS r ON s.store_id = r.store_id  
5 GROUP BY r.region;
```

Write down the line number of the bug. Then write the fix.

Exercise 3. Put in order (level 2)

Order the checks for a Copilot-generated transformation.

- Run it on a small dataset with a hand-worked answer
- Compare the result with a second method on real data

- Keep the known-data and comparison checks as tests
- Read every line: names, grain, join types

Exercise 4. Write the function (level 3)

Write `mismatches(candidate: Callable[[list[float]], float | None], reference: Callable[[list[float]], float | None], cases: list[list[float]]) -> list[int]`. Run both functions on each case. Return the indexes of cases where the candidate raises an exception, or where the results differ. Two `None` results match. Floats match within `math.isclose` defaults. Import `Callable` from `collections.abc`.

Starter code:

```
1 def mismatches(candidate, reference, cases):
2     return [i for i, case in enumerate(cases) if candidate(case) !=
           reference(case)]
```

Exercise 5. Architecture design (level 4)

Copilot wrote a PySpark step that deduplicates survey respondents and computes weighted reach per wave. Plan its verification before it goes into the silver pipeline.

Options: 8 rows: a duplicate respondent, a null weight, a wave with zero total weight, A 1% random sample of production.

1. Known dataset
2. Second method
3. After it passes

Sources

- Copilot for Data Engineering and Data Science is preview; always review AI-generated code; output might be inaccurate, incomplete or based on outdated library syntax; chat pane uses an approval diff. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/copilot-notebooks-overview> (updated 2026-04-01))
- With ANSI mode on, Spark division by zero raises `DIVIDE_BY_ZERO`; `F.try_divide` returns null instead. (verified: Local run on PySpark 4.x (practice venv), September 2026; error message names `try_divide`.)

C9 Governance and security

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Domains, OneLake security, sensitivity labels, Purview and lineage.

C9-L1 Domains and endorsement, finding the right data

Concept

A tenant with 400 workspaces needs two things before anyone trusts its data: a way to find data by business area, and a way to tell the authoritative copy from the rest. Fabric gives you **domains** and **endorsement**. Microsoft Learn shows neither with a preview label (checked September 2026).

Domains group workspaces by business area, such as Finance or Marketing. You can split a domain into **subdomains**.

- You associate **workspaces** with a domain. Every item in the workspace inherits the domain.
- Users filter the **OneLake catalog** by domain to find relevant data.
- Some tenant settings can be **delegated** to domain level, so each area sets its own rules.
- **Domain assignment does not grant or restrict access.** Workspace roles and item permissions decide who sees what.

Three roles manage domains. **Fabric admins** create domains and name domain admins. **Domain admins**, ideally the business owners, edit their domain and associate workspaces. **Domain contributors** are workspace admins allowed to assign their own workspaces to the domain.

Endorsement labels an item with one of three badges:

Badge	Meaning	Who applies it
Promoted	The creators think it is ready to share	Anyone with write access to the item
Certified	An authorised reviewer confirms it meets the organisation's standards	Only users a Fabric admin names
Master data	The authoritative source for core data such as customer lists	Only users a Fabric admin names

Any item except a Power BI dashboard can be promoted or certified. Master data applies only to items that hold data, such as lakehouses and semantic

models. Certification and master data work only after a Fabric admin enables them.

Endorsed items get a badge and appear first in some lists. That helps people pick the right copy. It does not stop anyone using an uncertified one.

Why it matters

A company had five semantic models called "Sales" across four workspaces. Two finance directors presented different quarterly revenue to the board. Both had used a model called "Sales". Certifying one model, retiring the others, and filing the finance workspaces under a Finance domain made the right model the first result anyone found.

Worked example: weak code

1	Workspaces: 400, no domains
1	Semantic models named "Sales": 5, none endorsed
2	Access plan: put finance workspaces in the Finance domain to restrict them

- 1. **Line 2.** Nothing tells users which copy is authoritative.
- 2. **Line 3.** Domains do not control access. Workspace roles do.

Worked example: strong code

1	Domains: Finance, Marketing, Operations (subdomain Finance > Tax)
2	Domain admins: business owners. Domain contributors: finance workspace admins
1	Certified: finance_sales_model (by a Fabric-admin-approved certifier)
4	Master data: crm_lakehouse, the single home of dim_customer
2	Access: workspace roles and OneLake security (lesson C9-L2)

- 1. **Line 3.** One authoritative model, labelled by an authorised reviewer.
- 2. **Line 5.** Discovery and access are separate controls.

Domains help people find data. Endorsement tells them which copy to trust. Neither controls access.

Recap card

- **What a domain does** Groups workspaces by business area for discovery and delegated settings. Does not control access.
- **Domain roles** Fabric admin creates domains. Domain admin edits one. Domain contributor, a workspace admin, assigns their workspaces.
- **Endorsement badges** Promoted (creators), Certified (named reviewers), Master data (named users, data items only). Dashboards get none.

- **Does endorsement block use of other copies?** No. It adds a badge and ranks the item first in some lists.

Practice

Exercise 1. Multiple choice (level 1)

You associate a workspace with the Finance domain. What changes for users outside finance?

- a) They lose access to the workspace
- b) Nothing about access. The workspace's items show under Finance in the OneLake catalog
- c) They gain read access to the items
- d) They must request access from the domain admin

Why?

- r1) Domain assignment does not affect item visibility or access. Workspace roles and item permissions do.
- r2) Domains only apply to Power BI items.

Exercise 2. Multiple choice (level 2)

A data engineer with Contributor access to a lakehouse wants to mark it as **Certified**. What happens?

- a) It works: Contributors can certify their items
- b) They can request certification. Only users a Fabric admin specifies can certify
- c) Certification is automatic once a data quality test passes
- d) Lakehouses cannot be certified

Why?

- r1) Any user can request certification, but only users specified by a Fabric administrator can certify.
- r2) Only workspace Admins can certify.

Exercise 3. Fill the gap (level 2)

Complete the domain roles.

- 1 Creates domains and names domain admins: Fabric ____ (1)
- 2 Assigns their own workspace to a domain: domain ____ (2), who must be a workspace admin

Exercise 4. Write the function (level 3)

Write `allowed_badges(item_type: str) -> set[str]` for the endorsement rules. Use the badge names "promoted", "certified" and "master_data". A "dashboard" gets no badges. Data items ("lakehouse", "warehouse", "semantic_model", "kql_database") get all three. Every other item type gets promoted and certified. Keep the data item types in a module constant.

Starter code:

```
1 def allowed_badges(item_type):
2     return {"promoted", "certified", "master_data"}
```

Exercise 5. Architecture design (level 4)

Organise a tenant where finance has 30 workspaces, three copies of the customer table, and complaints that nobody finds the right sales model.

Options: A Finance domain with finance workspaces associated, finance leads as domain admins, A naming convention only: prefix finance workspaces with fin_.

1. Grouping
2. The customer table
3. Keeping non-finance users out of finance data

Sources

- Workspaces associate with domains and items inherit; OneLake catalog filters by domain; tenant settings can be delegated; domain assignment doesn't affect visibility or access; roles Fabric admin, domain admin, domain contributor (must be workspace admin); subdomains. (verified: <https://learn.microsoft.com/en-us/fabric/governance/domains> (updated 2025-05-01))
- Badges Promoted, Certified, Master data; any item except dashboards can be promoted (by users with write permission) or certified; only admin-specified users certify or label master data; master data only for data items; must be enabled by a Fabric admin. (verified: <https://learn.microsoft.com/en-us/fabric/governance/endorsement-overview> (updated 2024-07-11))

C9-L2 OneLake security, roles down to rows and columns

Concept

OneLake security controls who reads which data in OneLake: tables, folders, schemas, rows and columns. Every supported engine enforces it (Microsoft Learn, checked September 2026).

Roles. Each role has a permission (Read or ReadWrite), the data it covers, and its members: Entra users or groups. The model is **deny by default** and supports only **grant** roles. There are no deny roles. A lakehouse starts with a **DefaultReader** role that gives users with the ReadAll permission access to its data. Remove or narrow it before your own roles mean anything.

Who it restricts. OneLake security governs **Viewers** and users with Read permission on the item. Workspace Admins, Members and Contributors already have data access through their workspace role. Make analysts Viewers.

Row and column security. A role can filter rows (RLS) and hide columns (CLS) on a table. Within one role, access is the **intersection** of its table, row and column rules. Across roles, access is the **union**: row filters combine with OR, so a user in roles for `city = 'Leeds'` and `city = 'York'` sees both cities. Two exceptions:

- In the **SQL analytics endpoint**, CLS uses deny semantics: if any of the user's roles hides a column, the endpoint blocks it.
- `SELECT *` on a table with hidden columns returns the allowed columns in Spark, but **fails** in the SQL analytics endpoint.

Keep row and column rules that belong together in **the same role**. OneLake rejects a combination where two roles allow different columns of a table and either role also filters its rows. Roles with ReadWrite cannot carry RLS or CLS.

Engines. RLS and CLS are GA in Spark notebooks, lakehouses, Direct Lake on OneLake semantic models, and the SQL analytics endpoint in **user's identity** access mode. Eventhouse supports RLS only, in preview.

Why it matters

A team added a OneLake security role limiting analysts to UK rows of `fact_sales`, then found the analysts still saw every country. Two causes. The analysts were Contributors on the workspace, so roles did not restrict them. And the DefaultReader role was still in place, granting full reads to every

Viewer with ReadAll. Making analysts Viewers and removing DefaultReader fixed both.

Worked example: weak code

- ❶ Analysts' workspace role: Contributor
- ❷ Roles: DefaultReader (unchanged), uk_sales (fact_sales where country = 'UK')
- ❸ Salary hiding: role hr_cols hides salary, role hr_rows filters team = 'Ops'

1. **Line 1.** Contributors already have data access. OneLake roles do not restrict them.
2. **Line 2.** DefaultReader still grants full reads. The union of roles gives everything.
3. **Line 3.** Row and column rules split across roles. Keep them in one role.

Worked example: strong code

- ❶ Analysts' workspace role: Viewer
- ❷ Roles: DefaultReader removed. uk_sales (fact_sales where country = 'UK')
- ❸ Salary hiding: role hr_ops (dim_employee where team = 'Ops', columns without salary)
- ❹ SQL analytics endpoint: user's identity access mode

1. **Line 1.** Viewers are the users OneLake security governs.
2. **Line 3.** Rows and columns in one role: the intersection applies.
3. **Line 4.** The endpoint enforces OneLake security only in this mode.

Readers are Viewers, defaults are removed, related row and column rules share a role, and every engine runs in a mode that enforces them.

Recap card

- **OneLake security combines roles how?** Within a role: intersection of table, row and column rules. Across roles: union. Row filters OR together.
- **Whom does OneLake security restrict?** Viewers and Read-permission users. Admins, Members and Contributors already have data access.
- **SELECT * with a hidden column** Spark returns allowed columns. SQL analytics endpoint returns an error.
- **Engines enforcing OneLake RLS and CLS (GA)** Spark, lakehouse, Direct Lake on OneLake, SQL analytics endpoint in user's identity mode. Eventhouse: RLS only, preview.

Practice

Exercise 1. Multiple choice (level 1)

A user belongs to two OneLake security roles on `fact_sales`: one filters `region = 'North'`, the other `region = 'South'`. Which rows do they see?

- a) North and South rows
- b) No rows: the filters contradict each other
- c) North rows only: the first role wins
- d) All rows

Why?

- r1) Roles combine as a union. Row filters from different roles combine with OR.
- r2) OneLake picks the most permissive single role.

Exercise 2. Multiple choice (level 2)

An analyst is a **Contributor** on the workspace. A OneLake security role limits them to UK rows. What do they see in Spark?

- a) UK rows only
- b) All rows: Contributors already have data access through their workspace role
- c) No rows
- d) An error

Why?

- r1) OneLake security governs Viewers and users with Read permission. Admins, Members and Contributors already have data access.
- r2) Spark ignores OneLake security.

Exercise 3. Multiple choice (level 2)

A table hides the `salary` column from a user. They run `SELECT * FROM dim_employee` in Spark, then in the SQL analytics endpoint. What happens?

- a) Spark returns the allowed columns. The endpoint returns an error
- b) Both return the allowed columns
- c) Both return an error
- d) Both return salary as null

Why?

- r1) Microsoft Learn documents that CLS behaves differently: Spark shows only allowed columns, the endpoint blocks the query.
- r2) Spark does not support column security.

Exercise 4. Write the function (level 3)

Write

`visible_columns(roles: list[dict[str, set[str]]], table: str, engine: Literal["spark", "sql_endpoint"]) -> set[str]`. Each role maps table names to the columns it allows. Consider only roles that include `table`. If none do, return an empty set (deny by default). For `"spark"`, return the union of the allowed columns. For `"sql_endpoint"`, return the intersection, because a column hidden by any role is blocked. Raise `ValueError` for any other engine.

Starter code:

```
1 def visible_columns(roles, table, engine):
2     columns = set()
3     for role in roles:
4         columns |= role[table]
5     return columns
```

Exercise 5. Architecture design (level 4)

HR analysts must see employee rows for their own region, without the `salary` column, in Spark and through a Direct Lake report.

Options: Viewer, Contributor.

1. Analysts' workspace role
2. Role design
3. Semantic model mode
4. DefaultReader role

Sources

- Role components (Read/ReadWrite, data, members); deny by default; grant roles only; DefaultReader role; default roles and ReadWrite matter only for Viewers since Admin/Member/Contributor have write access; ReadWrite roles cannot contain RLS/CLS. (verified: <https://learn.microsoft.com/en-us/fabric/onelake/security/data-access-control-model> (updated 2026-08-28); [get-started-security](https://learn.microsoft.com/en-us/fabric/onelake/security/get-started-security) (updated 2026-09-01))
- Within a role OLS $\cap$ RLS $\cap$ CLS; across roles union; RLS combines with OR; SQL analytics endpoint intersects CLS; roles with different columns plus RLS on the same table unsupported. (verified: <https://learn.microsoft.com/>)

en-us/fabric/onelake/security/data-access-control-model (updated 2026-08-28))

- **SELECT *** with hidden columns: Spark shows allowed columns, SQL analytics endpoint returns an error. (verified: <https://learn.microsoft.com/en-us/fabric/onelake/security/table-column-row-security> (updated 2026-08-19))
- RLS/CLS GA in lakehouse, Direct Lake on OneLake, Spark notebooks, SQL analytics endpoint in user's identity mode; Eventhouse RLS only in public preview; third-party engines preview. (verified: <https://learn.microsoft.com/en-us/fabric/onelake/security/read-secured-data> (updated 2026-08-28))

C9-L3 Sensitivity labels, Purview and lineage

Concept

Access control decides who opens an item. **Sensitivity labels** say how sensitive its content is, and travel with it. Fabric uses labels from **Microsoft Purview Information Protection**. You need a Purview licence, and labelling Power BI items also needs a Pro or Premium Per User licence (Microsoft Learn, checked September 2026).

How labels spread:

- **Downstream inheritance.** A label on a semantic model flows to the reports and dashboards built on it. It never overwrites a label someone applied manually, and never replaces a label with a less restrictive one.
- **Default and mandatory labelling.** Admins can apply a default label to new items, or require users to pick one before they save.
- **Export.** Labels, and the protection a Purview policy attaches, stay with data exported to Excel, PDF and PowerPoint.

Protection policies link a label to access. When an item carries the label, users and groups named in the policy keep the permissions they already have. Everyone else is blocked. The user who last applied the label is not blocked.

Purview also gives you:

- **Unified Catalog:** metadata about Fabric items, visible in Purview.
- **Data loss prevention (DLP):** detects sensitive data uploaded to OneLake and raises policy tips and alerts. It also restricts access to sensitive assets in KQL databases, SQL databases and warehouses.

- **Audit:** compliance admins track label activity.

Lineage shows how items connect. Every workspace has a lineage view that any user with a workspace role can open. It covers all items in the workspace, plus external sources one step upstream. **Impact analysis** starts from one item and lists the downstream items and workspaces a change might affect, and lets you notify their contacts.

Run impact analysis before a schema change. A label names the risk. A protection policy acts on it.

Why it matters

An engineer renamed a column in a lakehouse table on a Friday. Monday morning, 14 reports in three other workspaces showed errors, and a finance close was delayed. Impact analysis on the table would have listed all 14 reports and their workspaces, and notified their owners before the change.

Worked example: weak code

```
1 Change: rename cust_id to customer_id in fact_sales
❶ Checks: none, tested in the dev lakehouse notebook
❷ Label on hr_salaries model: none, "the workspace is private"
```

1. **Line 2.** Your notebook works. The 14 downstream reports break.
2. **Line 3.** Exports to Excel leave the workspace. A label and policy would travel with them.

Worked example: strong code

```
1 Change: rename cust_id to customer_id in fact_sales
❶ Checks: impact analysis lists 14 reports in 3 workspaces. Contacts notified. Views keep cust_id for one release
❷ Label on hr_salaries model: Highly Confidential, with a protection policy for the HR group
❸ Downstream: reports inherit Highly Confidential. Excel exports keep the label
```

1. **Line 2.** Know who breaks, tell them, and give them time.
2. **Line 3.** The policy blocks anyone outside the HR group, even if someone shares the item by mistake.
3. **Line 4.** Protection follows the data out of Fabric.

Run impact analysis before changes. Label sensitive items, attach protection policies, and let inheritance carry the label downstream.

Recap card

- **Downstream label inheritance rules** Flows to downstream items. Never overwrites a manual label. Never downgrades.
- **Protection policy** Tied to a label. Named users and groups keep existing permissions. Everyone else is blocked.
- **Lineage view** Every workspace has one. All items in the workspace plus external sources one step upstream.
- **Impact analysis** Downstream items and workspaces a change might affect, with a way to notify their contacts.

Practice

Exercise 1. Multiple choice (level 1)

A report built on a **Confidential** semantic model has a **Highly Confidential** label that someone applied by hand. The model's label changes to **General**. What happens to the report's label?

- a) It becomes General
- b) It stays Highly Confidential
- c) It becomes Confidential
- d) It is removed

Why?

- r1) Downstream inheritance never overwrites a manually applied label, and never downgrades a label.
- r2) Reports never inherit labels.

Exercise 2. Multiple choice (level 2)

You plan to change a column type in a lakehouse table. Which tool lists every downstream item and workspace that the change might break?

- a) Impact analysis
- b) The OneLake catalog domain filter
- c) Sensitivity labels
- d) The Capacity Metrics app

Why?

- r1) Impact analysis shows affected downstream items and workspaces and can notify their contacts.
- r2) Impact analysis blocks the change until owners approve.

Exercise 3. Fill the gap (level 2)

Complete the protection policy behaviour.

- 1 Item labelled Highly Confidential, policy allows the HR group.
- 2 HR group member with Viewer access: ____ (1) their Viewer access.
- 3 Finance user with Viewer access: ____ (2).

Exercise 4. Write the function (level 3)

Write `inherit_label(current: str | None, manual: bool, upstream: str, order: list[str]) -> str | None`. It returns the downstream item's label after its upstream item's label becomes `upstream`. `order` lists labels from least to most restrictive. Rules: a manual label is kept. An unlabelled item takes `upstream`. Otherwise the item takes `upstream` only if it is more restrictive than `current`. Raise `ValueError` if a label is not in `order`.

Starter code:

```
1 def inherit_label(current, manual, upstream, order):  
2     return upstream
```

Exercise 5. Architecture design (level 4)

Protect a payroll semantic model and its reports. Analysts outside HR sometimes receive shared links by mistake, and HR exports summaries to Excel.

Options: Highly Confidential, with reports inheriting it downstream, No label, the workspace is restricted.

1. Label on the payroll model
2. Access control for mis-shared links
3. Before renaming a payroll column

Sources

- Labels come from Purview Information Protection; Purview licence plus Pro or PPU needed to label Power BI items; labels persist on export to Excel, PDF, PowerPoint; default and mandatory labelling exist. (verified: <https://learn.microsoft.com/en-us/fabric/governance/information-protection> (updated 2026-07-13))
- Downstream inheritance never overwrites manual labels or downgrades to a less restrictive label. (verified: <https://learn.microsoft.com/en-us/fabric/governance/service-security-sensitivity-label-downstream-inheritance> (updated 2026-07-13))

- Protection policies: users in the policy retain existing permissions, others blocked; label issuer exempt. (verified: <https://learn.microsoft.com/en-us/fabric/governance/protection-policies-overview> (updated 2025-12-03))
- Purview Unified Catalog, DLP on OneLake uploads and access restriction for KQL, SQL database, warehouse; audit of label activity. (verified: <https://learn.microsoft.com/en-us/fabric/governance/microsoft-purview-fabric> (updated 2026-02-09))
- Every workspace has a lineage view accessible to any workspace role; shows workspace items and external sources one step upstream; impact analysis lists downstream items and workspaces and notifies contacts. (verified: <https://learn.microsoft.com/en-us/fabric/governance/lineage> (updated 2023-11-06); impact-analysis (updated 2023-05-23))

C10 DevOps and cost

Roadmap: Months 16 to 18: MLOps and deployment; Current work. Git integration, deployment pipelines and capacity monitoring.

C10-L1 Git integration for Fabric workspaces

Concept

Git integration connects a Fabric workspace to a branch and folder in a Git repository. Supported providers are Azure DevOps, GitHub and GitHub Enterprise, all cloud-based only. Some item types are still in preview, so check the supported-items list for each one you use (Microsoft Learn, checked September 2026).

What Git holds. Item definitions, not data. Each item gets its own folder of definition files. For a lakehouse, Git tracks metadata only: its name, shortcuts and OneLake security roles. Tables and files are not tracked, and Git operations never overwrite them. Rebuilding a workspace from Git gives you empty lakehouses. Your pipelines must reload the data.

Two directions:

- **Commit** exports workspace changes to the branch.
- **Update** imports branch changes into the workspace and overwrites the workspace version.

The source control panel shows each item's state. If an item changed in both places, you have a **conflict** to resolve before you continue.

Who does what. Only a workspace **Admin** connects or disconnects the workspace. Members and Contributors commit and update, limited by their permissions in the Git repository. Viewers see no Git information.

One workspace, one branch. A workspace is a shared runtime. Anything one person changes there affects everyone using it. Microsoft's recommended pattern is a **separate workspace per developer**, each connected to its own feature branch. You develop in your workspace, commit to your branch, and open a pull request into the main branch. The shared workspace connected to main then updates from Git.

Alternatively, develop in client tools such as VS Code for notebooks or Power BI Desktop for reports, and push to your branch.

Keep notebooks small and parameterised, and keep configuration in files your pipeline reads. Reviewers read diffs. A 900-line notebook with hard-coded paths produces diffs nobody can review.

Why it matters

Three engineers shared one development workspace connected to main. One of them updated from Git while another had an hour of uncommitted notebook edits. The update overwrote the workspace version, and the edits were lost. With a workspace and branch per developer, each person's work stays isolated until it merges.

Worked example: weak code

- ❶ Workspace: `sales_dev`, connected to `main`, used by 3 engineers
- ❷ Process: edit in `sales_dev`, commit to `main` when it works
- ❸ Recovery plan: restore the lakehouse from Git

1. **Line 1.** A shared runtime: one person's update overwrites another's work.
2. **Line 2.** No review before `main`.
3. **Line 3.** Git holds lakehouse metadata only. Tables and files are not in Git.

Worked example: strong code

- ❶ Workspaces: `sales_dev_ana` (branch `feature/ana-dedupe`), `sales_dev_raj` (branch `feature/raj-gold`)
- ❷ Process: commit to your feature branch, pull request into `main`, review, merge
- ❸ Shared: `sales_integration` connected to `main`, updated from Git after each merge
- ❹ Recovery plan: redeploy items from Git, then rerun ingestion pipelines to reload data

1. **Line 1.** One workspace per developer, one branch per workspace.
2. **Line 2.** Changes reach `main` through review.
3. **Line 4.** Items come from Git. Data comes from pipelines.

One workspace per developer on its own branch, pull requests into `main`, and a data reload plan, because Git holds definitions, not data.

Recap card

- **What Git integration stores for a lakehouse** Metadata only: name, shortcuts, OneLake security roles. Not tables or files.
- **Commit vs Update** Commit exports workspace changes to Git. Update imports Git changes and overwrites the workspace version.

- **Who connects a workspace to Git?** Admin only. Members and Contributors commit and update. Viewers see no Git information.
- **Recommended development process** One workspace per developer on its own branch. Pull requests into main.

Practice

Exercise 1. Multiple choice (level 1)

A lakehouse is connected to Git. What does the repository hold for it?

- a) The Delta tables and files
- b) Metadata only, such as its name, shortcuts and OneLake security roles
- c) Nothing: lakehouses are not supported
- d) A full backup, restorable to any commit

Why?

- r1) Microsoft Learn: only lakehouse metadata is tracked. Git operations never overwrite table or file data.
- r2) Git has a size limit that excludes tables.

Exercise 2. Multiple choice (level 2)

You click **Update** in a workspace. A notebook has an uncommitted edit in the workspace and a different change in Git. What does Fabric report?

- a) A conflict you must resolve
- b) Nothing: Git wins silently
- c) Nothing: the workspace wins silently
- d) The two versions merge automatically

Why?

- r1) An item changed in both the workspace and the branch is a conflict, shown in the source control panel.
- r2) Fabric merges notebooks cell by cell.

Exercise 3. Multiple choice (level 2)

Which workspace role connects a workspace to a Git repository?

- a) Admin
- b) Member
- c) Contributor
- d) Viewer

Why?

- r1) Connect and disconnect are Admin operations. Members and Contributors commit and update.
- r2) Only Admins have Git accounts.

Exercise 4. Write the function (level 3)

Your team's rule for the source control panel: resolve conflicts first, then commit your changes, then update from Git. Write `next_step(statuses: dict[str, str]) -> str`. Values are "synced", "uncommitted", "update_required" or "conflict". Return "resolve" if any item is in conflict, else "commit" if any is uncommitted, else "update" if any needs an update, else "none". Raise `ValueError` for an unknown status. Use a constant for the allowed statuses.

Starter code:

```
1 def next_step(statuses):
2     if "uncommitted" in statuses.values():
3         return "commit"
4     return "update"
```

Exercise 5. Architecture design (level 4)

Set up source control for a team of four engineers building notebooks and pipelines in Fabric.

Options: One workspace per engineer, each on its own feature branch, One shared dev workspace connected to main.

1. Development workspaces
2. Path to main
3. After a workspace is rebuilt from Git

Sources

- Providers Azure DevOps, GitHub, GitHub Enterprise (cloud only); some item types preview; workspace structure and subfolders preserved. (verified: <https://learn.microsoft.com/en-us/fabric/cicd/git-integration/intro-to-git-integration> (updated 2026-07-21))
- Commit exports and overwrites Git content; update overwrites workspace content; Admin connects; Member/Contributor commit and update per Git role; Viewer sees no Git info. (verified: <https://learn.microsoft.com/en-us/fabric/cicd/git-integration/git-integration-process> (updated 2026-06-15))

- A workspace connects to a single branch; recommended separate workspace per developer (branched workspace) or client tools. (verified: <https://learn.microsoft.com/en-us/fabric/cicd/git-integration/manage-branches> (updated 2026-06-15))
- Lakehouse Git tracks metadata only (name, shortcuts, data access roles); tables and files not tracked and never overwritten. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/lakehouse-git-deployment-pipelines> (updated 2026-05-08))
- A conflict occurs when the same item changes in both the workspace and the repository; Commit is disabled until it is resolved. (verified: <https://learn.microsoft.com/en-us/fabric/cicd/git-integration/conflict-resolution> (updated 2026-06-15))

C10-L2 Deployment pipelines, dev to test to production

Concept

Git integration (lesson C10-L1) versions your items. **Deployment pipelines** move them between environments. A pipeline has **2 to 10 stages**, with Development, Test and Production as the default three. Each stage is a workspace (Microsoft Learn, checked September 2026).

What a deployment copies. When you deploy from one stage to the next, Fabric copies item content and metadata to the target stage. It does not copy **data**. After deploying a semantic model, you refresh it in the target stage. After deploying a lakehouse, you load its tables there.

Overwrite, not replace. If the item already exists in the target stage, Fabric overwrites its content. The item's **ID, URL and permissions stay the same**, so reports and links in production keep working.

Deployment rules give each stage its own settings:

- **Data source rules:** point a semantic model at the production database, not the test one.
- **Parameter rules:** set a parameter's value per stage.
- **Default lakehouse rules:** attach a notebook to the stage's own lakehouse.

You define a rule in the **target** stage, on the item. It applies from the next deployment into that stage.

Autobinding. If a deployed item depends on another item, such as a report on a semantic model, and the linked item exists in the target stage, Fabric binds to the target copy automatically.

Compare shows which items differ between two stages before you deploy. Use it every time. The deployment pipelines REST APIs automate the same steps from a build system.

Many teams combine the tools. Developers work in Git-connected workspaces and merge through pull requests. The merged branch updates the Development stage, and deployment pipelines promote tested content to Test and Production. Microsoft's CI/CD guidance also describes Git-only release options, so pick one and document it.

Why it matters

A team deployed a notebook from Test to Production. It still pointed at the test lakehouse, so for a week the production job wrote its gold tables into Test. The dashboard in Production read stale tables and nobody noticed. A default lakehouse rule on the notebook in the Production stage would have attached it to the production lakehouse on every deployment.

Worked example: weak code

```
1 Stages: Development, Production
❶ Notebook nb_gold: default lakehouse lh_test (hard-coded)
3 Semantic model sales_model: data source test SQL database
❷ Before deploying: nothing
```

1. **Line 2.** Production runs against test data.
2. **Line 4.** No compare step: you learn what changed after it breaks.

Worked example: strong code

```
1 Stages: Development, Test, Production
❶ Rule in Production on nb_gold: default lakehouse = lh_prod
3 Rule in Production on sales_model: data source = production SQL database
❷ Before deploying: compare Test with Production, review the differing items
❸ After deploying: refresh sales_model, run the load pipeline
```

1. **Line 2.** Rules live in the target stage and apply on every deployment.
2. **Line 4.** Know what changes before it changes.
3. **Line 5.** Deployment copies items, not data.

Stage-specific settings belong in deployment rules. Compare before you deploy. Reload data after, because deployments copy items, not data.

Recap card

- **Deployment pipeline stages** 2 to 10 stages, default Development, Test, Production. One workspace per stage.
- **What a deployment copies** Item content and metadata. Not data: refresh models and reload tables after.
- **Deployment rule types** Data source, parameter, default lakehouse. Defined in the target stage on the item.
- **Deploying over an existing item** Content is overwritten. ID, URL and permissions stay the same.

Practice

Exercise 1. Multiple choice (level 1)

You deploy a semantic model from Test to Production for the first time. What must you do next?

- a) Nothing: the data is copied with it
- b) Refresh the model in Production, because data is not copied between stages
- c) Recreate the report permissions
- d) Reconnect the workspace to Git

Why?

- r1) Microsoft Learn: a semantic model refresh is required because data isn't copied from one stage to another.
- r2) Deployment pipelines always clear data in Production.

Exercise 2. Multiple choice (level 2)

Where do you define a rule that points a notebook at `lh_prod` in the Production stage?

- a) In the Production stage, on the notebook, as a default lakehouse rule
- b) In the Development stage, on the notebook
- c) In the notebook's code, with an if statement on the workspace name
- d) In the Git repository

Why?

- r1) Rules are defined in the target stage on the item, and apply from the next deployment into that stage.
- r2) Rules are global to the pipeline.

Exercise 3. Put in order (level 2)

Order a release from Test to Production.

- Check row counts and a known report figure in Production
- Check the deployment rules in the Production stage
- Compare Test with Production and review differences
- Deploy Test to Production
- Refresh semantic models and run load pipelines in Production

Exercise 4. Write the function (level 3)

Model deployment rules. Write `deployed_settings(source: dict[str, str], rules: dict[str, dict[str, str]], stage: str) -> dict[str, str]`. `source` holds an item's settings in the source stage. `rules` maps a stage name to the settings its rules override. Return the settings in `stage`: the source values, with that stage's rules applied. Do not modify `source`. Raise `ValueError` if a rule names a setting the item does not have.

Starter code:

```
1 def deployed_settings(source, rules, stage):
2     source.update(rules[stage])
3     return source
```

Exercise 5. Architecture design (level 4)

Design the release process for a team that uses Git integration and wants a Test stage signed off before Production.

Options: Workspace connected to main, updated from Git after each merged pull request, Engineers edit directly in the Development stage workspace.

1. Development stage
2. Promotion to Test and Production
3. After each Production deployment

Sources

- Deployment pipelines have 2 to 10 stages, default three. (verified: <https://learn.microsoft.com/en-us/fabric/cicd/deployment-pipelines/intro-to-deployment-pipelines> (updated 2026-07-17))
- Data isn't copied between stages, semantic model refresh required; existing items overwritten with ID, URL and permissions unchanged; autobinding to linked items in the target stage; REST APIs available. (verified: <https://learn.microsoft.com/en-us/fabric/cicd/deployment-pipelines/understand-the-deployment-process> (updated 2025-12-15))

- Rule types: data source, parameter, default lakehouse; defined in the target stage; apply from the next deployment. (verified: <https://learn.microsoft.com/en-us/fabric/cicd/deployment-pipelines/create-rules> (updated 2025-12-15))
- CI/CD options include Git-based release options and deployment pipelines; many organisations use a hybrid. (verified: <https://learn.microsoft.com/en-us/fabric/cicd/manage-deployment> (updated 2026-03-18))

C10-L3 Capacity, smoothing and throttling

Concept

Every Fabric operation draws capacity units (CUs) from its workspace's capacity (lesson C1-L1). Fabric lets operations **burst** above the capacity's size, then **smooths** the cost over future time (Microsoft Learn, checked September 2026).

Smoothing. Fabric measures usage in 30-second **timepoints**, 2,880 per day.

- **Interactive** operations, such as report queries, are smoothed over at least 5 minutes and up to 64 minutes.
- **Background** operations, such as scheduled Spark jobs, pipeline runs and model refreshes, are smoothed over **24 hours**.

Example: an F2 has 2 CUs, or 48 CU hours per day. A background job that uses 1 CU hour adds $1 / 48$, about 2.1%, to every timepoint for the next 24 hours.

Throttling starts when smoothed usage borrows too much future capacity:

Future capacity already used	What happens
Up to 10 minutes	Overage protection: no throttling
Over 10, up to 60 minutes	Interactive jobs are delayed by 20 seconds
Over 60 minutes, up to 24 hours	Interactive jobs are rejected
Over 24 hours	All requests are rejected, background too

Background jobs keep running well after users feel the delay. A capacity deep in debt takes hours to recover, because only idle capacity pays it back.

Surge protection lets capacity admins reject background jobs earlier, before the capacity reaches deep throttling. At workspace level, it caps a

workspace's share of CU use over a rolling 24 hours, unless you tag the workspace **Mission Critical**.

The **Capacity Metrics app** shows what happened. A capacity admin installs it and shares it. The **Health** page flags throttling and rejections across your capacities. The **Compute** page gives 14 days of utilisation by item and operation.

Practical rules: schedule heavy background jobs away from business hours, keep production on its own capacity, and check the Compute page before you buy a larger SKU.

Why it matters

A team scheduled a full rebuild of 40 gold tables at 08:00 on an F16. By 09:30 the capacity had borrowed more than an hour of future use, and Power BI reports started failing for every user. The rebuild itself kept running. Moving it to 02:00, and making it incremental, kept the morning reports responsive without a larger SKU.

Worked example: weak code

- ❶ Capacity: F16, shared by production reports and data science experiments
- ❷ Schedule: full rebuild of gold tables at 08:00 on weekdays
- ❸ Monitoring: users report slow dashboards

1. **Line 1.** Experiments draw from the same CUs as production.
2. **Line 2.** Heavy background work just before peak interactive use.
3. **Line 3.** You find out from users, after throttling has started.

Worked example: strong code

- 1 Capacity: F16 for production only. Experiments on a separate F4
- ❶ Schedule: incremental gold build at 02:00
- ❷ Surge protection: background rejection threshold set on the F16
- ❸ Monitoring: Capacity Metrics app Health page checked daily, Compute page reviewed weekly

1. **Line 2.** Less work, smoothed over hours when few users are active.
2. **Line 3.** Background jobs are stopped before users are rejected.
3. **Line 4.** You see throttling in the data before users feel it.

Know which of your jobs are background, run them off-peak, protect production with surge protection, and read the Metrics app before you resize.

Recap card

- **Smoothing windows** Interactive: 5 to 64 minutes. Background: 24 hours. Timepoints are 30 seconds, 2,880 a day.
- **Throttling stages** Up to 10 min: none. 10 to 60: 20 s interactive delay. 60 min to 24 h: interactive rejection. Over 24 h: all rejected.
- **Surge protection** Capacity level: earlier background rejection. Workspace level: cap a workspace's CU share over 24 hours, except Mission Critical.
- **Capacity Metrics app** Installed by a capacity admin. Health page for throttling. Compute page for 14 days of usage.

Practice

Exercise 1. Multiple choice (level 1)

Over what period does Fabric smooth the CU cost of a background Spark job?

- a) 24 hours
- b) 5 minutes
- c) The job's own duration
- d) It is not smoothed

Why?

- r1) Microsoft Learn: background operations are smoothed over a 24-hour period.
- r2) Spark jobs run for exactly 24 hours.

Exercise 2. Multiple choice (level 2)

A capacity has used 90 minutes of future capacity. What happens to a user opening a Power BI report, and to a scheduled pipeline?

- a) The report is rejected. The pipeline still runs
- b) Both run with a 20-second delay
- c) Both are rejected
- d) Both run normally

Why?

- r1) Between 60 minutes and 24 hours, interactive jobs are rejected. Background rejection starts only above 24 hours.
- r2) Pipelines always have priority.

Exercise 3. Fill the gap (level 2)

A background job uses 2 CU hours on an F4. Complete the smoothing calculation.

- ```
1 F4 capacity per day: 4 CUs x 24 hours = ____ (1) CU hours
2 Share of every timepoint for 24 hours: 2 / ____ (1) = about ____ (2)%
```

### Exercise 4. Write the function (level 3)

Write `throttle_stage(future_minutes: float) -> str` returning the throttling stage for minutes of future capacity already used: "none" up to 10, "interactive\_delay" above 10 up to 60, "interactive\_rejection" above 60 up to 1,440 (24 hours), and "background\_rejection" above that. Raise `ValueError` for negative input. Keep the limits in constants.

Starter code:

```
1 def throttle_stage(future_minutes):
2 if future_minutes < 10:
3 return "none"
4 if future_minutes < 60:
5 return "interactive_delay"
6 return "interactive_rejection"
```

### Exercise 5. Architecture design (level 4)

Your F32 throttles every weekday between 08:30 and 10:00. The Compute page shows a 06:00 full refresh of a large semantic model and data science notebooks running all day on the same capacity.

Options: Make it incremental and move it to 01:00, Keep it: refreshes are background and smoothed.

1. The 06:00 full refresh
2. Data science notebooks
3. Protection for report users

### Sources

- Timepoints 30 s, 2,880 per day; interactive smoothed 5 to 64 minutes; background 24 hours; F2 1 CU hour example ~2.1%; throttling stages at 10 min, 60 min, 24 h with 20 s interactive delay. (verified: <https://learn.microsoft.com/en-us/fabric/enterprise/throttling> (updated 2026-08-14))
- Surge protection: capacity-level background rejection threshold; workspace-level CU cap over rolling 24 h with Mission Critical exclusion.

(verified: <https://learn.microsoft.com/en-us/fabric/enterprise/surge-protection> (updated 2026-05-27))

- Capacity Metrics app installed by capacity admin; Health page; Compute page 14-day view. (verified: <https://learn.microsoft.com/en-us/fabric/enterprise/metrics-app> (updated 2026-01-14))

## Track C best-practice checklist

- ☐ Check the status (GA or preview) of every Fabric feature before you depend on it.
- ☐ One copy of each dataset. Use shortcuts or mirroring before you copy.
- ☐ Lakehouse for Spark and files. Warehouse for T-SQL writes and multi-table transactions.
- ☐ The SQL analytics endpoint is read-only for table data. Write through Spark or a Warehouse.
- ☐ Semantic models use star schemas. DAX measures state their filter context.
- ☐ Direct Lake models are checked for fallback to DirectQuery.
- ☐ Every Copilot output is read, run on known data and compared against a second method.
- ☐ Workspaces map to domains. Access uses roles and OneLake security, not item sharing.
- ☐ Items live in Git. Changes move through deployment pipelines.
- ☐ Capacity use is monitored, and heavy jobs are scheduled off peak.

## Track D: Finding errors in code

Find and fix errors fast, including errors in code that Claude and Copilot write.

# D1 Reading tracebacks and stack traces

Roadmap: Every month. Find the root cause in Python and Spark errors.

## D1-L1 Read a Python traceback from the bottom up

### Concept

A traceback lists the calls that led to an error, oldest first. The line that matters most is at the **bottom**. Read it in this order:

1. **The last line** names the exception type and message: `NameError: name 'uuid' is not defined`. That is what went wrong.
2. **The frame just above it** shows where: file, line number, function, and the source line. Python 3.11 and later add `^` or `~` marks under the exact expression.
3. **Walk up** until you reach the first frame in **your own code**. Frames inside `site-packages` or the standard library are rarely the cause. The call from your code that fed them bad input usually is.

Here is your reference `log_pipeline_run`, cut down, on Python 3.13:

```
1 Traceback (most recent call last):
2 File "observability.py", line 7, in <module>
3 main()
4 ~~~~^^
5 File "observability.py", line 5, in main
6 log_run()
7 ~~~~~~^^
8 File "observability.py", line 2, in log_run
9 return str(uuid.uuid4())
10 ~~~~~^
11 NameError: name 'uuid' is not defined. Did you forget to import 'uuid'?
```

Python 3.12 added the "Did you forget to import" hint. Python 3.11 prints the same error without it.

**Chained exceptions** print two tracebacks. The joining sentence tells you how they relate:

- **"The above exception was the direct cause of the following exception"** means someone wrote `raise NewError(...) from err`. The first traceback is the root cause.
- **"During handling of the above exception, another exception occurred"** means a second error happened inside an `except` block. Your error handler has its own bug. Fix both.



When you catch an exception to add context, use `raise ... from err`. The original traceback stays attached, so the next person sees both.

Never read only the first line of a traceback, and never read only the message. `KeyError: 'row_count_max'` names the missing key. The frame tells you which dictionary, and so which config file, lacks it.

## Why it matters

A contract validator failed with a 60-line traceback. The on-call engineer read the top frames, which pointed at the job scheduler, and restarted the job three times. The last line said `KeyError: 'row_count_max'`, and the frame above it named `_check_row_count`. The new tracker's contract file lacked one key. Reading from the bottom would have found it in a minute.

## Worked example: weak code

```
1 def load_contract(path: str) -> dict[str, object]:
2 """Load a contract file."""
3 try:
4 with open(path) as f:
5 return json.load(f)
6 except Exception:
7 raise ValueError("bad contract")
```

1. **Line 6.** Catches everything, including a typo in your own code.
2. **Line 7.** No path, no cause. The traceback prints 'During handling...', and the reader must guess which error was the real one.

## Worked example: strong code

```
1 import json
2
3
4 def load_contract(path: str) -> dict[str, object]:
5 """Load a contract file.
6
7 Raises:
8 ValueError: If the file is not valid JSON.
9 """
10 try:
11 with open(path, encoding="utf-8") as f:
12 return json.load(f)
13 except json.JSONDecodeError as err:
14 raise ValueError(f"contract {path} is not valid JSON: line {err.lineno}") from err
```

1. **Line 13.** Catch only the error you expect. A missing file still raises `FileNotFoundError` with its own traceback.

2. **Line 14.** The message names the file and line. `from err` keeps the original traceback as the direct cause.

Read the last line, then the frame above it, then walk up to your own code. When you re-raise, add context and chain with `from err`.

## Recap card

- **Order for reading a traceback** Last line (what), frame above it (where), then up to the first frame in your own code.
- **'The above exception was the direct cause...'** Someone chained with `raise ... from err`. The first traceback is the root cause.
- **'During handling of the above exception...'** The except block raised a new error. Fix the handler and the original error.
- **Re-raising with context** Catch the narrow exception, add facts to the message, and chain with `from err`.

## Practice

### Exercise 1. Multiple choice (level 1)

A traceback is 40 lines long. Where do you start reading?

- a) The last line: exception type and message
- b) The first line after 'Traceback (most recent call last)'
- c) The longest frame
- d) The first frame inside site-packages

### Why?

- r1) Frames run oldest first. The exception and the frame that raised it are at the bottom.
- r2) Python prints the most important frame last to save space.

### Exercise 2. Spot the bug (level 2)

This traceback came from Python 3.13. Tap the frame line in **your own code** where you would start fixing. Your code is in `/repo/pipelines/`.

```
1 Traceback (most recent call last):
2 File "/repo/pipelines/run_qa.py", line 31, in <module>
3 report = check_weights(df, "final_weight")
4 File "/repo/pipelines/qa_statistical.py", line 19, in check_weights
5 "mean": round(stats["mean"], 4),
6 ~~~~~^~~~~~
7 TypeError: type NoneType doesn't define __round__ method
```

Write down the line number of the bug.

### Exercise 3. Multiple choice (level 2)

A traceback shows a `KeyError`, then the sentence "During handling of the above exception, another exception occurred", then a `NameError`. What does that tell you?

- a) The except block that handled the `KeyError` has its own bug
- b) The code used `raise ... from err`
- c) Only the `NameError` matters
- d) Python retried the code

#### Why?

- r1) The second exception was raised while handling the first, without an explicit chain. Both need a fix.
- r2) `NameError` always hides a `KeyError`.

### Exercise 4. Write the function (level 3)

Write `first_own_frame(traceback_text: str, project_dir: str) -> tuple[str, int] | None`. Return the file path and line number of the **last** frame line (`File "<path>", line <n>, in <name>`) whose path starts with `project_dir`. That is the first frame in your code when you read from the bottom. Return `None` if no frame is in your code. For chained tracebacks, use the last traceback in the text.

Starter code:

```
1 def first_own_frame(traceback_text, project_dir):
2 for line in traceback_text.splitlines():
3 if project_dir in line:
4 return line
```

### Exercise 5. Refactor (level 4)

Refactor `row_count_bounds` so a missing key raises `ValueError` naming the contract id and the missing key, chained to the original `KeyError` with `from`. Add type hints and a docstring.

Starter code:

```
1 def row_count_bounds(contract):
2 try:
3 return contract["row_count_min"], contract["row_count_max"]
4 except Exception:
5 raise ValueError("bad contract")
```

## Sources

- Python 3.11+ prints caret or tilde markers under the failing expression; 3.12 and 3.13 add 'Did you forget to import' to NameError for standard modules; 3.11 does not. (verified: Local runs on CPython 3.11.15, 3.12.3 and 3.13.12, September 2026.)
- raise ... from err prints 'The above exception was the direct cause of the following exception'; an error inside an except block prints 'During handling of the above exception, another exception occurred'. (verified: Local runs on CPython 3.11.15 and 3.13.12, September 2026.)
- round(None, 4) raises TypeError: type NoneType doesn't define **round** method. (verified: Local run on CPython 3.13.12, September 2026.)

## D1-L2 Spark errors, the condition and the real cause

---

### Concept

A Spark error message is long, but it has a fixed shape. In PySpark 4.0 (Databricks Runtime 17.3 LTS) most errors start with an **error condition** in square brackets and end with a **SQLSTATE**:

```
1 [UNRESOLVED_COLUMN.WITH_SUGGESTION] A column, variable, or function
2 parameter with name `final_wieght` cannot be resolved. Did you mean one
3 of the following? [`final_weight`, `age`, `wave_id`, `respondent_id`].
4 SQLSTATE: 42703
```

Read three things:

- **The condition** names the problem. Search for it in the Databricks error conditions reference.
- **The SQLSTATE class**, its first two characters. Class **42** is a syntax error or access rule violation: your code names something that does not exist. Class **22** is a data exception: the code is valid but a value is not.
- **The origin line**. Runtime errors show `== DataFrame ==` and `"cast" was called from ...` with your file and line.

**When the error appears.** Transformations are lazy, so a data error appears at the **action** (`collect`, `count`, `write`), not at the line that caused it. The origin line points you back. In classic Spark, a missing column fails at once, when you call `select`.

**Python UDF errors** arrive as `PythonException` wrapping a normal Python traceback. Read that inner traceback bottom up (lesson D1-L1).

**Shadowed built-ins** give a misleading message. Your `data_contracts.py` imports `sum` from `pyspark.sql.functions`. Then `sum(1 for c in checks)` raises `[NOT_COLUMN_OR_STR] Argument col should be a Column or str, got generator`. The code never meant to call Spark. Import functions as `F`.

**ANSI mode** is on by default in Spark 4.0 and Databricks Runtime 17.0 and above. A cast of `'abc'` to `INT` now raises `[CAST_INVALID_INPUT]` instead of returning null. Use `try_cast` where bad values are expected, and count the nulls it produces.

In code, catch `pyspark.errors.AnalysisException` and read `err.getCondition()` and `err.getSqlState()`, not the message text.

## Why it matters

After a runtime upgrade to Databricks Runtime 17.3 LTS, a silver job that had run for two years failed with

`[CAST_INVALID_INPUT] The value 'n/a' ... cannot be cast to "INT"`. The engineer's first thought was a Spark bug. The `SQLSTATE` started with 22, a data exception: the data had always held `'n/a'`, and the old runtime had turned it into null without a word. The new error exposed a data quality problem that was two years old.

## Worked example: weak code

```
❶ from pyspark.sql.functions import col, sum
2
❷ silver = bronze.withColumn("age", col("age").cast("int"))
4 passed = sum(1 for c in checks if c.status == "PASS")
5 silver.write.mode("append").saveAsTable("research_silver.nat_respondents")
```

1. **Line 1.** Shadows Python's `sum`. The next line fails with `NOT_COLUMN_OR_STR`.
2. **Line 3.** Under ANSI mode, one bad value fails the whole write, several lines later.

## Worked example: strong code

```
❶ from pyspark.sql import functions as F
2
❷ silver = bronze.withColumn("age", F.col("age").try_cast("int"))
❸ bad_ages = bronze.filter(F.col("age").isNotNull() &
 F.col("age").try_cast("int").isNull()).count()
5 passed = sum(1 for c in checks if c.status == "PASS")
6 silver.write.mode("append").saveAsTable("research_silver.nat_respondents")
```

1. **Line 1.** Python's `sum` stays Python's `sum`.
2. **Line 3.** Bad values become null on purpose.
3. **Line 4.** Then you count them, so nothing disappears silently.

Read the condition, the SQLSTATE class and the origin line. Class 42 means fix the code. Class 22 means look at the data.

## Recap card

- **Shape of a Spark 4 error message** [ERROR\_CONDITION] message ... SQLSTATE: xxxxx, often with an origin line under == DataFrame ==.
- **SQLSTATE class 42 vs 22** 42: syntax error or access rule violation, fix the code. 22: data exception, look at the data.
- **Why a data error appears at write()** Transformations are lazy. The action runs the plan. The origin line shows where the expression was written.
- **NOT\_COLUMN\_OR\_STR ... got generator** Python's sum is shadowed by pyspark.sql.functions.sum. Import functions as F.

## Practice

### Exercise 1. Multiple choice (level 1)

A job fails with [UNRESOLVED\_COLUMN.WITH\_SUGGESTION] ...

namefinal\_wieghtcannot be resolved. Did you mean one of the following?

[final\_weight, ...]. SQLSTATE: 42703. What do you fix?

- a) The column name in the code: final\_wieght is a typo for final\_weight
- b) The data: some rows lack final\_weight
- c) The cluster: restart it
- d) The table permissions

### Why?

- r1) Class 42 errors mean the code names something that does not exist. The message suggests the right name.
- r2) Spark always suggests the column you meant.

### Exercise 2. Multiple choice (level 2)

pass\_rate in your data\_contracts.py runs sum(1 for c in self.checks if c.status == "PASS") and fails with [NOT\_COLUMN\_OR\_STR] Argument col should be a Column or str, got generator. What is the cause?

- a) sum is PySpark's sum, imported by name, which shadows Python's built-in
- b) self.checks is a DataFrame
- c) Generators are not allowed in dataclasses

- d) The status column is missing

## Why?

- r1) `from pyspark.sql.functions import sum` replaces the built-in in that module, so the generator goes to Spark's sum.
- r2) PySpark disables Python built-ins in notebooks.

## Exercise 3. Spot the bug (level 3)

This job fails on line 5 with `[CAST_INVALID_INPUT]` The value `'n/a'` of the type `"STRING"` cannot be cast to `"INT"` and `"cast"` was called from ... line 3. Tap the line that causes it, then fix it so bad ages become null. Keep functions as `F`.

```
1 from pyspark.sql import functions as F
2
3 silver = bronze.withColumn("age", F.col("age").cast("int"))
4 silver = silver.withColumn("wave_id", F.lit(wave_id))
5 silver.write.mode("append").saveAsTable("research_silver.nat_respondents")
```

Write down the line number of the bug. Then write the fix.

## Exercise 4. Write the function (level 3)

Write `trriage(message: str) -> tuple[str | None, str]` for Spark error messages. Extract the error condition, the text in the first `[...]` at the start of the message, or `None`. Then classify by the SQLSTATE at the end (`SQLSTATE: 42703`): class `"42"` gives `"code"`, class `"22"` gives `"data"`, anything else or no SQLSTATE gives `"unknown"`.

Starter code:

```
1 def triage(message):
2 if "42" in message:
3 return None, "code"
4 return None, "data"
```

## Exercise 5. Architecture design (level 4)

Write the first-response steps for your team's runbook when a Spark job fails.

Options: The error condition, the SQLSTATE and the origin line, The Java stack trace from the top.

1. First thing to read
2. SQLSTATE class 22 after a runtime upgrade
3. Handling errors in code

## Sources

- PySpark 4.0.4 error texts: UNRESOLVED\_COLUMN.WITH\_SUGGESTION with SQLSTATE 42703; CAST\_INVALID\_INPUT with SQLSTATE 22018 and an origin line; PythonException wrapping a Python traceback; NOT\_COLUMN\_OR\_STR for a generator passed to shadowed sum; getCondition() and getSqlState() on AnalysisException; Column.try\_cast exists. (verified: Local runs on PySpark 4.0.4, September 2026.)
- ANSI mode is enabled by default in Apache Spark 4.0 and Databricks Runtime 17.0 and above. (verified: <https://docs.databricks.com/aws/en/sql/language-manual/sql-ref-ansi-compliance> (checked 2026-09-25))
- SQLSTATE class 22 is data exception; class 42 is syntax error or access rule violation. (verified: <https://docs.databricks.com/aws/en/error-messages/sqlstates> (checked 2026-09-25))

## D1-L3 Py4J and Spark Connect errors

---

### Concept

PySpark talks to the Spark engine in one of two ways, and each shapes the errors you see. Facts below were checked on PySpark 4.0.4 and the Databricks docs in September 2026.

**Spark Classic** runs your Python next to a JVM driver through **Py4J**. PySpark converts most engine errors into Python classes such as `AnalysisException`. Some still arrive raw, as `py4j.protocol.Py4JJavaError`:

```
1 Py4JJavaError: An error occurred while calling
2 z:org.apache.spark.api.python.PythonRDD.collectAndServe.
3 : org.apache.spark.SparkException: Job aborted due to stage failure ...
4 Caused by: org.apache.spark.api.python.PythonException: Traceback ...
5 File "/repo/job.py", line 5, in <lambda>
6 ZeroDivisionError: division by zero
```

Skip the Java frames. Search for the **last** `Caused by:` line and the Python traceback inside it. That is your bug.

**Spark Connect** sends a query plan to a remote server over gRPC. Databricks **serverless compute uses Spark Connect**, and supports only Spark Connect APIs. Three differences matter when you debug:

- **Analysis is deferred.** Classic Spark rejects a misspelt column when you call `select`. Spark Connect accepts the `select` and raises `AnalysisException` at the action, or when you read `df.columns` or `df.schema`. A try around the `select` catches nothing.



- **Error classes live in** `pyspark.errors.exceptions.connect`. They subclass `SparkConnectGrpcException` and the familiar base classes, so `except pyspark.errors.AnalysisException` still works. The message ends with a `JVM stacktrace:` section from the server.
- **JVM-only attributes fail.** `spark.sparkContext` raises `[JVM_ATTRIBUTE_NOT_SUPPORTED]`. RDD code, such as `sc.parallelize`, does not run on serverless.

Two habits follow. Catch `pyspark.errors` classes, never `Py4JJavaError`, so the same code works on both. And place error handling around the **action**, where both modes raise, not around the transformation.

## Why it matters

A notebook wrapped `df.select(col(name))` in `try/except AnalysisException` to skip columns missing from older waves. It worked on a classic cluster. After the team moved the job to serverless, the `except` never fired: the `select` succeeded, and the `write` failed with the same `AnalysisException` two cells later. Checking `name in df.columns` before the `select` works in both modes.

## Worked example: weak code

```

1 from py4j.protocol import Py4JJavaError
2 from pyspark.sql import functions as F
3
4 try:
5 selected = df.select(F.col(optional_column))
6 except Py4JJavaError:
7 selected = df
8 selected.write.mode("append").saveAsTable(target_table)

```

1. **Line 1.** `Py4J` exists only in Spark Classic. Serverless has no `Py4J`.
2. **Line 5.** Under Spark Connect this line never raises. The error arrives at the `write`.

## Worked example: strong code

```

1 from pyspark.errors import AnalysisException
2 from pyspark.sql import functions as F
3
4 columns = [c for c in (key_column, optional_column) if c in df.columns]
5 selected = df.select(*[F.col(c) for c in columns])
6 try:
7 selected.write.mode("append").saveAsTable(target_table)
8 except AnalysisException as err:

```

```
9 raise RuntimeError(f"write to {target_table} failed:
 {err.getCondition()}") from err
```

1. **Line 4.** Check the schema explicitly. Works the same in Classic and Connect.
2. **Line 8.** A `pyspark.errors` class, caught around the action, where both modes raise.

Find the last "Caused by" in a Py4J error. Under Spark Connect, expect errors at the action, and catch `pyspark.errors` classes.

## Recap card

- **Finding the cause in a Py4JJavaError** Search for the last 'Caused by:'. A Python UDF error holds a Python traceback there.
- **Spark Connect and missing columns** Analysis is deferred. The error comes at the action or when you read `df.columns` or `df.schema`.
- **Exception class that works on Classic and Connect** `pyspark.errors.AnalysisException` (and other `pyspark.errors` classes). Not `Py4JJavaError`.
- **Serverless compute and the Spark API** Spark Connect only. No RDD APIs. `spark.sparkContext` raises `JVM_ATTRIBUTE_NOT_SUPPORTED`.

## Practice

### Exercise 1. Multiple choice (level 1)

A raw `Py4JJavaError` fills 200 lines, mostly Java frames. Where is the cause of a Python UDF failure?

- a) After the last 'Caused by:', in the Python traceback it contains
- b) In the first Java frame
- c) In the line 'An error occurred while calling'
- d) It is not shown. You must read the executor logs

### Why?

- r1) The JVM wraps the Python worker's exception. The innermost cause holds the Python traceback.
- r2) Java always prints the cause first.

### Exercise 2. Predict the output (level 2)

This runs on **serverless** (Spark Connect). The DataFrame has no column `region`. What does it print?

```
1 from pyspark.errors import AnalysisException
2 from pyspark.sql import functions as F
```

```

3
4 try:
5 out = df.select(F.col("region"))
6 print("select ok")
7 except AnalysisException:
8 print("caught")

```

### Exercise 3. Multiple choice (level 2)

Which `except` clause catches a missing-column error on both a classic cluster and serverless?

- a) `except pyspark.errors.AnalysisException`
- b) `except py4j.protocol.Py4JJavaError`
- c) `except pyspark.errors.exceptions.connect.SparkConnectGrpcException`
- d) `except KeyError`

#### Why?

- r1) Both modes raise a subclass of `pyspark.errors.AnalysisException`.
- r2) `AnalysisException` is a `Py4J` class.

### Exercise 4. Write the function (level 3)

Write `python_cause(message: str) -> str | None`. Given the text of a `Py4J` error, return the last line of the Python traceback in the **last** `Caused by:` `org.apache.spark.api.python.PythonException` section: the first non-blank line that does not start with whitespace, `Traceback` or `File`, after that marker. Return `None` if there is no such section.

Starter code:

```

1 def python_cause(message):
2 return message.splitlines()[-1]

```

### Exercise 5. Architecture design (level 4)

A notebook written for a classic cluster moves to serverless. Plan the changes to its error handling.

Options: Check `name in df.columns` before selecting, Wrap each select in `try/except`.

1. Optional columns in older waves
2. Exception class to catch
3. A helper that uses `sc.parallelize`

## Sources

- Classic: raw Py4JJavaError for RDD Python failures with 'Caused by: org.apache.spark.api.python.PythonException' holding the Python traceback; eager analysis at select. Connect: analysis deferred to action or schema access; errors in `pyspark.errors.exceptions.connect` subclass `SparkConnectGrpcException` and are caught by `pyspark.errors.AnalysisException`; message includes 'JVM stacktrace:'; `spark.sparkContext` raises `JVM_ATTRIBUTE_NOT_SUPPORTED`. (verified: Local runs on PySpark 4.0.4, classic `local[1]` and `remote('local[1]')`, September 2026.)
- Serverless compute supports only Spark Connect APIs, not RDD APIs, and defers analysis and name resolution to execution time. (verified: <https://docs.databricks.com/aws/en/compute/serverless/limitations> and <https://docs.databricks.com/aws/en/spark/connect-vs-classic> (checked 2026-09-25))

# D2 Debugging tools

Roadmap: Every month. breakpoint(), pdb, logging levels, Spark UI and explain().

## D2-L1 breakpoint() and pdb, stop and look

### Concept

A traceback shows where code failed. A **debugger** shows the values at that moment. Python ships one: **pdb**.

**Stop at a line.** Put `breakpoint()` in your code. When Python reaches it, pdb opens a prompt, `(Pdb)`, with every local variable in scope. Setting the environment variable `PYTHONBREAKPOINT=0` turns every `breakpoint()` call into a no-op, so a forgotten one cannot hang a scheduled job. Remove them before you commit anyway.

**Stop where it crashed.** Run a script under pdb with post-mortem on:

```
1 python -m pdb -c continue run_qa.py
2 ...
3 ZeroDivisionError: float division by zero
4 Uncaught exception. Entering post mortem debugging
5 > /repo/run_qa.py(2)weighted_pct()
6 -> ratio = reach / base
7 (Pdb) p reach, base
8 (5.0, 0.0)
```

You land in the frame that raised, with its variables intact. Checked on Python 3.13.

The commands you use most:

| Command                                    | Does                                   |
|--------------------------------------------|----------------------------------------|
| <code>p expr</code> , <code>pp expr</code> | Print, or pretty-print, a value        |
| <code>w</code>                             | Show the stack, newest frame last      |
| <code>u</code> , <code>d</code>            | Move up to the caller, or down again   |
| <code>n</code>                             | Run the next line, stepping over calls |
| <code>s</code>                             | Step into a call                       |
| <code>c</code>                             | Continue to the next breakpoint        |
| <code>l</code>                             | List source around the current line    |
| <code>q</code>                             | Quit                                   |

A useful pattern: after a crash, `p` the inputs of the failing frame. If they look wrong, `u` to the caller and `p` what it passed. Repeat until you find the first frame where a value went bad. That is the frame to fix.

**In Databricks notebooks**, the built-in interactive debugger gives breakpoints, stepping and a variable view for Python cells. It runs on serverless and on recent runtimes, after you turn it on in your editor settings.

**In Spark code** a debugger stops only the driver. Code inside a UDF runs on executors, where your breakpoint never fires. Test UDF logic as a plain Python function first.

## Why it matters

A weighted-reach function returned 0.0 for one subgroup. An analyst added eleven print statements over two hours and still missed it. A post-mortem `pdb` session on a forced failure showed in one command that `base` was the unweighted count, passed in by a caller that swapped two arguments. `u`, then `p`, found the caller.

## Worked example: weak code

```
1 def weighted_pct(reach: float, base: float) -> float:
2 """Weighted percentage."""
3 print("DEBUG reach", reach)
4 print("DEBUG base", base)
5 breakpoint()
6 return round(reach / base * 100, 1)
```

1. **Line 3.** Print debugging leaves noise in job logs and shows only what you guessed to print.
2. **Line 5.** Committed by mistake, it waits for input in a job that nobody is watching.

## Worked example: strong code

```
1 def weighted_pct(reach: float, base: float) -> float | None:
2 """Weighted percentage, or None when the base is zero."""
3 if base == 0:
4 return None
5 return round(reach / base * 100, 1)
6
7 # To investigate a failure, run: python -m pdb -c continue run_qa.py
8 # then: p reach, base u p rows
```

1. **Line 3.** The fix the debugging session pointed to: guard the division.
2. **Line 7.** Debug from outside the code. Nothing to remove before you commit.

Let the crash open the debugger, print the failing frame's inputs, then walk up until the values go wrong.

## Recap card

- **Open pdb where a script crashed** `python -m pdb -c continue script.py`. You land in the frame that raised.
- **pdb: print, stack, up, next, step, continue** `p`, `w`, `u`, `n`, `s`, `c`.
- **PYTHONBREAKPOINT=0** Makes `breakpoint()` a no-op. A safety net for jobs, not a reason to commit breakpoints.
- **Why a breakpoint in a Spark UDF never fires** UDFs run on executors. The debugger stops only the driver. Test UDF logic as plain Python.

## Practice

### Exercise 1. Multiple choice (level 1)

A scheduled job runs with `PYTHONBREAKPOINT=0`. The code contains `breakpoint()`. What happens at that line?

- a) Nothing: the call does nothing and the job continues
- b) The job waits for input forever
- c) The job fails with a `RuntimeError`
- d) `pdb` opens and quits at once

### Why?

- r1) `PYTHONBREAKPOINT=0` makes `breakpoint()` a no-op.
- r2) Scheduled jobs never run `pdb`.

### Exercise 2. Fill the gap (level 2)

You are in a post-mortem `pdb` session in `weighted_pct`. Complete the commands.

```
1 (Pdb) ____ (1) reach, base # print the inputs
2 (Pdb) ____ (2) # move to the caller's frame
3 (Pdb) p rows
```

### Exercise 3. Put in order (level 2)

Order a debugging session for a function that returns a wrong number.

- Stop in the function with `breakpoint()` or post-mortem `pdb`
- Print the inputs, walk up with `u` until a value is first wrong
- Reproduce with the smallest input that gives the wrong number
- Fix that frame and add a test with the small input

- Remove the breakpoint and run the tests

### Exercise 4. Write the function (level 3)

Post-mortem debugging reads the local variables of the frame that raised. Write `crash_locals(fn: Callable[..., object], *args: object) -> dict[str, object] | None` that calls `fn(*args)`. If it raises, return a copy of the local variables of the **innermost** frame of the traceback (the one pdb lands in). If it does not raise, return `None`. Import `Callable` from `collections.abc`.

Starter code:

```
1 def crash_locals(fn, *args):
2 try:
3 fn(*args)
4 except Exception:
5 return locals()
```

### Exercise 5. Spot the bug (level 3)

This function is ready for a scheduled job, except for one line. Tap it, then remove it and anything else left from debugging.

```
1 def unweighted_base(rows: list[dict[str, object]]) -> int:
2 """Count in-scope respondents."""
3 in_scope = [r for r in rows if r["in_scope"] is True]
4 breakpoint()
5 return len(in_scope)
```

Write down the line number of the bug. Then write the fix.

## Sources

- `python -m pdb -c continue` enters post-mortem at the raising frame and prints 'Uncaught exception. Entering post mortem debugging'; `PYTHONBREAKPOINT=0` disables `breakpoint()`; commands `p`, `w`, `u`, `d`, `n`, `s`, `c`, `l`, `q`. (verified: Local runs on CPython 3.13.12, September 2026.)
- Databricks notebooks have a built-in interactive debugger for Python, available on serverless and recent runtimes, turned on in editor settings. (verified: <https://docs.databricks.com/aws/en/notebooks/debugger> (checked 2026-09-25))
- ruff rule T100 flags debugger calls such as `breakpoint()`. (verified: ruff 0.16.9 run in this repo reports 'T100 Trace found: `breakpoint` used'.)



## D2-L2 Logging levels that help at 3 a.m.

### Concept

A job that failed overnight leaves you its logs and nothing else. Good logs answer three questions fast: what ran, with what inputs, and where it stopped. Python's `logging` module does this better than `print`.

#### One logger per module:

```
1 import logging
2
3 logger = logging.getLogger(__name__)
```

**Levels**, lowest to highest, with their numbers:

| Level    | Number | Use it for                                  |
|----------|--------|---------------------------------------------|
| DEBUG    | 10     | Detail you need only while investigating    |
| INFO     | 20     | Normal progress: step started, rows written |
| WARNING  | 30     | Unexpected but handled: a WARN check        |
| ERROR    | 40     | A step failed                               |
| CRITICAL | 50     | The pipeline cannot continue                |

A logger passes on records at or above its effective level. The root logger's default is **WARNING**, so `logger.info(...)` prints nothing until you configure logging. Checked on Python 3.13:

- `logging.basicConfig(level=logging.INFO)` configures the root logger **once**. A second call does nothing, unless you pass `force=True`.
- `logger.exception("...")` inside an `except` block logs at **ERROR** and appends the traceback.
- Pass values as arguments: `logger.info("wrote %d rows to %s", n, table)`. The string is built only if the record is emitted. ruff rule G004 flags f-strings in logging calls.

**Logging is not control flow.** `logger.critical(...)` records a message. It does not stop anything. Your standard says a **CRITICAL** check **raises**. Log, then raise.

**Logging is not your run log.** Text logs are for people debugging. The Delta run log table, with one structured row per run, is for queries and dashboards. Write both.

Every log line in a pipeline should carry the identifiers that let you find the data: tracker, wave, table, run id.

## Why it matters

A silver job logged "Processing..." and "Done" at INFO for every wave, and nothing else. When wave W14 produced 30% fewer rows, the logs could not say which step dropped them. Adding one INFO line per step, with the row count, tracker and wave, turned the next investigation from a day of reruns into a two-minute log search.

## Worked example: weak code

```
1 import logging
2
3 def run_checks(results: list[dict[str, str]], wave_id: str) -> None:
4 """Report check results."""
5 for r in results:
6 if r["status"] == "FAIL" and r["severity"] == "CRITICAL":
7 logging.critical(f"check failed: {r['check']}")
8 print("checked", r["check"])
```

1. **Line 7.** Logs CRITICAL and carries on. The pipeline continues with bad data. The f-string is built even when filtered (G004).
2. **Line 8.** print has no level, no logger name and no wave id.

## Worked example: strong code

```
1 import logging
2
3 logger = logging.getLogger(__name__)
4
5
6 class CriticalCheckFailed(RuntimeError):
7 """A CRITICAL data check failed and the pipeline must stop."""
8
9
10 def run_checks(results: list[dict[str, str]], wave_id: str) -> None:
11 """Log each check result and stop on a CRITICAL failure."""
12 for r in results:
13 logger.info("wave %s check %s status %s", wave_id, r["check"],
14 r["status"])
15 if r["status"] == "FAIL" and r["severity"] == "CRITICAL":
16 logger.critical("wave %s CRITICAL check %s failed", wave_id,
17 r["check"])
18 raise CriticalCheckFailed(f"{r['check']} failed for wave
19 {wave_id}")
```

1. **Line 3.** A named logger per module. The name shows up in every record.
2. **Line 13.** Identifiers in every line, formatted only when emitted.
3. **Line 16.** Log, then raise. The level describes. The exception stops.

One named logger per module, levels that match severity, identifiers in every line, and a raise after every CRITICAL.

## Recap card

- **Logging levels and numbers** DEBUG 10, INFO 20, WARNING 30, ERROR 40, CRITICAL 50. Default effective level: WARNING.
- **basicConfig called twice** The second call does nothing, unless `force=True`.
- **logger.exception** Logs at ERROR with the traceback. Call it inside an except block.
- **Does logger.critical stop the pipeline?** No. Log, then raise.

## Practice

### Exercise 1. Predict the output (level 1)

Nothing has configured logging. What does this print?

```
1 import logging
2
3 logger = logging.getLogger("qa")
4 logger.info("rows written")
5 logger.warning("weight sum low")
```

### Exercise 2. Multiple choice (level 2)

A notebook calls `logging.basicConfig(level=logging.WARNING)` in its first cell, then `logging.basicConfig(level=logging.DEBUG)` later. What is the root level after the second call?

- a) WARNING: the second call does nothing
- b) DEBUG
- c) An error is raised
- d) Both handlers are active

### Why?

- r1) `basicConfig` does nothing if the root logger already has handlers, unless `force=True`.
- r2) Logging levels can never be lowered.

### Exercise 3. Fill the gap (level 2)

Log the failure with its traceback, then stop the pipeline.

```
1 try:
2 validate(df)
3 except ContractError:
4 logger.__(1)("contract failed for wave %s", wave_id)
5 __(2)
```

## Exercise 4. Write the function (level 3)

Write `log_check(logger: logging.Logger, check: str, status: str, severity: str, wave_id: str) -> None`. Log at INFO for "PASS", WARNING for "WARN", ERROR for "FAIL" with non-CRITICAL severity, and CRITICAL for "FAIL" with "CRITICAL" severity, then raise `RuntimeError`. Every message contains the wave id and check name. Use %-style arguments, not f-strings, in the logging calls. Raise `ValueError` for an unknown status.

Starter code:

```
1 import logging
2
3
4 def log_check(logger, check, status, severity, wave_id):
5 logger.info(f"{check} {status}")
```

## Exercise 5. Architecture design (level 4)

Set the logging for a silver job that runs nightly on Databricks.

Options: INFO, with DEBUG switched on by a job parameter when investigating, DEBUG always.

1. Level in production
2. Each step's log line
3. Run history for dashboards

## Sources

- Default effective level WARNING; unconfigured output is the bare message; basicConfig second call no-op unless force=True; logger.exception logs at ERROR with traceback; level numbers 10 to 50. (verified: Local run on CPython 3.13.12, September 2026.)
- ruff rule G004 flags f-strings in logging calls. (verified: ruff 0.16.9 run in this repo reports 'G004 Logging statement uses f-string'.)

## D2-L3 explain(), the Spark UI and the query profile

---

### Concept

A slow Spark job often shows why in its **plan**. Read it before you tune anything.

`df.explain()` prints the physical plan without running the query.

`df.explain("formatted")` numbers each operator and lists its details below.

The other modes are "simple" (the default), "extended", "codegen" and

"cost". Read the tree from the **bottom up**: the scan is at the bottom, the result at the top.

What to look for (examples from PySpark 4.0.4):

| In the plan                                 | Means                                                                     |
|---------------------------------------------|---------------------------------------------------------------------------|
| PartitionFilters: [(wave_id = 3)]           | Only matching partitions are read. Good.                                  |
| PushedFilters: [GreaterThan(w, 1.0)]        | The filter runs inside the file reader. Good.                             |
| PushedFilters: [] on a filtered query       | Nothing was pushed down. Check the filter's column and type.              |
| Exchange hashpartitioning(...)              | A shuffle: data moves across the cluster. Each one costs time.            |
| BroadcastHashJoin                           | The small side is copied to every executor. No shuffle of the large side. |
| SortMergeJoin                               | Both sides are shuffled and sorted. Normal for two large tables.          |
| CartesianProduct or BroadcastNestedLoopJoin | Every row meets every row. Usually a missing join condition.              |
| BatchEvalPython                             | A Python UDF. Rows leave the JVM for a Python worker.                     |

`AdaptiveSparkPlan isFinalPlan=false` means adaptive query execution will adjust the plan at run time, for example switching a join to broadcast.

`explain()` shows the starting plan.

**After the run**, look at what happened. On classic compute, the **Spark UI** shows jobs, stages, tasks and the SQL tab with row counts per operator. On **serverless**, the Spark UI is not available. Use the **query profile**, which shows each operator with its time spent, rows and memory peak, and a Top operators panel.

In the UI or profile, compare **rows out** of each join with rows in. A join whose output exceeds its larger input is duplicating rows (lesson D3-L1). Look for one task far slower than the rest of its stage: that is skew.

### Why it matters

A gold job's runtime grew from 4 minutes to 70 over a month. Nobody had changed the code. `explain()` showed `BroadcastNestedLoopJoin`: a join on a range condition, fine while the dimension held 50 rows, painful at 4,000. The

plan named the problem in one line. Rewriting the range join as an equi-join on a precomputed band key brought it back to 5 minutes.

### Worked example: weak code

```
1 from pyspark.sql import functions as F
2
3 joined = respondents.crossJoin(regions).filter(
4 F.col("postcode_area") == F.col("area_code")
5)
6 joined.write.mode("overwrite").saveAsTable("research_gold.nat_regions")
```

1. **Line 3.** A cross join first, then a filter. Check the plan before you trust the optimiser to rewrite it.
2. **Line 6.** Runs without ever reading the plan.

### Worked example: strong code

```
1 from pyspark.sql import functions as F
2
3 joined = respondents.join(
4 F.broadcast(regions), respondents["postcode_area"] ==
5 regions["area_code"], "left"
6)
7 joined.explain("formatted")
8 # Expect BroadcastHashJoin, no CartesianProduct, one scan of respondents.
9 joined.write.mode("overwrite").saveAsTable("research_gold.nat_regions")
```

1. **Line 4.** An explicit equi-join with the small side broadcast.
2. **Line 6.** Read the plan once, when you write the code, and again when data volumes change.

Read the plan bottom up, look for pushed filters, shuffles, join types and Python UDFs, then confirm rows per operator in the Spark UI or query profile.

### Recap card

- **Read a physical plan in which direction?** Bottom up: scan at the bottom, result at the top.
- **PartitionFilters and PushedFilters** Filters applied while reading. Empty PushedFilters on a filtered query means no pushdown.
- **Plan operators that signal trouble** CartesianProduct or BroadcastNestedLoopJoin (cross join), BatchEvalPython (Python UDF), many Exchange (shuffles).
- **Spark UI on serverless** Not available. Use the query profile: time spent, rows and memory per operator.

## Practice

### Exercise 1. Multiple choice (level 1)

A job's plan contains `CartesianProduct`. What is the most likely cause?

- a) A join without an equality condition, so every row meets every row
- b) A broadcast join
- c) A Python UDF
- d) Partition pruning

#### Why?

- r1) A Cartesian product pairs every row with every row. It comes from a cross join or a join condition Spark cannot use as an equi-join.
- r2) Spark uses Cartesian products for all small tables.

### Exercise 2. Spot the bug (level 2)

This is the formatted plan for `spark.read.parquet(path).filter(F.col("w") > 1.0).groupBy("wave_id").agg(F.sum("w"))`. Tap the line that shows a shuffle.

```
1 AdaptiveSparkPlan (6)
2 +- HashAggregate (5)
3 +- Exchange (4)
4 +- HashAggregate (3)
5 +- Filter (2)
6 +- Scan parquet (1)
```

Write down the line number of the bug.

### Exercise 3. Multiple choice (level 2)

Your job runs on **serverless** compute. Where do you see time spent and rows per operator after it runs?

- a) The query profile
- b) The Spark UI's SQL tab
- c) `df.explain('cost')`
- d) The driver logs

#### Why?

- r1) The Spark UI is not available on serverless. Databricks points you to the query profile.
- r2) The query profile is the only tool that works in Databricks.

### Exercise 4. Write the function (level 3)

Write `plan_warnings(plan: str, max_exchanges: int = 2) -> list[str]` that reads the text of `df.explain()` and returns warnings, in this order, each at most once: "cross join" if the plan contains `CartesianProduct` or `BroadcastNestedLoopJoin`, "python udf" if it contains `BatchEvalPython` or `ArrowEvalPython`, and "many shuffles" if the word `Exchange` appears more than `max_exchanges` times, not counting `BroadcastExchange`.

Starter code:

```
1 def plan_warnings(plan, max_exchanges=2):
2 warnings = []
3 if "Cartesian" in plan:
4 warnings.append("cross join")
5 return warnings
```

### Exercise 5. Architecture design (level 4)

A gold job on classic compute slowed from 5 to 40 minutes. Plan the investigation.

Options: Run `explain('formatted')` and read the joins and exchanges, Double the cluster size.

1. First step
2. Checking what happened at run time
3. One task in a stage takes 30 minutes, the rest 20 seconds

### Sources

- `explain` modes `simple`, `extended`, `codegen`, `cost`, `formatted`; plans show `PartitionFilters`, `PushedFilters`, `Exchange hashpartitioning`, `BroadcastHashJoin`, `SortMergeJoin`, `CartesianProduct`, `BatchEvalPython`, `AdaptiveSparkPlan isFinalPlan=false`. (verified: Local runs on PySpark 4.0.4, September 2026.)
- The Spark UI is not available on serverless; use the query profile, which shows time spent, rows and memory per operator and a Top operators panel. (verified: <https://docs.databricks.com/aws/en/compute/serverless/limitations> and <https://docs.databricks.com/aws/en/sql/user/queries/query-profile> (checked 2026-09-25))



# D3 Silent errors

Roadmap: Every month. Join explosion, nulls in filters, type coercion, timezone drift, mutable defaults.

## D3-L1 Join row explosion and nulls in filters

### Concept

Some bugs raise no error. The job succeeds and the numbers are wrong. Two of the most common come from joins and from nulls.

**Join row explosion.** A left join of respondents to a region lookup should keep one row per respondent. If the lookup holds a key twice, say `N` for "North" and "North (old)", every respondent in `N` appears twice. Weighted totals double for that region. Nothing fails.

Check the grain around every join:

- Count rows before and after. A many-to-one join must not change the count.
- In pandas, pass `validate="many_to_one"`. A duplicate key then raises `MergeError: Merge keys are not unique in right dataset`. Checked on pandas 3.0.2.
- In Spark, count duplicate keys in the lookup before joining:  
`dim.groupBy(key).count().filter("count > 1").`

**Nulls behave differently in pandas and Spark.** Checked on pandas 3.0.2 and PySpark 4.0.4:

| Operation                                                        | pandas                             | Spark and SQL                          |
|------------------------------------------------------------------|------------------------------------|----------------------------------------|
| Filter <code>region != "N"</code>                                | Keeps rows where region is missing | <b>Drops</b> rows where region is null |
| Join on a null key                                               | <b>Matches</b> null to null        | Never matches a null key               |
| <code>x NOT IN (subquery)</code> where the subquery holds a null | Not applicable                     | Returns <b>no rows</b>                 |

In Spark and SQL, a comparison with null gives null, not True or False, and a filter keeps only True. So `F.col("region_code") != "N"` silently drops every respondent with no region. Say what you mean:

- Keep nulls: `(F.col("region_code") != "N") | F.col("region_code").isNull()`, or `~F.col("region_code").eqNullSafe("N")`.
- Exclude a list that might contain nulls: use `NOT EXISTS` or a `left_anti` join, not `NOT IN`.

In pandas, the surprise runs the other way: null keys match each other. Drop or fill null keys before a merge unless you want that.

## Why it matters

A regional reach table showed the North at 118% weighted reach. The region lookup had gained a second row for code `N` when a boundary changed. The left join doubled every northern respondent. A row count check before and after the join, or `validate="many_to_one"`, would have stopped the job at the join.

## Worked example: weak code

```
1 def add_region(respondents: pd.DataFrame, regions: pd.DataFrame) ->
 pd.DataFrame:
2 """Attach region names."""
3 merged = respondents.merge(regions, on="region_code", how="left")
4 return merged[merged["region"] != "Unknown"]
```

1. **Line 3.** No grain check. A duplicate code doubles rows, and null codes match a null row in regions.
2. **Line 4.** A filter that also changes the row count, mixed into the join step.

## Worked example: strong code

```
1 def add_region(respondents: pd.DataFrame, regions: pd.DataFrame) ->
 pd.DataFrame:
2 """Attach region names. One output row per respondent.
3
4 Raises:
5 pandas.errors.MergeError: If a region code appears twice in
 regions.
6 """
7 lookup = regions.dropna(subset=["region_code"])
8 merged = respondents.merge(lookup, on="region_code", how="left",
9 validate="many_to_one")
10 if len(merged) != len(respondents):
11 raise ValueError(f"join changed rows: {len(respondents)} ->
 {len(merged)}")
```

```
11 return merged
```

1. **Line 7.** pandas matches null keys to each other. Drop them from the lookup.
2. **Line 8.** pandas raises MergeError if the lookup has a duplicate key.
3. **Line 9.** An explicit grain check. It also works where validate is not available.

Check row counts across every join. Know how your engine treats nulls in filters and keys, and write the null case explicitly.

## Recap card

- **Check for join row explosion** Row counts before and after. pandas: `validate='many_to_one'`. Spark: count duplicate keys in the lookup first.
- **col != 'x' with nulls: pandas vs Spark** pandas keeps missing values. Spark and SQL drop nulls, because the comparison is null.
- **Null join keys: pandas vs Spark** pandas merge matches null to null. Spark and SQL never match null keys.
- **NOT IN with a null in the list** Returns no rows. Use NOT EXISTS or a `left_anti` join.

## Practice

### Exercise 1. Predict the output (level 1)

What does this print? (pandas 3)

```
1 import pandas as pd
2
3 resp = pd.DataFrame({"id": [1, 2, 3], "code": ["N", "S", "N"]})
4 regions = pd.DataFrame({"code": ["N", "N", "S"], "name": ["North", "North
 (old)", "South"]})
5 print(len(resp.merge(regions, on="code", how="left")))
```

### Exercise 2. Multiple choice (level 2)

In PySpark, `respondents.filter(F.col("region_code") != "N")` runs on 1,000 rows: 300 are N, 50 are null. How many rows remain?

- a) 650
- b) 700
- c) 350
- d) 1,000

### Why?

- r1) `null != 'N'` evaluates to null, and a filter keeps only rows where the condition is true.
- r2) Spark drops nulls from every DataFrame before filtering.

### Exercise 3. Multiple choice (level 2)

`SELECT * FROM respondents WHERE region_code NOT IN (SELECT code FROM excluded_regions)` returns no rows. `excluded_regions` holds 'N' and one null. Why?

- a) `x NOT IN (... , NULL)` is never true, because `x <> NULL` is null
- b) `NOT IN` is not supported in Spark SQL
- c) The subquery returns too many rows
- d) Every respondent is in region N

#### Why?

- r1) `NOT IN` expands to `x <> 'N' AND x <> NULL`. The second term is null, so the whole condition is never true.
- r2) Subqueries ignore nulls.

### Exercise 4. Write the function (level 3)

Write

```
safe_left_join(facts: pd.DataFrame, lookup: pd.DataFrame, key: str) ->
pd.DataFrame. Drop rows with a null key from lookup first. If a key appears
more than once in the remaining lookup, raise ValueError naming the
duplicated keys in sorted order. Return the left join, which must have exactly
len(facts) rows. Null keys in facts must get missing lookup values, not
match anything.
```

Starter code:

```
1 import pandas as pd
2
3
4 def safe_left_join(facts, lookup, key):
5 return facts.merge(lookup, on=key, how="left")
```

### Exercise 5. Spot the bug (level 3)

This step should remove test respondents (`panel == "test"`) and keep everyone else, including respondents with no panel recorded. 4% of real respondents vanish. Tap the line, then fix it.

```
1 from pyspark.sql import functions as F
2
3 live = respondents.filter(F.col("panel") != "test")
4 live.write.mode("overwrite").saveAsTable("research_silver.nat_live")
```

Write down the line number of the bug. Then write the fix.

## Sources

- pandas 3.0.2: `validate='many_to_one'` raises `MergeError` 'Merge keys are not unique in right dataset; not a many-to-one merge'; `!=` keeps missing values; merge matches null keys. (verified: Local run on pandas 3.0.2 (Pyodide pin), September 2026.)
- PySpark 4.0.4: `!=` filter drops nulls; `eqNullSafe` keeps them when negated; null join keys do not match; `NOT IN` with a null in the subquery returns no rows; `NOT EXISTS` and `left_anti` do not. (verified: Local run on PySpark 4.0.4, September 2026.)

## D3-L2 Type coercion and timezone drift

---

### Concept

Two more silent errors: values that change type, and times that change meaning with the zone.

**Type coercion.** Checked on pandas 3.0.2, PySpark 4.0.4 and Python 3.13:

- `pd.read_csv` turns `respondent_id 0123` into the integer `123`. The leading zero is gone, and joins to the panel file fail. Pass `dtype={"respondent_id": "string"}` for identifiers.
- An integer column with one missing value becomes `float64` in pandas. `1` becomes `1.0`, and string keys built from it read `"1.0"`.
- Strings sort as text: `sorted(["10", "9", "2"])` gives `['10', '2', '9']`. Wave "W10" sorts before "W2". Pad your ids, or sort on a number.
- Spark's `df.union(other)` matches columns **by position**, not by name. A delivery with `wave_id` and `respondent_id` swapped unions without error, with the values in the wrong columns. Use `unionByName`.
- Under ANSI mode, Spark casts a string column compared with a number: `F.col("age") > 9` raises `CAST_INVALID_INPUT` on the value `'x'`. Cast explicitly with `try_cast` first.

**Timezone drift.** A timestamp without a zone is a guess.

- `datetime.utcnow()` returns a **naive** datetime and is deprecated from Python 3.12. Comparing naive with aware raises `TypeError: can't compare offset-naive and offset-aware datetimes`. Use `datetime.now(timezone.utc)`.
- Spark's `TIMESTAMP` type is an instant, displayed in the **session time zone** (`spark.sql.session.timeZone`). `F.to_date` on `2026-06-30T23:30:00Z` gives

30 June in UTC and **1 July** in Europe/London. A wave cut-off at midnight moves respondents between waves.

- The same text without an offset, `'2026-06-30 23:30:00'`, parses to instants 4 hours apart in UTC and America/New\_York sessions.
- Some local times do not exist. pandas refuses to localise `2026-03-29 01:30` to Europe/London: the clocks skipped that hour.

Store instants in UTC. Convert to local time only for display, or for a business rule that names its zone, such as "fieldwork closes at midnight UK time".

## Why it matters

A wave closed at midnight UK time on 30 June. The pipeline ran with the session time zone set to UTC and assigned every interview completed between 23:00 and 00:00 UTC on 30 June to wave W14. In British Summer Time those interviews happened after midnight, in W15. 340 respondents moved wave. Both waves' reach figures shifted by up to 1.5 points.

## Worked example: weak code

```
1 from datetime import datetime
2
3 def in_wave(completed_at: datetime, close_at: datetime) -> bool:
4 """Was the interview completed before the wave closed?"""
5 return completed_at < close_at
6
7 ① close = datetime(2026, 6, 30, 23, 59, 59)
8 ② in_wave(datetime.utcnow(), close)
```

1. **Line 7.** Naive: is this UK time or UTC? The code cannot tell, and nor can a reader.
2. **Line 8.** Deprecated since Python 3.12, and naive too.

## Worked example: strong code

```
1 from datetime import datetime, timedelta, timezone
2
3 ① BST = timezone(timedelta(hours=1), "BST")
4
5 def in_wave(completed_at: datetime, close_at: datetime) -> bool:
6 """Was the interview completed before the wave closed?
7
8 Raises:
9 ValueError: If either datetime is naive.
10 """
11
12 ② if completed_at.tzinfo is None or close_at.tzinfo is None:
13 raise ValueError("datetimes must carry a time zone")
14 return completed_at < close_at
```

```
15 close = datetime(2026, 7, 1, 0, 0, tzinfo=BST)
3 in_wave(datetime.now(timezone.utc), close)
```

1. **Line 3.** The business rule names its zone. In production, use `zoneinfo.ZoneInfo("Europe/London")` to handle the clock change.
2. **Line 11.** Refuse naive values instead of guessing.
3. **Line 16.** Aware UTC now. Aware datetimes compare correctly across zones.

Read identifiers as strings, union by name, and keep every timestamp aware and stored in UTC.

## Recap card

- **pd.read\_csv and leading zeros** Default parsing makes 0123 into 123. Pass `dtype='string'` for identifiers.
- **Spark union vs unionByName** `union` matches columns by position, silently. `unionByName` matches by name.
- **datetime.utcnow()** Deprecated from Python 3.12 and naive. Use `datetime.now(timezone.utc)`.
- **Spark TIMESTAMP and the session time zone** An instant, shown and converted in `spark.sql.session.timeZone`. `to_date` can move a row to another day.

## Practice

### Exercise 1. Predict the output (level 1)

What does this print?

```
1 waves = ["W10", "W9", "W2"]
2 print(sorted(waves))
```

### Exercise 2. Multiple choice (level 2)

A Spark session has `spark.sql.session.timeZone` set to `Europe/London`. An interview's `completed_at` is the instant `2026-06-30T23:30:00Z`. What does `F.to_date("completed_at")` return?

- a) 2026-07-01
- b) 2026-06-30
- c) It depends on the cluster's operating system clock
- d) NULL

### Why?

- r1) **TIMESTAMP** is an instant. `to_date` uses the session time zone, where 23:30 UTC is 00:30 on 1 July in BST.

- r2) Spark always converts to local time before storing.

### Exercise 3. Fill the gap (level 2)

Replace the deprecated call with an aware UTC timestamp.

```
1 from datetime import datetime, timezone
2
3 run_timestamp = datetime.__(1)__(__(2))
```

### Exercise 4. Write the function (level 3)

Write `load_delivery(csv_text: str) -> pd.DataFrame`. Read the CSV with `respondent_id` and `postcode_area` as strings, keeping leading zeros. Convert `final_weight` to float. If any non-empty weight fails to parse, raise `ValueError` naming how many. Empty weights stay missing.

Starter code:

```
1 import io
2
3 import pandas as pd
4
5
6 def load_delivery(csv_text):
7 return pd.read_csv(io.StringIO(csv_text))
```

### Exercise 5. Spot the bug (level 3)

A redelivery arrives with its columns in a different order. After this step, some respondent ids read W14. Tap the line, then fix it.

```
1 from pyspark.sql import functions as F
2
3 combined = original.union(redelivery)
4 combined = combined.withColumn("loaded_at", F.current_timestamp())
```

Write down the line number of the bug. Then write the fix.

## Sources

- pandas 3.0.2: `read_csv` drops leading zeros by default; int with None becomes float64; `tz_localize` rejects a nonexistent DST time. (verified: Local run on pandas 3.0.2, September 2026.)
- Python 3.13: `utcnow()` raises a `DeprecationWarning` and is naive; comparing naive and aware datetimes raises `TypeError`. (verified: Local run on CPython 3.13.12, September 2026.)
- PySpark 4.0.4: `union` matches by position; `to_date` depends on session time zone; naive strings parse relative to the session zone; ANSI casts a



string compared with an int and raises `CAST_INVALID_INPUT` on 'x'.  
(verified: Local run on PySpark 4.0.4, September 2026.)

## D3-L3 Mutable defaults and the truthiness of zero

---

### Concept

Two Python habits produce wrong results with no error message. Both appear in code written by people and by AI assistants, including your reference pipeline.

**Mutable default arguments.** Python evaluates a default value once, when it defines the function. Every call that relies on the default shares that one object. Write `def add_check(result, results=[])`, and each call appends to the same list. The second validation run returns the first run's checks as well.

The fix is a `None` sentinel:

```
1 def add_check(result: str, results: list[str] | None = None) -> list[str]:
2 results = [] if results is None else results
3 results.append(result)
4 return results
```

ruff flags the pattern as rule B006. Treat that warning as a bug.

**Truthiness of numbers.** `if value:` is `False` for `None`, and also for `0`, `0.0`, empty strings and empty containers. When zero is a valid measurement, such as a 0% reach, a zero weight or zero rows, a truthiness test treats it as missing.

Your reference `check_wave_stability` does this:

```
1 if current_pct and prior_pct:
```

A metric that falls from 12% to 0% skips the comparison. The largest possible drop produces no warning.

Test for the case you mean. For missing, write `is None`. For zero, write `== 0`.

```
1 if current_pct is not None and prior_pct is not None:
```

The same rule covers type hints. `error_message: str = None` claims the value is always a string. Write `str | None = None`, and mypy then makes you handle the `None` case.

### Why it matters

A questionnaire routing error removes one question for a whole wave. Its weighted percentage drops from 12% to 0.0%. The wave-over-wave stability check runs `if current_pct and prior_pct:`, treats 0.0 as missing and skips

the metric. QA reports PASS. The dashboard publishes 0% reach for that question, and a client asks why the channel "lost its whole audience" in a week.

The check existed to catch exactly this change. One truthiness test disabled it for the one value that mattered most.

### Worked example: weak code

```
❷ def compare_waves(current: dict, prior: dict, tolerance: float = 0.10,
 flags=[]):
 2 """Flag metrics whose share moved more than tolerance between
 waves."""
 3 for metric, current_pct in current.items():
 4 prior_pct = prior.get(metric)
 ❸ if current_pct and prior_pct:
 6 if abs(current_pct - prior_pct) > tolerance:
 7 flags.append(metric)
 8 return flags
```

1. **Line 1.** Bare `dict` hides the value type. The values are floats or `None`.
2. **Line 1.** The default list is created once. Flags from one call leak into the next.
3. **Line 5.** A share of 0.0 is falsy, so the biggest possible drop is never compared.

### Worked example: strong code

```
1 def compare_waves(
 ❶ current: dict[str, float | None],
 3 prior: dict[str, float | None],
 4 tolerance: float = 0.10,
 5) -> list[str]:
 6 """Flag metrics whose share moved more than tolerance between waves.
 7
 8 A metric is skipped only when a share is missing. A share of 0.0 is
 compared.
 9 """
 ❷ flags: list[str] = []
 11 for metric, current_pct in current.items():
 12 prior_pct = prior.get(metric)
 ❸ if current_pct is None or prior_pct is None:
 14 continue
 15 if abs(current_pct - prior_pct) > tolerance:
 16 flags.append(metric)
 17 return flags
```

1. **Line 2.** The hint says a share can be missing, so mypy checks every use.
2. **Line 10.** A new list on every call. No state leaks between runs.
3. **Line 13.** Skips only missing values. 0.0 goes on to the comparison.

Both fixes are one line. Both bugs pass every test that uses only non-zero values and a single call.

## Recap card

- **When does Python evaluate a default argument?** Once, when it defines the function. A list or dict default is shared by every call.
- **Safe pattern for an optional list argument** Default to None. Inside: `results = [] if results is None else results.`
- **What is wrong with** `if current_pct and prior_pct: ?` 0.0 is falsy, so a real share of zero is treated as missing. Use `is not None`.
- **Type hint for an optional string that defaults to None** `str | None = None`. Never `str = None`.

## Practice

### Exercise 1. Predict the output (level 1)

What does this print?

```
1 def add_flag(flag, flags=[]):
2 flags.append(flag)
3 return flags
4
5 add_flag("null_weights")
6 print(add_flag("row_count"))
```

### Exercise 2. Multiple choice (level 2)

`current_pct = 0.0` and `prior_pct = 0.12`. The tolerance is 0.10. What does `check_wave_stability` from the reference code do with this metric?

- a) Flags it as WARN, because the change is 0.12.
- b) Marks it PASS.
- c) Leaves it out of the results.
- d) Raises a ZeroDivisionError.

### Why?

- r1) `if current_pct and prior_pct:` is False when `current_pct` is 0.0, so the block never runs.
- r2) The function filters out metrics below 1%.

### Exercise 3. Fill the gap (level 3)

Fix both signatures. The error message is optional, and the flags list must not be shared between calls.

```
1 def log_run(status: str, error_message: ____ (1) = None) -> None: ...
2
```

```
3 def collect(flag: str, flags: list[str] | None = ____ (2)) ->
 list[str]: ...
```

### Exercise 4. Spot the bug (level 3)

This check should warn when a share moves by more than the tolerance. A drop from 12% to 0% produces no warning. Tap the line with the bug, then fix the function.

```
1 def stability_status(current_pct: float | None, prior_pct: float | None,
 tolerance: float = 0.10) -> str:
2 """Return WARN, PASS, or SKIP when a share is missing."""
3 if current_pct and prior_pct:
4 change = abs(current_pct - prior_pct)
5 return "WARN" if change > tolerance else "PASS"
6 return "SKIP"
```

Write down the line number of the bug. Then write the fix.

### Exercise 5. Write the function (level 4)

Write `compare_waves(current, prior, tolerance=0.10)`. Both arguments map metric names to a share or `None`. Return the metrics whose share changed by more than `tolerance`, in the order of `current`. Skip a metric only when either share is `None` or the metric is missing from `prior`. Calling the function twice must not carry results over.

Starter code:

```
1 def compare_waves(
2 current: dict[str, float | None],
3 prior: dict[str, float | None],
4 tolerance: float = 0.10,
5) -> list[str]:
6 ...
```

### Exercise 6. Refactor (level 5)

A validator collects check results with the helper below. The second contract validation in a job reports the first contract's failures too. Refactor `add_result` so each validation starts empty, keep the signature compatible for callers that pass a list, and add type hints and a docstring.

Starter code:

```
1 def add_result(name, passed, results=[]):
2 results.append({"check": name, "passed": passed})
3 return results
```

## Sources

- ruff rule B006 flags mutable default arguments. (verified: ruff 0.16.9 run in this repo reports B006 on `def f(x=[]):`.)

# D4 Data errors

Roadmap: Every month. Row counts across layers, distribution drift and weight totals.

## D4-L1 Row counts across layers

### Concept

Every layer of a medallion pipeline changes the row count for a reason. Write the reasons down as rules, then check them on every run. A count that breaks a rule is a bug, even when no step failed.

For one tracker and one wave:

| Check              | Rule                                                                          |
|--------------------|-------------------------------------------------------------------------------|
| Bronze vs delivery | Bronze rows for the delivery = rows in the delivered file or manifest         |
| Silver vs bronze   | Silver rows = distinct respondent keys in bronze. Never more than bronze rows |
| Gold vs silver     | Sum of gold <code>unweighted_n</code> = silver rows in scope                  |
| Redelivery         | After MERGE, silver rows = rows before + rows inserted                        |

**Where the numbers come from.** Count each layer with one `agg` pass, not repeated `count()` calls on an uncached `DataFrame`: `df.agg(F.count("*"), F.count_distinct("respondent_id"))`. For writes, Delta records metrics in the table history. `DESCRIBE HISTORY` shows `numOutputRows` for a write, and for a MERGE `numSourceRows`, `numTargetRowsInserted` and `numTargetRowsUpdated`. Checked on delta-spark 4.0.1.

**Filter both sides the same way.** Your reference `get_unprocessed_waves` compares bronze waves **for one tracker** with silver waves **for all trackers**. When another tracker already has wave W14 in silver, this tracker's W14 is never loaded. Every reconciliation query filters bronze, silver and gold on the same tracker and wave.

**Store the result.** Write each reconciliation as a structured row in the run log: layer, expected, actual, status. A failed CRITICAL rule halts the pipeline before gold is published.

Counts catch lost and duplicated rows. They do not catch wrong values in the right number of rows. Lessons D4-L2 and D4-L3 cover those.

## Why it matters

Wave W14 for one tracker never reached silver. The job logged success every night for three weeks, because `get_unprocessed_waves` found W14 in silver, loaded by a different tracker. A rule "silver rows for tracker and wave > 0 when bronze rows > 0" would have failed on the first night.

## Worked example: weak code

```
1 bronze_n = bronze.count()
2 silver_n = silver.count()
3 if silver_n > 0:
4 print("silver loaded")
```

1. **Line 1.** All trackers and all waves. The number means nothing for this load.
2. **Line 3.** Any row from any tracker passes this check.

## Worked example: strong code

```
1 from pyspark.sql import functions as F
2
3 def layer_counts(df, tracker_id: str, wave_id: str) -> dict[str, int]:
4 """Rows and distinct respondents for one tracker and wave, in one
5 pass."""
6 row = (
7 df.filter((F.col("tracker_id") == tracker_id) & (F.col("wave_id")
8 == wave_id))
9 .agg(F.count("*").alias("rows"),
10 F.count_distinct("respondent_id").alias("keys"))
11 .first()
12)
13 return {"rows": row["rows"], "keys": row["keys"]}
14
15 bronze_c = layer_counts(bronze, "NAT", "W14")
16 silver_c = layer_counts(silver, "NAT", "W14")
17 if silver_c["rows"] != bronze_c["keys"]:
18 raise RuntimeError(f"silver {silver_c['rows']} != bronze keys
19 {bronze_c['keys']}")
```

1. **Line 6.** The same tracker and wave filter on every layer.
2. **Line 7.** Both counts in one pass over the data.
3. **Line 14.** A written rule, checked every run. It halts before gold.

Write the expected relationship between layers as rules, compute counts in one pass on identical filters, and halt when a CRITICAL rule fails.

## Recap card

- **Row-count rules: silver vs bronze** Silver rows = distinct respondent keys in bronze for the tracker and wave.

- **Rows after a MERGE** Rows before + numTargetRowsInserted (minus any deletes). Updates add nothing.
- **Counting efficiently** One agg pass: F.count('\*') and F.count\_distinct(key). Not repeated count() calls.
- **The get\_unprocessed\_waves bug** Bronze filtered on tracker, silver not. Filter both sides on the same keys.

## Practice

### Exercise 1. Multiple choice (level 1)

Bronze holds 10,400 rows for tracker NAT, wave W14, with 10,000 distinct respondents. Silver deduplicates on respondent and wave. How many silver rows do you expect?

- a) 10,000
- b) 10,400
- c) Any number up to 10,400
- d) It depends on the weights

#### Why?

- r1) Silver's grain is one row per respondent per wave, so it holds one row per distinct key.
- r2) Silver always drops 4% of rows.

### Exercise 2. Multiple choice (level 2)

A redelivery MERGE into silver reports `numSourceRows = 500`, `numTargetRowsInserted = 120`, `numTargetRowsUpdated = 380`. Silver had 10,000 rows before. How many rows does it hold now?

- a) 10,120
- b) 10,500
- c) 10,380
- d) 10,000

#### Why?

- r1) Updates replace existing rows. Only inserts add rows.
- r2) MERGE never removes rows, so every source row is added.

### Exercise 3. Fill the gap (level 2)

Compute rows and distinct respondents in one pass.

```
1 row = df.agg(
```



```

2 F.count("*").alias("rows"),
3 F.____(1)("respondent_id").alias("keys"),
4).____(2)()

```

### Exercise 4. Write the function (level 3)

Write `reconcile(delivered: int, bronze_rows: int, bronze_keys: int, silver_rows: int, gold_n: int, in_scope: int) -> list[dict[str, object]]`. Return one record per rule, in this order, each `{"rule": name, "expected": int, "actual": int, "status": "PASS" or "FAIL"}: "bronze_vs_delivery"` (`bronze_rows` equals `delivered`), `"silver_vs_bronze_keys"` (`silver_rows` equals `bronze_keys`), `"gold_vs_silver_scope"` (`gold_n` equals `in_scope`). Raise `ValueError` if any count is negative, or if `in_scope` exceeds `silver_rows`.

Starter code:

```

1 def reconcile(delivered, bronze_rows, bronze_keys, silver_rows, gold_n,
 in_scope):
2 return [{"rule": "all", "status": "PASS" if silver_rows > 0 else
 "FAIL"}]

```

### Exercise 5. Spot the bug (level 4)

This is your reference `get_unprocessed_waves`, cut down. Tracker NAT's wave W14 is never loaded when another tracker already has W14 in silver. Tap the line with the bug, then fix it.

```

1 bronze_waves = (
2 spark.table(bronze_table).filter(F.col("tracker_id") == tracker_id)
3 .select("wave_id").distinct()
4)
5 silver_waves = spark.table(silver_table).select("wave_id").distinct()
6 new_waves = [r.wave_id for r in
 bronze_waves.subtract(silver_waves).collect()]

```

Write down the line number of the bug. Then write the fix.

## Sources

- Delta MERGE history records `numSourceRows`, `numTargetRowsInserted`, `numTargetRowsUpdated`; writes record `numOutputRows`. (verified: Local run of `DESCRIBE HISTORY` on `delta-spark 4.0.1` with `PySpark 4.0.4`, September 2026.)

## D4-L2 Distribution drift between waves

### Concept

Row counts can be perfect while the data inside has changed. A routing error drops one question. A supplier recodes region. A weighting run uses last year's targets. Each shows up as a **shift in distributions** between waves.

Three checks, from simple to sensitive:

- **Category shares.** For each categorical variable, compare each category's weighted share with the prior wave. Flag a move larger than a tolerance, say 5 percentage points. This is what your `check_wave_stability` does, once its truthiness bug is fixed (lesson D3-L3).
- **New and missing categories.** A category present in one wave and absent in the other is a recode until proven otherwise. `region = "Northern"` appearing beside `"North"` is a supplier change, not a population change.
- **Population stability index (PSI)** summarises a whole distribution in one number:

```
1 PSI = sum over categories of (actual - expected) x ln(actual / expected)
```

where `expected` and `actual` are shares in the prior and current wave. PSI is 0 when nothing moved. A category with a share of zero makes the logarithm undefined, so replace zero shares with a small floor, such as 0.0001, before computing. A common rule of thumb reads PSI below 0.1 as stable and above 0.25 as a large shift. Those cut-offs are convention, not a standard: calibrate them on your own history of waves.

**Guard the divisions.** Every share divides by a total. An empty wave, or a variable with no valid answers, gives a total of zero. Your reference `_check_distributions` divides by `total` with no guard and fails on an empty delivery. Check the total first and report "no data" as its own status.

Drift is a question, not a verdict. Real change happens. The check makes a person look before the numbers are published.

### Why it matters

A supplier renamed region code `"North"` to `"Northern"` in wave W21. Silver loaded both values, the region lookup matched only `"North"`, and the northern subgroup fell from 1,900 respondents to 12. Every row count passed. A new-and-missing category check on `region` would have flagged `"Northern"` on the first run.

## Worked example: weak code

```
1 def share_changes(prior: dict[str, int], current: dict[str, int]) -> dict[str, float]:
2 """Change in share per category."""
3 p_total = sum(prior.values())
4 c_total = sum(current.values())
5 return {k: current[k] / c_total - prior[k] / p_total for k in prior}
```

1. **Line 4.** An empty wave makes this 0, and the division fails.
2. **Line 5.** Loops over prior only. A new category in current is ignored, and a missing one raises `KeyError`.

## Worked example: strong code

```
1 def share_changes(prior: dict[str, int], current: dict[str, int]) -> dict[str, float]:
2 """Change in share per category, over the union of categories.
3
4 Raises:
5 ValueError: If either wave has no observations.
6 """
7 p_total = sum(prior.values())
8 c_total = sum(current.values())
9 if p_total == 0 or c_total == 0:
10 raise ValueError("a wave has no observations")
11 categories = sorted(prior.keys() | current.keys())
12 return {
13 k: current.get(k, 0) / c_total - prior.get(k, 0) / p_total for k
14 in categories
15 }
```

1. **Line 9.** An empty wave is reported, not divided by.
2. **Line 11.** New and missing categories both appear, with a share of 0 on the side where they are absent.

Compare shares over the union of categories, flag new and missing values, summarise with PSI, and guard every total.

## Recap card

- **Three drift checks** Category shares vs prior wave, new and missing categories, and a PSI summary.
- **PSI formula** Sum of (actual - expected) x ln(actual / expected) over shares, with zero shares floored.
- **PSI thresholds 0.1 and 0.25** A rule of thumb, not a standard. Calibrate on your own wave history.
- **Percentage points vs relative change** 30% to 45% is +15 points and +50% relative. Say which a tolerance uses.

## Practice

### Exercise 1. Multiple choice (level 1)

Wave W20 has region values North, South, East. Wave W21 has Northern, South, East. Row counts match. What is the most likely cause?

- a) A recode by the supplier
- b) The population of the North moved
- c) A weighting error
- d) Normal sampling variation

### Why?

- r1) A category that vanishes while a similar new one appears is a label change until proven otherwise.
- r2) Suppliers always rename categories between waves.

### Exercise 2. Predict the output (level 2)

What does this print?

```
1 prior = {"yes": 30, "no": 70}
2 current = {"yes": 45, "no": 55}
3 change = current["yes"] / 100 - prior["yes"] / 100
4 print(round(change * 100))
```

### Exercise 3. Fill the gap (level 2)

Complete the PSI term for one category, with zero shares floored.

```
1 import math
2
3 e = max(expected_share, FLOOR)
4 a = max(actual_share, FLOOR)
5 term = (a - e) * math.__(1) (a / ____(2))
```

### Exercise 4. Write the function (level 3)

Write `psi(expected: dict[str, int], actual: dict[str, int], floor: float = 0.0001) -> float`. Both map categories to counts. Use the union of categories, a missing category counting 0. Convert counts to shares, replace shares below `floor` with `floor`, and return the sum of  $(a - e) * \ln(a / e)$ . Raise `ValueError` if either total is 0.

Starter code:

```
1 import math
2
3
4 def psi(expected, actual, floor=0.0001):
```

```
5 return sum((actual[k] - expected[k]) * math.log(actual[k] /
 expected[k]) for k in expected)
```

### Exercise 5. Spot the bug (level 3)

This is your reference `_check_distributions` logic, as a plain function. It fails on an empty delivery. Tap the line, then fix it: return `"NO_DATA"` when the total is 0.

```
1 def distribution_status(value_count: int, total: int, min_pct: float,
 max_pct: float) -> str:
2 """PASS if the value's share is within bounds, else WARN."""
3 pct = value_count / total
4 return "PASS" if min_pct <= pct <= max_pct else "WARN"
```

Write down the line number of the bug. Then write the fix.

### Sources

- $PSI = \sum (a - e) \ln(a / e)$  over category shares; zero shares need a floor. (verified: Definition checked by the exercise tests in this repo; the formula follows from its terms.)
- PSI thresholds of 0.1 and 0.25 are an informal industry convention. (unverified: Widely quoted in credit-scoring practice; no official standard located. Taught as a rule of thumb only.)

## D4-L3 Weight totals and survey checks

### Concept

Survey weights carry the whole estimate. A wrong weight gives a plausible but wrong percentage in every table. Check them on every wave, before gold.

### Totals and range:

- No null weights, and none zero or negative, unless your design allows zero.
- The weight sum matches its target: the population total for grossing weights, or the sample size for weights normalised to mean 1. Your contract's `weight_sum_min` and `weight_sum_max` encode this.
- The ratio of largest to smallest weight stays below an agreed limit. One respondent standing for 40 others dominates their subgroup.

**Efficiency.** Unequal weights cost precision. Kish's approximation measures it:

```
1 deff = n x sum(w^2) / (sum w)^2 = 1 + CV^2
2 n_eff = (sum w)^2 / sum(w^2) = n / deff
```

Here CV is the population coefficient of variation of the weights. Weights 0.5, 1.0, 1.5 and 3.0 give a deff of 1.39 and an effective sample of 2.88 out of 4. Track `n_eff` per subgroup and wave. A sudden fall means the weighting struggled, often because a quota cell was nearly empty.

**Spark details that break these checks.** Checked on PySpark 4.0.4:

- `F.stddev` is the **sample** standard deviation. It returns null for a single row, and the Kish identity uses the population version, `F.stddev_pop`.
- On an empty `DataFrame`, `mean` and `stddev` return null. Your reference `check_weight_distribution` then calls `round(None, 4)` and raises `TypeError`.
- `F.sum(F.when(F.col("reach_flag") == 1, 1))` is **null**, not 0, for a subgroup where nobody was reached. Your gold table then shows a blank instead of 0% reach. Wrap it in `F.coalesce(..., F.lit(0))`.

Decide what each result means for the run. A null weight is CRITICAL: halt. A deff above your usual range is WARN: log it and have someone look.

## Why it matters

A weighting run used last year's population targets for one age band. The weight sum stayed within tolerance overall, so the contract check passed. The 55+ band's effective sample fell from 800 to 190, and its reach estimate jumped 9 points. A per-subgroup `n_eff` check with a WARN threshold would have flagged it before the dashboard refreshed.

## Worked example: weak code

```
1 stats = df.agg(F.mean(w).alias("mean"), F.stddev(w).alias("sd")).first()
2 cv = stats["sd"] / stats["mean"] if stats["mean"] else None
3 result = {"mean": round(stats["mean"], 4), "cv": round(cv, 4) if cv else None}
```

1. **Line 1.** Sample standard deviation, and null for a single row.
2. **Line 3.** `round(None, 4)` raises `TypeError` on an empty wave. `if cv` treats a CV of 0 as missing.

## Worked example: strong code

```
1 row = df.agg(
2 F.count(w).alias("n"),
3 F.sum(w).alias("total"),
4 F.sum(F.col(w) * F.col(w)).alias("total_sq"),
5).first()
```

```

2 if row["n"] == 0 or row["total"] is None or row["total"] == 0:
7 raise ValueError("no weights to check")
8 deff = row["n"] * row["total_sq"] / row["total"] ** 2
9 n_eff = row["total"] ** 2 / row["total_sq"]

```

1. **Line 4.** Kish needs only n, sum and sum of squares, in one pass.

2. **Line 6.** Empty and zero totals are reported, never divided by.

Check nulls, range, sum and efficiency per wave and subgroup. Mind Spark's nulls: sample stddev, empty aggregates and sums over no matches.

## Recap card

- **Kish design effect and effective sample**  $deff = n \times \text{sum}(w^2) / (\text{sum } w)^2 = 1 + CV^2$ .  $n\_eff = (\text{sum } w)^2 / \text{sum}(w^2)$ .
- **F.stddev in Spark** Sample standard deviation, null for one row. Use `F.stddev_pop` for  $1 + CV^2$ .
- **sum(when(cond, 1)) with no matches** Null, not 0. Wrap in `coalesce(..., lit(0))` or add `otherwise(0)`.
- **Severity for a null weight** CRITICAL: halt before gold.

## Practice

### Exercise 1. Multiple choice (level 1)

In gold, `F.sum(F.when(F.col("reach_flag") == 1, 1))` gives what for a subgroup where no one was reached?

- a) NULL
- b) 0
- c) An error
- d) The subgroup is dropped

### Why?

- r1) when without otherwise gives null for non-matching rows, and the sum of only nulls is null.
- r2) Spark sums return null for any group with a zero.

### Exercise 2. Predict the output (level 2)

What does this print?

```

1 w = [1.0, 1.0, 1.0, 1.0]
2 n = len(w)
3 deff = n * sum(x * x for x in w) / sum(w) ** 2
4 print(deff)

```

### Exercise 3. Fill the gap (level 2)

Make the unweighted reach count 0 when no one in the subgroup was reached.

```
1 F. ____ (1) (F.sum(F.when(F.col("reach_flag") == 1, 1)),
 F.lit(____ (2))).alias("unweighted_reach")
```

### Exercise 4. Write the function (level 3)

Write

`weight_report(weights: list[float | None], max_ratio: float = 20.0) -> dict[str, object]`. Return `{"status": ..., "n": ..., "n_eff": ..., "deff": ...}`. Status is "CRITICAL" if the list is empty or any weight is None or not positive (and then `n_eff` and `deff` are None), "WARN" if max/min weight exceeds `max_ratio`, otherwise "PASS". Use `n_eff = (sum w)^2 / sum(w^2)` and `deff = n / n_eff`, each rounded to 4 decimals.

Starter code:

```
1 def weight_report(weights, max_ratio=20.0):
2 n_eff = sum(weights) ** 2 / sum(w * w for w in weights)
3 return {"status": "PASS", "n": len(weights), "n_eff": n_eff, "deff":
 len(weights) / n_eff}
```

### Exercise 5. Architecture design (level 4)

Set severities for weight checks in the silver QA step.

Options: CRITICAL: halt before gold, WARN: drop those rows and continue.

1. Any null `final_weight`
2. Subgroup `deff` 40% above its usual range
3. Standard deviation for the CV

### Sources

- Kish  $deff = n \sum(w^2) / (\sum w)^2 = 1 + CV^2$  (population CV);  $n_{eff} = (\sum w)^2 / \sum(w^2)$ ; example weights give 1.3889 and 2.88. (verified: Computed locally on CPython 3.13, September 2026; identity follows algebraically.)
- PySpark 4.0.4: `sum(when(...))` over no matches is null; `stddev` is the sample version and null for one row; mean and `stddev` of an empty DataFrame are null. (verified: Local run on PySpark 4.0.4, September 2026.)



# D5 Errors typical of AI-generated code

Roadmap: Every month. Invented APIs, deprecated calls, missing edge cases, self-confirming tests.

## D5-L1 Invented functions, parameters and syntax

### Concept

AI assistants predict plausible code, so they sometimes produce a function, parameter or statement that does not exist, or exists only on another platform. Your reference code has one: `unity_catalog.sql` uses `CREATE ROW ACCESS POLICY`, which is Snowflake syntax. Databricks uses a SQL UDF with `ALTER TABLE ... SET ROW FILTER`.

What invented code does when it runs, checked on PySpark 4.0.4 and pandas 3.0.2:

| Invented code                                          | Result                                                                                      |
|--------------------------------------------------------|---------------------------------------------------------------------------------------------|
| <code>spark.read.delta(path)</code>                    | <code>AttributeError: 'DataFrameReader' object has no attribute 'delta'</code>              |
| <code>df.drop_duplicates(keep="latest")</code>         | <code>ValueError: keep must be either "first", "last" or False</code>                       |
| <code>pd.read_csv(path, encodng="utf-8")</code>        | <code>TypeError: ... unexpected keyword argument 'encodng'. Did you mean 'encoding'?</code> |
| <code>CREATE ROW ACCESS POLICY ...</code> in Spark SQL | <code>[PARSE_SYNTAX_ERROR] Syntax error at or near 'ROW'</code>                             |
| <code>F.col("a").isnotnull()</code>                    | <code>TypeError: 'Column' object is not callable</code>                                     |

The last one is sly. The real method is `isNotNull`. PySpark treats an unknown attribute on a column as a **struct field lookup**, so `F.col("a").isnotnull` builds a valid expression for a field named `isnotnull`. Only the call fails. Without the parentheses, the typo would fail later, at analysis, far from the line that caused it.

Most invented names fail loudly, which makes them cheap to catch **if you run the code**. The expensive cases run: a parameter that exists but means something else, or a function from a newer version than your runtime.

## Check every API an assistant gives you:

1. Find it in the official docs **for your version**: Databricks Runtime 17.3 LTS, PySpark 4.0, pandas 3.0.
2. In a notebook, confirm it exists: `hasattr(F, "try_divide")`, `help(pd.DataFrame.merge)`, or `inspect.signature(...)`.
3. For SQL, check the statement in the Databricks SQL reference, not a general SQL tutorial. Syntax from Snowflake, T-SQL and PostgreSQL looks alike and differs.

Treat an unfamiliar name as unverified until a step confirms it.

## Why it matters

An assistant proposed `CREATE ROW ACCESS POLICY` to restrict survey rows by region. The engineer pasted it into the migration script without running it. The deployment failed in production at 2 a.m. with a parse error, and the table stayed unrestricted until morning. Checking the statement against the Databricks SQL reference would have taken two minutes.

## Worked example: weak code

```
❶ CREATE ROW ACCESS POLICY region_policy
2 AS (region STRING) RETURNS BOOLEAN ->
❷ is_member('analysts_' || region);
4 ALTER TABLE research_silver.nat_respondents ADD ROW ACCESS POLICY
region_policy ON (region);
```

1. **Line 1.** Snowflake syntax. Databricks fails with `PARSE_SYNTAX_ERROR`.
2. **Line 3.** `is_member` checks workspace groups. Unity Catalog guidance favours `is_account_group_member`.

## Worked example: strong code

```
❶ CREATE FUNCTION research_governance.region_filter(region STRING)
2 RETURN is_account_group_member(concat('analysts_', region));
3
4 ALTER TABLE research_silver.nat_respondents
❷ SET ROW FILTER research_governance.region_filter ON (region);
```

1. **Line 1.** A SQL UDF returning a boolean, in a three-level name.
2. **Line 5.** The Databricks statement for row filters (Unity Catalog).

Look up every function, parameter and statement for your exact version, and run it on small data before it reaches a pipeline.

## Recap card

- **CREATE ROW ACCESS POLICY** Snowflake syntax. Databricks: SQL UDF plus ALTER TABLE ... SET ROW FILTER.
- **F.col('a').isNotNull()** Unknown attribute becomes a struct field lookup, then 'Column' object is not callable. The method is isNotNull.
- **Three ways to confirm an API exists** Official docs for your version, hasattr or help() in a notebook, inspect.signature for parameters.
- **The costly invented-API case** Code that runs: a real parameter with a different meaning, or a function from a newer version.

## Practice

### Exercise 1. Multiple choice (level 1)

An assistant writes `spark.read.delta("/Volumes/research/raw/nat")` for Databricks. What happens?

- a) `AttributeError: DataFrameReader has no attribute 'delta'`
- b) It reads the Delta table
- c) It reads the files as Parquet
- d) A warning, then the read succeeds

### Why?

- r1) The reader has `format('delta').load(path)` and `table(name)`, but no `delta()` method.
- r2) Delta tables cannot be read by path.

### Exercise 2. Predict the output (level 2)

What is printed? (PySpark 4.0)

```
1 from pyspark.sql import functions as F
2
3 expr = F.col("region").isNotNull
4 print(type(expr).__name__)
```

### Exercise 3. Spot the bug (level 2)

An assistant wrote this Databricks SQL to mask the postcode column for non-analysts. Tap the line that does not exist in Databricks, then rewrite it with the Databricks statement.

```
1 CREATE FUNCTION research_governance.mask_postcode(p STRING)
2 RETURN CASE WHEN is_account_group_member('analysts') THEN p ELSE '***'
END;
```

```
3 CREATE MASKING POLICY postcode_mask ON research_silver.nat_respondents
 (postcode);
```

Write down the line number of the bug. Then write the fix.

### Exercise 4. Write the function (level 3)

Write `unknown_params(func: Callable[..., object], params: list[str]) -> list[str] | None`. Use `inspect.signature` to return the names in `params` that `func` does not accept, in the order given. If `func` accepts `**kwargs`, return `None`: you cannot tell from the signature alone.

Starter code:

```
1 def unknown_params(func, params):
2 return []
```

### Exercise 5. Architecture design (level 4)

Set the team rule for APIs suggested by an assistant.

Options: Link the official docs page for our version in the pull request, Accept it if the assistant is confident.

1. Before code with a new function or parameter is merged
2. SQL statements such as row filters and masks
3. A function you cannot find in the docs

### Sources

- PySpark 4.0.4: `spark.read.delta` and `DataFrameWriter.saveAsDelta` raise `AttributeError`; `F.col('a').isnotnull` is a `Column` and calling it raises `TypeError`; `CREATE ROW ACCESS POLICY` raises `PARSE_SYNTAX_ERROR`. (verified: Local run on PySpark 4.0.4, September 2026.)
- pandas 3.0.2: `drop_duplicates(keep='latest')` raises `ValueError`; `read_csv(encodng=...)` raises `TypeError` with a 'Did you mean' hint. (verified: Local run on pandas 3.0.2 with CPython 3.13, September 2026.)
- Databricks row filters use a SQL UDF and `ALTER TABLE ... SET ROW FILTER`; column masks use `ALTER TABLE ... ALTER COLUMN ... SET MASK`. (verified: <https://docs.databricks.com/aws/en/data-governance/unity-catalog/filters-and-masks/manually-apply> (checked 2026-09-25))

# D5-L2 Deprecated APIs and version mismatches

## Concept

An assistant learned from years of code written for older versions. It often suggests an API that worked in 2021 and has since been deprecated, renamed or removed. Your reference `observability.py` uses `datetime.utcnow()`, deprecated from Python 3.12.

Examples that matter for this course, checked in September 2026:

| Old code                            | Current state                                              | Use instead                             |
|-------------------------------------|------------------------------------------------------------|-----------------------------------------|
| <code>datetime.utcnow()</code>      | Deprecated, Python 3.12                                    | <code>datetime.now(timezone.utc)</code> |
| <code>df.append(other)</code>       | Removed from pandas (gone in 3.0)                          | <code>pd.concat([df, other])</code>     |
| <code>df.applymap(f)</code>         | Removed in pandas 3.0                                      | <code>df.map(f)</code>                  |
| <code>df["a"][0] = x</code>         | Never updates <code>df</code> under pandas 3 Copy-on-Write | <code>df.loc[0, "a"] = x</code>         |
| <code>import distutils</code>       | Removed in Python 3.12                                     | <code>packaging, setuptools</code>      |
| <code>input_file_name()</code>      | Not supported from Databricks Runtime 17.3 LTS             | <code>_metadata.file_path</code>        |
| <code>mssparkutils</code> in Fabric | Being retired                                              | <code>notebookutils</code>              |

The chained assignment row is the dangerous one. `pandas 3.0.2` emits a `ChainedAssignmentError` warning and leaves the `DataFrame` unchanged. If warnings are hidden, the update simply does not happen.

**Version mismatches** cause the opposite problem: code for a **newer** version than yours. A function added in `PySpark 4.1` fails on `Databricks Runtime 17.3 LTS`, which ships `Spark 4.0`. `Pyodide` in this app runs `pandas 3.0.2`. Your laptop might run `3.0.6`.

Defend in three ways:

- **State the versions** in every prompt, and in the project: `VERSIONS.md`, `pyproject.toml` pins, the runtime in the job definition.
- **Show warnings** in tests. Run `pytest` with `-W error::DeprecationWarning` for your own code, so a deprecated call fails the build.
- **Compare versions as numbers**, never as strings. `"3.10" < "3.9"` is `True` as text.

## Why it matters

An assistant fixed a data cleaning function by adding `df["weight"][mask] = 1.0` to cap weights. Tests passed on a developer laptop with pandas 1.5. In the pandas 3 environment, the cap never happened. Weights above 5 reached gold unchanged for four waves. The warning had been filtered out in the job's logging configuration.

## Worked example: weak code

```
1 import pandas as pd
2
3 def cap_weights(df: pd.DataFrame, cap: float) -> pd.DataFrame:
4 """Cap weights at cap."""
5 df["final_weight"][df["final_weight"] > cap] = cap
6 return df.applymap(lambda v: v)
```

1. **Line 5.** Chained assignment. Under pandas 3 Copy-on-Write it never changes df.
2. **Line 6.** `applymap` no longer exists in pandas 3.0: `AttributeError`.

## Worked example: strong code

```
1 import pandas as pd
2
3 def cap_weights(df: pd.DataFrame, cap: float) -> pd.DataFrame:
4 """Return a copy with final_weight capped at cap. Written for pandas
5 3.0."""
6 out = df.copy()
7 out.loc[out["final_weight"] > cap, "final_weight"] = cap
8 return out
```

1. **Line 4.** The docstring states the target version.
2. **Line 6.** One `.loc` call: row mask and column together. Works in every pandas version.

Assume suggested code targets some other version. Check deprecations for yours, turn warnings into test failures, and compare versions numerically.

## Recap card

- **Chained assignment in pandas 3** `df['a'][i] = x` never updates df (Copy-on-Write). Use `df.loc[i, 'a'] = x`.
- **pandas 3 removals** `DataFrame.append` (use `pd.concat`) and `applymap` (use `DataFrame.map`).
- **input\_file\_name() on DBR 17.3 LTS** Not supported. Use `_metadata.file_path`.
- **Comparing version strings** Never as text: `'3.10' < '3.9'`. Compare integer tuples or use `packaging.version`.

## Practice

### Exercise 1. Multiple choice (level 1)

An assistant adds `F.input_file_name()` to your bronze job on Databricks Runtime 17.3 LTS. What do you use instead?

- a) The `_metadata.file_path` column
- b) Keep `input_file_name()`, it is fine
- c) `os.path.basename` on the DataFrame
- d) A Python UDF reading the path

#### Why?

- r1) Databricks release notes state `input_file_name` is no longer supported from 17.3 LTS. The `file_metadata` column replaces it.
- r2) `input_file_name` was renamed `file_path`.

### Exercise 2. Predict the output (level 2)

What does this print with pandas 3.0?

```
1 import pandas as pd
2
3 df = pd.DataFrame({"w": [1.0, 9.0]})
4 df["w"][1] = 5.0
5 print(df["w"].tolist())
```

### Exercise 3. Fill the gap (level 2)

Update this code from pandas 1.x to pandas 3.0.

```
1 combined = pd.____(1)([waves, new_wave], ignore_index=True)
2 labels = combined.____(2)(str.strip)
```

### Exercise 4. Write the function (level 3)

Write `version_at_least(installed: str, required: str) -> bool` comparing dotted versions such as `"3.0.2"` and `"3.10"` numerically. Ignore anything after the numeric part of each segment (`"4.0.4rc1"` counts as `4.0.4`). Missing segments count as 0, so `"3.0"` equals `"3.0.0"`. Raise `ValueError` if a segment has no leading digits.

Starter code:

```
1 def version_at_least(installed, required):
2 return installed >= required
```

## Exercise 5. Spot the bug (level 3)

This function caps weights, but under pandas 3 the weights come back unchanged. Tap the line, then fix it without changing the caller's DataFrame.

```
1 import pandas as pd
2
3 def cap_weights(df: pd.DataFrame, cap: float) -> pd.DataFrame:
4 """Return a copy with final_weight capped at cap."""
5 out = df.copy()
6 out["final_weight"][out["final_weight"] > cap] = cap
7 return out
```

Write down the line number of the bug. Then write the fix.

## Sources

- pandas 3.0.2: no DataFrame.append or applymap; DataFrame.map exists; default string dtype 'str'; chained assignment emits ChainedAssignmentError and leaves the frame unchanged. Python 3.13: distutils missing; utcnow deprecated. (verified: Local run on pandas 3.0.2 with CPython 3.13.12, September 2026.)
- input\_file\_name is no longer supported from Databricks Runtime 17.3 LTS. (verified: <https://docs.databricks.com/aws/en/release-notes/runtime/17.3lts> (checked 2026-09-25))
- mssparkutils is renamed NotebookUtils and will be retired. (verified: <https://learn.microsoft.com/en-us/fabric/data-engineering/notebook-utilities> (see lesson C5-L1))

## D5-L3 Five review checks for AI-written code

---

### Concept

Code from an assistant usually runs. The problems hide in what it leaves out. Review it with five checks.

**1. Missing edge cases.** Assistants write the happy path. Ask what happens with an empty DataFrame, a zero total, a null, a duplicate key, a single row. Your reference `check_weight_distribution` fails on an empty wave for exactly this reason (lesson D4-L3).

**2. Self-confirming tests.** A test that computes its expected value with the same formula as the code proves nothing:

```
1 def test_share():
2 assert share(3, 4) == 3 / 4 # repeats the code
```



If the formula is wrong, the test is wrong the same way. Write expected values by hand, from small data you understand: `assert share(3, 4) == 0.75`, `assert share(0, 0) is None`.

**3. Secrets in code.** Assistants fill in placeholders that look like real tokens, or keep the one you pasted into the prompt. A token in a notebook ends up in Git history. Read secrets at run time: `dbutils.secrets.get(scope, key)` on Databricks, `notebookutils.credentials.getSecret(vault_url, name)` in Fabric, or an environment variable on your laptop.

**4. Over-engineering.** A 12-line task arrives as an abstract base class, a factory and a config object. Every extra layer is code to test and review. Ask for the simplest version that meets the requirement, then add structure when a second use case appears.

### 5. The wrong filter context in DAX.

`CALCULATE(SUM(fact_reach[weighted_reach]), ALL(fact_reach))` removes **every** filter on the fact table, including wave and region slicers. A "share of total" measure then divides by the all-time, all-region total. Remove only the filter the measure means to ignore: `ALL(dim_brand)` for a share across brands.

Each check takes a minute. Each catches a class of bug that runs without error.

## Why it matters

An assistant wrote a reach function and its tests in one reply. All 12 tests passed. Each test computed its expected value with the same expression as the function, including the same missing `is None` check. A reviewer replaced three expected values with hand-worked numbers, and two tests failed at once: the function returned 0.0 where it should have returned None.

## Worked example: weak code

```
1 API_TOKEN = "paste-your-token-here-EXAMPLE"
2
3 def share(part: float, total: float) -> float:
4 return part / total
5
6 def test_share():
7 assert share(3, 4) == 3 / 4
```

- Line 1.** A secret in code. The real token replaces the placeholder and lands in Git.
- Line 4.** No guard for a zero total.

3. **Line 7.** The expected value repeats the code. It cannot catch a wrong formula.

## Worked example: strong code

```
1 import os
2
3 def share(part: float, total: float) -> float | None:
4 """part / total, or None when total is 0."""
5 return None if total == 0 else part / total
6
7 def api_token() -> str:
8 """Read the token at run time, never from source code."""
9 return os.environ["SURVEY_API_TOKEN"]
10
11 def test_share():
12 assert share(3, 4) == 0.75
13 assert share(0, 0) is None
```

1. **Line 5.** The edge case is handled and documented.

2. **Line 9.** The secret stays outside the code. On Databricks use `dbutils.secrets.get`.

3. **Line 13.** Hand-written expected values, including the edge case.

Review AI code for edge cases, self-confirming tests, secrets, needless structure and, in DAX, filters removed too widely.

## Recap card

- **Five review checks for AI code** Edge cases, self-confirming tests, secrets, over-engineering, DAX filter context.
- **Self-confirming test** Expected value computed with the code's own formula. Use hand-worked values instead.
- **Where secrets live** Secret scopes or Key Vault, read at run time. Never in code or prompts.
- **ALL(fact\_table) in a share measure** Removes every filter, including wave and region slicers. Use ALL on the one dimension you mean.

## Practice

### Exercise 1. Multiple choice (level 1)

An assistant writes `assert weighted_pct(r, b) == r / b * 100` as the test for `weighted_pct`. What is wrong with it?

- a) It repeats the implementation, so a wrong formula passes
- b) It should use `pytest.approx`
- c) Tests must not use multiplication

- d) Nothing: it checks the output

### Why?

- r1) A test's expected value must come from outside the code: small data worked by hand.
- r2) Tests should never call the function they test.

### Exercise 2. Spot the bug (level 2)

An assistant produced this Databricks helper. Tap the line that must never be committed, then fix it to read the value from the secret scope `survey-api`, key `token`.

```
1 import requests
2
3 API_TOKEN = "tok_EXAMPLE_do_not_use_0000"
4
5 def fetch_wave(wave_id: str) -> dict[str, object]:
6 """Fetch wave metadata from the survey platform."""
7 headers = {"Authorization": f"Bearer {API_TOKEN}"}
8 return requests.get(f"https://survey.example/api/waves/{wave_id}",
9 headers=headers, timeout=30).json()
```

Write down the line number of the bug. Then write the fix.

### Exercise 3. Multiple choice (level 2)

A report has a wave slicer set to W14. An assistant's measure is `Share = DIVIDE(SUM(fact_reach[weighted_reach]),`

`CALCULATE(SUM(fact_reach[weighted_reach]), ALL(fact_reach)))`. What is the denominator?

- a) Weighted reach for all waves, all regions and all brands
- b) Weighted reach for W14 across all brands
- c) Weighted reach for W14 for the current brand
- d) An error, because ALL needs a column

### Why?

- r1) `ALL(fact_reach)` removes every filter on the fact table, including the wave slicer.
- r2) `CALCULATE` always ignores slicers.

### Exercise 4. Write the function (level 3)

An assistant wrote `weighted_mean(values, weights)` as `sum(v * w for v, w in zip(values, weights)) / sum(weights)`. Rewrite it with the edge cases it missed. Return `None` if the lists are empty or the weight total is 0. Skip pairs

where the value or the weight is `None`. Raise `ValueError` if the lists differ in length or any weight is negative.

Starter code:

```
1 def weighted_mean(values, weights):
2 return sum(v * w for v, w in zip(values, weights)) / sum(weights)
```

## Exercise 5. Refactor (level 4)

An assistant answered "flag weights above a cap" with the code below.

Replace it with one function, `flag_high_weights(weights: list[float], cap: float) -> list[int]`, returning the indexes of weights above `cap`. No classes.

Starter code:

```
1 from abc import ABC, abstractmethod
2
3
4 class WeightRule(ABC):
5 @abstractmethod
6 def check(self, weight):
7 ...
8
9
10 class CapRule(WeightRule):
11 def __init__(self, cap):
12 self.cap = cap
13
14 def check(self, weight):
15 return weight > self.cap
16
17
18 class RuleFactory:
19 @staticmethod
20 def create(kind, **kwargs):
21 if kind == "cap":
22 return CapRule(**kwargs)
23 raise ValueError(kind)
24
25
26 def flag_high_weights(weights, cap):
27 rule = RuleFactory.create("cap", cap=cap)
28 return [i for i, w in enumerate(weights) if rule.check(w)]
```

## Sources

- `dbutils.secrets.get(scope, key)` reads a secret and values are redacted in notebook output. (verified: <https://docs.databricks.com/aws/en/security/secrets/> (checked 2026-09-25))

- In DAX, ALL(table) inside CALCULATE removes all filters on that table; ALL on one dimension keeps other filters. (verified: Taught and tested in lesson C6-L2; DAX CALCULATE and ALL semantics per Microsoft Learn DAX reference.)

## D6 Verifying AI output

Roadmap: Every month. A workflow to check AI code before you trust it.

### D6-L1 Read before you run, and check each API

---

#### Concept

Track D so far showed what goes wrong in AI-written code. Module D6 is the routine that catches it. Step one happens before the code runs at all.

**Why read first.** Running code is the fastest check for invented names, but running is not free. A suggested cell might overwrite a table, drop a column or read a whole year of data. Reading takes two minutes. Restoring a bronze table takes a morning, if you can do it at all.

#### The read-through:

1. **Restate the goal** in one sentence before you look at the code. "Append wave W14 deliveries to bronze with audit columns." You now have something to check against.
2. **List every external name:** imports, functions, methods, parameters, tables and columns. An `ast` walk lists the calls for you (exercise 4).
3. **Check each name** against the docs for your version and against the table schema. `DESCRIBE TABLE`, `df.columns`, `hasattr`, `inspect.signature` and `help()` answer most questions in seconds.
4. **Find side effects:** `mode("overwrite")`, `DROP`, `DELETE`, `TRUNCATE`, `VACUUM`, writes to production names. Each needs a reason that matches the goal.
5. **Check the scope:** filters on tracker and wave, no `collect()` on large data, no loop that calls `count()` per item.

Then run it, on development data first, in a dev schema: `research_dev`, not `research_bronze`.

**Where assistants go wrong most** in this course's work: the write mode, three-level table names, the filter keys, and parameters that exist but mean something else. `mode("overwrite")` on an append-only bronze table replaces every earlier delivery. It runs without error.

A read-through does not prove code correct. It removes the mistakes that are visible on the page, so testing (lesson D6-L2) can focus on the ones that are not.

## Why it matters

An assistant rewrote a bronze ingestion cell to "simplify" it. The new cell used `mode("overwrite")`. It ran in production and succeeded, replacing 38 waves of append-only deliveries with one. Delta time travel recovered the table because retention had not yet expired. A two-minute read, looking for side effects, would have found the one word that mattered.

## Worked example: weak code

```
❶ # Suggested by an assistant: "load the new delivery"
2 delivery = spark.read.format("csv").option("header", True).load(path)
❷ delivery.write.mode("overwrite").saveAsTable("research_bronze.nat_raw")
```

1. **Line 1.** No stated goal to check against.
2. **Line 3.** Overwrite on an append-only table replaces every earlier delivery. It runs cleanly.

## Worked example: strong code

```
❶ # Goal: append the W14 delivery to bronze with audit columns. Bronze is
 # append-only.
2 # Checked: format("csv").option("header", True) and saveAsTable (PySpark
 # 4.0 API reference).
3 from pyspark.sql import functions as F
4
5 delivery = spark.read.format("csv").option("header", True).load(path)
6 delivery = delivery.withColumn("source_file",
 F.col("_metadata.file_path"))
❷ delivery.write.mode("append").saveAsTable("research_dev.nat_raw")
```

1. **Line 1.** The goal, written first, makes the write mode checkable.
2. **Line 7.** Append, and into a dev schema until the run has been checked.

Restate the goal, list and check every name, hunt for side effects, then run on development data.

## Recap card

- **Five steps of the read-through** Restate the goal, list names, check each name, find side effects, run on dev data.
- **Side effects to hunt for** `mode('overwrite')`, `DROP`, `DELETE`, `TRUNCATE`, `VACUUM`, writes to production names.
- **Quick existence checks in a notebook** `hasattr` for names, `inspect.signature(...).parameters` for parameters, `DESCRIBE TABLE` or `df.columns` for columns.
- **What a read-through does not do** Prove correctness. It clears visible mistakes so tests can find the rest.

## Practice

### Exercise 1. Multiple choice (level 1)

An assistant gives you a cell that writes to `research_bronze.nat_raw`. What do you check first?

- a) The write mode and the target, against the goal and the table's rules
- b) Whether the code is formatted with ruff
- c) How long it will take
- d) Nothing: run it and read the result

### Why?

- r1) Side effects on a production table cannot always be undone. Bronze is append-only.
- r2) Assistants always pick the wrong write mode.

### Exercise 2. Put in order (level 2)

Order the read-through of an AI-written cell.

- Restate the goal in one sentence
- Find side effects: write modes, DROP, DELETE
- Check each name against the docs for your version and the schema
- Run it on development data in a dev schema
- List every function, parameter, table and column

### Exercise 3. Fill the gap (level 2)

In a notebook, confirm that a suggested function and parameter exist.

```
1 import inspect
2 from pyspark.sql import functions as F
3
4 ____ (1)(F, "try_divide")
5 "validate" in inspect.____ (2)(pd.DataFrame.merge).parameters
```

### Exercise 4. Write the function (level 3)

Write `called_names(source: str) -> list[str]` that parses Python source with `ast` and returns the dotted names of every call, such as `spark.read.format` or `F.col`, unique and sorted. Ignore calls whose target is not a name or attribute chain (for example a call on a subscript). This gives you the checklist for step 2 of the read-through.

Starter code:

```
1 def called_names(source):
```



```
2 return [line for line in source.splitlines() if "(" in line]
```

### Exercise 5. Spot the bug (level 3)

Goal: append this wave's delivery to bronze, which is append-only. Tap the line that breaks the goal, then fix it.

```
1 from pyspark.sql import functions as F
2
3 delivery = spark.read.format("csv").option("header", True).load(path)
4 delivery = delivery.withColumn("ingestion_timestamp",
5 F.current_timestamp())
5 delivery.write.mode("overwrite").saveAsTable("research_bronze.nat_raw")
```

Write down the line number of the bug. Then write the fix.

### Sources

- PySpark DataFrameReader.format/option/load, DataFrameWriter.mode and saveAsTable, and the `_metadata.file_path` column exist as used. (verified: Local runs on PySpark 4.0.4 (see D1 and D5 claims); `_metadata.file_path` per Databricks Auto Loader docs (A8-L2).)

## D6-L2 Small data with a known answer, failing test first

### Concept

After the read-through, run the code where you already know the right answer. That means small data you built by hand, and a test written **before** you accept the code.

**Build the known-answer data.** Five to ten rows, chosen so every rule in the requirement matters:

- a normal case you can compute in your head,
- a duplicate key, a null, a zero total, an empty group,
- a value exactly on a boundary, such as a weight equal to the cap.

Write the expected output by hand, next to the input. If you cannot work out the answer by hand, the requirement is not clear enough yet, and neither the assistant nor you should write code.

**Write the test first, and watch it fail.** The order matters:

- Write the test with the hand-worked expected values.
- Run it against the current code, or a stub. It must **fail**. A test that passes before the feature exists tests nothing.
- Accept or write the code. Run the test. It passes.

4. Break the code on purpose, say change `>` to `>=`. The test must fail again.  
If it still passes, it does not cover that rule.

Step 4 is a quick form of mutation testing. It tells you whether your tests would catch the mistakes an assistant makes.

**For Spark code**, run the same idea on a local `SparkSession`.

`pyspark.testing.assertDataFrameEqual(actual, expected)` compares two DataFrames. In PySpark 4.0 it ignores row order by default (check `checkRowOrder=False`) and compares floats with `rto1=1e-05`.  
`assertSchemaEqual` compares schemas.

Keep the known-answer tests in the repository. They check every later change, including the next suggestion from an assistant.

## Why it matters

An assistant wrote a deduplication step for silver and a test for it. The test passed. A reviewer changed the dedup key from `(respondent_id, wave_id)` to `respondent_id` alone, a real bug, and the test still passed. Its fixture had only one wave. Adding a respondent who appears in two waves made the test fail on the broken key, as it should.

## Worked example: weak code

```
1 def test_share():
2 result = share(30, 120)
3 assert result is not None
4 assert 0 <= result <= 1
```

1. **Line 3.** Passes for almost any implementation, including wrong ones.
2. **Line 4.** A range, not an answer. 0.3 would pass as well as 0.25.

## Worked example: strong code

```
1 import pytest
2
3 @pytest.mark.parametrize(
4 ("part", "total", "expected"),
5 [
6 (30, 120, 0.25), # normal case, worked by hand
7 (0, 50, 0.0), # zero share is a real value
8 (5, 0, None), # zero total has no share
9 (50, 50, 1.0), # boundary
10],
11)
12 def test_share(part: float, total: float, expected: float | None) -> None:
13 assert share(part, total) == expected
```

1. **Line 6.** An exact answer you computed yourself.

2. **Line 8.** The edge case the requirement names.

Hand-built data, hand-worked answers, a test that fails first, and a deliberate break to prove the test bites.

## Recap card

- **Known-answer data** 5 to 10 hand-built rows: normal case, duplicate, null, zero total, empty group, boundary.
- **Test-first order** Write the test, watch it fail, accept the code, watch it pass, break the code on purpose, watch it fail.
- **assertDataFrameEqual defaults (PySpark 4.0)** `checkRowOrder=False`, `rtol=1e-05`, `atol=1e-08`.
- **Why break code on purpose** If the test still passes, it misses that rule. A quick form of mutation testing.

## Practice

### Exercise 1. Multiple choice (level 1)

You write a test for a feature that does not exist yet, and it passes. What does that tell you?

- a) The test does not check the feature
- b) The feature already works
- c) The test is well written
- d) The code is fast

### Why?

- r1) A test must fail when the behaviour it checks is missing. Passing without the feature means it asserts something else.
- r2) Tests always fail the first time they run.

### Exercise 2. Multiple choice (level 2)

Your tests pass. You change `w > cap` to `w >= cap` in the code on purpose, and the tests still pass. What do you add?

- a) A case with a weight exactly equal to the cap
- b) More cases with weights far above the cap
- c) A timeout on the test
- d) Nothing: the tests pass

## Why?

- r1) > and >= differ only on the boundary, so only a boundary case can tell them apart.
- r2) Mutations always need more data.

### Exercise 3. Fill the gap (level 2)

Compare a Spark result with the expected rows in a local test, ignoring row order.

```
1 from pyspark.testing import ____ (1)
2
3 expected = spark.createDataFrame([("W01", 0.25)], "wave_id string, share
 double")
4 assertDataFrameEqual(result, expected, ____ (2)=False)
```

### Exercise 4. Write the function (level 3)

Write the known-answer cases before the code. Write `share_cases()` -> `list[tuple[float, float, float | None]]` returning `(part, total, expected)` cases for `share(part, total)`, which returns `part / total`, or `None` when `total` is 0. Your cases must all pass against a correct `share`, and each of three wrong versions must fail at least one of them: one that raises on a zero total, one that returns 0.0 on a zero total, and one that returns a percentage (times 100).

Starter code:

```
1 def share_cases():
2 return [(1.0, 2.0, 0.5)]
```

### Exercise 5. Architecture design (level 4)

An assistant offers a silver deduplication step keyed on respondent and wave. Plan the known-answer test.

Options: A duplicate in one wave, the same respondent in two waves, a null key, an unchanged respondent, 1,000 rows copied from production.

1. Fixture rows
2. Expected output
3. Proving the test bites

## Sources

- `pyspark.testing.assertDataFrameEqual` signature in 4.0.4 has `checkRowOrder=False`, `rtol=1e-05`, `atol=1e-08`; `assertSchemaEqual` exists. (verified: inspect.signature on PySpark 4.0.4, September 2026.)

## D6-L3 Ask for assumptions, then compare a second approach

---

### Concept

Every piece of code assumes things about its input. Code from an assistant assumes things nobody wrote down. The last two steps of the verification routine bring them into the open.

**Ask for the assumptions.** After the assistant writes code, ask: "List every assumption this code makes about its input data, one per line." Typical answers for a pipeline step:

- `respondent_id` is unique within a wave.
- `final_weight` is never null and always positive.
- `wave_id` looks like W plus two digits.
- Timestamps are in UTC.
- The lookup table has one row per key.

Then **check each one against the real data**, and turn the important ones into checks that run with the pipeline. An assumption you have not checked is a bug waiting for the wave where it stops being true. The list is also a review aid: if an assumption surprises you, the code does something you did not ask for.

The assistant can miss assumptions too. Add your own from the D3 and D4 lessons: nulls in filters, duplicate keys in joins, empty groups, time zones.

**Compare a second approach.** Compute the same result by an independent route, and compare:

- Spark SQL against the PySpark DataFrame version.
- pandas on a small extract against Spark on the same rows.
- A total from the gold table against a direct aggregate of silver.

"Independent" matters. Asking the same assistant to rewrite the same logic tends to repeat the same mistake. Write the second version yourself, or derive it from a different table.

When the two disagree, do not pick the one you like. Run both on your known-answer data (lesson D6-L2). The one that matches the hand-worked answer is right, and the other has a bug worth understanding.

## Why it matters

An assistant's gold aggregation and an analyst's hand-written SQL disagreed by 0.8 points of reach for one region. On the known-answer fixture, the SQL matched the hand calculation. The assistant's version used an inner join to the region lookup, which dropped 212 respondents with a region code missing from the lookup. Asked for its assumptions, the assistant had listed "every region code exists in the lookup". Nobody had checked it.

## Worked example: weak code

```
1 # Assistant's code, accepted after it ran
❶ gold = silver.join(regions,
 "region_code").groupBy("region").agg(F.sum("final_weight"))
3 gold.write.mode("overwrite").saveAsTable("research_gold.nat_region_totals")
```

1. **Line 2.** An inner join assumes every code is in the lookup. Nobody listed or checked that.

## Worked example: strong code

```
1 # Assumption: every region_code in silver exists in regions. Checked:
❶ missing = silver.join(regions, "region_code", "left_anti").count()
3 if missing:
4 raise ValueError(f"{missing} respondents have a region code missing
 from the lookup")
5
6 gold = silver.join(regions, "region_code",
 "left").groupBy("region").agg(F.sum("final_weight").alias("w"))
7 # Second approach: the grand total must equal silver's weight total.
8 gold_total = gold.agg(F.sum("w")).first()[0]
9 silver_total = silver.agg(F.sum("final_weight")).first()[0]
❷ if abs(gold_total - silver_total) > 1e-6 * silver_total:
11 raise ValueError(f"gold {gold_total} != silver {silver_total}")
```

1. **Line 2.** The assumption, turned into a check that runs every time.
2. **Line 10.** An independent total. A dropped or duplicated row makes them disagree.

Make the assistant state its assumptions, check them on real data, keep the important ones as checks, and confirm results by an independent route.

## Recap card

- **Prompt for assumptions** 'List every assumption this code makes about its input data, one per line.' Then check each on real data.
- **Common data assumptions to check** Unique keys per grain, non-null positive weights, id formats, UTC timestamps, one lookup row per key.

- **An independent second approach** SQL vs DataFrame code, pandas vs Spark on an extract, gold totals vs silver aggregates. Not the same assistant rewriting itself.
- **When two methods disagree** Run both on known-answer data. The one matching the hand-worked answer is right.

## Practice

### Exercise 1. Multiple choice (level 1)

An assistant lists "respondent\_id is unique within a wave" as an assumption of its code. What do you do?

- a) Check it on the real data, and add a uniqueness check if it matters
- b) Accept it: the assistant knows the data
- c) Delete duplicates so the assumption holds
- d) Ignore it: assumptions are documentation only

### Why?

- r1) An assumption is a claim about data. Only the data can confirm it, and it can stop holding in a later wave.
- r2) Assistants list false assumptions on purpose.

### Exercise 2. Predict the output (level 2)

Two methods compute a region's weighted total. What does this print?

```
1 method_a = 1250.0 # assistant's PySpark, inner join
2 method_b = 1310.0 # your SQL, left join
3 known_answer = 1310.0
4 print("a" if method_a == known_answer else "b")
```

### Exercise 3. Fill the gap (level 2)

Check the assumption "every region\_code in silver exists in regions".

```
1 missing = silver.join(regions, "region_code", "____(1)".____(2)()
```

### Exercise 4. Write the function (level 3)

Write `violated_assumptions(rows: list[dict[str, object]]) -> list[str]` that checks three assumptions and returns the names of those that fail, in this order: "unique\_respondent\_per\_wave" (no two rows share `respondent_id` and `wave_id`), "positive\_weight" (every `final_weight` is a number above 0, not None), "wave\_format" (every `wave_id` is a string matching `W` followed by exactly two digits). An empty list of rows violates nothing.

Starter code:

```
1 def violated_assumptions(rows):
2 return []
```

### Exercise 5. Architecture design (level 4)

Choose the second approach for verifying an assistant's PySpark gold aggregation of weighted reach by region and wave.

Options: Your own Spark SQL query over silver, written without looking at the PySpark code, Ask the same assistant to rewrite it in SQL.

1. Independent method
2. The two results differ for one region
3. After the fix

### Sources

- left\_anti joins return left rows with no match; used here to check a lookup assumption. (verified: Local run on PySpark 4.0.4 (see D3-L1 claims), September 2026.)



## D7 Review practice

Roadmap: Every month. 25 AI-style snippets with seeded bugs and a hidden answer key.

### D7-L1 Review practice: Python and pandas

---

#### Concept

This lesson is practice. Each exercise shows a short snippet in the style an assistant writes: tidy, plausible and typed. Each holds one seeded bug, drawn from your reference pipeline and from the lessons in Track D. The answer key stays hidden until you submit.

#### How to review each snippet:

1. Read the prompt's goal first. The bug is a gap between the goal and the code.
2. Walk the lines in order. For each, ask: what does this assume, and what input breaks it?
3. Tap the line you would change. Then write the fix, and run it against the checker.

#### What to look for in Python and pandas code:

- Names used but never imported, and imports that shadow built-ins.
- Deprecated calls: `datetime.utcnow()`, chained assignment, `applymap`.
- Type hints that lie: `str = None`.
- Truthiness where 0, 0.0 or an empty string is a valid value.
- Mutable default arguments.
- Joins that drop or duplicate rows, and nulls that behave unexpectedly in filters.
- Text compared where numbers were meant, such as wave ids.

Your score counts two things: bugs found (the right line) and fixes proposed (a fix that passes the checker). Aim to find the bug before you open a hint. When a hint helps, note which one: that is the check you skipped.

Work at your own pace. A snippet you get wrong is worth more than one you get right, as long as you read the explanation.

## Why it matters

Reviewing is a separate skill from writing. Engineers who review AI output well do it by habit: goal first, then line by line, asking what each line assumes. Short, deliberate practice on seeded bugs builds that habit faster than waiting for real bugs to appear in production.

### Worked example: weak code

```
1 def latest_wave(wave_ids: list[str]) -> str:
2 """Return the most recent wave."""
3 return max(wave_ids)
```

1. **Line 3.** Text comparison: `max(['W9', 'W10'])` is 'W9'.

### Worked example: strong code

```
1 def latest_wave(wave_ids: list[str]) -> str:
2 """Return the most recent wave. Wave ids look like 'W' plus a
 number."""
3 return max(wave_ids, key=lambda w: int(w[1:]))
```

1. **Line 3.** Compares the numbers, so W10 beats W9.

Goal first, then line by line: what does each line assume, and which input breaks it?

## Recap card

- **Review routine for a snippet** Goal first, then each line: what does it assume, and which input breaks it?
- **Python bugs seeded in this lesson** Missing import, mutable default, `utcnow`, `str = None`, truthiness of 0.0, inner join, text max.
- **When truthiness is fine** Testing a list or dict for emptiness. Not for numbers where 0 is valid.

## Practice

### Exercise 1. Spot the bug (level 1)

Goal: return a new unique run id for the run log. Find the bug and fix it.

```
1 def new_run_id() -> str:
2 """Return a new unique run id."""
3 return str(uuid.uuid4())
```

Write down the line number of the bug. Then write the fix.

### Exercise 2. Predict the output (level 1)

Goal: each call starts with an empty list of warnings. What does this print?

```

1 def add_warning(msg, warnings=[]):
2 warnings.append(msg)
3 return warnings
4
5 add_warning("weight sum low")
6 print(len(add_warning("null ids")))

```

### Exercise 3. Spot the bug (level 2)

Goal: record the run time as an aware UTC timestamp. Find the bug and fix it.

```

1 from datetime import datetime, timezone
2
3 def run_timestamp() -> datetime:
4 """Current time in UTC for the run log."""
5 return datetime.utcnow()

```

Write down the line number of the bug. Then write the fix.

### Exercise 4. Spot the bug (level 2)

Goal: `error_message` is optional. The signature must say so for mypy. Find the bug and fix the signature.

```

1 def log_run(
2 pipeline_name: str,
3 status: str,
4 error_message: str = None,
5) -> None:
6 """Append one record to the run log."""

```

Write down the line number of the bug. Then write the fix.

### Exercise 5. Spot the bug (level 3)

Goal: count respondents with a recorded weight. A weight of 0.0 is recorded and valid for this count. Find the bug and fix it.

```

1 def count_weighted(weights: list[float | None]) -> int:
2 """Number of respondents with a recorded weight."""
3 recorded = [w for w in weights if w]
4 return len(recorded)

```

Write down the line number of the bug. Then write the fix.

### Exercise 6. Multiple choice (level 3)

Goal: keep every respondent except those in region N, including respondents with no region. Does `df[df["region"] != "N"]` meet the goal in pandas?

- a) Yes: in pandas, `!=` keeps missing regions
- b) No: missing regions are dropped
- c) No: it raises on missing values

- d) Only if the column is a string dtype

### Why?

- r1) pandas treats missing != 'N' as True, so those rows stay. Spark would drop them.
- r2) pandas fills missing values with empty strings before comparing.

### Exercise 7. Spot the bug (level 4)

Goal: attach region names to every respondent. Respondents whose code is missing from the lookup must stay, with no name. Find the bug and fix it.

```
1 import pandas as pd
2
3 def add_region(resp: pd.DataFrame, regions: pd.DataFrame) -> pd.DataFrame:
4 """Attach region names. One row per respondent."""
5 lookup = regions.dropna(subset=["code"])
6 return resp.merge(lookup, on="code", how="inner",
7 validate="many_to_one")
```

Write down the line number of the bug. Then write the fix.

### Exercise 8. Spot the bug (level 5)

Goal: return the most recent wave from ids like `W9`, `W10`, `W11`. Find the bug and fix it.

```
1 def latest_wave(wave_ids: list[str]) -> str:
2 """Return the most recent wave id."""
3 if not wave_ids:
4 raise ValueError("no waves")
5 return max(wave_ids)
```

Write down the line number of the bug. Then write the fix.

### Sources

- Snippet behaviours (NameError without import, mutable default sharing, utcnow naive and deprecated, pandas != keeps missing, max of strings) match local runs. (verified: Local runs on CPython 3.13.12 and pandas 3.0.2, September 2026 (see D1, D3 and D5 claims).)

## D7-L2 Review practice: PySpark and Spark SQL

### Concept

Nine snippets of PySpark and Spark SQL, each with one seeded bug. Most come straight from your reference pipeline's known defects. The answer key stays hidden until you submit.

Use the same routine as lesson D7-L1: read the goal, walk the lines, ask what each line assumes. Spark adds its own questions:

- **Imports.** Is `sum`, `min`, `max`, `round` or `filter` imported by name from `pyspark.sql.functions`? Is `functions` as `F` used instead?
- **Filters.** Does a comparison with a nullable column drop nulls? Is a boolean compared with `== True`?
- **Grain.** Does every `groupBy` include each grain column the docstring names? Does every join keep the grain?
- **Aggregates.** Can a `sum(when(...))` return null? Is a division guarded?
- **Passes over the data.** Are there several `count()` calls on an uncached `DataFrame`, where one `agg` would do?
- **Writes.** Does a MERGE source hold one row per key? Is a rename silently skipped when the source column is missing?
- **SQL.** Does `NOT IN` meet a nullable subquery?

For each snippet, tap the line, then write the fix. The checker tests the fix statically: it looks for the pattern the standard requires. Your score counts bugs found and fixes proposed.

When you finish, compare your list of misses with the checklist above. The item you missed most is the one to add to your personal review routine.

## Why it matters

Spark code fails quietly more often than Python code. A shadowed import fails loudly. A missing grain column, a null sum or a duplicate MERGE key can pass every run and still publish wrong numbers. Practice on the defects from your own pipeline trains your eye for the ones that will appear in its next version.

### Worked example: weak code

```
❶ from pyspark.sql.functions import col, sum
2
❷ metrics = silver.filter(col("in_scope") ==
 True).groupBy("wave_id").agg(sum("final_weight"))
```

1. **Line 1.** Shadows Python's `sum` for the whole module.
2. **Line 3.** `== True` trips ruff E712. Filter on the column directly.

### Worked example: strong code

```
❶ from pyspark.sql import functions as F
2
```

```
2 metrics = silver.filter(F.col("in_scope")).groupBy("wave_id").agg(F.sum("final_weight"))
```

1. **Line 1.** Python built-ins keep their meaning.
2. **Line 3.** A boolean column is already a condition.

For Spark code, check imports, null behaviour in filters and sums, grain in groupBy and joins, passes over data and MERGE keys.

## Recap card

- **Spark review checklist** Imports, null behaviour in filters and sums, grain in groupBy, passes over data, MERGE keys, silent renames, NOT IN.
- **MERGE with repeated source keys** Fails when several source rows match one target row. Deduplicate the source on the merge keys first.
- **withColumnRenamed on a missing column** A silent no-op. Check `df.columns` and raise.

## Practice

### Exercise 1. Spot the bug (level 1)

Goal: a gold aggregation module that also uses Python's built-in `sum` in a helper. Find the bug and fix the import.

```
1 from pyspark.sql.functions import col, count, sum, when
2
3 def total_checks(results: list[dict[str, str]]) -> int:
4 """Number of passed checks."""
5 return sum(1 for r in results if r["status"] == "PASS")
```

Write down the line number of the bug. Then write the fix.

### Exercise 2. Spot the bug (level 1)

Goal: keep in-scope respondents. Find the line that breaks the standard and fix it.

```
1 from pyspark.sql import functions as F
2
3 in_scope = silver.filter(F.col("in_scope") == True)
4 in_scope.write.mode("overwrite").saveAsTable("research_dev.nat_in_scope")
```

Write down the line number of the bug. Then write the fix.

### Exercise 3. Spot the bug (level 2)

Goal: drop test panel respondents and keep everyone else, including respondents with no panel recorded. Find the bug and fix it.

```
1 from pyspark.sql import functions as F
```

```

2
3 live = respondents.filter(F.col("panel_type") != "test")

```

Write down the line number of the bug. Then write the fix.

#### Exercise 4. Multiple choice (level 2)

Goal: gold shows 0 unweighted reach for a subgroup where nobody was reached. The reference code uses `sum(when(col("reach_flag") == 1, 1)).alias("unweighted_reach")`. What does gold show?

- a) NULL
- b) 0
- c) The subgroup's row count
- d) An error

#### Why?

- r1) when without otherwise is null for non-matching rows, and a sum of nulls is null.
- r2) Spark drops zero sums.

#### Exercise 5. Spot the bug (level 3)

Goal (from the docstring): reach metrics at tracker x wave x subgroup grain. Find the bug and fix it.

```

1 from pyspark.sql import functions as F
2
3 metrics = (
4 silver.filter(F.col("in_scope"))
5 .groupBy("wave_id", "age_group", "gender", "region")
6 .agg(F.sum("final_weight").alias("weighted_base"))
7)

```

Write down the line number of the bug. Then write the fix.

#### Exercise 6. Spot the bug (level 3)

Goal: table health counts for the run log, computed efficiently. Find the line that shows the problem and rewrite the counts as one aggregation.

```

1 from pyspark.sql import functions as F
2
3 df = spark.table(table_name)
4 health = {
5 "row_count": df.count(),
6 "null_ids": df.filter(F.col("respondent_id").isNull()).count(),
7 "null_weights": df.filter(F.col("final_weight").isNull()).count(),
8 }

```

Write down the line number of the bug. Then write the fix.

## Exercise 7. Spot the bug (level 4)

Goal: MERGE a redelivery into silver on respondent and wave. Redeliveries sometimes repeat a respondent. Find the bug and fix it.

```
1 from delta.tables import DeltaTable
2
3 target = DeltaTable.forName(spark, "research_silver.nat_respondents")
4 (target.alias("t")
5 .merge(delivery.alias("s"), "t.respondent_id = s.respondent_id AND
6 t.wave_id = s.wave_id")
7 .whenMatchedUpdateAll()
8 .whenNotMatchedInsertAll()
9 .execute())
```

Write down the line number of the bug. Then write the fix.

## Exercise 8. Predict the output (level 4)

Goal: list respondents outside the excluded regions. `excluded` holds 'N' and one NULL. How many rows does this return from 1,000 respondents, 300 of them in N?

```
1 SELECT COUNT(*) FROM respondents
2 WHERE region_code NOT IN (SELECT code FROM excluded);
```

## Exercise 9. Spot the bug (level 5)

Goal: rename delivery columns to canonical names. If a source column is missing, stop with an error naming it. Find the bug and fix it.

```
1 def rename_columns(bronze_df, column_map: dict[str, str]):
2 """Rename source columns to canonical names."""
3 for src, tgt in column_map.items():
4 bronze_df = bronze_df.withColumnRenamed(src, tgt)
5 return bronze_df
```

Write down the line number of the bug. Then write the fix.

## Sources

- Snippet behaviours (shadowed sum NOT\_COLUMN\_OR\_STR, != drops nulls, sum(when) null, NOT IN with NULL returns no rows, withColumnRenamed no-op, MERGE multiple-source-rows failure) match local runs. (verified: Local runs on PySpark 4.0.4 and delta-spark 4.0.1 (see D1, D3, D4 claims and practice/tests/test\_a3\_l1\_merge.py), September 2026.)



## D7-L3 Review practice: Databricks SQL and DAX

---

### Concept

Eight snippets of Databricks SQL and DAX, each with one seeded bug. SQL and DAX look simple, which makes their bugs easy to read past. The answer key stays hidden until you submit.

#### What to look for in Databricks SQL:

- **Syntax from another platform.** `CREATE ROW ACCESS POLICY` and `CREATE MASKING POLICY` are Snowflake. Databricks uses a SQL UDF with `SET ROW FILTER` or `SET MASK`.
- **Grants to individuals.** Your standard grants to groups, never to one user's email.
- **Group checks.** `is_member()` checks workspace groups. Unity Catalog guidance favours `is_account_group_member()`.
- **Names.** Two-part names resolve against the current catalog. Use three-level names: `catalog.schema.table`.
- **Ratios.** An `AVG` of percentages weights every subgroup equally. Recompute ratios from their sums.
- **MERGE sources.** One row per key, or the statement fails.

#### What to look for in DAX:

- **Division.** `/` returns infinity or an error on a zero denominator. `DIVIDE` returns BLANK.
- **Filter context.** `ALL(fact_table)` removes every filter on the table, wave and region slicers included. Name the one dimension whose filter you mean to remove.

Review each snippet with the same routine: goal, then line by line. For SQL, also check each statement against the Databricks SQL reference in your head: does this statement exist here? The checker tests your fix against the pattern the standard requires. Your score counts bugs found and fixes proposed.

This lesson closes the practice module. The capstone review in Phase 6 uses the same skills on a whole pull request.

### Why it matters

Security SQL and report measures carry outsized risk. A grant to one person survives their move to another team. A row filter written in the wrong dialect

leaves a table open. A DAX measure with the wrong filter context shows a plausible number to every viewer of a dashboard. Reviewing these few lines carefully protects many readers.

### Worked example: weak code

```
1 SELECT region, AVG(weighted_reach_pct) AS reach_pct
2 FROM research_gold.nat_metrics
3 GROUP BY region;
```

1. **Line 1.** An average of subgroup percentages. A subgroup of 12 counts as much as one of 1,200.

### Worked example: strong code

```
1 SELECT region,
2 SUM(weighted_reach) / NULLIF(SUM(weighted_base), 0) AS reach_pct
3 FROM research_gold.nat_metrics
4 GROUP BY region;
```

1. **Line 2.** Total over total, with the zero base guarded.

For SQL, check the dialect, grants, group functions, names and ratios. For DAX, check DIVIDE and the exact filter each ALL removes.

### Recap card

- **SQL review checklist** Dialect, grants to groups, is\_account\_group\_member, three-level names, ratios from sums, one source row per MERGE key.
- **DAX review checklist** DIVIDE for every ratio. ALL on the one dimension you mean, never the whole fact table.
- **Deduplicating a MERGE source in SQL** ROW\_NUMBER() OVER (PARTITION BY keys ORDER BY ingestion\_timestamp DESC), keep rn = 1.

### Practice

#### Exercise 1. Spot the bug (level 1)

Goal: restrict research\_silver.nat\_respondents rows by region with a Databricks row filter. Find the bug and fix it.

```
1 CREATE FUNCTION research_governance.region_filter(region STRING)
2 RETURN is_account_group_member(concat('analysts_', region));
3 CREATE ROW ACCESS POLICY region_policy ON research_silver.nat_respondents
 (region);
```

Write down the line number of the bug. Then write the fix.

## Exercise 2. Spot the bug (level 1)

Goal: let the analysts read the gold metrics table, following the standard. Find the bug and fix it.

```
1 GRANT USE CATALOG ON CATALOG research TO `analysts`;
2 GRANT USE SCHEMA ON SCHEMA research.research_gold TO `analysts`;
3 GRANT SELECT ON TABLE research.research_gold.nat_metrics TO
 `ana.lopez@example.com`;
```

Write down the line number of the bug. Then write the fix.

## Exercise 3. Spot the bug (level 2)

Goal: a column mask that shows postcodes only to members of the account-level group `analysts`, as Unity Catalog guidance recommends. Find the bug and fix it.

```
1 CREATE FUNCTION research_governance.mask_postcode(p STRING)
2 RETURN CASE WHEN is_member('analysts') THEN p ELSE '***' END;
3 ALTER TABLE research_silver.nat_respondents ALTER COLUMN postcode SET
 MASK research_governance.mask_postcode;
```

Write down the line number of the bug. Then write the fix.

## Exercise 4. Multiple choice (level 2)

Goal: a scheduled query must read production gold. It runs `SELECT * FROM research_gold.nat_metrics`. The job's session has `research_dev` as its current catalog. What does it read?

- a) `research_dev.research_gold.nat_metrics`, if it exists
- b) `research.research_gold.nat_metrics`, because gold is production
- c) An error: two-part names are not allowed
- d) Both tables, unioned

### Why?

- r1) A two-part name resolves against the current catalog. Only a three-level name is unambiguous.
- r2) Databricks always prefers dev catalogs.

## Exercise 5. Spot the bug (level 3)

Goal: weighted reach percentage per region from the gold table, which has one row per subgroup with `weighted_reach` and `weighted_base`. Find the bug and fix it, guarding a zero base.

```
1 SELECT region,
2 AVG(weighted_reach / weighted_base) AS reach_pct
3 FROM research.research_gold.nat_metrics
```

```
4 GROUP BY region;
```

Write down the line number of the bug. Then write the fix.

### Exercise 6. Fill the gap (level 3)

Goal: a reach percentage measure that shows an empty cell, not infinity, where the base is zero. The assistant wrote `RETURN reach_w / base_w`.

Complete the fix.

```
1 Reach % =
2 VAR reach_w = SUM (fact_reach[weighted_reach])
3 VAR base_w = SUM (fact_reach[weighted_base])
4 RETURN ____ (1) (reach_w, ____ (2))
```

### Exercise 7. Spot the bug (level 4)

Goal: each brand's share of weighted reach across all brands, respecting the wave and region slicers. Find the bug and fix it.

```
1 Brand share =
2 VAR this_brand = SUM (fact_reach[weighted_reach])
3 VAR all_brands = CALCULATE (SUM (fact_reach[weighted_reach]), ALL
 (fact_reach))
4 RETURN DIVIDE (this_brand, all_brands)
```

Write down the line number of the bug. Then write the fix.

### Exercise 8. Spot the bug (level 5)

Goal: MERGE a redelivery into silver. The delivery can repeat a respondent within a wave. Keep the row with the latest `ingestion_timestamp`. Find the bug and fix it.

```
1 MERGE INTO research.research_silver.nat_respondents AS t
2 USING research.research_bronze.nat_raw_w14 AS s
3 ON t.respondent_id = s.respondent_id AND t.wave_id = s.wave_id
4 WHEN MATCHED THEN UPDATE SET *
5 WHEN NOT MATCHED THEN INSERT *;
```

Write down the line number of the bug. Then write the fix.

## Sources

- Databricks row filters and masks, `is_account_group_member`, three-level namespace and grants follow the docs verified in Track A. (verified: <https://docs.databricks.com/aws/en/data-governance/unity-catalog/filters-and-masks/manually-apply> and related pages (checked 2026-09-25))
- `SELECT * EXCEPT (col)` is supported in Databricks SQL. (verified: <https://docs.databricks.com/aws/en/sql/language-manual/sql-ref-syntax-qry-star> (checked 2026-09-25))

## D8 Prompting for correct code

Roadmap: Every month. State versions, constraints, expected outputs and tests.

### D8-L1 State versions and constraints in every prompt

---

#### Concept

Most wrong code from an assistant answers a question you did not ask. It targets a version you do not run, a platform you do not use, or a standard you did not mention. Track D showed the results: `CREATE ROW ACCESS POLICY`, `datetime.utcnow()`, `applymap`. The cheapest fix happens before the code exists: a better prompt.

**Name the versions.** The assistant cannot see your environment. Say it:

```
1 Environment: Databricks Runtime 17.3 LTS (Spark 4.0, Python 3.12),
2 PySpark DataFrame API, Delta tables in Unity Catalog.
3 Local tests: pandas 3.0, pytest 9.
```

**Name the platform and dialect.** "SQL" means nothing. Write "Databricks SQL" or "Fabric Warehouse T-SQL". The same goes for "a notebook": Databricks or Fabric?

**State the constraints your standard sets.** Paste the few that apply, not the whole file:

- Import `functions` as `F`. Never import `sum`, `min` or `max` by name.
- Type hints and a Google-style docstring on every function.
- Bronze is append-only. Silver uses `MERGE` on `(respondent_id, wave_id)`.
- Three-level table names. No secrets in code.
- Test for `None` with `is None`.

**State what not to do.** "Do not use RDDs, Python UDFs or `collect()` on the full table." Negative constraints stop the common wrong turns.

**Ask for a list of assumptions** at the end of the answer (lesson D6-L3).

Keep a prompt header with these lines in your repository, next to `CLAUDE.md`, and paste it at the start of each coding conversation. A prompt that states versions and constraints does not guarantee correct code. It removes the mistakes that come from guessing, so your review can focus on the rest.

## Why it matters

An engineer asked for "a PySpark function to cap survey weights". The answer used `pandas_udf` with a type-hint style from Spark 3.0 docs, imported `max` from `pyspark.sql.functions`, and wrote to a two-part table name. The same request with a four-line header naming Databricks Runtime 17.3 LTS, the import rule and three-level names produced code that passed review on the first read.

## Worked example: weak code

❶ Write a PySpark function that caps survey weights at 5 and saves the result.

1. **Line 1.** No version, no platform, no standard, no table rules. The assistant guesses all of them.

## Worked example: strong code

❶ Environment: Databricks Runtime 17.3 LTS (Spark 4.0, Python 3.12), Unity Catalog.

2 Standards: from `pyspark.sql` import `functions` as `F`; type hints and Google docstrings;

❷ no Python UDFs, RDDs or `collect()` on full tables; three-level table names.

4 Task: write `cap_weights(df, cap: float) -> DataFrame` that caps `final_weight` at `cap`.

5 Do not write to any table. Return the `DataFrame`.

❸ Finish by listing the assumptions your code makes about the input.

1. **Line 1.** Versions and platform, stated once.
2. **Line 3.** Negative constraints close the common wrong turns.
3. **Line 6.** Makes hidden assumptions reviewable.

Every coding prompt names versions, platform and dialect, the constraints from your standard, and what not to do.

## Recap card

- **Four things every coding prompt states** Versions and platform, the dialect, the relevant standards, and what not to do.
- **Why name the SQL dialect** Syntax differs by platform. 'SQL' alone invites Snowflake or T-SQL statements.
- **Where to keep the prompt header** In the repository next to `CLAUDE.md`, updated when versions change.

## Practice

### Exercise 1. Multiple choice (level 1)

You ask for "SQL to add a row filter by region". Which line most improves the prompt?

- a) Target: Databricks SQL, Unity Catalog, Databricks Runtime 17.3 LTS
- b) Please write clean code
- c) Make it efficient
- d) Explain your answer in detail

#### Why?

- r1) Row filter syntax differs by platform. Naming the dialect rules out Snowflake's CREATE ROW ACCESS POLICY.
- r2) Longer prompts always give better answers.

### Exercise 2. Put in order (level 2)

Order the parts of a coding prompt header.

- Environment: runtime, language and library versions
- What not to do: no UDFs, no RDDs, no secrets
- Standards that apply: imports, typing, table rules
- The task, with inputs, outputs and side effects
- Ask for a list of assumptions

### Exercise 3. Fill the gap (level 2)

Complete the environment line for this course's Databricks lessons.

```
1 Environment: Databricks Runtime ____ (1) LTS (Spark ____ (2), Python 3.12), Unity Catalog.
```

### Exercise 4. Write the function (level 3)

Write `unstated_versions(prompt: str, libraries: list[str]) -> list[str]`. Return the libraries, in the order given, that the prompt does not mention followed by a version number, such as `pandas 3.0`, `PySpark 4.0` or `Python 3.12`. Match library names case-insensitively. A version is digits and dots after optional whitespace and an optional `v`. A library mentioned without a version counts as unstated.

Starter code:

```
1 def unstated_versions(prompt, libraries):
2 return [lib for lib in libraries if lib not in prompt]
```

## Exercise 5. Architecture design (level 4)

Build the prompt header your team pastes into every coding conversation for the survey pipeline.

Options: Databricks Runtime 17.3 LTS, Spark 4.0, Python 3.12, pandas 3.0 for local tests, Use the latest versions of everything.

1. Versions line
2. Standards
3. Where the header lives

## Sources

- Databricks Runtime 17.3 LTS ships Spark 4.0.0 and Python 3.12.3; Runtime 18 is not yet LTS. (verified: <https://docs.databricks.com/aws/en/release-notes/runtime/17.3lts> and /18 (checked 2026-09-25))

## D8-L2 State expected outputs and tests in the prompt

---

### Concept

"Write a function that computes weighted reach" leaves the assistant to decide what correct means. It will decide something plausible, and you will find out later whether it matches your meaning. Put your meaning in the prompt instead: examples with expected outputs, the edge cases, and the tests the code must pass.

**Give a small input and the exact output.** A table of five rows and the result you worked out by hand removes most ambiguity at once:

```
1 Input (respondent_id, reach_flag, final_weight):
2 R1, 1, 2.0 R2, 0, 1.0 R3, 1, 1.0 R4, null, 1.0
3 Expected: weighted_reach = 3.0, weighted_base = 4.0, pct = 0.75
4 R4 has no reach answer: exclude it from the base.
```

**List the edge cases and their answers.** Empty input returns None. A zero base returns None. A share of 0.0 is valid. A duplicate key raises ValueError. Each line is a decision the assistant would otherwise make for you.

**Paste the tests, or ask for them in a fixed form.** Two options:

- Paste your own failing tests (lesson D6-L2) and ask for code that makes them pass without changing them.
- Ask for tests first, "using only hand-worked expected values, not the function's formula", review them, then ask for the code.



**State the output format.** Return type, column names and order, rounding, and whether the function writes anywhere. "Return a DataFrame with columns wave\_id, weighted\_reach\_pct. Do not write to a table."

A prompt with examples and tests turns review from "does this look right?" into "do the tests pass, and are they the tests I wrote?" The second question takes seconds.

## Why it matters

Asked for "weighted reach by wave", an assistant included respondents with a null reach answer in the base. Another conversation, given the same task with a five-row example that excluded the null, produced code that matched on the first try. The example took two minutes to write and settled a question that had caused a four-point difference in the dashboard.

## Worked example: weak code

1 Write a function to compute the weighted reach percentage by wave.

1. **Line 1.** Nulls? Empty waves? Output shape? The assistant decides all of it.

## Worked example: strong code

```
1 Task: weighted_reach_pct(rows) -> float | None.
2 rows: list of (reach_flag: int | None, final_weight: float).
3
4 Example: [(1, 2.0), (0, 1.0), (1, 1.0), (None, 1.0)] -> 0.75
5 (the None row is excluded from the base).
6 Edge cases: [] -> None; all reach_flag None -> None; [(0, 1.0)] -> 0.0.
7
8 1 Make these tests pass without editing them:
9
10 assert weighted_reach_pct([(1, 2.0), (0, 1.0), (1, 1.0), (None, 1.0)])
11 == 0.75
12
13 assert weighted_reach_pct([]) is None
14
15 assert weighted_reach_pct([(0, 1.0)]) == 0.0
```

1. **Line 3.** A worked example settles how nulls are treated.

2. **Line 6.** The acceptance criteria, stated before any code exists.

Give a worked example, the edge cases with answers, the output format and the tests. Then review against them.

## Recap card

- **Four things that specify correct code in a prompt** A worked example, edge cases with answers, the output format, and fixed tests.
- **Instruction that protects your tests** 'Make these tests pass without editing them. If a test seems wrong, say so.'

- **Why work the example out yourself** If you cannot compute the answer by hand, the rule is not clear enough to ask for code.

## Practice

### Exercise 1. Multiple choice (level 1)

Which addition to a prompt best settles how null reach answers are treated?

- a) A five-row example with a null and the expected result worked by hand
- b) 'Handle nulls properly'
- c) 'Be careful with edge cases'
- d) 'Use best practices'

### Why?

- r1) An example states the rule precisely. Words like 'properly' leave the rule to the assistant.
- r2) Assistants ignore instructions written in words.

### Exercise 2. Predict the output (level 2)

Work out the expected value for your prompt's example. Rows are `(reach_flag, final_weight)`, and rows with a `None` flag are excluded from the base. What does this print?

```
1 rows = [(1, 2.0), (0, 1.0), (1, 1.0), (None, 1.0)]
2 valid = [(f, w) for f, w in rows if f is not None]
3 print(sum(w for f, w in valid if f == 1) / sum(w for _, w in valid))
```

### Exercise 3. Fill the gap (level 2)

Complete the instruction that stops an assistant from weakening your tests.

```
1 Make these tests pass ____ (1) editing them. If a test seems wrong,
 ____ (2) instead of changing it.
```

### Exercise 4. Write the function (level 3)

Write `examples_to_asserts(func_name: str, examples:`

`list[tuple[tuple[object, ...], object]]) -> list[str]` that turns worked examples into assert lines for a prompt. Each example is `(args, expected)`. Produce `assert func(arg1, arg2) == expected` using `repr` for every argument and value, except that an expected `None` produces `assert func(...) is None`.

Starter code:

```
1 def examples_to_asserts(func_name, examples):
2 return [f"assert {func_name}({args}) == {expected}" for args,
 expected in examples]
```

## Exercise 5. Architecture design (level 4)

Plan the prompt for an assistant to write a silver deduplication function.

Options: Six input rows with two duplicates and the expected four output rows, No example: the assistant knows what deduplication means.

1. Example in the prompt
2. Rule for duplicates
3. Tests

## Sources

- Generated assert lines using repr and 'is None' execute as valid Python. (verified: Checked by the exercise tests in this repo, September 2026.)

## D8-L3 Iterate on a failing answer

---

### Concept

The first answer often fails a test. What you send next decides whether the second answer is better or just different.

**Send the evidence, not a feeling.** "It doesn't work" invites a rewrite from scratch. Send:

- the **exact failing test** and its assertion,
- the **full error**, or the actual versus expected values,
- the **smallest input** that shows the problem,
- what you already checked.

```
1 test_zero_share fails:
2 assert compare_waves({"a": 0.0}, {"a": 0.12}) == ["a"]
3 got [] instead of ["a"]
4 Line 6 uses `if current_pct and prior_pct:`. I think 0.0 is treated as
 missing.
5 Fix only this. Do not change the tests or other behaviour.
```

**One problem per round.** Ask for a fix to one failure at a time, and say what must not change. Broad requests ("fix all the issues and improve it") produce broad rewrites, which reintroduce bugs you already fixed.

**Watch the tests, not only the code.** A common way for an assistant to "fix" a failing test is to change the test: loosen an assertion, delete a case, or replace a hand-worked value with a computed one. Diff the test file after every answer. Any removed or changed assertion needs your approval.

**Know when to stop.** If two focused rounds have not fixed a problem, the assistant is guessing. Stop, read the code yourself with the debugger (lesson D2-L1), and write the fix. Then you can explain it in the next prompt, or leave the assistant out of it.

**Start fresh when the thread goes wrong.** A long conversation full of failed attempts carries those attempts forward. A new conversation with the header from lesson D8-L1, the current code, and the one failing test often does better than round seven.

## Why it matters

After four rounds of "still failing, please fix", an assistant's version of a stability check passed every test. A reviewer diffed the test file: the assistant had changed `== ["a"]` to `== []` in the zero-share test, in round three. The real bug, if `current_pct` and `prior_pct`, was still there. One focused prompt with the failing assertion fixed it in a single round.

## Worked example: weak code

❶ Still not working. Please fix everything and make it better.

1. **Line 1.** No test, no error, no scope. The assistant rewrites everything, including what already worked.

## Worked example: strong code

❶ test\_missing\_metric fails with KeyError: 'new' at line 7, `prior\_pct = prior[metric]`.

2 Input: compare\_waves({"new": 0.4}, {}). Expected: [].

3 Fix only the lookup so a metric missing from prior is skipped.

❷ Do not edit the tests. Keep every other behaviour the same.

1. **Line 1.** The exact test, error and line.
2. **Line 4.** Scope limits stop regressions and test tampering.

Send the failing test, the error and the smallest input. Fix one thing per round, diff the tests, and stop after two failed rounds.

## Recap card

- **What to send when AI code fails a test** The failing assertion, actual vs expected, the smallest input, what you checked, and 'fix only this'.
- **The test-tampering check** Diff the test file after every answer. Removed or changed assertions need your approval.
- **When to stop iterating** After two focused rounds without progress: debug it yourself, or restart in a fresh conversation.

## Practice

### Exercise 1. Multiple choice (level 1)

An assistant's function fails one test. What do you send next?

- a) The failing test, the actual and expected values, and 'fix only this'
- b) 'It doesn't work, try again'
- c) 'Rewrite it more carefully'
- d) The whole project so it has more context

### Why?

- r1) Precise evidence and a narrow scope lead to a targeted fix instead of a rewrite.
- r2) Assistants only respond to test output.

### Exercise 2. Spot the bug (level 2)

This follow-up prompt has one line that invites a rewrite instead of a fix. Tap it, then rewrite the prompt so it limits the scope.

```
1 test_zero_share fails: expected ["a"], got [].
2 Input: compare_waves({"a": 0.0}, {"a": 0.12}).
3 While you are at it, refactor the whole module and improve the naming.
```

Write down the line number of the bug. Then write the fix.

### Exercise 3. Put in order (level 2)

Order your steps after an assistant's code fails a test.

- If two rounds fail, debug and fix it yourself
- Diff both the code and the test file in the answer
- Collect the failing assertion, the error and the smallest input
- Run the full test suite
- Send them with 'fix only this, do not edit the tests'

### Exercise 4. Write the function (level 3)

Write `removed_asserts(before: str, after: str) -> list[str]`. Given two versions of a test file, return the `assert` statements present in `before` but missing from `after`, compared after stripping whitespace at both ends, in the order they appear in `before`. An `assert` that was changed counts as removed. Only lines that start with `assert` after stripping count.

Starter code:

```
1 def removed_asserts(before, after):
```

### Exercise 5. Architecture design (level 4)

Round three of a conversation still fails the same test, and the thread holds four earlier attempts. Decide the next steps.

Options: Debug it yourself with pdb on the failing test, Send 'still failing' again.

1. Next move
2. If you use the assistant again
3. Before accepting any answer

### Sources

- The stop-after-two-rounds rule is this course's working guideline, not a published standard. (verified: Course guideline stated as such in the lesson text.)

## Track D best-practice checklist

- ☐ Read a traceback from the bottom up. Find the first frame in your own code.
- ☐ Reproduce the error on the smallest data that shows it.
- ☐ Check row counts before and after every join.
- ☐ Test numbers with `is None`, never with truthiness, when zero is a valid value.
- ☐ Never use a mutable object as a default argument.
- ☐ Check every function and parameter an AI suggests against the official docs for your version.
- ☐ Run AI code on a small input with a known answer before you trust it.
- ☐ Write the failing test first. A test that only repeats the code proves nothing.
- ☐ Ask the model to list its assumptions. Check each one.
- ☐ State versions, constraints, expected outputs and tests in every prompt.

# Appendix: solutions and answer keys

Attempt each exercise before you read its solution. Track D answer keys are here too.

## A1-L1 Single-purpose functions and clear names

---

### A1-L1 exercise 1

**Hints.** (1) Look at the naming table in the concept. (2) Classes and constants have their own styles. (3) Functions: lower case, words joined by underscores, verb first.

**Answer.** b) `check_table_health` Because: r1) Functions use snake\_case and start with a verb.

Functions use `verb_noun` in snake\_case. PascalCase is for classes and UPPER\_SNAKE\_CASE for constants, so the style alone tells a reader what kind of object they are looking at.

#### Common wrong answers.

- a: PascalCase marks a class, such as `CheckResult`. A function uses snake\_case.
- c: camelCase is the Java and JavaScript convention. PEP 8 and your standard use snake\_case for functions.
- d: UPPER\_SNAKE\_CASE marks a constant, such as `MAX_WEIGHT_CV`.
- b|r2: Right name, wrong reason. Name length has no effect on speed. The convention signals what kind of object it is.

### A1-L1 exercise 2

**Hints.** (1) Tables follow `[tracker]_[content]`. (2) Bronze holds raw deliveries, silver holds respondents, gold holds metrics. (3) The job pattern is `[tracker]_[layer]_pipeline`.

#### Answer.

```
1 BRONZE_TABLE = "research_bronze.nat_raw"
2 SILVER_TABLE = "research_silver.nat_respondents"
3 GOLD_TABLE = "research_gold.nat_metrics"
4 JOB_NAME = "nat_silver_pipeline"
```

`nat_raw`, `nat_respondents` and `nat_metrics` name what each table holds. The schema carries the layer. Jobs use `nat_silver_pipeline`.



### Common wrong answers.

- `g1 = nat_bronze`: The layer is already in the schema name, `research_bronze`. The table name says what it holds: `nat_raw`.
- `g4 = silver_nat_pipeline`: Jobs start with the tracker: `[tracker]_[layer]_pipeline`. Notebooks start with the layer.
- `g2 = nat_silver`: Silver tables are named after their grain: one row per respondent, so `nat_respondents`.

### A1-L1 exercise 3

**Hints.** (1) What does the word `get` promise the caller? (2) Which line changes a table? (3) Delete line 6. Deleting data belongs in its own, clearly named function.

**Answer.** The bug is on line 6.

Fixed version:

```
1 def get_unprocessed_waves(spark: SparkSession, bronze: str, silver: str,
2 tracker_id: str) -> list[str]:
3 """Return wave ids in bronze but not yet in silver for one tracker."""
4 bronze_waves = spark.table(bronze).filter(F.col("tracker_id") ==
5 tracker_id).select("wave_id").distinct()
6 silver_waves = spark.table(silver).filter(F.col("tracker_id") ==
7 tracker_id).select("wave_id").distinct()
8 waves = [row.wave_id for row in
9 bronze_waves.subtract(silver_waves).collect()]
10 return waves
```

Line 6 deletes rows from silver inside a function named `get_`. A caller who only wants the list of waves loses data. Remove the line. If you need that clean-up, give it its own function, such as `delete_waves()`, and call it from the orchestrator.

### Common wrong answers.

- Line 5: `collect()` on a short list of wave ids is fine. Look for a line that changes data.
- Line 4: Filtering silver by tracker is correct. It fixes a known defect in the reference code.

### A1-L1 exercise 4

**Hints.** (1) Each function takes `df: DataFrame` and returns a `DataFrame`. (2) Move line 3 into `filter_in_scope` and line 4 into `aggregate_base`. (3) Drop `== True`, and leave `spark.table` and `.write` out of both functions.

**Answer.**

```

1 from pyspark.sql import DataFrame
2 from pyspark.sql import functions as F
3
4
5 def filter_in_scope(df: DataFrame) -> DataFrame:
6 """Keep respondents flagged as in scope."""
7 return df.filter(F.col("in_scope"))
8
9
10 def aggregate_base(df: DataFrame) -> DataFrame:
11 """Weighted base at wave x region grain."""
12 return df.groupBy("wave_id",
13 "region").agg(F.sum("final_weight").alias("weighted_base"))

```

Pure transforms take a DataFrame and return one. Reading and writing stay in the orchestrator, so the transforms run in a unit test with a handful of rows.

## A1-L1 exercise 5

**Hints.** (1) Values that change between environments are config. (2) Anything that touches the catalog is a read or a write. (3) The orchestrator only sequences the other functions.

### Answer.

- Table names are configuration. Load them once and pass them in.
- `spark.table(...)` is a read. It takes the session and the table name as parameters.
- Summing weights is a pure transform.
- The orchestrator calls each step and logs the run. It holds no business logic.

Config holds names and thresholds. Reads touch the catalog. Transforms turn one DataFrame into another. The orchestrator calls them in order and logs the outcome.

### Common wrong answers.

- `table_name = orchestrator`: A table name typed into the orchestrator is still hardcoded. Put it in config, so dev and prod runs point at different tables.
- `weights = orchestrator`: If the orchestrator holds the aggregation, you cannot test it without running the whole pipeline. Make it a transform.
- `spark_table = transform`: A transform receives a DataFrame. Reading a table needs a session and a catalog, which is the read function's job.

## A1-L2 Type hints and docstrings a reviewer can trust

---

### A1-L2 exercise 1

**Hints.** (1) The default is `None`. Does the hint allow `None`? (2) Python 3.10 added the `X | Y` union syntax. (3) `str | None = None`.

**Answer.** b) `def log_run(error_message: str | None = None) -> None:`  
Because: r1) The hint must include `None`, because `None` is a value the parameter takes.

`str | None = None` says the parameter is a string or `None`, and defaults to `None`. Python ignores hints at run time. `mypy` and readers rely on them.

#### Common wrong answers.

- a: `str = None` claims the value is always a string. `mypy` strict mode rejects the implicit `Optional`.
- c: No hint at all. `mypy --strict` rejects unannotated parameters.
- d: This makes the type `None` and the default the class `str`. It is backwards.
- b|r2: Right signature, wrong reason. Python ignores hints at run time. They exist for tools and readers.

### A1-L2 exercise 2

**Hints.** (1) What does `for k in "id"` loop over? (2) A string is a sequence of characters. (3) Two characters, two conditions, joined by `AND`.

#### Answer.

```
1 t.i = s.i AND t.d = s.d
```

A string is iterable, so the loop yields `i` and `d`. The condition names two columns that do not exist. With `merge_keys: list[str]`, `mypy` flags the call before it runs.

#### Common wrong answers.

- `t.id = s.id`: That is what the caller intended. Iterating a string yields its characters, `i` then `d`.
- `TypeError`: Python raises no error. A string is iterable, which is why the bug is silent until Spark runs the `MERGE`.

### A1-L2 exercise 3

**Hints.** (1) The type for a fixed set of values lives in `typing`. (2) Its name means the values are written out literally. (3) `Literal`, and `None` for the return type.

**Answer.**

```
1 from typing import Literal
2
3 def record_check(name: str, status: Literal["PASS", "WARN", "FAIL"]) ->
 None:
4 ...
```

`Literal["PASS", "WARN", "FAIL"]` restricts the value to those three strings. `-> None` states the function returns nothing.

#### Common wrong answers.

- `g1b = Enum`: An Enum is a class you define. Inline string choices use `Literal[...]` from `typing`.
- `g2 = void`: Python has no `void`. A function that returns nothing is annotated `-> None`.
- `g1b = str`: `str` accepts any string, so `"pass"` or `"FAILED"` would slip through. `Literal` lists the allowed values.

### A1-L2 exercise 4

**Hints.** (1) Check the type first with `isinstance(keys, str)`, then check for an empty list. (2) Build each clause with an f-string that uses `target` and `source`. (3) `" AND ".join(f"{target}.{k} = {source}.{k}" for k in keys)`.

**Answer.**

```
1 def build_merge_condition(keys: list[str], target: str = "t", source: str
 = "s") -> str:
2 """Build a MERGE ON condition that matches every key column.
3
4 Args:
5 keys: Columns that identify one target row.
6 target: Alias of the target table.
7 source: Alias of the source DataFrame.
8
9 Returns:
10 The condition, for example "t.id = s.id AND t.wave = s.wave".
11
12 Raises:
13 TypeError: If keys is a string instead of a list.
14 ValueError: If keys is empty.
15 """
16 if isinstance(keys, str):
```

```

17 raise TypeError("keys must be a list of column names, not a
 string")
18 if not keys:
19 raise ValueError("keys must name at least one column")
20 return " AND ".join(f"{target}.{k} = {source}.{k}" for k in keys)

```

The hint `list[str]` lets mypy catch a string at the call site. The run-time `TypeError` catches callers that skip mypy, such as a notebook. The `ValueError` stops an empty condition, which would match nothing and insert every row.

## A1-L2 exercise 5

**Hints.** (1) Import `Literal` from `typing` and name the return type. (2) `change` is `float | None`. `tolerance` is `float`. (3) Keep `if change is None`. Truthiness would turn a change of 0.0 into SKIP.

### Answer.

```

1 from typing import Literal
2
3 ChangeStatus = Literal["SKIP", "WARN", "PASS"]
4
5
6 def classify_change(change: float | None, tolerance: float = 0.10) ->
 ChangeStatus:
7 """Classify a wave-over-wave change in a share.
8
9 Args:
10 change: Absolute change in the share, or None when a share is
 missing.
11 tolerance: Largest change that still passes.
12
13 Returns:
14 SKIP when change is None, WARN above tolerance, otherwise PASS.
15 """
16 if change is None:
17 return "SKIP"
18 return "WARN" if change > tolerance else "PASS"

```

The `Literal` return type means mypy rejects a typo such as `\ "PASSED\ "` at every caller. `float | None` records that a change can be missing. The `is None` test keeps a real change of 0.0 as `PASS`.

## A1-L3 Config over hardcoded values, and PEP 8 with ruff

### A1-L3 exercise 1

**Hints.** (1) What changes when the same code runs in dev and in prod? (2) Logic stays the same. Names and limits change. (3) The catalog differs between environments.

**Answer.** b) The catalog name, `research_prod` or `research_dev` Because: r1) It changes between environments while the logic stays the same.

The catalog name varies by environment. The formula, the step order and the write strategy are logic, which belongs in reviewed, tested code.

#### Common wrong answers.

- a: The formula is business logic. It is the same in every environment, so it belongs in a tested function.
- c: The step order is the pipeline's design. It lives in the orchestrator and changes only with a code review.
- d: MERGE versus append is a design decision about idempotence, not a setting.
- b|r2: Right choice, wrong reason. Length is irrelevant. Configuration holds what varies between environments, trackers or contract versions.

### A1-L3 exercise 2

**Hints.** (1) What does `frozen=True` do to attribute assignment? (2) The exception fires before the value changes. (3) Two lines are printed.

#### Answer.

```
1 blocked
2 research_dev
```

`frozen=True` makes assignment raise `FrozenInstanceError`, so `blocked` prints and the catalog stays `research_dev`. No code can switch a run to production halfway through.

#### Common wrong answers.

- `research_prod`: A frozen dataclass rejects assignment after creation, so the value stays `research_dev`.
- `blocked research_prod`: The assignment raised before it changed anything. The value stays `research_dev`.

### A1-L3 exercise 3

**Hints.** (1) One command lints, one formats. (2) The lint command's name means inspecting code. (3) `ruff check .` and `ruff format --check .`

**Answer.**

```
1 ruff check .
2 ruff format --check .
```

`ruff check .` reports lint errors such as F821 and E712. `ruff format --check .` exits non-zero when any file would be reformatted, which fails CI without touching the code.

**Common wrong answers.**

- `g1 = lint`: `ruff` has no `lint` command. Linting runs with `ruff check`.
- `g2 = fmt`: The formatter command is `ruff format`. Add `--check` so CI reports instead of rewriting.
- `g1 = format`; `g2 = check`: Swapped. `ruff check` lints, `ruff format --check` verifies formatting.

### A1-L3 exercise 4

**Hints.** (1) Declare three fields on a frozen dataclass, then add a `@property` for the table name. (2) Index the dict with `raw["key"]`. `.get` with a default hides a missing key. (3) Check `0 <= tolerance <= 1` before you build the object.

**Answer.**

```
1 from dataclasses import dataclass
2
3
4 @dataclass(frozen=True)
5 class RunConfig:
6 """Settings for one pipeline run."""
7
8 catalog: str
9 tracker: str
10 stability_tolerance: float
11
12 @property
13 def silver_table(self) -> str:
14 """Three-level name of the silver table."""
15 return f"{self.catalog}.research_silver.{self.tracker}_respondents"
16
17
18 def load_run_config(raw: dict[str, object]) -> RunConfig:
19 """Build a validated RunConfig from parsed JSON.
20
```

```

21 Args:
22 raw: Parsed configuration.
23
24 Returns:
25 The frozen configuration.
26
27 Raises:
28 KeyError: If a required key is missing.
29 ValueError: If stability_tolerance is outside 0 to 1.
30 """
31 tolerance = float(str(raw["stability_tolerance"]))
32 if not 0 <= tolerance <= 1:
33 raise ValueError(f"stability_tolerance
34 {tolerance} is outside 0 to 1")
35 return RunConfig(
36 catalog=str(raw["catalog"]),
37 tracker=str(raw["tracker"]),
38 stability_tolerance=tolerance,
39)

```

Indexing with `raw["key"]` fails at load with the key's name. Defaults such as `raw.get("\catalog\", \"research_prod\")` are worse than no config: a missing value silently points the run at production.

## A1-L3 exercise 5

**Hints.** (1) Search the function for a catalog name. (2) Every table name must come from `cfg`. (3) Add `scope_table` to `RunConfig` and read `cfg.scope_table`.

**Answer.** The bug is on line 4.

Fixed version:

```

1 def weighted_base_by_wave(spark: SparkSession, cfg: RunConfig) ->
 DataFrame:
2 """Weighted base per wave from the configured silver table."""
3 silver = spark.table(cfg.silver_table)
4 scope = spark.table(cfg.scope_table)
5 in_scope = silver.join(scope, "region").filter(F.col("in_scope"))
6 return
 in_scope.groupBy("wave_id").agg(F.sum("final_weight").alias("weighted_base"))

```

Line 4 hardcodes a production table, so every environment reads production scope rules. Add `scope_table` to the config and read it from `cfg`.

## Common wrong answers.

- Line 3: Line 3 reads the silver table from config, which is correct. Look for a name typed into the code.
- Line 5: The join and filter are fine. The problem is which table the join reads.



## A2-L1 Bronze: keep every delivery exactly as received

---

### A2-L1 exercise 1

**Hints.** (1) What is bronze for? (2) Which layer cleans and validates? (3) Bronze keeps everything, labelled with where it came from.

**Answer.** b) Append them as received, with the audit columns. Because: r1) Bronze records deliveries as received. Silver and the data contract decide what to do with bad rows.

Bronze records every delivery as received. The contract check flags the null weights, and silver decides how to handle them. Nothing is lost.

#### Common wrong answers.

- a: Dropping them in bronze destroys the evidence. You can no longer tell the supplier which rows were wrong.
- c: Imputing a weight in bronze invents data. The weighted totals would look plausible and be wrong.
- d: Accept or reject is the contract check's decision, after the delivery is recorded. Bronze records first.
- b|r2: Right action, wrong reason. Delta columns accept nulls unless you add a NOT NULL constraint.

### A2-L1 exercise 2

**Hints.** (1) What type does inferSchema pick for a column of digits? (2) Integers have no leading zeros. (3) The value prints as an int.

#### Answer.

```
1 123
```

inferSchema reads the column as an integer, so 00123 becomes 123. Read bronze without inferSchema and every column stays a string, zeros included.

#### Common wrong answers.

- 00123: That is what you get without inferSchema. With it, Spark sees digits, picks an integer type, and drops the leading zeros.
- '00123': The value is no longer a string. inferSchema made the column an integer.

## A2-L1 exercise 3

**Hints.** (1) Bronze never removes what is already there. (2) The timestamp function's name describes the current moment. (3) `current_timestamp` and `append`.

**Answer.**

```
1 audited = raw.withColumn("ingestion_timestamp", F.current_timestamp())
2 audited.write.format("delta").mode("append").saveAsTable(table)
```

`F.current_timestamp()` stamps each load. `mode("append")` adds rows and leaves earlier deliveries untouched.

**Common wrong answers.**

- `g2 = overwrite`: Overwrite deletes every earlier delivery. Bronze is append-only.
- `g1 = now`: There is no `F.now`. The PySpark function is `F.current_timestamp()`.
- `g2 = merge`: `merge` is not a write mode, and bronze never updates rows. Use `append`.

## A2-L1 exercise 4

**Hints.** (1) `df.withColumns({...})` adds several columns in one call (PySpark 3.3 and later). (2) Constants go through `F.lit()`. The timestamp comes from `F.current_timestamp()`. (3) Add columns only. No filter, no cast.

**Answer.**

```
1 from pyspark.sql import DataFrame
2 from pyspark.sql import functions as F
3
4
5 def add_audit_columns(df: DataFrame, ingestion_id: str, source_file: str,
6 delivery_version: int) -> DataFrame:
7 """Add the four bronze audit columns to a delivery.
8
9 Args:
10 df: The delivery as read, all columns as strings.
11 ingestion_id: One UUID for this load.
12 source_file: Path of the delivery file.
13 delivery_version: Supplier's version number for the delivery.
14
15 Returns:
16 The delivery with ingestion_id, source_file, ingestion_timestamp
17 and delivery_version.
18 """
19 return df.withColumns({
20 "ingestion_id": F.lit(ingestion_id),
21 "source_file": F.lit(source_file),
```

```

20 "ingestion_timestamp": F.current_timestamp(),
21 "delivery_version": F.lit(delivery_version),
22 })

```

`withColumns` adds all four audit columns in one step. Everything else stays as delivered: no filter, no cast.

## A2-L1 exercise 5

**Hints.** (1) What happens to existing rows when this job writes? (2) One word on line 5 decides it. (3) Use `mode("append")`.

**Answer.** The bug is on line 5.

Fixed version:

```

1 def ingest_to_bronze(spark: SparkSession, path: str, table: str,
 delivery_version: int) -> None:
2 """Append one delivery file to bronze as received."""
3 raw = spark.read.option("header", True).csv(path)
4 audited = add_audit_columns(raw, str(uuid.uuid4()), path,
 delivery_version)
5 audited.write.format("delta").mode("append").saveAsTable(table)

```

`overwrite` replaces the table, so each delivery erases the one before. Bronze must append.

### Common wrong answers.

- Line 3: Reading without `inferSchema` is correct for bronze. Look at how the rows are written.
- Line 4: A new UUID per load is correct: each delivery gets its own `ingestion_id`.

## A2-L1 exercise 6

**Hints.** (1) Bronze never changes what the supplier sent. (2) Silver has one clean row per respondent per wave. (3) Gold holds aggregates.

**Answer.**

- Bronze stores deliveries as received.
- Silver renames, casts and resolves versions, so it holds one row per respondent per wave.
- Gold aggregates for reporting, with the grain stated.

Bronze keeps evidence, silver makes it clean and unique, gold makes it ready to report.

### Common wrong answers.

- `cast = bronze`: Casting in bronze loses the original value when the cast fails or guesses wrong. Cast in silver, with checks.
- `dedupe = bronze`: Bronze keeps every delivery, duplicates included. Silver picks the newest version.
- `reach = silver`: Silver holds respondent-level rows. Aggregates for reporting belong in gold.

## A2-L2 Silver: canonical names, checked types, one row per respondent per wave

---

### A2-L2 exercise 1

**Hints.** (1) Does the frame have a column called `wgt`? (2) What does `withColumnRenamed` do when the source column is missing? (3) It returns the frame unchanged.

#### Answer.

```
1 ['respondent_id', 'weight']
```

`withColumnRenamed` is a no-op when the source column is missing. No exception, no warning. Check the columns yourself before you rename.

### Common wrong answers.

- `['respondent_id', 'final_weight']`: There is no column called `wgt`, so nothing is renamed. `weight` keeps its name.
- `AnalysisException`: No error. `withColumnRenamed` silently ignores a missing column, which is exactly the danger.

### A2-L2 exercise 2

**Hints.** (1) Which Spark version does Runtime 17.3 LTS use? (2) What changed about casts in Spark 4.0? (3) ANSI mode is on by default.

**Answer.** b) The job fails with `CAST_INVALID_INPUT`. Because: r1) Spark 4.0 enables ANSI mode by default, and ANSI casts raise on malformed input.

Spark 4.0 turns ANSI mode on, so a malformed cast raises

`CAST_INVALID_INPUT`. `try_cast` returns `NULL` instead, which you then count.

### Common wrong answers.

- `a`: That was the behaviour before ANSI mode became the default. On Spark 4.0 use `try_cast` to get `NULL`.
- `c`: Casting never drops rows. It either raises or returns a value.

- `d`: Spark never turns malformed text into 0.0. It raises, or with `try_cast` returns NULL.
- `b|r2`: Right outcome, wrong reason. Delta accepts NULLs unless a NOT NULL constraint says otherwise. The error comes from the ANSI cast.

### A2-L2 exercise 3

**Hints.** (1) The safe variant of cast starts with a word meaning attempt. (2) It exists as a SQL function and, since PySpark 4.0, as a Column method. (3) `try_cast`.

**Answer.**

```
1 df = df.withColumns({c: F.col(c).try_cast(t) for c, t in
 type_map.items()})
```

`Column.try_cast` returns NULL where the value cannot be cast. Count those NULLs and let a data-quality check decide what to do.

**Common wrong answers.**

- `g1 = cast`: Under ANSI mode on Spark 4.0, `cast` raises on the first bad value. You need the variant that returns NULL.
- `g1 = astype`: `astype` is an alias of `cast`, so it raises under ANSI mode too.

### A2-L2 exercise 4

**Hints.** (1) Find the missing columns with a set difference. (2) `missing = sorted(set(column_map) - set(df.columns))`. (3) Raise if `missing` is not empty, then `return df.withColumnsRenamed(column_map)`.

**Answer.**

```
1 from pyspark.sql import DataFrame
2
3
4 def rename_to_canonical(df: DataFrame, column_map: dict[str, str]) ->
 DataFrame:
5 """Rename supplier columns to registry names.
6
7 Args:
8 df: A bronze delivery.
9 column_map: Supplier name to canonical name.
10
11 Returns:
12 The frame with canonical names.
13
14 Raises:
15 ValueError: If any source column in column_map is missing from df.
16 """
17 missing = sorted(set(column_map) - set(df.columns))
```

```

18 if missing:
19 raise ValueError(f"missing source columns: {missing}")
20 return df.withColumnsRenamed(column_map)

```

The set difference lists every missing source column in one go. Only then is the rename safe, because `withColumnsRenamed` would skip missing columns without telling you.

## A2-L2 exercise 5

**Hints.** (1) What is silver's grain? (2) Which columns does the deduplication treat as the key? (3) The key must include `wave_id`.

**Answer.** The bug is on line 5.

Fixed version:

```

1 def prepare_silver(df: DataFrame, column_map: dict[str, str], type_map:
 dict[str, str]) -> DataFrame:
2 """Rename, cast and keep the newest version per respondent and
 wave."""
3 renamed = rename_to_canonical(df, column_map)
4 typed = cast_to_registry_types(renamed, type_map)
5 return latest_per_key(typed, ["respondent_id", "wave_id"],
 "delivery_version")

```

Deduplicating on `respondent_id` alone keeps one wave per respondent and throws the others away. The key must match the grain: respondent and wave.

### Common wrong answers.

- Line 3: The rename is checked and correct. Look at which keys the deduplication uses.
- Line 4: `try_cast` is the right choice for silver. The rows vanish in a later step.

## A2-L2 exercise 6

**Hints.** (1) Checks on source names come before anything renames them. (2) A CRITICAL result must stop the run before the write. (3) The MERGE needs a source with one row per key.

**Answer.**

1. Check that every mapped source column exists
2. Rename to canonical names
3. `try_cast` to registry types
4. Count cast failures and apply the WARN or CRITICAL threshold
5. Keep the newest version per respondent and wave

## 6. MERGE into silver

Check, rename, cast, count failures, deduplicate, merge. Each step depends on the one before, and nothing reaches silver until the checks pass.

### Common wrong answers.

- After the rename, the source names are gone. Check the source columns first.
- Count cast failures before the MERGE. A CRITICAL failure must stop bad data reaching silver.
- Deduplicating after the MERGE is too late. A repeated key fails the MERGE, or inserts twice when the target has no match.

## A2-L3 Gold: state the grain, group by all of it, guard every ratio

---

### A2-L3 exercise 1

**Hints.** (1) List the grain columns in the docstring. (2) Compare them with the `groupBy`. (3) `tracker_id` is missing.

**Answer.** The bug is on line 3.

Fixed version:

```
1 def base_by_subgroup(df: DataFrame) -> DataFrame:
2 """Weighted base at tracker x wave x region grain."""
3 return df.groupBy("tracker_id", "wave_id",
4 "region").agg(F.sum("final_weight").alias("weighted_base"))
```

The grain names three columns, and the `groupBy` has two. Without `tracker_id`, rows from different trackers fold into one.

### Common wrong answers.

- Line 2: The docstring is the contract, and it is right. The code has to match it.

### A2-L3 exercise 2

**Hints.** (1) What does `F.when(cond, 1)` return when `cond` is false? (2) What is the sum of a column that holds only NULL? (3) NULL prints as None in Python.

**Answer.**

```
1 None
```

`F.when` without `otherwise` yields NULL for the unmatched row, and the sum of only NULLs is NULL. Power BI shows a blank instead of 0. Wrap the sum in `F.coalesce(..., F.lit(0))`.

### Common wrong answers.

- 0: `F.when` without `otherwise` returns NULL for non-matching rows, and sum of only NULLs is NULL. Wrap it in `F.coalesce(..., F.lit(0))`.
- 1: The `when` condition is false for the only row, so there is nothing to add.

## A2-L3 exercise 3

**Hints.** (1) The first function returns its first non-null argument. (2) The second returns a value only when a condition holds. (3) `coalesce` and `when`.

### Answer.

```
1 F.coalesce(F.sum(reached), F.lit(0.0)).alias("weighted_reach")
2 ...
3 F.when(F.col("weighted_base") > 0, F.col("weighted_reach") /
 F.col("weighted_base"))
```

`F.coalesce(x, 0)` turns a NULL sum into 0. `F.when(base > 0, ratio)` returns NULL when the base is 0, because there is no `otherwise`.

### Common wrong answers.

- `g1 = ifnull`: `ifnull` works in SQL, but the standard DataFrame function for this is `F.coalesce`, which takes any number of columns.
- `g2 = where`: `where` filters rows. You want a conditional value per row: `F.when`.
- `g1 = nvl`: `nvl` is a SQL function name. In the DataFrame API use `F.coalesce`.

## A2-L3 exercise 4

**Hints.** (1) Build the reached-weight expression with `F.when`, then aggregate with `groupBy(*grain)`. (2) Wrap each sum in `F.coalesce(..., F.lit(0.0))`. (3) Add the percentage with `F.when(F.col("weighted_base") > 0, reach / base)`.

### Answer.

```
1 from pyspark.sql import DataFrame
2 from pyspark.sql import functions as F
3
4
5 def reach_metrics(df: DataFrame, grain: list[str]) -> DataFrame:
6 """Weighted reach, one row per grain key.
7
8 Args:
9 df: In-scope respondents with reach_flag and final_weight.
```



```

10 grain: Columns that define one output row.
11
12 Returns:
13 unweighted_n, weighted_base, weighted_reach and weighted_reach_pct
14 per grain key.
15 """
16 reached = F.when(F.col("reach_flag") == 1, F.col("final_weight"))
17 metrics = df.groupBy(*grain).agg(
18 F.count("*").alias("unweighted_n"),
19 F.coalesce(F.sum("final_weight"),
20 F.lit(0.0)).alias("weighted_base"),
21 F.coalesce(F.sum(reached), F.lit(0.0)).alias("weighted_reach"),
22)
23 return metrics.withColumn(
24 "weighted_reach_pct",
25 F.when(F.col("weighted_base") > 0, F.col("weighted_reach") /
26 F.col("weighted_base")),
27)

```

`groupBy(*grain)` takes the grain as data, so the same function serves every gold table. Coalesce turns empty sums into 0, and the guarded ratio returns NULL for an empty base.

## A2-L3 exercise 5

**Hints.** (1) Three of these are ordinary custom properties or comments. (2) One property key has a special meaning in Spark. (3) Ownership is set with ALTER TABLE ... OWNER TO.

**Answer.** c) ALTER TABLE research\_gold.nat\_metrics SET TBLPROPERTIES ('owner' = 'data.engineering@example.org') Because: r1) owner is a reserved table property, so Spark rejects setting it this way.

Spark reserves owner and rejects it with UNSUPPORTED\_FEATURE.SET\_TABLE\_PROPERTY. Your reference unity\_catalog.sql sets it, so that script fails. Use ALTER TABLE ... OWNER TO for ownership, and a comment or custom property for the grain.

### Common wrong answers.

- a: COMMENT ON TABLE is valid and a good home for the grain.
- b: A custom property such as grain is valid. Tools can read it with SHOW TBLPROPERTIES.
- d: A custom property such as layer is valid.
- c|r2: Right statement, wrong reason. Property values can be any string. The key owner is reserved, and Spark raises UNSUPPORTED\_FEATURE.SET\_TABLE\_PROPERTY.

## A3-L1 Idempotent MERGE, and why you deduplicate the source first

---

### A3-L1 exercise 1

**Hints.** (1) Count how many source rows match the one target row. (2) Ask yourself which of the two source rows Delta should use for the update. (3) Delta raises `DELTA_MULTIPLE_SOURCE_ROW_MATCHING_TARGET_ROW_IN_MERGE` in this case.

**Answer.** b) The MERGE fails with an error. Because: r1) Two source rows match one target row, and Delta cannot tell which should win.

Delta requires at most one source row per matched target row. With two, the update is ambiguous, so the `MERGE` fails with

`DELTA_MULTIPLE_SOURCE_ROW_MATCHING_TARGET_ROW_IN_MERGE`. Deduplicate the source first.

#### Common wrong answers.

- a: Delta has no notion of 'last' in an unordered DataFrame. Row order in Spark is not guaranteed, so Delta refuses to guess.
- d: Spark gives no row order without an explicit sort, so 'first' has no meaning. Delta refuses to guess.
- c: Both rows match the existing target row, so neither goes to the insert clause. Duplicate inserts happen only when no target row matches.
- b|r2: Right outcome, wrong reason. Delta does not enforce primary keys. The error comes from the ambiguous match, not from a key constraint.
- b|r3: Right outcome, wrong reason. Caching has nothing to do with it. The error comes from two source rows matching one target row.

### A3-L1 exercise 2

**Hints.** (1) Two clauses and one action. (2) The clause names start with `whenMatched` and `whenNotMatched`, and both copy every column. (3) The copy-every-column versions end in `ALL`. The action that runs a MERGE is not `collect`.

#### Answer.

```
1 (
2 DeltaTable.forName(spark,
3 "research_silver.nat_respondents").alias("t")
4 .merge(source.alias("s"), "t.respondent_id = s.respondent_id AND
5 t.wave_id = s.wave_id")
```

```
4 .whenMatchedUpdateAll()
5 .whenNotMatchedInsertAll()
6 .execute()
7)
```

`whenMatchedUpdateAll()` copies every source column onto the matched target row. `whenNotMatchedInsertAll()` inserts rows with no match. `execute()` runs the MERGE.

### Common wrong answers.

- `g3 = collect: collect()` is a DataFrame action. A MERGE builder runs with `execute()`.
- `g1 = whenMatchedUpdate: whenMatchedUpdate()` needs a `set=` mapping of columns. With no arguments you want `whenMatchedUpdateAll()`.
- `g2 = whenNotMatchedBySourceDelete: whenNotMatchedBySourceDelete()` deletes target rows that are missing from the source. With a one-wave source it would delete every other wave.

### A3-L1 exercise 3

**Hints.** (1) Work through the first run: which clause fires for each row? (2) On the second run, does any source row lack a match? (3) Matched rows update in place. Updates do not add rows.

### Answer.

```
1 3
```

Run one inserts 3 rows. Run two matches all 3 and updates them with identical values. The count stays 3. That is idempotence.

### Common wrong answers.

- 6: 6 is what an append would give. On the second run every source row matches an existing target row, so the MERGE updates instead of inserting.
- 0: The first run inserts all 3 rows, because none of them matches an empty table.

### A3-L1 exercise 4

**Hints.** (1) The grain is respondent per wave. Which columns identify one row? (2) Compare the deduplication keys with the MERGE condition. (3) A `w2` row for respondent `R1` matches the `w1` row for `R1`, and overwrites it.

**Answer.** The bug is on line 6.

Fixed version:

```

1 def upsert_respondents(spark: SparkSession, source: DataFrame) -> None:
2 """MERGE one wave of respondents into silver. Grain: respondent per
 wave."""
3 latest = latest_per_key(source, ["respondent_id", "wave_id"],
4 "delivery_version")
5 (
6 DeltaTable.forName(spark,
7 "research_silver.nat_respondents").alias("t")
8 .merge(latest.alias("s"), "t.respondent_id = s.respondent_id AND
9 t.wave_id = s.wave_id")
10 .whenMatchedUpdateAll()
11 .whenNotMatchedInsertAll()
12 .execute()
13)

```

The condition matches on `respondent_id` alone. A wave W2 row for a respondent matches that respondent's W1 row and overwrites it, including `wave_id`. Match on the full grain: `t.respondent_id = s.respondent_id AND t.wave_id = s.wave_id`.

### Common wrong answers.

- Line 3: Line 3 deduplicates on both key columns, which matches the grain. Look at what the MERGE matches on.
- Line 7: `whenMatchedUpdateAll()` is right for a redelivery. The problem is which rows count as a match.

## A3-L1 exercise 5

**Hints.** (1) You need a window, a rank and a filter. (2)

`Window.partitionBy(*keys).orderBy(F.col(order_col).desc())` sorts each key's rows newest first. (3) Add `F.row_number().over(window)` as a column, keep rank 1, then drop the helper column.

### Answer.

```

1 from pyspark.sql import DataFrame, Window
2 from pyspark.sql import functions as F
3
4
5 def latest_per_key(df: DataFrame, keys: list[str], order_col: str) ->
 DataFrame:
6 """Keep the row with the highest order_col for each key.
7
8 Args:
9 df: Source rows, possibly with repeated keys.
10 keys: Columns that identify one target row.
11 order_col: Column that ranks versions. Highest wins.
12
13 Returns:
14 One row per key.

```

```

15 """
16 window = Window.partitionBy(*keys).orderBy(F.col(order_col).desc())
17 return (
18 df.withColumn("_rank", F.row_number().over(window))
19 .filter(F.col("_rank") == 1)
20 .drop("_rank")
21)

```

A window partitioned by the keys and ordered by `order_col` descending puts the newest row first in each key. `row_number()` numbers them 1, 2, 3. Keep 1, drop the helper column.

### A3-L1 exercise 6

**Hints.** (1) Nothing unchecked should reach silver. (2) The MERGE needs a source with one row per key. (3) The run log records the outcome of everything before it.

#### Answer.

1. Validate the bronze delivery against its data contract
2. Rename and cast columns from the Variable Registry
3. Keep the newest row per respondent and wave
4. MERGE into silver on `respondent_id` and `wave_id`
5. Run the table health check on silver
6. Write one record to the pipeline run log

Check the delivery, map it to canonical names, deduplicate, merge, confirm the table is healthy, then log the run with its final status.

#### Common wrong answers.

- Deduplicating after the MERGE is too late. The MERGE has already failed, or has already inserted the repeated rows.
- The contract describes the delivery as received. Check it before you rename, or the column names in the contract no longer match.
- Log last, so the run record includes the health check result and the final status.

## A3-L2 Schema enforcement is on by default. Evolve on purpose.

### A3-L2 exercise 1

**Hints.** (1) Is schema evolution on or off by default? (2) The table has no region column. (3) Enforcement rejects the write.

**Answer.** c) The append fails with a schema mismatch. Because: r1) Schema enforcement rejects writes with columns the table does not have.

Enforcement is on by default, so a write with an extra column fails. You opt in to evolution per write, or change the table with `ALTER TABLE`.

### Common wrong answers.

- a: That needs `mergeSchema`. Without it, enforcement rejects the new column.
- b: An append does not drop columns. A MERGE with `updateAll` does, which makes it the dangerous case.
- d: A rescue column is an Auto Loader feature (`_rescued_data`), not a Delta write behaviour.
- c|r2: Right outcome, wrong reason. Delta schemas can change, with `mergeSchema`, `withSchemaEvolution` or `ALTER TABLE`. The default is to enforce.

### A3-L2 exercise 2

**Hints.** (1) Does a MERGE fail on an extra source column the way an append does? (2) `updateAll` and `insertAll` work on the target's columns. (3) The extra column is ignored without an error.

### Answer.

```
1 ['id', 'w']
```

Without schema evolution, `updateAll` and `insertAll` use the target's columns and quietly ignore `region`. Your own column check is the only guard here.

### Common wrong answers.

- `['id', 'w', 'region']`: Only with `.withSchemaEvolution()` on the builder. Without it, `updateAll` and `insertAll` copy the target's columns and ignore `region`.
- `AnalysisException`: No error. The MERGE succeeds and the extra source column is silently ignored. That is why this case is dangerous.

### A3-L2 exercise 3

**Hints.** (1) The builder method names what it enables. (2) It exists on `DeltaMergeBuilder` in `delta-spark 4.0`. (3) `withSchemaEvolution`.

### Answer.

```
1 (
2 DeltaTable.forName(spark, table).alias("t")
```

```

3 .merge(source.alias("s"), condition)
4 .withSchemaEvolution()
5 .whenMatchedUpdateAll()
6 .whenNotMatchedInsertAll()
7 .execute()
8)

```

`withSchemaEvolution()` turns evolution on for this MERGE only, which is safer than a session-wide setting.

### Common wrong answers.

- `g1 = mergeSchema`: `mergeSchema` is a write option for appends and overwrites. The MERGE builder uses `withSchemaEvolution()`.
- `g1 = option`: The MERGE builder has no `option` method for this. Call `withSchemaEvolution()`.

### A3-L2 exercise 4

**Hints.** (1) Two set differences, one in each direction. (2) Wrap each in `sorted()`. (3) `assert_no_drift` calls `schema_drift` and raises if either list is non-empty.

### Answer.

```

1 def schema_drift(actual: list[str], expected: list[str]) -> dict[str,
 list[str]]:
2 """Compare delivered columns with the registry."""
3
4 Args:
5 actual: Columns in the delivery.
6 expected: Canonical columns from the registry.
7
8 Returns:
9 Sorted unexpected and missing column names.
10 """
11 return {
12 "unexpected": sorted(set(actual) - set(expected)),
13 "missing": sorted(set(expected) - set(actual)),
14 }
15
16
17 def assert_no_drift(actual: list[str], expected: list[str]) -> None:
18 """Raise ValueError when the delivery's columns differ from the
 registry."""
19 drift = schema_drift(actual, expected)
20 if drift["unexpected"] or drift["missing"]:
21 raise ValueError(f"schema drift: {drift}")

```

Checking both directions matters because Delta treats them differently: an extra column fails an append, but a missing column lands as NULL without an error.

## A3-L2 exercise 5

**Hints.** (1) Which line changes a setting? (2) What does that setting do to every MERGE in the session? (3) Remove the `conf.set`. Evolve deliberately with `withSchemaEvolution()` when a change is approved.

**Answer.** The bug is on line 3.

Fixed version:

```
1 def upsert_silver(spark: SparkSession, source: DataFrame, table: str,
2 keys: list[str]) -> None:
3 """MERGE a deduplicated source into silver."""
4 condition = " AND ".join(f"t.{k} = s.{k}" for k in keys)
5 (
6 DeltaTable.forName(spark, table).alias("t")
7 .merge(source.alias("s"), condition)
8 .whenMatchedUpdateAll()
9 .whenNotMatchedInsertAll()
10 .execute()
11)
```

Line 3 turns on automatic schema evolution for the whole Spark session, so every MERGE adds any new source column. Remove it. When you approve a new column, use `ALTER TABLE ... ADD COLUMN` or `.withSchemaEvolution()` on that one MERGE.

### Common wrong answers.

- Line 8: `updateAll` is fine. It only adds columns when evolution is switched on, and something switched it on.
- Line 4: The condition is correct. Look for a setting that changes how every MERGE treats new columns.

## A3-L3 Time travel, RESTORE, OPTIMIZE and VACUUM

### A3-L3 exercise 1

**Hints.** (1) The old version's files still exist. (2) Which command makes an old version current? (3) `RESTORE TABLE ... TO VERSION AS OF n`.

**Answer.** b) `RESTORE TABLE` to the version before the bad write. Because: r1) The earlier version's files are still on storage, and `RESTORE` makes that version current again.

`RESTORE` points the table back at the old version's files, as a new version. History is preserved, so you can inspect the bad write later.



### Common wrong answers.

- a: It works, but it takes hours. RESTORE takes seconds while the old files still exist.
- c: VACUUM deletes old files, which are the ones you need to recover. It makes things worse.
- d: OPTIMIZE compacts files in the current version. It does not change the data.
- b|r2: Right action, wrong reason. RESTORE writes a new version. The bad write stays in history.

### A3-L3 exercise 2

**Hints.** (1) The time travel clause has three words. (2) The same clause appears in the second statement. (3) `VERSION AS OF` and `RESTORE`.

#### Answer.

```
1 SELECT * FROM research_gold.nat_metrics VERSION AS OF 41;
2 RESTORE TABLE research_gold.nat_metrics TO VERSION AS OF 41;
```

`VERSION AS OF n` reads a past version. `RESTORE TABLE ... TO VERSION AS OF n` makes it current.

### Common wrong answers.

- g1 = `AT VERSION`: The Delta syntax is `VERSION AS OF`.
- g2 = `ROLLBACK`: `ROLLBACK` undoes a transaction. Delta's command to make an old version current is `RESTORE`.

### A3-L3 exercise 3

**Hints.** (1) Does `RESTORE` delete history? (2) Every commit adds one version. (3) The next version after 3.

#### Answer.

```
1 4
```

`RESTORE` is a commit like any other, so it becomes version 4 with the contents of version 1. Versions 2 and 3 remain in history.

### Common wrong answers.

- 1: `RESTORE` does not rewind the version counter. It writes a new version whose content matches version 1.
- 3: Version 3 is the `DELETE`. Every write, `RESTORE` included, adds a new version.

### A3-L3 exercise 4

**Hints.** (1) Which command deletes files? (2) What retention does it use? (3) Line 3. Run VACUUM with the default 7-day retention and drop line 2.

**Answer.** The bug is on line 3.

Fixed version:

```
1 OPTIMIZE research_silver.nat_respondents ZORDER BY (wave_id);
2 VACUUM research_silver.nat_respondents;
3 ANALYZE TABLE research_silver.nat_respondents COMPUTE STATISTICS;
```

RETAIN 0 HOURS deletes every file no longer in the current version, so no earlier version is readable. Line 2 disables the check that would have blocked it. Run VACUUM with default retention, or set the retention properties to the window you need.

#### Common wrong answers.

- Line 1: OPTIMIZE compacts files and keeps history. It does not block time travel.
- Line 2: Line 2 disables the guard, which is bad, but on its own it deletes nothing. The line that removes the files is the problem.

### A3-L3 exercise 5

**Hints.** (1) Which property decides how long VACUUM keeps old files? (2) The log must cover the same window. (3) VACUUM still runs, with default arguments.

#### Answer.

- Set both retention properties to 30 days. From Databricks Runtime 18.0, the log retention must be at least the deleted-file retention.
- Keep running a plain VACUUM t. It respects the 30-day property and stops storage growing without limit.

Both retention properties at 30 days keep every version from the last month readable. A plain VACUUM then cleans up only what is older.

#### Common wrong answers.

- deleted = d7: With 7 days, VACUUM removes files after a week, so versions older than 7 days are unreadable. The policy needs 30.
- vacuum = none: Without VACUUM, storage grows with every write. Keep it, and let the retention property protect 30 days.
- vacuum = zero: RETAIN 0 HOURS overrides the property and deletes history immediately.

- log = d7: The log must keep entries at least as long as the data files, or the versions cannot be reconstructed.

## A4-L1 The three-level namespace, one catalog per environment

---

### A4-L1 exercise 1

**Hints.** (1) Which level can you lock a whole environment behind? (2) Grants on the top level control everything below it. (3) The catalog.

**Answer.** a) Catalog Because: r1) Permissions and storage attach at catalog level, so dev principals can be kept out of prod entirely.

One catalog per environment gives a hard boundary. A dev job without privileges on `research_prod` cannot read or write it by accident.

#### Common wrong answers.

- b: Schemas are better used for layers. Putting environments there mixes dev and prod under one catalog's permissions.
- c: A prefix is a naming habit, not a boundary. Anyone with rights on the schema can write both.
- d: A column puts dev and prod rows in the same table. One bad filter mixes them.
- a|r2: Right level, wrong reason. The point is isolation: grants on `research_prod` do not reach `research_dev`, and the reverse.

### A4-L1 exercise 2

**Hints.** (1) `catalog.schema.table`. (2) The environment is the catalog. (3) `research_prod`, then `research_gold`.

#### Answer.

```
1 GOLD = "research_prod.research_gold.nat_metrics"
```

`research_prod.research_gold.nat_metrics`: environment, layer, table.

#### Common wrong answers.

- `g1 = research_gold`; `g2 = research_prod`: Swapped. The catalog comes first and holds the environment. The schema holds the layer.
- `g1 = hive_metastore`: `hive_metastore` is the legacy store. Production tables live in a Unity Catalog catalog such as `research_prod`.

## A4-L1 exercise 3

**Hints.** (1) Count the parts in the table name. (2) What decides the catalog for a two-part name? (3) Add `cfg.catalog` as the first part.

**Answer.** The bug is on line 3.

Fixed version:

```
1 def write_metrics(df: DataFrame, cfg: RunConfig) -> None:
2 """Write gold metrics for the configured environment."""
3 target = f"{cfg.catalog}.research_gold.{cfg.tracker}_metrics"
4 df.write.format("delta").mode("append").saveAsTable(target)
```

A two-part name resolves against the session's current catalog, which here is production. Build a three-part name from `cfg.catalog`.

### Common wrong answers.

- Line 4: Append is right for this job. The question is which table it writes to.

## A4-L1 exercise 4

**Hints.** (1) Keep the suffixes in a dict constant. (2) Check for a dot in catalog and for an unknown layer before building the name. (3)

```
f"{catalog}.research_{layer}.{tracker}_{LAYER_SUFFIX[layer]}"
```

### Answer.

```
1 LAYER_SUFFIX = {"bronze": "raw", "silver": "respondents", "gold":
2 "metrics"}
3
4 def table_name(catalog: str, layer: str, tracker: str) -> str:
5 """Three-level Unity Catalog name for a tracker's table in one layer.
6
7 Args:
8 catalog: Environment catalog, for example research_dev.
9 layer: bronze, silver or gold.
10 tracker: Tracker short name, for example nat.
11
12 Returns:
13 The fully qualified table name.
14
15 Raises:
16 ValueError: If the layer is unknown or the catalog contains a dot.
17 """
18 if "." in catalog:
19 raise ValueError(f"catalog must be a single name, got {catalog}/
20 r}")
21 if layer not in LAYER_SUFFIX:
22 raise ValueError(f"unknown layer {layer!r}")
23 return f"{catalog}.research_{layer}.{tracker}_{LAYER_SUFFIX[layer]}"
```

One function owns the naming convention. It rejects inputs that would produce a wrong name instead of silently creating a new table.

### A4-L1 exercise 5

**Hints.** (1) The strongest boundary goes at the top. (2) Layers group tables with the same contract. (3) Tables hold one dataset at one grain.

**Answer.**

- Catalog per environment: `research_dev`, `research_prod`.
- Schema per layer, plus `research_governance` for the run log, registry and contracts.
- Table per tracker and grain: `nat_respondents`.

Environment, layer, table: each level answers one question, and the full name tells a reader all three.

**Common wrong answers.**

- `layer = catalog`: A catalog per layer multiplies catalogs by environments, and grants must then be repeated everywhere. Use schemas for layers.
- `tracker = schema`: A schema per tracker scatters the layers. Keep layers as schemas, and trackers as tables within them.
- `env = schema`: Schemas under one catalog share its permissions and storage. Environments need the stronger boundary of separate catalogs.

## A4-L2 Grant to groups, grant the least, grant at the right level

---

### A4-L2 exercise 1

**Hints.** (1) The namespace goes catalog, schema, table. (2) Each USE privilege opens one level. (3) The data privilege comes last.

**Answer.**

1. GRANT USE CATALOG ON CATALOG `research_prod` TO `research_analysts`
2. GRANT USE SCHEMA ON SCHEMA `research_prod.research_gold` TO `research_analysts`
3. GRANT SELECT ON SCHEMA `research_prod.research_gold` TO `research_analysts`

Reading a table needs USE CATALOG, USE SCHEMA and SELECT. Granting SELECT on the schema makes every current and future table readable.

### Common wrong answers.

- The grants all work in any order, but read them outermost first: catalog, then schema, then data. This order is inside out.
- Catalog comes before schema. USE CATALOG opens the catalog, USE SCHEMA opens the schema inside it.

### A4-L2 exercise 2

**Hints.** (1) How many privileges does a read need? (2) Each parent level must be usable. (3) USE CATALOG and USE SCHEMA are missing.

**Answer.** b) The query fails, because they lack USE CATALOG and USE SCHEMA. Because: r1) Reading needs USE on the parent catalog and schema as well as SELECT on the table.

A read needs USE CATALOG, USE SCHEMA and SELECT. With SELECT alone the query fails with a permission error.

### Common wrong answers.

- a: SELECT on the table is necessary but not enough. The catalog and schema must be usable too.
- c: Missing privileges raise a permission error. They never return an empty result.
- d: Masks are unrelated. The query fails before any column is read.
- b|r2: Right outcome, wrong reason. SELECT can be granted on a table. The missing pieces are USE CATALOG and USE SCHEMA.

### A4-L2 exercise 3

**Hints.** (1) Which principal is not a group? (2) An email address names one person. (3) Create a group such as research\_silver\_readers and grant to it.

**Answer.** The bug is on line 3.

Fixed version:

```
1 GRANT USE CATALOG ON CATALOG research_prod TO `research_analysts`;
2 GRANT USE SCHEMA, SELECT ON SCHEMA research_prod.research_gold TO
 `research_analysts`;
3 GRANT USE CATALOG ON CATALOG research_prod TO `research_silver_readers`;
4 GRANT USE SCHEMA, SELECT ON SCHEMA research_prod.research_silver TO
 `research_silver_readers`;
```

Line 3 grants to one person. Grant to a group instead, and add Sam to the group. When Sam moves teams, one membership change removes the access.

### Common wrong answers.

- Line 2: That grant goes to a group at schema level, which is the pattern to follow.
- Line 1: USE CATALOG to a group is correct and needed.

### A4-L2 exercise 4

**Hints.** (1) Two statements are enough. (2) One for the catalog, one for the schema. (3) `GRANT USE SCHEMA, SELECT, MODIFY ON SCHEMA research_prod.research_silver TO ...`.

### Answer.

```
1 GRANT USE CATALOG ON CATALOG research_prod TO `research-pipeline-sp`;
2 GRANT USE SCHEMA, SELECT, MODIFY ON SCHEMA research_prod.research_silver
 TO `research-pipeline-sp`;
```

USE CATALOG opens the catalog. USE SCHEMA, SELECT and MODIFY on the schema cover every table in it, now and in future. Nothing else is needed.

### A4-L2 exercise 5

**Hints.** (1) Who writes, and who only reads? (2) Does anyone in the list need raw deliveries? (3) Groups, not users, for every grant.

### Answer.

- Analysts read gold: USE CATALOG, USE SCHEMA and SELECT on the gold schema.
- The pipeline writes silver: the same plus MODIFY on the silver schema.
- Analysts have no reason to see raw deliveries, so no bronze access.
- Administration goes through an admin group, never through grants on user accounts.

Least privilege: people read the layer they report from, pipelines write the layers they own, and admin rights live in a group.

### Common wrong answers.

- analysts = `modify`: Reports only read. MODIFY would let an analyst change gold outside the pipeline.
- analysts\_bronze = `select`: Bronze holds raw deliveries, which can include fields gold never exposes. Give access only when there is a need.
- admin = `user_all`: Grants to user accounts outlive the people's roles. Use a group.

## A4-L3 Row filters and column masks in Unity Catalog

### A4-L3 exercise 1

**Hints.** (1) The filter is a SQL function. (2) You attach it with ALTER TABLE. (3) SET ROW FILTER ... ON (column).

**Answer.** b) ALTER TABLE t SET ROW FILTER f ON (col) Because: r1) Databricks row filters are SQL UDFs bound to a table with SET ROW FILTER. Create a SQL UDF that returns BOOLEAN, then bind it with ALTER TABLE ... SET ROW FILTER f ON (col).

#### Common wrong answers.

- a: That is Snowflake syntax. Databricks raises a parse error on it.
- c: GRANT has no WHERE clause. Grants control the whole object.
- d: A view filters only readers who use the view. Anyone with SELECT on the table bypasses it. A row filter applies to the table itself.
- b|r2: Right statement, wrong reason. You can add a row filter to an existing table with ALTER TABLE.

### A4-L3 exercise 2

**Hints.** (1) The function checks account-level group membership. (2) The column clause has two words. (3) is\_account\_group\_member and SET MASK.

#### Answer.

```
1 CREATE FUNCTION research_prod.research_governance.mask_postcode(postcode
 STRING)
2 RETURN CASE WHEN is_account_group_member('pii_readers') THEN postcode
 ELSE '***' END;
3
4 ALTER TABLE research_prod.research_silver.nat_respondents
5 ALTER COLUMN postcode SET MASK
 research_prod.research_governance.mask_postcode;
```

is\_account\_group\_member('pii\_readers') checks the account group. ALTER COLUMN ... SET MASK f binds the mask to the column.

#### Common wrong answers.

- g1 = is\_member: is\_member checks workspace-local groups. Unity Catalog uses account-level groups: is\_account\_group\_member.
- g1 = current\_user: current\_user() returns a name, not a group check. Masks should test group membership.
- g2 = SET ROW FILTER: Row filters apply to the table. A column uses SET MASK.



### A4-L3 exercise 3

**Hints.** (1) Which line names one identity? (2) Access reviews look at groups and grants, not at function bodies. (3) Replace the email check with a group check.

**Answer.** The bug is on line 3.

Fixed version:

```
1 CREATE OR REPLACE FUNCTION
 research_prod.research_governance.respondent_filter(tracker_id STRING)
2 RETURN is_account_group_member('research_approved_analysts')
3 OR is_account_group_member('research_pipelines');
```

Line 3 grants access to one named identity inside the function body, where no access review looks. Put the service principal in a group and test the group.

#### Common wrong answers.

- Line 2: Testing an account group is the right pattern. Look for a named person or identity.
- Line 1: A governance schema is the right home for the function.

### A4-L3 exercise 4

**Hints.** (1) The function returns TRUE for rows the reader may see. (2) Group members see everything, so the group check stands alone on one side of OR. (3) Then bind with ALTER TABLE ... SET ROW FILTER ... ON (region).

**Answer.**

```
1 CREATE OR REPLACE FUNCTION
 research_prod.research_governance.region_filter(region STRING)
2 RETURN is_account_group_member('research_all_regions') OR region =
 'NATIONAL';
3
4 ALTER TABLE research_prod.research_gold.nat_metrics
5 SET ROW FILTER research_prod.research_governance.region_filter ON
 (region);
```

group OR region = 'NATIONAL' is TRUE for every row when the reader is in the group, and only for national rows otherwise. AND would hide everything from people outside the group. The function does nothing until SET ROW FILTER binds it.

### A4-L3 exercise 5

**Hints.** (1) Whole table, some rows, one column. (2) Grants are all or nothing on an object. (3) Filters remove rows. Masks change values.

## Answer.

- No access at all: do not grant SELECT. Grants decide who reads the table.
- Some rows only: a row filter decides which rows each reader sees.
- All rows, hidden values: a column mask decides what each reader sees in one column.

Grants control access to the object, filters control rows, masks control column values. Combine them: grant SELECT to analysts, filter trackers, mask postcodes.

## Common wrong answers.

- `team = filter`: A filter that returns nothing still lets the team query the table and see its structure. Withhold the grant.
- `tracker = mask`: A mask changes values in a column. The requirement removes whole rows.
- `postcode = filter`: Filtering would remove the respondents entirely. The requirement keeps the rows and hides one column.

## A5-L1 Lazy evaluation, and why every `count()` re-runs the plan

---

### A5-L1 exercise 1

**Hints.** (1) Which call returns a plain Python value instead of a DataFrame? (2) Transformations return DataFrames. (3) `count()` returns an int, so it must run the plan.

**Answer.** c) `n = df3.count()` Because: r1) `count()` is an action. The other lines are transformations that only extend the plan.

`count()` returns a number, so Spark must execute the plan. The other lines return DataFrames and only describe work.

## Common wrong answers.

- a: `filter` is a transformation. It returns a new DataFrame without reading any data.
- b: `withColumn` is a transformation. Nothing runs yet.
- d: `groupBy().agg()` is still a transformation. It runs only when an action such as `show` or `write` asks for the result.
- c|r2: Right line, wrong reason. Every line runs as Python. Only actions make Spark execute the plan.

## A5-L1 exercise 2

**Hints.** (1) How many actions are there? (2) Does Spark keep results between actions without cache? (3) Two full runs over 1,000 rows.

**Answer.**

```
1 2000
```

Nothing is kept between actions, so each `count()` runs the UDF on all 1,000 rows. This was measured locally on PySpark 4.0.4.

### Common wrong answers.

- `1000`: Each action re-runs the whole plan, UDF included. Two counts mean two full evaluations.
- `500`: The filter comes after the UDF, so the UDF runs on all 1,000 rows each time.
- `0`: Defining `df` runs nothing, but each `count()` is an action that executes the plan.

## A5-L1 exercise 3

**Hints.** (1) The `DataFrame` method that takes several aggregate expressions. (2) The function that counts rows where a condition is true. (3) `agg` and `count_if`.

**Answer.**

```
1 stats = df.agg(
2 F.count("*").alias("rows"),
3 F.count_if(F.col("final_weight").isNull()).alias("null_weights"),
4).collect()[0]
```

`df.agg(...)` evaluates every expression in one pass. `F.count_if(cond)` counts rows where `cond` is true.

### Common wrong answers.

- `g1 = select`: `select` with aggregate functions also works, but `agg` is the idiomatic call for a single aggregated row.
- `g2 = count`: `F.count` counts non-null values of an expression, and `isNull()` is never null. Every row would count. Use `count_if`.
- `g2 = sum`: `F.sum` of a boolean fails. You would need `.cast("int")` first. `count_if` says what you mean.

## A5-L1 exercise 4

**Hints.** (1) Every value can be an aggregate expression. (2) `F.count("*")`, `F.count_if(...)`, `F.min(...)`, `F.max(...)`. (3) Put all five in one `df.agg(...)`, collect once, and return `row.asDict()`.

**Answer.**

```
1 def table_health(df: DataFrame) -> dict[str, float | int | None]:
2 """Row count, null counts and weight range in one pass over df."""
3 row = df.agg(
4 F.count("*").alias("rows"),
5 F.count_if(F.col("respondent_id").isNull()).alias("null_ids"),
6 F.count_if(F.col("final_weight").isNull()).alias("null_weights"),
7 F.min("final_weight").alias("min_weight"),
8 F.max("final_weight").alias("max_weight"),
9).collect()[0]
10 return row.asDict()
```

Five aggregate expressions in one `agg` produce one row in one scan. The original ran five separate jobs over the same data.

## A5-L1 exercise 5

**Hints.** (1) Simple aggregates can share one pass. (2) Different shapes need the data stored somewhere reusable. (3) Serverless has no cache APIs.

**Answer.**

- Combine every count and sum into one `agg`. One scan answers them all.
- For checks that need different shapes, write the expensive result to a Delta table once and read it back.
- Serverless compute does not support cache APIs, so `cache()` is not an option there.

Combine what you can into one action. Store an expensive result in Delta when several differently shaped queries need it. Keep `cache()` for classic clusters.

### Common wrong answers.

- `several_counts = repeat`: Each action re-runs the join. Four checks would run the expensive join four times.
- `many_shapes = cache`: On serverless compute, cache APIs are not supported. Materialise to a Delta table instead.
- `cache_serverless = use`: Databricks documents that DataFrame and SQL cache APIs are not supported on serverless compute.

## A5-L2 Built-in functions over Python UDFs

---

### A5-L2 exercise 1

**Hints.** (1) Which option never leaves the JVM? (2) Are trim and upper available as built-ins? (3) `F.trim` and `F.upper`.

**Answer.** c) `F.upper(F.trim(F.col("s")))` Because: r1) Built-ins run inside the JVM with no data transfer to Python.

Built-ins run in Spark's generated JVM code on the executors. UDFs of either kind pay to move data to Python and back.

#### Common wrong answers.

- a: A Python UDF moves every value to a Python process and back. It measured about six times slower than a built-in.
- b: A pandas UDF beats a Python UDF because it uses Arrow batches, but it still leaves the JVM. A built-in wins.
- d: `collect()` pulls 50 million values onto one machine. It is the slowest option and can crash the driver.
- c|r2: Right choice, wrong reason. Built-ins run on the executors, like everything else. They win by staying in the JVM.

### A5-L2 exercise 2

**Hints.** (1) What does `\s+` match? (2) What do built-ins return for NULL input? (3) Whitespace removed, then upper-cased. NULL stays NULL.

#### Answer.

```
1 ['AB1', None]
```

The regex removes all whitespace, `upper` capitalises, and the NULL row stays NULL. No special handling needed.

#### Common wrong answers.

- `[' AB 1 ', None]`: `regexp_replace` with `\s+` removes every run of whitespace, including the inner space.
- `PythonException`: Built-ins handle NULL. Only a Python UDF would raise on the None row.
- `['AB1', '']`: NULL input gives NULL output, not an empty string.

### A5-L2 exercise 3

**Hints.** (1) Padding on the left has a function named after that side. (2) Pulling a capture group out of a string uses a regex function. (3) `lpad` and `regexp_extract`.

**Answer.**

```
1 df = df.withColumns({
2 "respondent_id": F.lpad(F.col("respondent_id"), 8, "0"),
3 "region": F.regexp_extract(F.col("area_code"), r"^[A-Z]{2}" -, 1),
4 })
```

`F.lpad(col, 8, "0")` adds leading zeros. `F.regexp_extract(col, pattern, 1)` returns capture group 1, or an empty string when there is no match.

**Common wrong answers.**

- `g1 = rpad`: `rpad` pads on the right. Leading zeros go on the left: `lpad`.
- `g2 = regexp_replace`: `regexp_replace` substitutes text. To pull out a capture group, use `regexp_extract` with the group index.
- `g2 = substring`: `substring` takes positions, not a pattern, and cannot check the hyphen. `regexp_extract` with group 1 matches the rule exactly.

### A5-L2 exercise 4

**Hints.** (1) The function takes a Column and returns a Column expression. (2) `F.when(age < 35, "16-34").when(age < 55, "35-54")`. (3) Finish with `.when(age.isNull(), "55+")`, not `.otherwise`.

**Answer.**

```
1 from pyspark.sql import Column
2 from pyspark.sql import functions as F
3
4
5 def age_band(age: Column) -> Column:
6 """16-34, 35-54 or 55+. NULL age gives NULL."""
7 return (
8 F.when(age < 35, "16-34")
9 .when(age < 55, "35-54")
10 .when(age.isNull(), "55+")
11)
```

`F.when` chains evaluate in order. Without `otherwise`, an unmatched row is `NULL`, and `age < 35` is `NULL` for a `NULL` age, so the final `isNull()` guard keeps `NULL` ages `NULL`.

## A5-L2 exercise 5

**Hints.** (1) What does the UDF receive for a NULL postcode? (2) `None.replace` raises `AttributeError`. (3) Replace the UDF with `F.upper(F.regexp_replace(...))`.

**Answer.** The bug is on line 6.

Fixed version:

```
1 from pyspark.sql import DataFrame
2 from pyspark.sql import functions as F
3
4
5 def clean_postcodes(df: DataFrame) -> DataFrame:
6 """Remove spaces from postcodes and upper-case them. NULL stays
7 NULL."""
8 return df.withColumn("postcode",
9 F.upper(F.regexp_replace(F.col("postcode"), r"\s+", "")))
```

For a NULL postcode the UDF receives `None`, and `None.replace` raises inside the Python worker, failing the job. Built-ins pass NULL through, and run faster.

### Common wrong answers.

- Line 4: The decorator makes it a UDF, which is slow, but the crash happens where a value is used.
- Line 11: The call is fine. The failure is inside the function, on a NULL value.

## A5-L3 Joins, broadcast, partitions and reading the plan

### A5-L3 exercise 1

**Hints.** (1) Which side is small? (2) What is the most expensive step in a join of two large tables? (3) Copy the small side to every executor.

**Answer.** b) BroadcastHashJoin, with the lookup sent to every executor  
Because: r1) Broadcasting the small side avoids shuffling the large side.

With a tiny lookup, broadcast it. The large table stays where it is, and each executor joins locally.

### Common wrong answers.

- a: A sort-merge join shuffles all 40 million rows across the network to meet 200 lookup rows. Broadcasting the lookup avoids that.
- c: That moves the join into a slow Python UDF. A broadcast join does the same job in the JVM.
- d: A cross join multiplies 40 million by 200 before filtering: 8 billion rows.

- `b|r2`: Right strategy, wrong reason. Broadcasting costs memory on every executor. It wins by avoiding the shuffle, and only when the broadcast side is small.

## A5-L3 exercise 2

**Hints.** (1) The hint names the strategy. (2) An equality join with a broadcast side uses a hash table. (3) Broadcast + Hash + Join.

**Answer.**

```
1 BroadcastHashJoin
```

`F.broadcast` makes Spark build a hash table from `regions` on every executor. The plan shows `BroadcastHashJoin` with `LeftOuter`.

**Common wrong answers.**

- `SortMergeJoin`: The broadcast hint tells Spark to broadcast `regions`, so the plan uses `BroadcastHashJoin`.
- `BroadcastNestedLoopJoin`: That node appears for joins without an equality condition. This join is on `region_code`, so it is a hash join.

## A5-L3 exercise 3

**Hints.** (1) The function names the join strategy. (2) Keep all rows from the left side. (3) The explain mode that splits the plan into numbered, detailed steps.

**Answer.**

```
1 joined = fact.join(F.broadcast(regions), on="region_code", how="left")
2 joined.explain(mode="formatted")
```

`F.broadcast(regions)` forces a broadcast hash join, `how="left"` keeps every fact row, and `mode="formatted"` prints a detailed plan with pushed filters.

**Common wrong answers.**

- `g2 = inner`: An inner join drops fact rows whose region code is missing from the lookup, without any warning. Use a left join.
- `g3 = verbose`: explain modes are simple, extended, codegen, cost and formatted. `formatted` gives the per-step details.
- `g1 = cache`: Caching does not change the join strategy. `F.broadcast` does.

## A5-L3 exercise 4

**Hints.** (1) How many partitions does the write use? (2) One partition means one task. (3) Remove `.repartition(1)`.

**Answer.** The bug is on line 4.



Fixed version:

```
1 def write_gold(metrics: DataFrame, table: str) -> None:
2 """Write gold metrics for the wave."""
3 metrics.write.format("delta").mode("append").saveAsTable(table)
```

repartition(1) funnels every row into one task, so 15 of 16 cores sit idle. Remove it. For file sizes, rely on OPTIMIZE, liquid clustering or predictive optimization.

### Common wrong answers.

- Line 6: Append is the right mode here. The slowness comes from how many tasks do the writing.
- Line 5: Delta is the right format. Look at the partition count.

### A5-L3 exercise 5

**Hints.** (1) `fact.join(F.broadcast(regions), on=..., how=...)`. (2) A left join keeps unknown codes, with region NULL. (3)

`enriched.agg(F.count_if(F.col("region").isNull())).collect()[0][0]`.

### Answer.

```
1 from pyspark.sql import DataFrame
2 from pyspark.sql import functions as F
3
4
5 def enrich_with_regions(fact: DataFrame, regions: DataFrame) -> DataFrame:
6 """Left-join fact rows to the small regions lookup, broadcasting the
 lookup."""
7 return fact.join(F.broadcast(regions), on="region_code", how="left")
8
9
10 def unmatched_codes(enriched: DataFrame) -> int:
11 """Rows whose region_code had no match in the lookup."""
12 return
13 int(enriched.agg(F.count_if(F.col("region").isNull())).collect()[0][0])
```

The broadcast left join keeps every fact row and never shuffles the fact table. Counting NULL regions afterwards turns silently unknown codes into a number your QA can check.

## A6-L1 Data contracts with severity levels and structured results

---

### A6-L1 exercise 1

**Hints.** (1) Would the problem corrupt totals or the grain? (2) Or is it something a person should look at, with no immediate harm? (3) Keys and weights are CRITICAL. Labels and small shifts are WARN.

**Answer.**

- Duplicate keys break the silver grain and the MERGE: CRITICAL.
- A weight total out of range makes every weighted figure wrong: CRITICAL.
- A few unexpected labels are worth investigating but do not corrupt totals: WARN.
- A small distribution shift can be real change: WARN, reviewed by a person.

CRITICAL is for problems that make published numbers wrong or break the table. WARN is for anomalies that need a human look.

**Common wrong answers.**

- keys = warn: Duplicate keys fail the MERGE or break the grain. Silver must not receive them.
- weights = warn: Every weighted estimate depends on the weight total. Publishing with wrong weights is worse than publishing late.
- category = critical: Twelve odd labels should not stop the whole wave. Log them and follow up with the supplier.

### A6-L1 exercise 2

**Hints.** (1) How many round trips with the supplier do you want? (2) Which failures must stop the run? (3) Store all results, then raise once.

**Answer.** b) All three results stored, then one exception naming both CRITICAL failures. Because: r1) Evaluate everything, record everything, then act on severity once.

Run every check, store every result, then raise one exception listing all CRITICAL failures. The WARN is stored and reviewed.

**Common wrong answers.**

- a: Stopping at the first failure hides the second, and the supplier needs another round trip.
- c: Printed text is lost. Results must be records in a table.
- d: Two CRITICAL failures must stop the run before silver.

- `b|r2`: Right outcome, wrong reason. Here you deliberately raise late, after all checks have run, so one run reports everything.

### A6-L1 exercise 3

**Hints.** (1) The typing construct for exact values. (2) The dataclass option that forbids changing fields after creation. (3) `Literal` and `frozen`.

**Answer.**

```
1 Status = Literal["PASS", "WARN", "FAIL", "SKIP"]
2
3
4 @dataclass(frozen=True)
5 class CheckResult:
6 check_name: str
7 status: Status
```

`Literal[...]` lists the allowed strings. `frozen=True` means no code can change a result after the check records it.

**Common wrong answers.**

- `g1 = Union`: `Union` combines types, not values. A fixed set of strings is a `Literal`.
- `g2 = slots`: `slots=True` saves memory but still allows changes. `frozen=True` makes a result immutable once recorded.

### A6-L1 exercise 4

**Hints.** (1) Collect all CRITICAL failures into a list first. (2) Raise once, joining their names into the message. (3) Return the WARN results when nothing CRITICAL failed.

**Answer.**

```
1 class ContractViolation(Exception):
2 """Raised when one or more CRITICAL checks fail."""
3
4
5 def enforce(results: list[dict[str, str]]) -> list[dict[str, str]]:
6 """Act on contract results.
7
8 Args:
9 results: Every check result for one delivery.
10
11 Returns:
12 The WARN results, for logging.
13
14 Raises:
15 ContractViolation: If any CRITICAL check failed. The message names
16 each one.
```

```

16 """
17 critical = [r for r in results if r["status"] == "FAIL" and
18 r["severity"] == "CRITICAL"]
19 if critical:
20 names = ", ".join(f"{r['check_name']}" (expected {r['expected']},
21 got {r['actual']})" for r in critical)
22 raise ContractViolation(f"{len(critical)} critical check(s)
23 failed: {names}")
24 return [r for r in results if r["status"] == "WARN"]

```

Collecting before raising means one run reports every CRITICAL failure. Filtering on both status and severity keeps a MEDIUM failure from stopping the pipeline.

## A6-L1 exercise 5

**Hints.** (1) Which `sum` does line 11 call? (2) An import can replace a built-in name for the whole module. (3) Import `functions` as `F` instead of importing `sum` by name.

**Answer.** The bug is on line 2.

Fixed version:

```

1 from dataclasses import dataclass, field
2 from pyspark.sql import functions as F
3
4
5 @dataclass
6 class ContractValidationResult:
7 checks: list = field(default_factory=list)
8
9 @property
10 def pass_rate(self) -> float:
11 passed = sum(1 for c in self.checks if c.status == "PASS")
12 return passed / len(self.checks) if self.checks else 0

```

Line 2 imports Spark's `sum`, which replaces Python's built-in for the whole module. Line 11 then passes a generator to a Spark column function, which fails. Import the module as `F` and use `F.sum` where you need Spark's version. This is defect 2 in your reference README.

## Common wrong answers.

- Line 11: The generator and the built-in `sum` are correct. But which `sum` does this line call?
- Line 12: The empty-list guard is fine. The failure happens one line earlier, caused by an import.

## A6-L1 exercise 6

**Hints.** (1) Three behaviours: keep, drop, fail. (2) The requirement says never reach the table, but keep going. (3) Drop.

**Answer.** b) `@dp.expect_or_drop("valid_id", "respondent_id IS NOT NULL")`

Because: r1) `expect_or_drop` removes violating rows and records how many were dropped.

`expect` keeps and records, `expect_or_drop` drops and records, `expect_or_fail` stops. Choose the one that matches the severity of the rule.

### Common wrong answers.

- a: `expect` keeps violating rows and only records metrics. NULL keys would reach the table.
- c: `expect_or_fail` stops the whole update on the first violation. The requirement keeps valid rows flowing.
- d: A later filter lets bad rows land first, and loses the violation metrics.
- b|r2: Right decorator, wrong reason. Failing the update is `expect_or_fail`. `expect_or_drop` drops the rows and keeps going.

## A6-L2 Statistical checks that survive empty, single-row and zero data

---

### A6-L2 exercise 1

**Hints.** (1) `F.stddev` is an alias for `stddev_samp`. (2) The sample formula divides by  $n - 1$ . (3) With  $n = 1$  the result is NULL, printed as None.

**Answer.**

```
1 None
```

`F.stddev` computes the sample standard deviation, dividing by  $n - 1$ . For one row that is undefined, so the result is NULL. Your checks must expect it.

### Common wrong answers.

- 0.0: `F.stddev` is the sample standard deviation, which divides by  $n - 1$ . With one row that is undefined, so Spark returns NULL.
- 1.2: That is the mean. The spread of a single value is undefined for the sample formula.

### A6-L2 exercise 2

**Hints.** (1) Which mode is on by default in Spark 4.0? (2) Does that mode tolerate division by zero? (3) It raises, and the message suggests `try_divide`.

**Answer.** c) The query fails with `DIVIDE_BY_ZERO`. Because: r1) ANSI mode is on by default in Spark 4.0, and it raises on division by zero, `DOUBLE` included.

Under ANSI mode, division by zero raises `DIVIDE_BY_ZERO`, even for doubles. Use `F.try_divide(a, b)` or `F.when(b > 0, a / b)`.

### Common wrong answers.

- a: Without ANSI mode some engines return Infinity for doubles. Spark 4.0 with ANSI on raises `DIVIDE_BY_ZERO`.
- b: `NULL` is what `F.try_divide` returns. The plain `/` raises under ANSI mode.
- d: No engine returns 0.0 for a division by zero. Spark 4.0 raises.
- c|r2: Right outcome, wrong reason. The error comes from the ANSI division, before anything is written.

### A6-L2 exercise 3

**Hints.** (1) The safe division function starts with `try`. (2) CV = standard deviation / mean. (3) `try_divide` and `mean`.

### Answer.

```
1 cv = F.try_divide(F.stddev("final_weight"), F.mean("final_weight"))
```

`F.try_divide(F.stddev(w), F.mean(w))` gives the CV, or `NULL` when the mean is 0 or `NULL`.

### Common wrong answers.

- `g1 = divide`: There is no `F.divide`. The safe version is `F.try_divide`, which returns `NULL` for a zero divisor.
- `g2 = stddev`: The CV divides the standard deviation by the mean.

### A6-L2 exercise 4

**Hints.** (1) Check `n` first, then `cv` with `is None`. (2) Only compare `cv` with `max_cv` once you know it is a number. (3) Return the status and a short reason in a tuple.

### Answer.

```
1 def check_weight_spread(stats: dict[str, float | int | None], max_cv:
 float = 1.5) -> tuple[str, str]:
2 """Classify the weight coefficient of variation.
3
4 Args:
5 stats: n (row count) and cv (None when undefined).
6 max_cv: Largest CV that passes.
```

```

7
8 Returns:
9 (status, reason), where status is SKIP, WARN or PASS.
10 """
11 n = stats["n"]
12 cv = stats["cv"]
13 if n is None or n < 2:
14 return ("SKIP", "fewer than 2 rows")
15 if cv is None:
16 return ("SKIP", "cv undefined")
17 if cv > max_cv:
18 return ("WARN", f"cv {cv:.4f} above {max_cv}")
19 return ("PASS", f"cv {cv:.4f}")

```

Ordering the checks means the comparison only ever sees a real number. `is None` keeps a CV of 0.0 as a pass.

## A6-L2 exercise 5

**Hints.** (1) What is `0.0` and `0.12` in Python? (2) The block is skipped for a share of exactly zero. (3) Use `is not None` on both shares.

**Answer.** The bug is on line 3.

Fixed version:

```

1 def compare_share(metric: str, current_pct: float | None, prior_pct:
 float | None, tolerance: float = 0.10) -> dict[str, object] | None:
2 """Return a stability result, or None when a share is missing."""
3 if current_pct is not None and prior_pct is not None:
4 change = abs(current_pct - prior_pct)
5 return {"metric": metric, "change_pp": round(change, 4),
6 "status": "WARN" if change > tolerance else "PASS"}
7 return None

```

`0.0` is falsy, so the largest possible drop skipped the comparison. This is defect 8 in your reference README. Test for missing values with `is not None`.

## Common wrong answers.

- Line 5: The rounding and status logic are fine once reached. The share of 0.0 never reaches them.
- Line 4: The subtraction is correct. The problem is the condition guarding it.

## A6-L3 Reconcile counts and totals between layers

### A6-L3 exercise 1

**Hints.** (1) What does silver keep when a key appears twice? (2) What does bronze keep? (3) Count distinct keys in bronze.

**Answer.** b) The number of distinct (respondent\_id, wave\_id) keys in bronze for W14. Because: r1) Silver keeps one row per key, the newest version, so it matches the number of distinct keys across all deliveries.

Silver holds one row per (respondent\_id, wave\_id). The matching bronze figure is the number of distinct keys across all deliveries of the wave.

**Common wrong answers.**

- a: Bronze holds both deliveries, so its row count is roughly double. Silver keeps one row per key.
- c: Only if both deliveries hold exactly the same keys. Distinct keys is correct in every case.
- d: A redelivery can add or drop respondents. Silver keeps every key ever delivered for the wave, at its newest version.
- b|r2: Right figure, wrong reason. Bronze is append-only and keeps every delivery.

**A6-L3 exercise 2**

**Hints.** (1) Can 0.1 be stored exactly as a binary float? (2) What is 0.1 + 0.2 printed in full? (3) isclose allows a tiny relative difference.

**Answer.**

```
1 False True
```

0.1 + 0.2 is 0.30000000000000004, so == is False. math.isclose allows a relative difference of 1e-9, so it is True. Weight totals need the same treatment.

**Common wrong answers.**

- True True: 0.1 and 0.2 have no exact binary representation, so their sum is 0.30000000000000004. Exact equality fails.
- False False: math.isclose compares with a relative tolerance of 1e-9 by default, so it treats the two values as equal.

**A6-L3 exercise 3**

**Hints.** (1) math.isclose(expected, actual, rel\_tol=1e-9, abs\_tol=1e-6). (2) Format numbers with an f-string: f"{x:.2f}". (3) difference is actual - expected, as a float.

**Answer.**

```
1 import math
2
3
```



```

4 def reconciliation_result(name: str, expected: float, actual: float) ->
 dict[str, object]:
5 """Compare two layer totals and return a CRITICAL check result.
6
7 Args:
8 name: Check name, for example weight_total_silver_gold.
9 expected: Total from the upstream layer.
10 actual: Total from the downstream layer.
11
12 Returns:
13 A check result with status PASS or FAIL and the difference.
14 """
15 ok = math.isclose(expected, actual, rel_tol=1e-9, abs_tol=1e-6)
16 return {
17 "check_name": name,
18 "status": "PASS" if ok else "FAIL",
19 "severity": "CRITICAL",
20 "expected": f"{expected:.2f}",
21 "actual": f"{actual:.2f}",
22 "difference": actual - expected,
23 }

```

`rel_tol` absorbs rounding in large totals. `abs_tol` covers totals near zero, where any relative tolerance is meaningless.

## A6-L3 exercise 4

**Hints.** (1) What does bronze hold after a redelivery? (2) Silver has one row per key. Which bronze figure matches that? (3) Count distinct respondent IDs within the wave.

**Answer.** The bug is on line 3.

Fixed version:

```

1 def bronze_silver_keys(bronze: DataFrame, silver: DataFrame, wave_id:
 str) -> tuple[int, int]:
2 """Keys in bronze and rows in silver for one wave."""
3 b = bronze.filter(F.col("wave_id") ==
4 wave_id).agg(F.count_distinct("respondent_id")).collect()[0][0]
5 s = silver.filter(F.col("wave_id") == wave_id).count()
6 return b, s

```

Bronze keeps every delivery, so its raw count grows with each redelivery. Count distinct keys for the wave. Within one wave, distinct `respondent_id` is the same as distinct `(respondent_id, wave_id)`.

## Common wrong answers.

- Line 4: Silver has one row per key, so its row count is the right figure. The bronze side is wrong.

## A6-L3 exercise 5

**Hints.** (1) Each layer is checked before the next reads it. (2) Nothing is published before every check passes. (3) Record the evidence, then publish.

**Answer.**

1. MERGE the wave into silver
2. Reconcile bronze distinct keys with silver rows
3. Build gold metrics for the wave
4. Reconcile silver in-scope totals with gold totals
5. Store reconciliation results with the run log
6. Refresh the Power BI semantic model

Load, check, load, check, record, publish. A failed reconciliation stops the sequence before clients see anything.

**Common wrong answers.**

- Reconcile silver before building gold, so a bad silver load never feeds gold.
- Store the results before publishing, so the published wave has a record of passing.
- Publishing before the silver-to-gold reconciliation lets unreconciled numbers reach clients.

## A7-L1 Structured logging with the logging module, not print

---

### A7-L1 exercise 1

**Hints.** (1) What is the default level threshold? (2) Is INFO above or below WARNING? (3) Only the warning appears, as the bare message.

**Answer.**

```
1 3 unexpected labels
```

With no configuration, the effective level is WARNING and Python's last-resort handler prints just the message. Configure a level and format in the entry point to see INFO lines with context.

**Common wrong answers.**

- `loaded 3 unexpected labels`: Without configuration, loggers pass WARNING and above. The INFO line is dropped.

- `WARNING:silver:3 unexpected labels: That format comes from logging.basicConfig()`. Without it, Python's last-resort handler prints the message only.

## A7-L1 exercise 2

**Hints.** (1) The module's own name is a built-in variable. (2) One logger method adds the traceback automatically. (3) A bare statement re-raises the current exception.

**Answer.**

```
1 logger = logging.getLogger(__name__)
2
3 try:
4 rows = do_merge(df)
5 except Exception:
6 logger.exception("merge failed")
7 raise
```

`getLogger(__name__)` names the logger after the module. `logger.exception` logs at ERROR with the traceback. A bare `raise` keeps the original exception and fails the job.

**Common wrong answers.**

- `g1 = "root"`: A hardcoded name loses the module path. `__name__` names the logger after the module automatically.
- `g2 = error`: `error` logs at ERROR but without the traceback, unless you add `exc_info=True`. `exception` includes it.
- `g3 = pass`: `pass` swallows the failure, so the job reports success. Re-raise with a bare `raise`.

## A7-L1 exercise 3

**Hints.** (1) What does the scheduler see when the function returns normally? (2) Returning 0 looks like an empty but successful merge. (3) Replace `return 0` with `raise`.

**Answer.** The bug is on line 8.

Fixed version:

```
1 def merge_wave(df, run_id: str, wave_id: str) -> int:
2 """MERGE one wave. Logs failures."""
3 log = logging.LoggerAdapter(logger, {"run_id": run_id, "wave_id":
4 wave_id})
5 try:
6 return do_merge(df)
7 except Exception:
8 log.exception("merge failed")
```

Returning 0 converts a failure into an apparently successful empty run. Log it, then `raise`, so the job fails and failure notifications fire.

### Common wrong answers.

- Line 7: Logging the exception with its traceback is right. The problem is what happens next.
- Line 5: The MERGE call is fine. Look at how its failure is handled.

### A7-L1 exercise 4

**Hints.** (1) `logging.LoggerAdapter(logger, {...})` attaches context. (2) Branch on `rows == 0`. (3) `log.warning("merge wrote 0 rows")` and `log.info("merge finished rows=%d", rows)`.

### Answer.

```

1 import logging
2
3
4 def wave_logger(run_id: str, wave_id: str) -> logging.LoggerAdapter:
5 """Module logger with run and wave context on every line."""
6 return logging.LoggerAdapter(logging.getLogger("pipelines.silver"),
7 {"run_id": run_id, "wave_id": wave_id})
8
9 def log_merge_result(log: logging.LoggerAdapter, rows: int) -> None:
10 """INFO for a normal merge, WARNING for an empty one."""
11 if rows == 0:
12 log.warning("merge wrote 0 rows")
13 else:
14 log.info("merge finished rows=%d", rows)

```

The adapter puts run and wave IDs on every line without repeating them in each call. A separate WARNING for an empty merge makes the suspicious case stand out.

### A7-L1 exercise 5

**Hints.** (1) What would you want to see on a normal night? (2) What should stand out without stopping the run? (3) Bulky detail only when debugging.

### Answer.

- Normal milestones: INFO.
- A result that is possible but suspicious: WARNING.
- A failure, logged with `log.exception` before re-raising: ERROR.
- Bulky diagnostic detail: DEBUG, switched on only when needed.

Levels let one log serve two audiences: INFO and above for the nightly check, DEBUG for the engineer troubleshooting.

### Common wrong answers.

- zero = `info`: Zero rows usually means an upstream problem. At INFO it disappears among normal lines.
- fail = `warning`: A failed MERGE stops the run. That is ERROR, with the traceback.
- plan = `info`: A full plan on every run floods the log. Keep it at DEBUG.

## A7-L2 A pipeline run log table with an explicit schema

---

### A7-L2 exercise 1

**Hints.** (1) Where does Spark get the column types here? (2) What type can it infer from None alone? (3) It raises.

**Answer.** b) It fails with CANNOT\_DETERMINE\_TYPE. Because: r1) Spark infers each column's type from the data, and a column holding only None gives it nothing to infer from.

With no schema, Spark infers types from values. A column of only None has no type to infer, so it raises CANNOT\_DETERMINE\_TYPE. Declare a StructType.

### Common wrong answers.

- a: Only with an explicit schema. Inference sees only None and cannot choose STRING.
- c: Spark does not drop columns during inference. It raises.
- d: Inference fails before any table is involved.
- b|r2: Right outcome, wrong reason. None is fine when the schema is declared. The problem is inferring a type from nothing but None.

### A7-L2 exercise 2

**Hints.** (1) The module that generates unique ids. (2) The class that provides utc. (3) `datetime.now(timezone.utc)`.

### Answer.

```
1 import uuid
2 from datetime import datetime, timezone
3
4 run_id = str(uuid.uuid4())
5 started_at = datetime.now(timezone.utc)
```

`import uuid` and `from datetime import datetime, timezone` supply every name. `datetime.now(timezone.utc)` gives a timezone-aware UTC timestamp.

### Common wrong answers.

- `g3 = utcnow`: `datetime.utcnow()` is deprecated from Python 3.12 and returns a naive datetime. Use `datetime.now(timezone.utc)`.
- `g2 = date`: You need `timezone`, for `timezone.utc`.

### A7-L2 exercise 3

**Hints.** (1) What does Spark infer from a column that holds only None? (2) Where can you tell Spark the types instead? (3) Build a StructType and pass it to `createDataFrame` with a tuple row.

**Answer.** The bug is on line 3.

Fixed version:

```
1 def run_record(spark: SparkSession, run_id: str, rows: int, status: str,
2 error_message: str | None) -> DataFrame:
3 """One run log row."""
4 schema = T.StructType([
5 T.StructField("run_id", T.StringType(), False),
6 T.StructField("rows_processed", T.LongType(), False),
7 T.StructField("status", T.StringType(), False),
8 T.StructField("error_message", T.StringType(), True),
9])
10 return spark.createDataFrame([(run_id, rows, status, error_message)],
11 schema)
```

Without a schema, the successful case has `error_message=None` only, and inference fails. A `StructType` declares the type, so None is simply NULL.

### Common wrong answers.

- Line 7: Passing None is correct for a success. The problem is how Spark decides the column's type.

### A7-L2 exercise 4

**Hints.** (1) Declare seven StructFields in order, with `error_message` nullable. (2) Build the row as a tuple in the same order. (3)

```
spark.createDataFrame([row], RUN_LOG_SCHEMA).
```

**Answer.**

```
1 from datetime import datetime, timezone
2
3 from pyspark.sql import DataFrame, SparkSession
4 from pyspark.sql import types as T
5
6 RUN_LOG_SCHEMA = T.StructType([
```

```

7 T.StructField("run_id", T.StringType(), False),
8 T.StructField("pipeline_name", T.StringType(), False),
9 T.StructField("rows_processed", T.LongType(), False),
10 T.StructField("status", T.StringType(), False),
11 T.StructField("error_message", T.StringType(), True),
12 T.StructField("started_at", T.TimestampType(), False),
13 T.StructField("finished_at", T.TimestampType(), False),
14])
15
16
17 def run_record(spark: SparkSession, run_id: str, pipeline_name: str,
18 rows_processed: int, status: str, started_at: datetime, error_message:
19 str | None = None) -> DataFrame:
20 """One run log row typed by RUN_LOG_SCHEMA.
21
22 Args:
23 spark: Active session.
24 run_id: Unique id of the run.
25 pipeline_name: Job name.
26 rows_processed: Rows written by the run.
27 status: SUCCESS or FAILED.
28 started_at: UTC start time.
29 error_message: Error text for a failed run.
30
31 Returns:
32 A one-row DataFrame ready to append to the run log table.
33 """
34 finished_at = datetime.now(timezone.utc)
35 row = (run_id, pipeline_name, rows_processed, status, error_message,
36 started_at, finished_at)
37 return spark.createDataFrame([row], RUN_LOG_SCHEMA)

```

The module-level schema makes every record the same shape, whatever its values. Tuples in schema order are the simplest row format.

## A7-L2 exercise 5

**Hints.** (1) The id must exist before any work is logged. (2) try, except, finally is Python's fixed order. (3) finally runs whether the work succeeded or raised.

### Answer.

1. Create `run_id` and `started_at`
2. try: run the layer and read `rows_processed` from the commit metrics
3. Set status `SUCCESS`
4. except: set status `FAILED` and `error_message`, `log.exception`, then raise
5. finally: append the run record to the run log table

`finally` runs after both paths, so every run is recorded with its true status, and the re-raise in `except` still fails the job.

### Common wrong answers.

- The except block belongs before finally in the try statement, and the record must see the FAILED status the except block sets.
- Create the run id first, so log lines during the work carry it.

## A7-L3 Alert on failures, and on the silence that hides them

---

### A7-L3 exercise 1

**Hints.** (1) Does the run fail? (2) When does `on_failure` send? (3) No failure, no email.

**Answer.** b) Nothing. Because: r1) The task returns normally, so the run succeeds and `on_failure` never fires.

Catching and not re-raising turns a failure into a success, so the failure notification never fires. Log, record FAILED, then re-raise.

### Common wrong answers.

- a: `on_failure` fires when the run fails. Swallowing the exception made the run succeed.
- c: There is no warning notification for a caught exception. The run simply succeeds.
- d: Only if `on_success` were configured. It is not.
- b|r2: Right outcome, wrong reason. Notifications work for every task type. The run succeeded because the exception was swallowed.

### A7-L3 exercise 2

**Hints.** (1) The event key names what happened to the run. (2) Send to the on-call group address. (3) `on_failure` and `data-oncall@example.org`.

**Answer.**

```
1 email_notifications:
2 on_failure:
3 - data-oncall@example.org
```

`email_notifications.on_failure` lists the addresses to notify when a run fails. Use a group address so the alert survives holidays and team changes.

### Common wrong answers.

- `g1 = on_error`: The key is `on_failure`. Other keys include `on_start` and `on_success`.



- `g2 = sam.jones@example.org`: Send alerts to a group address. One person's inbox goes quiet when they are away.

### A7-L3 exercise 3

**Hints.** (1) Which values does the table hold? (2) Does 'FAIL' equal 'FAILED'? (3) Use `status = 'FAILED'`.

**Answer.** The bug is on line 3.

Fixed version:

```
1 SELECT pipeline_name, run_id, error_message
2 FROM research_prod.research_governance.pipeline_run_log
3 WHERE status = 'FAILED'
4 AND finished_at >= current_timestamp() - INTERVAL 1 DAY;
```

String comparison is exact, so 'FAIL' never matches 'FAILED'. Typed statuses in the writer (Literal) and a test query against real rows catch this.

#### Common wrong answers.

- Line 4: The one-day window is fine. Compare the status literal with what the table stores.

### A7-L3 exercise 4

**Hints.** (1) Track every pipeline name, starting with no success. (2) Keep the latest SUCCESS finished\_at per pipeline. (3) Stale when there is no success, or the latest is before now - max\_age\_hours.

#### Answer.

```
1 from datetime import datetime, timedelta
2
3
4 def stale_pipelines(runs: list[dict[str, object]], now: datetime,
5 max_age_hours: int = 26) -> list[str]:
6
7 """Pipelines whose latest SUCCESS is older than max_age_hours, or that
8 never succeeded.
9
10 Args:
11 runs: Run log rows with pipeline_name, status and finished_at.
12 now: Current time, timezone-aware.
13 max_age_hours: Largest acceptable age of the latest success.
14
15 Returns:
16 Sorted pipeline names.
17 """
18
19 cutoff = now - timedelta(hours=max_age_hours)
20 latest: dict[str, datetime | None] = {}
21 for r in runs:
```

```

18 name = str(r["pipeline_name"])
19 latest.setdefault(name, None)
20 finished = r["finished_at"]
21 if r["status"] == "SUCCESS" and isinstance(finished, datetime):
22 current = latest[name]
23 if current is None or finished > current:
24 latest[name] = finished
25 return sorted(name for name, last in latest.items() if last is None
 or last < cutoff)

```

Recording every pipeline first means one that never succeeded still appears. Comparing only the latest success stops a single later failure from raising a false alarm.

## A7-L3 exercise 5

**Hints.** (1) Which problems make the job fail? (2) Which problems leave the job successful, or never started? (3) Is every change worth waking someone?

**Answer.**

- Exceptions fail the job: the failure notification handles them, provided nothing swallows them.
- A job that never runs cannot fail: detect the missing SUCCESS with a query.
- A success that did nothing: query for rows\_processed = 0.
- Normal variation is not an alert. Alerts that fire on noise get ignored.

Failures that raise are covered by job notifications. Silence and empty successes need scheduled queries. Noise needs no alert at all.

**Common wrong answers.**

- paused = job: A paused job never runs, so it never fails, and the failure notification never fires.
- zero\_rows = job: The job succeeds, so no failure notification fires. Only a query on the run log sees the zero.
- slow = query: A one-minute change is noise. Alerting on it trains people to ignore alerts.

## A8-L1 Process only the waves you have not processed, correctly

### A8-L1 exercise 1

**Hints.** (1) Which silver waves does the function compare against? (2) Is W14 in silver, for any tracker? (3) Both of A's waves appear in silver.

**Answer.** b) None Because: r1) Tracker B's W14 in silver makes tracker A's W14 look done.

Without a tracker filter on silver, any tracker's wave counts as done. A's W14 never loads, and nothing fails.

### Common wrong answers.

- a: That is the correct answer to the question the function should ask. Because silver is not filtered by tracker, B's W14 hides A's.
- c: W13 is in silver for tracker A, so it is correctly excluded. W14 is wrongly excluded too.
- d: W13 is already loaded for tracker A.
- b|r2: Right outcome, wrong reason. subtract removes only matching values. Here every value matches, because silver mixes trackers.

### A8-L1 exercise 2

**Hints.** (1) The Catalog method that checks whether a table exists. (2) The join type that keeps left rows without a match. (3) `tableExists` and `left_anti`.

### Answer.

```
1 if not spark.catalog.tableExists(silver):
2 return sorted(r.wave_id for r in source.collect())
3 new = source.join(done, "wave_id", "left_anti")
```

`spark.catalog.tableExists` covers the first run. A `left_anti` join keeps bronze waves with no matching silver wave.

### Common wrong answers.

- `g2 = left`: A left join keeps every bronze wave, matched or not. You want the ones with no match: `left_anti`.
- `g2 = left_semi`: `left_semi` keeps the waves that do match: the ones already done. You want the opposite.
- `g1 = exists`: The Catalog method is `tableExists`.

### A8-L1 exercise 3

**Hints.** (1) Compare lines 3 and 4. (2) Which side lacks the tracker filter? (3) Add `.filter(F.col("tracker_id") == tracker_id)` to silver.

**Answer.** The bug is on line 4.

Fixed version:

```
1 def unprocessed_waves(spark: SparkSession, bronze: str, silver: str,
 tracker_id: str) -> List[str]:
```

```

2 """Waves in bronze for this tracker that silver does not hold for this
 tracker."""
3 source = spark.table(bronze).filter(F.col("tracker_id") ==
 tracker_id).select("wave_id").distinct()
4 done = spark.table(silver).filter(F.col("tracker_id") ==
 tracker_id).select("wave_id").distinct()
5 return sorted(r.wave_id for r in source.join(done, "wave_id",
 "left_anti").collect())

```

Silver is read for all trackers, so another tracker's wave hides this tracker's wave. Filter both sides by the same keys.

### Common wrong answers.

- Line 3: Bronze is filtered by tracker, which is right. Compare it with how silver is read.
- Line 5: The left anti join is correct once both sides cover the same tracker.

### A8-L1 exercise 4

**Hints.** (1) Build the bronze waves for the tracker first. (2) Return them all if silver does not exist. (3) Otherwise filter silver the same way and left-anti-join.

### Answer.

```

1 from pyspark.sql import SparkSession
2 from pyspark.sql import functions as F
3
4
5 def unprocessed_waves(spark: SparkSession, bronze: str, silver: str,
 tracker_id: str) -> list[str]:
6 """Wave ids in bronze for this tracker that silver does not hold for
 this tracker.
7
8 Args:
9 spark: Active session.
10 bronze: Bronze table name.
11 silver: Silver table name. May not exist yet.
12 tracker_id: Tracker to scope both sides to.
13
14 Returns:
15 Sorted wave ids to process.
16 """
17 source = spark.table(bronze).filter(F.col("tracker_id") ==
 tracker_id).select("wave_id").distinct()
18 if not spark.catalog.tableExists(silver):
19 return sorted(r.wave_id for r in source.collect())
20 done = spark.table(silver).filter(F.col("tracker_id") ==
 tracker_id).select("wave_id").distinct()
21 return sorted(r.wave_id for r in source.join(done, "wave_id",
 "left_anti").collect())

```

Scoping both sides by tracker, handling the missing table and anti-joining gives the exact set of new waves. Sorting keeps runs deterministic.

### A8-L1 exercise 5

**Hints.** (1) What does the detection compare? (2) Is W14 on both sides? (3) Redeliveries need a signal beyond the wave ID.

**Answer.** b) It is skipped, because W14 is already in silver. Because: r1) The detection compares wave IDs only, and W14 exists on both sides.

Wave-ID detection cannot see redeliveries. Compare the latest ingestion timestamp per wave, or read bronze incrementally with change data feed or Auto Loader.

#### Common wrong answers.

- a: The detection never looks at ingestion\_id. It compares wave IDs, and W14 is in silver.
- c: Nothing runs for W14, so nothing can fail.
- d: W14 is not selected at all.
- b|r2: Right outcome, wrong reason. Bronze appends every redelivery. The detection is what skips it.

## A8-L2 Auto Loader for file ingestion into bronze

---

### A8-L2 exercise 1

**Hints.** (1) Where does a stream keep its progress? (2) It is a separate path, given in the writeStream options. (3) checkpointLocation.

**Answer.** b) The checkpoint location Because: r1) Auto Loader records discovered files in a key-value store in the checkpoint and resumes from it.

The checkpoint holds a RocksDB store of discovered files and the stream's progress. On restart, Auto Loader continues from it, so files are neither skipped nor repeated.

#### Common wrong answers.

- a: The schema location stores the inferred schema, not the list of processed files.
- c: \_rescued\_data holds values that did not fit the schema. It says nothing about which files were loaded.
- d: Auto Loader never reads the target table to decide what to load. The checkpoint does that.

- b|r2: Right answer, wrong reason. Auto Loader leaves source files in place. It records them in the checkpoint.

## A8-L2 exercise 2

**Hints.** (1) The source format has the same prefix as its options. (2) Where the inferred schema is stored. (3) The trigger that processes everything available, then stops.

**Answer.**

```
1 raw = (
2 spark.readStream.format("cloudFiles")
3 .option("cloudFiles.format", "csv")
4 .option("cloudFiles.schemaLocation", schema_path)
5 .load(landing)
6)
7 raw.writeStream.option("checkpointLocation",
8 checkpoint).trigger(availableNow=True).toTable(table)
```

`format("cloudFiles")` selects Auto Loader. `cloudFiles.schemaLocation` stores the schema. `trigger(availableNow=True)` runs the stream as a batch.

**Common wrong answers.**

- `g1 = autoLoader`: The source format is named `cloudFiles`.
- `g3 = once`: `once=True` processes everything in a single micro-batch. `availableNow=True` also stops when done, but can split a large backlog into several batches, which Databricks recommends for Auto Loader.
- `g2 = checkpointLocation`: The checkpoint goes on the writer. The reader option for the schema store is `cloudFiles.schemaLocation`.

## A8-L2 exercise 3

**Hints.** (1) Which mode applies when you do not provide a schema? (2) How does that mode pick up a new column? (3) Fail once, restart with the new schema.

**Answer.** b) The stream stops with `UnknownFieldException` after recording the new schema. The next start includes region. Because: r1) The default mode without a schema is `addNewColumns`, which evolves by failing once and restarting.

Without a schema, the default is `addNewColumns`: the stream stops with `UnknownFieldException`, the schema location gains the column, and a restart continues. Jobs with retries make that restart automatic.

### Common wrong answers.

- a: Nothing is dropped silently. In addNewColumns mode the stream stops so the new column is picked up.
- c: That describes the `rescue` mode. The default without a schema is `addNewColumns`.
- d: Auto Loader never rebuilds the target. It evolves the schema for new files.
- b|r2: Right outcome, wrong reason. It does change the schema: it records the evolved schema, then needs a restart.

### A8-L2 exercise 4

**Hints.** (1) How many streams share the same checkpoint? (2) A checkpoint records one stream's progress. (3) Give each stream its own checkpoint path.

**Answer.** The bug is on line 1.

Fixed version:

```
1 def start_ingest(spark: SparkSession, landing: str, table: str,
 schema_path: str, checkpoint: str) -> StreamingQuery:
2 """Load new files from landing into table. checkpoint must be unique
 to this stream."""
3 raw =
 spark.readStream.format("cloudFiles").option("cloudFiles.format",
 "csv").option("cloudFiles.schemaLocation", schema_path).load(landing)
4 return raw.writeStream.option("checkpointLocation",
 checkpoint).trigger(availableNow=True).toTable(table)
```

Both streams write their progress to the same checkpoint, so each reads the other's state and skips files. One checkpoint per stream, passed in from config.

### Common wrong answers.

- Line 6: The read is fine, with its own schema location. The problem is on the write side.

### A8-L2 exercise 5

**Hints.** (1) Latency or cost? (2) The checkpoint must be durable and private to the stream. (3) Bronze keeps evidence.

**Answer.**

- availableNow on a schedule: incremental processing at batch cost.
- A dedicated checkpoint per stream in a Unity Catalog volume, so it survives and is governed.

- Keep `_rescued_data` in bronze and count it in the contract checks: it is evidence of schema drift.

Scheduled `availableNow`, a durable per-stream checkpoint in a volume, and rescued data counted as a quality signal.

### Common wrong answers.

- `trigger = continuous`: A continuous stream keeps compute running all day for a few files. Scheduled `availableNow` costs far less.
- `checkpoint = tmp`: Temporary storage can be cleared. Losing the checkpoint reloads every file.
- `checkpoint = shared`: Streams must not share a checkpoint. They would skip each other's files.
- `rescued = drop`: Dropping `_rescued_data` in bronze destroys evidence. Bronze keeps everything.

## A8-L3 Change data feed, so downstream layers read only what changed

---

### A8-L3 exercise 1

**Hints.** (1) An update records two versions of the row. (2) Which one comes after the change? (3) `post` = after.

**Answer.** c) `update_postimage` Because: r1) An update records the row before (pre-image) and after (post-image). The post-image holds the new values.

An update writes `update_preimage` (old values) and `update_postimage` (new values). Use the post-image, and drop the pre-image.

### Common wrong answers.

- a: `insert` marks a new row. An update of an existing row produces a pre-image and a post-image.
- b: The pre-image is the row before the update, with the old weight.
- d: `delete` marks a removed row.
- c | r2: Right row type, wrong reason. Updates appear as `update_preimage` and `update_postimage`, not delete plus insert.



## A8-L3 exercise 2

**Hints.** (1) Which versions does `table_changes('c1', 2)` read? (2) How many rows does one updated row produce? (3) Two for the update, one for the delete.

**Answer.**

```
1 3
```

From version 2: `update_preimage` and `update_postimage` for id 1, and `delete` for id 2. Three rows. This was run on delta-spark 4.0.1.

**Common wrong answers.**

- 2: The update produces two rows, pre-image and post-image, plus one delete row: three.
- 5: Version 1's inserts come before the starting version 2, so they are not included.

## A8-L3 exercise 3

**Hints.** (1) The read option names what you are reading. (2) `startingVersion` is inclusive. (3) `readChangeFeed` and `last_processed_version + 1`.

**Answer.**

```
1 changes = (
2 spark.read.option("readChangeFeed", "true")
3 .option("startingVersion", last_processed_version + 1)
4 .table(silver)
5)
```

`readChangeFeed` switches the read to the change feed. Starting at `last_processed_version + 1` reads each change exactly once.

**Common wrong answers.**

- `g2 = last_processed_version`: `startingVersion` is inclusive. Starting at the last processed version reads its changes twice.
- `g1 = enableChangeDataFeed`: `delta.enableChangeDataFeed` is the table property. The read option is `readChangeFeed`.
- `g2 = 0`: Version 0 rereads all history every run, and fails if CDF was enabled later.

## A8-L3 exercise 4

**Hints.** (1) Filter out `update_preimage` first. (2) Window: partition by keys, order by `_commit_version` descending. (3) Keep `row_number` 1, drop the helper column.

## Answer.

```
1 from pyspark.sql import DataFrame, Window
2 from pyspark.sql import functions as F
3
4
5 def latest_state(changes: DataFrame, keys: list[str]) -> DataFrame:
6 """One row per key: its newest change, excluding update pre-images.
7
8 Args:
9 changes: Rows read with readChangeFeed.
10 keys: Columns identifying one row of the source table.
11
12 Returns:
13 The newest insert, update_postimage or delete per key.
14 """
15 window =
16 Window.partitionBy(*keys).orderBy(F.col("_commit_version").desc())
17 return (
18 changes.filter(F.col("_change_type") != "update_preimage")
19 .withColumn("_rank", F.row_number().over(window))
20 .filter(F.col("_rank") == 1)
21 .drop("_rank")
22)
```

Pre-images share a commit version with their post-images, so without the filter a pre-image can win the ranking and bring back old values. Ranking by version keeps the newest change per key.

## A8-L3 exercise 5

**Hints.** (1) Which change types are not deletes? (2) One of them holds the old values. (3) Keep insert and update\_postimage only.

**Answer.** The bug is on line 4.

Fixed version:

```
1 def current_rows(spark: SparkSession, table: str, last_processed_version:
2 int) -> DataFrame:
3 """Inserted and updated rows since the last run, with their new
4 values."""
5 changes = spark.read.option("readChangeFeed",
6 "true").option("startingVersion", last_processed_version +
7 1).table(table)
8 return changes.filter(F.col("_change_type").isin(["insert",
9 "update_postimage"]))
```

Excluding only deletes still keeps update\_preimage, the old version of each updated row. Name the change types you want: insert and update\_postimage.

### Common wrong answers.

- Line 3: The read starts after the last processed version, which is correct.

## A9-L1 pytest basics, and a SparkSession fixture shared by every test

---

### A9-L1 exercise 1

**Hints.** (1) How long does a SparkSession take to start? (2) How often should it start in one test run? (3) Once.

**Answer.** b) session Because: r1) Spark takes seconds to start, and tests can safely share one local session.

`scope="session"` starts Spark once for the whole run. Tests stay independent as long as each creates its own data and tables.

### Common wrong answers.

- a: Function scope rebuilds the fixture for every test: 40 Spark start-ups.
- c: Same cost as function scope, plus duplicated set-up code in every test.
- d: Class scope rebuilds per test class. Session scope builds once.
- b|r2: Right scope, wrong reason. Fixtures of any scope can yield. Session scope wins on start-up cost.

### A9-L1 exercise 2

**Hints.** (1) How many values does `parametrize` receive? (2) Each value becomes its own test. (3) 3 + 1.

**Answer.**

```
1 4
```

`parametrize` generates one test per value, so `pytest` reports `test_wave_loads[W13]`, `[W14]`, `[W15]` and `test_empty_wave_list`: four.

### Common wrong answers.

- 2: Each `parametrize` value is a separate test: three for `test_wave_loads`, plus one.
- 3: Three from the parametrized test, plus `test_empty_wave_list`: four.

### A9-L1 exercise 3

**Hints.** (1) The decorator that declares a fixture. (2) The widest scope. (3) The keyword that pauses a function and resumes it later.

**Answer.**

```

1 @pytest.fixture(scope="session")
2 def spark() -> Iterator[SparkSession]:
3 session = SparkSession.builder.master("local[2]").getOrCreate()
4 yield session
5 session.stop()

```

`@pytest.fixture(scope="session")` builds the fixture once. `yield` gives the session to the tests, and the line after it runs at the end.

### Common wrong answers.

- `g3 = return`: `return` ends the function, so `session.stop()` never runs. `yield` hands over the session and resumes for teardown.
- `g2 = module`: Module scope starts Spark once per test file. Session scope starts it once for the run.

### A9-L1 exercise 4

**Hints.** (1) Pick inputs where the weighted and unweighted answers differ. (2) `[1, 0]` with weights `[3.0, 1.0]`: weighted 0.75, unweighted 0.5. (3) Add `assert weighted_share([1, 0], [0.0, 0.0]) is None`.

### Answer.

```

1 def test_weighted_share() -> None:
2 assert weighted_share([1, 0], [3.0, 1.0]) == 0.75
3 assert weighted_share([1, 0], [0.0, 0.0]) is None

```

A test proves something only if a plausible bug makes it fail. Known answers chosen to separate the right rule from the wrong one, plus the edge case, catch both bugs.

### A9-L1 exercise 5

**Hints.** (1) What is the default fixture scope? (2) How many times does Spark start in 45 tests? (3) `@pytest.fixture(scope="session")`.

**Answer.** The bug is on line 5.

Fixed version:

```

1 import pytest
2 from pyspark.sql import SparkSession
3
4
5 @pytest.fixture(scope="session")
6 def spark():
7 session = SparkSession.builder.master("local[2]").getOrCreate()
8 yield session
9 session.stop()

```

The default function scope starts and stops Spark for every test, at several seconds each. Session scope starts it once.

### Common wrong answers.

- Line 9: Stopping the session in teardown is right. The problem is how often the fixture runs.
- Line 7: The builder is fine. How many times is it called in the run?

## A9-L2 A fast local SparkSession and readable test data

---

### A9-L2 exercise 1

**Hints.** (1) The default suits clusters, not tests. (2) It is a round number in the hundreds. (3) 200.

#### Answer.

```
1 200
```

The default is 200, which suits large data on a cluster. On test data it creates hundreds of empty tasks per shuffle. Set it to 2.

#### Common wrong answers.

- 2: Two is what you should set for tests. The default is 200.
- auto: Databricks can manage shuffle partitions automatically, but open-source Spark's default setting is 200.

### A9-L2 exercise 2

**Hints.** (1) The local master takes the thread count in brackets. (2) Match the shuffle partitions to the threads. (3) The helper from the delta package ends with `_pip`.

#### Answer.

```
1 builder = (
2 SparkSession.builder.master("local[2]")
3 .config("spark.sql.shuffle.partitions", "2")
4 .config("spark.sql.extensions",
5 "io.delta.sql.DeltaSparkSessionExtension")
6 .config("spark.sql.catalog.spark_catalog",
7 "org.apache.spark.sql.delta.catalog.DeltaCatalog")
8)
9 spark = configure_spark_with_delta_pip(builder).getOrCreate()
```

`local[2]` runs two threads in-process, two shuffle partitions avoid empty tasks, and `configure_spark_with_delta_pip` adds the Delta JAR that matches your delta-spark version.

## Common wrong answers.

- `g1 = local[1]`: One thread hides bugs that only appear with several partitions. Use two.
- `g1 = yarn`: yarn needs a cluster. Unit tests run in-process with `local[...]`.
- `g3 = configure_delta`: The delta-spark helper is `configure_spark_with_delta_pip`.

## A9-L2 exercise 3

**Hints.** (1) Build the defaults dict inside the function. (2) Compare the override keys with the default keys. (3) Raise `KeyError` for unknown keys, then update and return.

**Answer.**

```
1 def respondent(**overrides: object) -> dict[str, object]:
2 """A valid silver respondent row. Override only the fields a test is
3 about.
4
5 Raises:
6 KeyError: If an override names an unknown field.
7
8 """
9 row: dict[str, object] = {"respondent_id": "R1", "wave_id": "W1",
10 "final_weight": 1.0, "in_scope": True}
11 unknown = set(overrides) - set(row)
12 if unknown:
13 raise KeyError(f"unknown fields: {sorted(unknown)}")
14 row.update(overrides)
15 return row
```

A fresh dict per call keeps tests independent. Rejecting unknown fields turns a typo in a test into an immediate error instead of a test that checks nothing.

## A9-L2 exercise 4

**Hints.** (1) Column names alone leave Spark to infer types. (2) What can it infer from `[None, None]`? (3) Pass a DDL string with the types.

**Answer.** The bug is on line 4.

Fixed version:

```
1 def weights_df(spark: SparkSession, weights: list[float | None]) ->
2 DataFrame:
3 """One respondent per weight."""
4 rows = [(f"R{i}", w) for i, w in enumerate(weights)]
5 return spark.createDataFrame(rows,
6 "respondent_id STRING, final_weight DOUBLE")
```

A list of names gives Spark no types, so it infers them. With only `None` weights there is nothing to infer from. A DDL schema string fixes the types for every test.

**Common wrong answers.**

- Line 3: Building tuples is fine. The problem is that Spark has to guess the types.

**A9-L2 exercise 5**

**Hints.** (1) Spark runs on the JVM. (2) Tests depend on the fixture and the builders. (3) Infrastructure first, then data helpers, then tests.

**Answer.**

1. Install Java 17 or 21
2. Create a virtual environment and install pyspark, delta-spark and pytest
3. Write a session-scoped spark fixture in `conftest.py`
4. Write test data builders with explicit schemas
5. Write tests with known answers against the builders

Java, then the Python environment, then the shared fixture, then data builders, then the tests that use them.

**Common wrong answers.**

- Either order installs fine, but Spark needs Java before anything runs: check it first so the first test run does not fail.
- Tests need the spark fixture to create DataFrames. Write it before the tests.

## **A9-L3 Compare DataFrames properly, and test that MERGE is idempotent**

---

**A9-L3 exercise 1**

**Hints.** (1) What is the default for `checkRowOrder`? (2) How are floats compared? (3) Both differences are tolerated.

**Answer.** a) Passes Because: r1) It ignores row order by default and compares floats with a relative tolerance.

`assertDataFrameEqual` ignores row order unless you set `checkRowOrder=True`, and compares floats within `rtol=1e-5`. Tested on PySpark 4.0.4.

### Common wrong answers.

- `b`: `checkRowOrder` is `False` by default, so order is ignored.
- `c`: Floats compare with `rtol=1e-5` by default, so `0.30000000000000004` equals `0.3`.
- `d`: `checkRowOrder=True` would make it stricter. The default ignores order.
- `a|r2`: Right outcome, wrong reason. It compares rows too. `assertSchemaEqual` compares only schemas.

### A9-L3 exercise 2

**Hints.** (1) Sorting handles the order. (2) What is `0.1 + 0.2` exactly? (3) Exact equality on floats.

#### Answer.

```
1 False
```

Sorting fixes the order but not floating-point rounding. That is why a helper with tolerances beats hand-written comparisons.

### Common wrong answers.

- `True`: Sorting fixes the order, but `0.1 + 0.2` is `0.30000000000000004`, so the tuples differ.

### A9-L3 exercise 3

**Hints.** (1) The module name says what it is for. (2) The parameter name says what it checks. (3) `testing` and `checkRowOrder`.

#### Answer.

```
1 from pyspark.testing import assertDataFrameEqual, assertSchemaEqual
2
3 assertDataFrameEqual(top_regions, expected, checkRowOrder=True)
```

`pyspark.testing` provides the helpers. `checkRowOrder=True` makes order part of the comparison, which matters for ranked output.

### Common wrong answers.

- `g1 = sql`: The test helpers live in `pyspark.testing`.
- `g2 = ordered`: The parameter is `checkRowOrder`.

### A9-L3 exercise 4

**Hints.** (1) `load()`, `snapshot`, `load()` again. (2) The snapshot must be materialised with `collect()`, or it re-reads the changed table. (3) `assertDataFrameEqual(spark.table(table), first)`.

#### Answer.



```

1 from collections.abc import Callable
2
3 from pyspark.sql import SparkSession
4 from pyspark.testing import assertDataFrameEqual
5
6
7 def check_idempotent(spark: SparkSession, load: Callable[[], None],
8 table: str) -> None:
9 """Assert that running load twice leaves table as one run did.
10
11 Raises:
12 AssertionError: If the second run changed the table.
13 """
14 load()
15 first = spark.table(table).collect()
16 load()
17 assertDataFrameEqual(spark.table(table), first)

```

`collect()` freezes the first state. A `DataFrame` would be lazy and re-read the table after the second run, so the comparison would always pass.

`assertDataFrameEqual` then compares without caring about order.

## A9-L3 exercise 5

**Hints.** (1) Does a `groupBy` return rows in a fixed order? (2) `collect()` lists keep that order. (3) Use `assertDataFrameEqual`, which ignores order.

**Answer.** The bug is on line 5.

Fixed version:

```

1 def test_gold_metrics(spark: SparkSession) -> None:
2 silver = spark.createDataFrame([("R1", "W1", "N", 1.0), ("R2", "W1",
3 "S", 1.0)], "respondent_id STRING, wave_id STRING, region STRING,
4 final_weight DOUBLE")
5 expected = spark.createDataFrame([("N", 1.0), ("S", 1.0)], "region
6 STRING, weighted_base DOUBLE")
7 actual = aggregate_base(silver).select("region", "weighted_base")
8 assertDataFrameEqual(actual, expected)

```

Aggregation output has no guaranteed order, so the two lists match only when the partitions happen to line up. `assertDataFrameEqual` compares the rows as a multiset.

## Common wrong answers.

- Line 4: Selecting the columns to compare is fine. The flakiness comes from how they are compared.

## A10-L1 Git folders, one branch per change, and no secrets in the repository

---

### A10-L1 exercise 1

**Hints.** (1) Which path includes review and tests? (2) What should be the only way into production? (3) A merge, then automation.

**Answer.** b) Automation, after a pull request merges into the deployment branch Because: r1) Production changes must pass review and CI, and automation makes that the only path.

Production Git folders hold the deployment branch and are updated only by automation after a merge, so every production change was reviewed and tested.

#### Common wrong answers.

- a: A developer's own testing skips review and CI. Production should only change through a merge.
- c: Direct edits in production drift from Git and bypass every check.
- d: User folders hold work in progress on personal branches. Production needs a stable deployment branch.
- b|r2: Right answer, wrong reason. People can edit Git folders. The rule is that production ones change only through automation.

### A10-L1 exercise 2

**Hints.** (1) The dbutils module named after what it stores. (2) The key names the secret inside the scope. (3) `secrets` and `supplier_api_key`.

#### Answer.

```
1 token = dbutils.secrets.get(scope="research", key="supplier_api_key")
```

`dbutils.secrets.get(scope=..., key=...)` reads the secret at run time, and Databricks redacts it in notebook output.

#### Common wrong answers.

- `g1 = widgets`: Widgets hold job parameters, which appear in the job UI and run history. Secrets come from `dbutils.secrets`.
- `g1 = fs`: `dbutils.fs` reads files. Reading a key from a file puts it back on storage in plain text.

### A10-L1 exercise 3

**Hints.** (1) Which line holds a credential? (2) Where should the credential come from instead? (3) Delete it and take the token as a parameter.

**Answer.** The bug is on line 4.

Fixed version:

```
1 import requests
2
3 SUPPLIER_URL = "https://supplier.example.org/api/deliveries"
4
5
6 def fetch_deliveries(token: str) -> requests.Response:
7 """Fetch the delivery list. token comes from dbutils.secrets.get at
 run time."""
8 return requests.get(SUPPLIER_URL, headers={"Authorization": f"Bearer {token}"}, timeout=30)
```

Line 4 commits a credential. Make the function take `token: str` and have the notebook read it with `dbutils.secrets.get`. Rotate the leaked token, because Git history keeps it.

#### Common wrong answers.

- Line 3: A public API URL is not a secret. It could even move to config, but it is not what blocks the review.
- Line 9: Line 9 uses the token. The leak is where the token is written.

### A10-L1 exercise 4

**Hints.** (1) Nothing reaches main without review. (2) CI runs on the pull request. (3) Deployment is the last step, done by automation.

#### Answer.

1. Create a feature branch from main in your Git folder
2. Change the code and add or update tests
3. Push and open a pull request
4. CI runs lint, type checks and tests
5. A colleague reviews and approves
6. Merge to main; automation deploys to production

Branch, change, pull request, CI, review, merge and automatic deployment. Every production change passes the same gates.

#### Common wrong answers.

- Deploying before review and CI means production runs unchecked code.
- Create the branch first, so your changes never touch main directly.

## A10-L1 exercise 5

**Hints.** (1) What must stay out of logs and history? (2) What changes per environment? (3) What is specific to one machine?

**Answer.**

- Credentials: a secret scope, read with `dbutils.secrets.get`.
- Code: Git, reviewed and tested.
- Environment choice: a job parameter or target setting, so the same code deploys to dev and prod.
- Local environments: ignored by Git. Pin dependencies in `pyproject.toml` instead.

Secrets in a secret scope, code in Git, environment choice as a parameter, local environments ignored.

**Common wrong answers.**

- `token = param`: Job parameters show in the job definition and run history. Credentials belong in a secret scope.
- `token = repo`: Git history keeps every committed secret for good.
- `catalog = repo`: Hardcoding the catalog in code means dev runs write to prod. Pass it in per environment.
- `env = commit`: A virtual environment is machine-specific and large. Commit the pinned dependency list instead.

## A10-L2 Declarative Automation Bundles and parameterised jobs

---

### A10-L2 exercise 1

**Hints.** (1) Was the rename breaking? (2) Did you need to change your files? (3) Only the name changed.

**Answer.** b) The name. The `databricks bundle` CLI and existing configuration keep working unchanged. Because: r1) Databricks describes the rename as non-breaking.

The rename is non-breaking. The CLI, `databricks.yml` and all existing configuration work as before.

**Common wrong answers.**

- a: The CLI command group is still `databricks bundle`.
- c: The configuration file is still `databricks.yml`.

- `d`: Targets are unchanged and remain the way to describe environments.
- `b|r2`: Right answer, wrong reason. It is the same product under a new name.

## A10-L2 exercise 2

**Hints.** (1) The mode values are full words. (2) The prod target sets the same variable as dev. (3) `development`, `production` and `catalog`.

**Answer.**

```

1 targets:
2 dev:
3 mode: development
4 default: true
5 variables:
6 catalog: research_dev
7 prod:
8 mode: production
9 variables:
10 catalog: research_prod

```

`mode: development` prefixes names and pauses schedules. `mode: production` adds production checks. Each target sets the same variable to its own value.

**Common wrong answers.**

- `g1 = dev`: The mode values are `development` and `production`. `dev` is only the target's name.
- `g3 = catalog_name`: The variable is declared as `catalog`. A different name here sets nothing, and `prod` falls back to the default.

## A10-L2 exercise 3

**Hints.** (1) Which value is the same in every target? (2) Where should environment-specific values come from? (3) Use `${var.catalog}`.

**Answer.** The bug is on line 7.

Fixed version:

```

1 resources:
2 jobs:
3 nat_silver_pipeline:
4 name: nat_silver_pipeline
5 parameters:
6 - name: catalog
7 default: ${var.catalog}
8 tasks:
9 - task_key: silver
10 notebook_task:
11 notebook_path: ./notebooks/silver_nat_transform.py

```

A literal default of `research_prod` applies to every target, dev included. `{var.catalog}` takes the value each target sets.

### Common wrong answers.

- Line 11: A relative notebook path inside the bundle is correct.
- Line 6: Declaring a catalog parameter is right. Its default value is the problem.

### A10-L2 exercise 4

**Hints.** (1) Call `widgets.get("catalog")` and `widgets.get("wave_id")`. (2) Check the catalog against an allowed set. (3) Raise `ValueError` for a bad catalog or an empty `wave_id`. Never fill in a default.

### Answer.

```
1 ALLOWED_CATALOGS = {"research_dev", "research_prod"}
2
3
4 def read_job_parameters(widgets: object) -> dict[str, str]:
5 """Read and validate the job parameters passed to this notebook.
6
7 Args:
8 widgets: dbutils.widgets, or any object with get(name).
9
10 Returns:
11 catalog and wave_id.
12
13 Raises:
14 ValueError: If the catalog is not allowed or wave_id is empty.
15 """
16 get = getattr(widgets, "get")
17 catalog = str(get("catalog"))
18 wave_id = str(get("wave_id")).strip()
19 if catalog not in ALLOWED_CATALOGS:
20 raise ValueError(f"catalog {catalog!r} is not one of
21 {sorted(ALLOWED_CATALOGS)}")
22 if not wave_id:
23 raise ValueError("wave_id is empty")
24 return {"catalog": catalog, "wave_id": wave_id}
```

Parameters arrive as strings from outside the code. Validating them at the start turns a mistyped catalog into a clear failure, instead of a run against the wrong environment.

### A10-L2 exercise 5

**Hints.** (1) Catch configuration errors first. (2) You can only run what is deployed. (3) Then look at the results.

### Answer.

1. databricks bundle validate -t dev
2. databricks bundle deploy -t dev
3. databricks bundle run -t dev nat\_silver\_pipeline
4. Check the run log and QA results for the dev run

Validate, deploy, run, then check the evidence. The same sequence runs in CI for prod, without the manual run.

### Common wrong answers.

- Validate first. It catches YAML and reference errors before anything is deployed.
- run starts the deployed job. Deploy the change first, or you run the old version.

## A10-L3 CI: lint, type-check, test and validate on every pull request

---

### A10-L3 exercise 1

**Hints.** (1) Who should own production deployments? (2) What identity is not a person? (3) A service principal.

**Answer.** b) A service principal, authenticated with OAuth M2M or workload identity federation Because: r1) A service principal is not tied to a person, and its access can be scoped to exactly what deployment needs.

A service principal gives automation its own identity with only the rights it needs. Databricks recommends OAuth M2M for it, or workload identity federation with no stored secret.

### Common wrong answers.

- a: A personal token stops working when that person leaves, and it carries all their access.
- c: Pull request authors should not need production rights. The deploy runs after merge, as automation.
- d: Databricks requires authentication for every API call.
- b | r2: Right identity, wrong reason. Personal tokens can technically deploy. The problem is who owns the access.

## A10-L3 exercise 2

**Hints.** (1) OAuth client credentials have two parts. (2) They match the secret names on the right. (3) `CLIENT_ID` and `CLIENT_SECRET`.

**Answer.**

```
1 env:
2 DATABRICKS_HOST: ${vars.DATABRICKS_HOST}
3 DATABRICKS_CLIENT_ID: ${secrets.DATABRICKS_CLIENT_ID}
4 DATABRICKS_CLIENT_SECRET: ${secrets.DATABRICKS_CLIENT_SECRET}
```

`DATABRICKS_CLIENT_ID` and `DATABRICKS_CLIENT_SECRET`, with `DATABRICKS_HOST`, let the CLI authenticate as the service principal.

**Common wrong answers.**

- `g1 = TOKEN`: `DATABRICKS_TOKEN` is for personal access tokens. OAuth M2M uses the client ID and secret.
- `g2 = PASSWORD`: Service principals authenticate with a client secret, not a password.

## A10-L3 exercise 3

**Hints.** (1) Which setting stops a failing step from failing the job? (2) It sits under the `pytest` step. (3) Remove `continue-on-error: true`.

**Answer.** The bug is on line 6.

Fixed version:

```
1 steps:
2 - uses: actions/checkout@v4
3 - run: pip install -e . --group dev
4 - run: ruff check .
5 - run: pytest
```

`continue-on-error: true` makes the step pass whatever `pytest` returns. Remove it. If tests are flaky, fix or quarantine them in code, and keep CI honest.

**Common wrong answers.**

- Line 5: Running `pytest` is right. The line after it changes what a failure means.
- Line 4: The lint step is fine and can fail the job.

## A10-L3 exercise 4

**Hints.** (1) Which checks take seconds? (2) Which starts a JVM? (3) Which talks to the workspace?



**Answer.**

1. `ruff check .` and `ruff format --check .`
2. `mypy`
3. `pytest` with a local `SparkSession`
4. `databricks bundle validate -t dev`

Lint and types take seconds, tests take a minute or two with Spark, and validation needs the workspace. Failing fast saves CI time and gives quicker feedback.

**Common wrong answers.**

- Spark tests take the longest. Run the seconds-long checks first, so a typo fails the job quickly.
- Validation needs workspace credentials and a network call. Run the local checks first.

**A10-L3 exercise 5**

**Hints.** (1) Checks belong before review and merge. (2) Production changes only after a merge. (3) Credentials never live in files.

**Answer.**

- Checks on every pull request and on main.
- Deploy to prod only after a merge, needing the checks.
- Protect main with required checks.
- Authenticate as a service principal: federation first, OAuth M2M second, secrets never in files.

Checks gate every pull request, protection makes them binding, deployment follows a merge, and CI authenticates as a service principal without secrets in the repository.

**Common wrong answers.**

- `trigger_deploy = every_push`: Deploying every branch push puts unreviewed code in production.
- `branch_protection = none`: Without required checks, CI reports problems but cannot stop a merge.
- `auth = pat`: A token in a workflow file is committed to Git, and ties deployment to one person.
- `trigger_checks = nightly`: Nightly checks find problems after they merge. Run them on the pull request.

## B1-L1 One pyproject.toml and a src layout

---

### B1-L1 exercise 1

**Hints.** (1) Which table never ships in the built package? (2) It was added by PEP 735. (3) `[dependency-groups]`.

**Answer.** b) `[dependency-groups]` Because: r1) Dependency groups (PEP 735) are named lists that are not part of the built distribution.

`[dependency-groups]` holds development-only requirements. `pip install --group dev (pip 25.1+)` installs them.

#### Common wrong answers.

- a: `[project] dependencies` install for every user of the package, in production too.
- c: That table configures how pytest runs. It does not install anything.
- d: A comment installs nothing and drifts from reality.
- b|r2: Right table, wrong reason. Groups can hold any development tools: tests, linting, docs, notebooks.

### B1-L1 exercise 2

**Hints.** (1) The runtime key is plural. (2) The development table has a hyphenated name. (3) `dependencies` and `dependency-groups`.

#### Answer.

```
1 [project]
2 name = "churn"
3 requires-python = ">=3.11"
4 dependencies = ["pandas>=3.0,<3.1", "scikit-learn>=1.8,<1.10"]
5
6 [dependency-groups]
7 dev = ["pytest>=9", "ruff>=0.16", "mypy>=2.3"]
```

`[project].dependencies` lists what users need to run the package.

`[dependency-groups]` lists what developers need to work on it.

#### Common wrong answers.

- `g1 = requires`: The runtime list under `[project]` is called `dependencies`.
- `g2 = tool.dev`: `[tool.*]` tables configure tools. Development requirements go in `[dependency-groups]`.
- `g2 = optional-dependencies`: Optional dependencies (extras) ship with the package for users to opt into. Development-only tools belong in `[dependency-groups]`.

### B1-L1 exercise 3

**Hints.** (1) `tomllib.loads(text)` returns a dict. (2) Use `.get` with defaults for the project table and the dependencies list. (3) Keep entries that contain none of `= < > ~ !`.

**Answer.**

```
1 import tomllib
2
3 CONSTRAINT_CHARS = set("=<>~!")
4
5
6 def unpinning(pyproject_text: str) -> list[str]:
7 """Runtime dependencies with no version constraint.
8
9 Args:
10 pyproject_text: Contents of pyproject.toml.
11
12 Returns:
13 Sorted requirement strings without =, <, >, ~ or !.
14 """
15 project = tomllib.loads(pyproject_text).get("project", {})
16 deps = project.get("dependencies", [])
17 return sorted(d for d in deps if not CONSTRAINT_CHARS & set(d))
```

`tomllib` ships with Python 3.11+. Checking characters is enough to spot a missing constraint, and `.get` handles projects that declare no dependencies.

### B1-L1 exercise 4

**Hints.** (1) Why can Python not see `src/churn`? (2) What makes a package importable everywhere in the environment? (3) An editable install.

**Answer.** b) Run `pip install -e .` in the project's virtual environment. Because: r1) The `src` layout is designed so that only the installed package is importable, which makes tests use it like any user would.

`pip install -e .` installs the package in editable mode, so imports resolve to `src/churn` and edits take effect without reinstalling.

**Common wrong answers.**

- a: Path hacks in every file defeat the point of the layout and break when files move.
- c: That removes the protection the layout gives: tests would import the working directory instead of the installed package.
- d: Copies drift. Tests must exercise the real package.
- b|r2: Right fix, wrong reason. `pytest` can import from folders. The `src` layout deliberately hides the package until it is installed.

## B1-L1 exercise 5

**Hints.** (1) What must more than one program import? (2) What should CI run on every change? (3) What must never be cloned?

**Answer.**

- Shared logic in the package, imported by notebook and jobs.
- One-off exploration in notebooks, which import the package.
- Known-answer checks in tests/.
- Real data in governed storage, never in Git. Keep only small anonymised samples in the repository.

Package for shared logic, notebooks for exploration, tests for checks, governed storage for data.

**Common wrong answers.**

- `features = notebook`: The scoring job cannot import a notebook cell, so it would copy it, and the copies would drift.
- `extract = src`: Client data in the repository leaks with every clone. Keep it in governed storage.
- `known_answer = notebook`: A check in a notebook runs only when someone reruns the notebook. Put it in tests/ so CI runs it.

## B1-L2 One virtual environment per project, and python -m pip

---

### B1-L2 exercise 1

**Hints.** (1) How many Pythons might be on your PATH? (2) Which pip does `python -m pip` use? (3) The one belonging to that python.

**Answer.** b) It installs into the interpreter that `python` runs, whatever is first on PATH. Because: r1) `python -m pip` runs pip as a module of that exact interpreter.

`python -m pip` ties the install to the interpreter you will run. Bare `pip` is whichever pip script PATH finds first.

**Common wrong answers.**

- a: Speed is the same. The benefit is installing into the right interpreter.
- c: It installs into the current interpreter's environment, which inside a venv is the project only.
- d: Resolution is identical. Only the target interpreter is guaranteed.

- `bin`: Right answer, wrong reason. Bare pip still works. It may simply belong to a different Python.

## B1-L2 exercise 2

**Hints.** (1) What does `sys.prefix` point to inside a venv? (2) And `sys.base_prefix`? (3) They are different folders.

**Answer.**

```
1 True
```

`sys.prefix` is the virtual environment, `sys.base_prefix` is the Python it was created from. They differ only inside a venv.

**Common wrong answers.**

- `False`: Inside a venv, `sys.prefix` points at the venv folder and `sys.base_prefix` at the base installation, so they differ.

## B1-L2 exercise 3

**Hints.** (1) The standard library module that creates environments. (2) The flag for an editable install. (3) `venv` and `-e`.

**Answer.**

```
1 python -m venv .venv
2 source .venv/bin/activate
3 python -m pip install -e . --group dev
```

`python -m venv .venv` creates the environment. `pip install -e .` installs the project so edits under `src/` take effect immediately.

**Common wrong answers.**

- `g1 = virtualenv`: `virtualenv` is a separate package. The standard library module is `venv`.
- `g2 = -r: -r` reads a requirements file. Editable installs of the project use `-e .`.

## B1-L2 exercise 4

**Hints.** (1) Inside a venv the two prefixes differ. (2) Raise when they are equal. (3) Read them inside the function body.

**Answer.**

```
1 import sys
2
3
4 def require_virtualenv() -> None:
5 """Stop when not running inside a virtual environment.
```

```

6
7 Raises:
8 RuntimeError: If sys.prefix equals sys.base_prefix.
9 """
10 if sys.prefix == sys.base_prefix:
11 raise RuntimeError("activate the project's virtual environment
first")

```

Equal prefixes mean the system or base interpreter is running. Reading them at call time keeps the check correct in tests and in any process that changes them.

## B1-L2 exercise 5

**Hints.** (1) Can one environment hold both constraints? (2) Where will the next developer look? (3) CI must match how each environment is used.

### Answer.

- Two environments, because the constraint is real and tested.
- Two dependency groups in `pyproject.toml`, with the reason written next to the pin.
- Two CI jobs, so each environment is tested exactly as it is used.

When two tools conflict, give each its own pinned environment, record why in `pyproject.toml`, and test both in CI.

### Common wrong answers.

- `envs = one_latest`: PySpark 4.0.4 fails with pandas 3.x: `assertDataFrameEqual` raises `ImportError`. One environment cannot satisfy both.
- `envs = one_old`: That works, but the analysis code then runs on an older pandas than intended. Separate environments keep each part current.
- `ci = one_job`: Installing both groups together recreates the conflict in CI.
- `declare = readme`: Chat messages are lost. Record the pin and its reason in `pyproject.toml`.

## B1-L3 ruff and mypy in daily work

### B1-L3 exercise 1

**Hints.** (1) How long is each list? (2) What does `zip` do when lengths differ? (3) It stops at the shortest.

### Answer.

```
1 [(1, 2.0), (0, 1.0)]
```

`zip` stops at the shortest input, so the third respondent disappears. `ruff`'s B905 asks for `strict=True`, which raises `ValueError` on unequal lengths (Python 3.10+).

### Common wrong answers.

- `ValueError`: Only `zip(..., strict=True)` raises. Plain `zip` stops at the shorter input without a word.
- `[(1, 2.0), (0, 1.0), (1, None)]`: That is `itertools.zip_longest`. `zip` drops the extra element.

## B1-L3 exercise 2

**Hints.** (1) When can share be `None`? (2) Is that case possible in real data? (3) Handle it.

**Answer.** b) Handle the `None` case before the multiplication. Because: `r1` `mypy` found a real path where the value is `None`. Handling it removes the crash.

`mypy` found a reachable crash. Handle `None` explicitly, for example by returning `"n/a"`.

### Common wrong answers.

- a: The `ignore` hides the warning and keeps the crash.
- c: The hint would then lie, and `mypy` would flag the `return None` instead. The crash stays.
- d: Weakening the checker hides this bug and every future one like it.
- b|r2: Right fix, wrong reason. This is a real crash on zero weight, not style.

## B1-L3 exercise 3

**Hints.** (1) The flag that applies safe fixes. (2) The formatter subcommand. (3) The type checker's command.

### Answer.

```
1 ruff check --fix .
2 ruff format .
3 mypy
```

`ruff check --fix .` applies safe fixes, `ruff format .` formats, and `mypy` type-checks with the settings in `pyproject.toml`.

### Common wrong answers.

- `g1 = --unsafe-fixes`: Unsafe fixes can change behaviour. Use `--fix` for the safe ones, and review the rest by hand.

- `g3 = python -m pytest`: Tests matter, but the step after formatting in this routine is the type checker.

### B1-L3 exercise 4

**Hints.** (1) Annotate: `flags: list[int]`, `weights: list[float]`. (2) `zip(flags, weights, strict=True)`. (3) In report, store the share and return 'n/a' when it is None.

**Answer.**

```

1 def weighted_share(flags: list[int], weights: list[float]) -> float |
 None:
2 """Weighted share of flagged rows, or None when the total weight is
 0."""
3 total = sum(weights)
4 if total == 0:
5 return None
6 return sum(w for f, w in zip(flags, weights, strict=True) if f ==
 1) / total
7
8
9 def report(flags: list[int], weights: list[float]) -> str:
10 """Share as a percentage string, or 'n/a' when undefined."""
11 share = weighted_share(flags, weights)
12 if share is None:
13 return "n/a"
14 return f"{share * 100:.1f}%"

```

The types make the None case visible, `strict=True` stops silent truncation, and report handles None on purpose.

### B1-L3 exercise 5

**Hints.** (1) Prefer real type information. (2) When there is none, narrow the exception. (3) Never global.

**Answer.**

- Stubs exist: install them in the dev group, and mypy checks your use of the library.
- No stubs: silence only that module with a targeted override.
- Never switch it off globally, or you lose the check for libraries that do ship types.

Install stubs where they exist, override per module where they do not, and keep the global check on.



### Common wrong answers.

- `stubs_exist = override`: Stubs exist, so installing them gives you real checks for free.
- `many = global`: A global setting hides missing stubs you could have installed, for every library, forever.

## B2-L1 Pure functions, side effects at the edges

---

### B2-L1 exercise 1

**Hints.** (1) Does the result depend on anything except the arguments? (2) Does the function change anything outside itself? (3) Only one answers no to both.

**Answer.** b) `def share(n, total): return n / total if total else None`  
Because: r1) Its result depends only on its arguments, and it changes nothing outside itself.

`share` depends only on `n` and `total` and changes nothing. The others mutate, read the clock, or write a file.

### Common wrong answers.

- `a`: `append` mutates the caller's list, a side effect.
- `c`: It reads the clock, so the same input gives a different result each call.
- `d`: Writing a file is a side effect. It belongs at the edge.
- `b|r2`: Right function, wrong reason. Purity is about inputs and side effects, not length.

### B2-L1 exercise 2

**Hints.** (1) Is `df` a copy or the caller's object? (2) Does column assignment change the object it is called on? (3) It does.

### Answer.

```
1 ['x', 'flag']
```

`df` is the same object as `original`. Assigning a column changes it, so the caller sees `flag` appear. Use `df.assign(flag=...)` to return a new frame.

### Common wrong answers.

- `['x']`: Copy-on-Write protects against changes through derived objects, but `df["flag"] = ...` changes `df` itself, and `df` is the caller's object.

### B2-L1 exercise 3

**Hints.** (1) Add a parameter for the year. (2) Use `df.assign(...)`, which returns a new frame. (3) `return df.assign(age=as_of_year - df["birth_year"])`.

**Answer.**

```
1 import pandas as pd
2
3
4 def add_age(df: pd.DataFrame, as_of_year: int) -> pd.DataFrame:
5 """Return a copy of df with age = as_of_year - birth_year."""
6 return df.assign(age=as_of_year - df["birth_year"])
```

The year is an input, so tests fix it. `assign` returns a new frame, so callers keep their data.

### B2-L1 exercise 4

**Hints.** (1) When is the default list created? (2) Is it shared between calls? (3) Build a fresh list inside the function.

**Answer.** The bug is on line 1.

Fixed version:

```
1 def flag_low_bases(bases: dict[str, float], minimum: float) -> list[str]:
2 """Names of subgroups whose base is below minimum."""
3 return [name for name, base in bases.items() if base < minimum]
```

The default list is created once and shared by every call, a hidden state that makes the function impure. Build the list inside, here with a comprehension.

**Common wrong answers.**

- Line 5: `append` is where the state leaks, but the shared list is created on line 1.

### B2-L1 exercise 5

**Hints.** (1) Does the step touch anything outside Python's memory? (2) Does its result depend only on its arguments? (3) I/O at the edges, logic in the core.

**Answer.**

- Reading and writing, logging and choosing the reference date: edge.
- Features and scoring, given all inputs as arguments: pure core.

A thin edge does I/O. A pure core does the thinking and gets the tests.

### Common wrong answers.

- `read = core`: Reading storage depends on the world outside the function. Keep it at the edge.
- `write = core`: Writing and logging are side effects. Keep them at the edge.
- `features = edge`: Feature logic given its inputs is deterministic. Make it pure so it can be tested and reused.

## B2-L2 Dataclasses as configuration objects

---

### B2-L2 exercise 1

**Hints.** (1) What does `frozen=True` promise? (2) Is the failure silent or loud? (3) The exception is a subclass of `AttributeError`.

**Answer.** b) Raises `dataclasses.FrozenInstanceError`. Because: `r1` `frozen=True` makes the generated `setattr` raise on every field assignment. Frozen dataclasses raise `dataclasses.FrozenInstanceError` on assignment. Build a changed copy with `dataclasses.replace`.

### Common wrong answers.

- a: That is the behaviour without `frozen=True`. Frozen instances reject assignment.
- c: Assignment never returns a copy. Use `dataclasses.replace(cfg, min_base=30)` for that.
- d: It fails loudly, which is the point: settings cannot drift during a run.
- b | r2: Right answer, wrong reason. Without `frozen=True` every field, int or not, accepts assignment.

### B2-L2 exercise 2

**Hints.** (1) Does a dataclass check or convert types? (2) What does the generated `repr` show for a string? (3) It uses `repr()` of each value, so strings keep their quotes.

### Answer.

```
1 Settings(min_base='30')
```

The hint `int` is for your editor and `mypy`. At runtime the dataclass stores `'30'`. Check or convert in `__post_init__`.

### Common wrong answers.

- `Settings(min_base=30)`: Dataclasses do not convert values. The string `'30'` is stored as it is.

- `TypeError`: Type hints are not checked at runtime. The call succeeds and stores a string.

### B2-L2 exercise 3

**Hints.** (1) A plain `= []` raises `ValueError` at class definition. (2) You need a function that builds a fresh list per instance. (3) The keyword is `default_factory`.

**Answer.**

```
1 from dataclasses import dataclass, field
2
3
4 @dataclass
5 class Run:
6 waves: list[str] = field(default_factory=list)
```

`field(default_factory=list)` calls `list()` for each new instance, so no two instances share one list.

**Common wrong answers.**

- `g1 = default: field(default=list)` stores the list class itself, not a new list.

### B2-L2 exercise 4

**Hints.** (1) `ReportConfig(**raw)` passes each key as a keyword argument. (2) The generated `__init__` raises `TypeError` on an unknown or missing keyword. (3) Check `isinstance(self.min_base, int)` and `self.min_base > 0` in `__post_init__`.

**Answer.**

```
1 from dataclasses import dataclass
2
3
4 @dataclass(frozen=True)
5 class ReportConfig:
6 """Settings for one reporting run."""
7
8 weight_col: str
9 min_base: int = 50
10
11 def __post_init__(self) -> None:
12 if not isinstance(self.min_base, int) or self.min_base <= 0:
13 raise ValueError(f"min_base must be a positive int, got {self.min_base!r}")
14
15
16 def load_config(raw: dict[str, object]) -> ReportConfig:
17 """Build a ReportConfig from a parsed YAML or JSON mapping.
18
```

```

19 Args:
20 raw: Keys must match ReportConfig fields.
21
22 Returns:
23 A validated, frozen config.
24
25 Raises:
26 TypeError: On an unknown or missing key.
27 ValueError: On an invalid min_base.
28 """
29 return ReportConfig(**raw) # type: ignore[arg-type]

```

Unpacking the dict into the constructor lets the generated `__init__` reject unknown and missing keys. `__post_init__` catches bad values. `frozen=True` stops later edits.

## B2-L2 exercise 5

**Hints.** (1) Which shape rejects a misspelt key? (2) When would you rather see a bad value: now or in hour three? (3) Globals make functions impure (lesson B2-L1).

### Answer.

- A frozen dataclass: named fields, unknown keys rejected, no edits during the run.
- Check once at load time in **post\_init**, so a bad file fails before any work starts.
- Pass the object as a parameter, so tests build their own config.

Build a frozen, validated object at the edge and pass it into the pure core.

### Common wrong answers.

- `shape = dict`: A dict accepts any key, including typos, and any function can change it.
- `shape = globals`: Reading a file at import is a module-level side effect, and every test shares the same values.
- `when = use`: Checks spread across functions fail late, after hours of work, and drift apart.
- `pass = import`: A shared import is a global. Tests cannot give one function a different config.

## B2-L3 When a class earns its place

---

### B2-L3 exercise 1

**Hints.** (1) Which task has two phases? (2) What must survive between those phases? (3) Learned frequencies must be stored between fit and transform.

**Answer.** b) Learn category frequencies on training data and encode test data with them. Because: r1) State learned in one call must be reused in a later call.

Fit then transform needs stored state, the classic case for a class, as in scikit-learn's encoders.

#### Common wrong answers.

- a: One step with no state. A function taking the mapping is enough.
- c: One step with no state. A function is enough.
- d: Settings with no behaviour belong in a dataclass, not a full class.
- b|r2: Right answer, wrong reason. Size does not decide it. State across calls does.

### B2-L3 exercise 2

**Hints.** (1) What is in self.seen after the first call? (2) Which of 2, 3 and 4 is not in it? (3) Only 4.

#### Answer.

```
1 [4]
```

The object keeps state between calls. Reuse it on a new wave and it drops rows the new wave needs.

#### Common wrong answers.

- [2, 3, 4]: self.seen still holds 1, 2 and 3 from the first call, so 2 and 3 are dropped.

### B2-L3 exercise 3

**Hints.** (1) Move self.col into a parameter. (2) Compute the total once and test it for zero. (3) `df.assign(share=0.0)` when the total is zero.

#### Answer.

```
1 import pandas as pd
2
3
4 def add_share(df: pd.DataFrame, col: str) -> pd.DataFrame:
```

```

5 """Return df with share = col / column total, or 0.0 when the total is
 zero."""
6 total = df[col].sum()
7 if total == 0:
8 return df.assign(share=0.0)
9 return df.assign(share=df[col] / total)

```

The class held one value and did one thing. A function with two parameters is shorter, and the zero guard stops a NaN column.

## B2-L3 exercise 4

**Hints.** (1) Which state should transform read? (2) Does transform change that state? (3) Refitting inside transform discards what fit learned.

**Answer.** The bug is on line 11.

Fixed version:

```

1 class Scaler:
2 """Standardise with the mean and sd learned in fit."""
3
4 def fit(self, values: list[float]) -> "Scaler":
5 n = len(values)
6 self.mean_ = sum(values) / n
7 self.sd_ = (sum((v - self.mean_) ** 2 for v in values) / n) ** 0.5
8 return self
9
10 def transform(self, values: list[float]) -> list[float]:
11 return [(v - self.mean_) / self.sd_ for v in values]

```

Refitting in `transform` throws away the learned state, the only reason the class exists. It also leaks test data into the scaling.

## Common wrong answers.

- Line 6: The mean formula is correct. The problem is when it is recomputed.
- Line 12: The scaling formula is correct. It uses values that were just overwritten on line 11.

## B2-L3 exercise 5

**Hints.** (1) Does it hold data, do something, or both? (2) Do several operations share the same data? (3) Delete the class: what breaks?

**Answer.**

- Recoding is one stateless step: a function.
- Settings with no behaviour: a frozen dataclass.

- A survey design shares the same data across several operations and must keep weights valid: a class.

Match the form to the need: function for a step, dataclass for settings, class for shared state with behaviour.

### Common wrong answers.

- `recode = class`: One step, no state that outlives the call. A class adds construction with no benefit.
- `settings = class`: No behaviour to attach. A frozen dataclass gives fields, equality and immutability for free.
- `design = function`: Every function would need weights, strata and clusters passed again. A class keeps them together and valid.

## B3-L1 Dtypes in pandas 3, nullable and string

### B3-L1 exercise 1

**Hints.** (1) Can NumPy's int64 hold NaN? (2) What does pandas use for NaN by default? (3) A float.

**Answer.** c) float64 Because: r1) NumPy int64 cannot hold a missing value, so pandas stores NaN in a float column.

NaN is a float, so a missing value upcasts integers to float64. Ask for Int64 to keep integers.

### Common wrong answers.

- a: int64 has no way to represent a missing value.
- b: Int64 is used only when you ask for it, for example with `dtype="Int64"` or `convert_dtypes()`.
- d: Numbers with None infer a numeric dtype, not object.
- c|r2: Right answer, wrong reason. `pd.Series([4, 2])` is int64. The None forces the float.

### B3-L1 exercise 2

**Hints.** (1) What dtype does `read_csv` infer for 01234? (2) What does the blank row add? (3) A number with a missing value is float64.

### Answer.

```
1 1234.0
```

`read_csv` reads 01234 as a number and the blank as NaN, so the column is float64 and the leading zero is gone. Pass `dtype={"postcode": "str"}`.



### Common wrong answers.

- 01234: Without a dtype map, read\_csv infers a number. The blank row makes it float64.
- 1234: The blank row forces float64, so the value prints with .0.

### B3-L1 exercise 3

**Hints.** (1) IDs are labels, not quantities. (2) Which integer dtype holds pd.NA? (3) str and Int64.

#### Answer.

```
1 df = pd.read_csv(path, dtype={"respondent_id": "str", "household_size":
 "Int64"})
```

str keeps IDs as written. Int64 is the nullable integer dtype.

### Common wrong answers.

- g2 = int64: Lower-case int64 is NumPy's integer and fails on a blank value. Use capital-I Int64.
- g1 = int: IDs are labels. Reading them as numbers drops leading zeros.

### B3-L1 exercise 4

**Hints.** (1) read\_csv takes a dtype mapping. (2) IDs as str, counts as Int64. (3) pd.read\_csv(io.StringIO(text), dtype=PANEL\_DTYPES).

#### Answer.

```
1 import io
2
3 import pandas as pd
4
5 PANEL_DTYPES: dict[str, str] = {"respondent_id": "str", "household_size":
 "Int64"}
6
7
8 def read_panel(text: str) -> pd.DataFrame:
9 """Read panel CSV text with explicit dtypes.
10
11 Args:
12 text: CSV content with respondent_id and household_size.
13
14 Returns:
15 A frame with str IDs and nullable integer counts.
16 """
17 return pd.read_csv(io.StringIO(text), dtype=PANEL_DTYPES)
```

The dtype map fixes both columns at load time. In a real pipeline it comes from config, not code.

## B3-L1 exercise 5

**Hints.** (1) What does `mean()` do with `<NA>`? (2) What does it do with 0? (3) Filling blanks with 0 adds fake answers.

**Answer.** The bug is on line 7.

Fixed version:

```
1 import pandas as pd
2
3
4 def mean_household(df: pd.DataFrame) -> float:
5 """Mean household size among respondents who answered."""
6 sizes = df["household_size"].astype("Int64")
7 return float(sizes.mean())
```

`mean()` skips `<NA>` by default. Filling with 0 turns non-answers into households of size zero and pulls the mean down.

### Common wrong answers.

- Line 6: `Int64` is right: it keeps blanks as `<NA>`, which `mean()` skips. The problem is what happens next.
- Line 8: `mean()` is fine on its own. It skips `<NA>`, but by then the blanks are zeros.

## B3-L2 Vectorise instead of apply

### B3-L2 exercise 1

**Hints.** (1) Look for the option with no per-row Python code. (2) A vectorised call takes whole columns as arguments. (3) `np.where(cond, a, b)` picks from `a` or `b` element by element in compiled code.

**Answer.** c) `np.where(df["reach_flag"] == 1, df["final_weight"], 0.0)`

Because: r1) It evaluates the condition and the choice on whole arrays in compiled code.

`np.where` takes whole arrays and runs one compiled loop. The other three options call Python code once per row.

### Common wrong answers.

- `a: apply(axis=1)` calls the lambda once per row in Python. It measured about 500 times slower than `np.where` on one million rows.
- `b:` A list comprehension beats `apply`, but it still runs a Python loop over every value.
- `d: iterrows()` builds a Series for each row. It is the slowest option here.

- `c|r2`: Right choice, wrong reason. Length has nothing to do with speed. `np.where` works on whole arrays in compiled code.
- `c|r3`: Right choice, wrong reason. There is no `lambda`, and `np.where` runs on one core. It is fast because it avoids a Python call per row.

## B3-L2 exercise 2

**Hints.** (1) `None` in a float Series becomes `NaN`. (2) What does any comparison with `NaN` return? (3) Every comparison with `NaN` is `False`.

**Answer.**

```
1 [False, True, False]
```

`None` becomes `NaN` in a float64 Series, and `NaN > 1` is `False`. So the missing value is treated as "not above 1". This is the missing-value trap from the concept.

**Common wrong answers.**

- `[False, True, None]`: A comparison on a float Series returns plain booleans. `NaN > 1` evaluates to `False`, not to a missing value.
- `[False, True, nan]`: Comparisons do not propagate `NaN` on float64 data. They return `False`, which is why missing values get silently mislabelled.
- `[True, True, False]`: `0.0 > 1` is `False`. Only 1.5 exceeds 1.

## B3-L2 exercise 3

**Hints.** (1) The function takes a condition, a value for True and a value for False. (2) Its name is the SQL keyword for filtering rows. (3) `np.where(cond, value_if_true, value_if_false)`, with zero as the false value.

**Answer.**

```
1 df["weighted_reach"] = np.where(df["reach_flag"] == 1,
 df["final_weight"], 0.0)
```

`np.where(cond, df["final_weight"], 0.0)` keeps the weight where the flag is 1 and returns 0.0 elsewhere, as one float array.

**Common wrong answers.**

- `g1 = select`: `np.select` takes a list of conditions and a list of choices. For one condition and two outcomes, use `np.where`.
- `g1 = if`: There is no `np.if`. The vectorised if-else in NumPy is `np.where`.
- `g2 = None`: `None` turns the result into an object array. Use `0.0` so the column stays numeric.

## B3-L2 exercise 4

**Hints.** (1) Compute the total weight first, and return early when it is zero. (2) `df['final_weight'].where(df['reach_flag'] == 1, 0.0)` keeps weights for reached rows only. (3) Sum the masked weights, divide by the total, and wrap the result in `float()`.

**Answer.**

```
1 import numpy as np
2 import pandas as pd
3
4
5 def weighted_reach_pct(df: pd.DataFrame) -> float | None:
6 """Weighted share of respondents with reach_flag == 1.
7
8 Args:
9 df: Rows with reach_flag and final_weight.
10
11 Returns:
12 The weighted share, or None when the total weight is zero.
13 """
14 total = df["final_weight"].sum()
15 if total == 0:
16 return None
17 reached = df["final_weight"].where(df["reach_flag"] == 1, 0.0).sum()
18 return float(reached / total)
```

Sum the weights once for the denominator and return `None` when it is zero, which covers the empty frame too. Mask weights by the flag, sum, divide and convert to `float`.

## B3-L2 exercise 5

**Hints.** (1) `pd.cut` turns numbers into labelled bands and keeps missing values missing. (2) Bins `[-np.inf, 35, 55, np.inf]` with `right=False` give `[.., 35)`, `[35, 55)`, `[55, ..)`. (3) Wrap it in `df.assign(age_group=pd.cut(...))` so you return a new frame.

**Answer.**

```
1 import numpy as np
2 import pandas as pd
3
4 AGE_BINS = [-np.inf, 35, 55, np.inf]
5 AGE_LABELS = ["16-34", "35-54", "55+"]
6
7
8 def add_age_group(df: pd.DataFrame) -> pd.DataFrame:
9 """Return a copy of df with an age_group column. Missing ages stay
10 missing."""
11 return df.assign(
```

```

11 age_group=pd.cut(df["age"], bins=AGE_BINS, labels=AGE_LABELS,
12 right=False)

```

`pd.cut` with left-closed bins (`right=False`) matches the original `<` comparisons exactly. It keeps `NaN` as missing, which fixes a bug the original had. `assign` returns a new frame.

### B3-L2 exercise 6

**Hints.** (1) What is `NaN >= 60`? (2) `np.where` has only two outcomes, so a missing value gets one of them. (3) Build the band as a Series, then keep it only where `minutes.notna()`.

**Answer.** The bug is on line 4.

Fixed version:

```

1 def add_viewing_band(df: pd.DataFrame) -> pd.DataFrame:
2 """Band weekly minutes into heavy (60 or more) and light. Missing
 minutes stay missing."""
3 minutes = df["minutes"]
4 band = pd.Series(np.where(minutes >= 60, "heavy", "light"),
5 index=df.index)
6 return df.assign(viewing_band=band.where(minutes.notna()))

```

`NaN >= 60` is `False`, so `np.where` returns `"light"` for missing minutes. Wrap the result in a Series on the same index and apply `.where(minutes.notna())` to restore the missing values.

### Common wrong answers.

- Line 3: Line 3 only selects the column. The missing value is lost where the condition is evaluated.
- Line 5: `assign` is right. It writes whatever `band` holds, and `band` is already wrong.

## B3-L3 Method chaining and Copy-on-Write

### B3-L3 exercise 1

**Hints.** (1) What does `df[mask]` return under Copy-on-Write? (2) Where does the second step write? (3) To a copy that is thrown away.

**Answer.** b) Nothing. pandas emits a `ChainedAssignmentError` warning. Because: r1) The first indexing step returns a new object, so the write lands on a temporary copy.

Two indexing steps mean two objects. The write hits the temporary one. Write in one step with `.loc`, or use `assign`.

### Common wrong answers.

- `a`: Under Copy-on-Write the assignment goes to a temporary copy. Use `df.loc[mask, "weight"] = 0`.
- `c`: `ChainedAssignmentError` is a warning, not an exception. The code keeps running with `df` unchanged.
- `d`: The filter only selects rows. And the write never reaches `df` anyway.
- `b|r2`: Right answer, wrong reason. `df.loc[mask, col] = value` still changes `df`.

### B3-L3 exercise 2

**Hints.** (1) Is `s` a view you can write through? (2) When does Copy-on-Write copy? (3) On the first write to shared data.

#### Answer.

```
1 [1, 2, 3]
```

`s` shares memory with `df` until you write to it. The write triggers a copy, so `df` keeps its values.

### Common wrong answers.

- `[100, 2, 3]`: Under Copy-on-Write, `s` behaves as a separate object. Writing to it copies the data first, so `df` is unchanged.

### B3-L3 exercise 3

**Hints.** (1) Which column names does each step need? (2) Which step depends on which rows are present? (3) Rename, filter, normalise, band.

#### Answer.

1. `.rename(columns={"wt": "weight"})`
2. `.loc[lambda d: d["age"] >= 18]`
3. `.assign(weight=lambda d: d["weight"] / d["weight"].mean())`
4. `.pipe(add_age_band, bins=BINS, labels=LABELS)`

Each step sees the output of the step above. Rename first so later steps use canonical names. A step that aggregates, such as a mean, must follow the filter that defines its population.

### Common wrong answers.

- The normalising step reads the weight column, which does not exist until the rename.
- Normalising before the filter averages over minors too, so adult weights no longer average 1.0.

### B3-L3 exercise 4

**Hints.** (1) Filter with `.loc[lambda d: ...]`. (2) Cap with `Series.clip(upper=cap)`. (3) Join them: `df.loc[...].assign(weight=lambda d: d["weight"].clip(upper=cap))`.

**Answer.**

```
1 import pandas as pd
2
3
4 def cap_weights(df: pd.DataFrame, cap: float) -> pd.DataFrame:
5 """Keep adults and cap their weights at cap."""
6 return df.loc[lambda d: d["age"] >= 18].assign(
7 weight=lambda d: d["weight"].clip(upper=cap)
8)
```

`loc` then `assign` builds a new frame in two steps, with no temporary copies to write to.

### B3-L3 exercise 5

**Hints.** (1) How many indexing steps does line 7 take before the `=`? (2) Where does the value land? (3) Use `.loc` in one step, or `Series.where` inside `assign`.

**Answer.** The bug is on line 7.

Fixed version:

```
1 import pandas as pd
2
3
4 def zero_invalid(df: pd.DataFrame) -> pd.DataFrame:
5 """Set weight to 0 where valid is False."""
6 return df.assign(weight=df["weight"].where(df["valid"], 0.0))
```

`out[mask]["weight"] = 0.0` writes to a temporary copy.

`df.assign(weight=df["weight"].where(df["valid"], 0.0))` keeps weights where valid and puts 0.0 elsewhere.

**Common wrong answers.**

- Line 6: `copy()` is harmless. Under Copy-on-Write it is also cheap until a write.
- Line 8: Returning `out` is fine. `out` never received the zeros.

## B4-L1 Polars expressions and lazy frames

---

### B4-L1 exercise 1

**Hints.** (1) `scan_` versus `read_`. (2) When does a `LazyFrame` run? (3) On `collect()`.

**Answer.** b) Nothing is read yet. It holds a query plan. Because: r1) A `LazyFrame` records steps and runs them only on `collect()`.

`scan_parquet` builds a plan. `collect()` optimises and runs it, pushing the filter to the reader.

#### Common wrong answers.

- a: That is what `pl.read_parquet` followed by filter does. `scan_parquet` is lazy.
- c: Close, but no rows are read until `collect()`. Then only w9 rows are read.
- d: Nothing is written. A `LazyFrame` is a plan in memory.
- b|r2: Right answer, wrong reason. Polars reads Parquet itself.

### B4-L1 exercise 2

**Hints.** (1) Which values are in group w2? (2) What is the sum of no non-null values? (3) Polars returns the additive identity.

**Answer.**

```
1 [3.0, 0.0]
```

The sum of an empty set of values is 0 in Polars. A real zero and "no data" look the same, so report `null_count()` alongside totals.

#### Common wrong answers.

- `[3.0, None]`: Polars sums an all-null group to 0, not null. Check `null_count()` when that matters.
- `[3.0]`: Every group appears in the output, including w2.

### B4-L1 exercise 3

**Hints.** (1) A window over groups that keeps every row. (2) The method that runs a `LazyFrame`. (3) `over` and `collect`.

**Answer.**

```
1 out = (
2 lf.with_columns(share=pl.col("weight") /
3 pl.col("weight").sum().over("wave"))
4 .collect()
5)
```



`.over("wave")` broadcasts the wave total to each row. `collect()` runs the plan.

### Common wrong answers.

- `g1 = group_by`: `group_by` collapses rows. `over` computes the group total and keeps every row.
- `g2 = execute`: A `LazyFrame` runs with `collect()`.

### B4-L1 exercise 4

**Hints.** (1) Aggregate to wave and brand first. (2) Get the wave total with `.sum().over("wave")`. (3) Guard with `pl.when(total > 0).then(...).otherwise(None)`.

### Answer.

```
1 import polars as pl
2
3
4 def brand_share(df: pl.DataFrame) -> pl.DataFrame:
5 """Weighted share of each brand within its wave. Grain: one row per
6 wave and brand.
7
8 Args:
9 df: Columns wave, brand, weight.
10
11 Returns:
12 Columns wave, brand, share. share is null when the wave total is
13 0.
14 """
15 return (
16 df.group_by("wave", "brand")
17 .agg(brand_weight=pl.col("weight").sum())
18 .with_columns(total=pl.col("brand_weight").sum().over("wave"))
19 .with_columns(
20 share=pl.when(pl.col("total") > 0)
21 .then(pl.col("brand_weight") / pl.col("total"))
22 .otherwise(None)
23)
24 .select("wave", "brand", "share")
25 .sort("wave", "brand")
26)
```

Aggregate to the output grain, broadcast the wave total with `over`, and guard the division so a zero total gives null instead of NaN.

### B4-L1 exercise 5

**Hints.** (1) What does `pl.len()` count? (2) Does it count rows with a null score? (3) `mean()` skips nulls.

**Answer.** The bug is on line 8.

Fixed version:

```
1 import polars as pl
2
3
4 def mean_score(df: pl.DataFrame) -> pl.DataFrame:
5 """Mean score per region among respondents who answered."""
6 return (
7 df.group_by("region")
8 .agg(mean_score=pl.col("score").mean())
9 .sort("region")
10)
```

`pl.len()` counts every row, including skipped answers, so the denominator is too big. `pl.col("score").mean()` divides by the non-null count.

### Common wrong answers.

- Line 7: Grouping by region is right. The formula inside `agg` is the problem.
- Line 9: Sorting does not change values.

## B4-L2 DuckDB, SQL over files and DataFrames

---

### B4-L2 exercise 1

**Hints.** (1) Which options build a new SQL string from the value? (2) Which one keeps the SQL fixed? (3) The placeholder.

**Answer.** b) `con.execute("... WHERE region = ?", [region])` Because: r1) The value is sent as data, separate from the SQL text.

A ? placeholder keeps the SQL text fixed and passes the value as data.

### Common wrong answers.

- a: An f-string pastes the value into the SQL. A quote in the value changes the query.
- c: Concatenation pastes the value into the SQL, the same risk as an f-string.
- d: % formatting is string pasting too, not a parameter.
- b|r2: Right answer, wrong reason. Safety comes from sending the value separately.

### B4-L2 exercise 2

**Hints.** (1) How many rows does an aggregate without GROUP BY return? (2) What is SQL sum over no rows? (3) NULL, returned as None inside a tuple.

## Answer.

```
1 (None,)
```

One row, holding NULL. Check `row[0]` is `None` before using the value.

## Common wrong answers.

- `(0.0,)`: SQL sum over zero rows is NULL, which DuckDB returns as `None`.
- `None`: An aggregate without GROUP BY always returns one row. `fetchone()` gives a tuple holding `None`.

## B4-L2 exercise 3

**Hints.** (1) The DuckDB table function for Parquet files. (2) The positional placeholder. (3) `read_parquet` and `?`.

## Answer.

```
1 con.execute(
2 "SELECT count(*) FROM read_parquet('extracts/*.parquet') WHERE wave
3 = ?",
4 [wave],
5).fetchone()
```

`read_parquet` with a glob reads all files as one table. `?` binds the first value in the list.

## Common wrong answers.

- `g2 = '{wave}'`: That is f-string pasting, and the string is not even an f-string here. Use a `?` placeholder.
- `g1 = read_csv`: The files are Parquet. `read_csv` would fail to parse them.

## B4-L2 exercise 4

**Hints.** (1) Replace the f-string with a `?` placeholder and a list of values. (2) What does `sum` return when no rows match? (3) Return `0.0` when `row[0]` is `None`, else `float(row[0])`.

## Answer.

```
1 import duckdb
2
3 REGION_TOTAL_SQL = "SELECT sum(weight) FROM responses WHERE region = ?"
4
5
6 def region_total(con: duckdb.DuckDBPyConnection, region: str) -> float:
7 """Total weight for one region.
8
9 Args:
10 con: Connection holding a responses table.
11 region: Region name, passed as a query parameter.
```

```

12
13 Returns:
14 The weight total, or 0.0 when no rows match.
15 """
16 row = con.execute(REGION_TOTAL_SQL, [region]).fetchone()
17 if row is None or row[0] is None:
18 return 0.0
19 return float(row[0])

```

The placeholder keeps the SQL fixed, so quotes in the value are just characters. The None check turns SQL's NULL into the 0.0 the caller expects.

## B4-L2 exercise 5

**Hints.** (1) Can a placeholder stand for a table name? (2) What stops arbitrary text becoming SQL? (3) Globals make testing hard (lesson B2-L1).

### Answer.

- Values such as the wave go in as parameters.
- Identifiers such as table names cannot be parameters. Check them against an allowlist before building the SQL.
- Pass one connection in, so tests supply an in-memory database.

Parameters for values, an allowlist for identifiers, and a connection passed in.

### Common wrong answers.

- `wave = fstring`: A quote in the wave changes the SQL. Use a parameter.
- `table = param`: Placeholders bind values, not identifiers. `FROM ?` is not valid SQL for a table name.
- `table = raw`: Any text the user types becomes SQL. Check it against the allowed tables first.
- `con = default`: The default connection is a shared global. Tests cannot swap in their own.

## B4-L3 Choosing pandas, Polars, DuckDB or Spark

### B4-L3 exercise 1

**Hints.** (1) Can the whole dataset sit in RAM? (2) Which tools aggregate while reading? (3) Scan-based engines on one machine.

**Answer.** b) DuckDB or Polars lazy, aggregating while reading Because: r1) Both read Parquet in chunks and push filters and aggregation into the scan, so the full data never sits in memory.

DuckDB and Polars lazy read in chunks and aggregate as they go. The result is small.

### Common wrong answers.

- **a**: pandas loads the whole selection into memory first. 40 GB does not fit in 16 GB.
- **c**: A cluster works, but a one-off aggregate on one machine does not need one.
- **d**: Manual chunking re-implements what DuckDB and Polars do for you, with more room for mistakes.
- **b|r2**: Right answer, wrong reason. Age is not the point. Out-of-memory scanning is.

### B4-L3 exercise 2

**Hints.** (1) Where is `respondent_id` after `set_index`? (2) Does Polars have an index? (3) The default is `include_index=False`.

### Answer.

```
1 ['score']
```

The ID lives in the pandas index, and `from_pandas` drops the index by default. Call `reset_index()` first.

### Common wrong answers.

- `['respondent_id', 'score']`: Polars has no index. `from_pandas` drops it unless you pass `include_index=True`.

### B4-L3 exercise 3

**Hints.** (1) DuckDB relations have `.df()`, `.pl()` and `.fetchall()`. (2) Polars converts to pandas with a `to_` method. (3) `pl` and `to_pandas`.

### Answer.

```
1 top_brands = con.sql("SELECT brand, sum(weight) AS total FROM responses
 GROUP BY brand").pl()
2 for_model = top_brands.to_pandas()
```

`.pl()` returns a Polars DataFrame from DuckDB. `.to_pandas()` converts Polars to pandas through Arrow.

### Common wrong answers.

- `g1 = df`: `.df()` returns pandas. The first step asks for Polars.
- `g2 = df`: Polars DataFrames convert with `to_pandas()`.

## B4-L3 exercise 4

**Hints.** (1) How much data exists after line 7? (2) Where should the filter run? (3) Scan, filter, aggregate, collect.

**Answer.** The bug is on line 7.

Fixed version:

```
1 import polars as pl
2
3
4 def wave_mean(path: str, wave: str) -> float:
5 """Mean weight for one wave."""
6 result = (
7 pl.scan_parquet(path)
8 .filter(pl.col("wave") == wave)
9 .select(pl.col("weight").mean())
10 .collect()
11 .item()
12)
13 return float(result)
```

Line 7 reads every wave eagerly, then copies it into pandas. A lazy scan pushes the wave filter into the reader and returns one number.

### Common wrong answers.

- Line 8: The filter is correct, but by then every wave is already in memory twice.

## B4-L3 exercise 5

**Hints.** (1) Where does the data live? (2) How big is it next to memory? (3) Which library does the next step need?

### Answer.

- Small data feeding the Python modelling stack: pandas.
- SQL across many local files: DuckDB.
- Data already in the lakehouse, scheduled: Spark.
- Large local files with a transform pipeline: Polars lazy.

Size against memory, location of the data, and the next library in line decide the engine.

### Common wrong answers.

- model = spark: 50,000 rows fit in memory many times over. Spark adds start-up time and a cluster for no gain.
- nightly = pandas: Pulling lakehouse tables into one pandas process moves the data away from the compute that already sits next to it.

- transform = `pandas`: `pandas` would load most of 25 GB into 32 GB of RAM. A lazy scan keeps memory in check.

## B5-L1 DataFrame schemas with pandera

---

### B5-L1 exercise 1

**Hints.** (1) Without lazy, when does validation stop? (2) What does the exception hold with lazy? (3) `failure_cases`, one row per failure.

**Answer.** b) Every failure is collected and raised together as `SchemaErrors`. Because: r1) You see all problems in one run, with `failure_cases` listing each one.

Without `lazy`, the first failing check raises `SchemaError`. With it, `pandera` runs every check and raises `SchemaErrors` with all failures.

#### Common wrong answers.

- a: Validation runs at once. lazy refers to collecting errors, not delaying them.
- c: `pandera` raises. Dropping rows needs an explicit choice, not a flag on `validate`.
- d: All checks run. lazy changes how failures are reported.
- b|r2: Right answer, wrong reason. lazy runs every check, so it is not faster.

### B5-L1 exercise 2

**Hints.** (1) Which values fail `gt(0)`? (2) Does lazy stop at the first? (3) All failing values are listed.

**Answer.**

```
1 [-2.0, -3.0]
```

Both negative weights fail. `lazy=True` lists both in `failure_cases`.

#### Common wrong answers.

- `[-2.0]`: With `lazy=True` `pandera` reports every failing value, not only the first.
- `[1, 2]`: `failure_case` holds the failing values. The row positions are in the `index` column.

### B5-L1 exercise 3

**Hints.** (1) A Column argument about repeated values. (2) A schema argument about extra columns. (3) `unique` and `strict`.

**Answer.**

```
1 schema = pa.DataFrameSchema(
2 {"respondent_id": pa.Column(str, unique=True), "score":
3 pa.Column(int)},
4 strict=True,
5)
```

`unique=True` enforces the grain. `strict=True` fails on any unexpected column.

**Common wrong answers.**

- `g2 = coerce`: `coerce` converts dtypes. `strict` rejects unknown columns.
- `g1 = nullable`: `nullable` allows nulls. `unique` rejects duplicates.

### B5-L1 exercise 4

**Hints.** (1) Each column is `pa.Column(dtype, checks..., options)`. (2) Columns reject nulls unless `nullable=True`. (3) `strict` and `coerce` go on the `DataFrameSchema`, `lazy` on `validate`.

**Answer.**

```
1 import pandas as pd
2 import pandera.pandas as pa
3
4 SCORES = pa.DataFrameSchema(
5 {
6 "respondent_id": pa.Column(str, unique=True),
7 "nps": pa.Column(int, pa.Check.in_range(0, 10)),
8 "comment": pa.Column(str, nullable=True),
9 },
10 strict=True,
11 coerce=True,
12)
13
14
15 def validate_scores(df: pd.DataFrame) -> pd.DataFrame:
16 """Validate an NPS extract. Grain: one row per respondent.
17
18 Raises:
19 pa.errors.SchemaErrors: With every failing check.
20 """
21 return SCORES.validate(df, lazy=True)
```

The schema names every column, enforces the grain with `unique`, bounds the score, and allows nulls only where they are expected.



## B5-L1 exercise 5

**Hints.** (1) What does `lazy=True` raise? (2) Is `SchemaErrors` a subclass of `SchemaError`? (3) Even if it were, what does the `except` branch return?

**Answer.** The bug is on line 12, 13.

Fixed version:

```
1 import pandas as pd
2 import pandera.pandas as pa
3
4 WEIGHTS = pa.DataFrameSchema({"weight": pa.Column(float,
5 pa.Check.gt(0))}, strict=True)
6
7 def load_weights(path: str) -> pd.DataFrame:
8 """Read and validate weights.
9
10 Raises:
11 pa.errors.SchemaErrors: When any check fails.
12 """
13 return WEIGHTS.validate(pd.read_csv(path), lazy=True)
```

Two faults. `lazy=True` raises `SchemaErrors`, which `except SchemaError` does not catch, so the handler is dead code for value failures. And when it does catch, it returns the unvalidated frame. Let the error raise, so the run stops.

### Common wrong answers.

- Line 4: The schema is sound: positive, non-null weights, no extra columns.
- Line 11: Validating with `lazy=True` is right. What happens when it fails?

## B5-L2 pydantic models for config and input records

### B5-L2 exercise 1

**Hints.** (1) What is the default value of the extra setting? (2) Does `ignore` warn? (3) Change it with `model_config`.

**Answer.** b) The key is ignored silently. Because: r1) The default extra setting is `"ignore"`. Set `extra="forbid"` to reject unknown keys.

Default `extra="ignore"` drops unknown keys. Config models should use `extra="forbid"`.

### Common wrong answers.

- a: Only with `ConfigDict(extra="forbid")`. The default ignores extras.
- c: pydantic does not guess spellings. `minbase` is a different key.
- d: No warning. That is why `extra="forbid"` matters for config files.

- `br2`: Right answer, wrong reason. It checks keys when you set `extra="forbid"`.

## B5-L2 exercise 2

**Hints.** (1) Is "30" accepted in lax mode? (2) Does `model_copy` run validators? (3) It does not.

**Answer.**

```
1 -1
```

Lax mode converts "30" to 30. `model_copy(update=...)` then sets -1 without validation. Rebuild with `model_validate` instead.

**Common wrong answers.**

- 30: `model_copy` applies the update. It just does not validate it.
- `ValidationError`: `model_copy(update=...)` skips validation, so `gt=0` is never checked.

## B5-L2 exercise 3

**Hints.** (1) The extra setting that raises. (2) The config key that blocks assignment. (3) Less than or equal.

**Answer.**

```
1 class Response(BaseModel):
2 model_config = ConfigDict(extra="forbid", frozen=True)
3
4 respondent_id: str
5 nps: int = Field(ge=0, le=10)
```

`extra="forbid"`, `frozen=True` and `le=10` encode the rules in the model.

**Common wrong answers.**

- `g1 = ignore`: ignore is the default and drops unknown keys silently.
- `g3 = lt`: `lt=10` excludes 10. NPS allows 10, so use `le`.

## B5-L2 exercise 4

**Hints.** (1) Bounds go in `Field(gt=...)`. (2) Unknown keys and assignment are controlled in `model_config`. (3) `ConfigDict(extra="forbid", frozen=True)` and `model_validate(raw)`.

**Answer.**

```
1 from pydantic import BaseModel, ConfigDict, Field
2
3
4 class WeightingConfig(BaseModel):
5 """Settings for one weighting run."""
```

```

6
7 model_config = ConfigDict(extra="forbid", frozen=True)
8
9 weight_col: str
10 min_base: int = Field(default=50, gt=0)
11 cap: float = Field(default=5.0, gt=1.0)
12
13
14 def load_config(raw: dict[str, object]) -> WeightingConfig:
15 """Validate a parsed YAML or JSON mapping.
16
17 Raises:
18 pydantic.ValidationError: On any invalid, missing or unknown key.
19 """
20 return WeightingConfig.model_validate(raw)

```

Every rule lives in the model: required fields, bounds, no extras, no changes. `model_validate` reports all failures at once.

## B5-L2 exercise 5

**Hints.** (1) Which pydantic methods validate? (2) Does `model_copy`? (3) Rebuild with `model_validate` from `model_dump()`.

**Answer.** The bug is on line 14.

Fixed version:

```

1 from pydantic import BaseModel, ConfigDict, Field
2
3
4 class Cfg(BaseModel):
5 """Run settings."""
6
7 model_config = ConfigDict(extra="forbid", frozen=True)
8
9 min_base: int = Field(default=50, gt=0)
10
11
12 def with_min_base(cfg: Cfg, value: int) -> Cfg:
13 """Return a validated copy of cfg with a new min_base."""
14 return Cfg.model_validate(**cfg.model_dump(), "min_base": value)

```

`model_copy(update=...)` skips validation. Dump, merge and `model_validate` so the bound applies to the new value.

## Common wrong answers.

- Line 9: The bound is right. It is never checked on the copy.
- Line 7: `frozen=True` is right. `model_copy` builds a new object, so `frozen` does not stop it.

## B5-L3 Validate at the boundaries, with severity

---

### B5-L3 exercise 1

**Hints.** (1) Where does data you do not control arrive? (2) Where do you make promises to someone else? (3) In and out.

**Answer.** b) After reading, and before writing. Because: r1) Those are the boundaries where data changes owner. Tests cover the pure middle.

Validate what you receive and what you deliver. Test the logic in between.

#### Common wrong answers.

- a: Checks inside a tested pure core add noise. The core's inputs are already validated.
- c: A bad input would pass through every transform first, and the error would name the wrong step.
- d: Unit tests check your code on test data. They say nothing about this week's supplier file.
- b|r2: Right answer, wrong reason. The point is ownership boundaries, not speed.

### B5-L3 exercise 2

**Hints.** (1) Is 0.0 truthy? (2) Which branch runs? (3) The fallback.

#### Answer.

```
1 no data
```

0.0 is falsy, so the best possible result is reported as missing data. Test with `is not None`.

#### Common wrong answers.

- `pass: 0.0 is falsy, so if share:` skips the comparison. Use `if share is not None:`.

### B5-L3 exercise 3

**Hints.** (1) Check input before using it. (2) Check output before delivering it. (3) Log before raising.

#### Answer.

1. Read the supplier file
2. Run input checks and collect results
3. Log all results, then raise if any CRITICAL failed
4. Run the pure transforms

5. Run output checks on the result

6. Write the client table

Read, check, enforce, transform, check, write. Each boundary gets its checks before the data moves on.

### Common wrong answers.

- Input checks must run before the transforms, so bad input never reaches them.
- Output checks after writing are too late. The client already has the table.

### B5-L3 exercise 4

**Hints.** (1) Compute n once. The share is None when n is 0. (2) Build one QaResult per check. Use `share is not None` and `share <= max_missing`. (3) enforce filters on `severity == "CRITICAL"` and not passed.

### Answer.

```
1 from dataclasses import dataclass
2 from typing import Literal
3
4 import pandas as pd
5
6 Severity = Literal["CRITICAL", "WARN"]
7
8
9 @dataclass(frozen=True)
10 class QaResult:
11 """One check outcome."""
12
13 check: str
14 severity: Severity
15 passed: bool
16 observed: float | None
17
18
19 def run_checks(df: pd.DataFrame, min_rows: int, max_missing: float) ->
20 list[QaResult]:
21 """Evaluate input checks. Never raises.
22
23 Args:
24 df: Input with a weight column.
25 min_rows: Minimum acceptable row count.
26 max_missing: Largest acceptable share of null weights.
27
28 Returns:
29 One QaResult per check.
30 """
31 n = len(df)
32 share = None if n == 0 else float(df["weight"].isna().sum()) / n
```

```

32 return [
33 QaResult("min_rows", "CRITICAL", n >= min_rows, float(n)),
34 QaResult("missing_weight", "WARN", share is not None and share <=
max_missing, share),
35]
36
37
38 def enforce(results: list[QaResult]) -> None:
39 """Stop the run when any CRITICAL check failed.
40
41 Raises:
42 RuntimeError: Naming each failed CRITICAL check.
43 """
44 failed = [r.check for r in results if r.severity == "CRITICAL" and
not r.passed]
45 if failed:
46 raise RuntimeError(f"critical checks failed: {failed}")

```

`run_checks` measures and never raises, so every result reaches the log.  
`enforce` decides, and only CRITICAL failures stop the run.

## B5-L3 exercise 5

**Hints.** (1) Does the failure make every output wrong? (2) Or does it affect one cell, handled by a known rule? (3) Grain and weights are foundations.

### Answer.

- Duplicate respondents and broken weights make every figure wrong: CRITICAL.
- A small base is a reporting flag, handled by suppression: WARN.
- A rise in missing age is worth a look but does not invalidate the run: WARN.

CRITICAL for failures that invalidate the whole output. WARN for local issues with a known handling rule.

### Common wrong answers.

- `grain = warn`: Duplicates double-count respondents in every total. Nothing downstream is correct.
- `weights_sum = warn`: If weights do not sum to the sample, every weighted figure is off. Stop the run.
- `base = critical`: Small bases are expected in some cells. Flag and suppress them. Do not stop the whole run.

## B6-L1 Seeds and random generators

---

### B6-L1 exercise 1

**Hints.** (1) Which option has no shared global state? (2) Which one can a test control? (3) `default_rng`.

**Answer.** b) `rng = np.random.default_rng(42)`, passed to the functions that need it Because: r1) Each Generator owns its state, so no other code can change the stream it gives your function.

`default_rng(seed)` gives an independent, seeded generator. Pass it in so each function's randomness is explicit.

#### Common wrong answers.

- a: This seeds a global generator that every library shares. Any extra call anywhere shifts your numbers.
- c: The standard library seed does not affect NumPy's generators.
- d: Unseeded runs cannot be reproduced or debugged.
- b|r2: Right answer, wrong reason. `np.random.seed` still exists as legacy API. Its global state is the problem.

### B6-L1 exercise 2

**Hints.** (1) Does the generator's state change after a draw? (2) Is the second draw from the same position? (3) It continues the stream.

#### Answer.

```
1 False
```

The generator advances. Two draws differ, and rerunning the script repeats both.

#### Common wrong answers.

- True: One generator moves on after each draw. Only a fresh generator with the same seed repeats the first draw.

### B6-L1 exercise 3

**Hints.** (1) The function that returns a Generator. (2) The Generator method that makes child generators. (3) `default_rng` and `spawn`.

#### Answer.

```
1 rng = np.random.default_rng(20260925)
2 worker_rngs = rng.spawn(4)
```

`default_rng` builds the parent. `spawn(4)` returns four independent children.

## Common wrong answers.

- `g1 = seed: np.random.seed` sets the global generator and returns `None`.

## B6-L1 exercise 4

**Hints.** (1) `hashlib.sha256(data)` needs bytes:

`respondent_id.encode("utf-8")`. (2) `.hexdigest()` returns hex text. `int(text, 16)` reads it. (3) Finish with `% n_buckets`.

## Answer.

```
1 import hashlib
2
3
4 def bucket(respondent_id: str, n_buckets: int) -> int:
5 """Stable bucket for an ID, the same in every process.
6
7 Args:
8 respondent_id: The ID to bucket.
9 n_buckets: Number of buckets.
10
11 Returns:
12 An integer from 0 to n_buckets - 1.
13 """
14 digest = hashlib.sha256(respondent_id.encode("utf-8")).hexdigest()
15 return int(digest, 16) % n_buckets
```

SHA-256 of the ID's bytes is the same everywhere, so the bucket is too.

`hash()` is salted per process for strings.

## B6-L1 exercise 5

**Hints.** (1) Which generator does `np.random.choice` use? (2) Who else can move that generator? (3) Use a local `default_rng(seed)`.

**Answer.** The bug is on line 6, 9.

Fixed version:

```
1 import numpy as np
2
3
4 def bootstrap_means(values: list[float], n_boot: int, seed: int) ->
 list[float]:
5 """Bootstrap means of values, reproducible from seed."""
6 rng = np.random.default_rng(seed)
7 data = np.asarray(values)
8 return [float(rng.choice(data, size=len(data), replace=True).mean())
9 for _ in range(n_boot)]
```



`np.random.seed` and `np.random.choice` use one global generator that any imported code advances. A local `default_rng(seed)` belongs to this function alone.

### Common wrong answers.

- Line 10: Taking the mean is fine. The draws come from a shared global generator.

## B6-L2 Record the data, config and code behind every result

---

### B6-L2 exercise 1

**Hints.** (1) Which option changes when one value in the file changes? (2) Which stays the same when the file is copied? (3) A content hash.

**Answer.** c) The SHA-256 of the file's bytes Because: r1) Any change to the content changes the hash, and identical content gives the same hash everywhere.

A content hash is tied to the bytes, not to the name or the file system.

### Common wrong answers.

- a: Suppliers resend changed files under the same name.
- b: Copying a file changes its modified time without changing content, and edits can keep it.
- d: A file with changed values and the same number of rows has the same count.
- c|r2: Right answer, wrong reason. Length does not matter. Content identity does.

### B6-L2 exercise 2

**Hints.** (1) What does `sort_keys` do? (2) Which comes first alphabetically? (3) c before m.

### Answer.

```
1 {"cap": 5.0, "min_base": 50}
```

Sorted keys make the text, and so the hash, independent of the order you wrote the config in.

### Common wrong answers.

- `{"min_base": 50, "cap": 5.0}`: `sort_keys=True` orders keys alphabetically: cap before min\_base.

## B6-L2 exercise 3

**Hints.** (1) Hashes work on bytes. (2) Python 3.11 added a hashlib helper for file objects. (3) `rb` and `file_digest`.

**Answer.**

```
1 with path.open("rb") as f:
2 digest = hashlib.file_digest(f, "sha256").hexdigest()
```

Binary mode gives the exact bytes. `hashlib.file_digest` reads them in chunks.

**Common wrong answers.**

- `g1 = r`: Text mode decodes and may change line endings. Hash the raw bytes with `rb`.

## B6-L2 exercise 4

**Hints.** (1) Open the file in binary mode. (2) `hashlib.file_digest(f, "sha256")` or `hashlib.sha256(path.read_bytes())`. (3) For the config, `json.dumps(config, sort_keys=True).encode("utf-8")`.

**Answer.**

```
1 import hashlib
2 import json
3 from pathlib import Path
4
5
6 def file_sha256(path: Path) -> str:
7 """SHA-256 of a file's bytes, read in chunks.
8
9 Args:
10 path: File to hash.
11
12 Returns:
13 The hex digest.
14 """
15 with path.open("rb") as f:
16 return hashlib.file_digest(f, "sha256").hexdigest()
17
18
19 def config_sha256(config: dict[str, object]) -> str:
20 """SHA-256 of a config, independent of key order.
21
22 Args:
23 config: A JSON-serialisable mapping.
24
25 Returns:
26 The hex digest of the sorted JSON text.
27 """
28 text = json.dumps(config, sort_keys=True)
```

Hash content, not names, and serialise configs in a fixed key order before hashing.

## B6-L2 exercise 5

**Hints.** (1) What changes without changing its name? (2) What would you compare between two runs? (3) Store it where the next analyst will look.

### Answer.

- Content hash of an immutable input.
- Hash of the loaded config, so defaults and key order are handled.
- Git commit plus a dirty flag.
- Stored with the output and in the run log, where comparisons are easy.

Record identities, not names: content hashes and commits. Keep them where a diff between runs is one query away.

### Common wrong answers.

- `data = name`: A resent file keeps its name. Only a content hash shows the data changed.
- `config = file_path`: The file at that path changes later, and defaults are not in it.
- `code = notebook`: The notebook changes under the same name. The commit pins the code.
- `store = chat`: Chat is hard to search and compare. Keep manifests with the output.

## B6-L3 Pinned dependencies and lock files

---

### B6-L3 exercise 1

**Hints.** (1) How many versions does a range accept? (2) What about packages you did not list, such as numpy under pandas? (3) Exact versions plus hashes.

**Answer.** a) The exact version of every direct and indirect dependency, with file hashes. Because: r1) Every install then gets identical packages, on every machine.

Ranges declare what you support. The lock pins what you run, down to every indirect dependency.

### Common wrong answers.

- `b`: Neither file is generated from imports. The lock records what the resolver chose.
- `c`: Locks affect installation, not import speed.
- `d`: A range accepts many versions. The lock records exactly one of each, including indirect dependencies.
- `a|r2`: Right answer, wrong reason. The benefit is identical installs, not skipped work.

### B6-L3 exercise 2

**Hints.** (1) Does `uv.lock` include polars? (2) What does `--locked` assert? (3) That the lock would stay unchanged.

#### Answer.

```
1 The lockfile at `uv.lock` needs to be updated, but `--locked` was
 provided. To update the lockfile, run `uv lock`.
```

`--locked` asserts the lock matches `pyproject.toml`. It does not, so the command exits with an error and tells you to run `uv lock`.

### Common wrong answers.

- `Installed polars`: `--locked` forbids changing the lock. The build fails instead of resolving on the runner.

### B6-L3 exercise 3

**Hints.** (1) `pip-compile` has a flag that adds sha256 lines. (2) `pip install` has a flag that enforces them. (3) `generate-hashes` and `require-hashes`.

#### Answer.

```
1 pip-compile --generate-hashes requirements.in
2 python -m pip install --require-hashes -r requirements.txt
```

Hashes turn a pin into a guarantee: a different file with the same version number is rejected.

### Common wrong answers.

- `g1 = upgrade`: `--upgrade` moves pins to newer versions. The task asks for hashes.

### B6-L3 exercise 4

**Hints.** (1) The lock respects the range in `pyproject.toml`. (2) Tests come before the commit. (3) Branch, bound, lock, test, commit.

**Answer.**

1. Create a branch
2. Raise the bound in `pyproject.toml` if the new version is outside the range
3. `uv lock --upgrade-package scikit-learn`
4. Run the test suite and compare key outputs with the last release
5. Commit `pyproject.toml` and `uv.lock` with the reason

Widen the range, re-lock one package, prove it with tests, then commit both files together.

**Common wrong answers.**

- Locking before raising the bound resolves within the old range, so the upgrade you want is not allowed yet.
- Test before you commit, so a broken upgrade never lands.

**B6-L3 exercise 5**

**Hints.** (1) Which file says what you support, and which says what you run? (2) Where do indirect dependencies get pinned? (3) What makes CI fail on a stale lock?

**Answer.**

- Ranges in `pyproject.toml` say what the code supports.
- The committed lock pins every package with hashes.
- `uv sync --locked` installs exactly that, and fails on a stale lock.

Ranges for intent, a lock for fact, and a locked install everywhere the code runs.

**Common wrong answers.**

- `declare = none`: No bounds means any future major version is acceptable, which is rarely true.
- `declare = exact`: Exact pins in `pyproject.toml` miss indirect dependencies and block security fixes. Pin exactly in the lock.
- `lock = freeze`: `freeze` captures whatever is on one laptop, one-off installs included, with no hashes.
- `install = fresh`: Fresh resolution picks up new releases silently. That is how a job breaks with no code change.

## B7-L1 Profile a dataset before you analyse it

---

### B7-L1 exercise 1

**Hints.** (1) Which call reports the maximum? (2) What would the maximum be? (3) 99.

**Answer.** b) `df["score"].describe()` Because: r1) `describe()` reports min and max, so a max of 99 stands out.

`describe()` gives count, mean, min, quartiles and max. A max of 99 on a 1 to 10 scale is a code, not a score.

#### Common wrong answers.

- a: `head()` shows five rows. The 99s are likely elsewhere.
- c: 99 is a valid integer. The dtype looks fine.
- d: The shape counts rows and columns, not values.
- b|r2: Right answer, wrong reason. `describe()` reports values. It does not clean them.

### B7-L1 exercise 2

**Hints.** (1) Does `nunique` count None? (2) Which b does `duplicated()` mark? (3) `keep="first"` is the default.

**Answer.**

```
1 2 1
```

Two distinct non-missing values. One repeat, because the first b is not marked.

#### Common wrong answers.

- 3 1: `nunique()` excludes missing values by default. Only a and b count.
- 2 2: `duplicated()` keeps the first b unmarked. Only the second b is True.

### B7-L1 exercise 3

**Hints.** (1) Which `keep` value marks every copy? (2) The mean of True/False is the share of True. (3) `False` and `mean`.

**Answer.**

```
1 dupes = df.duplicated(subset=["respondent_id", "wave"], keep=False).sum()
2 missing = df.isna().mean()
```

`keep=False` flags all copies. `isna().mean()` gives the missing share per column.

## Common wrong answers.

- `g1 = "first": keep="first"` leaves one copy unmarked. `keep=False` marks all of them.
- `g2 = sum: sum()` gives counts. The mean of a boolean column is the share.

## B7-L1 exercise 4

**Hints.** (1) Loop over columns and build a dict per column. (2)

`pd.api.types.is_numeric_dtype(s)` tells you whether min and max apply. (3)

`pd.DataFrame.from_dict(rows, orient="index")` builds the table.

## Answer.

```
1 import pandas as pd
2
3
4 def profile(df: pd.DataFrame) -> pd.DataFrame:
5
6 """Profile each column: dtype, missing share, distinct values, numeric
7 range.
8
9 Args:
10 df: Any DataFrame.
11
12 Returns:
13 One row per input column.
14 """
15 rows = {}
16 for col in df.columns:
17 s = df[col]
18 numeric = pd.api.types.is_numeric_dtype(s)
19 rows[col] = {
20 "dtype": str(s.dtype),
21 "missing_share": float(s.isna().mean()),
22 "n_unique": int(s.nunique()),
23 "min": float(s.min()) if numeric else None,
24 "max": float(s.max()) if numeric else None,
25 }
26 return pd.DataFrame.from_dict(rows, orient="index")
```

A per-column table is easy to save and to compare with the next delivery's profile.

## B7-L1 exercise 5

**Hints.** (1) Which rows does `uplicated()` mark by default? (2) Is the first b marked? (3) `keep=False`.

**Answer.** The bug is on line 6.

Fixed version:

```

1 import pandas as pd
2
3
4 def rows_with_duplicate_ids(df: pd.DataFrame) -> pd.DataFrame:
5 """Every row whose respondent_id appears more than once."""
6 mask = df.duplicated(subset=["respondent_id"], keep=False)
7 return df[mask]

```

The default `keep="first"` leaves the first copy unmarked. `keep=False` marks both.

### Common wrong answers.

- Line 7: Filtering with the mask is right. The mask misses the first copy.

## B7-L2 Plot before you summarise

---

### B7-L2 exercise 1

**Hints.** (1) What does one number summarise? (2) Remember Anscombe. (3) Shape needs a plot.

**Answer.** c) Very little about its shape. Because: r1) Anscombe's quartet has identical correlations with a line, a curve and an outlier-driven fit.

Correlation measures linear association. Curves, clusters and single points give the same value. Plot first.

### Common wrong answers.

- a: Anscombe's set II is a curve with the same correlation.
- b: Same correlation, different shapes: that is the point of the quartet.
- d: Only some of the sets have an outlier. The number cannot tell you which.
- c|r2: Right answer, wrong reason. Correlation is useful. It does not describe shape.

### B7-L2 exercise 2

**Hints.** (1) These are Anscombe's sets I and III. (2) What correlation does the quartet share? (3) 0.816.

### Answer.

```
1 True
```

Both correlations round to 0.82. Only a plot shows that set III is a perfect line bent by one point.



### Common wrong answers.

- False: Both are 0.82 to two places. Set III is a straight line plus one outlier, and still matches set I.

### B7-L2 exercise 3

**Hints.** (1) The pyplot function that returns (fig, ax). (2) Axes methods that set things start with set\_. (3) subplots and set\_xlabel.

#### Answer.

```
1 fig, ax = plt.subplots()
2 ax.hist(values, bins=30)
3 ax.set_xlabel("Monthly spend (GBP)")
```

plt.subplots() returns a figure and axes. Axes setters are set\_title, set\_xlabel and set\_ylabel.

### Common wrong answers.

- g2 = xlabel: plt.xlabel is the global pyplot function. On an Axes object the method is set\_xlabel.

### B7-L2 exercise 4

**Hints.** (1) Start with fig, ax = plt.subplots(). (2) ax.hist(values, bins=bins), then the three set\_ calls. (3) Return fig.

#### Answer.

```
1 import matplotlib.pyplot as plt
2 from matplotlib.figure import Figure
3
4
5 def plot_distribution(values: list[float], title: str, xlabel: str, bins:
 int = 30) -> Figure:
6 """Histogram of values with labelled axes.
7
8 Args:
9 values: Numbers to plot.
10 title: Figure title.
11 xlabel: X axis label, with units.
12 bins: Number of histogram bins.
13
14 Returns:
15 The figure. The caller shows or saves it.
16 """
17 fig, ax = plt.subplots(figsize=(6, 4))
18 ax.hist(values, bins=bins)
19 ax.set_title(title)
20 ax.set_xlabel(xlabel)
21 ax.set_ylabel("Respondents")
22 return fig
```

Returning a labelled figure makes the plot testable and lets the caller choose to show or save it.

### B7-L2 exercise 5

**Hints.** (1) How many variables does the question involve? (2) Is time part of the question? (3) Missingness is its own variable.

**Answer.**

- Codes and spikes: a histogram.
- A relationship and its outliers: a scatter plot.
- A break in a series: a line over waves.
- Where answers are missing: a bar chart of missing share by column and device.

Match the plot to the question: one variable, a pair, a series or missingness.

**Common wrong answers.**

- codes = scatter: A spike at 999 in one variable shows best as a tall, isolated bar in a histogram.
- driver = hist: A histogram shows one variable. Influential points need both variables: a scatter plot.
- break = hist: A histogram pools the waves. A line over waves shows the step at wave 9.

## B7-L3 Write assumptions down, then turn them into checks

---

### B7-L3 exercise 1

**Hints.** (1) What happens when the analyst leaves? (2) What happens when the supplier changes the code? (3) Written, applied, checked.

**Answer.** b) In the assumptions log, in config for the handling, and in a QA check. Because: r1) Written, applied and checked in code, it survives staff changes and new deliveries.

The log states it, config applies it, and a check detects when it stops being true.

**Common wrong answers.**

- a: A comment explains one line. It does not run, and nothing fails when the coding changes.
- c: Personal notes leave with the person.

- `d`: New analysts and new pipelines do not know. That is how 99s reach a mean.
- `b|r2`: Right answer, wrong reason. The point is surviving change, not compliance.

## B7-L3 exercise 2

**Hints.** (1) Which columns does `df.replace` touch? (2) Is 99 a valid age? (3) It is blanked anyway, and the column becomes float.

**Answer.**

```
1 [nan, 40.0, 30.0]
```

Whole-frame `replace` blanks the 99-year-old too. Recode per column with a mapping of column to codes.

**Common wrong answers.**

- `[99, 40, 30]`: replace on the whole frame changes every column, including a real age of 99.

## B7-L3 exercise 3

**Hints.** (1) Evidence before belief. (2) Belief before code. (3) Handling before the check that tests its result.

**Answer.**

1. Observe the coding in the questionnaire and the profile
2. Write the assumption, evidence and owner in the log
3. Implement the handling in code, driven by config
4. Add a QA check that runs on every delivery
5. Update `last_verified` when the check or the source is reviewed

Observe, write, implement, check, review. Each step leaves a record the next analyst can follow.

**Common wrong answers.**

- Code written before the evidence encodes a guess.
- The check runs on recoded data, so the handling comes first.

## B7-L3 exercise 4

**Hints.** (1) Work column by column: `df[col].isin(vals)` marks the codes. (2) `Series.mask(cond)` sets values to missing where `cond` is True. (3) Wrap it in `df.assign(**{col: ... for col, vals in codes.items()})`.

**Answer.**

```

1 import pandas as pd
2
3
4 def recode_missing(df: pd.DataFrame, codes: dict[str, list[int]]) ->
 pd.DataFrame:
5 """Set listed codes to missing, per column.
6
7 Args:
8 df: Input frame.
9 codes: Column name to the values that mean missing in that column.
10
11 Returns:
12 A new frame. Unlisted columns are unchanged.
13 """
14 return df.assign(**{col: df[col].mask(df[col].isin(vals)) for col,
 vals in codes.items()})

```

Per-column masking changes only what the assumption covers, and `assign` returns a new frame.

## B7-L3 exercise 5

**Hints.** (1) Is it about rows, values or people? (2) Rows: structure. Values: coding. People: scope. (3) Completes are about who is in, not how answers are coded.

### Answer.

- Grain: structure, checked with `duplicated(keep=False)` on the key.
- Region codes: coding, checked with `isin(lookup)`.
- Age limit and completes only: scope, checked with a filter count that must be zero.

Naming the kind of assumption tells you the kind of check to write.

### Common wrong answers.

- `grain = coding`: Grain is about how rows relate to each other. Check it with key uniqueness.
- `region = scope`: Valid region codes are a coding rule: every value must appear in the lookup.
- `adults = coding`: Age is coded correctly. The rule is about who is included: scope.

## B8-L1 Survey weights and the effective sample size

---

### B8-L1 exercise 1

**Hints.** (1) What do unequal weights do to precision? (2) Which number measures the remaining information? (3) Kish's  $n_{\text{eff}}$ .

**Answer.** b) 71 Because: r1) Unequal weights reduce precision, and  $n_{\text{eff}}$  measures the information that remains.

Use  $n_{\text{eff}}$  for standard errors, intervals and tests. Report the raw base alongside.

#### Common wrong answers.

- a: Raw  $n$  overstates precision when weights vary, so intervals come out too narrow.
- c: The weight total depends on the scaling of the weights, not on the information in the sample.
- d: That difference has no statistical meaning.
- b|r2: Right answer, wrong reason.  $n_{\text{eff}}$  is the correct amount of information, not a cautious guess.

### B8-L1 exercise 2

**Hints.** (1) Which respondents answered 1, and what are their weights? (2) Sum of those weights over the total weight. (3)  $5.5 / 8$ .

#### Answer.

```
1 0.6875 0.6
```

Weighted:  $5.5 / 8 = 0.6875$ . Unweighted:  $3 / 5 = 0.6$ .

#### Common wrong answers.

- 0.6 0.6: `np.average` with weights gives  $(0.5 + 1 + 4) / 8 = 0.6875$ .
- 1.1 0.6: That is  $\text{sum}(y * w) / n$ . A weighted mean divides by  $\text{sum}(w) = 8$ .

### B8-L1 exercise 3

**Hints.** (1) Sum of weights, squared, over the sum of squared weights. (2) Design effect = raw  $n / n_{\text{eff}}$ . (3) 2 and `len(w)`.

#### Answer.

```
1 n_eff = w.sum() ** 2 / (w ** 2).sum()
2 deff = len(w) / n_eff
```

$n_{\text{eff}} = (\sum w)^2 / \sum w^2$ .  $\text{deff} = n / n_{\text{eff}}$ . Equal weights give 1.

## Common wrong answers.

- `g2 = w.sum()`: The design effect compares the raw count with `n_eff`:  
`len(w) / n_eff`.

## B8-L1 exercise 4

**Hints.** (1) Drop rows missing either column first. (2) Divide by the weight total, and return None when it is 0. (3) `n_eff = w.sum() ** 2 / (w ** 2).sum()`.

## Answer.

```
1 import pandas as pd
2
3
4 def weighted_share(df: pd.DataFrame, value: str, weight: str) -> float |
 None:
5 """Weighted mean of value over rows with both value and weight
6 present.
7
8 Returns:
9 The weighted mean, or None for no rows or a zero weight total.
10
11 """
12 d = df[[value, weight]].dropna()
13 total = d[weight].sum()
14 if len(d) == 0 or total == 0:
15 return None
16 return float((d[value] * d[weight]).sum() / total)
17
18 def effective_n(weights: pd.Series) -> float | None:
19 """Kish effective sample size, (sum w)^2 / sum(w^2), ignoring missing
20 weights.
21
22 Returns:
23 n_eff, or None when there are no non-zero weights.
24
25 """
26 w = weights.dropna()
27 sum_sq = (w**2).sum()
28 if sum_sq == 0:
29 return None
30 return float(w.sum() ** 2 / sum_sq)
```

Weighted means divide by the weight total. `n_eff` tells you how much information the weighted sample holds.

## B8-L1 exercise 5

**Hints.** (1) Which describes the population? (2) Which measures information? (3) Which tells readers how many answered?

### Answer.

- Chart: the weighted share, which estimates the population.
- Interval:  $n_{\text{eff}}$ , so precision reflects the weights.
- Footnote: the raw base, so readers know how many people answered.
- Suppression:  $n_{\text{eff}}$ , so a heavily weighted small group is not reported as if it were large.

Weighted for estimates,  $n_{\text{eff}}$  for precision and suppression, raw  $n$  for transparency.

### Common wrong answers.

- estimate = `unweighted`: The unweighted share describes the sample, not the population.
- ci = `raw_n`: Raw  $n$  overstates precision when weights vary.
- suppress = `raw_n`: A group of 60 with  $n_{\text{eff}}$  25 carries the information of 25 people. Compare  $n_{\text{eff}}$  with the threshold.

## B8-L2 Confidence intervals and effect sizes

---

### B8-L2 exercise 1

**Hints.** (1) What is  $p(1-p)$  when  $p = 0$ ? (2) What is the half-width then? (3) Zero.

**Answer.** a)  $[0, 0]$  Because: r1) With  $p = 0$  the standard error is 0, so the interval collapses.

Wald claims certainty from 20 people. Wilson gives  $[0, 0.161]$ .

### Common wrong answers.

- b: That is the Wilson interval. Wald has zero width at  $p = 0$ .
- c:  $p(1-p)$  is 0, so the half-width is 0, not 0.05.
- d:  $\text{sqrt}(0)$  is 0. No error, just a misleading interval.
- a|r2: Right answer, wrong reason. Wilson gives a sensible interval for 20. Wald fails at the edge.

### B8-L2 exercise 2

**Hints.** (1) How large is each group? (2) What happens to standard errors as  $n$  grows? (3) Tiny differences become significant.

### Answer.

1 True

$p = 0.025$ . The difference is 0.5 points, Cohen's  $h$  about 0.01. Significant is not the same as important.

### Common wrong answers.

- `False`: With 100,000 per group, a 0.5-point difference gives  $p = 0.025$ . Significant, and still tiny.

### B8-L2 exercise 3

**Hints.** (1)  $\alpha$  is 1 minus the confidence level. (2) The method named after Edwin Wilson. (3) 0.05 and `wilson`.

### Answer.

```
1 low, high = proportion_confint(count=40, nobs=100, alpha=0.05,
 method="wilson")
```

`alpha=0.05, method="wilson"` gives (0.309, 0.498).

### Common wrong answers.

- $g1 = 0.95$ :  $\alpha$  is the error rate. A 95% interval uses  $\alpha=0.05$ .
- $g2 = \text{normal}$ : normal is the Wald interval, which fails near 0 and 1.

### B8-L2 exercise 4

**Hints.** (1) Guard first:  $n \leq 0$  or  $p$  outside  $[0, 1]$  returns `None`. (2) Compute `denom`, `centre` and `half` in that order. (3) Clip with `max(0.0, ...)` and `min(1.0, ...)`.

### Answer.

```
1 import math
2
3
4 def wilson_interval(p: float, n: float, z: float = 1.959964) ->
 tuple[float, float] | None:
5 """Wilson score interval for a share.
6
7 Args:
8 p: Observed share, weighted or not.
9 n: Sample size, or n_eff for a weighted share.
10 z: Normal quantile. 1.959964 gives 95%.
11
12 Returns:
13 (low, high) clipped to [0, 1], or None for invalid input.
14 """
15 if n <= 0 or not 0 <= p <= 1:
16 return None
17 denom = 1 + z**2 / n
18 centre = (p + z**2 / (2 * n)) / denom
19 half = z / denom * math.sqrt(p * (1 - p) / n + z**2 / (4 * n**2))
```



```
20 return (max(0.0, centre - half), min(1.0, centre + half))
```

Wilson shifts the centre towards 0.5 and widens the interval when information is thin, so it never collapses to zero width.

### B8-L2 exercise 5

**Hints.** (1) Is 0.04 bigger than 0.2? (2) What does a reader need to judge size? (3) Set the threshold before looking.

**Answer.**

- Headline: no meaningful change. 0.04 is far below the agreed 0.2.
- Detail: the difference and its interval, next to the threshold, so readers see the size.
- Rule: decide the meaningful size in advance and judge against it.

Size decides meaning. Significance only says the size is not zero.

**Common wrong answers.**

- headline = `sig_up`: True, and misleading. With 90,000 respondents almost any change is significant.
- detail = `p_only`: A p-value says nothing about size. Show the difference and its interval.
- rule = `p_rule`: In large samples this flags trivial changes every quarter.

## B8-L3 Many tests, and the false positives they bring

### B8-L3 exercise 1

**Hints.** (1) What is the chance one test is not significant? (2) And all 20? (3)  $1 - 0.95^{*20}$ .

**Answer.** b) About 64% Because:  $1 - 0.95^{*20} = 0.64$ .

The chance none of 20 fire is  $0.95^{*20} = 0.36$ , so at least one fires with probability 0.64.

**Common wrong answers.**

- a: 5% is the rate for one test. The family has 20 chances.
- c:  $20 \times 5\%$  is the expected number of false positives (1), not a probability.
- d: True nulls still produce  $p < 0.05$  in 5% of tests, by definition.
- b|r2: Right number, shaky reasoning. The exact calculation is 1 minus the chance of no false positives.

## B8-L3 exercise 2

**Hints.** (1) Holm:  $0.01 \times 5$ ,  $0.02 \times 4$ , ... (2) BH: compare the  $k$ -th smallest  $p$  with  $k \times 0.05 / 5$ . (3) BH is less strict.

**Answer.**

```
1 1 4
```

Holm rejects only 0.01 (adjusted 0.05). BH rejects the four  $p$ -values up to 0.04, controlling the share of false discoveries instead.

**Common wrong answers.**

- 4 4: Holm multiplies the smallest  $p$  by 5, the next by 4, and so on:  $0.02 \times 4 = 0.08$  fails.
- 1 1: Benjamini-Hochberg compares  $p_{(k)}$  with  $k/m \times \alpha$ :  $0.04 \leq 4/5 \times 0.05$ , so four pass.

## B8-L3 exercise 3

**Hints.** (1) The `statsmodels` function in `stats.multitest`. (2) The method name is the author's surname. (3) `multipletests` and `holm`.

**Answer.**

```
1 reject, p_adj, _, _ = multipletests(p_values, alpha=0.05, method="holm")
```

`multipletests(..., method="holm")` returns reject flags and adjusted  $p$ -values.

**Common wrong answers.**

- `g2 = fdr_bh`: `fdr_bh` is Benjamini-Hochberg. Holm's method is "holm".

## B8-L3 exercise 4

**Hints.** (1) `order = sorted(range(m), key=lambda i: p_values[i])` gives positions from smallest  $p$ . (2) For the  $k$ -th position, `running = max(running, min(1.0, (m - k) * p))`. (3) Store at `adjusted[i]`, the original position.

**Answer.**

```
1 def holm_adjust(p_values: list[float]) -> list[float]:
2 """Holm step-down adjusted p-values, in input order.
3
4 Args:
5 p_values: Raw p-values for one family of tests.
6
7 Returns:
8 Adjusted p-values, same order as the input.
9 """
10 m = len(p_values)
11 order = sorted(range(m), key=lambda i: p_values[i])
```

```

12 adjusted = [0.0] * m
13 running = 0.0
14 for k, i in enumerate(order):
15 running = max(running, min(1.0, (m - k) * p_values[i]))
16 adjusted[i] = running
17 return adjusted

```

Holm scales the smallest p by m, the next by m - 1, and forces the sequence to be non-decreasing. Keeping positions lets the caller match results to tests.

## B8-L3 exercise 5

**Hints.** (1) Is the goal to avoid any false positive, or a high share of them? (2) How many tests are in the family? (3) One test needs no adjustment.

**Answer.**

- Few decisive tests: Holm, which controls any false positive in the family.
- Screening many cells: Benjamini-Hochberg, which controls the share of false leads.
- A single pre-registered test: no family, so no adjustment.

Match the error rate to the decision: family-wise for confirmation, false discovery rate for screening.

**Common wrong answers.**

- screen = `holm`: Holm across 400 tests is so strict it flags almost nothing. For screening, control the false discovery rate.
- confirm = `none`: Three tests driving one decision form a family. Without adjustment the error rate rises to about 14%.
- single = `holm`: One test is a family of one. Holm leaves its p-value unchanged.

## B9-L1 Power and sample size, before fieldwork

### B9-L1 exercise 1

**Hints.** (1) What does power 0.25 mean? (2) How often would a real effect be missed? (3) Three times in four.

**Answer.** b) Very little. The test was unlikely to detect the effect even if it were real. Because: r1) Power 0.25 means a real effect of that size is missed 75% of the time.

The test could not answer the question. Report it as inconclusive and, next time, size the sample first.

### Common wrong answers.

- a: An underpowered test misses real effects most of the time. No significance is not no effect.
- c: Nothing in the result points to a negative effect.
- d: Power is not a confidence level about the effect being zero.
- b|r2: Right answer, wrong reason. A well-powered non-significant result is informative.

### B9-L1 exercise 2

**Hints.** (1) The effect sits squared in the denominator. (2) Halve the effect. (3)  $(1/0.5)^2$ .

#### Answer.

```
1 4
```

Half the effect, about four times the sample: 1,534 becomes roughly 6,000 per group.

### Common wrong answers.

- 2: n grows with the inverse square of the effect. Half the effect needs about four times the sample.

### B9-L1 exercise 3

**Hints.** (1) The method that solves for whichever argument you leave out. (2) Which input is 0.8? (3) `solve_power` and `power`.

#### Answer.

```
1 h = proportion_effectsize(0.45, 0.40)
2 n = NormalIndPower().solve_power(effect_size=h, alpha=0.05, power=0.8)
```

`solve_power` leaves `nobs1` unset and solves for it: 1,533.05, so 1,534 per group.

### Common wrong answers.

- `g1 = power`: `power()` returns power for a given n. `solve_power` finds the missing quantity, here n.

### B9-L1 exercise 4

**Hints.** (1) Two-sided: `z_a = NormalDist().inv_cdf(1 - alpha / 2)`. (2) Square  $(\text{top} / (\text{p2} - \text{p1}))$ , then multiply by `design_effect`. (3) `math.ceil` last.

#### Answer.

```
1 import math
2 from statistics import NormalDist
```

```

3
4
5 def n_per_group(
6 p1: float, p2: float, alpha: float = 0.05, power: float = 0.8,
7 design_effect: float = 1.0
8) -> int:
9 """Sample per group to detect p1 vs p2 with a two-sided test.
10
11 Args:
12 p1: Baseline share.
13 p2: Share under the minimum detectable effect.
14 alpha: Two-sided significance level.
15 power: Chance of detecting the effect if real.
16 design_effect: Multiplier for weighted or clustered samples.
17
18 Returns:
19 Respondents per group, rounded up.
20
21 Raises:
22 ValueError: If p1 equals p2 or a share is outside (0, 1).
23 """
24 if p1 == p2 or not (0 < p1 < 1 and 0 < p2 < 1):
25 raise ValueError("shares must differ and lie strictly between 0
26 and 1")
27 z_a = NormalDist().inv_cdf(1 - alpha / 2)
28 z_b = NormalDist().inv_cdf(power)
29 p_bar = (p1 + p2) / 2
30 top = z_a * math.sqrt(2 * p_bar * (1 - p_bar)) + z_b * math.sqrt(p1 *
31 (1 - p1) + p2 * (1 - p2))
32 return math.ceil((top / (p2 - p1)) ** 2 * design_effect)

```

The formula turns the four decisions into a number. The design effect scales it for the information weights remove.

## B9-L1 exercise 5

**Hints.** (1) Two groups, 1,534 each, times 1.5. (2) Does a region need less sample to detect the same effect? (3) Decide before fieldwork.

### Answer.

- Overall: about 4,600 needed with the design effect. 4,000 is close but short.
- Regions: each needs the full sample for the same MDE.
- Settle the trade-off with the client before fieldwork, in writing.

Power maths turns a budget conversation into explicit trade-offs, agreed before any data exists.

### Common wrong answers.

- total = plenty: Two cells of 1,534, times a design effect of 1.5, need about 4,600.
- regions = quarter: A quarter of the sample cannot detect the same 5-point lift. Each region is its own test.
- decision = run: Underpowered regional reads produce noise dressed as findings.

## B9-L2 Analyse an A/B test the way it was randomised

---

### B9-L2 exercise 1

**Hints.** (1) Is a 1% gap on 100,000 users chance? (2) Which users went missing? (3) Find the cause first.

**Answer.** b) Stop and find why B lost users before trusting any result. Because: r1) A split that differs from the design means something removed users unevenly, and they may differ in outcome.

SRM means the arms no longer differ only by chance. Find the cause, fix it, and rerun.

### Common wrong answers.

- a: With 100,000 users, a 1% gap is far beyond chance. The missing users may be the story.
- c: Reweighting assumes the missing users resemble the rest. They rarely do.
- d: Random deletion balances counts but not the kind of user lost from B.
- b|r2: Right answer, wrong reason. The concern is a specific cause, not a p-value rule.

### B9-L2 exercise 2

**Hints.** (1) Same proportions, ten times the users. (2) The statistic grows with n for the same share. (3)  $p = 0.23$ , then  $p = 0.00015$ .

**Answer.**

|              |
|--------------|
| 1 True False |
|--------------|

The same imbalance in share becomes decisive as the sample grows. Check SRM on counts, not on how the percentages look.

### Common wrong answers.

- **True True**: The same 50.6% split is chance at 10,000 users ( $p = 0.23$ ) but not at 100,000 ( $p = 0.00015$ ).
- **False False**: At 10,000 users, 60 either side of 5,000 is well within chance:  $p = 0.23$ .

### B9-L2 exercise 3

**Hints.** (1) Validity before results. (2) Units before estimates. (3) Primary before secondary.

#### Answer.

1. Check the split against the design
2. Aggregate to one outcome per randomised user
3. Estimate the primary metric's lift with its interval
4. Compare the interval with the planned minimum effect
5. Report secondary metrics as exploratory

Trust, then units, then the one planned answer, then everything else labelled as exploration.

### Common wrong answers.

- Checking the split after seeing the lift invites excuses. Check it first.
- Estimate on per-user outcomes, so aggregate before estimating.

### B9-L2 exercise 4

**Hints.** (1) Validate inputs first. (2) Expected counts:  $n * \text{share}$  and  $n * (1 - \text{share})$ . (3) Sum  $(\text{observed} - \text{expected})^2 / \text{expected}$ , then  $\text{erfc}(\sqrt{\text{stat} / 2})$ .

#### Answer.

```
1 import math
2
3
4 def srm_p_value(n_a: int, n_b: int, expected_share_a: float = 0.5) ->
 float:
5 """Sample ratio mismatch p-value: chi-square with 1 degree of freedom.
6
7 Args:
8 n_a: Users observed in arm A.
9 n_b: Users observed in arm B.
10 expected_share_a: Planned share of users in A.
11
12 Returns:
13 The p-value. Below 0.001 signals a mismatch to investigate.
14
15 Raises:
```

```

16 ValueError: On negative counts, a zero total or an invalid share.
17 """
18 if n_a < 0 or n_b < 0 or n_a + n_b == 0 or not 0 < expected_share_a <
19 1:
20 raise ValueError("counts must be non-negative with a positive
21 total, share in (0, 1)")
22 n = n_a + n_b
23 e_a, e_b = n * expected_share_a, n * (1 - expected_share_a)
24 stat = (n_a - e_a) ** 2 / e_a + (n_b - e_b) ** 2 / e_b
25 return math.erfc(math.sqrt(stat / 2))

```

The statistic compares observed with planned counts. For 1 degree of freedom the chi-square p-value equals `erfc(sqrt(stat / 2))`, so no SciPy is needed.

## B9-L2 exercise 5

**Hints.** (1) Which users does line 6 drop? (2) Is the same rule applied to A? (3) Keep everyone as assigned.

**Answer.** The bug is on line 6.

Fixed version:

```

1 import pandas as pd
2
3
4 def conversion_by_arm(users: pd.DataFrame) -> dict[str, float]:
5 """Conversion rate per assigned arm. One row per randomised user."""
6 return users.groupby("arm")["converted"].mean().to_dict()

```

Filtering B on `saw_new_page` keeps its most engaged users and drops the rest, while A keeps everyone. The arms are no longer comparable. Analyse everyone as assigned.

## Common wrong answers.

- Line 7: Grouping by arm is right. The rows reaching it are the problem.

## B9-L3 Peeking, novelty and other ways tests mislead

### B9-L3 exercise 1

**Hints.** (1) How many chances does noise get? (2) One look: 5%. Ten looks: more. (3) The simulation gave 18.4%.

**Answer.** b) About 18% Because: r1) Each look is another chance for noise to cross the threshold.

Repeated looks multiply the chances of a false positive. Fix the horizon or use a planned sequential design.



### Common wrong answers.

- `a`: 5% holds for one planned look. Ten looks gave 18.4% in the simulation.
- `c`: Too high for 10 looks. Unlimited looking pushes it higher still.
- `d`: Identical arms still produce  $p < 0.05$  by chance. That is what a 5% level means.
- `b|r2`: Right number, wrong reason. More users do not create significance when there is no effect.

### B9-L3 exercise 2

**Hints.** (1) Half of 3.1 is 1.55. (2) Is 0.4 below that? (3) Yes.

**Answer.**

```
1 novelty suspected
```

A lift that fades fast is often novelty. Report the later weeks, or run longer before deciding.

### Common wrong answers.

- `stable`: 0.4 is less than half of 3.1, so the effect is fading.

### B9-L3 exercise 3

**Hints.** (1) Both, not either. (2) The plan field for duration. (3) `and` and `min_days`.

**Answer.**

```
1 return n_per_arm_now >= plan.n_per_arm and days_run >= plan.min_days
```

Sample size gives power. Duration covers weekly cycles and novelty. The analysis needs both.

### Common wrong answers.

- `g1 = or`: or lets a big day-2 sample through before a full weekly cycle. Both conditions must hold.

### B9-L3 exercise 4

**Hints.** (1) Draw conversions for all users at once, then take cumulative sums along each row. (2) Pick the columns at the end of each look: `np.arange(1, looks + 1) * n_per_look - 1`. (3) A simulation counts if `(np.abs(z) > z_crit).any(axis=1)`.

**Answer.**

```
1 import numpy as np
2
3
```

```

4 def peeking_false_positive_rate(
5 rng: np.random.Generator,
6 n_sims: int,
7 n_per_look: int,
8 looks: int,
9 rate: float,
10 z_crit: float = 1.959964,
11) -> float:
12 """Share of A/A simulations declared significant at any of the looks.
13
14 Args:
15 rng: Seeded generator.
16 n_sims: Number of simulated tests.
17 n_per_look: New users per arm between looks.
18 looks: Number of analyses.
19 rate: True conversion rate in both arms.
20 z_crit: Two-sided threshold for each look.
21
22 Returns:
23 The false positive rate across simulations.
24 """
25 total = looks * n_per_look
26 a = (rng.random((n_sims, total)) < rate).cumsum(axis=1)
27 b = (rng.random((n_sims, total)) < rate).cumsum(axis=1)
28 idx = np.arange(1, looks + 1) * n_per_look - 1
29 n = idx + 1
30 xa, xb = a[:, idx], b[:, idx]
31 pooled = (xa + xb) / (2 * n)
32 se = np.sqrt(pooled * (1 - pooled) * 2 / n)
33 diff = (xb - xa) / n
34 z = np.divide(diff, se, out=np.zeros_like(diff), where=se > 0)
35 return float((np.abs(z) > z_crit).any(axis=1).mean())

```

Cumulative sums give every look's data from one draw. Counting a simulation when any look crosses the threshold shows how peeking inflates false positives.

## B9-L3 exercise 5

**Hints.** (1) What should people see before the answer is ready? (2) How many looks, at what threshold? (3) Who shares the treatment?

**Answer.**

- Daily views of health, not of the answer.
- One planned analysis, or a group-sequential design agreed before launch.
- Randomise by store, because visitors in the same store share prices.

Separate monitoring from deciding, plan the looks, and randomise at the level where treatment is shared.

### Common wrong answers.

- dashboard = `lift`: A daily lift invites stopping on noise. Show health metrics until the horizon.
- looks = `daily`: 28 looks at 5% each inflate the false positive rate far above 5%.
- unit = `visitor`: Visitors in one store see the same shelf price. Randomising visitors mixes arms within a store.

## B10-L1 Target leakage, features that know the answer

---

### B10-L1 exercise 1

**Hints.** (1) When is the prediction made? (2) Which value is measured after it? (3) March, during the month you predict.

**Answer.** c) Number of support calls in March Because: r1) March calls happen after the prediction is made, and may be caused by leaving.

Features must come from before the prediction date. March calls are not available on 1 March.

### Common wrong answers.

- a: Tenure on 1 March is known on 1 March.
- b: February is before the prediction date. Safe to use.
- d: Plan type on 1 March is known at prediction time.
- c|r2: Right answer, wrong reason. February calls are fine. Timing decides.

### B10-L1 exercise 2

**Hints.** (1) Does `days_inactive` separate the classes perfectly? (2) And `contract_months_left`, in reverse? (3) An AUC of 0 is as informative as 1.

**Answer.**

```
1 1.0 1.0
```

Both separate the classes perfectly, one in reverse. A screen must check `max(auc, 1 - auc)`.

### Common wrong answers.

- `1.0 0.0`: b is 0.0, a perfect inverse relationship. `max(b, 1 - b)` turns it into 1.0.

### B10-L1 exercise 3

**Hints.** (1) Everything is judged against the prediction date. (2) Timing, then data-driven checks. (3) Explanation needs a suspect.

**Answer.**

1. Write down the prediction date and what the model will see then
2. Find the date each feature value becomes known
3. Drop features known only after the prediction date
4. Run the single-feature AUC screen on what remains
5. Explain any high-AUC feature with the data owner, or drop it

Define the moment of prediction, check each feature's timing, then use the screen as a second line of defence.

**Common wrong answers.**

- Screening first finds the obvious leaks but misses timing leaks with modest AUC. Start from the prediction date.
- Screen first, then explain what the screen flags.

### B10-L1 exercise 4

**Hints.** (1) `X.select_dtypes("number").columns` limits the loop to numeric features. (2) Filter rows with `X[col].notna()` and check `y[known].nunique() == 2`. (3) Flag when `max(auc, 1 - auc) >= threshold`.

**Answer.**

```
1 import pandas as pd
2 from sklearn.metrics import roc_auc_score
3
4
5 def leakage_screen(X: pd.DataFrame, y: pd.Series, threshold: float =
 0.95) -> list[str]:
6 """Numeric features that alone separate the classes suspiciously well.
7
8 Args:
9 X: Candidate features.
10 y: Binary target aligned with X.
11 threshold: AUC strength at or above which a feature is flagged.
12
13 Returns:
14 Suspect column names, in column order.
15 """
16 suspects = []
17 for col in X.select_dtypes("number").columns:
18 known = X[col].notna()
19 if known.sum() == 0 or y[known].nunique() < 2:
20 continue
```

```

21 auc = roc_auc_score(y[known], X.loc[known, col])
22 if max(auc, 1 - auc) >= threshold:
23 suspects.append(col)
24 return suspects

```

The screen looks for any single feature that nearly separates the classes. Missingness itself can leak, as `late_only` shows: filled with zeros it would look perfect.

## B10-L1 exercise 5

**Hints.** (1) Would you know it at the moment of the transaction? (2) Can it be recomputed from earlier data? (3) Who fills it in, and when?

**Answer.**

- Amount: known at transaction time.
- Chargeback: a consequence of fraud, filed weeks later.
- Transaction count: useful, but must count only transactions before this one.
- Reviewer: assigned only to suspected cases, a proxy for the label.

Keep what is known at prediction time, rebuild what can be made point in time, drop the rest.

**Common wrong answers.**

- `chargeback = keep`: A chargeback is filed after the fraud is found. It is the label in disguise.
- `txn_count = keep`: Computed on the extract date, it includes transactions after this one. Rebuild it point in time.
- `reviewer = keep`: Reviewers are assigned to flagged cases. The feature encodes the labelling process.

## B10-L2 Fit on training data only

---

### B10-L2 exercise 1

**Hints.** (1) Which step computes something from the data? (2) Which uses only a fixed constant? (3) 2.54 cm per inch.

**Answer.** b) Converting height from inches to centimetres Because: r1) A unit conversion uses a fixed constant. It learns nothing from the rows.

Fixed transformations are safe anywhere. Anything that computes a statistic from rows must be fitted on training rows only.

### Common wrong answers.

- a: The median is learned from the rows, including test rows.
- c: Selection by target correlation learns from test labels. It is the strongest leak of the four.
- d: Mean and sd are learned from the rows.
- b|r2: Right answer, wrong reason. Timing does not make a step safe. Learning nothing from the data does.

### B10-L2 exercise 2

**Hints.** (1) Did the selection see the rows used for testing? (2) Can it pick features that fit those labels by chance? (3) Far above 0.5.

#### Answer.

```
1 0.85
```

With 5,000 features, 20 fit any 100 labels by chance. Selecting them on all rows leaks the test labels. Inside the folds, the same code scores 0.48.

### Common wrong answers.

- 0.5: That is the honest answer for noise. Selection on all rows has seen the test labels, so CV reports 0.85.

### B10-L2 exercise 3

**Hints.** (1) Which rows does the scaler learn from? (2) What must the test set not do? (3) `X_train` and `transform`.

#### Answer.

```
1 scaler = StandardScaler().fit(X_train)
2 X_train_s = scaler.transform(X_train)
3 X_test_s = scaler.transform(X_test)
```

Fit once on training rows. Transform everything with the same fitted object.

### Common wrong answers.

- `g2 = fit_transform`: `fit_transform` on the test set refits the scaler on test rows.
- `g1 = x`: Fitting on all rows leaks test statistics into training.

### B10-L2 exercise 4

**Hints.** (1) `fit_prep`: median first, then fill, then mean and std. (2) `apply_prep` reads only from `params`. (3) Guard `params["sd"] == 0`.

#### Answer.

```
1 import pandas as pd
```

```

2
3
4 def fit_prep(train: pd.DataFrame) -> dict[str, float]:
5 """Learn imputation and scaling parameters for income from training
 rows.
6
7 Returns:
8 median, mean and sd of income after median imputation.
9 """
10 median = float(train["income"].median())
11 filled = train["income"].fillna(median)
12 return {"median": median, "mean": float(filled.mean()), "sd":
 float(filled.std())}
13
14
15 def apply_prep(df: pd.DataFrame, params: dict[str, float]) ->
 pd.DataFrame:
16 """Impute and scale income with parameters learned from training
 rows."""
17 filled = df["income"].fillna(params["median"])
18 if params["sd"] == 0:
19 return df.assign(income=0.0)
20 return df.assign(income=(filled - params["mean"]) / params["sd"])

```

All learning sits in `fit_prep`. `apply_prep` is a pure function of the data and the learned parameters, so test rows can never influence them.

## B10-L2 exercise 5

**Hints.** (1) Which rows does `both` contain? (2) Whose labels feed the rates? (3) Group the training rows only, and give unseen regions the training mean.

**Answer.** The bug is on line 6, 7.

Fixed version:

```

1 import pandas as pd
2
3
4 def encode_region(train: pd.DataFrame, test: pd.DataFrame) ->
 tuple[pd.DataFrame, pd.DataFrame]:
5
6 """Replace region with its training churn rate. Unseen regions get the
 training mean."""
7 rates = train.groupby("region")["churned"].mean()
8 prior = float(train["churned"].mean())
9 return (
10 train.assign(region=train["region"].map(rates)),
11 test.assign(region=test["region"].map(rates).fillna(prior)),

```

Concatenating before grouping puts test labels into the encoding. Learn rates from training rows, and fall back to the training mean for regions the

training data never saw. Inside training, cross-fitted encodings reduce the remaining leak, as scikit-learn's TargetEncoder does.

**Common wrong answers.**

- Line 8: Mapping both sets with one table is right. The table was built from test labels.

## B10-L3 Time-safe features and point-in-time joins

---

### B10-L3 exercise 1

**Hints.** (1) How will the model be used? (2) Can it see the future in production? (3) Then not in validation either.

**Answer.** b) Train on months up to a cut-off, test on later months Because: r1) It matches production, where the model predicts a future it has not seen. A time split reproduces the production situation: learn from the past, predict the future.

**Common wrong answers.**

- a: A random split trains on rows from after the test rows, which production never allows.
- c: Backwards in time: the model would learn from the future of its test set.
- d: Training accuracy rewards memorising and says nothing about new data.
- b|r2: Right answer, wrong reason. The point is mirroring the direction of time, not data quality.

### B10-L3 exercise 2

**Hints.** (1) Is there a score on the prediction date itself? (2) What is the default of `allow_exact_matches`? (3) True.

**Answer.**

1 9

The same-day score matches by default. If it was recorded after the prediction was made, it leaks. Use `allow_exact_matches=False`.

**Common wrong answers.**

- 3: `allow_exact_matches` defaults to True, so the score from the same day matches. Set it to False to exclude it.



### B10-L3 exercise 3

**Hints.** (1) The pandas join for 'last value before'. (2) Exclude exact timestamp matches. (3) `merge_asof` and `False`.

**Answer.**

```
1 pd.merge_asof(preds, scores, left_on="as_of", right_on="scored_at",
2 by="customer", allow_exact_matches=False)
```

`merge_asof` takes the last earlier match per by group.

`allow_exact_matches=False` makes "earlier" strict.

**Common wrong answers.**

- `g1 = merge`: A plain merge joins every score for the customer, past and future.

### B10-L3 exercise 4

**Hints.** (1) Sort by customer and date first. (2) Running total with `groupby(...).cumsum()`, minus the current spend. (3) `groupby(...).cumcount()` counts earlier rows from 0.

**Answer.**

```
1 import pandas as pd
2
3
4 def add_prior_spend(orders: pd.DataFrame) -> pd.DataFrame:
5 """Spend and order count before each order, per customer.
6
7 Args:
8 orders: Columns customer, ordered_at, spend.
9
10 Returns:
11 Orders sorted by customer and date, with prior_spend and
12 prior_orders.
13 """
14 o = orders.sort_values(["customer", "ordered_at"])
15 return o.assign(
16 prior_spend=o.groupby("customer")["spend"].cumsum() - o["spend"],
17 prior_orders=o.groupby("customer").cumcount(),
18)
```

Sorting fixes the time order. Subtracting the current spend and counting from 0 make both features strictly prior.

### B10-L3 exercise 5

**Hints.** (1) Does the merge filter scores by date? (2) Which score does `tail(1)` pick? (3) The latest overall, even from the future. Use `merge_asof`.

**Answer.** The bug is on line 6, 7.

Fixed version:

```
1 import pandas as pd
2
3
4 def latest_nps(preds: pd.DataFrame, scores: pd.DataFrame) -> pd.DataFrame:
5 """Most recent NPS strictly before each prediction date."""
6 joined = pd.merge_asof(
7 preds.sort_values("as_of"),
8 scores.sort_values("scored_at"),
9 left_on="as_of",
10 right_on="scored_at",
11 by="customer",
12 allow_exact_matches=False,
13)
14 return joined[["customer", "as_of", "nps"]].sort_values(["customer",
15 "as_of"]).reset_index(drop=True)
```

A merge on customer followed by "latest" picks scores from after the prediction date. `merge_asof` with `allow_exact_matches=False` picks the last strictly earlier score.

### Common wrong answers.

- Line 8: Selecting and sorting columns is fine. The rows chosen before it include future scores.

## B11-L1 Train, validation and test, split by the right unit

---

### B11-L1 exercise 1

**Hints.** (1) What did the test set do here? (2) Is the best of 40 lucky? (3) Use validation for choosing.

**Answer.** b) An optimistic estimate, because the test set chose the model. Because: r1) Picking the best of 40 on the same data selects for luck on that data.

Choose with validation or CV. Keep the test set for one final measurement.

### Common wrong answers.

- a: It would be, if the test set were used once. Selection on it adds optimism.
- c: Selection favours settings that got lucky on this data, so the bias is upwards.
- d: Test rows were not used for fitting, so it differs from the training score.
- b|r2: Right answer, wrong reason. Used once, a test set is unbiased.

## B11-L1 exercise 2

**Hints.** (1) `test_size` counts groups, not rows. (2) How many respondents is 25% of 4? (3) Each has 2 rows.

**Answer.**

```
1 2 1
```

One of four respondents goes to test, with both rows. No respondent is split across sets.

**Common wrong answers.**

- `2 2`: A group split moves whole groups. 25% of 4 respondents is 1 respondent, with both of their rows.
- `1 1`: The test respondent brings every row: 2 rows.

## B11-L1 exercise 3

**Hints.** (1) The argument that preserves class proportions. (2) The argument that fixes the seed. (3) `stratify` and `random_state`.

**Answer.**

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
 stratify=y, random_state=42)
```

`stratify=y` preserves the class balance. A fixed `random_state` makes the split repeatable.

**Common wrong answers.**

- `g1 = shuffle`: `shuffle` controls whether rows are shuffled. `stratify` keeps class shares.

## B11-L1 exercise 4

**Hints.** (1) `GroupShuffleSplit(n_splits=1, test_size=..., random_state=seed)`. (2) `next(splitter.split(df, groups=df[group_col]))` gives index arrays. (3) Select with `df.iloc[...]`.

**Answer.**

```
1 import pandas as pd
2 from sklearn.model_selection import GroupShuffleSplit
3
4
5 def split_by_group(
6 df: pd.DataFrame, group_col: str, test_size: float, seed: int
7) -> tuple[pd.DataFrame, pd.DataFrame]:
8 """Split rows so no group appears in both train and test.
9
10 Args:
```

```

11 df: Rows to split.
12 group_col: Column identifying the repeated unit, such as
13 respondent.
14 test_size: Share of groups in the test set.
15 seed: Random seed for reproducibility.
16
17 Returns:
18 """
19 splitter = GroupShuffleSplit(n_splits=1, test_size=test_size,
20 random_state=seed)
21 train_idx, test_idx = next(splitter.split(df, groups=df[group_col]))
22 return df.iloc[train_idx], df.iloc[test_idx]

```

Group splitting keeps each respondent on one side, so validation measures performance on new people, as production will.

### B11-L1 exercise 5

**Hints.** (1) What will be new in production? (2) Are rows independent? (3) Is time part of the use?

**Answer.**

- Panel members will be new in production: group by member.
- Independent rows with a rare class: stratify on the target.
- Forecasting: train on the past, test on later months.

The split mirrors production: new people, rare classes kept, or the future.

**Common wrong answers.**

- panel = stratified: Members' waves would land on both sides. The model would recognise people, not patterns.
- fraud = group: Rows are independent here, so grouping adds nothing. The risk is a test set with too few frauds.
- forecast = stratified: A random split lets the model learn from months after the ones it is tested on.

## B11-L2 Pipeline and ColumnTransformer

### B11-L2 exercise 1

**Hints.** (1) What does cross\_val\_score do with the estimator in each fold? (2) Which rows does the scaler see when fitted? (3) Only that fold's training rows.

**Answer.** a) The Pipeline refits the scaler on each fold's training rows only. Because: r1) cross\_val\_score clones and fits the whole pipeline per fold.

Each fold gets a fresh copy of the pipeline, fitted on that fold's training rows. The test rows are only transformed.

### Common wrong answers.

- `b`: Test rows are scaled, with the parameters learned from that fold's training rows.
- `c`: Fitting once on all rows is exactly the leak the pipeline avoids.
- `d`: Pipelines only run the steps you put in them.
- `a` | `r2`: Right answer, wrong reason. The scaler learns mean and scale, per fold.

### B11-L2 exercise 2

**Hints.** (1) Was region e seen in fit? (2) What is the default `handle_unknown`? (3) "error".

### Answer.

```
1 ValueError
```

The default raises on unseen categories. Use `handle_unknown="ignore"` for production inputs.

### Common wrong answers.

- `encoded`: The default `handle_unknown` is "error". An unseen category raises `ValueError`.

### B11-L2 exercise 3

**Hints.** (1) The `handle_unknown` value that does not raise. (2) The remainder value that keeps columns. (3) `ignore` and `passthrough`.

### Answer.

```
1 prep = ColumnTransformer(
2 [{"cat", OneHotEncoder(handle_unknown="ignore"), ["region"]}],
3 remainder="passthrough",
4)
```

`handle_unknown="ignore"` avoids errors on new categories.

`remainder="passthrough"` keeps other columns. Use it on purpose, not by accident.

### Common wrong answers.

- `g2 = drop`: `drop` is the default, and it discards every unlisted column.

## B11-L2 exercise 4

**Hints.** (1) `make_pipeline(SimpleImputer(strategy="median"), StandardScaler())` for numbers. (2) `make_pipeline(SimpleImputer(strategy="most_frequent"), OneHotEncoder(handle_unknown="ignore"))` for categories. (3) `Pipeline([("prep", prep), ("clf", LogisticRegression(max_iter=1000))])`.

### Answer.

```
1 from sklearn.compose import ColumnTransformer
2 from sklearn.impute import SimpleImputer
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.pipeline import Pipeline, make_pipeline
5 from sklearn.preprocessing import OneHotEncoder, StandardScaler
6
7
8 def build_pipeline(numeric: list[str], categorical: list[str]) ->
 Pipeline:
9 """Preprocessing and classifier as one estimator.
10
11 Args:
12 numeric: Numeric feature columns.
13 categorical: Categorical feature columns.
14
15 Returns:
16 An unfitted Pipeline with steps prep and clf.
17 """
18 prep = ColumnTransformer(
19 [
20 ("num", make_pipeline(SimpleImputer(strategy="median"),
21 StandardScaler()), numeric),
22 (
23 "cat",
24 make_pipeline(SimpleImputer(strategy="most_frequent"),
25 OneHotEncoder(handle_unknown="ignore")),
26 categorical,
27),
28],
29)
30 return Pipeline([("prep", prep), ("clf",
31 LogisticRegression(max_iter=1000))])
```

One estimator holds every learned step. Cross-validation refits it per fold, and production runs exactly what was validated.

## B11-L2 exercise 5

**Hints.** (1) Which columns does the ColumnTransformer list? (2) What happens to the others by default? (3) `remainder="drop"`.

**Answer.** The bug is on line 9.

Fixed version:

```
1 from sklearn.compose import ColumnTransformer
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.pipeline import Pipeline
4 from sklearn.preprocessing import OneHotEncoder
5
6
7 def build(categorical: list[str]) -> Pipeline:
8 """Encode categorical columns and fit a classifier on every
9 feature."""
10 prep = ColumnTransformer(
11 [("cat", OneHotEncoder(handle_unknown="ignore"), categorical)],
12 remainder="passthrough",
13)
14 return Pipeline([("prep", prep), ("clf",
15 LogisticRegression(max_iter=1000))])
```

`ColumnTransformer` drops unlisted columns by default. Pass them through on purpose, or better, list and preprocess them explicitly.

### Common wrong answers.

- Line 10: The classifier uses whatever prep outputs. prep never outputs tenure.

## B11-L3 Cross-validation and honest hyperparameter search

---

### B11-L3 exercise 1

**Hints.** (1) How was the best chosen? (2) On which data was it scored? (3) Selection on the same folds.

**Answer.** b) The mean CV score of the best setting, optimistic because it was selected as the best Because: r1) Selecting the maximum of many noisy scores favours settings that were lucky on these folds.

`best_score_` is the winner's CV score. Use nested CV or an untouched test set for the estimate you report.

### Common wrong answers.

- a: It is the maximum over 80 candidates on the same folds, so it carries selection optimism.
- c: GridSearchCV does not use a separate test set.
- d: It is a CV score on held-out folds, not a training score.
- b|r2: Right answer, wrong reason. The search is fine. Its best score needs an outer check.

## B11-L3 exercise 2

**Hints.** (1) How does `make_pipeline` name its steps? (2) Lower-case class names. (3) Parameter keys are `step__param`.

**Answer.**

```
1 False True
```

`make_pipeline` names steps `standardscaler` and `logisticregression`. Use `Pipeline([("clf", ...)])` when you want short, stable names for your grids.

**Common wrong answers.**

- `True True`: `make_pipeline` names steps after their class in lower case. There is no step called `clf`.
- `True False`: Step names come from `make_pipeline`: `logisticregression`, so the key is `logisticregression__C`.

## B11-L3 exercise 3

**Hints.** (1) `step__param`. (2) The argument for the metric. (3) `clf__C` and `scoring`.

**Answer.**

```
1 search = GridSearchCV(pipe, {"clf__C": [0.01, 0.1, 1, 10]}, cv=inner,
 scoring="roc_auc")
```

Pipeline parameters are addressed as `step__param`. `scoring` sets the metric the search optimises.

**Common wrong answers.**

- `g1 = C`: Inside a Pipeline, prefix the parameter with the step name and two underscores: `clf__C`.
- `g1 = clf_C`: Two underscores separate the step name from the parameter: `clf__C`.

## B11-L3 exercise 4

**Hints.** (1) Build the Pipeline, then the `GridSearchCV` around it. (2) Pass the search object, unfitted, to `cross_val_score`. (3) Both loops use `scoring="roc_auc"`.

**Answer.**

```
1 import numpy as np
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.model_selection import GridSearchCV, StratifiedKFold,
 cross_val_score
4 from sklearn.pipeline import Pipeline
5 from sklearn.preprocessing import StandardScaler
```



```

6
7
8 def nested_cv_auc(X: np.ndarray, y: np.ndarray, c_values: list[float],
 seed: int) -> float:
9 """Nested CV estimate of ROC AUC for a tuned logistic regression.
10
11 Args:
12 X: Features.
13 y: Binary target.
14 c_values: Candidate regularisation strengths.
15 seed: Base seed for the inner and outer splitters.
16
17 Returns:
18 Mean outer-fold ROC AUC of the whole search.
19 """
20 pipe = Pipeline([("scale", StandardScaler()), ("clf",
 LogisticRegression(max_iter=1000))])
21 inner = StratifiedKFold(n_splits=3, shuffle=True, random_state=seed)
22 outer = StratifiedKFold(n_splits=5, shuffle=True, random_state=seed +
 1)
23 search = GridSearchCV(pipe, {"clf__C": c_values}, cv=inner,
 scoring="roc_auc")
24 return float(cross_val_score(search, X, y, cv=outer,
 scoring="roc_auc").mean())

```

The outer loop treats the search as one model. Its folds never influence which C is chosen, so the mean is an honest estimate.

## B11-L3 exercise 5

**Hints.** (1) Which data never influences a choice? (2) Which metric is tuned? (3) Which number is honest?

### Answer.

- Hold out a test set at the start.
- Search on the rest with CV, a fixed seed and the reported metric.
- Report one test-set score, plus the fold spread from CV as a guide to uncertainty.

Keep one set of data out of every decision, and report its score once.

### Common wrong answers.

- `search = best_test`: Choosing on the test set makes its score optimistic.
- `report = best_score`: `best_score_` is the best of many tries on the same folds.
- `holdout = use_all`: Without a holdout, use nested CV for the estimate. Otherwise the number is optimistic.

## B12-L1 Choose the metric from the decision

---

### B12-L1 exercise 1

**Hints.** (1) What does a constant prediction score? (2) Compare with 1 - prevalence. (3) 0.976.

**Answer.** b) Nothing useful yet. Predicting all negatives scores the same. Because: r1) Accuracy is dominated by the 97.6% majority class.

With 2.4% positives, 0.976 is the no-skill baseline. Look at recall, precision and average precision.

#### Common wrong answers.

- a: A constant 'negative' prediction scores 0.976 too.
- c: That would be recall. Accuracy mixes both classes.
- d: Accuracy says nothing about probability calibration.
- b|r2: Right answer, wrong reason. Accuracy works for balanced classes with equal costs.

### B12-L1 exercise 2

**Hints.** (1) A random ranking's precision at every recall level is about the prevalence. (2) What is the prevalence? (3) About 5%.

**Answer.**

```
1 0.05
```

The no-skill average precision equals the prevalence. Always report it beside the model's score.

#### Common wrong answers.

- 0.5: 0.5 is the random baseline for ROC AUC. For average precision the baseline is the prevalence, here 5%.

### B12-L1 exercise 3

**Hints.** (1) np.sort sorts ascending. (2) Count from the end. (3) -200, then compare with threshold.

**Answer.**

```
1 threshold = np.sort(scores)[-200]
2 flagged = scores >= threshold
```

The threshold comes from capacity, not from 0.5. Ties at the threshold can flag a few more than 200.

## Common wrong answers.

- `g1 = 200`: `np.sort` is ascending. The 200th largest value is at index -200.

## B12-L1 exercise 4

**Hints.** (1) `flagged = scores >= threshold`. `tp` counts flagged positives. (2) `precision = tp / flagged count`. `recall = tp / positive count`. (3) Return `None` when the denominator is 0.

## Answer.

```
1 import numpy as np
2 from sklearn.metrics import average_precision_score, roc_auc_score
3
4
5 def evaluate(y_true: np.ndarray, scores: np.ndarray, threshold: float) ->
 dict[str, float | None]:
6 """Ranking metrics, threshold metrics and the prevalence baseline.
7
8 Args:
9 y_true: Binary labels.
10 scores: Model scores, higher means more likely positive.
11 threshold: Flag cases with score >= threshold.
12
13 Returns:
14 prevalence, roc_auc, average_precision, precision, recall.
15 """
16 flagged = scores >= threshold
17 tp = int((flagged & (y_true == 1)).sum())
18 n_flagged = int(flagged.sum())
19 n_pos = int((y_true == 1).sum())
20 return {
21 "prevalence": float(y_true.mean()),
22 "roc_auc": float(roc_auc_score(y_true, scores)),
23 "average_precision": float(average_precision_score(y_true,
24 scores)),
25 "precision": tp / n_flagged if n_flagged > 0 else None,
26 "recall": tp / n_pos if n_pos > 0 else None,
27 }
```

Counts make precision and recall transparent. Returning `None` for an empty denominator avoids a misleading 0.0 and a crash.

## B12-L1 exercise 5

**Hints.** (1) Is capacity fixed? (2) Are costs known and unequal? (3) Is there a threshold yet?

## Answer.

- Fraud with fixed review capacity: precision in the top 150, the share of analyst time spent on real fraud.

- Comparing rankers: ROC AUC, with average precision beside it for rare positives.
- Unequal, known costs: total cost at the threshold, which picks the threshold too.

Metrics follow the decision: capacity, costs, or pure ranking.

### Common wrong answers.

- fraud = accuracy: At 1% prevalence, flagging nothing is 99% accurate.
- credit = auc: AUC ignores the threshold and the cost difference, which decide profit.
- compare = cost: Without a threshold or agreed costs, compare rankings first.

## B12-L2 Calibration, when the probability itself matters

---

### B12-L2 exercise 1

**Hints.** (1) What does AUC measure? (2) What would a calibrated model show in the 0.8 to 0.9 bin? (3) About 85% observed.

**Answer.** a) It ranks well but is poorly calibrated. Because: r1) AUC measures ordering only. The probabilities can be far off while the order is right.

Good ranking, poor calibration. Fine for a top-k list, wrong for expected-value sums.

### Common wrong answers.

- b: Calibrated would mean about 85% churn in that bin, not 20%.
- c: Nothing contradicts the AUC. Ranking and calibration are separate properties.
- d: 0.85 AUC is good ranking.
- a|r2: Right answer, wrong reason. AUC says nothing about calibration.

### B12-L2 exercise 2

**Hints.** (1) Brier = mean of  $(p - y)^2$ . (2) For  $p = 0.25$ : three rows of 0.0625 and one of 0.5625. (3) For  $p = 0.5$  every row is 0.25.

### Answer.

|   |        |      |
|---|--------|------|
| 1 | 0.1875 | 0.25 |
|---|--------|------|

The prevalence is the best constant prediction. A model worse than 0.1875 here is worse than knowing only the base rate.

### Common wrong answers.

- 0.25 0.25: Predicting the prevalence (0.25) gives  $(3 \times 0.0625 + 0.5625) / 4 = 0.1875$ .

### B12-L2 exercise 3

**Hints.** (1) The scikit-learn wrapper in `sklearn.calibration`. (2) Platt scaling's other name. (3) `CalibratedClassifierCV` and `sigmoid`.

### Answer.

```
1 calibrated = CalibratedClassifierCV(model, method="sigmoid",
 cv=5).fit(X_train, y_train)
```

`CalibratedClassifierCV` fits the model and a calibration map on separate folds, so calibration is learned on data the model did not train on.

### Common wrong answers.

- `g2 = isotonic`: Isotonic works, but overfits small samples. The prompt asks for `sigmoid`.

### B12-L2 exercise 4

**Hints.** (1) `edges = np.linspace(0.0, 1.0, n_bins + 1)`. (2) Group a `DataFrame` of `bin`, `y` and `p` by `bin`. (3) Named aggregation: `mean` of `p`, `mean` of `y`, `size`.

### Answer.

```
1 import numpy as np
2 import pandas as pd
3
4
5 def calibration_table(y: np.ndarray, p: np.ndarray, n_bins: int) ->
 pd.DataFrame:
6 """Mean predicted probability and observed rate per equal-width bin.
7
8 Args:
9 y: Binary outcomes.
10 p: Predicted probabilities.
11 n_bins: Number of equal-width bins on [0, 1].
12
13 Returns:
14 One row per non-empty bin: mean_predicted, observed, n.
15 """
16 edges = np.linspace(0.0, 1.0, n_bins + 1)
17 idx = np.searchsorted(edges[1:-1], p, side="right")
18 df = pd.DataFrame({"bin": idx, "y": y, "p": p})
19 table = df.groupby("bin").agg(mean_predicted=("p", "mean"),
 observed=("y", "mean"), n=("y", "size"))
20 return table.reset_index(drop=True)
```

Observed is the positive rate **within** each bin, the number the predicted mean should match. `groupby` leaves empty bins out.

## B12-L2 exercise 5

**Hints.** (1) Does the use read the order or the number? (2) What does summing probabilities assume? (3) What do balanced weights do to probabilities?

**Answer.**

- A top-k list needs order, not calibrated values.
- A count forecast sums probabilities, so they must be calibrated.
- Reweighting distorts probabilities. Recalibrate on held-out data and check the Brier score against the constant baseline.

Deliver ranking where order is enough. Deliver calibrated probabilities where the number is used.

**Common wrong answers.**

- `callist = calibrated`: Calibration does not change the order, so it does not change who is in the top 500.
- `forecast = ranking`: A sum of probabilities is only right if the probabilities are right.
- `training = balanced_only`: In this lesson's example, balanced weights moved the mean probability from 0.024 to 0.42.

## B12-L3 Imbalanced classes and error analysis by segment

---

### B12-L3 exercise 1

**Hints.** (1) Is the problem the ranking or the cut-off? (2) What does 0.5 mean at 2% prevalence? (3) Choose the threshold from the decision.

**Answer.** b) Lower the threshold to match review capacity or costs. Because: r1) The model's ranking may be fine. Only the cut-off is wrong for rare positives.

At 2% prevalence few cases score above 0.5. Pick the threshold that matches capacity or cost, and check precision and recall there.

**Common wrong answers.**

- a: Never resample the test set. Its balance must match production.
- c: A different model does not fix a threshold chosen by default.

- `d`: Accuracy would look excellent while the model finds nothing.
- `b|r2`: Right answer, wrong reason. A lower threshold trades precision for recall. It does not improve the model.

## B12-L3 exercise 2

**Hints.** (1) Recall is found over positives. (2) Overall pools the counts. (3) 84 / 120.

**Answer.**

```
1 email 0.8
2 chat 0.2
3 overall 0.7
```

An overall 0.7 hides a segment at 0.2. Pooled rates weight large segments more, so small segments disappear from the headline.

**Common wrong answers.**

- `email 0.8 chat 0.2 overall 0.5`: Overall recall pools the counts: 84 of 120 is 0.7, not the mean of the two rates.

## B12-L3 exercise 3

**Hints.** (1) What must exist before training? (2) What do segment metrics depend on? (3) Reading examples comes last.

**Answer.**

1. Split train and test, stratified, test left at its natural balance
2. Train the model normally
3. Choose the threshold from capacity or cost on validation data
4. Break test results down by segment, with counts
5. Read the most confident false positives and false negatives

Split, train, threshold, then look for where and why the model fails.

**Common wrong answers.**

- Split before training, so the test set never influences the model.
- Segment precision and recall depend on the threshold. Choose it first.

## B12-L3 exercise 4

**Hints.** (1) Loop over `df.groupby(segment, sort=True)`. (2) `tp` counts rows where `y` and `flagged` are both 1, within the segment. (3) Build a list of dicts and return `pd.DataFrame(rows)`.

**Answer.**

```
1 import pandas as pd
```

```

2
3
4 def metrics_by_segment(df: pd.DataFrame, segment: str) -> pd.DataFrame:
5 """Count, positives, precision and recall per segment.
6
7 Args:
8 df: Columns y and flagged (0/1) plus the segment column.
9 segment: Name of the segment column.
10
11 Returns:
12 One row per segment, sorted by segment.
13 """
14 rows = []
15 for name, g in df.groupby(segment, sort=True):
16 tp = int((g["y"] == 1) & (g["flagged"] == 1)).sum()
17 n_flag = int(g["flagged"].sum())
18 n_pos = int(g["y"].sum())
19 rows.append(
20 {
21 segment: name,
22 "n": len(g),
23 "positives": n_pos,
24 "precision": tp / n_flag if n_flag > 0 else None,
25 "recall": tp / n_pos if n_pos > 0 else None,
26 }
27)
28 return pd.DataFrame(rows)

```

Per-segment counts and rates show where the model fails. Missing values for empty denominators stop a segment with no positives looking like a failure.

## B12-L3 exercise 5

**Hints.** (1) Which set must look like production? (2) What does line 7 do to it? (3) Return test unchanged.

**Answer.** The bug is on line 7, 8.

Fixed version:

```

1 import pandas as pd
2
3
4 def prepare_sets(train: pd.DataFrame, test: pd.DataFrame, seed: int) ->
5 tuple[pd.DataFrame, pd.DataFrame]:
6 """Oversample training positives. The test set keeps its natural
7 balance."""
8 pos_train = train[train["y"] == 1].sample(frac=5.0, replace=True,
9 random_state=seed)
10 return pd.concat([train, pos_train]), test

```

Oversampling the test set inflates positives and duplicates them, so recall is measured on a world that does not exist. Resample training data only.



### Common wrong answers.

- Line 6: Oversampling training positives is a legitimate choice. The problem is the other set.

## B13-L1 Permutation importance on held-out data

---

### B13-L1 exercise 1

**Hints.** (1) How many split points does a unique ID offer? (2) Which data is impurity importance computed on? (3) Training data.

**Answer.** b) Impurity importance favours features with many distinct values, measured on training data. Because: r1) Many split points let trees memorise training rows, which impurity importance rewards.

Unique values give trees endless ways to split training rows. Impurity importance counts those splits. Permutation on held-out data does not.

### Common wrong answers.

- a: Permutation importance on held-out data would show a drop if it did. In this lesson's example it was 0.004.
- c: More trees make the bias more stable, not smaller.
- d: Trees are not affected by scaling.
- b|r2: Right answer, wrong reason. Forests handle integers. Cardinality causes the bias.

### B13-L1 exercise 2

**Hints.** (1) What does the model still have when only age is shuffled? (2) The same information via birth\_year. (3) So the single drop is small.

**Answer.**

1 True

Correlated features share credit. Permute them as a group to see the importance of the information they carry.

### Common wrong answers.

- False: With birth\_year intact, the model still has the information, so shuffling age alone costs little.

### B13-L1 exercise 3

**Hints.** (1) Held-out rows. (2) The spread across repeats. (3) `x_test` and `importances_std`.

## Answer.

```
1 result = permutation_importance(model, X_test, y_test, scoring="roc_auc",
 n_repeats=10, random_state=0)
2 print(result.importances_mean, result.importances_std)
```

Held-out rows show what generalises. The standard deviation shows whether a drop is distinguishable from zero.

## Common wrong answers.

- `g1 = X_train`: On training data, memorised features look important. Use held-out rows.

## B13-L1 exercise 4

**Hints.** (1) Base AUC first, on the unshuffled data. (2) `X.assign(**{col: rng.permutation(X[col].to_numpy())})` shuffles a copy. (3) Average the drops with `np.mean`.

## Answer.

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.base import ClassifierMixin
4 from sklearn.metrics import roc_auc_score
5
6
7 def permutation_drop(model: ClassifierMixin, X: pd.DataFrame, y:
 pd.Series, col: str, rng: np.random.Generator, n_repeats: int) ->
 float:
8 """Mean fall in ROC AUC when col is shuffled.
9
10 Args:
11 model: Fitted classifier with predict_proba.
12 X: Held-out features.
13 y: Held-out labels.
14 col: Column to shuffle.
15 rng: Seeded generator.
16 n_repeats: Number of shuffles.
17
18 Returns:
19 Mean of base AUC minus shuffled AUC.
20 """
21 base = roc_auc_score(y, model.predict_proba(X)[: , 1])
22 drops = []
23 for _ in range(n_repeats):
24 shuffled = X.assign(**{col: rng.permutation(X[col].to_numpy())})
25 drops.append(base - roc_auc_score(y, model.predict_proba(shuffled)
26 [: , 1]))
27 return float(np.mean(drops))
```

Shuffling keeps the column's distribution but breaks its link with the target. The fall in AUC is how much the model relied on it.

### B13-L1 exercise 5

**Hints.** (1) Which method uses held-out data? (2) What happens when one of two twins is shuffled? (3) Does reliance imply cause?

**Answer.**

- Permutation importance on held-out data, with the metric the model is judged on.
- Correlated features grouped, so shared information is not split into two small numbers.
- Words that describe the model, not causes in the world.

Deliver reliance, measured honestly and worded honestly.

**Common wrong answers.**

- method = `impurity`: Impurity importance favours high-cardinality features and reflects training data.
- correlated = `separate`: Each would look small, because the other carries the same information.
- wording = `cause`: Importance shows what the model uses. Causes need a different design, such as an experiment.

## B13-L2 SHAP values, and what they do not tell you

---

### B13-L2 exercise 1

**Hints.** (1) What does SHAP attribute? (2) The model's output. (3) Not the outcome in the world.

**Answer.** b) The model uses the feature to push churn scores up. Because: r1) SHAP values attribute the model's output. They say nothing about intervening in the world.

SHAP describes the model's reliance. Test interventions separately.

**Common wrong answers.**

- `a`: Causal claims need causal evidence, such as an experiment.
- `c`: Nothing in a SHAP value indicates measurement error.
- `d`: Large contributions are normal for useful features.
- `b` | `r2`: Right answer, wrong reason. The values are real attributions of the model's output.

## B13-L2 exercise 2

**Hints.** (1) Add the base value and all contributions. (2) That is in log-odds. (3)  $p = 1 / (1 + \exp(-\log\_odds))$ .

**Answer.**

```
1 -0.2 0.45
```

Additivity holds in log-odds for this model. The probability is 0.45 after the logistic transform, and the contributions do not add up in probability space.

**Common wrong answers.**

- `-0.2 -0.2`: The contributions sum to log-odds. Convert with the logistic function to get a probability.
- `0.8 0.69`: Include the base value:  $-1.2 + 0.8 + 0.5 - 0.3 = -0.2$ .

## B13-L2 exercise 3

**Hints.** (1) The explainer for tree ensembles. (2) The attribute with one number per feature per row. (3) `TreeExplainer` and `values`.

**Answer.**

```
1 explainer = shap.TreeExplainer(model)
2 explanation = explainer(X_test)
3 contributions = explanation.values
```

Calling the explainer returns a `shap.Explanation`. `.values` holds contributions, `.base_values` the baseline.

**Common wrong answers.**

- `g2 = base_values`: `base_values` is the starting point. The per-feature contributions are `.values`.

## B13-L2 exercise 4

**Hints.** (1) Check the lengths first. (2) `sorted(..., key=lambda pair: abs(pair[1]), reverse=True)`. (3) Slice the first `k`.

**Answer.**

```
1 def top_contributions(names: list[str], values: list[float], k: int) ->
 list[tuple[str, float]]:
2 """The k largest contributions to one prediction, by absolute size.
3
4 Args:
5 names: Feature names.
6 values: SHAP values for one row, same order as names.
7 k: Number of pairs to return.
8
9 Returns:
```

```

10 (name, value) pairs, largest absolute value first, sign kept.
11
12 Raises:
13 ValueError: If names and values differ in length.
14 """
15 if len(names) != len(values):
16 raise ValueError("one SHAP value per feature is required")
17 ranked = sorted(zip(names, values, strict=True), key=lambda pair:
18 abs(pair[1]), reverse=True)
19 return ranked[:k]

```

A large negative contribution matters as much as a large positive one. Rank by size, report with sign.

### B13-L2 exercise 5

**Hints.** (1) Do log-odds add as percentages? (2) Local or global? (3) Model or world?

**Answer.**

- Name the output space and convert only the final score.
- Give the local top contributions for this case, not the global ranking.
- Describe the model's reasoning, not a claim about the person.

Local, space-aware and modest wording keeps explanations accurate.

**Common wrong answers.**

- space = `mixed`: Log-odds contributions do not add up in probability points. Labelling them so is wrong.
- content = `global`: The global ranking says what the model uses in general, not why this case was declined.
- wording = `cause`: SHAP shows what moved the model's score. It is not evidence about the person.

## B13-L3 Performance by subgroup, and the fairness trade-offs

---

### B13-L3 exercise 1

**Hints.** (1) Which other features correlate with sex? (2) How would you measure a gap without the column? (3) Keep it for measurement where lawful.

**Answer.** b) They can persist through proxies, and you lose the ability to measure them if you discard the column. Because: r1) Other features correlate with sex, and measurement needs the attribute.

Exclusion is not neutrality. Keep the attribute for measuring gaps, and decide separately whether the model uses it.

### Common wrong answers.

- a: Proxies such as occupation or purchase history can carry the same information.
- c: They may change either way. The point is you can no longer see them.
- d: Removing a feature says nothing about calibration.
- b|r2: Right answer, wrong reason. Removing a feature does change the model, just not necessarily the disparity.

### B13-L3 exercise 2

**Hints.** (1)  $TPR = tp / \text{positives}$ . (2)  $FPR = fp / \text{negatives}$ . (3) 150 / 390.

### Answer.

```
1 A 0.65 0.38
2 B 0.15 0.03
```

Same threshold, very different error rates, because the groups' base rates differ.

### Common wrong answers.

- A 0.65 0.15 B 0.15 0.03: FPR divides false positives by negatives:  $150 / 390 = 0.38$ .

### B13-L3 exercise 3

**Hints.** (1) True positives are flagged positives. (2) False positives are flagged negatives. (3) pos and neg.

### Answer.

```
1 tpr = (flag & pos).sum() / pos.sum()
2 fpr = (flag & neg).sum() / neg.sum()
```

Both rates condition on the true class: TPR on positives, FPR on negatives.

### Common wrong answers.

- g2 = pos: False positives are flagged negatives: flag & neg.

### B13-L3 exercise 4

**Hints.** (1) Loop over `df.groupby(group, sort=True)`. (2) Boolean masks: pos, neg and flag. (3) Guard each denominator.

### Answer.

```
1 import pandas as pd
2
```

```

3
4 def group_rates(df: pd.DataFrame, group: str) -> pd.DataFrame:
5 """Selection rate, TPR and FPR per group, with counts.
6
7 Args:
8 df: Columns y and flagged (0/1) and the group column.
9 group: Name of the group column.
10
11 Returns:
12 One row per group, sorted by group.
13 """
14 rows = []
15 for name, g in df.groupby(group, sort=True):
16 pos, neg = g["y"] == 1, g["y"] == 0
17 flag = g["flagged"] == 1
18 rows.append(
19 {
20 group: name,
21 "n": len(g),
22 "positives": int(pos.sum()),
23 "selection_rate": float(flag.mean()),
24 "tpr": float((flag & pos).sum() / pos.sum()) if pos.sum()
25 > 0 else None,
26 "fpr": float((flag & neg).sum() / neg.sum()) if neg.sum()
27 > 0 else None,
28 }
29)
30 return pd.DataFrame(rows)

```

Rates conditioned on the true class show who bears which errors. Counts beside them show how much evidence there is.

### B13-L3 exercise 5

**Hints.** (1) Who is harmed, and by which error? (2) Is the score used as a number? (3) Match the property to the harm.

#### Answer.

- Help should reach genuine cases in every group at the same rate: TPR.
- Wrong referrals harm people, so wrong flags should not fall more on one group: FPR.
- Prices use the probability itself, which must mean the same in every group: calibration.

No single property fits all uses. Pick the one that protects against the harm this decision causes, and document why.

#### Common wrong answers.

- support = fpr: An unneeded offer of help costs little. Missing people who need it is the harm.

- `investigate = tpr`: The harm is to people wrongly referred. Balance false positive rates.
- `pricing = tpr`: Pricing uses probabilities as numbers. They must be calibrated in each group.

## B14-L1 Track every run with MLflow

---

### B14-L1 exercise 1

**Hints.** (1) Params and metrics are stored differently. (2) Which one is numeric? (3) Metrics.

**Answer.** b) `str` Because: r1) MLflow stores parameters as strings, so compare and convert carefully.

Parameters are strings. `run.data.params["C"]` is `'1.0'`. Convert before comparing numerically.

**Common wrong answers.**

- a: Metrics are floats. Parameters come back as strings: `'1.0'`.
- c: `1.0` is a float in Python, and MLflow stores it as the string `'1.0'`.
- d: MLflow does not use Decimal for parameters.
- b|r2: Right answer, wrong reason. Databases store floats. MLflow chooses strings for parameters.

### B14-L1 exercise 2

**Hints.** (1) Can a run's parameter change after logging? (2) What would that do to reproducibility? (3) MLflow refuses.

**Answer.**

```
1 MLflowException
```

"Changing param values is not allowed." Start a new run for new settings.

**Common wrong answers.**

- C is 2.0: Parameters are immutable within a run. Changing a logged value raises `MLflowException`.
- C is 1.0: The second call does not pass silently. It raises.

### B14-L1 exercise 3

**Hints.** (1) The function for one tag. (2) MLflow 3 names logged models with a keyword. (3) `set_tag` and `name`.

**Answer.**



```

1 mlflow.set_tag("git_commit", commit)
2 mlflow.sklearn.log_model(model, name="model")

```

Tags hold run metadata such as commits. In MLflow 3, `log_model` takes `name=` for the model's name in the run.

### Common wrong answers.

- `g1 = log_param`: A commit is metadata about the run, not a setting of the model. Use a tag.

### B14-L1 exercise 4

**Hints.** (1) Open the run with a `with` block. (2) `log_params` for settings, `set_tags` for commit and data hash, `log_metric` for the score. (3) `log_model` with `name="model"`, then return `run.info.run_id`.

### Answer.

```

1 import mlflow
2 import numpy as np
3 from sklearn.linear_model import LogisticRegression
4
5
6 def tracked_fit(X_train: np.ndarray, y_train: np.ndarray, c: float, seed:
 int, commit: str, data_hash: str, cv_auc: float) -> str:
7 """Fit a logistic regression and record the run in MLflow.
8
9 Returns:
10 The MLflow run ID.
11 """
12 with mlflow.start_run(run_name=f"lr-c{c}") as run:
13 mlflow.log_params({"C": c, "seed": seed})
14 mlflow.set_tags({"git_commit": commit, "input_sha256": data_hash})
15 model = LogisticRegression(C=c, random_state=seed,
16 max_iter=1000).fit(X_train, y_train)
17 mlflow.log_metric("cv_auc", cv_auc)
18 mlflow.sklearn.log_model(model, name="model")
19 return run.info.run_id

```

A run that records settings, code, data, score and model can be found, compared and reproduced.

### B14-L1 exercise 5

**Hints.** (1) Does it change the result? (2) Is it a number to compare? (3) Is it a file?

### Answer.

- Seed: a parameter, because it changes the result.
- Commit: a tag, metadata that links the run to code.

- AUC: a metric, numeric and comparable across runs.
- Plot: an artifact file.

Parameters for settings, tags for identity, metrics for numbers, artifacts for files.

### Common wrong answers.

- seed = `tag`: The seed is a setting that changes the output. Log it as a parameter so searches can filter on it.
- auc = `param`: Parameters are strings. Metrics are numeric and sortable.
- commit = `metric`: A commit hash is not a number to compare. Tag it.

## B14-L2 Registry versions and aliases, not stages

---

### B14-L2 exercise 1

**Hints.** (1) Which reference changes without editing the job? (2) What replaced stages? (3) Aliases.

**Answer.** b) `models:/churn_model@champion` Because: r1) The alias moves on promotion and rollback, so the job's code never changes.

`@champion` resolves to whichever version currently holds the alias.

### Common wrong answers.

- a: A fixed version number means editing the job for every release.
- c: Stages are deprecated since MLflow 2.9 and not supported in Unity Catalog.
- d: Files have no version history, run link or access control.
- b|r2: Right answer, wrong reason. Speed is the same. Indirection is the benefit.

### B14-L2 exercise 2

**Hints.** (1) Follow each assignment in order. (2) The last line is a rollback. (3) Both point at 3.

### Answer.

```
1 {'champion': '3', 'previous_champion': '3'}
```

Promotion saved the old champion, and rollback restored it in one assignment. Versions 3 and 5 are both unchanged.

## Common wrong answers.

- `{'champion': '5', 'previous_champion': '3'}`: The last line rolls back: champion takes previous\_champion's version, 3.

## B14-L2 exercise 3

**Hints.** (1) The `MlflowClient` method that sets an alias. (2) The URI separator for aliases. (3) `set_registered_model_alias` and `@`.

## Answer.

```
1 client.set_registered_model_alias("churn_model", "champion", "4")
2 model = mlflow.pyfunc.load_model("models:/churn_model@champion")
```

`set_registered_model_alias(name, alias, version)` moves the label. `models:/name@alias` loads through it.

## Common wrong answers.

- `g2 = /`: `models:/name/X` reads X as a version or a stage. Aliases use `@`.

## B14-L2 exercise 4

**Hints.** (1) Resolve the champion through its alias, then read both scores from their runs. (2) Return False early when `cand_score < champ_score + min_gain`. (3) Set `previous_champion` before moving champion.

## Answer.

```
1 from __future__ import annotations
2
3 from typing import TYPE_CHECKING
4
5 if TYPE_CHECKING:
6 from mlflow import MlflowClient
7
8
9 def promote_if_better(client: MlflowClient, name: str, candidate: str,
10 metric: str, min_gain: float) -> bool:
11 """Move the champion alias to candidate if it beats the champion by
12 min_gain.
13
14 Args:
15 client: MLflow client, or any object with the same methods.
16 name: Registered model name.
17 candidate: Candidate version number, as a string.
18 metric: Run metric to compare, higher is better.
19 min_gain: Smallest improvement that justifies promotion.
20
21 Returns:
22 True if the alias moved.
23 """
24 champ = client.get_model_version_by_alias(name, "champion")
```

```

23 champ_score = client.get_run(champ.run_id).data.metrics[metric]
24 cand_score = client.get_run(client.get_model_version(name,
candidate).run_id).data.metrics[metric]
25 if cand_score < champ_score + min_gain:
26 return False
27 client.set_registered_model_alias(name, "previous_champion",
champ.version)
28 client.set_registered_model_alias(name, "champion", candidate)
29 return True

```

Promotion is a comparison plus two alias moves. The margin guards against noise, and the saved alias makes rollback trivial.

## B14-L2 exercise 5

**Hints.** (1) How are Unity Catalog objects named? (2) What replaced stages? (3) Grant to groups (lesson A4-L2).

**Answer.**

- Three-level names place the model in a catalog and schema, governed like tables.
- Jobs load by alias. Unity Catalog models do not support stages.
- Grant the right to move aliases to a release group, not to every trainer.

Govern models like tables: three-level names, aliases for releases, group grants for who can release.

**Common wrong answers.**

- pointer = stage: Stages are deprecated and not supported for models in Unity Catalog.
- who = anyone: Moving the alias changes production. Limit it to a group, as with any production change.
- name = plain: Unity Catalog models use catalog.schema.model names.

## B14-L3 Reproduce a run from its record

---

### B14-L3 exercise 1

**Hints.** (1) What does a different hash mean? (2) Can you reproduce on different data? (3) Find the original.

**Answer.** b) Stop. You would be reproducing on different data. Find the original input. Because: r1) Any changed byte changes the hash, so the data is not the data the run used.

A hash mismatch means different data. Reproduction needs the original input, which is why raw inputs are kept immutable (lesson B6-L2).

### Common wrong answers.

- a: Same name, different content. The hash exists for exactly this check.
- c: Rewriting the record destroys the evidence of what the run used.
- d: Averaging runs on the wrong data does not recover the right result.
- b|r2: Right answer, wrong reason. The hash is reliable. That is why the mismatch matters.

### B14-L3 exercise 2

**Hints.** (1) What does + do with two strings? (2) And with two ints? (3) Concatenation, then addition.

#### Answer.

```
1 71 8
```

Unconverted parameters behave like text. Convert them back to their types before use.

### Common wrong answers.

- 8 8: '7' + '1' joins strings: '71'. Convert with int() first.

### B14-L3 exercise 3

**Hints.** (1) Which step provides the lock file? (2) Which check should stop you early? (3) Compare last.

#### Answer.

1. Check out the logged git commit
2. Build a fresh environment from the lock file at that commit
3. Load the input and verify its SHA-256 against input\_sha256
4. Convert the logged parameters back to their types
5. Rerun and compare the metric within a tolerance

Code, then environment, then data, then settings, then the comparison.

### Common wrong answers.

- The lock file that defines the environment lives at the commit. Check it out first.
- Verify the data before rerunning, or a mismatch wastes the run and confuses the comparison.

### B14-L3 exercise 4

**Hints.** (1) Find missing keys with a set difference. (2) Raise KeyError with the sorted list. (3) {key: cast(params[key]) for key, cast in types.items()}.

## Answer.

```
1 def parse_params(params: dict[str, str], types: dict[str, type]) ->
 dict[str, object]:
2 """Convert logged string parameters back to their types.
3
4 Args:
5 params: Parameters as MLflow returns them.
6 types: Expected keys and the type to convert each to.
7
8 Returns:
9 Typed parameters for the expected keys.
10
11 Raises:
12 KeyError: If an expected key is missing.
13 ValueError: If a value cannot be converted.
14 """
15 missing = sorted(set(types) - set(params))
16 if missing:
17 raise KeyError(f"run is missing parameters: {missing}")
18 return {key: cast(params[key]) for key, cast in types.items()}
```

Explicit types turn the run record back into the exact settings. Failing on a missing key beats silently rerunning with a default.

## B14-L3 exercise 5

**Hints.** (1) Is 0.004 rounding? (2) What does the record not contain? (3) Fix the record, not the check.

## Answer.

- A gap of 0.004 is a real difference, not rounding.
- Compare every recorded identity, then look for settings that were never logged.
- Log what was missing, so the next check passes for the right reason.

A failed reproduction points at a gap in the record. Close the gap, keep the check.

## Common wrong answers.

- verdict = fine: 0.004 AUC is larger than many model improvements you would ship. Something differs.
- next = relog: Editing the old run erases the evidence of what it produced.
- fix = widen: A wide tolerance hides incomplete records instead of fixing them.

## B15-L1 Sentence embeddings for open-ended answers

### B15-L1 exercise 1

**Hints.** (1) What defines the meaning of a direction? (2) The model that produced it. (3) Use one model for everything you compare.

**Answer.** b) No. Re-encode both years with the same model. Because: r1) Each model defines its own space. The same position means different things in each.

Re-encode all texts with one model and record which one.

#### Common wrong answers.

- a: Matching dimensions do not make the spaces the same.
- c: Normalising fixes length, not the meaning of directions.
- d: Averaging vectors from different spaces produces meaningless numbers.
- b|r2: Right answer, wrong reason. The issue is consistency, not quality.

### B15-L1 exercise 2

**Hints.** (1) Dot product:  $3 \times 1 + 4 \times 0$ . (2) Lengths: 5 and 1. (3)  $3 / 5$ .

#### Answer.

```
1 3.0 0.6
```

A dot product equals cosine similarity only for unit-length vectors. Normalise or divide by the lengths.

#### Common wrong answers.

- 0.6 0.6: The raw dot product is 3.0. Only after dividing by the lengths (5 and 1) do you get 0.6.

### B15-L1 exercise 3

**Hints.** (1) One length per row. (2) `keepdims` makes the shape broadcast. (3) 1 and `norms`.

#### Answer.

```
1 norms = np.linalg.norm(matrix, axis=1, keepdims=True)
2 unit = matrix / norms
```

`axis=1` gives each row's length, and `keepdims=True` keeps a column shape that divides row-wise. Guard rows with zero length.

## Common wrong answers.

- `g1 = 0`: `axis=0` gives one norm per column. Each row needs its own length: `axis=1`.

## B15-L1 exercise 4

**Hints.** (1) Divide dot products by both lengths. (2) Start similarities at `-inf`, fill only rows with non-zero length. (3) `np.argsort(-sims, kind="stable")[:k]`, then convert to `int`.

## Answer.

```
1 import numpy as np
2
3
4 def top_k_similar(query: np.ndarray, matrix: np.ndarray, k: int) ->
 list[int]:
5 """Row indices of the k rows most cosine-similar to query.
6
7 Args:
8 query: Vector of length d.
9 matrix: Array of shape (n, d).
10 k: Number of indices to return.
11
12 Returns:
13 Indices, most similar first. Zero-length rows rank last.
14
15 Raises:
16 ValueError: If query has zero length.
17 """
18 q_norm = np.linalg.norm(query)
19 if q_norm == 0:
20 raise ValueError("query vector has zero length")
21 norms = np.linalg.norm(matrix, axis=1)
22 sims = np.full(len(matrix), -np.inf)
23 ok = norms > 0
24 sims[ok] = (matrix[ok] @ query) / (norms[ok] * q_norm)
25 return [int(i) for i in np.argsort(-sims, kind="stable")[:k]]
```

Cosine similarity ignores length, so a long vector pointing elsewhere does not win. Zero rows are kept but ranked last.

## B15-L1 exercise 5

**Hints.** (1) One space for all comparisons. (2) What makes a cached vector reusable? (3) Mixing models breaks comparisons.

## Answer.

- One multilingual model puts all languages in one space. Test it on your data.



- Store the identity of text and model, so vectors can be reused and audited.
- Reuse vectors and keep the model fixed. Changing model means re-encoding everything.

Fix the model, record identities, and reuse vectors only within one model.

### Common wrong answers.

- `model = per_lang`: Different models give different spaces. Answers in different languages could not be compared.
- `model = english`: An English-only model maps other languages poorly. Test before you trust it.
- `update = newest`: New and old vectors would come from different spaces.
- `store = vector_only`: Without the model name and text hash, nobody knows which vectors can be reused.

## B15-L2 HDBSCAN, clusters with a noise label

---

### B15-L2 exercise 1

**Hints.** (1) What does -1 mean in HDBSCAN? (2) Where do those points sit? (3) In sparse regions.

**Answer.** b) Those answers are in sparse regions and share no dense theme with enough others. Because: r1) -1 is the noise label for points that no dense cluster claims.

Noise answers are real answers without a shared theme. Report their share and read some of them.

### Common wrong answers.

- a: Noise is a result, not a failure.
- c: -1 is not a cluster. Counting it as one mixes unrelated answers.
- d: Duplicates are close together, so they would usually form or join a cluster.
- b|r2: Right answer, wrong reason. The noise share depends on the data and the settings.

### B15-L2 exercise 2

**Hints.** (1) Remove -1 before counting clusters. (2) Noise: 2 of 10. (3) Largest cluster: 5 of the 8 clustered points.

**Answer.**

```
1 2 0.2 0.625
```

Two clusters, 20% noise, and the largest cluster holds 62.5% of clustered answers.

### Common wrong answers.

- 3 0.2 0.5: -1 is noise, not a cluster. Two clusters, and the largest (5 of 8) is 0.625 of clustered points.

### B15-L2 exercise 3

**Hints.** (1) The smallest group you would call a theme. (2) Copy the input instead of changing it. (3) `min_cluster_size` and `True`.

### Answer.

```
1 labels = HDBSCAN(min_cluster_size=20, copy=True).fit_predict(vectors)
```

`min_cluster_size` sets the smallest theme. `copy=True` protects your array and silences the scikit-learn 1.8 warning.

### Common wrong answers.

- `g1 = n_clusters`: HDBSCAN has no `n_clusters`. The main setting is `min_cluster_size`.

### B15-L2 exercise 4

**Hints.** (1) Handle the empty array first. (2)

`clustered = labels[labels != -1]`. (3) `np.unique(clustered, return_counts=True)` gives cluster sizes.

### Answer.

```
1 import numpy as np
2
3
4 def summarise_labels(labels: np.ndarray) -> dict[str, float | int | None]:
5 """Cluster count, noise share and largest cluster share.
6
7 Args:
8 labels: Cluster labels, -1 for noise.
9
10 Returns:
11 n_clusters, noise_share (of all points) and largest_share (of
12 clustered points).
13 """
14 n = len(labels)
15 if n == 0:
16 return {"n_clusters": 0, "noise_share": None, "largest_share":
17 None}
18 clustered = labels[labels != -1]
19 sizes = np.unique(clustered, return_counts=True)[1]
```

```

18 return {
19 "n_clusters": int(len(sizes)),
20 "noise_share": float((labels == -1).mean()),
21 "largest_share": float(sizes.max() / len(clustered)) if
 len(clustered) > 0 else None,
22 }

```

Noise is reported, not clustered. Guarded denominators keep all-noise and empty results honest.

## B15-L2 exercise 5

**Hints.** (1) What does an unbalanced result tell you? (2) Which setting controls the smallest theme? (3) Reproducibility (lesson B6-L2).

### Answer.

- Read before tuning: the samples show whether the big cluster is one theme or a mix.
- Adjust `min_cluster_size` (and `min_samples`) on purpose, one change at a time.
- Record every setting, so the themes can be reproduced next quarter.

Diagnose with samples, tune deliberately, and record everything that shaped the clusters.

### Common wrong answers.

- `diagnose = publish`: One dominant cluster and 40 tiny ones rarely reflect real themes. Look first.
- `tune = kmeans`: k-means forces all 45% noise into themes, inflating their shares.
- `record = final`: Without the settings, nobody can rebuild or compare the themes.

## B15-L3 Validate clusters before you name them

---

### B15-L3 exercise 1

**Hints.** (1) Would the same clusters appear on a new sample? (2) Can people name them? (3) Three checks, not one.

**Answer.** b) No. Also check stability across reruns and read samples from each cluster. Because: r1) Separation in one run says nothing about whether clusters recur or mean anything.

Separation, stability and meaning together. One number is a start, not a conclusion.

### Common wrong answers.

- a: No threshold turns one number into proof. Check stability and meaning too.
- c: 1 is almost never reached on real data. Good separation is not the only test.
- d: A dashboard layout is not evidence about the data.
- b|r2: Right answer, wrong reason. Silhouette is useful. It is not sufficient.

### B15-L3 exercise 2

**Hints.** (1) Does ARI care what the labels are called? (2) Are the groupings the same in the first case? (3) The second is worse than chance.

#### Answer.

```
1 1.0 -0.5
```

ARI compares groupings, not names. It can go below 0 when agreement is worse than chance.

### Common wrong answers.

- 0.0 -0.5: ARI ignores label names. Swapping 0 and 1 gives the same grouping: 1.0.
- 1.0 0.0: The second grouping splits every pair the first one keeps together, which scores below chance: -0.5.

### B15-L3 exercise 3

**Hints.** (1) Exclude the noise label. (2) Use the same mask for labels. (3) `!=` and `keep`.

#### Answer.

```
1 keep = labels != -1
2 score = silhouette_score(X[keep], labels[keep])
```

One mask for both arrays keeps points and labels aligned.

### Common wrong answers.

- `g1 ==`: That keeps only the noise. Keep points whose label is not -1.

### B15-L3 exercise 4

**Hints.** (1) Mask out -1 first. (2) Count clusters among the remaining labels. (3) `silhouette_score` needs  $2 \leq \text{clusters} \leq n - 1$ .

#### Answer.

```
1 import numpy as np
2 from sklearn.metrics import silhouette_score
```

```

3
4
5 def silhouette_without_noise(X: np.ndarray, labels: np.ndarray) -> float
 | None:
6 """Silhouette over clustered points only.
7
8 Args:
9 X: Feature matrix.
10 labels: Cluster labels, -1 for noise.
11
12 Returns:
13 The silhouette score, or None when it is undefined.
14 """
15 keep = labels != -1
16 kept = labels[keep]
17 n_clusters = len(np.unique(kept))
18 if n_clusters < 2 or n_clusters >= len(kept):
19 return None
20 return float(silhouette_score(X[keep], kept))

```

Noise is not a cluster, so it stays out. Guarding the cluster count turns an exception into a clear "undefined".

## B15-L3 exercise 5

**Hints.** (1) What does noise do to silhouette? (2) Would the themes come back on new data? (3) Which examples are honest?

### Answer.

- Silhouette on clustered points, with noise reported beside it.
- Stability across reruns, naming the themes that do not recur.
- Random verbatims, so readers see what a theme typically contains.

Honest methods show separation, stability and representative evidence.

### Common wrong answers.

- separation = `sil_all`: Noise is scattered by definition. Counting it as a cluster drags the score down and mixes two messages.
- stability = `skip`: A clustering can be well separated in one run and different in the next.
- examples = `best`: Hand-picked examples make every theme look cleaner than it is.

## B16-L1 Choose the device once, move model and data to it

---

### B16-L1 exercise 1

**Hints.** (1) Is `Tensor.to` in place? (2) What happens to an unassigned return value? (3) Reassign.

**Answer.** b) Where it was before. `.to()` on a tensor returns a new tensor. Because: r1) `Tensor.to` returns a new tensor. Only `nn.Module.to` changes the object in place.

`x = x.to(device)`. For modules, `model.to(device)` is in place and returns the same model.

#### Common wrong answers.

- a: That is true for models, not tensors. Write `x = x.to(device)`.
- c: The new tensor is discarded because it is not assigned.
- d: `x` is unchanged. The moved copy is what disappears.
- b|r2: Right answer, wrong reason. `.to()` works on tensors, it just does not modify them.

### B16-L1 exercise 2

**Hints.** (1) Which call is in place? (2) `Module.to`. (3) `Tensor.to` returns a copy.

**Answer.**

```
1 cpu meta
```

The asymmetry catches everyone once. Reassign tensors, and do not rely on reassigning models.

#### Common wrong answers.

- `meta meta`: `Tensor.to` returns a new tensor. `t` itself stays on the CPU.
- `cpu cpu`: `Module.to` moves parameters in place, so the model is on meta.

### B16-L1 exercise 3

**Hints.** (1) The device-agnostic method. (2) NumPy needs CPU memory. (3) `to` and `cpu`.

**Answer.**

```
1 x = x.to(device)
2 preds = model(x).detach().cpu().numpy()
```

`.to(device)` keeps code portable. `.cpu()` before `.numpy()` avoids an error on GPU tensors.

## Common wrong answers.

- `g1 = cuda: .cuda()` hardcodes the GPU. `.to(device)` works on any machine.

## B16-L1 exercise 4

**Hints.** (1) Check `prefer_gpu` in every GPU branch. (2) CUDA before MPS. (3) Fall through to `cpu`.

## Answer.

```
1 def pick_device_name(prefer_gpu: bool, cuda_ok: bool, mps_ok: bool) ->
 str:
2 """Name of the device to use: cuda, then mps, then cpu.
3
4 Args:
5 prefer_gpu: False forces the CPU, for tests and CI.
6 cuda_ok: Result of torch.cuda.is_available().
7 mps_ok: Result of torch.backends.mps.is_available().
8
9 Returns:
10 "cuda", "mps" or "cpu".
11 """
12 if prefer_gpu and cuda_ok:
13 return "cuda"
14 if prefer_gpu and mps_ok:
15 return "mps"
16 return "cpu"
```

Passing availability in makes the decision a pure function (lesson B2-L1): easy to test on a laptop with no GPU.

## B16-L1 exercise 5

**Hints.** (1) Which line tries to move `x`? (2) Does `Tensor.to` change `x`? (3) Reassign it.

**Answer.** The bug is on line 12.

Fixed version:

```
1 from collections.abc import Iterable
2
3 import torch
4
5
6 def validate(model: torch.nn.Module, loader: Iterable[tuple[torch.Tensor,
 torch.Tensor]], device: torch.device) -> float:
7 """Accuracy on the validation loader."""
8 model.eval()
9 correct, total = 0, 0
10 with torch.inference_mode():
11 for x, y in loader:
12 x = x.to(device)
```

```

13 preds = model(x).argmax(dim=1)
14 correct += int((preds.cpu() == y).sum())
15 total += len(y)
16 return correct / total if total > 0 else float("nan")

```

On the CPU, `x` and the model already share a device, so the bug hides. On the GPU, `x` never moved. Reassign: `x = x.to(device)`.

### Common wrong answers.

- Line 14: `preds.cpu()` is right: `y` stays on the CPU, so the comparison happens there.
- Line 8: `model.eval()` is correct for validation.

## B16-L2 A training loop that does what you think

### B16-L2 exercise 1

**Hints.** (1) Does backward overwrite or add? (2) It adds. (3) `zero_grad` resets.

**Answer.** b) Gradients accumulate across batches, so steps use the sum of all earlier gradients. Because: r1) `backward()` adds into `.grad`. Only `zero_grad()` resets it.

Call `optimizer.zero_grad()` before each backward pass.

### Common wrong answers.

- a: PyTorch accumulates on purpose, for gradient accumulation. Clearing is your job.
- c: No error: training runs with growing, wrong updates.
- d: It keeps changing, with updates that grow each batch.
- b|r2: Right answer, wrong reason. The optimizer uses the accumulated gradients.

### B16-L2 exercise 2

**Hints.** (1) The gradient of  $2w$  is 2. (2) How many backward calls? (3) They add up.

**Answer.**

```
1 6.0
```

Three accumulations of 2.0. With `zero_grad()` each iteration it would print 2.0.



### Common wrong answers.

- 2.0: Each backward adds 2.0 to w.grad. Three calls without zeroing give 6.0.

### B16-L2 exercise 3

**Hints.** (1) Data first. (2) Clear gradients before computing new ones. (3) Compute, then apply.

#### Answer.

1. `x, y = x.to(device), y.to(device)`
2. `optimizer.zero_grad()`
3. `loss = loss_fn(model(x), y)`
4. `loss.backward()`
5. `optimizer.step()`

Move, clear, forward, backward, step. `zero_grad` can also sit at the end of the iteration, as long as it runs between steps.

### Common wrong answers.

- `zero_grad` after backward erases the gradients the step needs.
- step before backward uses stale or empty gradients.

### B16-L2 exercise 4

**Hints.** (1) Split the list: before the window and the window. (2) Compare the window's best with the earlier best. (3) Subtract `min_delta` from the earlier best.

#### Answer.

```
1 def should_stop(val_losses: list[float], patience: int, min_delta: float
 = 0.0) -> bool:
2 """Early-stopping rule on validation losses, one per epoch.
3
4 Args:
5 val_losses: Validation loss per epoch, oldest first.
6 patience: Epochs allowed without improvement.
7 min_delta: Smallest decrease that counts as improvement.
8
9 Returns:
10 True when training should stop.
11
12 Raises:
13 ValueError: If patience is below 1.
14 """
15 if patience < 1:
16 raise ValueError("patience must be at least 1")
```

```

17 if len(val_losses) <= patience:
18 return False
19 best_before = min(val_losses[:-patience])
20 return min(val_losses[-patience:]) > best_before - min_delta

```

Comparing best against best ignores noisy single epochs. `min_delta` stops training from continuing on improvements too small to matter.

## B16-L2 exercise 5

**Hints.** (1) What does `CrossEntropyLoss` expect as input? (2) What does forward return? (3) Logits, not probabilities.

**Answer.** The bug is on line 13.

Fixed version:

```

1 import torch
2 from torch import nn
3
4
5 class Classifier(nn.Module):
6 """Two-layer classifier. Train with nn.CrossEntropyLoss."""
7
8 def __init__(self, n_in: int, n_classes: int) -> None:
9 super().__init__()
10 self.net = nn.Sequential(nn.Linear(n_in, 32), nn.ReLU(),
11 nn.Linear(32, n_classes))
12
13 def forward(self, x: torch.Tensor) -> torch.Tensor:
14 return self.net(x)

```

Softmax squeezes outputs into `[0, 1]`, and `CrossEntropyLoss` then applies log-softmax again. The loss and its gradients are wrong. Apply softmax only when you need probabilities for reporting.

## Common wrong answers.

- Line 10: The architecture is fine. The output passed to the loss is the problem.

## B16-L3 Reproducible PyTorch runs and safe checkpoints

### B16-L3 exercise 1

**Hints.** (1) What does weights-only loading allow? (2) What does a state dict contain? (3) Tensors and plain values.

**Answer.** b) Save `model.state_dict()` instead, rebuild the model class in code, and load the weights into it. Because: r1) State dicts hold only tensors and plain values, which load safely under the weights-only default.

Save `state_dict()`, keep the model class in version-controlled code, and load weights into a fresh instance.

### Common wrong answers.

- `a`: That loads arbitrary pickles, which can run code. Only do it for files you created and trust, and prefer fixing how you save.
- `c`: Downgrading removes a safety default and freezes you on an old version.
- `d`: The extension changes nothing about the content.
- `b|r2`: Right answer, wrong reason. Objects can be saved. Loading them safely is the problem.

### B16-L3 exercise 2

**Hints.** (1) Which generator does `torch.rand` use here? (2) Do `g1` and `g2` share state with it? (3) No.

### Answer.

```
1 True
```

Dedicated generators isolate the shuffle order from other random calls. That is why `DataLoaders` get their own.

### Common wrong answers.

- `False`: `torch.rand(100)` uses the global generator. `g1` and `g2` are separate and unaffected.

### B16-L3 exercise 3

**Hints.** (1) The Generator method that seeds it. (2) The `DataLoader` argument that takes it. (3) `manual_seed` and `generator`.

### Answer.

```
1 loader_gen = torch.Generator().manual_seed(seed)
2 loader = DataLoader(ds, batch_size=64, shuffle=True, generator=loader_gen)
```

`torch.Generator().manual_seed(seed)` returns the generator, ready to pass as `generator=`.

### Common wrong answers.

- `g2 = seed`: `DataLoader` has no seed argument. Pass a `torch.Generator` as `generator=`.

## B16-L3 exercise 4

**Hints.** (1) Three global generators: random, NumPy, torch. (2) One dedicated generator for the loader. (3) `return torch.Generator().manual_seed(seed)`.

**Answer.**

```
1 import random
2
3 import numpy as np
4 import torch
5
6
7 def seed_everything(seed: int) -> torch.Generator:
8 """Seed Python, NumPy and PyTorch global state at the entry point.
9
10 Args:
11 seed: The run's seed, logged with the run.
12
13 Returns:
14 A seeded generator to pass to shuffling DataLoaders.
15 """
16 random.seed(seed)
17 np.random.seed(seed)
18 torch.manual_seed(seed)
19 return torch.Generator().manual_seed(seed)
```

Seeding every global generator at the entry point covers code you do not control. The returned generator keeps batch order independent of everything else.

## B16-L3 exercise 5

**Hints.** (1) What must be restored to continue exactly? (2) Which two checkpoints serve different needs? (3) Keep the safe default.

**Answer.**

- State dicts plus epoch, RNG state and best score: everything needed to resume exactly.
- Save the last epoch for resuming, and the best epoch for release.
- Load with the safe default into a model defined in code.

Resume from state, release the best epoch, and never unpickle files you did not create.

**Common wrong answers.**

- `content = object`: A pickled object ties the file to the exact code layout and fails the weights-only default.

- `when = end`: An interruption at hour 11 loses everything, and the last epoch is rarely the best.
- `load = unsafe`: Unpickling files from shared storage can run arbitrary code.

## B17-L1 Unit tests for data code that catch real bugs

### B17-L1 exercise 1

**Hints.** (1) Which buggy values would each assertion let through? (2) Is 0.0 falsy? (3) Only one rejects 0.0.

**Answer.** b) `assert share(3, 0) is None` Because: r1) It fails for every result except None, including 0.0.

Write assertions that pass only for the correct value.

#### Common wrong answers.

- a: 0.0 is falsy too, so a buggy 0.0 would pass.
- c: That asserts the bug.
- d: Almost any wrong value passes.
- b|r2: Right answer, wrong reason. Precision matters, not speed.

### B17-L1 exercise 2

**Hints.** (1) What are the dtypes? (2) Does `==` compare dtypes? (3) `assert_frame_equal` does.

**Answer.**

```
1 False True
2 AssertionError
```

`==` compares values only. `assert_frame_equal` also checks dtypes, which catches a column that silently became float.

#### Common wrong answers.

- `True True equal`: Values match, but dtypes differ (int64 and float64). `assert_frame_equal` checks dtypes by default.

### B17-L1 exercise 3

**Hints.** (1) The pytest marker for a table of cases. (2) Float comparison with tolerance. (3) `parametrize` and `approx`.

**Answer.**

```

1 @pytest.mark.parametrize(("part", "total", "expected"), [(1, 4, 0.25),
 (0, 5, 0.0), (5, 5, 1.0)])
2 def test_share_values(part, total, expected):
3 assert share(part, total) == pytest.approx(expected)

```

`parametrize` runs the test once per row. `approx` compares floats with a sensible tolerance.

### Common wrong answers.

- `g1 = parameterize`: `pytest` spells it `parametrize`, without the extra `e`.

### B17-L1 exercise 4

**Hints.** (1) A test must fail for wrong values, not only for crashes. (2) Compare with the exact expected value: `== 0.25`. (3) Test for `None` with `is None`.

### Answer.

```

1 def share(part: float, total: float) -> float | None:
2 """part / total, or None when total is 0."""
3 if total == 0:
4 return None
5 return part / total
6
7
8 def test_share_basic() -> None:
9 assert share(1, 4) == 0.25
10
11
12 def test_share_zero_total() -> None:
13 assert share(3, 0) is None

```

The checker plays the role of a mutation test: it breaks the code and checks your tests notice. Precise assertions catch both bugs.

### B17-L1 exercise 5

**Hints.** (1) Where does the expected value come from? (2) Which inputs break data code? (3) Floating point.

### Answer.

- Tiny explicit inputs with hand-computed answers.
- Every edge the function promises to handle.
- `approx` for floats that come from arithmetic.

Independent expected values, every edge, and tolerant float comparison make tests that catch bugs.

### Common wrong answers.

- `normal = sample`: An expected value produced by the function under test only checks that it agrees with itself.
- `edges = skip`: Line coverage is not behaviour coverage. Edge branches are where bugs ship.
- `compare = exact`:  $0.1 + 0.2 \neq 0.3$  in floating point. Use `approx`.

## B17-L2 Property-based tests with Hypothesis

---

### B17-L2 exercise 1

**Hints.** (1) What makes a failure easy to debug? (2) A small input. (3) Shrinking.

**Answer.** b) A shrunk input: the simplest example it found that still fails. Because: r1) Shrinking turns a messy random failure into a minimal case you can read.

Shrinking reduced a weighted-mean failure to one value with one zero weight.

### Common wrong answers.

- `a`: Hypothesis keeps simplifying the failing input before reporting it.
- `c`: It reports one minimal example, not the whole search.
- `d`: The failing example is shown, such as `pairs=[(0.0, 0.0)]`.
- `b|r2`: Right answer, wrong reason. It generates complex inputs, then shrinks the failure.

### B17-L2 exercise 2

**Hints.** (1) What special values can `st.floats()` produce? (2) Is `nan` equal to itself? (3) No.

**Answer.**

```
1 fail
```

The test fails on `x=nan`. Bound the strategy or pass `allow_nan=False` unless the function must handle `nan`.

### Common wrong answers.

- `pass`: `st.floats()` generates `nan` by default, and `nan == nan` is `False`.

### B17-L2 exercise 3

**Hints.** (1) The decorator that configures a test. (2) The floats argument for nan. (3) settings and allow\_nan.

**Answer.**

```
1 @settings(max_examples=500)
2 @given(st.floats(min_value=-1e6, max_value=1e6, allow_nan=False))
3 def test_round_trip(x):
4 assert parse(format(x)) == x
```

@settings(max\_examples=500) raises the search budget. allow\_nan=False keeps nan out when the function does not promise to handle it.

**Common wrong answers.**

- g2 = allow\_infinity: Bounds already exclude infinities. The gap is for nan.

### B17-L2 exercise 4

**Hints.** (1) Decorate with @given and the strategy from the prompt. (2) The test receives rows as its argument. (3) Compare dedupe applied twice with dedupe applied once.

**Answer.**

```
1 from hypothesis import given, strategies as st
2
3
4 @given(st.lists(st.tuples(st.sampled_from(["a", "b", "c"]),
5 st.integers(0, 5))))
6 def test_dedupe_idempotent(rows: list[tuple[str, int]]) -> None:
7 assert dedupe(dedupe(rows)) == dedupe(rows)
```

A small id alphabet and small timestamps make duplicates and ties likely. Idempotence is one property. Pair it with a grain property, one row per id, which catches a dedupe that keeps ties: that bug is idempotent, so idempotence alone misses it.

### B17-L2 exercise 5

**Hints.** (1) What must never change? (2) What must come back? (3) What can be compared?

**Answer.**

- Totals do not depend on row order.
- Saving and loading must not lose data or dtypes.
- A weighted mean stays within the range of its inputs.
- A rewrite must agree with the code it replaces.



Good properties come from what the function promises: invariance, round trips, bounds and agreement.

### Common wrong answers.

- `total = bounds`: A total can exceed every input value. Order invariance is the property that holds.
- `io = invariance`: For I/O the key promise is that nothing is lost: a round trip.
- `fast = bounds`: Bounds would pass many wrong rewrites. Compare with the trusted version.

## B17-L3 Measure before you optimise

---

### B17-L3 exercise 1

**Hints.** (1) What is the most a speed-up of each part can save? (2) Compare 36 with 1. (3) Follow the profile.

**Answer.** b) File reading. Because: r1) A 1-minute function can save at most 1 minute. Reading holds 90% of the time.

Fix the biggest measured cost first. Guesses about "slow loops" waste days.

### Common wrong answers.

- `a`: Even an infinitely fast weighting function saves only 1 minute of 40.
- `c`: Effort should follow cost. One of them holds 36 of 40 minutes.
- `d`: Hardware is a last resort once the code's own costs are known.
- `b|r2`: Right answer, wrong reason. Here the profile says I/O. Elsewhere it may not.

### B17-L3 exercise 2

**Hints.** (1) What does repeat control? (2) What does number control? (3) One total per round.

### Answer.

|     |
|-----|
| 1 5 |
|-----|

Five rounds, each the total time for 100 calls. Report `min(rounds) / 100` as the per-call time.

### Common wrong answers.

- `100`: number is the calls per round. repeat is the number of rounds returned.

- 500: repeat returns one total per round, not one per call.

### B17-L3 exercise 3

**Hints.** (1) Time including everything a function calls. (2) The Stats method that prints rows. (3) `cumulative` and `print_stats`.

**Answer.**

```
1 cProfile.run("main()", "profile.out")
2 pstats.Stats("profile.out").sort_stats("cumulative").print_stats(10)
```

Cumulative time points to the high-level function that owns the cost. Then drill into it.

**Common wrong answers.**

- `g1 = calls`: Sorting by call count shows busy functions, not expensive ones. Sort by cumulative time.

### B17-L3 exercise 4

**Hints.** (1) Validate repeat and number first. (2) Time each round with `perf_counter` before and after the inner loop. (3) Keep the minimum of `round_time / number`.

**Answer.**

```
1 import time
2 from collections.abc import Callable
3
4
5 def best_time(fn: Callable[[], object], repeat: int = 5, number: int =
6 10) -> float:
7 """Best per-call time over several rounds.
8 Args:
9 fn: Function to time, called with no arguments.
10 repeat: Number of rounds.
11 number: Calls per round.
12
13 Returns:
14 The smallest per-call time in seconds.
15
16 Raises:
17 ValueError: If repeat or number is below 1.
18 """
19 if repeat < 1 or number < 1:
20 raise ValueError("repeat and number must be at least 1")
21 best = float("inf")
22 for _ in range(repeat):
23 start = time.perf_counter()
24 for _ in range(number):
```

```

25 fn()
26 best = min(best, (time.perf_counter() - start) / number)
27 return best

```

The minimum over rounds estimates the code's own cost, because interference only adds time. Dividing by `number` gives a per-call figure you can compare.

## B17-L3 exercise 5

**Hints.** (1) What runs between start and the end of the round? (2) Is reading the file part of `add_features`? (3) Move setup out of the timed region.

**Answer.** The bug is on line 11.

Fixed version:

```

1 import time
2
3 import pandas as pd
4
5
6 def benchmark(path: str, rounds: int) -> float:
7 """Best time of add_features over rounds."""
8 df = pd.read_parquet(path)
9 best = float("inf")
10 for _ in range(rounds):
11 start = time.perf_counter()
12 add_features(df)
13 best = min(best, time.perf_counter() - start)
14 return best

```

The file read is setup, and it dominates the timing. Read once before the loop, then time only the function.

## Common wrong answers.

- Line 13: Taking the minimum is right. The timed region includes more than the function.

## B18-L1 Docstrings and READMEs that answer the next person's questions

### B18-L1 exercise 1

**Hints.** (1) What would surprise a caller? (2) What can the signature not say? (3) Edge behaviour.

**Answer.** b) "Returns None when the weights sum to 0." Because: r1) It tells callers about a behaviour the signature and name cannot show.

Document what a caller must know and cannot see: edge cases, units, grain.

**Common wrong answers.**

- `a`: The code already shows that. Docstrings explain what the code cannot.
- `c`: Ownership belongs in the README or CODEOWNERS, not every docstring.
- `d`: It adds nothing.
- `b|r2`: Right answer, wrong reason. Value comes from information, not length.

**B18-L1 exercise 2**

**Hints.** (1) Where does Python store a function's docstring? (2) In **doc**. (3) It prints the string.

**Answer.**

```
1 part / total, or None when total is 0.
```

The first string in the body is the docstring. `help()`, editors and documentation tools read it from `__doc__`.

**Common wrong answers.**

- `None`: A string literal as the first statement of a function becomes its **doc**.

**B18-L1 exercise 3**

**Hints.** (1) What does a newcomer need first? (2) Then how to run it. (3) Details after.

**Answer.**

1. Purpose and users
2. Quick start: install and run commands
3. Inputs and outputs, with grain and schedule
4. Configuration
5. Tests and checks
6. Owners and support

Purpose, run, data, settings, checks, contacts: the order a newcomer's questions arrive in.

**Common wrong answers.**

- Say what the project is before telling people to run it.
- The quick start is what most readers came for. Put it near the top.

## B18-L1 exercise 4

**Hints.** (1) Summary line, blank line, then Args:, Returns:, Raises:. (2) Indent each entry under its section. (3) Mention missing weights under weights.

**Answer.**

```
1 import pandas as pd
2
3
4 def cap_weights(weights: pd.Series, cap: float) -> pd.Series:
5 """Cap survey weights at a maximum value.
6
7 Args:
8 weights: One weight per respondent. Missing weights stay missing.
9 cap: Largest allowed weight, at least 1.
10
11 Returns:
12 A new Series with every weight above cap set to cap.
13
14 Raises:
15 ValueError: If cap is below 1.
16 """
17 if cap < 1:
18 raise ValueError("cap must be at least 1")
19 return weights.clip(upper=cap)
```

The docstring now tells a caller the bounds, the edge behaviour and the failure mode, none of which the signature shows.

## B18-L1 exercise 5

**Hints.** (1) Who reads it, and where? (2) Will the value change? (3) How local is the explanation?

**Answer.**

- Behaviour a caller must handle: the docstring.
- How to run the project: the README.
- Why a specific line is unusual: an inline comment next to it.
- A setting that may change: config (lesson A1-L3), described in the README.

Put information where its reader looks, and settings where they can change without editing text.

**Common wrong answers.**

- none\_meaning = `readme`: Callers read the function's docstring, often in their editor, not the README.

- `threshold = docstring`: A value that can change belongs in config. Hardcoding it in prose invites drift.
- `why_line = docstring`: A workaround on one line is explained best right beside that line.

## B18-L2 Model cards that state limits as clearly as scores

### B18-L2 exercise 1

**Hints.** (1) Which section says 'do not use for'? (2) Reuse is a scope question. (3) Out of scope.

**Answer.** b) Out-of-scope uses Because: r1) It names the uses the model was not built or validated for.

Out-of-scope uses turn assumptions in the builders' heads into a rule readers see.

#### Common wrong answers.

- a: Details identify the model. They do not say where it must not be used.
- c: Evaluation shows how it performs on its intended task, not on others.
- d: Monitoring tracks the model in use. It does not stop misuse.
- b|r2: Right answer, wrong reason. Length does not matter. Its purpose does.

### B18-L2 exercise 2

**Hints.** (1) Which keys are absent? (2) Which are present but blank? (3) Both count.

**Answer.**

```
1 ['out_of_scope', 'subgroups', 'limitations']
```

Empty sections are as unhelpful as missing ones. A CI check should reject both.

#### Common wrong answers.

- `['out_of_scope', 'subgroups']`: limitations exists but is empty. The check treats blank text as missing.

### B18-L2 exercise 3

**Hints.** (1) The baseline for average precision. (2) How much evidence stands behind the number? (3) Prevalence and row count.

**Answer.**

```
1 Evaluation: AUC 0.78, average precision 0.21 (prevalence 2.4%), test Oct
 to Dec 2026, 51,000 rows
```

Prevalence gives the baseline and the row count gives the evidence behind the metric.

### Common wrong answers.

- `g1 = accuracy`: Accuracy is another metric. Average precision needs the prevalence as its baseline.

### B18-L2 exercise 4

**Hints.** (1) Loop over required, not over card. (2) `card.get(name, "")` handles absent keys. (3) Use `.strip()` to treat whitespace as empty.

### Answer.

```
1 def missing_sections(card: dict[str, str], required: list[str]) ->
 list[str]:
2 """Required model card sections that are absent or blank.
3
4 Args:
5 card: Section name to text.
6 required: Section names that must have content.
7
8 Returns:
9 Missing section names, in the order of required.
10 """
11 return [name for name in required if not card.get(name, "").strip()]
```

A small, strict check turns the template into a gate: no model ships with an empty limitations section.

### B18-L2 exercise 5

**Hints.** (1) What must users know to use the model safely? (2) What happens to hidden weaknesses? (3) Who reviews, and when?

### Answer.

- Per-group rates with counts, so the weakest group is visible.
- The gap stated plainly with a mitigation.
- A team and a date that make review happen.

A card is useful when it makes weaknesses and responsibilities impossible to miss.

### Common wrong answers.

- `subgroups = overall`: An overall rate hides a group at half the recall (lesson B12-L3).

- limit = omit: Hiding a known weakness moves the harm onto the people the model serves.
- review = adhoc: Without a date and an owner, reviews do not happen.

## B18-L3 Decision logs, so the reasons outlive the meeting

---

### B18-L3 exercise 1

**Hints.** (1) Why keep past reasoning? (2) How would a reader know it no longer applies? (3) Supersede and link.

**Answer.** c) Add a new entry, and mark the old one superseded with a link to the new one. Because: r1) The log keeps the history of reasoning. Each entry records what was known when it was written.

Supersede, never rewrite. The chain of entries shows how and why thinking changed.

#### Common wrong answers.

- a: Editing erases why the first choice made sense at the time.
- b: Deleting loses the history that explains past results.
- d: Without links, readers find the old entry and think it still applies.
- c|r2: Right answer, wrong reason. The point is preserving reasoning, not storage.

### B18-L3 exercise 2

**Hints.** (1) Which entries are accepted? (2) Does the superseded entry point to an existing one? (3) Yes: D-007.

#### Answer.

```
1 ['D-007'] True
```

Current decisions are the accepted ones. Every superseded entry should point to an existing replacement.

#### Common wrong answers.

- ['D-003', 'D-007'] True: D-003 is superseded, so it is not a current decision.

### B18-L3 exercise 3

**Hints.** (1) Problem first. (2) Then the choices. (3) Then the pick and its cost.



## Answer.

1. ID and title
2. Date and status
3. Context: the problem and the evidence
4. Options considered
5. Decision
6. Consequences and review point

Context, options, decision, consequences: the order of the reasoning itself.

## Common wrong answers.

- State the problem and options before the choice, so the reader can judge it.
- Options come before the decision that chooses among them.

## B18-L3 exercise 4

**Hints.** (1) Use `entry.get(...)` so missing keys do not raise. (2) Check the rules in the order given and append a message for each failure. (3) For superseded entries, test membership in `log_ids`.

## Answer.

```
1 VALID_STATUS = {"proposed", "accepted", "superseded"}
2
3
4 def decision_problems(entry: dict[str, object], log_ids: set[str]) ->
 list[str]:
5 """Problems with one decision log entry, in rule order.
6
7 Args:
8 entry: Parsed decision entry.
9 log_ids: IDs of all entries in the log.
10
11 Returns:
12 Problem messages. Empty when the entry is complete.
13 """
14 problems = []
15 status = entry.get("status")
16 if status not in VALID_STATUS:
17 problems.append("invalid status")
18 options = entry.get("options")
19 if not isinstance(options, list) or len(options) < 2:
20 problems.append("fewer than two options")
21 consequences = entry.get("consequences")
22 if not isinstance(consequences, str) or not consequences.strip():
23 problems.append("no consequences")
24 if status == "superseded" and entry.get("superseded_by") not in
 log_ids:
```

```
25 problems.append("superseded without a valid replacement")
26 return problems
```

A validator in CI keeps the log useful: every entry has real options and consequences, and every superseded entry leads somewhere.

### B18-L3 exercise 5

**Hints.** (1) Where does the next analyst look? (2) When are the reasons clearest? (3) How does a reader meet the decision?

**Answer.**

- In the repository, versioned with the code it explains.
- Written with the change, while the reasons are fresh.
- Linked from where people meet the decision.

Records are useful when they are versioned, timely and linked.

**Common wrong answers.**

- where = `slides`: Slides are hard to find, hard to search and not versioned with the code.
- when = `later`: By the end, the options and reasons are half forgotten.
- link = `search`: People do not search for a record they do not know exists. Link it from the value.

## C1-L1 Tenant, capacity, workspace and item

---

### C1-L1 exercise 1

**Hints.** (1) What does a capacity provide? (2) Compute, measured in a unit. (3) Capacity units.

**Answer.** b) The capacity has 32 capacity units (CUs) of compute. Because: r1) For F SKUs the number equals the capacity units: F2 is 2 CUs, F64 is 64. F SKUs are named by their capacity units. F32 provides 32 CUs.

**Common wrong answers.**

- a: Capacities are not sized by user count. The number is compute in CUs.
- c: OneLake storage is billed separately from capacity compute.
- d: The SKU does not cap workspaces in this way. It measures compute.
- b|r2: Right answer, wrong reason. The number is exactly the CU count.

## C1-L1 exercise 2

**Hints.** (1) Which level holds organisation-wide settings? (2) Which level provides compute to workspaces? (3) Items live in workspaces.

**Answer.**

1. Tenant
2. Capacity
3. Workspace
4. Item, such as a lakehouse

Tenant settings apply to all capacities. Each workspace draws compute from its capacity. Items live in workspaces.

**Common wrong answers.**

- A workspace is assigned to a capacity, which provides its compute. Capacity sits above workspace.
- Capacities belong to a tenant, where tenant settings apply to everything.

## C1-L1 exercise 3

**Hints.** (1) Which SKU lets free-licence users view Power BI content? (2) What happens when two workloads share one pool of CUs? (3) Size and isolation.

**Answer.**

- F64 or larger, because smaller F SKUs require each viewer to hold a Pro, PPU or trial licence.
- Production only, so experiments cannot use up the compute the report needs.

Licensing and isolation both point to a dedicated F64 or larger for this report.

**Common wrong answers.**

- sku = f16: Below F64, viewers need Pro, Premium Per User or trial licences. 200 free-licence users could not open the report.
- share = sandbox: A heavy notebook in the sandbox consumes the same CUs and can slow or throttle the report.

## C1-L1 exercise 4

**Hints.** (1) Build a mapping from each valid SKU name to its CUs. (2) Check membership in the mapping, then the hours. (3) Multiply and return a float.

**Answer.**

```
1 F_SKUS: dict[str, int] = {f"F{n}": n for n in (2, 4, 8, 16, 32, 64, 128, 256, 512, 1024, 2048)}
```

```

2
3
4 def cu_hours(sku: str, hours: float) -> float:
5 """Capacity-unit hours for running an F SKU capacity.
6
7 Args:
8 sku: F SKU name such as "F64".
9 hours: Hours the capacity runs.
10
11 Returns:
12 CUs multiplied by hours.
13
14 Raises:
15 ValueError: For an unknown SKU or negative hours.
16 """
17 if sku not in F_SKUS:
18 raise ValueError(f"unknown F SKU: {sku!r}")
19 if hours < 0:
20 raise ValueError("hours must not be negative")
21 return float(F_SKUS[sku] * hours)

```

An explicit list of SKUs rejects typos such as F3 or f64 instead of pricing a capacity that does not exist.

## C1-L1 exercise 5

**Hints.** (1) Was the afternoon's compute already used? (2) What happens to usage Fabric smoothed over the coming hours? (3) It is settled at pause.

**Answer.** b) Smoothed and overage usage still owed is added to the Azure bill at once, and items on the capacity stop working until it resumes.

Because: r1) Fabric smooths usage over time. Pausing settles what is still owed and makes the capacity's items unavailable.

Pausing saves money from then on, but the usage already smoothed forward is billed at once, and every item on the capacity stops until it resumes.

### Common wrong answers.

- a: The work already happened. Pausing settles the remaining smoothed and overage usage on the bill.
- c: A paused capacity is not operational. Its items stop until you resume it.
- d: Pausing affects compute. It does not delete OneLake data.
- b|r2: Right answer, wrong reason. There is no double rate. Pausing settles usage already incurred.

## C1-L2 OneLake, one lake per tenant

---

### C1-L2 exercise 1

**Hints.** (1) Can you create a second OneLake? (2) What is a workspace in OneLake terms? (3) A folder.

**Answer.** b) One for the whole tenant Because: r1) Every tenant has a single OneLake. Workspaces organise it; they do not create new lakes.

One tenant, one OneLake. Workspaces and items are folders inside it.

#### Common wrong answers.

- a: You cannot create more OneLakes. Business units get workspaces and domains inside the one lake.
- c: Capacities provide compute. Storage lives in the tenant's single OneLake.
- d: Each workspace is a folder in OneLake, not a separate lake.
- b|r2: Right answer, wrong reason. It is a platform design, not a licensing unit.

### C1-L2 exercise 2

**Hints.** (1) The workspace goes before the @. (2) The item name carries its type in lower case. (3) Tables for managed Delta tables.

#### Answer.

```
1 abfss://sales_prod@onelake.dfs.fabric.microsoft.com/lh_sales.lakehouse/
 Tables/orders
```

Workspace before @, then `item.itemtype`, then `Tables/<table>`.

#### Common wrong answers.

- g3 = `Files`: Files holds unmanaged files. Managed Delta tables live under Tables.
- g1 = `lh_sales`: The part before @ is the workspace, not the item.

### C1-L2 exercise 3

**Hints.** (1) What character does the workspace name contain? (2) The abfs URI does not allow it. (3) Use the GUIDs, which never change.

**Answer.** The bug is on line 1.

Fixed version:

```
1 WORKSPACE_ID = CONFIG["workspace_id"]
2 LAKEHOUSE_ID = CONFIG["lakehouse_id"]
```

```
3 path = f"abfss://{WORKSPACE_ID}@onelake.dfs.fabric.microsoft.com/
 {LAKEHOUSE_ID}/Tables/orders"
```

The space in the workspace name breaks the abfs URI. Renaming the workspace would fix this job and break everything else that uses the old name. GUIDs fix both.

### Common wrong answers.

- Line 3: The format of line 3 is right for names. The value it receives on line 1 is the problem.

### C1-L2 exercise 4

**Hints.** (1) Check the item type against a fixed set. (2) Reject spaces in workspace and item names. (3) `relative.lstrip('/')` avoids a double slash.

### Answer.

```
1 ONELAKE_HOST = "onelake.dfs.fabric.microsoft.com"
2 ITEM_TYPES = frozenset({"lakehouse", "warehouse"})
3
4
5 def onelake_path(workspace: str, item: str, item_type: str, relative:
 str) -> str:
6 """abfss path to data in a OneLake item, addressed by name.
7
8 Args:
9 workspace: Workspace name, without spaces.
10 item: Item name, without spaces.
11 item_type: "lakehouse" or "warehouse".
12 relative: Path inside the item, such as Tables/orders.
13
14 Returns:
15 The abfss URI.
16
17 Raises:
18 ValueError: For an unknown item type or names with spaces.
19 """
20 if item_type not in ITEM_TYPES:
21 raise ValueError(f"unsupported item type: {item_type!r}")
22 if " " in workspace or " " in item:
23 raise ValueError("abfs URIs do not allow spaces; use GUIDs
 instead")
24 return f"abfss://{workspace}@{ONELAKE_HOST}/{item}.{item_type}/
 {relative.lstrip('/')}"
```

Validating inputs turns a confusing URI error at read time into a clear message at build time.

## C1-L2 exercise 5

**Hints.** (1) What changes and what does not? (2) Which endpoint guarantees the region? (3) Does Databricks need a copy to read Delta?

**Answer.**

- GUIDs survive renames and avoid URI problems with special characters.
- The regional endpoint keeps requests in the workspace's region, which the residency rule needs.
- OneLake speaks ADLS APIs and stores open Delta tables, so the job reads them without a second copy.

Stable identifiers, a regional endpoint, and no copies.

**Common wrong answers.**

- address = `names`: Names change, and names with spaces break abfs URIs.
- endpoint = `global`: Microsoft Learn warns that the global endpoint can route data out of your region.
- copy = `copy_nightly`: A nightly copy is a second dataset to govern and keep in sync. Read the Delta tables where they are.

## C1-L3 Workspace roles, granted to groups

---

### C1-L3 exercise 1

**Hints.** (1) Which rights sit only at the top of the role table? (2) Workspace-level settings. (3) Admin.

**Answer.** a) Admin only Because: r1) Connecting Git, like deleting the workspace or creating a workspace identity, is reserved for Admins.

Only Admins connect a workspace to Git, create a workspace identity, and update or delete the workspace.

**Common wrong answers.**

- `b`: Members can add people and reshare, but connecting Git is an Admin right.
- `c`: Contributors create and change items, but cannot connect the workspace to Git.
- `d`: Viewers cannot change anything in the workspace.
- `a` | `r2`: Right answer, wrong reason. It is a permission boundary, not a licence.

## C1-L3 exercise 2

**Hints.** (1) Which two read rights does the role table list? (2) Which one does Viewer have? (3) ReadData, the T-SQL one.

**Answer.** b) It fails: Viewers read lakehouse data with T-SQL, not through Spark or OneLake APIs. Because: r1) Viewer includes ReadData (T-SQL) but not ReadAll (Spark and OneLake APIs).

Point Viewers to the SQL analytics endpoint, or grant ReadAll on the item if Spark access is really needed.

### Common wrong answers.

- a: Viewer grants T-SQL reads only. Spark reads need ReadAll, which starts at Contributor.
- c: There is no row limit tied to the role. The read is refused.
- d: Viewers can open and read notebooks. They cannot run them or read data through Spark.
- b|r2: Right answer, wrong reason. Spark is available. The Viewer role lacks ReadAll.

## C1-L3 exercise 3

**Hints.** (1) Who manages, who writes, who reads? (2) Least privilege. (3) Groups, not individuals.

### Answer.

- Platform team: Admin, as a group, so management does not depend on one person.
- Deployment identity: Contributor, the only writer in production.
- Analysts: Viewer. They query with T-SQL and use reports.

One managing group, one automated writer, and readers as Viewers.

### Common wrong answers.

- analysts = contributor: Contributors can change and delete every item in production. Analysts need to read.
- deploy = admin: Deployment needs to write items, not manage people or delete the workspace. Contributor is enough.
- platform = viewer: Someone must be able to manage the workspace. Give that right to a group.



## C1-L3 exercise 4

**Hints.** (1) A dictionary from action to role. (2) Spark and OneLake reads (ReadAll) start at Contributor. (3) Indexing a dict raises KeyError for unknown keys.

**Answer.**

```
1 MIN_ROLE: dict[str, str] = {
2 "connect_git": "Admin",
3 "delete_workspace": "Admin",
4 "add_members": "Member",
5 "run_notebook": "Contributor",
6 "spark_read": "Contributor",
7 "tsql_read": "Viewer",
8 }
9
10
11 def min_role(action: str) -> str:
12 """Lowest Fabric workspace role that allows an action.
13
14 Raises:
15 KeyError: For an action not in the table.
16 """
17 return MIN_ROLE[action]
```

Encoding the role table as data lets access requests be checked automatically, and an unknown action fails loudly instead of defaulting to a guess.

## C1-L3 exercise 5

**Hints.** (1) Which line grants write access? (2) Do analysts need to change production items? (3) Also replace the individual with a group.

**Answer.** The bug is on line 4.

Fixed version:

```
1 workspace: sales_prod
2 roles:
3 Admin: [grp-fabric-platform-admins]
4 Contributor: [sp-sales-deployment]
5 Viewer: [grp-sales-analysts, grp-sales-readers]
```

Analysts as Contributors can change or delete production items. Make them Viewers, keep automation as the only writer, and replace the individual with a group.

**Common wrong answers.**

- Line 5: An individual assignment breaks the groups rule, but it only grants read access. Line 4 gives write access to production.

- Line 3: An admin group is the right pattern.

## C2-L1 Lakehouse Tables and Files, and schemas

---

### C2-L1 exercise 1

**Hints.** (1) Which format does the endpoint read? (2) Which folder holds managed tables? (3) Delta, in Tables.

**Answer.** b) Nothing. Only Delta tables in Tables appear in the endpoint. Because: r1) The endpoint exposes Delta tables only. Parquet and CSV cannot be queried through it.

Only Delta tables reach the SQL engines. Files is for raw and non-Delta data.

#### Common wrong answers.

- a: Files is not scanned for tables. Write a Delta table to Tables with `saveAsTable`.
- c: Refreshing metadata picks up Delta tables. It never turns a Parquet file into a table.
- d: Fabric does not convert it. You write the Delta table yourself.
- b|r2: Right answer, wrong reason. Files is visible. It is just not tables.

### C2-L1 exercise 2

**Hints.** (1) What fills in the missing parts of the name? (2) The default lakehouse and the default schema. (3) dbo.

**Answer.** a) `lh_survey.dbo.respondents` Because: r1) A bare table name resolves to the dbo schema of the default lakehouse.

`respondents` resolves to `lh_survey.dbo.respondents`. If the table lives in `sales`, name the schema.

#### Common wrong answers.

- b: Spark does not search schemas. Missing parts come from defaults: default lakehouse, dbo schema.
- c: Omitted parts are allowed and filled from defaults.
- d: Nothing is unioned. One table is resolved.
- a|r2: Right answer, wrong reason. Spark reads any schema you name. dbo is only the default.

## C2-L1 exercise 3

**Hints.** (1) The four-part name is workspace.lakehouse.schema.table. (2) The third part is a schema. (3) The employees table lives in schema hrm in Microsoft's example, or dbo by default.

**Answer.**

```
1 SELECT e.name, d.dept_name
2 FROM ops_ws.hr_lakehouse.hrm.employees AS e
3 JOIN global_ws.corporate_lakehouse.company.departments AS d
4 ON e.deptno = d.deptno
```

workspace.lakehouse.schema.table reaches any table you have permission to read, across lakehouses and workspaces.

### Common wrong answers.

- g1 = Tables: Tables is a OneLake folder, not a schema. The third part is the schema name, such as dbo or hrm.

## C2-L1 exercise 4

**Hints.** (1) Read with .schema(RESPONDENT\_SCHEMA) and the header option. (2) Write with .write.mode("append").saveAsTable(table). (3) Return the DataFrame.

**Answer.**

```
1 from pyspark.sql import DataFrame, SparkSession
2 from pyspark.sql import types as T
3
4 RESPONDENT_SCHEMA = T.StructType([
5 T.StructField("respondent_id", T.StringType(), False),
6 T.StructField("wave", T.IntegerType(), False),
7 T.StructField("weight", T.DoubleType(), True),
8])
9
10
11 def load_delivery(spark: SparkSession, path: str, table: str) ->
 DataFrame:
12 """Read a CSV delivery and append it to a managed Delta table.
13
14 Args:
15 spark: Active session.
16 path: CSV path, usually under Files/landing.
17 table: Target table, ideally lakehouse.schema.table.
18
19 Returns:
20 The DataFrame that was read.
21 """
22 df = spark.read.option("header",
 True).schema(RESPONDENT_SCHEMA).csv(path)
```

```
23 df.write.mode("append").saveAsTable(table)
24 return df
```

A managed Delta table in Tables serves Spark, the SQL analytics endpoint and Direct Lake from one copy.

### C2-L1 exercise 5

**Hints.** (1) Which area holds non-Delta files? (2) Which area do SQL and Power BI read? (3) Schemas organise tables by domain.

**Answer.**

- Raw deliveries in Files, as received.
- Cleaned and reporting data as Delta tables in Tables, under schemas such as `survey_silver` and `survey_gold`.

Files for raw, Tables for Delta, and schemas to keep the tables organised.

**Common wrong answers.**

- `raw = tables_dbo`: CSV files under Tables are not Delta tables. They end up in the Unidentified area.
- `gold = files`: Power BI's Direct Lake and the SQL endpoint read Delta tables in Tables, not files.
- `clean = tables_dbo`: It works, but everything lands in one flat schema. Named schemas keep domains apart as the lakehouse grows.

## C2-L2 Warehouse T-SQL, and where it differs from SQL Server

---

### C2-L2 exercise 1

**Hints.** (1) What does NOT ENFORCED mean? (2) Does the engine check the key on insert? (3) No.

**Answer.** b) Both rows are stored. The key is not enforced. Because: r1) NOT ENFORCED keys are optimiser hints. The engine does not check them.

Both rows land. Check uniqueness after each load with a `GROUP BY ... HAVING COUNT(*) > 1` query.

**Common wrong answers.**

- a: That is SQL Server behaviour for enforced keys. Fabric Warehouse keys are NOT ENFORCED.
- c: INSERT never replaces rows. You would need MERGE for upsert behaviour.

- d: Nothing locks. Both rows are simply stored.
- b|r2: Right answer, wrong reason. No background job removes duplicates. You must check.

## C2-L2 exercise 2

**Hints.** (1) Not clustered. (2) Not checked by the engine. (3) `NONCLUSTERED` and `NOT ENFORCED`.

**Answer.**

```
1 ALTER TABLE dbo.dim_customer
2 ADD CONSTRAINT pk_dim_customer PRIMARY KEY NONCLUSTERED (customer_id)
 NOT ENFORCED;
```

`PRIMARY KEY NONCLUSTERED (...) NOT ENFORCED` is the only accepted form.

**Common wrong answers.**

- g1 = `CLUSTERED`: Fabric Warehouse supports primary keys only as `NONCLUSTERED`.

## C2-L2 exercise 3

**Hints.** (1) Two SQL Server types have no Fabric Warehouse equivalent. (2) One for money, one for Unicode text. (3) decimal and varchar replace them.

**Answer.** The bug is on line 4, 5.

Fixed version:

```
1 CREATE TABLE dbo.fact_payment (
2 payment_id BIGINT NOT NULL,
3 customer_id INT NOT NULL,
4 amount DECIMAL(19, 4) NOT NULL,
5 note VARCHAR(500),
6 paid_at DATETIME2(6) NOT NULL
7);
```

`money` and `nvarchar` are not supported. Use `DECIMAL(19, 4)` and `VARCHAR(500)`. Varchar uses a UTF-8 collation, so it stores Unicode text.

**Common wrong answers.**

- Line 6: `datetime2` with up to 6 fractional digits is supported.
- Line 2: `BIGINT` is supported.

## C2-L2 exercise 4

**Hints.** (1) Count rows per `customer_id`. (2) Filter groups, not rows. (3) `HAVING COUNT(*) > 1`.

**Answer.**

```
1 SELECT customer_id, COUNT(*) AS n
2 FROM dbo.dim_customer
3 GROUP BY customer_id
4 HAVING COUNT(*) > 1;
```

With keys unenforced, this query is the guard. Any row it returns means a duplicate got in.

## C2-L2 exercise 5

**Hints.** (1) What no longer enforces uniqueness? (2) Which feature is not supported at all? (3) Use the documented alternatives.

### Answer.

- Keys become hints, so add a duplicate check that fails the load.
- Triggers are not supported. Do the audit write explicitly in the loading procedure.
- Map each type to its documented alternative, keeping precision and scale.

A port succeeds when every lost guarantee is replaced by an explicit step.

### Common wrong answers.

- `pk = trust`: Sources resend files and pipelines rerun. Without a check, duplicates load silently.
- `trigger = keep`: Fabric Warehouse does not support triggers. The CREATE TRIGGER statement fails.
- `types = string`: Strings lose numeric semantics and cost storage. Use decimal for amounts.

## C2-L3 Lakehouse, Warehouse or SQL analytics endpoint

### C2-L3 exercise 1

**Hints.** (1) The endpoint is described as read-only. Read-only for what? (2) It blocks changes to table data, not every kind of object. (3) Views, functions and stored procedures that read are allowed.

**Answer.** b) You can create a view over lakehouse tables in it. Because: r1) Views are T-SQL objects that read data. The endpoint allows objects that read, not statements that change table data.

The endpoint is read-only for table data, but you can create views, functions and stored procedures in it. Fabric creates it with the lakehouse, and it reads the same Delta tables in OneLake.

### Common wrong answers.

- a: The endpoint is read-only for table data. Write with Spark in the lakehouse, or move the table to a warehouse.
- c: Fabric creates the endpoint automatically with every lakehouse.
- d: It reads the same Delta tables in OneLake. There is no second copy. It syncs metadata, not data.
- b|r2: Right choice, wrong reason. A view is a T-SQL object in the endpoint. The Delta tables stay unchanged.

### C2-L3 exercise 2

**Hints.** (1) Ask who writes the data, and in which language. (2) Spark writers and files: one item. T-SQL writers: another. (3) T-SQL readers of lakehouse data need no new item at all.

#### Answer.

- Raw files and PySpark: **Lakehouse**. Spark writes, and files sit next to tables.
- T-SQL writes across several tables: **Warehouse**, for DML and multi-table transactions.
- T-SQL reads of lakehouse data, plus views: **SQL analytics endpoint**. No copy needed.

Spark writers use a lakehouse. T-SQL writers use a warehouse. T-SQL readers of lakehouse data use the endpoint that comes with the lakehouse.

### Common wrong answers.

- finance = lakehouse: The finance team writes with T-SQL. A lakehouse accepts writes from Spark, and its endpoint is read-only. Use a warehouse.
- finance = endpoint: The endpoint cannot change table data, so the stored procedures would fail on their first UPDATE or INSERT.
- raw = warehouse: A warehouse does not hold raw files, and the team writes with PySpark. Use a lakehouse.
- adhoc = warehouse: Copying the lakehouse tables into a warehouse makes a second copy to keep in sync. The endpoint reads them where they are.

### C2-L3 exercise 3

**Hints.** (1) Which statements change table data? (2) The endpoint accepts reads and objects that read. (3) UPDATE on a lakehouse table fails through the endpoint.

**Answer.** The bug is on line 7.

Fixed version:

```
1 -- Connected to the SQL analytics endpoint of lh_research
2 CREATE VIEW dbo.v_wave_reach AS
3 SELECT wave_id, SUM(weighted_reach) / SUM(weighted_base) AS reach_pct
4 FROM dbo.nat_metrics
5 GROUP BY wave_id;
6 GO
7 -- Fix wave W00 in the lakehouse with Spark, not here.
```

Line 7 is an UPDATE on a lakehouse table, which the read-only endpoint rejects. Make that correction in a Spark notebook against the lakehouse, and keep only the view here.

### Common wrong answers.

- Line 2: Creating a view is allowed on the endpoint. Look for a statement that changes table data.
- Line 3: The division could hit zero, but that is not what stops this script. Look for a write.

### C2-L3 exercise 4

**Hints.** (1) Aggregate both sums per wave, then divide. (2) `NULLIF(x, 0)` returns NULL when x is 0, and dividing by NULL gives NULL. (3) `SUM(weighted_reach) / NULLIF(SUM(weighted_base), 0) AS reach_pct`, then `GROUP BY wave_id`.

### Answer.

```
1 CREATE VIEW dbo.v_wave_reach AS
2 SELECT
3 wave_id,
4 SUM(weighted_reach) / NULLIF(SUM(weighted_base), 0) AS reach_pct
5 FROM dbo.nat_metrics
6 GROUP BY wave_id;
```

Sum both columns per wave and divide by `NULLIF(SUM(weighted_base), 0)`. A zero base becomes NULL, so the division returns NULL instead of an error.

### C2-L3 exercise 5

**Hints.** (1) Three-part names are item.schema.table. (2) The first part names the item that holds the table. (3) The table sits in the lakehouse `lh_research`.

### Answer.

```
1 INSERT INTO dbo.wave_summary (wave_id, weighted_base)
2 SELECT wave_id, SUM(weighted_base)
3 FROM lh_research.dbo.nat_metrics
4 GROUP BY wave_id;
```



A warehouse query reads a lakehouse table in the same workspace with `lh_research.dbo.nat_metrics`. The data stays in one place.

### Common wrong answers.

- `g1 = wh_finance`: `wh_finance` is the warehouse you are in. The table lives in the lakehouse, so name the lakehouse.
- `g1 = OneLake`: There is no database called OneLake. The first part of a three-part name is the item name: the lakehouse.

## C2-L3 exercise 6

**Hints.** (1) The query failed first and worked later, with no change in between. (2) How does the endpoint learn that a new table exists? (3) Look at the metadata sync between the lakehouse and its endpoint.

**Answer.** b) The endpoint had not yet synced the new table's metadata. Because: r1) The endpoint discovers new Delta tables through an automatic sync that usually takes under a minute but has no fixed limit.

The endpoint syncs lakehouse metadata automatically, usually in under a minute, but the lag varies. When a later step depends on a new table, add a metadata refresh step (portal, REST API) or a retry before the query. Check the current refresh options on Microsoft Learn.

### Common wrong answers.

- `a`: The query worked an hour later without any change to the table, so the format was fine.
- `c`: The sync runs automatically. A manual refresh speeds it up, but it is not required every time.
- `d`: Permissions do not change on a timer. The delay points to the metadata sync.
- `b`|`r2`: Right choice, wrong reason. The endpoint does not copy data. It syncs metadata about the Delta tables.

## C3-L1 OneLake shortcuts, and the delete that goes through

---

### C3-L1 exercise 1

**Hints.** (1) Is the shortcut the data or a pointer? (2) What does removing a pointer do? (3) Contrast it with deleting inside the pointer.

**Answer.** b) It stays. Deleting a shortcut deletes only the shortcut object. Because: r1) Shortcuts do not cascade deletes. The shortcut is an independent object.

The shortcut goes, the data stays. Deleting a file or folder inside the shortcut is different.

**Common wrong answers.**

- a: Removing the shortcut removes only the link. Deleting files inside it is what reaches the target.
- c: There is no copy to recycle. The data was never in OneLake.
- d: Nothing is copied when you delete a shortcut.
- b|r2: Right answer, wrong reason. Deleting files inside a shortcut can delete ADLS data.

**C3-L1 exercise 2**

**Hints.** (1) Where do the files inside a shortcut live? (2) In the target. (3) Your permissions there decide.

**Answer.** b) The 2023 folder is deleted in the target storage account. Because: r1) Deleting a file or folder within a shortcut deletes it in the target when you have permission there.

Inside a shortcut you are working on the target. Deletes go through.

**Common wrong answers.**

- a: There is no hidden view. The folder is the target folder.
- c: Shortcuts are not read-only when your identity can write to the target.
- d: The shortcut still works. It just points at data that no longer includes 2023.
- b|r2: Right answer, wrong reason. The delete happens immediately, not on a schedule.

**C3-L1 exercise 3**

**Hints.** (1) Which area holds managed tables? (2) Which format does the endpoint read? (3) Tables and Delta.

**Answer.**

1 Create the shortcut under Tables, pointing at data in Delta format.

A shortcut under Tables that points at Delta data is recognised as a table.

## Common wrong answers.

- `g1 = Files`: Shortcuts under Files can hold any format, but they do not become tables.

## C3-L1 exercise 4

**Hints.** (1) Check shortcuts first and refuse early. (2) Compare against `root + "/"` so partner does not match `partner_old`. (3) Allowlist owned prefixes with the trailing slash.

## Answer.

```
1 OWNED_PREFIXES = ("Files/landing/", "Files/tmp/")
2
3
4 def safe_to_delete(path: str, shortcut_roots: set[str]) -> bool:
5 """True only for owned paths that are not under a shortcut.
6
7 Args:
8 path: Lakehouse-relative path, such as Files/landing/w14.csv.
9 shortcut_roots: Paths of every shortcut in the lakehouse.
10
11 Returns:
12 Whether a clean-up job may delete the path.
13 """
14 for root in shortcut_roots:
15 base = root.rstrip("/")
16 if path == base or path.startswith(base + "/"):
17 return False
18 return path.startswith(OWNED_PREFIXES)
```

Refusing anything under a shortcut, and allowing only named folders, makes a clean-up job unable to reach someone else's data.

## C3-L1 exercise 5

**Hints.** (1) Copy or point? (2) Where do shortcuts become tables? (3) What stops a delete from going through?

## Answer.

- A shortcut avoids a second copy and a pipeline to maintain.
- Under Tables, Delta data becomes a table for SQL and Direct Lake.
- Caching cuts egress costs for repeated reads from S3.
- A read-only identity means no delete inside the shortcut can reach the bucket.

Point, place it where it becomes a table, cache cross-cloud reads, and connect read-only.

### Common wrong answers.

- access = `copy`: A copy doubles storage and needs a pipeline and monitoring. A shortcut reads in place.
- writes = `owner`: With write access, deleting inside the shortcut deletes marketing's data.
- location = `files`: Under Files it will not be recognised as a table.

## C3-L2 Mirroring operational databases into OneLake

---

### C3-L2 exercise 1

**Hints.** (1) What do shortcuts point at? (2) Is an Azure SQL database a folder of Delta files? (3) It must be replicated.

**Answer.** b) Mirroring: it replicates the database into OneLake continuously. Because: r1) Azure SQL stores data in a proprietary format, so Fabric must replicate it, which is what mirroring does.

Databases in proprietary formats come in through mirroring. Open-format data in storage comes in through shortcuts.

### Common wrong answers.

- a: Shortcuts point at storage (OneLake, ADLS, S3 and similar), not at database engines.
- c: A nightly export is a hand-built copy with lag and maintenance. Mirroring does it continuously.
- d: There is no storage folder of Delta files to point at. The database must be replicated.
- b|r2: Right answer, wrong reason. Shortcuts copy nothing. Mirroring is needed because the source is a database.

### C3-L2 exercise 2

**Hints.** (1) How many CUs does the capacity have? (2) 1 TB free per CU. (3) 45 - 32.

**Answer.**

1 32 13

32 TB are free on an F32. The remaining 13 TB of mirrored storage is billed as OneLake storage.

### Common wrong answers.

- 64 0: The example on Learn uses F64. This capacity has 32 CUs, so 32 TB are free.

### C3-L2 exercise 3

**Hints.** (1) The standard catalog views schema. (2) Tables, not views. (3) `INFORMATION_SCHEMA` and `BASE TABLE`.

#### Answer.

```
1 SELECT TABLE_SCHEMA, TABLE_NAME
2 FROM INFORMATION_SCHEMA.TABLES
3 WHERE TABLE_TYPE = 'BASE TABLE';
```

`INFORMATION_SCHEMA.TABLES` lists the tables the endpoint exposes. Compare it with the source list.

### Common wrong answers.

- `g2 = VIEW`: Views are not mirrored tables. Filter on `BASE TABLE`.

### C3-L2 exercise 4

**Hints.** (1) Loop over the source tables. (2) `mirror.get(table)` is `None` when absent, so one comparison covers both cases. (3) Sort the result.

#### Answer.

```
1 def missing_from_mirror(source: dict[str, int], mirror: dict[str, int]) -
 > list[str]:
2 """Source tables that are absent from the mirror or differ in row
 count.
3
4 Args:
5 source: schema.table to row count in the source database.
6 mirror: schema.table to row count on the mirror's endpoint.
7
8 Returns:
9 Sorted table names that need attention.
10 """
11 return sorted(t for t, n in source.items() if mirror.get(t) != n)
```

One comparison, `mirror.get(t) != n`, catches both missing tables and partial replication.

### C3-L2 exercise 5

**Hints.** (1) Is there a table limit? (2) Which key types block mirroring? (3) How would you notice a missing table?

**Answer.**

- Mirror what reporting needs, within the 1,000-table limit.
- Tables with datetime2(7) keys cannot be mirrored. Plan for them explicitly.
- Reconcile daily, because ineligible tables fail quietly.

Plan within the documented limits, then verify with data, not with status pages.

**Common wrong answers.**

- scope = `all`: Azure SQL Database mirroring supports at most 1,000 tables.
- keys = `ignore`: Those tables are not eligible. They will simply be missing.
- check = `trust`: A table that was never eligible does not show as failing. Compare lists and counts.

## C3-L3 Mirroring Azure Databricks Unity Catalog into Fabric

---

**C3-L3 exercise 1**

**Hints.** (1) What format are Unity Catalog managed tables stored in? (2) Can Fabric read that format directly? (3) Metadata plus shortcuts.

**Answer.** b) Only the catalog structure. Data is read through shortcuts. Because: r1) Unity Catalog tables are already open Delta files, so Fabric mirrors metadata and reads the files in place.

Only the catalog structure is mirrored. The data stays in the Databricks storage account.

**Common wrong answers.**

- `a`: That is database mirroring for proprietary sources. Unity Catalog mirroring has no data movement.
- `c`: There are no snapshots. Fabric reads the current Delta files through shortcuts.
- `d`: Fabric engines read the Delta files directly. No Databricks compute runs the queries.
- `b` | `r2`: Right answer, wrong reason. It is about format: Delta can be read in place.

### C3-L3 exercise 2

**Hints.** (1) Check the list of unsupported table types. (2) RLS/CLM policies. (3) The table is left out.

**Answer.** c) It is not mirrored: tables with RLS or CLM policies are unsupported. Because: r1) The limitations page lists tables with RLS/CLM policies as unsupported.

Filtered and masked tables are excluded. Decide what to publish for Fabric users instead.

#### Common wrong answers.

- a: Unity Catalog policies are not carried into Fabric, and such tables are not supported.
- b: Fabric does not quietly drop the filter. It does not mirror the table.
- d: Views are themselves unsupported for this mirroring.
- c|r2: Right answer, wrong reason. The table type is unsupported, not read-only.

### C3-L3 exercise 3

**Hints.** (1) Scope before you connect. (2) The connection's rights cap what is readable. (3) Access, then verify.

#### Answer.

1. Choose tables safe for the whole Fabric audience, in a dedicated schema
2. Create a connection identity with SELECT on that schema only
3. Create the mirrored Azure Databricks item for that schema
4. Grant Fabric access with workspace roles, SQL permissions or model RLS
5. Compare table lists and row counts with Databricks

Scope, least-privilege connection, mirror, Fabric access, then a data check.

#### Common wrong answers.

- Mirroring first exposes whatever the connection can read. Decide the scope first.
- Set up access before users arrive. Verification then confirms what they will see.

### C3-L3 exercise 4

**Hints.** (1) Allow a fixed set of table types rather than excluding a list. (2) Also exclude tables with policies. (3) Sort the names.

#### Answer.

```

1 MIRRORABLE_TYPES = frozenset({"MANAGED", "EXTERNAL"})
2
3
4 def mirror_eligible(tables: list[dict[str, object]]) -> list[str]:
5 """Unity Catalog tables that Fabric can mirror.
6
7 Args:
8 tables: Dicts with name, table_type and has_row_filter_or_mask.
9
10 Returns:
11 Sorted names of eligible tables.
12 """
13 return sorted(
14 str(t["name"])
15 for t in tables
16 if t["table_type"] in MIRRORABLE_TYPES and not
17 t["has_row_filter_or_mask"]
18)

```

An allowlist of table types plus the policy check predicts which tables will appear, so a missing table is expected rather than a surprise.

### C3-L3 exercise 5

**Hints.** (1) Where do Unity Catalog rules stop applying? (2) Whose credential reads the data? (3) Where does the RLS live?

**Answer.**

- The filtered table cannot be mirrored, so rebuild the regional rule where users consume the data.
- Scope the connection, because it reads every mirrored table for everyone.
- Keep managers on the report. Direct endpoint access would bypass the model's RLS.

Rebuild the rule where users are, scope the credential, and keep users on the governed path.

**Common wrong answers.**

- source = `unfiltered_all`: Every manager would see every region. The rule must be rebuilt in Fabric.
- connection = `broad`: The connection's credential serves all queries. Grant it only what the mirror needs.
- users = `viewer`: Viewers can query the endpoint with T-SQL and read rows the model's RLS would hide.



## C4-L1 Choosing how data gets into Fabric

---

### C4-L1 exercise 1

**Hints.** (1) Which layer is this? (2) Which tool tracks incremental state itself? (3) Copy job.

**Answer.** b) Copy job Because: r1) Copy job supports incremental copy and scheduling natively for bronze ingestion.

Bronze ingestion starts with Copy job: incremental copy without building a pipeline.

#### Common wrong answers.

- a: Mirroring replicates databases and catalogs. This is file ingestion into bronze.
- c: Dataflow Gen2 is for low-code transformation. Nothing needs transforming yet.
- d: It works, but you build and maintain the incremental state that Copy job tracks for you.
- b|r2: Right answer, wrong reason. Pipelines and shortcuts read ADLS too.

### C4-L1 exercise 2

**Hints.** (1) Replicate, stream, transform or orchestrate? (2) Which tool is event-driven? (3) Which tool coordinates other tools?

#### Answer.

- Curated database tables: mirroring.
- Low-latency events: Eventstream.
- Heavy transformation: Spark.
- Ordering, branching and alerts across steps: a pipeline.

Each tool has a lane. Stay in it and there is less to build and maintain.

#### Common wrong answers.

- erp = copyjob: A Copy job works, but mirroring replicates the whole database with free replication compute and no schedule to manage.
- iot = copyjob: Copy jobs run on schedules. Seconds-level events need Eventstream.
- orchestrate = spark: A notebook can call others, but retries, branches and alerts across steps are what pipelines are for.

### C4-L1 exercise 3

**Hints.** (1) Movement first. (2) Transformation next. (3) Orchestration wraps the rest.

**Answer.**

1. Replicate a whole database or catalog?
2. Copy tables or files into bronze?
3. Events that must arrive within seconds?
4. Transformation needed, and by whom?
5. Several steps that need ordering and alerts?

Choose movement tools, then transformation, then wrap them in orchestration only where needed.

**Common wrong answers.**

- Starting with a pipeline leads to building copy logic inside it. Choose the movement tools first, then orchestrate them.
- Decide how data arrives before how it is transformed.

### C4-L1 exercise 4

**Hints.** (1) A dictionary from need to tool. (2) Check membership before returning. (3) Put the need in the error message.

**Answer.**

```
1 TOOL_FOR_NEED: dict[str, str] = {
2 "replicate_database": "Mirroring",
3 "bronze_copy": "Copy job",
4 "low_latency_events": "Eventstream",
5 "code_transform": "Spark notebook",
6 "lowcode_transform": "Dataflow Gen2",
7 "orchestrate": "Pipeline",
8 }
9
10
11 def recommend(need: str) -> str:
12 """Fabric tool for a data movement or processing need.
13
14 Raises:
15 ValueError: If the need is not in the table.
16 """
17 if need not in TOOL_FOR_NEED:
18 raise ValueError(f"unknown need: {need}")
19 return TOOL_FOR_NEED[need]
```

An explicit table makes the team's convention reviewable, and a clear error beats a silent default.

## C4-L1 exercise 5

**Hints.** (1) Where does the bug live? (2) In hand-written state tracking. (3) Which tool does that for you?

**Answer.** b) Move the tables to Copy jobs with incremental copy, which track state themselves. Because: r1) Copy job manages the state of the last successful run, so there is no hand-written watermark logic to get wrong.

Removing hand-built incremental state removes the bugs it causes. Copy job tracks the last successful run itself.

### Common wrong answers.

- a: Retries repeat the same logic, including its watermark bug.
- c: More runs repeat the off-by-one error more often.
- d: That moves the hand-built state into another tool without removing it.
- b|r2: Right answer, wrong reason. Pipelines can copy incrementally, but you build the state tracking.

## C4-L2 Incremental copy with a watermark, or with CDC

---

### C4-L2 exercise 1

**Hints.** (1) What does a watermark look for? (2) Does a deleted row have a new value? (3) CDC captures deletes.

**Answer.** b) Nothing. Watermark-based incremental copy cannot detect deletes. Because: r1) A watermark only finds rows with a larger value. A deleted row has no value to find.

Watermarks only see rows that exist with a larger value. Use CDC when deletes matter.

### Common wrong answers.

- a: That needs CDC, which captures delete events.
- c: Soft deletes come from CDC with SCD Type 2. A watermark sees nothing.
- d: The job succeeds. The deleted row just stays in the destination.
- b|r2: Right answer, wrong reason. CDC-based incremental copy does replicate deletes.

### C4-L2 exercise 2

**Hints.** (1) Which rows have rv above 105? (2) Where does row 3 fall? (3) Below the watermark.

## Answer.

```
1 [4]
```

Only row 4 is copied. Row 3 is lost for good unless you reset to a full copy, which is why commit-order columns such as ROWVERSION, or CDC, are safer.

## Common wrong answers.

- [3, 4]: Row 3 has rv 104, below the watermark 105, so it is never copied. That is the late-row blind spot.

## C4-L2 exercise 3

**Hints.** (1) Which method captures deletes? (2) Which update method keeps one row per key? (3) CDC and Merge.

## Answer.

```
1 Incremental method: CDC
2 Update method: Merge
```

CDC replicates inserts, updates and deletes. Merge keeps one row per key.

## Common wrong answers.

- g2 = Append: Append adds each changed row again, so a key appears many times.
- g1 = Watermark: A watermark cannot see deletes.

## C4-L2 exercise 4

**Hints.** (1) Handle last\_wm None first. (2) Strictly greater than, to avoid re-copying the boundary. (3) Return ([], last\_wm) when nothing is new.

## Answer.

```
1 Row = dict[str, int]
2
3
4 def next_batch(rows: list[Row], last_wm: int | None) -> tuple[list[Row],
 int | None]:
5 """Rows above the watermark, and the new watermark.
6
7 Args:
8 rows: Source rows, each with an integer rv.
9 last_wm: Largest rv copied so far, or None before the first run.
10
11 Returns:
12 The rows to copy and the watermark to store.
13 """
14 batch = [r for r in rows if last_wm is None or r["rv"] > last_wm]
15 if not batch:
16 return [], last_wm
```

```
17 return batch, max(r["rv"] for r in batch)
```

Strictly greater avoids duplicates, and the empty-run guard keeps the watermark stable. The late-row gap remains, which is why CDC or ROWVERSION matters.

### C4-L2 exercise 5

**Hints.** (1) Deletes matter here. (2) One row per key. (3) How do you recover missed changes?

**Answer.**

- CDC captures inserts, updates and deletes, with no watermark blind spots.
- Merge keeps one current row per order.
- Reset to a full copy for the table when you know changes were missed.

CDC plus Merge keeps silver in step with the source, and a per-table reset repairs known gaps.

**Common wrong answers.**

- `method = wm_modified`: A watermark misses deletes and late rows.
- `update = append`: Append adds each changed order again, breaking one row per key.
- `recovery = wait`: Changes below the stored state are not picked up by later runs.

## C4-L3 Pipelines that fail fast, retry sensibly and alert

---

### C4-L3 exercise 1

**Hints.** (1) It is much longer than most jobs. (2) Half a day. (3) 12 hours.

**Answer.** c) 12 hours Because: r1) The documented default timeout is 12 hours. Seven days is the maximum you can set.

12 hours by default. Set each activity's timeout close to its real run time, with some headroom.

**Common wrong answers.**

- a: 30 seconds is the default retry interval. The default timeout is 12 hours.
- d: Seven days is the maximum, not the default.
- b: The default is 12 hours.
- c|r2: Right answer, wrong reason. Timeouts are unrelated to capacity pausing.

## C4-L3 exercise 2

**Hints.** (1) Four branches: success, failure, completion, skip. (2) Completion is not a failure. (3) Either outcome.

**Answer.** c) Whether the notebook succeeds or fails Because: r1) Upon Completion runs after success or failure. Upon Failure is the failure-only branch.

A failure alert on Upon Completion fires every night. Use Upon Failure.

### Common wrong answers.

- a: That is Upon Failure.
- b: That is Upon Success.
- d: That is Upon Skip.
- c|r2: Right answer, wrong reason. It refers to this activity finishing, either way.

## C4-L3 exercise 3

**Hints.** (1) Inside a string, expressions are wrapped in `@{ }`. (2) A whole value that must stay a number starts with `@` alone. (3) `@{ }` and `@`.

### Answer.

```
1 Folder path: landing/@{pipeline().parameters.wave}/
2 Max rows: @pipeline().parameters.max_rows
```

`@{...}` inside text interpolates to a string. A bare `@...` expression keeps the value's type.

### Common wrong answers.

- `g2 = @{:` String interpolation always returns a string. For a number, use the bare `@` expression.

## C4-L3 exercise 4

**Hints.** (1) A regular expression with `fullmatch` validates the shape. (2) Check each field's range. (3) Cap the total at  $7 \times 24 \times 3600$  seconds.

### Answer.

```
1 import re
2
3 TIMEOUT_PATTERN = re.compile(r"(\d+)\.(\d{2}):(\d{2}):(\d{2})")
4 MAX_SECONDS = 7 * 24 * 3600
5
6
7 def parse_timeout(value: str) -> int:
8 """Seconds in a pipeline activity timeout written as D.HH:MM:SS.
9
```

```

10 Raises:
11 ValueError: On a bad format, out-of-range fields, or more than
12 seven days.
13 """
14 match = TIMEOUT_PATTERN.fullmatch(value)
15 if match is None:
16 raise ValueError(f"timeout must be D.HH:MM:SS, got {value!r}")
17 days, hours, minutes, seconds = (int(g) for g in match.groups())
18 if hours > 23 or minutes > 59 or seconds > 59:
19 raise ValueError(f"field out of range in {value!r}")
20 total = days * 86400 + hours * 3600 + minutes * 60 + seconds
21 if total > MAX_SECONDS:
22 raise ValueError("timeout exceeds the seven-day maximum")
23 return total

```

A linter that parses timeouts can flag every activity still on the 12-hour default, or set far above its normal run time.

### C4-L3 exercise 5

**Hints.** (1) Which branch lets gold run after a failure? (2) Completion means either outcome. (3) Use Upon Success.

**Answer.** The bug is on line 3.

Fixed version:

```

1 Activity: nb_silver_orders
2 Timeout: 0.00:30:00
3 Activity: nb_gold_sales depends on nb_silver_orders: Upon Success
4 Activity: email_on_failure depends on nb_silver_orders: Upon Failure

```

Upon Completion runs gold after a failed silver, publishing stale or partial data. Chain data steps on Upon Success.

### Common wrong answers.

- Line 4: The alert is correctly on Upon Failure.
- Line 2: A 30-minute timeout is sensible for a 10-minute notebook.

## C5-L1 From Databricks notebooks to Fabric notebooks

### C5-L1 exercise 1

**Hints.** (1) Where do Fabric secrets live? (2) Which utility module handles credentials? (3) notebookutils.credentials.

**Answer.** b) notebookutils.credentials.getSecret("https://kv-sales.vault.azure.net/", "sql-password") Because: r1) Fabric reads secrets from Azure Key Vault through notebookutils.credentials.

`getSecret(vault_url, name)` reads from Azure Key Vault, with your identity's Get permission.

### Common wrong answers.

- a: dbutils is Databricks-only. Fabric has no secret scopes.
- c: Workspace settings do not hold secrets as environment variables. Use Key Vault.
- d: Spark configuration is not a secret store.
- b|r2: Right answer, wrong reason. It reads from Key Vault at runtime.

### C5-L1 exercise 2

**Hints.** (1) Look at `run()`'s signature. (2) `timeout_seconds` has a default. (3) 90.

**Answer.** b) It fails after 90 seconds, the default timeout. Because: r1) `run()` has `timeout_seconds` with a default of 90.

Pass `timeout_seconds` sized to the child's real run time with headroom.

### Common wrong answers.

- a: Only if you pass a longer `timeout_seconds`.
- c: 12 hours is the pipeline activity default. `notebook.run` defaults to 90 seconds.
- d: `run()` waits for the child and returns its exit value.
- b|r2: Right answer, wrong reason. The limit is time, not job count.

### C5-L1 exercise 3

**Hints.** (1) The Fabric utility module. (2) The name of `run()`'s timeout parameter. (3) `notebookutils` and `timeout_seconds`.

### Answer.

```
1 files = notebookutils.fs.ls("Files/landing")
2 value = notebookutils.notebook.run("nb_child", timeout_seconds=600)
```

`notebookutils` replaces `dbutils`, and `timeout_seconds` sets the child's time limit.

### Common wrong answers.

- `gl = dbutils`: `dbutils` does not exist in Fabric.
- `gl = mssparkutils`: `mssparkutils` still works but is being retired. Use `notebookutils`.



## C5-L1 exercise 4

**Hints.** (1) Check for `dbutils.secrets.first` and `raise`. (2) Replace the specific prefixes, including the trailing dot. (3) Leave anything else alone.

**Answer.**

```
1 REWRITES: tuple[tuple[str, str], ...] = (
2 ("dbutils.fs.", "notebookutils.fs."),
3 ("dbutils.notebook.", "notebookutils.notebook."),
4 ("mssparkutils.", "notebookutils."),
5)
6
7
8 def port_line(line: str) -> str:
9 """Rewrite Databricks and legacy Fabric utility calls to
10 notebookutils.
11 Raises:
12 ValueError: For secret lookups, which must move to Azure Key Vault
13 by hand.
14 """
15 if "dbutils.secrets." in line:
16 raise ValueError("secrets need a manual change: use
17 notebookutils.credentials.getSecret with Azure Key Vault")
18 for old, new in REWRITES:
19 line = line.replace(old, new)
20 return line
```

Mechanical rewrites are safe for `fs` and `notebook` calls. Secrets need a design change, so the tool stops and says so.

## C5-L1 exercise 5

**Hints.** (1) What differs between an interactive session and a pipeline run? (2) Inline commands. (3) Move the library to an environment.

**Answer.** The bug is on line 1.

Fixed version:

```
1 # pandera==0.33.1 is installed in the environment attached to this
2 notebook.
3 import pandera.pandas as pa
4 df = spark.read.table("lh_sales.silver.orders").toPandas()
5 SCHEMA.validate(df, lazy=True)
```

Inline installs are off by default in pipeline runs. Put the library in an environment (Full mode for pipelines), which also pins the version reproducibly.

**Common wrong answers.**

- Line 2: The import fails only because the install on line 1 did not run.

- Line 3: Reading the table is fine.

## C5-L2 Environments, resource profiles and V-Order

---

### C5-L2 exercise 1

**Hints.** (1) Which resource profile do new workspaces use? (2) What does it do to V-Order? (3) Off, for faster writes.

**Answer.** b) No. New workspaces default to the writeHeavy profile, with V-Order off. Because: r1) spark.sql.parquet.vorder.default is false in new workspaces to favour write performance.

New workspaces favour writes. Turn V-Order on for tables that are read far more than written.

#### Common wrong answers.

- a: That was true for older workspaces. New ones default to writeHeavy, with V-Order off.
- c: Files holds non-table data. V-Order applies to Delta table writes.
- d: There is no size threshold. It depends on settings.
- b | r2: Right answer, wrong reason. V-Order still exists. You turn it on where reads matter.

### C5-L2 exercise 2

**Hints.** (1) Which option is reproducible? (2) Which mode does Microsoft recommend for pipelines? (3) Full mode.

**Answer.** b) An environment in Full mode, attached to the notebooks or set as workspace default Because: r1) Full mode builds a reproducible snapshot, and Microsoft recommends it for pipelines and scheduled runs.

A Full mode environment pins the dependency snapshot for every scheduled run.

#### Common wrong answers.

- a: Inline installs are turned off by default in pipeline runs, and resolve dependencies differently from run to run.
- c: Quick mode installs at every session start, adding minutes, and suits development more than scheduled runs.
- d: That is an inline install by another route: session-scoped and hard to reproduce.

- b|r2: Right answer, wrong reason. Both modes install the same packages. Full mode is reproducible.

### C5-L2 exercise 3

**Hints.** (1) The table-level property starts with delta. (2) delta.parquet.vorder.enabled. (3) Set it to true.

**Answer.**

```
1 ALTER TABLE lh.gold.sales_by_region
2 SET TBLPROPERTIES ("delta.parquet.vorder.enabled" = "true");
```

The table property applies V-Order to future writes to this table only.

**Common wrong answers.**

- g1 = spark.sql.parquet.vorder.enable: That old session setting is removed in Runtime 1.3 and later, and it is not a table property.

### C5-L2 exercise 4

**Hints.** (1) The option belongs on df.write. (2) Its name is parquet.vorder.enabled. (3) Chain mode, option, saveAsTable.

**Answer.**

```
1 from pyspark.sql import DataFrame
2
3
4 def write_gold(df: DataFrame, table: str) -> None:
5 """Overwrite a read-heavy gold table with V-Order for this write only.
6
7 Args:
8 df: Gold rows.
9 table: Target table, such as lh.gold.sales_by_region.
10 """
11 df.write.mode("overwrite").option("parquet.vorder.enabled",
12 "true").saveAsTable(table)
```

A writer option scopes V-Order to one write, so ingestion elsewhere in the session keeps the fast default.

### C5-L2 exercise 5

**Hints.** (1) Reproducible libraries. (2) Most writes are ingestion. (3) Where are reads heavy?

**Answer.**

- A Full mode default environment gives every scheduled notebook the same pinned libraries.
- writeHeavy suits ingestion, which is most of the Spark work.

- V-Order on gold tables speeds up the reads Power BI makes, without slowing ingestion.

Pin libraries once, optimise for writes by default, and pay for read layout only where reads happen.

### Common wrong answers.

- `libs = inline`: Inline installs are off in pipeline runs and are not reproducible.
- `session = read_pbi`: That applies V-Order to every write, including bronze and silver, which pay the write cost for no reader.
- `gold = none_vorder`: Direct Lake reads gold constantly. V-Order is worth its write cost there.

## C5-L3 Delta tables shared between Databricks and Fabric

---

### C5-L3 exercise 1

**Hints.** (1) Which experiences list `TIMESTAMP_NTZ` as supported? (2) Spark does. (3) SQL and Direct Lake do not.

**Answer.** c) Fabric Spark Runtime 2.0 Because: r1) The interoperability matrix lists `TIMESTAMP_NTZ` as supported in Spark but not in the SQL endpoint or Direct Lake.

Use `TIMESTAMP` for columns that SQL or Power BI users need.

### Common wrong answers.

- a: The SQL analytics endpoint does not support `TIMESTAMP_NTZ`.
- b: Direct Lake does not support `TIMESTAMP_NTZ`.
- d: Support differs by experience. Check the matrix.
- c | r2: Right answer, wrong reason. Every listed experience reads Delta. This feature is the issue.

### C5-L3 exercise 2

**Hints.** (1) Which table feature does Runtime 2.0 add by default? (2) Deletion vectors. (3) Which reader lacks them?

**Answer.** b) A Python notebook (non-Spark) using the default Delta reader Because: r1) Runtime 2.0 writes deletion vectors by default, and Python notebooks do not support deletion vectors.

Runtime 2.0 writes reader 3, writer 7 with deletion vectors. Python notebooks do not support them.

### Common wrong answers.

- a: The SQL analytics endpoint supports deletion vectors.
- c: Direct Lake supports deletion vectors.
- d: The writer's own runtime supports everything it writes.
- b|r2: Right answer, wrong reason. Python notebooks read OneLake. Deletion vectors are the problem.

### C5-L3 exercise 3

**Hints.** (1) The Delta command that returns one row of table details. (2) DESCRIBE ... (3) DESCRIBE.

### Answer.

```
1 DESCRIBE DETAIL gold.events;
```

DESCRIBE DETAIL returns minReaderVersion, minWriterVersion and the table features.

### Common wrong answers.

- g1 = SHOW: SHOW TBLPROPERTIES lists properties. DESCRIBE DETAIL shows protocol versions and table features.

### C5-L3 exercise 4

**Hints.** (1) Store each reader's supported set. (2) Set difference gives what is missing. (3) Index the dict so unknown readers raise KeyError.

### Answer.

```
1 ALL_FEATURES = frozenset(
2 {"deletion_vectors", "column_mapping_name", "column_mapping_id",
3 "timestamp_ntz", "liquid_clustering"}
4)
5 SUPPORTED: dict[str, frozenset[str]] = {
6 "spark_2_0": ALL_FEATURES,
7 "sql_endpoint": ALL_FEATURES - {"column_mapping_id", "timestamp_ntz"},
8 "direct_lake": ALL_FEATURES - {"timestamp_ntz"},
9 "python_notebook": frozenset({"timestamp_ntz"}),
10 }
11
12 def unsupported(reader: str, features: set[str]) -> list[str]:
13 """Table features a Fabric reader does not support.
14
15 Raises:
16 KeyError: For an unknown reader.
```

```
17 """
18 return sorted(features - SUPPORTED[reader])
```

Encoding the matrix lets a release check refuse to publish a gold table with features its readers cannot handle.

### C5-L3 exercise 5

**Hints.** (1) Which readers need this table? (2) Which features do they lack? (3) Check before release.

**Answer.**

- `TIMESTAMP` is supported by both readers. `TIMESTAMP_NTZ` is not.
- Column mapping in `id` mode is not supported by the SQL endpoint.
- A feature check before release catches changes a Databricks upgrade introduces.

Design for the readers that matter and check features on every release.

**Common wrong answers.**

- `time = ntz`: Neither the SQL endpoint nor Direct Lake supports `TIMESTAMP_NTZ`.
- `mapping = id_mode`: The SQL analytics endpoint does not support column mapping in `id` mode.
- `gate = trust`: Open format does not mean every reader supports every feature.

## C6-L1 Star schemas for semantic models

---

### C6-L1 exercise 1

**Hints.** (1) What does one row mean? (2) Which columns say which entities it concerns? (3) The dimension keys.

**Answer.** b) The dimension keys it holds, and their values Because: r1) Dimension key columns set the dimensionality, and their values set the granularity.

One row per combination of the dimension keys: that is the grain.

**Common wrong answers.**

- `a`: Measures are the numbers at that grain. They do not define it.
- `c`: Row count follows from the grain and the data. It does not define it.
- `d`: Import, DirectQuery or Direct Lake changes how data is read, not the grain.

- b|r2: Right answer, wrong reason. Grain comes from the data design.

## C6-L1 exercise 2

**Hints.** (1) Which way should filters flow in a star? (2) Dimension to fact. (3) Bi-directional breaks that.

**Answer.** b) Minimise them. They can hurt performance and confuse report users. Because: r1) Bi-directional filters can slow queries and filter in ways users do not expect.

Single-direction one-to-many relationships keep filtering predictable.

### Common wrong answers.

- a: The guidance is the opposite: minimise them.
- c: Date tables filter facts one way, like other dimensions.
- d: The concern is behaviour and performance, not support.
- b|r2: Right answer, wrong reason. They still exist. Use them sparingly and on purpose.

## C6-L1 exercise 3

**Hints.** (1) What must be decided first? (2) Dimensions provide keys to facts. (3) Relationships before measures.

### Answer.

1. State the fact table's grain
2. Build dimension tables with unique keys in gold
3. Build the fact table with dimension keys and measure columns
4. Create one-to-many, single-direction relationships
5. Write DAX measures

Grain, dimensions, facts, relationships, then measures.

### Common wrong answers.

- Decide the grain before you build the fact table, or you will rebuild it.
- Measures depend on relationships to filter correctly.

## C6-L1 exercise 4

**Hints.** (1) Loop over relationships and build a name. (2) Two independent checks, not elif. (3) Append in the order given.

### Answer.

```
1 def model_warnings(relationships: list[dict[str, str]]) -> list[str]:
2 """Warnings for relationships that break star-schema guidance.
```

```

3
4 Args:
5 relationships: Dicts with from_table, to_table, cardinality and
 direction.
6
7 Returns:
8 Warning strings in input order.
9 """
10 warnings = []
11 for rel in relationships:
12 name = f"{rel['from_table']}-{rel['to_table']}"
13 if rel["cardinality"] == "many_to_many":
14 warnings.append(f"{name}: many-to-many")
15 if rel["direction"] == "both":
16 warnings.append(f"{name}: bi-directional")
17 return warnings

```

Independent checks report every problem. A review script like this runs on models exported with tools such as Tabular Editor.

## C6-L1 exercise 5

**Hints.** (1) What does Direct Lake on SQL support? (2) What does one row mean? (3) Which way should filters flow?

**Answer.**

- Compute age\_band upstream: Direct Lake on SQL endpoints does not support calculated columns.
- One row per respondent, brand and wave gives a clean fact grain.
- Single-direction relationships keep filters predictable.

Do the modelling work in gold, keep the grain clean, and filter one way.

**Common wrong answers.**

- where = calc\_col: Calculated columns are not supported in Direct Lake on SQL analytics endpoints.
- grain = resp: A list column mixes grains and cannot be filtered by brand.
- rel = both: Bi-directional filtering is rarely needed and can inflate results.

## C6-L2 DAX measures that state their filter context

### C6-L2 exercise 1

**Hints.** (1) What was DIVIDE designed for? (2) What is returned without an alternate result? (3) BLANK.

**Answer.** b) BLANK Because: r1) DIVIDE returns BLANK when the denominator is zero or BLANK, unless an alternate result is given.



BLANK keeps empty cells empty in visuals.

### Common wrong answers.

- a: DIVIDE was designed to handle division by zero without an error.
- c: Only if you pass 0 as the alternate result.
- d: That is what the / operator can return. DIVIDE returns BLANK.
- b|r2: Right answer, wrong reason. The / operator does divide by zero. DIVIDE handles it.

### C6-L2 exercise 2

**Hints.** (1) Which table does ALL name? (2) Filters from dimensions reach the fact table. (3) Remove only the brand filter.

**Answer.** b) Every filter on fact\_awareness, including wave and region  
Because: r1) ALL on the fact table removes all its filters, including those that reach it through dimensions.

Use ALL(dim\_brand) to compare against all brands within the same wave and region.

### Common wrong answers.

- a: That is ALL(dim\_brand). ALL(fact\_awareness) removes much more.
- c: ALL changes filter context for every filter on the table, not just slicers.
- d: Filter context comes from the visual and the page. ALL removes the table's filters wherever they come from.
- b|r2: Right answer, wrong reason. ALL removes filters from the table or column you give it.

### C6-L2 exercise 3

**Hints.** (1) The DAX keyword for a variable. (2) The safe division function. (3) VAR and DIVIDE.

### Answer.

```
1 Weighted aware % =
2 VAR aware_w = SUMX (fact_awareness, fact_awareness[aware] *
 fact_awareness[weight])
3 VAR base_w = SUM (fact_awareness[weight])
4 RETURN DIVIDE (aware_w, base_w)
```

Variables name each part. DIVIDE handles a zero base.

### Common wrong answers.

- g1 = LET: DAX variables use VAR ... RETURN.

## C6-L2 exercise 4

**Hints.** (1) Two variables, then RETURN. (2) ALL on the brand dimension only. (3) DIVIDE the two.

**Answer.**

```
1 Share of aware across brands =
2 VAR this_brand = [Weighted aware %]
3 VAR all_brands = CALCULATE ([Weighted aware %], ALL (dim_brand))
4 RETURN DIVIDE (this_brand, all_brands)
```

Removing only the brand filter keeps the comparison inside the same wave and region, and DIVIDE handles an empty denominator.

## C6-L2 exercise 5

**Hints.** (1) Does the value depend on the visual's filters? (2) Is it a property of one row? (3) Calculated columns are not an option here.

**Answer.**

- Age band is a property of each respondent: compute it in gold.
- Percentages depend on the visual's filters: measures.
- The base count also depends on filters: a measure, shown beside every percentage.

Row properties in gold, filter-dependent arithmetic in measures.

**Common wrong answers.**

- ageband = `calc_col`: Direct Lake on SQL analytics endpoints does not support calculated columns.
- share = `gold`: A stored percentage cannot re-aggregate correctly when users filter or group differently.
- base = `gold`: The base changes with every filter. It must be computed at query time.

## C6-L3 Direct Lake, framing and the fallback you did not plan

---

### C6-L3 exercise 1

**Hints.** (1) Does Direct Lake copy data? (2) What does it need to know about the tables? (3) Which version to read.

**Answer.** b) Frames the model: records which Delta table versions queries will use. Because: r1) Framing copies metadata only, takes seconds, and sets the point in time queries see.

Framing takes seconds and fixes the table versions queries see until the next frame.

**Common wrong answers.**

- a: Import copies data. Direct Lake refresh copies metadata only.
- c: The model reads Delta tables. It never writes them.
- d: Refresh does not change storage mode.
- b|r2: Right answer, wrong reason. They do refresh. The refresh is framing.

**C6-L3 exercise 2**

**Hints.** (1) Can a view be read as Delta files? (2) What does Automatic do when direct loading fails? (3) Fallback.

**Answer.** b) The query falls back to DirectQuery and runs more slowly. Because: r1) Views cannot be loaded directly from Delta, and Automatic falls back silently.

Automatic keeps the report working through DirectQuery, silently. Replace views with Delta tables.

**Common wrong answers.**

- a: That is DirectLakeOnly, or Direct Lake on OneLake, which never falls back.
- c: Fabric does not materialise views for you. Write a Delta table instead.
- d: A view has no Delta files to read.
- b|r2: Right answer, wrong reason. Fallback is slower.

**C6-L3 exercise 3**

**Hints.** (1) Which option errors when Direct Lake cannot be used? (2) Which is the default? (3) `DirectLakeOnly` and `Automatic`.

**Answer.**

- ```
1 Development model: Direct Lake behavior = DirectLakeOnly
2 Production model: Direct Lake behavior = Automatic
```

`DirectLakeOnly` exposes fallback causes during development. `Automatic` protects users in production.

Common wrong answers.

- `g1 = DirectQueryOnly`: `DirectQueryOnly` always uses `DirectQuery`. It measures fallback cost but does not reveal where fallback happens.

C6-L3 exercise 4

Hints. (1) A dict from SKU to (files, rows). (2) Strictly greater than the limit is a problem. (3) Index the dict so unknown SKUs raise `KeyError`.

Answer.

```
1 SMALL = (1000, 300)
2 GUARDRAILS: dict[str, tuple[int, float]] = {
3     "F2": SMALL, "F4": SMALL, "F8": SMALL, "F16": SMALL, "F32": SMALL,
4     "F64": (5000, 1500), "F128": (5000, 3000), "F256": (5000, 6000),
5 }
6
7
8 def guardrail_problems(sku: str, parquet_files: int, rows_millions:
9     float) -> list[str]:
10     """Direct Lake per-table guardrails a table exceeds on a SKU.
11     Raises:
12         KeyError: For a SKU not in the table.
13     """
14     max_files, max_rows = GUARDRAILS[sku]
15     problems = []
16     if parquet_files > max_files:
17         problems.append("files")
18     if rows_millions > max_rows:
19         problems.append("rows")
20     return problems
```

A nightly check against the guardrails catches small-file build-up before Direct Lake falls back or framing fails.

C6-L3 exercise 5

Hints. (1) How many files does each append add? (2) What merges small files? (3) OPTIMIZE after the write.

Answer. The bug is on line 3.

Fixed version:

```
1 from pyspark.sql import DataFrame
2
3
4 def publish_hourly(df: DataFrame, table: str) -> None:
5     """Append the latest hour, then compact small files for Direct
6     Lake."""
7     df.repartition(1).write.mode("append").saveAsTable(table)
8     df.sparkSession.sql(f"OPTIMIZE {table}")
```

Hourly appends add a file each hour. OPTIMIZE compacts them, keeping the table under the guardrail. On busy tables, run it on a schedule rather than every hour.

Common wrong answers.

- Line 2: The docstring is fine.
- Line 4: Returning None is harmless.

C7-L1 Eventstream, events in, routed out

C7-L1 exercise 1

Hints. (1) Which format does a lakehouse use for tables? (2) Delta. (3) Eventhouse is a different destination.

Answer. b) Delta Lake tables Because: r1) The lakehouse destination converts real-time events to Delta Lake format in lakehouse tables.

Lakehouse destinations write Delta tables. Eventhouse destinations write KQL tables.

Common wrong answers.

- a: Events become Delta tables, which Spark, SQL and Direct Lake can read.
- c: Events become Delta tables.
- d: KQL tables live in an eventhouse, a different destination.
- b|r2: Right answer, wrong reason. Eventhouse destinations store data in KQL databases.

C7-L1 exercise 2

Hints. (1) Writing to an item needs which role? (2) Lesson C1-L3. (3) Contributor.

Answer. b) Contributor or higher Because: r1) Microsoft Learn requires Contributor or higher on the destination workspace for eventhouse, lakehouse and Activator destinations.

Contributor or higher on the destination workspace.

Common wrong answers.

- a: Viewers cannot write to items (lesson C1-L3).
- c: Admin is more than needed.
- d: Access to the destination workspace is required.

- b|r2: Right answer, wrong reason. Contributor applies to the workspace where the role is granted.

C7-L1 exercise 3

Hints. (1) Source first. (2) Protect the raw data before filtering. (3) Verify at the end.

Answer.

1. Connect the source
2. Route every event to a raw table
3. Create derived streams with filters for each consumer
4. Connect each derived stream to its destination
5. Compare event counts at source and destinations

Source, raw landing, derived streams, destinations, then a count check.

Common wrong answers.

- Store the raw stream first, so nothing is lost while you design filters.
- Counts can only be compared once destinations are connected and receiving events.

C7-L1 exercise 4

Hints. (1) Start with `telemetry_raw` for every event. (2) Check for a missing key before reading it. (3) Add fault destinations after the raw one.

Answer.

```

1 def route(event: dict[str, object]) -> list[str]:
2     """Destinations for one telemetry event. Every event is kept raw.
3
4     Args:
5         event: Parsed event.
6
7     Returns:
8         Destination names in routing order.
9     """
10    destinations = ["telemetry_raw"]
11    if "event_type" not in event:
12        destinations.append("quarantine")
13    elif event["event_type"] == "fault":
14        destinations += ["activator_faults", "bronze.vehicle_faults"]
15    return destinations

```

Raw first, always. Malformed events are kept and flagged, not lost.

C7-L1 exercise 5

Hints. (1) Which source captures every change? (2) Which store answers queries in seconds? (3) Which store does Spark read?

Answer.

- The CDC source streams inserts, updates and deletes after an initial snapshot.
- An eventhouse serves low-latency KQL queries for dashboards.
- A lakehouse destination writes Delta tables for Spark.

One CDC source, two destinations, each suited to its consumer.

Common wrong answers.

- source = `poll`: Polling misses deletes and adds load. CDC streams every change.
- live = `lakehouse`: Lakehouse tables suit batch analysis. Live dashboards query an eventhouse with KQL.
- models = `custom`: Spark reads Delta tables in a lakehouse directly.

C7-L2 Eventhouse and KQL for time-series questions

C7-L2 exercise 1

Hints. (1) Does take sort? (2) Is its result consistent? (3) Use top for ordered results.

Answer. c) 10 arbitrary rows, possibly different on each run Because: r1) take guarantees nothing about which records it returns unless the data is sorted.

take is for browsing. top is for answers.

Common wrong answers.

- a: Use top 10 by event_time desc for the most recent rows.
- b: Ingestion order is not guaranteed either.
- d: take runs without error. It just promises nothing about which rows.
- c|r2: Right answer, wrong reason. There is no fixed seed. Results can change between runs.

C7-L2 exercise 2

Hints. (1) Which rows does summarize group? (2) Are there rows in those bins? (3) No.

Answer. b) No rows for those bins Because: r1) summarize groups existing rows. Bins with no rows do not appear. make-series fills them.

Charts draw a line straight across the gap. Use make-series with a default to show the outage.

Common wrong answers.

- a: That is make-series with default = 0. summarize only returns groups that exist.
- c: No rows are created at all.
- d: Empty bins are not an error.
- b|r2: Right answer, wrong reason. There are no rows to count, so no group is formed.

C7-L2 exercise 3

Hints. (1) A function for 'one hour before now'. (2) The operator that sorts and limits. (3) ago and top.

Answer.

```
1 telemetry_raw
2 | where event_time > ago(1h)
3 | top 10 by event_time desc
```

ago(1h) sets the window. top 10 by event_time desc returns the latest ten.

Common wrong answers.

- g2 = take: take does not accept a sort and returns arbitrary rows.

C7-L2 exercise 4

Hints. (1) Filter with where and ago first. (2) make-series fills missing steps with the default. (3) ... default = 0 on event_time step 1h by vehicle_id.

Answer.

```
1 telemetry_raw
2 | where event_time > ago(24h)
3 | make-series avg_rpm = avg(rpm) default = 0 on event_time step 1h by
   vehicle_id
```

make-series returns one array per vehicle with every hour present, so gaps show as zeros.

C7-L2 exercise 5

Hints. (1) Who needs the raw data, and for how long? (2) Which period is actually queried? (3) Aggregates are small.

Answer.

- Keep raw data as long as audits need, and no longer: 90 days.
- Cache the 30 days people query, to balance speed and cost.
- Keep small aggregates for years, where trends need them.

Retention follows need, cache follows query patterns, and aggregates outlive raw data.

Common wrong answers.

- `raw_ret = keep_default`: Ten years of raw telemetry costs storage nobody uses. Set retention from the need.
- `raw_ret = d7`: Audits need 90 days. Data past retention is removed automatically.
- `raw_cache = c90`: Caching data nobody queries raises cost without speeding up real queries.

C7-L3 Activator, rules that act on events

C7-L3 exercise 1

Hints. (1) Is this a change or a state condition? (2) How many times does the value cross 100 from below? (3) Once.

Answer. a) Once Because: r1) Increases above is a change condition. It fires when the value crosses the threshold from below.

Change conditions fire on transitions. State conditions fire on every matching event.

Common wrong answers.

- `b`: That is Is greater than, a state condition that fires on each matching event.
- `c`: The crossing from below is the trigger. No decrease is needed.
- `d`: Change conditions do not repeat on a timer.
- `a|r2`: Right answer, wrong reason. There is no daily limit of one. The rule fires once because there is one crossing.

C7-L3 exercise 2

Hints. (1) What arrives from a silent sensor? (2) Nothing. Which condition works on absence? (3) Heartbeat.

Answer. b) No presence of data Because: r1) No presence of data fires when no new events arrive within a set time. Other conditions need an event to evaluate.

Every value-based condition needs an event. Only the heartbeat condition sees silence.

Common wrong answers.

- a: No event arrives, so there is no value to compare.
- c: Becomes false needs a new event with the value false.
- d: Changes needs a new event with a different value.
- b|r2: Right answer, wrong reason. A silent sensor sends nothing at all.

C7-L3 exercise 3

Hints. (1) Which column identifies one freezer? (2) The occurrence option reads 'When it has been ... for'. (3) `freezer_id` and `true`.

Answer.

```
1 Object: freezer, keyed by freezer_id
2 Condition: temperature Is greater than -15
3 Occurrence: When it has been true for 10 minutes
```

Key by the object. The occurrence turns a state condition into one alert for a sustained problem.

Common wrong answers.

- `g1 = event_id`: Every event has its own ID, so no state carries across readings. Key by the freezer.

C7-L3 exercise 4

Hints. (1) Keep a dict from `object_id` to 'was above'. (2) Fire when the value is above and the object was not. (3) Update the object's state after every event.

Answer.

```
1 def activations(events: list[tuple[str, float]], threshold: float) ->
  list[int]:
2     """Indexes of events where an object's value crosses above a
    threshold.
3
4     Args:
5         events: (object_id, value) pairs in arrival order.
6         threshold: Value that must be exceeded.
7
8     Returns:
9         Event indexes where the object moved from at-or-below to above.
```

```

10     """
11     above: dict[str, bool] = {}
12     fired: list[int] = []
13     for index, (object_id, value) in enumerate(events):
14         is_above = value > threshold
15         if is_above and not above.get(object_id, False):
16             fired.append(index)
17         above[object_id] = is_above
18     return fired

```

Per-object state is what makes the rule mean "this oven crossed 80", not "some reading somewhere crossed 80".

C7-L3 exercise 5

Hints. (1) Which ID identifies one van? (2) How many readings arrive per minute? (3) Which condition works without events?

Answer.

- Key by `van_id`, so each van keeps its own state.
- A sustained state condition sends one message per incident, well under the 30 per rule per recipient per hour limit.
- A heartbeat rule detects silence, which no value-based condition sees.

Per-object keys, sustained conditions for noisy streams, and a heartbeat for silence.

Common wrong answers.

- `object = fleet`: One object mixes 400 vans. A cool reading from one van clears a hot one.
- `hot = every`: At one reading every 5 seconds, a hot van sends 12 messages a minute and hits the 30-per-hour limit in under 3 minutes.
- `silent = speed_zero`: A parked van reports speed 0. A silent van reports nothing.

C8-L1 Copilot in Fabric, where it runs and what it costs

C8-L1 exercise 1

Hints. (1) Copilot consumes CUs. Where do CUs come from? (2) Paid, not trial. (3) F2 or P1 and up.

Answer. b) A paid F2 or higher, or P1 or higher, in a supported region
Because: r1) Microsoft Learn lists a paid F2+ or P1+ capacity in a supported region, with tenant settings enabled.

Paid F2+ or P1+, a supported region, and tenant settings switched on.

Common wrong answers.

- a: Copilot stops working when a trial expires.
- c: F64 matters for free-licence viewers (lesson C1-L1). Copilot starts at F2.
- d: Copilot consumes capacity units from the workspace's capacity.
- b|r2: Right answer, wrong reason. Copilot depends on the capacity and tenant settings, not a Pro licence.

C8-L1 exercise 2

Hints. (1) Is Copilot one feature or many? (2) Check the release-status page per workload. (3) Record the date: status changes.

Answer. b) Notebooks preview, Data Factory GA, with the date you checked
Because: r1) Microsoft Learn's release-status page lists each Copilot experience separately.

Status is per experience. Record each one with the date you checked.

Common wrong answers.

- a: Copilot in notebooks is preview (Learn, September 2026).
- c: Copilot in Data Factory is GA (Learn, September 2026).
- d: Status differs by workload and changes over time. Record each one with a date.
- b|r2: Right answer, wrong reason. Status is per feature. Some Data Factory features are preview.

C8-L1 exercise 3

Hints. (1) Output and input rates differ. (2) Output is four times the input rate. (3) 400 and 400.

Answer.

```
1 input: 1,000 tokens x 100 / 1,000 = 100 CU seconds
2 output: 1,000 tokens x 400 / 1,000 = 400 CU seconds
3 total: 500 CU seconds
```

Output costs 400 CU seconds per 1,000 tokens. Long answers cost more than long questions.

Common wrong answers.

- g1 = 100: Output tokens cost more than input: 400 CU seconds per 1,000.

C8-L1 exercise 4

Hints. (1) Three rates: 100, 10 and 400 per 1,000 tokens. (2) Uncached input = input_tokens - cached_input_tokens. (3) Validate before you compute.

Answer.

```
1 INPUT_RATE = 100
2 CACHED_INPUT_RATE = 10
3 OUTPUT_RATE = 400
4 TOKENS_PER_UNIT = 1000
5
6
7 def copilot_cu_seconds(
8     input_tokens: int, output_tokens: int, cached_input_tokens: int = 0
9 ) -> float:
10     """CU seconds for one Copilot request at the published token rates.
11
12     Args:
13         input_tokens: All input tokens, cached or not.
14         output_tokens: Output tokens.
15         cached_input_tokens: Input tokens served from the prompt cache.
16
17     Returns:
18         Capacity unit seconds consumed.
19
20     Raises:
21         ValueError: If a count is negative or cached exceeds input.
22     """
23     if min(input_tokens, output_tokens, cached_input_tokens) < 0:
24         raise ValueError("token counts must not be negative")
25     if cached_input_tokens > input_tokens:
26         raise ValueError("cached tokens cannot exceed input tokens")
27     uncached = input_tokens - cached_input_tokens
28     total = (
29         uncached * INPUT_RATE
30         + cached_input_tokens * CACHED_INPUT_RATE
31         + output_tokens * OUTPUT_RATE
32     )
33     return total / TOKENS_PER_UNIT
```

Charge each token once, at its own rate. Put the rates in named constants so you update one place when prices change.

C8-L1 exercise 5

Hints. (1) Tenant settings take security groups. (2) Is Copilot one feature? (3) Who owns production logic?

Answer.

- Scope the tenant setting to the security group that needs Copilot.

- Record each experience with its status and date, since status differs by workload.
- Treat Copilot output as a draft: a person reviews, tests and owns it.

Scoped access, a dated status record, and human ownership of every output.

Common wrong answers.

- scope = `everyone`: Scope tenant settings to the groups that need them. Every request draws on the F16.
- record = `copilot_yes`: Status is per experience. 'Yes' tells an auditor nothing.
- prod = `direct`: Running without error is not correctness. Lesson C8-L3 shows how to check.

C8-L2 Data agents, questions in plain language

C8-L2 exercise 1

Hints. (1) Which query types does the agent generate? (2) Only read queries. (3) Read-only holds for Admins too.

Answer. b) The agent generates only read queries, so it does not delete data Because: r1) Microsoft Learn states the data agent generates only read queries and keeps read-only connections.

Data agents are read-only by design.

Common wrong answers.

- a: The agent never generates create, update or delete queries.
- c: Read-only applies to every user, whatever their role.
- d: There is no approval flow. The agent does not write.
- b|r2: Right answer, wrong reason. The agent uses the user's permissions. It is read-only by design, whatever those permissions allow.

C8-L2 exercise 2

Hints. (1) Whose credentials does the agent use? (2) The person asking. (3) Ana cannot read EU sales.

Answer. b) Answers only from data Ana has permission to read Because: r1) The agent uses the requesting user's credentials and permissions.

The agent never widens access. Grant permissions on the data, not on the agent.

Common wrong answers.

- a: The agent runs as the person asking, not as its owner.
- c: The agent applies the user's permissions. It does not add its own masking.
- d: The agent queries data. It does not invent estimates for data it cannot read.
- b|r2: Right answer, wrong reason. There is no automatic country filter. Ana's own permissions decide what she reaches.

C8-L2 exercise 3

Hints. (1) What do examples get validated against? (2) Context before testing. (3) Test before publishing.

Answer.

1. Select only the tables and columns the questions need
2. Write instructions that define business terms
3. Add example question and SQL pairs
4. Run the ground-truth question set and record accuracy
5. Publish and share the agent

Schema, instructions, examples, evaluation, then publish.

Common wrong answers.

- Measure accuracy before users rely on the agent.
- Example queries are validated against the selected schema. Select the tables first.

C8-L2 exercise 4

Hints. (1) Use `answer is None` to detect a missing number. (2) `math.isclose` handles the tolerance. (3) Divide by `len(expected)`.

Answer.

```
1 import math
2
3
4 def evaluate(
5     expected: dict[str, float], actual: dict[str, float | None], rel_tol:
6     float = 0.001
7 ) -> tuple[float, list[str]]:
8     """Accuracy of a data agent against a ground-truth set.
9     Args:
10         expected: Question id to ground-truth answer.
```

```

11     actual: Question id to the agent's answer, None if it gave no
12     number.
13     rel_tol: Relative tolerance for a match.
14     Returns:
15     Share of expected questions answered correctly, and sorted failed
16     ids.
17     Raises:
18     ValueError: If expected is empty.
19     """
20     if not expected:
21         raise ValueError("ground-truth set is empty")
22     failed = []
23     for question, truth in expected.items():
24         answer = actual.get(question)
25         if answer is None or not math.isclose(answer, truth,
26         rel_tol=rel_tol):
27             failed.append(question)
28     return (len(expected) - len(failed)) / len(expected), sorted(failed)

```

The denominator is the ground-truth set. A zero is an answer. Both mistakes inflate or deflate accuracy without an error.

C8-L2 exercise 5

Hints. (1) Which file types are supported? (2) Which source types accept example queries? (3) Two sources, same numbers: how does the agent choose?

Answer.

- Data agents do not support PDF, Word or text files.
- Example queries work on lakehouse, warehouse and KQL sources, not semantic models.
- Two sources hold the same numbers, so tell the agent which to use for which question.

Supported sources only, examples where they work, and routing instructions where sources overlap.

Common wrong answers.

- `pdfs = include`: Unstructured files such as PDFs are not supported data agent sources.
- `examples = model`: Semantic model sources do not support example queries.
- `routing = no_instr`: With two sources of the same numbers, the agent needs an instruction to route questions.

C8-L3 Verify what Copilot writes

C8-L3 exercise 1

Hints. (1) What does the engine check before running? (2) Does it know what answer you expected? (3) Syntax and names only.

Answer. b) The syntax parses and the referenced names exist Because: r1) An engine checks syntax and names before it runs a query. It does not know what answer you wanted.

A clean run is the start of verification, not the end.

Common wrong answers.

- a: An average of averages runs cleanly and returns a wrong number.
- c: A query that scans a whole table runs without error too.
- d: The engine does not know the grain you intended.
- b|r2: Right answer, wrong reason. Microsoft says to review all generated code. The engine did the checking, not Copilot.

C8-L3 exercise 2

Hints. (1) Does every store weigh the same in the result? (2) Should it? (3) Total sales over total baskets.

Answer. The bug is on line 2.

Fixed version:

```
1 SELECT r.region,  
2       SUM(s.total_sales) / NULLIF(SUM(s.basket_count), 0) AS avg_basket  
3 FROM dbo.store_summary AS s  
4 LEFT JOIN dbo.dim_region AS r ON s.store_id = r.store_id  
5 GROUP BY r.region;
```

Recompute ratios from their totals at the level you report.

Common wrong answers.

- Line 4: A left join keeps stores without a region. That line is fine.
- Line 5: Grouping by region matches the requirement.

C8-L3 exercise 3

Hints. (1) Cheapest check first. (2) Which check gives you a known right answer? (3) Keep what you built.

Answer.

1. Read every line: names, grain, join types
2. Run it on a small dataset with a hand-worked answer

3. Compare the result with a second method on real data
4. Keep the known-data and comparison checks as tests

Read, known data, second method, then keep the checks.

Common wrong answers.

- Read first. A wrong join or invented column is cheaper to catch on paper.
- Known data first: it tells you which answer is right when the two methods disagree.

C8-L3 exercise 4

Hints. (1) Wrap only the candidate call in try/except. (2) Handle None before math.isclose, which rejects None. (3) `expected is not got` is True when one is None and the other is not.

Answer.

```
1 import math
2 from collections.abc import Callable
3
4
5 def mismatches(
6     candidate: Callable[[list[float]], float | None],
7     reference: Callable[[list[float]], float | None],
8     cases: list[list[float]],
9 ) -> list[int]:
10     """Indexes of cases where a candidate function disagrees with a
11         reference.
12
13     Args:
14         candidate: Function under test, such as Copilot-generated code.
15         reference: Trusted implementation.
16         cases: Inputs to run through both.
17
18     Returns:
19         Case indexes where the candidate raised or returned a different
20         result.
21     """
22     failed: list[int] = []
23     for index, case in enumerate(cases):
24         expected = reference(case)
25         try:
26             got = candidate(case)
27         except Exception:
28             failed.append(index)
29             continue
30         if expected is None or got is None:
31             if expected is not got:
32                 failed.append(index)
33         elif not math.isclose(got, expected):
```

```
32         failed.append(index)
33     return failed
```

A comparison harness turns "looks right" into a list of failing cases. Keep it with the pipeline tests.

C8-L3 exercise 5

Hints. (1) Which dataset has an answer you know? (2) Is the second method independent? (3) What checks the next change?

Answer.

- Hand-built rows cover the cases that break code quietly. A random sample rarely contains them.
- A second, independent method on the same rows catches logic errors.
- A committed test checks every later change, including the next Copilot edit.

Known awkward data, an independent second method, and a committed test.

Common wrong answers.

- `data = sample`: You have no hand-worked answer for a random sample, and it seldom contains the awkward cases.
- `second = rerun`: The same code gives the same wrong answer twice. The second method must be independent.
- `keep = note`: A note does not run. A test fails when a later change breaks the logic.

C9-L1 Domains and endorsement, finding the right data

C9-L1 exercise 1

Hints. (1) What are domains for? (2) Finding data. (3) Access stays with roles and permissions.

Answer. b) Nothing about access. The workspace's items show under Finance in the OneLake catalog Because: r1) Domain assignment does not affect item visibility or access. Workspace roles and item permissions do.

Domains organise discovery. Roles and permissions control access.

Common wrong answers.

- `a`: Domains do not restrict access.
- `c`: Domains do not grant access either.

- **d**: Access requests go through workspace roles and item permissions, not domains.
- **b|r2**: Right answer, wrong reason. Every item in the workspace inherits the domain, not only Power BI items.

C9-L1 exercise 2

Hints. (1) Promoted and Certified differ in who applies them. (2) Who names the certifiers? (3) A Fabric admin.

Answer. b) They can request certification. Only users a Fabric admin specifies can certify Because: r1) Any user can request certification, but only users specified by a Fabric administrator can certify.

Promotion is self-service. Certification is an authorised review.

Common wrong answers.

- **a**: Contributors can promote. Certifying needs a name on the Fabric admin's list.
- **c**: Certification is a human review. No test applies it.
- **d**: Any item except a Power BI dashboard can be certified.
- **b|r2**: Right answer, wrong reason. The workspace Admin role does not grant certification rights. The Fabric admin's list does.

C9-L1 exercise 3

Hints. (1) Tenant-level role first. (2) The domain role for workspace admins. (3) **admin** and **contributor**.

Answer.

- 1 Creates domains and names domain admins: Fabric admin
- 2 Assigns their own workspace to a domain: domain contributor, who must be a workspace admin

Fabric admins create domains. Domain contributors, who are workspace admins, assign their own workspaces.

Common wrong answers.

- g2 = **admin**: Domain admins associate workspaces too, but the role for workspace admins assigning their own workspaces is domain contributor.

C9-L1 exercise 4

Hints. (1) Handle the dashboard first. (2) Start from promoted and certified. (3) Add **master_data** only for data item types.

Answer.

```

1 DATA_ITEM_TYPES = frozenset({"lakehouse", "warehouse", "semantic_model",
    "kql_database"})
2 NO_BADGE_TYPES = frozenset({"dashboard"})
3
4
5 def allowed_badges(item_type: str) -> set[str]:
6     """Endorsement badges an item type can receive.
7
8     Args:
9         item_type: Fabric item type, such as "lakehouse" or "report".
10
11     Returns:
12         Allowed badge names.
13     """
14     if item_type in NO_BADGE_TYPES:
15         return set()
16     badges = {"promoted", "certified"}
17     if item_type in DATA_ITEM_TYPES:
18         badges.add("master_data")
19     return badges

```

Encode the rules as constants and small branches, so a rule change touches one line.

C9-L1 exercise 5

Hints. (1) Which feature gives catalog filtering? (2) Which badge marks the single source of truth? (3) What controls access?

Answer.

- A domain gives filtering in the OneLake catalog and a home for delegated settings.
- One master data copy, with shortcuts for other uses (lesson C3-L1).
- Access comes from roles, never from domains.

Domains for discovery, master data for the authoritative copy, roles for access.

Common wrong answers.

- domain = `naming`: A prefix helps a little. A domain adds catalog filtering, delegated settings and named owners.
- customer = `promote_all`: Three promoted copies tell users nothing about which one to trust.
- restrict = `by_domain`: Domain assignment does not affect access.

C9-L2 OneLake security, roles down to rows and columns

C9-L2 exercise 1

Hints. (1) Within a role: intersection. Across roles? (2) Union. (3) Row filters combine with OR.

Answer. a) North and South rows Because: r1) Roles combine as a union. Row filters from different roles combine with OR.

Across roles, access adds up. Plan roles so their union is what you intend.

Common wrong answers.

- **b**: Across roles, filters combine with OR, not AND.
- **c**: Role order does not matter. Access is the union.
- **d**: Only rows allowed by at least one role.
- **a|r2**: Right answer, wrong reason. OneLake unions the roles. It does not pick one.

C9-L2 exercise 2

Hints. (1) Which users does OneLake security govern? (2) Viewers and Read-permission users. (3) Contributors have write access already.

Answer. b) All rows: Contributors already have data access through their workspace role Because: r1) OneLake security governs Viewers and users with Read permission. Admins, Members and Contributors already have data access.

Restrict readers by making them Viewers first.

Common wrong answers.

- **a**: Contributors are not restricted by OneLake security roles.
- **c**: The workspace role gives them full data access.
- **d**: The read succeeds with all rows.
- **b|r2**: Right answer, wrong reason. Spark enforces OneLake security (GA). The Contributor role is why it does not apply here.

C9-L2 exercise 3

Hints. (1) The engines behave differently. (2) Which one is stricter? (3) The SQL endpoint.

Answer. a) Spark returns the allowed columns. The endpoint returns an error
Because: r1) Microsoft Learn documents that CLS behaves differently: Spark shows only allowed columns, the endpoint blocks the query.

List columns explicitly in SQL queries that users with CLS will run.

Common wrong answers.

- b: The SQL analytics endpoint blocks queries that touch hidden columns.
- c: Spark succeeds and omits the hidden column.
- d: Hidden columns are removed or blocked, not nulled.
- a|r2: Right answer, wrong reason. Spark enforces CLS by dropping the hidden column.

C9-L2 exercise 4

Hints. (1) Collect the column sets from roles that include the table. (2) No grants: return an empty set. (3) `set().union(*grants)` for Spark, `set.intersection(*grants)` for the endpoint.

Answer.

```
1 from typing import Literal
2
3
4 def visible_columns(
5     roles: list[dict[str, set[str]]],
6     table: str,
7     engine: Literal["spark", "sql_endpoint"],
8 ) -> set[str]:
9     """Columns a user sees in one table across their OneLake security
10        roles.
11
12    Args:
13        roles: Each role maps table names to the columns it allows.
14        table: Table being queried.
15        engine: "spark" (union across roles) or
16        "sql_endpoint" (intersection).
17
18    Returns:
19        Visible column names. Empty if no role includes the table.
20
21    Raises:
22        ValueError: If the engine is not recognised.
23    """
24    if engine not in ("spark", "sql_endpoint"):
25        raise ValueError(f"unknown engine: {engine}")
26    grants = [role[table] for role in roles if table in role]
27    if not grants:
28        return set()
29    if engine == "spark":
```

```
28     return set().union(*grants)
29     return set.intersection(*grants)
```

The same roles give different column sets in different engines. Test the engine your users query.

C9-L2 exercise 5

Hints. (1) Who does OneLake security govern? (2) Where do related row and column rules go? (3) Which Direct Lake mode enforces OneLake security?

Answer.

- Viewers are governed by OneLake security. Contributors are not.
- Row and column rules for the same table belong in one role. Split roles with different columns and a row filter are rejected.
- Direct Lake on OneLake enforces OneLake security (GA).
- DefaultReader grants full reads, and roles combine by union, so it overrides your restrictions.

Viewers, one role per audience, a supported engine mode, and no leftover default grants.

Common wrong answers.

- `ws_role = contributor`: Contributors already have data access. OneLake roles do not restrict them.
- `roles = split`: Two roles with different columns where one filters rows is a combination OneLake does not support. Keep them in one role.
- `model = import`: Direct Lake on OneLake is on the list of engines that enforce OneLake security. An import model holds its own copy and needs its own security rules.
- `default = keep`: Roles combine as a union. DefaultReader's full access wins over your filters.

C9-L3 Sensitivity labels, Purview and lineage

C9-L3 exercise 1

Hints. (1) Was the report's label set by hand? (2) Can inheritance downgrade? (3) No to both.

Answer. b) It stays Highly Confidential Because: r1) Downstream inheritance never overwrites a manually applied label, and never downgrades a label.

Inheritance only raises or fills in labels. It never lowers them or overrides a person.

Common wrong answers.

- a: Inheritance never replaces a label with a less restrictive one, and never touches a manual label.
- c: The report's label was set by hand, so inheritance leaves it alone.
- d: Inheritance does not remove labels.
- b|r2: Right answer, wrong reason. Reports do inherit labels from their semantic model. This one was set manually.

C9-L3 exercise 2

Hints. (1) Which feature follows dependencies? (2) It starts from lineage. (3) Impact analysis.

Answer. a) Impact analysis Because: r1) Impact analysis shows affected downstream items and workspaces and can notify their contacts.

Lineage shows the graph. Impact analysis reads it from one item downwards.

Common wrong answers.

- b: Domains group items for discovery. They do not track dependencies.
- c: Labels classify sensitivity, not dependencies.
- d: The Metrics app shows capacity use (lesson C10-L3).
- a|r2: Right answer, wrong reason. Impact analysis informs and notifies. It does not block changes.

C9-L3 exercise 3

Hints. (1) Named users keep what they already have. (2) Everyone else? (3) keeps and blocked.

Answer.

- 1 Item labelled Highly Confidential, policy allows the HR group.
- 2 HR group member with Viewer access: keeps their Viewer access.
- 3 Finance user with Viewer access: blocked.

A protection policy never adds permissions. It removes access from everyone it does not name.

Common wrong answers.

- g2 = keeps: Users not named in the policy are blocked, even with a workspace role.

C9-L3 exercise 4

Hints. (1) Validate both labels against order first. (2) Unlabelled takes upstream. Manual keeps current. (3) Compare `order.index` values.

Answer.

```
1 def inherit_label(  
2     current: str | None, manual: bool, upstream: str, order: list[str]  
3 ) -> str | None:  
4     """Downstream label after the upstream label changes.  
5  
6     Args:  
7         current: The downstream item's label, or None.  
8         manual: Whether a person applied the current label.  
9         upstream: The upstream item's new label.  
10        order: Labels from least to most restrictive.  
11  
12    Returns:  
13        The downstream item's label after inheritance.  
14  
15    Raises:  
16        ValueError: If a label is not in order.  
17    """  
18    for label in (current, upstream):  
19        if label is not None and label not in order:  
20            raise ValueError(f"unknown label: {label}")  
21    if current is None:  
22        return upstream  
23    if manual:  
24        return current  
25    if order.index(upstream) > order.index(current):  
26        return upstream  
27    return current
```

Encoding the rule makes its asymmetry visible: labels move up by inheritance, never down.

C9-L3 exercise 5

Hints. (1) What travels with an Excel export? (2) What blocks a person with a link? (3) What lists downstream items?

Answer.

- A label on the model flows to its reports and stays on Excel exports.
- A protection policy blocks anyone outside the HR group, whatever they were sent.
- Impact analysis names what breaks before it breaks.

Label, protect, and check impact before change.

Common wrong answers.

- label = `unlabelled`: Workspace restrictions stop at the workspace. Labels travel with exports.

- policy = `trust`: A protection policy blocks access even when a link reaches the wrong person.
- change = `rename`: Fixing after the break costs a morning of errors across workspaces. Check first.

C10-L1 Git integration for Fabric workspaces

C10-L1 exercise 1

Hints. (1) Does Git store data or definitions? (2) Definitions. (3) Tables are not tracked.

Answer. b) Metadata only, such as its name, shortcuts and OneLake security roles Because: r1) Microsoft Learn: only lakehouse metadata is tracked. Git operations never overwrite table or file data.

Git versions your items. Pipelines rebuild your data.

Common wrong answers.

- a: Tables and files are not tracked in Git.
- c: Lakehouses are supported. Only their metadata is tracked.
- d: Git is not a data backup. Plan data recovery separately.
- b|r2: Right answer, wrong reason. Git integration tracks item definitions by design, not data.

C10-L1 exercise 2

Hints. (1) Both sides changed. (2) What does the source control panel show then? (3) A conflict.

Answer. a) A conflict you must resolve Because: r1) An item changed in both the workspace and the branch is a conflict, shown in the source control panel.

Commit often, and work in your own workspace, so conflicts stay rare.

Common wrong answers.

- b: Changes on both sides are flagged as a conflict first.
- c: The workspace does not win by default.
- d: There is no automatic merge. You resolve the conflict.
- a|r2: Right answer, wrong reason. Fabric does not merge notebooks. You choose how to resolve.

C10-L1 exercise 3

Hints. (1) Connecting changes the workspace itself. (2) Lesson C1-L3. (3) Admin.

Answer. a) Admin Because: r1) Connect and disconnect are Admin operations. Members and Contributors commit and update.

Admins connect. Members and Contributors commit and update, within their Git permissions.

Common wrong answers.

- **b:** Members commit and update. Connecting is for Admins (lesson C1-L3).
- **c:** Contributors commit and update. Connecting is for Admins.
- **d:** Viewers see no Git information.
- **a|r2:** Right answer, wrong reason. Git permissions matter for everyone. The workspace role decides who connects.

C10-L1 exercise 4

Hints. (1) Collect the set of statuses present. (2) Validate against the constant first. (3) Check in priority order: conflict, uncommitted, update_required.

Answer.

```
1 ALLOWED_STATUSES = frozenset({"synced", "uncommitted", "update_required",  
    "conflict"})  
2  
3  
4 def next_step(statuses: dict[str, str]) -> str:  
5     """Next source control action for a workspace.  
6  
7     Args:  
8         statuses: Item name to source control status.  
9  
10    Returns:  
11        "resolve", "commit", "update" or "none".  
12  
13    Raises:  
14        ValueError: If a status is not recognised.  
15    """  
16    present = set(statuses.values())  
17    unknown = present - ALLOWED_STATUSES  
18    if unknown:  
19        raise ValueError(f"unknown statuses: {sorted(unknown)}")  
20    if "conflict" in present:  
21        return "resolve"  
22    if "uncommitted" in present:  
23        return "commit"
```

```
24     if "update_required" in present:
25         return "update"
26     return "none"
```

Priority order encodes the safe sequence: never let an update overwrite work you have not committed.

C10-L1 exercise 5

Hints. (1) Where does work in progress live? (2) How do changes reach main? (3) What does Git hold for a lakehouse?

Answer.

- One workspace and branch per engineer isolates work in progress.
- Pull requests put a review between each change and main.
- Git holds lakehouse metadata only, so data comes back through pipelines.

Isolated workspaces, reviewed merges, and pipelines that rebuild data.

Common wrong answers.

- workspaces = `shared`: A shared workspace is a shared runtime. One person's update overwrites another's uncommitted work.
- merge = `direct`: Direct commits skip review.
- data = `nothing`: Tables and files are not tracked in Git.

C10-L2 Deployment pipelines, dev to test to production

C10-L2 exercise 1

Hints. (1) What does a deployment copy? (2) Items and metadata. (3) Not data.

Answer. b) Refresh the model in Production, because data is not copied between stages. Because: r1) Microsoft Learn: a semantic model refresh is required because data isn't copied from one stage to another.

Deploy, then refresh or reload in the target stage.

Common wrong answers.

- `a`: Data is not copied between stages.
- `c`: Deployment does not touch permissions. On later deployments, IDs, URLs and permissions stay the same.
- `d`: Deployment pipelines and Git connections are separate.
- `b` | `r2`: Right answer, wrong reason. Deployment does not clear data. It never copied any.

C10-L2 exercise 2

Hints. (1) Which stage needs the different value? (2) The target. (3) Default lakehouse rule.

Answer. a) In the Production stage, on the notebook, as a default lakehouse rule Because: r1) Rules are defined in the target stage on the item, and apply from the next deployment into that stage.

The target stage owns its settings.

Common wrong answers.

- b: Rules belong to the target stage.
- c: Branching on workspace names hard-codes environments into code. Use a deployment rule.
- d: Deployment rules are pipeline configuration, set per stage.
- a|r2: Right answer, wrong reason. Rules are per stage and per item, not global.

C10-L2 exercise 3

Hints. (1) Rules must exist before deployment. (2) Compare before deploy. (3) Data after deploy, checks last.

Answer.

1. Check the deployment rules in the Production stage
2. Compare Test with Production and review differences
3. Deploy Test to Production
4. Refresh semantic models and run load pipelines in Production
5. Check row counts and a known report figure in Production

Rules, compare, deploy, reload, verify.

Common wrong answers.

- Rules apply at deployment. Check them before you deploy.
- Verify after the data is loaded, or you check an empty model.

C10-L2 exercise 4

Hints. (1) `rules.get(stage, {})` handles stages without rules. (2) Compare rule keys with source keys. (3) `{**source, **stage_rules}` builds a new dict.

Answer.

```
1 def deployed_settings(  
2     source: dict[str, str], rules: dict[str, dict[str, str]], stage: str  
3 ) -> dict[str, str]:
```

```

4     """Item settings in a target stage after its deployment rules apply.
5
6     Args:
7         source: Settings in the source stage.
8         rules: Stage name to the settings its rules override.
9         stage: Target stage.
10
11     Returns:
12         A new dict of settings for the target stage.
13
14     Raises:
15         ValueError: If a rule names a setting the item does not have.
16     """
17     stage_rules = rules.get(stage, {})
18     unknown = set(stage_rules) - set(source)
19     if unknown:
20         raise ValueError(f"rules for unknown settings: {sorted(unknown)}")
21     return {**source, **stage_rules}

```

Settings per stage, validated, with the source left intact: the same contract deployment rules give you.

C10-L2 exercise 5

Hints. (1) How do changes reach Development? (2) Which tool keeps IDs and applies rules? (3) What does a deployment not copy?

Answer.

- Git and pull requests control what reaches Development.
- Deployment pipelines promote with per-stage rules and a compare step.
- Deployments copy items, not data, so reload and check.

Git for development, deployment pipelines for promotion, and a reload and check after each release.

Common wrong answers.

- dev = `edits`: Direct edits skip review and conflict with Git updates.
- promote = `copy`: Manual copies lose rules, compare and autobinding, and produce new IDs and URLs.
- post = `done`: Data is not copied between stages.

C10-L3 Capacity, smoothing and throttling

C10-L3 exercise 1

Hints. (1) Background or interactive? (2) Background jobs are long and heavy. (3) 24 hours.

Answer. a) 24 hours Because: r1) Microsoft Learn: background operations are smoothed over a 24-hour period.

Background work is spread over a day, so one heavy job raises every timepoint a little.

Common wrong answers.

- b: Five minutes is the minimum for interactive operations.
- c: Smoothing spreads cost beyond the job's run time.
- d: Background operations are smoothed, over 24 hours.
- a|r2: Right answer, wrong reason. The smoothing window is fixed. It does not depend on job length.

C10-L3 exercise 2

Hints. (1) Which band is 90 minutes in? (2) 60 minutes to 24 hours. (3) Interactive rejection.

Answer. a) The report is rejected. The pipeline still runs Because: r1) Between 60 minutes and 24 hours, interactive jobs are rejected. Background rejection starts only above 24 hours.

Users feel throttling first. Background jobs, often the cause, keep running.

Common wrong answers.

- b: The 20-second delay applies between 10 and 60 minutes.
- c: Background jobs are rejected only above 24 hours of future use.
- d: Above 60 minutes, interactive jobs are rejected.
- a|r2: Right answer, wrong reason. Background jobs have a later threshold, not priority.

C10-L3 exercise 3

Hints. (1) CUs times hours in a day. (2) Divide the job by the daily total. (3) 96 CU hours, about 2.1%.

Answer.

- | |
|--|
| 1 F4 capacity per day: $4 \text{ CUs} \times 24 \text{ hours} = 96 \text{ CU hours}$ |
| 2 Share of every timepoint for 24 hours: $2 / 96 = \text{about } 2.1\%$ |

2 / 96 is about 2.1%, the same share as the F2 example with half the job.

Common wrong answers.

- g1 = 4: Multiply CUs by 24 hours for the daily CU hours.

C10-L3 exercise 4

Hints. (1) Three constants: 10, 60 and 1,440 minutes. (2) Use `<=` for each limit. (3) Validate first.

Answer.

```
1 OVERAGE_PROTECTION_MINUTES = 10
2 INTERACTIVE_DELAY_MINUTES = 60
3 INTERACTIVE_REJECTION_MINUTES = 24 * 60
4
5
6 def throttle_stage(future_minutes: float) -> str:
7     """Fabric throttling stage for the future capacity already consumed.
8
9     Args:
10         future_minutes: Minutes of future capacity used by smoothed
11         operations.
12
13     Returns:
14         "none", "interactive_delay", "interactive_rejection" or
15         "background_rejection".
16
17     Raises:
18         ValueError: If future_minutes is negative.
19     """
20     if future_minutes < 0:
21         raise ValueError("future_minutes must not be negative")
22     if future_minutes <= OVERAGE_PROTECTION_MINUTES:
23         return "none"
24     if future_minutes <= INTERACTIVE_DELAY_MINUTES:
25         return "interactive_delay"
26     if future_minutes <= INTERACTIVE_REJECTION_MINUTES:
27         return "interactive_rejection"
28     return "background_rejection"
```

Boundaries are where bugs hide. Test each limit exactly.

C10-L3 exercise 5

Hints. (1) Where does a 06:00 job's cost land? (2) Who shares the capacity? (3) Which setting protects interactive users?

Answer.

- Smoothing spreads a 06:00 job over the next 24 hours, including the morning peak. Less work, earlier, helps.
- Separate capacities stop experiments from consuming production CUs.
- Surge protection rejects background work before users are rejected. Resize only if the Compute page still shows a shortfall.

Reduce and move background work, isolate experiments, then protect users with surge protection.

Common wrong answers.

- refresh = keep: Smoothing does not remove cost. It spreads it across the morning peak.
- ds = same: All-day notebooks add to every timepoint on the production capacity.
- protect = resize: Fix the workload first. A larger SKU doubles the bill and hides the cause.

D1-L1 Read a Python traceback from the bottom up

D1-L1 exercise 1

Hints. (1) Read the header: 'most recent call last'. (2) Where is the most recent call? (3) At the bottom.

Answer. a) The last line: exception type and message Because: r1) Frames run oldest first. The exception and the frame that raised it are at the bottom. The bottom holds what went wrong. The frames above say where and how you got there.

Common wrong answers.

- b: The first frame is the oldest call, often your entry point or a scheduler. The error happened at the other end.
- c: Length says nothing about relevance. Start at the bottom.
- d: Library frames rarely hold the cause. Find the last frame in your own code.
- a|r2: Right answer, wrong reason. The order is chronological: 'most recent call last'.

D1-L1 exercise 2

Hints. (1) Read the last line first. (2) Which frame is closest to the error and inside /repo/pipelines/? (3) The frame for qa_statistical.py, line 19.

Answer. The bug is on line 4.

stats["mean"] is None when the DataFrame is empty, and round(None, 4) raises TypeError. The frame for qa_statistical.py line 19 is where you add a guard.

Common wrong answers.

- Line 2: That is the entry point that called the function. The error happened deeper.
- Line 7: That line says what went wrong: `round()` received `None`. The frame above it says where.

D1-L1 exercise 3

Hints. (1) Which sentence does `from err` print? (2) This one is the other sentence. (3) The handler failed.

Answer. a) The except block that handled the `KeyError` has its own bug
Because: r1) The second exception was raised while handling the first, without an explicit chain. Both need a fix.

"During handling" means your error handler raised. "Direct cause" means someone chained on purpose.

Common wrong answers.

- b: `from err` prints 'The above exception was the direct cause of the following exception'.
- c: The `KeyError` started it. Fix the handler and the missing key.
- d: Python does not retry. It records that the handler failed.
- a|r2: Right answer, wrong reason. The two errors are unrelated types. The handler raised a new one.

D1-L1 exercise 4

Hints. (1) Split on 'Traceback (most recent call last):' and keep the last part. (2) Match frame lines with a regular expression. (3) Keep overwriting the result, so the last match wins.

Answer.

```
1 import re
2
3 FRAME_PATTERN = re.compile(r'^\s*File "(?P<path>[^\"]+)", line (?
    P<line>\d+)')
4
5
6 def first_own_frame(traceback_text: str, project_dir: str) -> tuple[str,
    int] | None:
7     """Find the frame to start from when reading a traceback bottom up.
8
9     Args:
10         traceback_text: Full traceback text, possibly chained.
11         project_dir: Path prefix of your own code.
12
```

```

13     Returns:
14         (path, line) of the last frame inside project_dir, or None.
15     """
16     last_block = traceback_text.split("Traceback (most recent call
last):")[-1]
17     found: tuple[str, int] | None = None
18     for line in last_block.splitlines():
19         match = FRAME_PATTERN.match(line)
20         if match and match["path"].startswith(project_dir):
21             found = (match["path"], int(match["line"]))
22     return found

```

Reading bottom up is the same as keeping the last matching frame. Library frames never match your prefix.

D1-L1 exercise 5

Hints. (1) Catch `KeyError` as `err`, nothing wider. (2) `err.args[0]` is the missing key. (3) End the raise with `from err`.

Answer.

```

1 def row_count_bounds(contract: dict[str, object]) -> tuple[object,
object]:
2     """Row count bounds from a data contract.
3
4     Args:
5         contract: Parsed contract with contract_id, row_count_min and
row_count_max.
6
7     Returns:
8         (row_count_min, row_count_max).
9
10    Raises:
11        ValueError: If a bound is missing.
12    """
13    try:
14        return contract["row_count_min"], contract["row_count_max"]
15    except KeyError as err:
16        contract_id = contract.get("contract_id", "<unknown>")
17        raise ValueError(f"contract {contract_id} is missing key
{err.args[0]}") from err

```

A narrow except, a message with the facts, and an explicit chain. The next reader sees both tracebacks joined by "direct cause".

D1-L2 Spark errors, the condition and the real cause

D1-L2 exercise 1

Hints. (1) What does SQLSTATE class 42 mean? (2) Compare the name in the message with the suggestions. (3) A typo.

Answer. a) The column name in the code: `final_wieght` is a typo for `final_weight` Because: r1) Class 42 errors mean the code names something that does not exist. The message suggests the right name.

Unresolved names are code errors. Read the suggestion list, then fix the name.

Common wrong answers.

- b: Missing values are nulls, not an unresolved column. The name itself is wrong.
- c: The error is in the query plan. A restart runs the same plan.
- d: A permission error has a different condition. This one says the name cannot be resolved.
- a|r2: Right answer, wrong reason. The suggestions are the closest names, not always the right one. Check them.

D1-L2 exercise 2

Hints. (1) Which sum is being called? (2) Look at the module's imports. (3) Import functions as F.

Answer. a) `sum` is PySpark's sum, imported by name, which shadows Python's built-in Because: r1) `from pyspark.sql.functions import sum` replaces the built-in in that module, so the generator goes to Spark's sum. Import `functions as F`. Python's `sum`, `min`, `max`, `round`, `abs` and `filter` then keep their meaning.

Common wrong answers.

- b: `self.checks` is a list of `CheckResult` objects. The function receiving it is the problem.
- c: Generators work anywhere. Look at the imports.
- d: No column is involved. The code never meant to call Spark.
- a|r2: Right answer, wrong reason. PySpark does not disable anything. The import replaced the name.

D1-L2 exercise 3

Hints. (1) Actions trigger the work. Where was the failing expression written? (2) The message names the origin: cast. (3) try_cast returns null instead of failing.

Answer. The bug is on line 3.

Fixed version:

```
1 from pyspark.sql import functions as F
2
3 silver = bronze.withColumn("age", F.col("age").try_cast("int"))
4 silver = silver.withColumn("wave_id", F.lit(wave_id))
5 silver.write.mode("append").saveAsTable("research_silver.nat_respondents")
```

The write triggers the plan, so the error surfaces there. The origin line points at the cast. Use try_cast, then count the nulls it creates.

Common wrong answers.

- Line 5: The error appears at the write because that is the action. The origin line says cast, on line 3.
- Line 4: Adding a literal cannot fail on 'n/a'.

D1-L2 exercise 4

Hints. (1) Match the condition with `^\s*\[[([A-Z0-9_.]+)\]\]`. (2) Search for `SQLSTATE:\s*([0-9A-Z]{5})`. (3) The class is the first two characters of the SQLSTATE.

Answer.

```
1 import re
2
3 CONDITION_PATTERN = re.compile(r"^\s*\[[([A-Z0-9_.]+)\]\]")
4 SQLSTATE_PATTERN = re.compile(r"SQLSTATE:\s*([0-9A-Z]{5})")
5 CLASS_LABELS = {"42": "code", "22": "data"}
6
7
8 def triage(message: str) -> tuple[str | None, str]:
9     """Error condition and a first-pass category for a Spark error
10     message.
11
12     Args:
13         message: The error message text.
14
15     Returns:
16         (condition or None, "code", "data" or "unknown").
17     """
18     condition_match = CONDITION_PATTERN.match(message)
19     condition = condition_match.group(1) if condition_match else None
20     state_match = SQLSTATE_PATTERN.search(message)
```

```

20     if state_match is None:
21         return condition, "unknown"
22     return condition, CLASS_LABELS.get(state_match.group(1)[:2],
    "unknown")

```

A first-pass triage sends class 42 to the code and class 22 to the data. Parse the structured parts, never loose substrings.

D1-L2 exercise 5

Hints. (1) Which part of the message is structured? (2) What did ANSI mode expose? (3) What stays stable across releases?

Answer.

- The condition and SQLSTATE give the category in seconds. The origin line gives the place.
- Class 22 means real bad data. Switching ANSI off hides it again.
- Conditions are stable identifiers. Message wording changes between releases.

Condition, SQLSTATE class, origin line. Then fix code or investigate data, depending on the class.

Common wrong answers.

- first = java: The JVM stack is Spark's internals. The condition and origin line point at your code.
- class22 = ansi_off: That returns to silent nulls. The bad values are still there, uncounted.
- catch = text: Message text changes between versions. Conditions do not.

D1-L3 Py4J and Spark Connect errors

D1-L3 exercise 1

Hints. (1) Wrapped exceptions nest. (2) Look for 'Caused by:'. (3) The last one.

Answer. a) After the last 'Caused by:', in the Python traceback it contains Because: r1) The JVM wraps the Python worker's exception. The innermost cause holds the Python traceback.

Search the message for the last "Caused by:" and read its Python traceback bottom up.

Common wrong answers.

- b: Java frames show Spark's internals, not your code.

- **c**: That line names the JVM method Python called, such as `collectAndServe`.
- **d**: The Python traceback is inside the message, after 'Caused by: `org.apache.spark.api.python.PythonException`'.
- **a|r2**: Right answer, wrong reason. The cause comes last, nested inside the outer exception.

D1-L3 exercise 2

Hints. (1) When does Spark Connect resolve column names? (2) At execution, or when you read the schema. (3) No action runs here.

Answer.

```
1 select ok
```

Spark Connect builds the plan without resolving names, so the `select` succeeds. `out.count()` or `out.columns` would raise.

Common wrong answers.

- **caught**: That is Spark Classic, which analyses eagerly. Spark Connect defers analysis to the action.

D1-L3 exercise 3

Hints. (1) Which package is shared by both modes? (2) `pyspark.errors`. (3) `AnalysisException`.

Answer. a) `except pyspark.errors.AnalysisException` Because: r1) Both modes raise a subclass of `pyspark.errors.AnalysisException`.

Catch the base classes in `pyspark.errors`. Your code then works in both modes.

Common wrong answers.

- **b**: `Py4J` exists only in Spark Classic, and PySpark converts most errors before you see them.
- **c**: That class exists only under Spark Connect.
- **d**: Column errors are not Python `KeyErrors`.
- **a|r2**: Right answer, wrong reason. It is a `pyspark.errors` class, with Classic and Connect subclasses.

D1-L3 exercise 4

Hints. (1) `message.rsplit(MARKER, 1)[1]` gives the text after the last marker. (2) Skip the rest of the marker line, then indented lines. (3) Java frames start with a tab, which is whitespace.

Answer.

```
1 MARKER = "Caused by: org.apache.spark.api.python.PythonException:"
2
3
4 def python_cause(message: str) -> str | None:
5     """Python exception line inside the last PythonException of a Py4J
6     error.
7
8     Args:
9         message: Full error text.
10
11     Returns:
12         The exception line, such as "ZeroDivisionError: division by zero",
13         or None.
14     """
15     if MARKER not in message:
16         return None
17     section = message.rsplit(MARKER, 1)[1]
18     for line in section.splitlines()[1:]:
19         if not line.strip() or line[0].isspace():
20             continue
21         if line.startswith(("Traceback", "File")):
22             continue
23         return line
24     return None
```

The innermost cause is last. `rsplit` with a limit of 1 finds it directly.

D1-L3 exercise 5

Hints. (1) When does Connect resolve names? (2) Which package exists in both modes? (3) What does serverless support?

Answer.

- Spark Connect defers analysis, so a try around a select catches nothing. An explicit schema check works in both modes.
- `pyspark.errors` classes exist in both modes. `Py4J` exists only in Classic.
- Serverless supports only Spark Connect APIs, not RDDs.

Explicit schema checks, `pyspark.errors` classes, and the `DataFrame` API only.

Common wrong answers.

- `optional = try_select`: Under Spark Connect the select succeeds and the error moves to the action.
- `except_class = py4j`: Serverless has no `Py4J`.
- `rdd = keep`: Spark RDD APIs are not supported on serverless compute.

D2-L1 breakpoint() and pdb, stop and look

D2-L1 exercise 1

Hints. (1) The variable controls what breakpoint() calls. (2) 0 means disabled. (3) The job continues.

Answer. a) Nothing: the call does nothing and the job continues Because: r1) PYTHONBREAKPOINT=0 makes breakpoint() a no-op.

Set PYTHONBREAKPOINT=0 in job environments as a safety net, and still remove breakpoints before committing.

Common wrong answers.

- b: That is what happens without the variable. With 0, the call is ignored.
- c: No error. The call returns immediately.
- d: pdb never starts.
- a|r2: Right answer, wrong reason. Without the variable, a scheduled job would stop at the breakpoint.

D2-L1 exercise 2

Hints. (1) Print is one letter. (2) Up is one letter. (3) p and u.

Answer.

```
1 (Pdb) p reach, base      # print the inputs
2 (Pdb) u                  # move to the caller's frame
3 (Pdb) p rows
```

p shows values. u moves to the caller, where the bad value came from.

Common wrong answers.

- g2 = d: d moves down, towards the newest frame. You are already there. Use u.

D2-L1 exercise 3

Hints. (1) What makes each run fast? (2) Find the first wrong value. (3) Clean up last.

Answer.

1. Reproduce with the smallest input that gives the wrong number
2. Stop in the function with breakpoint() or post-mortem pdb
3. Print the inputs, walk up with u until a value is first wrong
4. Fix that frame and add a test with the small input
5. Remove the breakpoint and run the tests

Small reproduction, stop, walk up to the first wrong value, fix with a test, clean up.

Common wrong answers.

- Reproduce first, with small data, so each run is fast and you know the right answer.
- Fix and add the test while you know the cause, then clean up.

D2-L1 exercise 4

Hints. (1) `err.__traceback__` is the outermost traceback entry. (2) Follow `tb_next` until it is `None`. (3) `tb.tb_frame.f_locals` holds that frame's variables.

Answer.

```
1 from collections.abc import Callable
2
3
4 def crash_locals(fn: Callable[..., object], *args: object) -> dict[str,
    object] | None:
5     """Local variables of the frame that raised, as pdb post-mortem shows
        them.
6
7     Args:
8         fn: Function to call.
9         *args: Arguments for fn.
10
11     Returns:
12         A copy of the innermost frame's locals, or None if fn did not
        raise.
13     """
14     try:
15         fn(*args)
16     except Exception as err:
17         tb = err.__traceback__
18         while tb is not None and tb.tb_next is not None:
19             tb = tb.tb_next
20         return dict(tb.tb_frame.f_locals) if tb is not None else None
21     return None
```

A traceback is a linked list from your call down to the raise. `pdb post-mortem` opens at its last entry.

D2-L1 exercise 5

Hints. (1) Which line would stop an unattended job? (2) A debugging call. (3) Line 4.

Answer. The bug is on line 4.

Fixed version:

```
1 def unweighted_base(rows: list[dict[str, object]]) -> int:
2     """Count in-scope respondents."""
3     in_scope = [r for r in rows if r["in_scope"] is True]
4     return len(in_scope)
```

A leftover `breakpoint()` blocks a job waiting for input that never comes. ruff rule T100 flags it.

Common wrong answers.

- Line 3: `is True` is correct for a Python bool flag in a dict. The problem is the next line.

D2-L2 Logging levels that help at 3 a.m.

D2-L2 exercise 1

Hints. (1) What is the root logger's default level? (2) WARNING. (3) Only the warning appears, as bare text.

Answer.

```
1 weight sum low
```

With no configuration, records at WARNING and above go to stderr as the bare message. INFO is dropped.

Common wrong answers.

- `rows written weight sum low`: The default effective level is WARNING, so the INFO record is dropped.
- `WARNING:qa:weight sum low`: That is basicConfig's default format. Without configuration, Python's last-resort handler prints the message only.

D2-L2 exercise 2

Hints. (1) Does basicConfig reconfigure? (2) Only the first call counts. (3) `force=True` changes that.

Answer. a) WARNING: the second call does nothing Because: r1) basicConfig does nothing if the root logger already has handlers, unless `force=True`.

Pass `force=True`, or set levels on named loggers with `setLevel`.

Common wrong answers.

- `b`: Only with `force=True`, or with `logger.setLevel`.
- `c`: It returns silently, which is why this surprises people.
- `d`: No second handler is added.

- a|r2: Right answer, wrong reason. setLevel or force=True lower the level fine.

D2-L2 exercise 3

Hints. (1) Which logger method adds the traceback? (2) A bare statement re-raises the current exception. (3) exception and raise.

Answer.

```
1 try:
2     validate(df)
3 except ContractError:
4     logger.exception("contract failed for wave %s", wave_id)
5     raise
```

logger.exception logs at ERROR with the traceback. A bare raise re-raises the same exception.

Common wrong answers.

- g1 = error: error logs the message without the traceback. exception adds it.
- g2 = pass: Swallowing the error lets the pipeline continue with invalid data.

D2-L2 exercise 4

Hints. (1) A dict from status to level keeps the branches short. (2) logger.log(level, msg, *args) logs at any level. (3) Handle the CRITICAL case first: log, then raise.

Answer.

```
1 import logging
2
3 STATUS_LEVELS = {"PASS": logging.INFO, "WARN": logging.WARNING, "FAIL":
4     logging.ERROR}
5
6 def log_check(
7     logger: logging.Logger, check: str, status: str, severity: str,
8     wave_id: str
9 ) -> None:
10     """Log one check result at a level matching its outcome.
11
12     Args:
13         logger: Module logger.
14         check: Check name.
15         status: "PASS", "WARN" or "FAIL".
16         severity: Severity of the check, such as "CRITICAL".
17         wave_id: Wave being checked.
```

```

17
18     Raises:
19         ValueError: If status is unknown.
20         RuntimeError: If a CRITICAL check failed.
21     """
22     if status not in STATUS_LEVELS:
23         raise ValueError(f"unknown status: {status}")
24     if status == "FAIL" and severity == "CRITICAL":
25         logger.critical("wave %s check %s failed (CRITICAL)", wave_id,
26                        check)
27         raise RuntimeError(f"CRITICAL check {check} failed for wave
28                        {wave_id}")
29     logger.log(STATUS_LEVELS[status], "wave %s check %s %s", wave_id,
30              check, status)

```

Levels describe. Exceptions stop. A CRITICAL failure needs both.

D2-L2 exercise 5

Hints. (1) What do you read at 3 a.m.? (2) Which identifiers find the data? (3) Where does history for dashboards live?

Answer.

- INFO keeps logs readable. A parameter enables DEBUG without a code change.
- Identifiers and counts let you find the step that dropped rows.
- Structured history belongs in a table. Text logs are for people.

INFO by default, identifiers everywhere, and a structured run log alongside the text.

Common wrong answers.

- level = debug: DEBUG every night buries the lines you need under detail you rarely read.
- content = done: 'Done' cannot tell you which wave or how many rows.
- record = parse: Log text changes. A run log table with a fixed schema is queryable.

D2-L3 explain(), the Spark UI and the query profile

D2-L3 exercise 1

Hints. (1) What does Cartesian mean? (2) Every pair of rows. (3) Check the join condition.

Answer. a) A join without an equality condition, so every row meets every row. Because: r1) A Cartesian product pairs every row with every row. It comes from a cross join or a join condition Spark cannot use as an equi-join. Find the join that produced it and give it an equality condition on the keys.

Common wrong answers.

- b: Broadcast joins show as BroadcastHashJoin.
- c: Python UDFs show as BatchEvalPython or ArrowEvalPython.
- d: Pruning shows as PartitionFilters on the scan.
- a|r2: Right answer, wrong reason. Small tables get BroadcastHashJoin when there is an equality condition.

D2-L3 exercise 2

Hints. (1) Which operator moves data across the cluster? (2) Its name means swap. (3) Exchange.

Answer. The bug is on line 3.

`Exchange hashpartitioning(wave_id, 200)` shuffles partial sums so each wave's rows meet on one executor for the final sum.

Common wrong answers.

- Line 6: The scan reads files. The shuffle moves data between executors, which is an Exchange.
- Line 2: HashAggregate aggregates. The final aggregate needs the data shuffled first, by the Exchange below it.

D2-L3 exercise 3

Hints. (1) Is the Spark UI available on serverless? (2) No. (3) Its replacement.

Answer. a) The query profile. Because: r1) The Spark UI is not available on serverless. Databricks points you to the query profile.

On serverless, use the query profile. On classic compute, the Spark UI's SQL tab shows the same kind of detail.

Common wrong answers.

- b: Correct for classic compute. Serverless has no Spark UI.
- c: explain shows the plan before running, not what happened.
- d: Logs do not show per-operator metrics.
- a|r2: Right answer, wrong reason. Classic compute has the Spark UI too.

D2-L3 exercise 4

Hints. (1) Keep the operator names in constants. (2) A negative lookbehind, `(?<!\Broadcast)`, excludes `BroadcastExchange`. (3) Append warnings in the stated order.

Answer.

```
1 import re
2
3 CROSS_JOIN_OPERATORS = ("CartesianProduct", "BroadcastNestedLoopJoin")
4 PYTHON_UDF_OPERATORS = ("BatchEvalPython", "ArrowEvalPython")
5 SHUFFLE_PATTERN = re.compile(r"(?<!\Broadcast)\bExchange\b")
6
7
8 def plan_warnings(plan: str, max_exchanges: int = 2) -> list[str]:
9     """Warnings for common problems in a Spark physical plan.
10
11     Args:
12         plan: Text of df.explain().
13         max_exchanges: Number of shuffles allowed before warning.
14
15     Returns:
16         Warnings in a fixed order, each at most once.
17     """
18     warnings: list[str] = []
19     if any(op in plan for op in CROSS_JOIN_OPERATORS):
20         warnings.append("cross join")
21     if any(op in plan for op in PYTHON_UDF_OPERATORS):
22         warnings.append("python udf")
23     if len(SHUFFLE_PATTERN.findall(plan)) > max_exchanges:
24         warnings.append("many shuffles")
25     return warnings
```

A plan linter in a test catches a cross join the day it appears, not the day the dimension grows.

D2-L3 exercise 5

Hints. (1) What is free to check? (2) Where are rows and task times? (3) What does one slow task mean?

Answer.

- The plan shows join types and shuffles for free, before you spend compute.
- Rows per operator show where rows multiply. Task times show skew.
- One very slow task among fast ones is the signature of a skewed key.

Plan first, then runtime metrics, then the data behind any outlier.

Common wrong answers.

- `first = size`: A bigger cluster makes a cross join or a skewed key slightly less slow, at double the cost.
- `runtime = rerun`: A rerun tells you it is still slow, not why.
- `skew = slow_node`: Skew follows the data, not the machine. It returns after a restart.

D3-L1 Join row explosion and nulls in filters

D3-L1 exercise 1

Hints. (1) How many rows in regions match code N? (2) How many respondents have code N? (3) $2 \times 2 + 1$.

Answer.

1 5

Each N respondent matches two lookup rows. 4 rows for N, 1 for S. The left join multiplied the grain.

Common wrong answers.

- `3`: A left join keeps every left row, but each one appears once per matching right row. Code N matches twice.
- `4`: Two respondents have code N, and each matches two rows: $2 \times 2 + 1 = 5$.

D3-L1 exercise 2

Hints. (1) What is null `!= 'N'` in SQL logic? (2) Null, which is not true. (3) $1,000 - 300 - 50$.

Answer. a) 650 Because: `r1) null != 'N'` evaluates to null, and a filter keeps only rows where the condition is true.

Three-valued logic: comparisons with null give null. Add `| F.col("region_code").isNull()` to keep them.

Common wrong answers.

- `b`: That is pandas behaviour. In Spark the 50 null rows are dropped too.
- `c`: That keeps N and null. The filter removes N.
- `d`: The filter removes at least the 300 N rows.
- `a|r2`: Right answer, wrong reason. Spark keeps nulls in general. Only the comparison yields null here.

D3-L1 exercise 3

Hints. (1) Expand NOT IN into ANDs. (2) What is $x \neq \text{NULL}$? (3) Use NOT EXISTS or a left_anti join.

Answer. a) $x \text{ NOT IN } (\dots, \text{NULL})$ is never true, because $x \neq \text{NULL}$ is null
Because: r1) NOT IN expands to $x \neq 'N'$ AND $x \neq \text{NULL}$. The second term is null, so the whole condition is never true.

Use NOT EXISTS or a left_anti join. Both treat a null in the exclusion list as "no match".

Common wrong answers.

- b: It runs. It just returns no rows when the list contains a null.
- c: Two rows is fine. The null is the problem.
- d: Even an S respondent fails the null comparison.
- a|r2: Right answer, wrong reason. The subquery keeps the null, and that is what breaks NOT IN.

D3-L1 exercise 4

Hints. (1) `lookup.dropna(subset=[key])` removes null keys. (2) `value_counts()` then `counts[counts > 1]` finds duplicates. (3) Check `len(joined) == len(facts)` before returning.

Answer.

```
1 import pandas as pd
2
3
4 def safe_left_join(facts: pd.DataFrame, lookup: pd.DataFrame, key: str) -
    > pd.DataFrame:
5     """Left join that keeps the grain of facts.
6
7     Args:
8         facts: Rows to keep, one per output row.
9         lookup: Attributes to attach, at most one row per key.
10        key: Join column.
11
12     Returns:
13         The joined frame with len(facts) rows.
14
15     Raises:
16         ValueError: If lookup holds a key more than once, or the grain
17         changes.
18     """
19     clean = lookup.dropna(subset=[key])
20     counts = clean[key].value_counts()
21     duplicated = sorted(counts[counts > 1].index.tolist())
22     if duplicated:
```

```

22     raise ValueError(f"duplicate keys in lookup: {duplicated}")
23     joined = facts.merge(clean, on=key, how="left")
24     if len(joined) != len(facts):
25         raise ValueError(f"join changed rows: {len(facts)} ->
        {len(joined)}")
26     return joined

```

Report duplicates, never drop them silently: which "North" is right is a business decision. Null keys leave the lookup before the merge.

D3-L1 exercise 5

Hints. (1) What is null != 'test' in Spark? (2) Null, so the row is dropped. (3) eqNullSafe treats null as a comparable value.

Answer. The bug is on line 3.

Fixed version:

```

1 from pyspark.sql import functions as F
2
3 live = respondents.filter(~F.col("panel").eqNullSafe("test"))
4 live.write.mode("overwrite").saveAsTable("research_silver.nat_live")

```

`~F.col("panel").eqNullSafe("test")` is true for nulls, so respondents with no panel stay.

Common wrong answers.

- Line 4: The write stores what the filter kept. The rows were lost one line earlier.

D3-L2 Type coercion and timezone drift

D3-L2 exercise 1

Hints. (1) Compare 'W10' and 'W2' character by character. (2) '1' < '2'. (3) W10 first.

Answer.

```
1 ['W10', 'W2', 'W9']
```

Text sorts by character. Pad ids (W02, W09, W10) or sort by the number: `sorted(waves, key=lambda w: int(w[1:]))`.

Common wrong answers.

- `['W2', 'W9', 'W10']`: That is numeric order. Strings compare character by character, and '1' comes before '2'.
- `['W9', 'W2', 'W10']`: Strings sort ascending by character. 'W1...' comes first.

D3-L2 exercise 2

Hints. (1) What is 23:30 UTC in London in summer? (2) BST is UTC+1. (3) 00:30 the next day.

Answer. a) 2026-07-01 Because: r1) `TIMESTAMP` is an instant. `to_date` uses the session time zone, where 23:30 UTC is 00:30 on 1 July in BST.

Date boundaries depend on the zone. Set the session time zone deliberately, or convert with a named zone in the rule itself.

Common wrong answers.

- `b`: That is the date in UTC. The session time zone here is London, one hour ahead in summer.
- `c`: Spark uses the session time zone setting, not the OS clock.
- `d`: The conversion succeeds. It just gives a different day.
- `a|r2`: Right answer, wrong reason. Spark stores the instant. It applies the zone when converting or displaying.

D3-L2 exercise 3

Hints. (1) Which method takes a time zone argument? (2) Pass the UTC zone object. (3) `now` and `timezone.utc`.

Answer.

```
1 from datetime import datetime, timezone
2
3 run_timestamp = datetime.now(timezone.utc)
```

`datetime.now(timezone.utc)` returns an aware datetime in UTC.

Common wrong answers.

- `g1 = utcnow`: `utcnow()` is deprecated from Python 3.12 and returns a naive datetime.

D3-L2 exercise 4

Hints. (1) Pass `dtype=` to `read_csv` for the id columns. (2) Read the weight as a string first, then `pd.to_numeric(..., errors="coerce")`. (3) Failed = present before, missing after.

Answer.

```
1 import io
2
3 import pandas as pd
4
5 ID_COLUMNS = {"respondent_id": "string", "postcode_area": "string"}
6
```

```

7
8 def load_delivery(csv_text: str) -> pd.DataFrame:
9     """Read a delivery CSV with identifiers as strings and weights as
    floats.
10
11     Args:
12         csv_text: CSV text with respondent_id, postcode_area and
    final_weight.
13
14     Returns:
15         The delivery as a DataFrame.
16
17     Raises:
18         ValueError: If any non-empty weight cannot be parsed.
19     """
20     df = pd.read_csv(io.StringIO(csv_text), dtype={**ID_COLUMNS,
    "final_weight": "string"})
21     raw = df["final_weight"]
22     parsed = pd.to_numeric(raw, errors="coerce")
23     failed = int((raw.notna() & parsed.isna()).sum())
24     if failed:
25         raise ValueError(f"{failed} weights could not be parsed")
26     return df.assign(final_weight=parsed.astype("float64"))

```

Read as text, convert explicitly, and count what the conversion lost. Silent coercion hides bad data.

D3-L2 exercise 5

Hints. (1) How does union match columns? (2) By position. (3) unionByName matches by name.

Answer. The bug is on line 3.

Fixed version:

```

1 from pyspark.sql import functions as F
2
3 combined = original.unionByName(redelivery)
4 combined = combined.withColumn("loaded_at", F.current_timestamp())

```

union pairs the first column with the first column, whatever their names.

unionByName pairs by name, and fails if a name is missing.

Common wrong answers.

- Line 4: Adding a timestamp cannot move ids into the wave column. The union already mixed them.

D3-L3 Mutable defaults and the truthiness of zero

D3-L3 exercise 1

Hints. (1) When is the default `[]` created? (2) Both calls use the default. Is it the same list? (3) The list is created once, when Python defines the function.

Answer.

```
1 ['null_weights', 'row_count']
```

The default list is created when Python defines the function. Both calls append to that one list, so the second call returns both flags.

Common wrong answers.

- `['row_count']`: That is what the author intended. The default list is created once, so the second call appends to the list the first call used.
- `['row_count', 'null_weights']`: `append` adds to the end. The first call's flag is already in the list.

D3-L3 exercise 2

Hints. (1) Evaluate `0.0` and `0.12` in your head. (2) `0.0` is falsy, like `None`. (3) The whole block, including `results.append`, is skipped.

Answer. c) Leaves it out of the results. Because: `r1` `if current_pct and prior_pct:` is `False` when `current_pct` is `0.0`, so the block never runs.

`0.0` and `0.12` evaluates to `0.0`, which is falsy. The block that compares and appends never runs, so the metric vanishes from the results and no warning appears.

Common wrong answers.

- a: That is what the check should do. Look at the `if` line before the comparison: `0.0` is falsy.
- b: Nothing is appended for this metric, so it cannot be marked PASS. The `if` skips it.
- d: The check subtracts, it does not divide. No error appears, which is the problem.
- c|r2: Right outcome, wrong reason. There is no 1% filter. The truthiness test treats `0.0` as missing.

D3-L3 exercise 3

Hints. (1) The first default is None, so the type must allow None. (2) The second default must be immutable. (3) `str | None` for the first gap and `None` for the second.

Answer.

```
1 def log_run(status: str, error_message: str | None = None) -> None: ...
2
3 def collect(flag: str, flags: list[str] | None = None) -> list[str]: ...
```

`error_message: str | None = None` states that the value may be missing.

`flags: list[str] | None = None` avoids a shared list. Create the list inside the function when it is `None`.

Common wrong answers.

- `g1 = str: str = None` claims the value is always a string. mypy in strict mode rejects it. Write `str | None`.
- `g2 = []`: A list default is created once and shared by every call. Use `None` and create the list inside the function.

D3-L3 exercise 4

Hints. (1) Which values make `if current_pct and prior_pct`: False? (2) 0.0 is one of them. (3) Replace the truthiness test with `is not None` on both values.

Answer. The bug is on line 3.

Fixed version:

```
1 def stability_status(current_pct: float | None, prior_pct: float | None,
2                       tolerance: float = 0.10) -> str:
3     """Return WARN, PASS, or SKIP when a share is missing."""
4     if current_pct is not None and prior_pct is not None:
5         change = abs(current_pct - prior_pct)
6         return "WARN" if change > tolerance else "PASS"
7     return "SKIP"
```

Line 3 treats 0.0 as missing. Use `if current_pct is not None and prior_pct is not None`: so that only missing shares skip the comparison.

Common wrong answers.

- Line 4: The subtraction is correct. The code never reaches it when a share is 0.0.
- Line 5: The comparison with the tolerance is correct. Look at what stops the code reaching it.

D3-L3 exercise 5

Hints. (1) Start with `flags: list[str] = []` inside the function. (2) Read the prior share with `prior.get(metric)`, which gives `None` for a missing metric. (3) Skip with `if current_pct is None or prior_pct is None: continue`, then compare with the tolerance.

Answer.

```
1 def compare_waves(
2     current: dict[str, float | None],
3     prior: dict[str, float | None],
4     tolerance: float = 0.10,
5 ) -> list[str]:
6     """Flag metrics whose share moved more than tolerance between waves.
7
8     A metric is skipped only when a share is missing. A share of 0.0 is
9     compared.
10    """
11    flags: list[str] = []
12    for metric, current_pct in current.items():
13        prior_pct = prior.get(metric)
14        if current_pct is None or prior_pct is None:
15            continue
16        if abs(current_pct - prior_pct) > tolerance:
17            flags.append(metric)
18    return flags
```

Create the list inside the function. Use `prior.get` so a new metric gives `None`. Skip only on `None`, then compare the absolute change with the tolerance.

D3-L3 exercise 6

Hints. (1) Default `results` to `None`. (2) Create a new list only when `results` is `None`. (3) `results or []` looks similar but replaces an empty list the caller passed. Use `is None`.

Answer.

```
1 def add_result(
2     name: str, passed: bool, results: list[dict[str, object]] | None =
3     None
4 ) -> list[dict[str, object]]:
5     """Append one check result and return the list.
6
7     Args:
8         name: Check name.
9         passed: Whether the check passed.
10        results: List to append to. A new list is created when omitted.
11
12    Returns:
13        The list with the new result.
14    """
```



```

14     results = [] if results is None else results
15     results.append({"check": name, "passed": passed})
16     return results

```

Default to `None`, create the list when `results` is `None`, then append and return. Avoid `results` or `[]`: an empty list the caller passed is falsy, so the function would append to a new list and the caller would never see the result.

D4-L1 Row counts across layers

D4-L1 exercise 1

Hints. (1) What is silver's grain? (2) One row per respondent per wave. (3) Count distinct keys in bronze.

Answer. a) 10,000 Because: r1) Silver's grain is one row per respondent per wave, so it holds one row per distinct key.

The grain turns a vague expectation into an exact rule: silver rows = distinct bronze keys.

Common wrong answers.

- **b:** That keeps the 400 duplicates. Silver deduplicates on the key.
- **c:** The grain fixes the number exactly. A vague rule catches nothing.
- **d:** Weights do not change row counts.
- **a|r2:** Right answer, wrong reason. The number follows the grain, not a fixed rate.

D4-L1 exercise 2

Hints. (1) Which metric adds rows? (2) Inserts. (3) $10,000 + 120$.

Answer. a) 10,120 Because: r1) Updates replace existing rows. Only inserts add rows.

Check after every MERGE: rows after = rows before + numTargetRowsInserted - rows deleted.

Common wrong answers.

- **b:** 380 source rows updated existing rows, so they add nothing.
- **c:** Updated rows were already counted in the 10,000.
- **d:** 120 new respondents were inserted.
- **a|r2:** Right answer, wrong reason. Updated rows replace existing ones and add nothing.

D4-L1 exercise 3

Hints. (1) Which function counts distinct values? (2) Then fetch the single result row. (3) `count_distinct` and `first`.

Answer.

```
1 row = df.agg(  
2     F.count("*").alias("rows"),  
3     F.count_distinct("respondent_id").alias("keys"),  
4 ).first()
```

One `agg` returns both numbers in one scan. `first()` brings the single row back to the driver.

Common wrong answers.

- `g1 = count`: `count` counts non-null values, duplicates included. You need distinct keys.

D4-L1 exercise 4

Hints. (1) Validate inputs first. (2) List the rules as (name, expected, actual) tuples. (3) Silver is compared with distinct bronze keys.

Answer.

```
1 def reconcile(  
2     delivered: int,  
3     bronze_rows: int,  
4     bronze_keys: int,  
5     silver_rows: int,  
6     gold_n: int,  
7     in_scope: int,  
8 ) -> list[dict[str, object]]:  
9     """Row-count rules across layers for one tracker and wave.  
10  
11     Args:  
12         delivered: Rows in the delivery manifest.  
13         bronze_rows: Bronze rows for the delivery.  
14         bronze_keys: Distinct respondent keys in bronze.  
15         silver_rows: Silver rows for the wave.  
16         gold_n: Sum of gold unweighted_n.  
17         in_scope: Silver rows in scope.  
18  
19     Returns:  
20         One structured record per rule.  
21  
22     Raises:  
23         ValueError: If inputs are impossible.  
24     """  
25     counts = (delivered, bronze_rows, bronze_keys, silver_rows, gold_n,  
26               in_scope)  
27     if min(counts) < 0:
```

```

27         raise ValueError("counts must not be negative")
28     if in_scope > silver_rows:
29         raise ValueError("in_scope cannot exceed silver_rows")
30     rules = [
31         ("bronze_vs_delivery", delivered, bronze_rows),
32         ("silver_vs_bronze_keys", bronze_keys, silver_rows),
33         ("gold_vs_silver_scope", in_scope, gold_n),
34     ]
35     return [
36         {"rule": name, "expected": expected, "actual": actual,
37          "status": "PASS" if expected == actual else "FAIL"}
38         for name, expected, actual in rules
39     ]

```

Structured records with expected and actual values go straight into the run log, and a dashboard can chart them.

D4-L1 exercise 5

Hints. (1) Compare the filters on both sides. (2) Which side has no tracker filter? (3) Line 5.

Answer. The bug is on line 5.

Fixed version:

```

1 bronze_waves = (
2     spark.table(bronze_table).filter(F.col("tracker_id") == tracker_id)
3     .select("wave_id").distinct()
4 )
5 silver_waves = spark.table(silver_table).filter(F.col("tracker_id") ==
6     tracker_id).select("wave_id").distinct()
7 new_waves = [r.wave_id for r in
8     bronze_waves.subtract(silver_waves).collect()]

```

Both sides of a comparison need the same keys. Silver waves for all trackers hide this tracker's missing waves.

Common wrong answers.

- Line 6: subtract is the right operation. Its right side holds every tracker's waves.
- Line 2: The bronze side is filtered correctly. The silver side is not.

D4-L2 Distribution drift between waves

D4-L2 exercise 1

Hints. (1) Compare the labels, not the counts. (2) North vanished, Northern appeared. (3) A recode.

Answer. a) A recode by the supplier Because: r1) A category that vanishes while a similar new one appears is a label change until proven otherwise. Check new and missing categories before comparing shares.

Common wrong answers.

- b: A population does not rename itself. Check the labels first.
- c: Weights change shares, not category names.
- d: Sampling variation moves shares. It does not create new labels.
- a|r2: Right answer, wrong reason. Most waves have no recode, which is why a check must catch the one that does.

D4-L2 exercise 2

Hints. (1) Both totals are 100. (2) $0.45 - 0.30$. (3) In points, times 100.

Answer.

```
1 15
```

A move from 30% to 45% is 15 percentage points, and a 50% relative increase. State which one a tolerance means.

Common wrong answers.

- 50: That is the relative change ($45/30 - 1$). The check uses percentage points: $45\% - 30\%$.
- 0: The shares differ by 0.15, which is 15 points.

D4-L2 exercise 3

Hints. (1) Natural logarithm. (2) Of actual over expected. (3) `log` and `e`.

Answer.

```
1 import math
2
3 e = max(expected_share, FLOOR)
4 a = max(actual_share, FLOOR)
5 term = (a - e) * math.log(a / e)
```

$(a - e) * \ln(a / e)$ is never negative, so every category adds to the index.

Common wrong answers.

- `g1 = log10`: PSI uses the natural logarithm, `math.log`.

D4-L2 exercise 4

Hints. (1) Totals first, and guard zero. (2) Loop over `expected.keys() | actual.keys()`. (3) Floor each share with `max(share, floor)`.

Answer.

```

1 import math
2
3
4 def psi(expected: dict[str, int], actual: dict[str, int], floor: float =
    0.0001) -> float:
5     """Population stability index between two categorical distributions.
6
7     Args:
8         expected: Category counts in the reference wave.
9         actual: Category counts in the current wave.
10        floor: Minimum share, so a zero share does not break the
        logarithm.
11
12    Returns:
13        PSI, 0 when the shares are identical.
14
15    Raises:
16        ValueError: If either wave has no observations.
17    """
18    e_total = sum(expected.values())
19    a_total = sum(actual.values())
20    if e_total == 0 or a_total == 0:
21        raise ValueError("a wave has no observations")
22    index = 0.0
23    for category in expected.keys() | actual.keys():
24        e = max(expected.get(category, 0) / e_total, floor)
25        a = max(actual.get(category, 0) / a_total, floor)
26        index += (a - e) * math.log(a / e)
27    return index

```

Shares, the union of categories and a floor. Without the union, the recode that matters most adds nothing to the index.

D4-L2 exercise 5

Hints. (1) What is total for an empty delivery? (2) Zero. (3) Check it before dividing.

Answer. The bug is on line 3.

Fixed version:

```

1 def distribution_status(value_count: int, total: int, min_pct: float,
    max_pct: float) -> str:
2
3     """PASS if the value's share is within bounds, WARN if not, NO_DATA if
        empty."""
4     if total == 0:
5         return "NO_DATA"
6     pct = value_count / total
7     return "PASS" if min_pct <= pct <= max_pct else "WARN"

```

Guard every division. An empty delivery is a status to report, not a crash.

Common wrong answers.

- Line 4: The bounds check is fine. The division above it fails first.

D4-L3 Weight totals and survey checks

D4-L3 exercise 1

Hints. (1) What does `when()` return when no otherwise is given? (2) Null. (3) Sum of nulls.

Answer. a) NULL Because: r1) when without otherwise gives null for non-matching rows, and the sum of only nulls is null.

Add `.otherwise(0)` inside, or `F.coalesce(..., F.lit(0))` outside.

Common wrong answers.

- b: That is what you want, and what `coalesce(..., 0)` gives you. The raw sum is null.
- c: No error: the null passes silently into the table.
- d: `groupBy` keeps the subgroup. Its sum is null.
- a|r2: Right answer, wrong reason. Zeros sum normally. Only an all-null input gives null.

D4-L3 exercise 2

Hints. (1) $\text{sum}(w^2) = 4$, $(\text{sum } w)^2 = 16$. (2) $4 \times 4 / 16$. (3) 1.0.

Answer.

```
1 1.0
```

Equal weights give `deff` 1 and `n_eff` = n. Any variation raises `deff` above 1.

Common wrong answers.

- 4.0: $4 \times 4 / 16 = 1$. Equal weights cost no precision.
- 0.0: `deff` is at least 1. Equal weights give exactly 1.

D4-L3 exercise 3

Hints. (1) Which function returns the first non-null argument? (2) The replacement value. (3) `coalesce` and `0`.

Answer.

```
1 F.coalesce(F.sum(F.when(F.col("reach_flag") == 1, 1)),
  F.lit(0)).alias("unweighted_reach")
```

`coalesce` returns its first non-null argument, so a null sum becomes 0.

Common wrong answers.

- `g1 = isnull`: `isnull` tests for null and returns a boolean. `coalesce` replaces it.

D4-L3 exercise 4

Hints. (1) Check emptiness and invalid weights first. (2) Test `w` is not `None` and `w > 0`, not truthiness. (3) `n_eff = total**2 / sum of squares`; `deff = n / n_eff`.

Answer.

```
1 def weight_report(weights: list[float | None], max_ratio: float = 20.0) -> dict[str, object]:
2     """Weight checks with Kish's effective sample size and design effect.
3
4     Args:
5         weights: One weight per respondent. None means missing.
6         max_ratio: Largest allowed max/min weight ratio before WARN.
7
8     Returns:
9         status ("PASS", "WARN" or "CRITICAL"), n, n_eff and deff.
10    """
11    n = len(weights)
12    valid = [w for w in weights if w is not None and w > 0]
13    if n == 0 or len(valid) != n:
14        return {"status": "CRITICAL", "n": n, "n_eff": None, "deff": None}
15    total = sum(valid)
16    n_eff = total**2 / sum(w * w for w in valid)
17    status = "WARN" if max(valid) / min(valid) > max_ratio else "PASS"
18    return {"status": status, "n": n, "n_eff": round(n_eff, 4), "deff": round(n / n_eff, 4)}
```

A missing or non-positive weight is a data error to report, never a row to drop quietly.

D4-L3 exercise 5

Hints. (1) Which result makes the estimate undefined? (2) Which result needs judgement? (3) Which standard deviation does Kish use?

Answer.

- A null weight makes every estimate for that respondent undefined: halt.
- A high `deff` is a signal for a person, not proof of an error: warn.
- Kish's identity uses the population standard deviation. Guard the empty case.

CRITICAL for undefined estimates, WARN for signals, and the right statistic for the formula.

Common wrong answers.

- `nulls = warn`: Dropping rows changes the sample silently. Halt and fix the weighting.
- `deff = critical`: Real populations shift. Halting on every efficiency change blocks legitimate waves.
- `sd = samp`: `F.stddev` is the sample version, null for one row, and does not match $1 + CV^2$.

D5-L1 Invented functions, parameters and syntax

D5-L1 exercise 1

Hints. (1) Is `delta()` a `DataFrameReader` method? (2) Check with `hasattr(spark.read, 'delta')`. (3) No.

Answer. a) `AttributeError: DataFrameReader has no attribute 'delta'`
Because: r1) The reader has `format('delta').load(path)` and `table(name)`, but no `delta()` method.

Plausible, and invented. `hasattr` or the PySpark API reference answers in seconds.

Common wrong answers.

- b: The method does not exist. Use `spark.read.format("delta").load(path)` or `spark.read.table(name)`.
- c: No read happens. The attribute lookup fails first.
- d: Python raises `AttributeError` at once.
- a|r2: Right answer, wrong reason. Reading by path works with `format('delta').load(path)`.

D5-L1 exercise 2

Hints. (1) What does `Column` do with an unknown attribute? (2) It assumes a struct field. (3) The result is another `Column`.

Answer.

```
1 Column
```

`F.col("region").isNotNull` means "the field `isNotNull` inside the struct `region`". Calling it raises `TypeError: 'Column' object is not callable`.

Common wrong answers.

- `method`: The real method is `isNotNull`, with capitals. `isNotNull` is treated as a struct field name.

- `AttributeError`: Column turns unknown attributes into field lookups, so no error appears here.

D5-L1 exercise 3

Hints. (1) Which statement comes from another platform? (2) Databricks binds masks with `ALTER TABLE`. (3) `ALTER COLUMN ... SET MASK`.

Answer. The bug is on line 3.

Fixed version:

```
1 CREATE FUNCTION research_governance.mask_postcode(p STRING)
2 RETURN CASE WHEN is_account_group_member('analysts') THEN p ELSE '***'
   END;
3 ALTER TABLE research_silver.nat_respondents ALTER COLUMN postcode SET
   MASK research_governance.mask_postcode;
```

Databricks column masks: a SQL UDF, then `ALTER TABLE ... ALTER COLUMN ... SET MASK`.

Common wrong answers.

- Line 1: A SQL UDF is exactly what Databricks column masks use.
- Line 2: `is_account_group_member` is the recommended check in Unity Catalog.

D5-L1 exercise 4

Hints. (1) `inspect.signature(func).parameters` maps names to Parameter objects. (2) `p.kind` is `inspect.Parameter.VAR_KEYWORD` marks `**kwargs`. (3) Keep the order of params.

Answer.

```
1 import inspect
2 from collections.abc import Callable
3
4
5 def unknown_params(func: Callable[..., object], params: list[str]) ->
   list[str] | None:
6     """Parameter names that a function does not accept.
7
8     Args:
9         func: Function to check.
10        params: Names suggested for it, for example by an assistant.
11
12    Returns:
13        Unknown names in order, or None if func takes **kwargs.
14    """
15    parameters = inspect.signature(func).parameters
16    if any(p.kind is inspect.Parameter.VAR_KEYWORD for p in
   parameters.values()):
```

```
17         return None
18     return [name for name in params if name not in parameters]
```

A signature check catches an invented parameter before the code runs. With `**kwargs`, only the docs can answer.

D5-L1 exercise 5

Hints. (1) What proves an API exists for your version? (2) Which reference covers Databricks SQL? (3) What does 'undocumented' guarantee?

Answer.

- A docs link for the right version turns "looks right" into "checked".
- SQL dialects share keywords and differ in statements. Only the platform reference counts.
- Undocumented behaviour can change or disappear without notice.

Docs for your version, the platform's own SQL reference, and no undocumented APIs.

Common wrong answers.

- `check = trust`: Assistants state invented APIs with the same confidence as real ones.
- `sql = generic`: `CREATE ROW ACCESS POLICY` appears in tutorials. It is Snowflake, not Databricks.
- `unverified = try`: Running is necessary, not sufficient. Undocumented behaviour has no guarantee.

D5-L2 Deprecated APIs and version mismatches

D5-L2 exercise 1

Hints. (1) Check the 17.3 LTS release notes. (2) File details live in a hidden metadata column. (3) `_metadata.file_path`.

Answer. a) The `_metadata.file_path` column Because: r1) Databricks release notes state `input_file_name` is no longer supported from 17.3 LTS. The file metadata column replaces it.

Select `F.col("_metadata.file_path")` when reading files.

Common wrong answers.

- `b`: It is no longer supported on 17.3 LTS and above.
- `c`: `os.path` works on strings, not on DataFrame columns.
- `d`: A UDF cannot see the source file. The metadata column can.

- `a|r2`: Right answer, wrong reason. It was not renamed. A different mechanism, the `_metadata` column, replaces it.

D5-L2 exercise 2

Hints. (1) What does pandas 3 do with chained assignment? (2) The intermediate object behaves as a copy. (3) The original keeps 9.0.

Answer.

```
1 [1.0, 9.0]
```

pandas 3 warns with `ChainedAssignmentError` and leaves `df` unchanged. Use `df.loc[1, "w"] = 5.0`.

Common wrong answers.

- `[1.0, 5.0]`: That was pandas 1.x behaviour. Under Copy-on-Write, `df['w']` behaves as a copy, so the original is unchanged.

D5-L2 exercise 3

Hints. (1) `DataFrame.append` is gone. What joins frames end to end? (2) `applymap` was renamed. (3) `concat` and `map`.

Answer.

```
1 combined = pd.concat([waves, new_wave], ignore_index=True)
2 labels = combined.map(str.strip)
```

`pd.concat` replaces `append`. `DataFrame.map` replaces `applymap`.

Common wrong answers.

- `g2 = applymap`: `applymap` was removed in pandas 3.0. `DataFrame.map` replaces it.

D5-L2 exercise 4

Hints. (1) Split on dots and read the leading digits of each part. (2) Compare tuples of ints. (3) Pad the shorter tuple with zeros.

Answer.

```
1 import re
2
3 SEGMENT = re.compile(r"^(\d+)")
4
5
6 def _parse(version: str) -> tuple[int, ...]:
7     """Numeric segments of a dotted version string."""
8     parts = []
9     for segment in version.split("."):
10         match = SEGMENT.match(segment)
11         if match is None:
```

```

12         raise ValueError(f"invalid version segment {segment!r} in
    {version!r}")
13         parts.append(int(match.group(1)))
14         return tuple(parts)
15
16
17 def version_at_least(installed: str, required: str) -> bool:
18     """True if installed is the same as or newer than required.
19
20     Args:
21         installed: Installed version, such as "3.0.2".
22         required: Minimum version, such as "3.0".
23
24     Returns:
25         Whether installed >= required, compared numerically.
26
27     Raises:
28         ValueError: If a version has a segment without digits.
29     """
30     a, b = _parse(installed), _parse(required)
31     width = max(len(a), len(b))
32     return a + (0,) * (width - len(a)) >= b + (0,) * (width - len(b))

```

Tuples of integers compare element by element, which is version order. In real projects, use `packaging.version.Version`.

D5-L2 exercise 5

Hints. (1) Two indexing steps in one assignment. (2) Chained assignment. (3) `out.loc[mask, 'final_weight'] = cap`.

Answer. The bug is on line 6.

Fixed version:

```

1 import pandas as pd
2
3 def cap_weights(df: pd.DataFrame, cap: float) -> pd.DataFrame:
4     """Return a copy with final_weight capped at cap."""
5     out = df.copy()
6     out.loc[out["final_weight"] > cap, "final_weight"] = cap
7     return out

```

Index once with `.loc[rows, column]`. Chained indexing sets values on a temporary copy.

Common wrong answers.

- Line 5: The copy is right: it keeps the caller's frame unchanged.

D5-L3 Five review checks for AI-written code

D5-L3 exercise 1

Hints. (1) Where does the expected value come from? (2) From the same formula. (3) Self-confirming.

Answer. a) It repeats the implementation, so a wrong formula passes
Because: r1) A test's expected value must come from outside the code: small data worked by hand.

Write `assert weighted_pct(25, 200) == 12.5` and `assert weighted_pct(1, 0) is None`: numbers you worked out yourself.

Common wrong answers.

- b: approx helps with floats, but the test still repeats the formula.
- c: Arithmetic is fine. Copying the implementation is the problem.
- d: It checks that the code equals itself.
- a|r2: Right answer, wrong reason. Tests call the function. The expected value must be independent.

D5-L3 exercise 2

Hints. (1) Which line would leak through Git history? (2) A literal credential. (3) Read it at run time from the scope.

Answer. The bug is on line 3.

Fixed version:

```
1 import requests
2
3 def fetch_wave(wave_id: str) -> dict[str, object]:
4     """Fetch wave metadata from the survey platform."""
5     token = dbutils.secrets.get("survey-api", "token")
6     headers = {"Authorization": f"Bearer {token}"}
7     return requests.get(f"https://survey.example/api/waves/{wave_id}",
8                         headers=headers, timeout=30).json()
```

Secrets come from a secret scope at run time. Databricks also redacts values read with `dbutils.secrets.get` in notebook output.

Common wrong answers.

- Line 7: The header format is right. The value it uses comes from a literal on line 3.
- Line 8: The request is fine, with a timeout even. The secret is the problem.

D5-L3 exercise 3

Hints. (1) What does ALL on a whole table remove? (2) Every filter on it. (3) Including the wave.

Answer. a) Weighted reach for all waves, all regions and all brands Because: r1) ALL(fact_reach) removes every filter on the fact table, including the wave slicer.

Name the dimension whose filter you mean to remove: ALL(dim_brand) keeps the wave and region filters.

Common wrong answers.

- b: That is what ALL(dim_brand) would give. ALL on the fact table removes the wave filter too.
- c: That is the numerator's context.
- d: ALL accepts a table. That is exactly the problem here.
- a|r2: Right answer, wrong reason. CALCULATE keeps slicer filters unless a modifier such as ALL removes them.

D5-L3 exercise 4

Hints. (1) Check lengths before zip: zip stops at the shorter list. (2) Filter pairs with `is not None`. (3) Guard the total before dividing.

Answer.

```
1 def weighted_mean(values: list[float | None], weights: list[float |  
    None]) -> float | None:  
2     """Weighted mean that skips missing pairs.  
3  
4     Args:  
5         values: Values, None for missing.  
6         weights: Weights, None for missing.  
7  
8     Returns:  
9         The weighted mean, or None if there is nothing to average.  
10  
11     Raises:  
12         ValueError: If lengths differ or a weight is negative.  
13     """  
14     if len(values) != len(weights):  
15         raise ValueError("values and weights differ in length")  
16     pairs = [(v, w) for v, w in zip(values, weights) if v is not None and  
    w is not None]  
17     if any(w < 0 for _, w in pairs):  
18         raise ValueError("weights must not be negative")  
19     total = sum(w for _, w in pairs)  
20     if total == 0:  
21         return None
```

```
22     return sum(v * w for v, w in pairs) / total
```

Five edge cases the one-liner missed: different lengths, None values, negative weights, empty input and a zero total.

D5-L3 exercise 5

Hints. (1) What does the code do, stripped of structure? (2) One comparison per weight. (3) A list comprehension with enumerate.

Answer.

```
1 def flag_high_weights(weights: list[float], cap: float) -> list[int]:
2     """Indexes of weights strictly above cap.
3
4     Args:
5         weights: Respondent weights.
6         cap: Largest allowed weight.
7
8     Returns:
9         Indexes of weights greater than cap.
10    """
11    return [i for i, w in enumerate(weights) if w > cap]
```

Three classes and a factory wrapped one comparison. Add structure when a second rule actually arrives.

D6-L1 Read before you run, and check each API

D6-L1 exercise 1

Hints. (1) What can this cell change? (2) The table. (3) Which rule does bronze follow?

Answer. a) The write mode and the target, against the goal and the table's rules Because: r1) Side effects on a production table cannot always be undone. Bronze is append-only.

Side effects first: write mode, target name, DROP, DELETE. Then the rest.

Common wrong answers.

- b: Formatting matters later. A destructive write matters now.
- c: Speed is secondary to what the cell changes.
- d: The result of an overwrite is a replaced table.
- a|r2: Right answer, wrong reason. Often the mode is right. You check because the cost of a wrong one is high.

D6-L1 exercise 2

Hints. (1) What do you check the code against? (2) Names before side effects. (3) Run last.

Answer.

1. Restate the goal in one sentence
2. List every function, parameter, table and column
3. Check each name against the docs for your version and the schema
4. Find side effects: write modes, DROP, DELETE
5. Run it on development data in a dev schema

Goal, names, checks, side effects, then a dev run.

Common wrong answers.

- Without a goal stated first, you have nothing to check the code against.
- Find side effects before running, even on dev data.

D6-L1 exercise 3

Hints. (1) A built-in that returns True or False for an attribute. (2) inspect's function for parameters. (3) `hasattr` and `signature`.

Answer.

```
1 import inspect
2 from pyspark.sql import functions as F
3
4 hasattr(F, "try_divide")
5 "validate" in inspect.signature(pd.DataFrame.merge).parameters
```

`hasattr` checks a name. `inspect.signature(...).parameters` checks a parameter.

Common wrong answers.

- `g1 = getattr`: `getattr` raises if the name is missing. `hasattr` returns True or False.

D6-L1 exercise 4

Hints. (1) Walk the tree with `ast.walk` and keep `ast.Call` nodes. (2) Build the dotted name recursively from `Attribute.value` down to a `Name`. (3) Return `None` when the chain hits anything else.

Answer.

```
1 import ast
2
3
4 def _dotted(node: ast.expr) -> str | None:
```



```

5     """Dotted name for a Name or Attribute chain, else None."""
6     if isinstance(node, ast.Name):
7         return node.id
8     if isinstance(node, ast.Attribute):
9         base = _dotted(node.value)
10        return None if base is None else f"{base}.{node.attr}"
11    return None
12
13
14 def called_names(source: str) -> list[str]:
15     """Sorted unique dotted names of every call in the source.
16
17     Args:
18         source: Python source code.
19
20     Returns:
21         Names such as "F.col" or "spark.read.format".
22     """
23     names = set()
24     for node in ast.walk(ast.parse(source)):
25         if isinstance(node, ast.Call):
26             name = _dotted(node.func)
27             if name is not None:
28                 names.add(name)
29     return sorted(names)

```

A list of every call is a checklist: each name gets a docs link or a question mark.

D6-L1 exercise 5

Hints. (1) Which line changes existing data? (2) Look at the write mode. (3) Append.

Answer. The bug is on line 5.

Fixed version:

```

1 from pyspark.sql import functions as F
2
3 delivery = spark.read.format("csv").option("header", True).load(path)
4 delivery = delivery.withColumn("ingestion_timestamp",
5     F.current_timestamp())
6 delivery.write.mode("append").saveAsTable("research_bronze.nat_raw")

```

One word, `overwrite`, turns an append into a replacement of all history.

Common wrong answers.

- Line 4: Adding an ingestion timestamp is part of the bronze standard.
- Line 3: The read is fine. The write replaces history.

D6-L2 Small data with a known answer, failing test first

D6-L2 exercise 1

Hints. (1) What should a test do when the behaviour is missing? (2) Fail. (3) This one did not.

Answer. a) The test does not check the feature Because: r1) A test must fail when the behaviour it checks is missing. Passing without the feature means it asserts something else.

Watch every new test fail once. Then you know it can.

Common wrong answers.

- b: The feature does not exist. The test passed anyway, so it is not checking it.
- c: A test that cannot fail is not well written.
- d: Speed is unrelated.
- a|r2: Right answer, wrong reason. Tests do not fail by default. This one should have.

D6-L2 exercise 2

Hints. (1) Where do $>$ and $>=$ differ? (2) Only at equality. (3) Add that case.

Answer. a) A case with a weight exactly equal to the cap Because: r1) $>$ and $>=$ differ only on the boundary, so only a boundary case can tell them apart.

A deliberate break shows which rule your tests miss. Boundaries are the usual gap.

Common wrong answers.

- b: Both operators agree far from the cap. Those cases cannot catch this change.
- c: Timing has nothing to do with the comparison.
- d: They pass with a known bug, so they are incomplete.
- a|r2: Right answer, wrong reason. One well-chosen case is enough: the boundary.

D6-L2 exercise 3

Hints. (1) The function you call on the last line. (2) The parameter about row order. (3) `assertDataFrameEqual` and `checkRowOrder`.

Answer.

```
1 from pyspark.testing import assertDataFrameEqual
```

```

2
3 expected = spark.createDataFrame([("W01", 0.25)], "wave_id string, share
    double")
4 assertDataFrameEqual(result, expected, checkRowOrder=False)

```

`checkRowOrder=False` is the default. Writing it out documents the intent.

Common wrong answers.

- `g2 = ignoreRowOrder`: The parameter is `checkRowOrder`. It is `False` by default in PySpark 4.0.

D6-L2 exercise 4

Hints. (1) A normal case whose share is not 0 or 1, so a percentage differs. (2) A zero total expecting `None`. (3) Work each expected value out by hand.

Answer.

```

1 def share_cases() -> list[tuple[float, float, float | None]]:
2     """Known-answer cases for share(part, total), worked out by hand.
3
4     Returns:
5         (part, total, expected) tuples covering a normal share, a zero
6         share and a zero total.
7     """
8     return [
9         (1.0, 4.0, 0.25),
10        (0.0, 10.0, 0.0),
11        (5.0, 0.0, None),

```

Good cases are chosen to break wrong code. A case where every version agrees, such as (1, 1, 1.0) for a percentage bug, adds nothing.

D6-L2 exercise 5

Hints. (1) Which rows exercise each rule? (2) Who decides the right answer? (3) How do you prove a test can fail?

Answer.

- Each fixture row exercises one rule of the requirement.
- Expected rows come from you, not from the code under test.
- A deliberate wrong key must fail, or the fixture misses the multi-wave case.

Designed rows, hand-written expectations, and a deliberate break.

Common wrong answers.

- `fixture = sample`: You cannot work out 1,000 rows by hand, and production rarely has every edge case.

- `expected = run`: That makes the test self-confirming: it records the code's output, bugs included.
- `prove = rerun`: A rerun gives the same result. Only a deliberate break shows the test can fail.

D6-L3 Ask for assumptions, then compare a second approach

D6-L3 exercise 1

Hints. (1) Who knows whether it is true? (2) The data. (3) Check it, then keep the check.

Answer. a) Check it on the real data, and add a uniqueness check if it matters Because: r1) An assumption is a claim about data. Only the data can confirm it, and it can stop holding in a later wave.

Check each assumption with a query, then keep the important ones as pipeline checks.

Common wrong answers.

- `b`: The assistant has never seen your data. It guessed.
- `c`: Silently deleting rows hides a delivery problem. Check first, then decide.
- `d`: An unchecked assumption is a future bug.
- `a|r2`: Right answer, wrong reason. The assumptions are usually sensible guesses. You check because guesses can be wrong.

D6-L3 exercise 2

Hints. (1) Which method equals the known answer? (2) 1310. (3) Method b.

Answer.

```
1 b
```

The known answer settles disagreements. Method a lost rows, most likely in its inner join.

Common wrong answers.

- `a`: Method a gives 1250, which does not match the hand-worked 1310.

D6-L3 exercise 3

Hints. (1) Which join keeps left rows without a match? (2) Then count them. (3) `left_anti` and `count`.

Answer.

```
1 missing = silver.join(regions, "region_code", "left_anti").count()
```

A left anti join returns silver rows with no match in regions. Zero means the assumption holds.

Common wrong answers.

- `g1 = inner`: An inner join keeps the matches. You want the rows with no match.

D6-L3 exercise 4

Hints. (1) Uniqueness uses the pair (respondent_id, wave_id). (2) Positive means > 0 and not None. (3) Use `re.fullmatch` so 'W145' does not pass.

Answer.

```
1 import re
2
3 WAVE_PATTERN = re.compile(r"W\d{2}")
4
5
6 def violated_assumptions(rows: list[dict[str, object]]) -> list[str]:
7     """Names of data assumptions that the rows break.
8
9     Args:
10         rows: Records with respondent_id, wave_id and final_weight.
11
12     Returns:
13         Violated assumption names in a fixed order.
14     """
15     violated: list[str] = []
16     keys = [(r["respondent_id"], r["wave_id"]) for r in rows]
17     if len(keys) != len(set(keys)):
18         violated.append("unique_respondent_per_wave")
19     weights = [r["final_weight"] for r in rows]
20     if any(w is None or not isinstance(w, (int, float)) or w <= 0 for w
21           in weights):
22         violated.append("positive_weight")
23     waves = [r["wave_id"] for r in rows]
24     if any(not isinstance(w, str) or WAVE_PATTERN.fullmatch(w) is None
25           for w in waves):
26         violated.append("wave_format")
27     return violated
```

Each assumption becomes one small check with a name. The names go into the run log, so a failure says which assumption broke.

D6-L3 exercise 5

Hints. (1) What makes a second method independent? (2) What decides who is right? (3) What catches the next regression?

Answer.

- Independent code avoids repeating the same reasoning, and the same mistake.
- Hand-worked answers decide disagreements.
- A cheap reconciliation total stays in the pipeline to catch the next regression.

Independent method, known answer as referee, and a lasting reconciliation check.

Common wrong answers.

- `route = rewrite`: The same assistant tends to repeat its own assumptions, such as the inner join.
- `disagree = average`: An average of a right and a wrong number is wrong.
- `keep = drop_sql`: The next change can break it again. Keep a cheap independent check.

D7-L1 Review practice: Python and pandas

D7-L1 exercise 1

Hints. (1) Where does `uuid` come from? (2) It is never imported. (3) Add `import uuid`.

Answer. The bug is on line 3.

Fixed version:

```
1 import uuid
2
3
4 def new_run_id() -> str:
5     """Return a new unique run id."""
6     return str(uuid.uuid4())
```

Reference defect 1. `NameError` on the first run. `ruff` rule `F821` flags undefined names before you run anything.

Common wrong answers.

- Line 1: The signature is fine. The body uses a name that was never imported.

D7-L1 exercise 2

Hints. (1) When is the default list created? (2) Once, at definition. (3) Both calls share it.

Answer.

```
1 2
```

Mutable default argument. Default to None and create the list inside.

Common wrong answers.

- 1: The default list is created once and shared, so the second call sees the first call's warning.

D7-L1 exercise 3

Hints. (1) Is the returned datetime aware? (2) `utcnow()` returns a naive one, and is deprecated. (3) `datetime.now(timezone.utc)`.

Answer. The bug is on line 5.

Fixed version:

```
1 from datetime import datetime, timezone
2
3 def run_timestamp() -> datetime:
4     """Current time in UTC for the run log."""
5     return datetime.now(timezone.utc)
```

Reference defect 6. Naive timestamps cannot be compared with aware ones and hide their zone.

Common wrong answers.

- Line 1: The import is right, and `timezone` is even imported already. It is just not used.

D7-L1 exercise 4

Hints. (1) Can `error_message` be None? (2) Then the type must allow None. (3) `str | None = None`.

Answer. The bug is on line 4.

Fixed version:

```
1 def log_run(
2     pipeline_name: str,
3     status: str,
4     error_message: str | None = None,
5 ) -> None:
6     """Append one record to the run log."""
```

Reference defect 7. `str = None` claims the value is always a string. mypy --strict rejects it.

Common wrong answers.

- Line 3: status is required and always a string. The optional parameter is the next one.

D7-L1 exercise 5

Hints. (1) Is 0.0 truthy? (2) No. (3) Use `is not None`.

Answer. The bug is on line 3.

Fixed version:

```
1 def count_weighted(weights: list[float | None]) -> int:
2     """Number of respondents with a recorded weight."""
3     recorded = [w for w in weights if w is not None]
4     return len(recorded)
```

Truthiness treats 0.0 like None. Your standard: test for None with `is None`.

Common wrong answers.

- Line 4: len is right. The filter above drops 0.0.

D7-L1 exercise 6

Hints. (1) Lesson D3-L1: nulls in filters. (2) pandas and Spark differ here. (3) pandas keeps them.

Answer. a) Yes: in pandas, `!=` keeps missing regions Because: r1) pandas treats missing `!= 'N'` as True, so those rows stay. Spark would drop them.

The same filter in PySpark would drop the null rows. Know which engine you are reviewing.

Common wrong answers.

- b: That is Spark and SQL behaviour. pandas keeps them.
- c: The comparison runs without error.
- d: The result is the same for object and str dtypes.
- a|r2: Right answer, wrong reason. pandas does not fill. Missing compares as not-equal.

D7-L1 exercise 7

Hints. (1) Which join type keeps unmatched left rows? (2) Inner drops them. (3) `how='left'`.

Answer. The bug is on line 6.

Fixed version:


```

1 import pandas as pd
2
3 def add_region(resp: pd.DataFrame, regions: pd.DataFrame) -> pd.DataFrame:
4     """Attach region names. One row per respondent."""
5     lookup = regions.dropna(subset=["code"])
6     return resp.merge(lookup, on="code", how="left",
                        validate="many_to_one")

```

An inner join quietly removes respondents with unknown codes. The grain promise, one row per respondent, needs a left join.

Common wrong answers.

- Line 5: Dropping null lookup keys is right: pandas would otherwise match null to null.

D7-L1 exercise 8

Hints. (1) What is `max(['W9', 'W10'])`? (2) 'W9': text comparison. (3) Use `key=lambda w: int(w[1:])`.

Answer. The bug is on line 5.

Fixed version:

```

1 def latest_wave(wave_ids: list[str]) -> str:
2     """Return the most recent wave id."""
3     if not wave_ids:
4         raise ValueError("no waves")
5     return max(wave_ids, key=lambda w: int(w[1:]))

```

Text sorts character by character. The truthiness test on line 3 is fine: an empty list is exactly what it should catch.

Common wrong answers.

- Line 3: `if not wave_ids` is correct for a list: an empty list is the case to catch.

D7-L2 Review practice: PySpark and Spark SQL

D7-L2 exercise 1

Hints. (1) Which sum does line 5 call? (2) The one imported on line 1. (3) Import functions as F.

Answer. The bug is on line 1.

Fixed version:

```

1 from pyspark.sql import functions as F
2
3 def total_checks(results: list[dict[str, str]]) -> int:
4     """Number of passed checks."""

```

```
5     return sum(1 for r in results if r["status"] == "PASS")
```

Reference defects 2 and 3. The helper raises NOT_COLUMN_OR_STR because it calls Spark's sum on a generator.

Common wrong answers.

- Line 5: This line is correct Python. It fails only because line 1 replaced sum.

D7-L2 exercise 2

Hints. (1) What does ruff E712 flag? (2) Comparison to True. (3) The column is already boolean.

Answer. The bug is on line 3.

Fixed version:

```
1 from pyspark.sql import functions as F
2
3 in_scope = silver.filter(F.col("in_scope"))
4 in_scope.write.mode("overwrite").saveAsTable("research_dev.nat_in_scope")
```

Reference defect 12. `F.col("in_scope")` is the condition.

Common wrong answers.

- Line 4: Overwriting a dev table is fine here. The standard issue is the comparison.

D7-L2 exercise 3

Hints. (1) What is null != 'test' in Spark? (2) Null, which the filter drops. (3) `eqNullSafe`.

Answer. The bug is on line 3.

Fixed version:

```
1 from pyspark.sql import functions as F
2
3 live = respondents.filter(~F.col("panel_type").eqNullSafe("test"))
```

Spark's three-valued logic drops null rows from a `!=` filter.

Common wrong answers.

- Line 1: The import follows the standard.

D7-L2 exercise 4

Hints. (1) What does `when()` give for non-matching rows? (2) Null. (3) Sum of nulls.

Answer. a) NULL Because: r1) when without otherwise is null for non-matching rows, and a sum of nulls is null.

Reference defect 16. Wrap with `F.coalesce(..., F.lit(0))`.

Common wrong answers.

- **b:** That is the goal. It needs `coalesce(..., lit(0))` or `otherwise(0)`.
- **c:** Non-matching rows contribute null, not 1.
- **d:** No error. The null passes silently.
- **a|r2:** Right answer, wrong reason. A sum of real zeros is 0. A sum of only nulls is null.

D7-L2 exercise 5

Hints. (1) Compare the `groupBy` with the stated grain. (2) Which grain column is missing? (3) `tracker_id`.

Answer. The bug is on line 5.

Fixed version:

```
1 from pyspark.sql import functions as F
2
3 metrics = (
4     silver.filter(F.col("in_scope"))
5     .groupBy("tracker_id", "wave_id", "age_group", "gender", "region")
6     .agg(F.sum("final_weight").alias("weighted_base"))
7 )
```

Reference finding 15. Without `tracker_id`, two trackers' respondents merge into one row per wave and subgroup.

Common wrong answers.

- Line 4: The in-scope filter follows the standard.
- Line 6: The aggregate is fine. The grouping misses a grain column.

D7-L2 exercise 6

Hints. (1) How many times is the table scanned? (2) Once per `count()`. (3) One agg with conditional sums.

Answer. The bug is on line 5, 6, 7.

Fixed version:

```
1 from pyspark.sql import functions as F
2
3 df = spark.table(table_name)
4 row = df.agg(
5     F.count("*").alias("row_count"),
6     F.sum(F.col("respondent_id").isNull().cast("int")).alias("null_ids"),
```

```

7      F.sum(F.col("final_weight").isNull().cast("int")).alias("null_weights"),
8    ).first()
9    health = row.asDict()

```

Reference defect 11. Three actions mean three scans of an uncached table. One `agg` does it in one.

Common wrong answers.

- Line 3: Reading the table once is right. The three actions below each scan it again.

D7-L2 exercise 7

Hints. (1) What happens when two source rows match one target row? (2) MERGE fails. (3) Deduplicate the source on the keys.

Answer. The bug is on line 5.

Fixed version:

```

1  from delta.tables import DeltaTable
2
3  source = delivery.dropDuplicates(["respondent_id", "wave_id"])
4  target = DeltaTable.forName(spark, "research_silver.nat_respondents")
5  (target.alias("t")
6    .merge(source.alias("s"), "t.respondent_id = s.respondent_id AND
7      t.wave_id = s.wave_id")
8    .whenMatchedUpdateAll()
9    .whenNotMatchedInsertAll()
10   .execute())

```

Reference finding 14. Delta MERGE fails when several source rows match one target row. Deduplicate first, ideally keeping the latest delivery.

Common wrong answers.

- Line 6: Updating matched rows is right. The failure comes from two source rows matching one target row.

D7-L2 exercise 8

Hints. (1) Expand NOT IN into AND terms. (2) `x <> NULL` is null. (3) No row passes.

Answer.

```

1  0

```

Use `NOT EXISTS` or a `left_anti` join when the list can hold nulls.

Common wrong answers.

- 700: That is what NOT EXISTS returns. NOT IN with a NULL in the list is never true.

D7-L2 exercise 9

Hints. (1) What does withColumnRenamed do with a missing column? (2) Nothing, silently. (3) Check bronze_df.columns first.

Answer. The bug is on line 4.

Fixed version:

```
1 def rename_columns(bronze_df, column_map: dict[str, str]):
2     """Rename source columns to canonical names. Raises if a source column
        is missing."""
3     missing = sorted(set(column_map) - set(bronze_df.columns))
4     if missing:
5         raise ValueError(f"delivery is missing columns: {missing}")
6     for src, tgt in column_map.items():
7         bronze_df = bronze_df.withColumnRenamed(src, tgt)
8     return bronze_df
```

Reference finding 18. A renamed column in a delivery vanishes without error, and every later step sees a null or a missing field.

Common wrong answers.

- Line 3: Looping over the map is fine. The rename inside it hides a missing column.

D7-L3 Review practice: Databricks SQL and DAX

D7-L3 exercise 1

Hints. (1) Does line 3 exist in Databricks SQL? (2) It is Snowflake syntax. (3) ALTER TABLE ... SET ROW FILTER.

Answer. The bug is on line 3.

Fixed version:

```
1 CREATE FUNCTION research_governance.region_filter(region STRING)
2     RETURN is_account_group_member(concat('analysts_', region));
3 ALTER TABLE research_silver.nat_respondents SET ROW FILTER
    research_governance.region_filter ON (region);
```

Reference defect 4. The UDF is right. The binding statement comes from another platform.

Common wrong answers.

- Line 1: A SQL UDF is exactly what Databricks row filters use.

D7-L3 exercise 2

Hints. (1) Who receives each grant? (2) One line names a person. (3) Grant to the group.

Answer. The bug is on line 3.

Fixed version:

```
1 GRANT USE CATALOG ON CATALOG research TO `analysts`;
2 GRANT USE SCHEMA ON SCHEMA research.research_gold TO `analysts`;
3 GRANT SELECT ON TABLE research.research_gold.nat_metrics TO `analysts`;
```

Grants to groups survive team changes. A grant to one email is forgotten when that person moves.

Common wrong answers.

- Line 1: USE CATALOG to the group is needed and correct.

D7-L3 exercise 3

Hints. (1) Which group function does the goal name? (2) Account-level groups. (3) is_account_group_member.

Answer. The bug is on line 2.

Fixed version:

```
1 CREATE FUNCTION research_governance.mask_postcode(p STRING)
2 RETURN CASE WHEN is_account_group_member('analysts') THEN p ELSE '***'
END;
3 ALTER TABLE research_silver.nat_respondents ALTER COLUMN postcode SET
MASK research_governance.mask_postcode;
```

Reference finding 22. is_member checks workspace-local groups. Account groups are the ones Unity Catalog uses.

Common wrong answers.

- Line 3: ALTER COLUMN ... SET MASK is the correct Databricks statement.

D7-L3 exercise 4

Hints. (1) How many parts does the name have? (2) Which part is missing? (3) The current catalog fills it in.

Answer. a) research_dev.research_gold.nat_metrics, if it exists Because: r1)

A two-part name resolves against the current catalog. Only a three-level name is unambiguous.

Use three-level names in jobs: `research.research_gold.nat_metrics`.

Common wrong answers.

- b: The name does not say which catalog. The session decides.
- c: Two-part names work. That is why they are risky.
- d: One table is resolved.
- a|r2: Right answer, wrong reason. The current catalog decides, whatever it is.

D7-L3 exercise 5

Hints. (1) Does each subgroup weigh the same in AVG? (2) Should it? (3) Total over total.

Answer. The bug is on line 2.

Fixed version:

```
1 SELECT region,
2     SUM(weighted_reach) / NULLIF(SUM(weighted_base), 0) AS reach_pct
3 FROM research.research_gold.nat_metrics
4 GROUP BY region;
```

Recompute ratios from sums at the level you report. NULLIF turns a zero base into a null result instead of an error.

Common wrong answers.

- Line 4: Grouping by region matches the goal.

D7-L3 exercise 6

Hints. (1) What does / return when base_w is 0? (2) Infinity or an error. (3) `DIVIDE` with the base as its second argument.

Answer.

```
1 Reach % =
2 VAR reach_w = SUM ( fact_reach[weighted_reach] )
3 VAR base_w = SUM ( fact_reach[weighted_base] )
4 RETURN DIVIDE ( reach_w, base_w )
```

DIVIDE returns BLANK on a zero or blank denominator, so the cell stays empty (lesson C6-L2).

Common wrong answers.

- g1 = IFERROR: IFERROR hides every error, not only a zero base. DIVIDE returns BLANK for a zero or blank denominator.

D7-L3 exercise 7

Hints. (1) Which filters does ALL(fact_reach) remove? (2) All of them. (3) Name the brand dimension.

Answer. The bug is on line 3.

Fixed version:

```
1 Brand share =  
2 VAR this_brand = SUM ( fact_reach[weighted_reach] )  
3 VAR all_brands = CALCULATE ( SUM ( fact_reach[weighted_reach] ), ALL  
    ( dim_brand ) )  
4 RETURN DIVIDE ( this_brand, all_brands )
```

The denominator must keep the wave and region filters and drop only the brand filter.

Common wrong answers.

- Line 4: DIVIDE is right. The denominator's context is wrong.

D7-L3 exercise 8

Hints. (1) Can the source hold two rows for one key? (2) Yes, and MERGE fails then. (3) Deduplicate in the USING clause with ROW_NUMBER.

Answer. The bug is on line 2.

Fixed version:

```
1 MERGE INTO research.research_silver.nat_respondents AS t  
2 USING (  
3     SELECT * EXCEPT (rn) FROM (  
4         SELECT *, ROW_NUMBER() OVER (PARTITION BY respondent_id, wave_id  
            ORDER BY ingestion_timestamp DESC) AS rn  
5         FROM research.research_bronze.nat_raw_w14  
6     ) WHERE rn = 1  
7 ) AS s  
8 ON t.respondent_id = s.respondent_id AND t.wave_id = s.wave_id  
9 WHEN MATCHED THEN UPDATE SET *  
10 WHEN NOT MATCHED THEN INSERT *;
```

Deduplicate the source on the merge keys, keeping the latest row, before MERGE. Reference finding 14 in SQL form.

Common wrong answers.

- Line 3: The join condition uses both merge keys. The source is the problem.

D8-L1 State versions and constraints in every prompt

D8-L1 exercise 1

Hints. (1) Which fact can the assistant not guess? (2) The platform. (3) Name the dialect.

Answer. a) Target: Databricks SQL, Unity Catalog, Databricks Runtime 17.3 LTS Because: r1) Row filter syntax differs by platform. Naming the dialect rules out Snowflake's CREATE ROW ACCESS POLICY.

The dialect decides the syntax. Name it every time.

Common wrong answers.

- b: Vague quality words do not change which dialect the assistant picks.
- c: Efficiency does not matter if the statement does not exist on your platform.
- d: An explanation of the wrong dialect is still the wrong dialect.
- a|r2: Right answer, wrong reason. Length does not help. Specific facts do.

D8-L1 exercise 2

Hints. (1) Context before the task. (2) Rules, then prohibitions. (3) Assumptions at the end.

Answer.

1. Environment: runtime, language and library versions
2. Standards that apply: imports, typing, table rules
3. What not to do: no UDFs, no RDDs, no secrets
4. The task, with inputs, outputs and side effects
5. Ask for a list of assumptions

Environment, standards, prohibitions, task, assumptions.

Common wrong answers.

- Context first. An assistant starts shaping the answer as soon as it reads the task.
- The assumptions request comes last, after the task it refers to.

D8-L1 exercise 3

Hints. (1) Check VERSIONS.md. (2) The LTS runtime this course pins. (3) 17.3 LTS with Spark 4.0.

Answer.

```
1 Environment: Databricks Runtime 17.3 LTS (Spark 4.0, Python 3.12), Unity Catalog.
```

Databricks Runtime 17.3 LTS ships Apache Spark 4.0.0 and Python 3.12.3.

Common wrong answers.

- `g1 = 18`: Runtime 18 is not LTS yet. This course targets 17.3 LTS.

D8-L1 exercise 4

Hints. (1) Escape the library name with `re.escape`. (2) Follow it with `\s*v?\d+(\.\d+)*`. (3) Use lookarounds so `pandas-stubs` does not match `pandas`.

Answer.

```
1 import re
2
3
4 def unstated_versions(prompt: str, libraries: list[str]) -> list[str]:
5     """Libraries a prompt names without a version number.
6
7     Args:
8         prompt: Prompt text.
9         libraries: Library names that need a stated version.
10
11     Returns:
12         Libraries without a version, in the order given.
13     """
14     missing = []
15     for lib in libraries:
16         pattern = rf"(?![\w-]){re.escape(lib)}(?:[\w-])\s*v?\d+(\.\d+)*"
17         if re.search(pattern, prompt, flags=re.IGNORECASE) is None:
18             missing.append(lib)
19     return missing
```

A small lint on your own prompts. Run it on the header you paste, and you never forget the versions line.

D8-L1 exercise 5

Hints. (1) Which versions will the code run on? (2) Which rules apply to this task? (3) Where do shared standards live?

Answer.

- Pinned versions match what the code will run on.
- A short list of relevant rules gets followed. A long file gets skimmed, by people and models.
- A header in the repository is shared, versioned and updated when the runtime changes.

Pinned versions, relevant rules, and a header kept in the repository.

Common wrong answers.

- versions = latest: The latest PySpark is not what Databricks Runtime 17.3 LTS runs.
- standards = whole: The rules that matter get lost among the ones that do not apply.
- storage = memory: Unwritten headers drift and differ between people.

D8-L2 State expected outputs and tests in the prompt

D8-L2 exercise 1

Hints. (1) Which option states the rule exactly? (2) An example with the answer. (3) Option a.

Answer. a) A five-row example with a null and the expected result worked by hand Because: r1) An example states the rule precisely. Words like 'properly' leave the rule to the assistant.

A worked example is a specification the assistant cannot misread.

Common wrong answers.

- b: Properly how? Excluded from the base, counted as 0, or an error?
- c: Which edge cases, and with what answers?
- d: Best practice does not say whether a null counts in the base.
- a|r2: Right answer, wrong reason. Instructions in words work when they are specific. An example is the most specific form.

D8-L2 exercise 2

Hints. (1) Which rows are valid? (2) Reach 3.0, base 4.0. (3) 0.75.

Answer.

1 0.75

Working the example out yourself is part of writing the prompt. If you cannot, the rule is not clear yet.

Common wrong answers.

- 0.6: That includes the None row in the base: 3 / 5. The rule excludes it: 3 / 4.

D8-L2 exercise 3

Hints. (1) Tests are fixed. (2) Disagreement goes to you, not into the test. (3) without and say so.

Answer.

```
1 Make these tests pass without editing them. If a test seems wrong, say so
  instead of changing it.
```

Fixed tests are the acceptance criteria. An assistant that edits them to pass has changed the requirement.

Common wrong answers.

- `g1 = by`: Editing the tests to pass is exactly what the instruction must forbid.

D8-L2 exercise 4

Hints. (1) `' , '.join(repr(a) for a in args)` builds the argument list. (2) Use `{expected!r}` for the expected value. (3) Special-case `None` with `is None`.

Answer.

```
1 def examples_to_asserts(
2     func_name: str, examples: list[tuple[tuple[object, ...], object]]
3 ) -> list[str]:
4     """Turn worked examples into assert lines for a prompt or a test file.
5
6     Args:
7         func_name: Name of the function under test.
8         examples: (args, expected) pairs worked out by hand.
9
10    Returns:
11        One assert statement per example.
12    """
13    lines = []
14    for args, expected in examples:
15        call = f"{func_name}({' , '.join(repr(a) for a in args)})"
16        if expected is None:
17            lines.append(f"assert {call} is None")
18        else:
19            lines.append(f"assert {call} == {expected!r}")
20    return lines
```

Worked examples become executable tests in one step. Paste them into the prompt and into the test file.

D8-L2 exercise 5

Hints. (1) What settles the key and the grain? (2) What happens with two rows for one key? (3) When are tests written?

Answer.

- A worked example fixes the grain and which columns form the key.

- An explicit tie rule removes a decision the assistant would otherwise make silently.
- Fixed tests written first are acceptance criteria, not a mirror of the code.

Example, explicit rules, and tests written first.

Common wrong answers.

- `example = none_ex`: Deduplicate on which key, keeping which row? The example answers both.
- `tie = any`: 'Any' makes the output depend on execution order, and it changes between runs.
- `tests = after`: Tests written after the code tend to confirm what the code does, bugs included.

D8-L3 Iterate on a failing answer

D8-L3 exercise 1

Hints. (1) What does the assistant need to know? (2) What failed, and how. (3) Test, values, scope.

Answer. a) The failing test, the actual and expected values, and 'fix only this' Because: r1) Precise evidence and a narrow scope lead to a targeted fix instead of a rewrite.

Evidence plus scope. The same rule you would follow when asking a colleague.

Common wrong answers.

- `b`: Without evidence, the assistant guesses what failed and often rewrites working parts.
- `c`: 'More carefully' gives no information about what is wrong.
- `d`: More context is not more evidence. Send the failing test and the smallest input.
- `a|r2`: Right answer, wrong reason. Assistants respond to anything. Evidence makes the response useful.

D8-L3 exercise 2

Hints. (1) Which line widens the task? (2) An open-ended refactor. (3) Replace it with a scope limit.

Answer. The bug is on line 3.

Fixed version:

```
1 test_zero_share fails: expected ["a"], got [].
2 Input: compare_waves({"a": 0.0}, {"a": 0.12}).
3 Fix only this failure. Do not edit the tests or change other behaviour.
```

One problem per round. Refactoring can come later, as its own request with its own tests.

Common wrong answers.

- Line 1: The failing test and values are exactly what the prompt needs.
- Line 2: The smallest failing input is useful evidence.

D8-L3 exercise 3

Hints. (1) What goes into the prompt? (2) What must you check before trusting a pass? (3) When do you stop asking?

Answer.

1. Collect the failing assertion, the error and the smallest input
2. Send them with 'fix only this, do not edit the tests'
3. Diff both the code and the test file in the answer
4. Run the full test suite
5. If two rounds fail, debug and fix it yourself

Evidence, focused prompt, diff, full suite, and a limit on rounds.

Common wrong answers.

- Diff before running. Passing tests mean little if the tests were edited.
- Collect the evidence first. It is what the prompt sends.

D8-L3 exercise 4

Hints. (1) Build a set of stripped lines from after. (2) Walk before in order and keep stripped assert lines not in that set. (3) Counting asserts misses a changed one.

Answer.

```
1 def removed_asserts(before: str, after: str) -> list[str]:
2     """Assert statements in the old test file that the new one lacks.
3
4     Args:
5         before: Test file text before the assistant's change.
6         after: Test file text after the change.
7
8     Returns:
9         Stripped assert lines from before that are missing in after, in
10        order.
11        """
12     kept = {line.strip() for line in after.splitlines()}
```

```

12     removed = []
13     for line in before.splitlines():
14         stripped = line.strip()
15         if stripped.startswith("assert") and stripped not in kept:
16             removed.append(stripped)
17     return removed

```

Run it on every answer that touches tests. A changed assertion needs your approval, never a silent accept.

D8-L3 exercise 5

Hints. (1) What is faster after two failed rounds? (2) What does a long failed thread carry? (3) What can make a suite green without a fix?

Answer.

- Two focused rounds without progress means the assistant is guessing. Your debugger is faster.
- A fresh thread drops the failed attempts and starts from your diagnosis.
- A green suite with edited tests proves nothing.

Stop and debug, restart clean with a diagnosis, and diff the tests every time.

Common wrong answers.

- `next = round4`: A fourth vague round tends to produce a fourth guess.
- `fresh = same_thread`: The failed attempts stay in context and pull later answers towards them.
- `tests = trust_green`: Green can mean the tests were weakened.