X2-L1 lab: Capstone 2: build and verify the NAT reach report

Platform: Fabric trial workspace

Run this in a Fabric trial workspace connected to an Azure Databricks workspace with Unity Catalog. Copilot needs a paid F2 or higher capacity (lesson C8-L1). Keep your verification log as a table with columns: output_id, prompt, output summary, check_method, check detail, result, checked_by, date.

1. In Databricks, create gold table `nat_metrics_shareable` with no row filter or column mask, holding only data every analyst may see. Run OPTIMIZE on it.
2. In Fabric, create a mirrored Azure Databricks catalog item for the gold schema. Confirm that nat_metrics_shareable appears and that any view or filtered table does not.
3. Create a semantic model in Direct Lake on OneLake mode over the mirrored tables. Build a star schema with dim_region and dim_wave.
4. Add measures for weighted reach and region share with DIVIDE. For each measure Copilot drafts, compare the value for one wave with a SQL total from gold, and log the result.
5. Give analysts the Viewer role. Create a OneLake security role per region that filters the fact table, and test it with a test account.
6. Draft the report pages with Copilot. Log each visual you keep, with the check you used.
7. Run log_problems (exercise 4) on your log. Fix every problem before sharing the report.
8. Write a one-page design note covering storage mode, row security and verification, and the decisions from exercise 5.

Notes